arXiv:2610.26734v1 [cs.LG] 19 Oct 2026

CAVA: Complementary-Label-Aware Data Valuation

Tan-Ha Mai
National Taiwan University

Hsuan-Tien Lin
National Taiwan University

Vu Nguyen
Amazon

Abstract

Complementary-label learning *(CLL) trains a classifier from labels that name a class an instance does* not *belong to. Such labels are cheap to collect from crowdworkers or vision-language models (VLMs), making them attractive for large-scale annotation. However, both sources can mistakenly assign the true class as the complementary label. Such annotation errors are particularly harmful in CLL because the model is explicitly trained to suppress the correct answer. Optimal transport (OT) data valuation identifies harmful samples by comparing training data with a clean validation set, but transferring it directly to CLL creates a semantic mismatch: label c means "not class c" in the training but "class c" in the validation. We show that this makes the OT label groups noisy mixtures whose behavior depends on the noise rate. The transferred score is informative only at sufficiently high noise, becomes uninformative at the crossover rate $1/M$ for M classes, and reverses below it, where practical CLL annotation typically operates. We propose* CAVA, *which introduces a* complementary label cost *and a* label margin *that isolates the contribution of each complementary label to the transport potential. The margin keeps a consistent direction across noise rates, detects diagonal noise with a parameter-free threshold of closed-form precision, and allows detected samples to be recycled as ordinary-label supervision. On synthetic and human-annotated benchmarks, CAVA separates diagonal noise almost perfectly where directly transferred OT valuation performs worse than random. With only a fraction of the clean validation data, it also improves every CLL objective we evaluate, showing that OT-based valuation becomes reliable once the semantics of complementary supervision are modeled explicitly.*

1. Introduction

Robust vision classifiers depend on large labeled datasets, but annotation remains expensive. Complementary-label learning (CLL) reduces this burden by using a simple form of weak supervision [33]: instead of assigning the correct class, each image is labeled with a class it does *not* belong to [4, 13]. Telling that a photo is *not a truck* may take only a glance, whereas deciding whether an object in fog is a ship or an airplane can require closer inspection. This asymmetry makes complementary labels relatively easy to collect at scale, first from crowdworkers [35] and increasingly from vision-language models (VLMs), which can annotate millions of images at very low cost [28, 42, 43]. CLL is therefore a practical option for large-scale annotation in label spaces of up to about 20 classes, where current CLL methods are workable [29], provided that errors in these labels can be handled reliably.

This condition matters because neither humans nor VLMs are perfect annotators. Both can occasionally assign the *true* class as the class an image does not belong to [28, 35]. CLL is particularly sensitive to this *diagonal noise* (Figure 8, Appendix A.2). A clean complementary label only rules out one incorrect class, whereas a diagonal label explicitly tells the learner to suppress the correct one. Even a small amount of such noise can destabilize training, and common techniques such as augmentation and label sharing can further propagate the error [23, 27]. A practical CLL pipeline therefore needs a denoising method that works at dataset scale, does not depend on a particular annotator, and provides some indication of how reliable its own decisions are.

Data valuation offers a natural way to do this [10]. Among existing approaches, optimal-transport (OT) valuation is attractive because it is scalable and largely independent of the downstream learner [7, 16, 17]. OT valuation compares the training set with a small clean validation set using a cost that combines feature distance and label information, then uses the resulting dual potentials to estimate how much each training sample contributes to the transport. In LAVA [16], for example, the label cost measures the distance between feature distributions associated with two label groups, one from training and one from validation. The difficulty in CLL is that these two sets use labels with different meanings: a training label c means "not class c", while a validation label c means "class c".

This semantic mismatch changes the valuation in a systematic way. Standard OT label costs implicitly treat samples carrying label c as examples of class c, but complementary labels violate this assumption by construction. The training

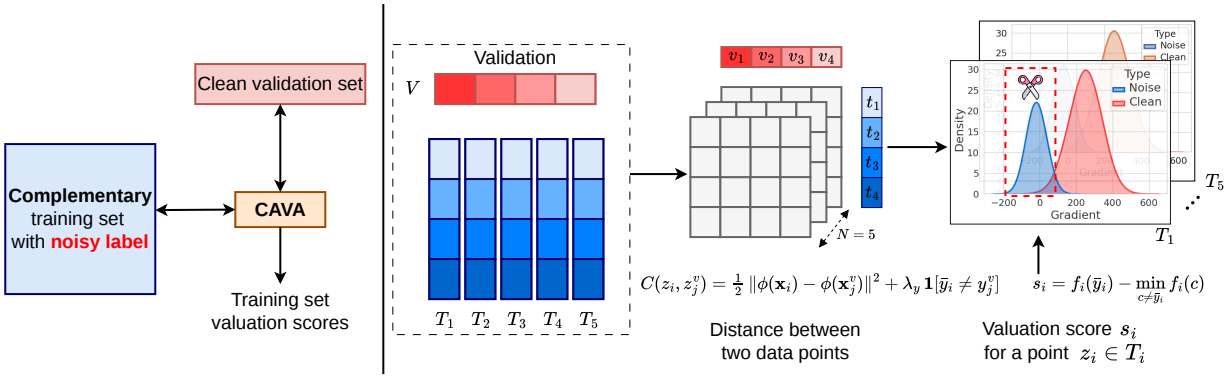

Figure 1. **Overview of CAVA.** The pool is split into tasks, and each task is matched to a clean validation set by entropic OT with the complementary label cost. Each sample is scored by its label margin; samples with a negative margin are flagged as diagonal noise and then removed or recycled.

group labeled c is mostly a mixture of the *other* classes; class c appears in that group only through annotation noise. As a result, the transferred label cost acquires an effective weight whose sign depends on the noise rate: it is positive above the crossover rate $1/M$ for M classes, zero at $1/M$, and negative below it (Figure 2, left). The resulting valuation is therefore informative when noise is frequent, uninformative at the crossover, and reversed when noise is rare, where it can rank corrupted samples as the most trustworthy ones. This distinction is easy to miss because standard synthetic settings often operate above the crossover, where direct transfer appears to work. Real annotation errors are typically much rarer and lie below it, where the transferred score can perform worse than random. The same issue also affects valuation methods that interpret a complementary label as an ordinary class label, including KNN-Shapley [15] and Data-OOB [19], which can similarly assign high value to diagonal noise (Figure 10 in Appendix C.4).

We address this problem with CAVA, a complementary-label-aware data valuation method for CLL (Figure 1). CAVA keeps the OT formulation but changes the two components through which the semantic mismatch enters. First, the *complementary label cost* is defined between clean classes, so noisy complementary-label groups do not determine the cost itself (Figure 2, right). Second, the *label margin* measures how much the observed complementary label contributes to the transport potential of a sample, using the same OT solution (Figure 3). For a diagonal-noise sample, the observed complementary label is its true class, which the clean validation set matches cheaply; its margin is negative whenever the validation set places the image in that class. Unlike the transferred score, the sign of the margin has the same meaning at every noise rate. It therefore gives an *adaptive cut*: the threshold is fixed at zero, the fraction of flagged samples adapts to the unknown noise rate without a pruning ratio, and the precision of this cut has a closed form in terms of the accuracy of the induced posterior classifier. CAVA uses that precision to decide what to do with flagged samples: they can be removed from the complementary-label pool or, when the detection is sufficiently reliable, recycled as ordinary-label supervision so that the detected error itself becomes useful training signal.

Our contributions are as follows.

- We identify a *flip law* for transferred data valuation under complementary labels: the effective label-cost weight changes sign with the noise rate, making existing OT valuation blind at $1/M$ and reversed below it, which is the regime relevant to real annotation errors.
- We propose CAVA, an efficient denoising method for CLL, and show that the sign of its label margin has a fixed meaning across noise rates and numbers of classes. Its zero threshold needs no pruning ratio and has a closed-form precision that can be used to control the risk of recycling flagged samples.
- On synthetic and human-annotated benchmarks, CAVA separates diagonal noise with AUC above 0.95 on label-free representations, and almost perfectly on a supervised one, in settings where transferred valuation performs worse than random; it improves a wide range of CLL objectives using only a fraction of the validation data, and scales to a one-million-image augmented pool. We also show that the detection is the agreement of a validation-induced posterior with the given label, which a fitted classifier matches, so the contribution lies in the corrected valuation and the recycling it enables, not in the transport solve.

2. Complementary-Label Learning

Problem setup. Let $\mathcal{X}$ be the input space and $\mathcal{Y} = \{1, \dots, M\}$ the label space. Each instance $\mathbf{x} \in \mathcal{X}$ has a true label $y \in \mathcal{Y}$ that is never observed during training. In complementary-label learning (CLL) [13], the annotation is a *complementary label* $\bar{y} \in \mathcal{Y}$ that names a class the instance does *not* belong to, so a correct one satisfies $\bar{y} \neq y$. The

learner receives a pool $\bar{\mathcal{D}} = \{(\mathbf{x}_i, \bar{y}_i)\}_{i=1}^N$ drawn i.i.d. from an unknown distribution and seeks a classifier $g : \mathcal{X} \to \mathbb{R}^M$ whose prediction $\arg\max_k g_k(\mathbf{x})$ recovers the true label.

Diagonal noise. Since an annotation may be wrong, we call $\bar{y}_i$ the *given label*. It is *valid* if $\bar{y}_i \neq y_i$, whichever wrong class it names, and *diagonal noise* if $\bar{y}_i = y_i$, the only case in which the exclusion itself is wrong. We write $\mathcal{N} = \{i : \bar{y}_i = y_i\}$ and $\eta = \mathbb{P}(\bar{y} = y)$ for the noise rate. No assumption is made on how valid labels are distributed; the analysis of Section 3 uses the uniform model of Assumption 1, which is also our synthetic protocol.

Learning objective. A CLL method minimizes the empirical counterpart of the complementary risk

$$R_{\mathrm{comp}}(g) = \mathbb{E}_{(\mathbf{x},\bar{y})}\big[\mathcal{L}_{\mathrm{comp}}\big(\bar{y}, g(\mathbf{x})\big)\big], \tag{2.1}$$

where $\mathcal{L}_{\mathrm{comp}}$ is a complementary-label loss such as the unbiased estimator [14], SCL [4], forward correction [39], or CPE [22]. Each of these losses decreases when predicted probability moves away from the class named by $\bar{y}$, so a diagonal-noise sample pushes the model away from its *true* class [27, 35].

Valuation setting. As in OT-based data valuation [16, 17], a small clean validation set $\mathcal{D}_v = \{(\mathbf{x}_j^v, y_j^v)\}_{j=1}^{N_v}$ with ordinary labels is available for valuation only, never for training. The task is to score every sample of $\bar{\mathcal{D}}$ from $\bar{\mathcal{D}}$, $\mathcal{D}_v$, and a fixed feature extractor ϕ (Section 3.1), so that the ranking separates $\mathcal{N}$ from the valid samples, and then to remove the flagged samples or recycle them with $\bar{y}_i$ as an ordinary label. The selection observes neither y_i nor $\mathcal{N}$; the only supervised component is the fixed feature extractor ϕ, whose provenance is stated in Appendix C.2.

3. CAVA: Complementary-Label Aware Data Valuation

We first describe the direct transfer of OT valuation to CLL, which we call TW-LAVA, and show why it fails (Sections 3.1 and 3.2). We then introduce the two components of CAVA, the complementary label cost and the label margin (Section 3.3), and the selection and recycling built on them (Sections 3.4 and 3.5). Proofs are in Appendix B.1.

3.1. TW-LAVA: OT valuation transferred to CLL

TW-LAVA keeps the label cost and the score of LAVA [16] and adds a task-wise decomposition that keeps each transport problem small. The pool $\bar{\mathcal{D}}$ is split into T equal-size random tasks $I_1, \dots, I_T \subseteq \{1, \dots, N\}$, and each task is valued against the validation set $\mathcal{D}_v$ independently. For task t, let $\mu_t = \sum_{i \in I_t} a_i \delta_{z_i}$ with $z_i = (\mathbf{x}_i, \bar{y}_i)$ and $\nu = \sum_{j=1}^{N_v} b_j \delta_{z_j^v}$ with $z_j^v = (\mathbf{x}_j^v, y_j^v)$, with uniform weights a_i and b_j. The ground cost combines a feature term and a label term,

$$C(z_i, z_j^v) = \lambda_{\mathbf{x}}\, C_{\mathrm{feat}}(\mathbf{x}_i, \mathbf{x}_j^v) + \lambda_y\, C_{\mathrm{label}}(\bar{y}_i, y_j^v), \tag{3.1}$$

with the feature cost

$$C_{\mathrm{feat}}(\mathbf{x}_i, \mathbf{x}_j^v) = \tfrac{1}{2}\big\|\phi(\mathbf{x}_i) - \phi(\mathbf{x}_j^v)\big\|_2^2 \tag{3.2}$$

of a fixed feature extractor ϕ, whose outputs are L2-normalized and rescaled to a median feature cost of one; we then set $\lambda_{\mathbf{x}} = 1$, so $\lambda_y > 0$ is the only cost weight. LAVA compares label groups across the two sets, which under complementary labels gives the *noisy-group label cost*

$$C_{\mathrm{label}}(\bar{y}_i, y_j^v) = \tfrac{1}{2} W_2^2\big(\bar{\mathcal{D}}_{\bar{y}_i}, \mathcal{D}^v_{y_j^v}\big) =: W^{\mathrm{grp}}(\bar{y}_i, y_j^v), \tag{3.3}$$

where $\bar{\mathcal{D}}_c$ is the feature distribution of the pool samples with given label c and $\mathcal{D}_k^v$ that of the validation samples of class k. Each task solves the entropic OT problem

$$W_{C,\varepsilon}(\mu_t, \nu) = \min_{\pi \in \Pi(a,b)} \sum_{i,j} \pi_{ij} C(z_i, z_j^v) + \varepsilon \sum_{i,j} \pi_{ij} \log \pi_{ij}, \tag{3.4}$$

with $\varepsilon = \varepsilon_{\mathrm{OT}} > 0$ and optimal pool-side and validation-side potentials $u^{(t),*}$ and $w^{(t),*}$. TW-LAVA scores sample i by the calibrated potential of LAVA,

$$v_i^{(t)} = u_i^{(t),*} - \frac{1}{|I_t| - 1} \sum_{k \in I_t,\, k \neq i} u_k^{(t),*}, \tag{3.5}$$

the change of the transport cost when sample i receives more mass.

3.2. Why TW-LAVA fails: a flip at $\eta = 1/M$

The noisy-group label cost assumes that the samples carrying label c are instances of class c. Under complementary labels they are not. If classes are balanced and a given label equals the true label with probability η and is otherwise uniform over the other $M - 1$ classes, the group with label c is the mixture

$$\bar{\mathcal{D}}_c = (1 - \alpha)\, \bar{P} + \alpha\, P_c, \qquad \alpha(\eta) = \frac{M\eta - 1}{M - 1}, \tag{3.6}$$

where P_k is the feature distribution of class k and $\bar{P}$ the average over classes. The group leans toward its own class when noise is frequent and away from it when noise is rare, and the label cost inherits this lean.

Theorem 1 (Flip law)**.** *Under this noise model, (i) at $\eta = 1/M$ the given labels leave the ground cost for any label divergence, so no score computed from the transport can separate diagonal noise from valid samples; (ii) for a squared mean-embedding label divergence,*

$$W^{\mathrm{grp}}(c, k) = \alpha(\eta)\, W_D^{\mathrm{cl}}(c, k) + r_c + t_k, \tag{3.7}$$

where W_D^{cl} is the same divergence between clean classes and the shifts r_c, t_k do not change the transport plan.

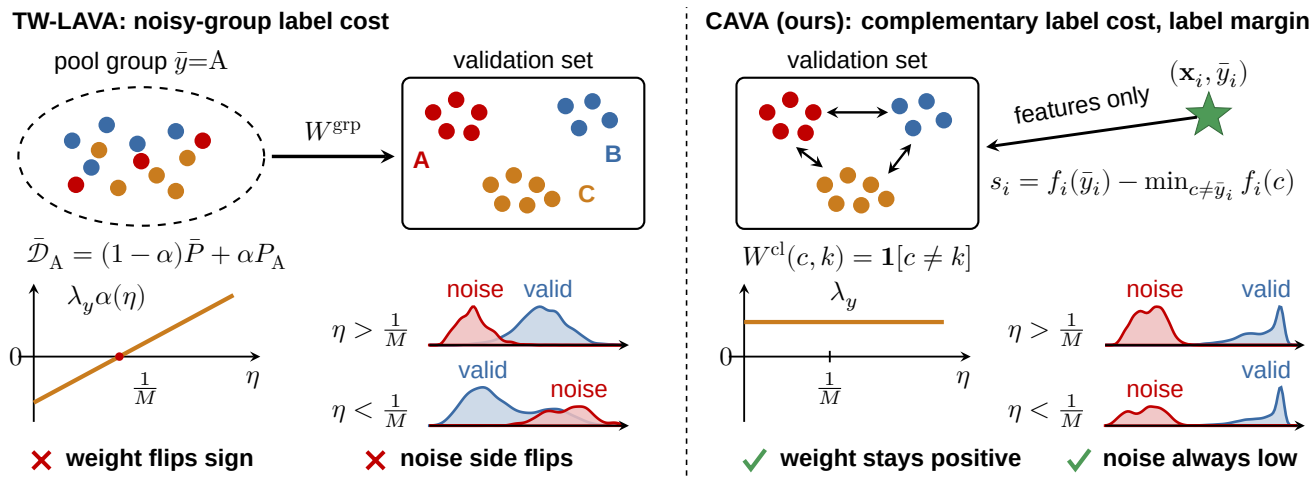

Figure 2. **From TW-LAVA to CAVA.** *Left*: TW-LAVA compares noisy label groups with clean classes, so its label weight $\lambda_y\alpha(\eta)$ changes sign at $\eta = 1/M$ and the noise switches sides (Theorem 1). *Right*: CAVA compares clean classes only, so the sign of its margin is fixed: negative exactly when the given label is the most likely class under the transport posterior (Theorem 2). Scores are measured on CIFAR-10 ($\eta = 25\%$, top) and CLCIFAR-10 (bottom).

At the level of the transport plan, TW-LAVA therefore uses a clean-class label cost with the *effective* weight $\lambda_y\alpha(\eta)$: positive above $\eta = 1/M$, zero at it, and negative below, where it rewards a given label that agrees with the validation class and reverses the score (the scores themselves coincide when the clean classes have equal radii; Appendix B.1). Part (ii) is exact for mean-embedding divergences; for the Sinkhorn class divergence of the pipeline we observe the same sign change with a smaller slope (Remark 1). Synthetic benchmarks sit above the crossover ($\alpha = 1/6$ at $\eta = 25\%$, $M = 10$), and human annotation errors sit below it ($\alpha < 0$ on CLCIFAR-10, whose labels are not uniform, so the law applies there only qualitatively). On CIFAR-10 the TW-LAVA score crosses chance exactly at $\eta = 1/M$ and is worse than random below it (Figure 9, Appendix B.1).

3.3. Complementary label cost and label margin

CAVA keeps the transport problem of TW-LAVA and changes the two ingredients that the flip runs through (Figure 2, right).

Complementary label cost. The label cost is defined between clean classes, so the noisy groups never enter the cost and the given label is only looked up:

$$C_{\mathrm{label}}(\bar{y}_i, y_j^v) = W^{\mathrm{cl}}(\bar{y}_i, y_j^v), \qquad W^{\mathrm{cl}}(c,k) = \mathbf{1}[c \neq k]. \tag{3.8}$$

The class-distance variant W_D^{cl} of Theorem 1 performs within 0.01 AUC of this 0/1 cost (Appendix B.1), and the 0/1 cost makes the analysis exact.

Label margin. The calibrated score of Eq. (3.5) measures the effect of the *mass* of a sample, so a typical image with a wrong label and an atypical image with a right label can score alike. We measure instead what the *label* contributes.

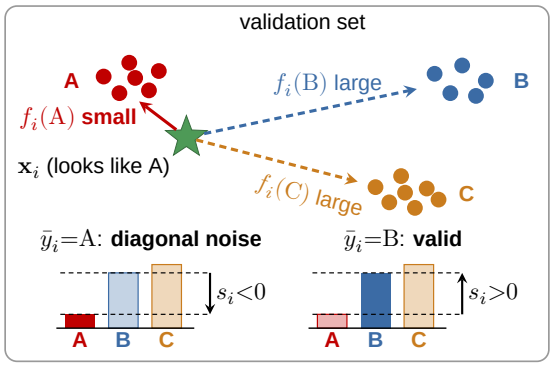

Figure 3. **The label margin.** The image looks like class A, so it is cheap to transport under label A and costly under B or C. The bars show $f_i(c)$, and the margin s_i compares the given label (solid bar) with the cheapest other label. Given A, $s_i < 0$ and the sample is flagged as diagonal noise; given B, $s_i > 0$ and the label is valid.

With the validation potential $w = w^{(t),*}$ of the problem under Eq. (3.8), the potential of sample i under label c is

$$f_i(c) = -\varepsilon \log \sum_j b_j \exp\Big(\tfrac{w_j - C_{\mathrm{feat}}(\mathbf{x}_i, \mathbf{x}_j^v) - \lambda_y \mathbf{1}[c \neq y_j^v]}{\varepsilon}\Big), \tag{3.9}$$

with $f_i(\bar{y}_i) = u_i^{(t),*}$, so no further transport problem is solved. The *label margin* compares the given label with its best alternative,

$$s_i = f_i(\bar{y}_i) - \min_{c \neq \bar{y}_i} f_i(c). \tag{3.10}$$

A diagonal-noise sample names its own class, the one label the validation set serves cheaply, so its margin is negative

whenever the transport places the image in that class (Figure 3).

Theorem 2 (Direction of the label margin)**.** *For the* $0/1$ *complementary label cost and any validation potential* w*,*

$$s_i = \psi\big(p_i(\bar{y}_i)\big) - \psi\Big(\max_{c \neq \bar{y}_i} p_i(c)\Big), \tag{3.11}$$

where $p_i(k) \propto \sum_{j:y_j^v=k} b_j\, e^{(w_j - C_{\mathrm{feat}}(\mathbf{x}_i, \mathbf{x}_j^v))/\varepsilon}$ *is the class posterior induced by the transport and* ψ *is strictly decreasing. Hence* $s_i \leq 0$ *if and only if* $\bar{y}_i$ *is the most likely class under* p_i*, for every noise rate and every* M*.*

The theorem fixes the sign of the margin relative to the posterior p_i, so $s_i \leq 0$ is a cut that needs no pruning ratio; how often diagonal noise is flagged is the accuracy of the posterior classifier on it (Theorem 3), which depends on the features and the validation set, not on η or M. The closed form also makes the score cheap: one softmax over the validation set gives p_i and then s_i. As $\varepsilon \to 0$ the margin reduces to a nearest-neighbor label check, while a finite ε keeps the transport weighting of the whole validation set.

3.4. Task-wise selection

CAVA selects within each task, which keeps every transport problem small and adapts the cut to each task. The *fixed cut* flags the fraction p of each task with the lowest margins, as prior OT valuation methods prune a fixed fraction [16, 17]. The *adaptive cut* needs no budget: its threshold is fixed at zero by Theorem 2, so it flags exactly the samples whose given label is the most likely class, and the fraction it removes adapts to the noise rate of each task without knowing η:

$$I_t^{\mathrm{rem}} = \{i \in I_t : \; s_i \leq 0\}, \qquad I^{\mathrm{keep}} = \bigcup_{t=1}^{T} \big(I_t \setminus I_t^{\mathrm{rem}}\big). \tag{3.12}$$

TW-LAVA has no such cut, because the sign of its score depends on whether η lies above or below $1/M$, which is unknown in practice; it is therefore used with the fixed cut. Without recycling, flagged samples are discarded and the classifier is trained on the retained pool,

$$\mathcal{L}_{\mathrm{sel}} = \frac{1}{|I^{\mathrm{keep}}|} \sum_{i \in I^{\mathrm{keep}}} \mathcal{L}_{\mathrm{comp}}\big(\bar{y}_i, g(\mathbf{x}_i)\big), \tag{3.13}$$

where $\mathcal{L}_{\mathrm{comp}}$ is any CLL loss; for the fixed cut, I^{keep} is the complement of the flagged tails.

3.5. Recycling gated by precision

A diagonal-noise label satisfies $\bar{y}_i = y_i$, so it is wrong as an exclusion but correct as an ordinary label. CAVA can therefore recycle the flagged set $I^{\mathrm{rec}} = \bigcup_t I_t^{\mathrm{rem}}$ into a cross-entropy branch and train

$$\begin{aligned}\mathcal{L}_{\mathrm{CAVA}} = {} & \frac{\kappa}{|I^{\mathrm{rec}}|} \sum_{i \in I^{\mathrm{rec}}} \mathcal{L}_{\mathrm{CE}}\big(\bar{y}_i, g(\mathbf{x}_i)\big) \\ & + \frac{1-\kappa}{|I^{\mathrm{keep}}|} \sum_{i \in I^{\mathrm{keep}}} \mathcal{L}_{\mathrm{comp}}\big(\bar{y}_i, g(\mathbf{x}_i)\big),\end{aligned} \tag{3.14}$$

with weight $\kappa \in [0, 1]$. A falsely recycled valid sample injects a wrong ordinary label, so recycling helps only when the flagged set is pure. The adaptive cut makes this purity predictable: by Theorem 2, a sample is flagged exactly when the posterior classifier $h_i = \arg\max_k p_i(k)$ agrees with its given label.

Theorem 3 (Precision-gated recycling)**.** *Let* $F = \{i : s_i \leq 0\}$ *and let* e *be the error rate of* h *against the true labels. Under the noise model of Section 3.2, if the errors of* h *are independent of the label draw (*h *depends on the given labels only through the validation potential* w*), the precision of* F *is*

$$\pi = \frac{\eta\,(1-e)}{\eta\,(1-e) + (1-\eta)\,e/(M-1)}, \tag{3.15}$$

which exceeds η *whenever* $e \leq 1 - 1/M$*. Conditionally on the selection, training the cross-entropy branch on* F *has excess true-label risk at most* $2B(1-\pi) + 2B|\pi - (1-e)| + \Delta_{\mathrm{gen}}(|F|, \delta)$ *with probability* $1-\delta$*, where* B *bounds the loss, the middle term is the distribution shift of the flagged set, and* Δ_{gen} *vanishes as* $|F|$ *grows (Theorem 4, Remark 2).*

A valid sample is flagged only when h errs *and* its wrong guess equals the given label, an extra factor $1/(M-1)$, which is why the flagged set is far purer than the base rate. Neither the bias term $2B(1-\pi)$ nor the shift term, which is about e when the flagged set is pure, shrinks with more data, so the precision and the recall of the posterior classifier decide the mode: recycling when it is high, selection only (Eq. (3.13)) otherwise. Eq. (3.15) needs only e, which can be estimated from the validation set; in our experiments we report the precision measured with the true labels after selection (Appendix C.3).

4. Experiments

4.1. Experimental setup

We evaluate CAVA on MNIST [20], CIFAR-10, CIFAR-20 [18], the real-world CLImage datasets CLCIFAR-10/20 and CLMIN-10/20 [35], an enlarged CIFAR-10-1M setting, and three further domains, FOOD15, RESISC45-20, and Places365-10, built from Food-101 [1], NWPU-RESISC45 [2], and Places365 [41] (Appendix C.1). All label spaces have at most 20 classes, the scale at which current CLL methods are workable: a complementary label excludes a single class, so its supervision weakens as the

Table 1. **Main results: test accuracy (%) with AutoAugment.** All methods score the pool against the same 2,000 validation images (20 per class on CLMIN), and accuracy is measured after the final epoch on the remaining test images. The fixed cut removes $p = 25\%$ of each task on the synthetic datasets and 5% (3% on CLCIFAR-20) on the human-labeled ones; the adaptive cut removes every sample with $s_i \leq 0$ and needs no ratio. Mean±std over three seeds; best per column and block in **bold**.

		Synthetic ($\eta = 25\%$)			**Real-world**			
Method	**Selection**	MNIST	CIFAR-10	CIFAR-20	CLCIFAR-10	CLCIFAR-20	CLMIN-10	CLMIN-20
(a) Selection only: SCL-NL on the retained pool, Eq. (3.13)								
No pruning	–	0.14±0.09	1.19±0.02	0.49±0.10	42.36±0.91	7.71±0.08	14.89±1.02	4.55±1.18
TW-LAVA	Fixed p	98.21±0.29	76.04±0.25	14.43±1.00	47.46±0.17	6.90±0.11	12.56±0.84	4.28±0.35
CAVA (ours)	Fixed p	99.20±0.10	**79.26**±0.25	**17.70**±1.38	79.63±0.33	**10.78**±1.25	**27.44**±2.69	**10.39**±0.67
CAVA (ours)	Adaptive cut	**99.47**±0.04	79.12±0.13	12.54±2.68	**80.00**±0.59	10.27±1.15	**27.44**±1.84	9.56±1.13
(b) Selection with recycling: CE+SCL-NL, Eq. (3.14)								
TW-LAVA	Fixed p	97.68±0.14	88.75±0.40	59.03±0.77	38.71±0.85	10.24±0.44	23.33±0.67	5.33±0.73
CAVA (ours)	Fixed p	99.36±0.07	91.57±0.14	**75.34**±0.17	81.84±0.15	47.71±0.38	**52.78**±0.19	35.22±1.11
CAVA (ours)	Adaptive cut	**99.66**±0.03	**91.97**±0.14	73.05±0.32	**85.69**±0.29	**48.86**±0.66	52.33±1.45	**36.11**±1.19

number of classes grows [29]. For synthetic experiments, each training example receives an observed complementary label $\bar{y}_i$: with probability η, we set $\bar{y}_i = y_i$ to create diagonal complementary-label noise; otherwise, we sample a valid complementary label uniformly from $\mathcal{Y} \setminus \{y_i\}$. The test set always uses clean ground-truth labels. Unless otherwise stated, every valuation method uses the same validation set of $N_v = 2{,}000$ clean images, a class-stratified subset of the test split. The validation images are used only to score the training samples: they never enter model training, model selection, or the training objective. Test accuracy is measured after the final training epoch on the remaining 8,000 test images, which the valuation never sees, so no quantity is selected on the evaluation data. CLMIN-10 and CLMIN-20 contain 64×64 ImageNet images, which we resize to 32×32; their test splits hold only 50 images per class, so 20 per class form the validation set and the other 30 per class the evaluation set. This assumption is shared by the valuation baselines and is examined in Appendix D.8. All valuation methods share one fixed feature extractor. The main tables use the supervised embedder of LAVA, which was trained with the clean labels of the training split and is therefore an upper bound on feature quality (Appendix C.2); Appendix C.5 repeats the detection and training experiments with two extractors that have seen no label of the pool, CLIP features and a network trained on the validation images only.

We compare against No Pruning and the transferred OT valuation TW-LAVA, and at small noise rates also against Random Pruning, SAVA [17], KNN-Shapley [15], Data-OOB [19], and small-loss selection [11]; LAVA itself enters the runtime and memory comparison (Appendix D.10). For CLL objectives, we evaluate FWD [39], CPE [22], OP [25], SCARCE [37], SCL-NL, and SCL-EXP [4]. All models use ResNet-18 [12] and are trained under the same protocol across compared methods. Full hyperparameters, seeds, single-run settings for expensive scalability experiments, and additional dataset details are provided in Appendix C. Table 2 fixes the names of the three OT valuation methods: LAVA is the original full-set method, TW-LAVA is its task-wise transfer to complementary labels (Section 3.1), and CAVA replaces the label cost and the score (Section 3.3).

Table 2. The three OT valuation methods compared in this paper. TW-LAVA differs from LAVA only in the task-wise decomposition; CAVA differs from TW-LAVA only in the label cost and the score.

Method	Valuation scope	Label cost	Score
LAVA [16]	full set (≤20K)	noisy-group	calibrated
TW-LAVA	task-wise, all data	noisy-group	calibrated
CAVA (ours)	task-wise, all data	complementary	label margin

4.2. Results and analysis

Main results. Table 1 compares CAVA with the TW-LAVA baseline under one validation set. With the same fixed cut, CAVA is higher than TW-LAVA in all 14 cells, and without selection SCL-NL collapses on the synthetic datasets (at most 1.19%). The gap is largest on the human-annotated datasets, whose error rates lie below $1/M$: there, TW-LAVA does not beat training on the full pool on CLCIFAR-20 and both CLMIN datasets, as the flip law predicts, while CAVA with recycling reaches 85.69%, 48.86%, 52.78%, and 36.11% on CLCIFAR-10, CLCIFAR-20, CLMIN-10, and CLMIN-20, against at most 38.71%, 10.24%, 23.33%, and 5.33% for TW-LAVA with recycling. The adaptive cut, which needs no pruning ratio, is best or tied in eight of the 14 CAVA comparisons and within one point of the fixed cut except on CIFAR-20, where it flags fewer noisy samples.

Beyond CIFAR. Figure 7 repeats the comparison on three further domains, food, aerial scenes, and places, at three

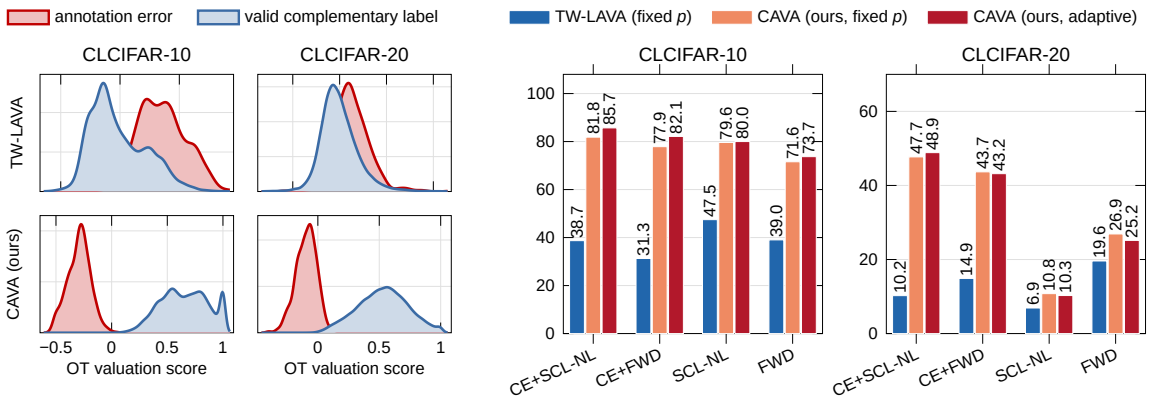

(a) OT valuation scores (b) Test accuracy (%)

Figure 4. **Real-world complementary labels.** (a) Score distributions: under TW-LAVA (top), annotation errors overlap with valid labels and sit on the wrong side; under CAVA (bottom), they separate on the low side. (b) Test accuracy of TW-LAVA and CAVA at the fixed cut of Table 1 (5% on CLCIFAR-10, 3% on CLCIFAR-20) and of CAVA with the adaptive cut, mean of three seeds; all use the same 2,000 validation images.

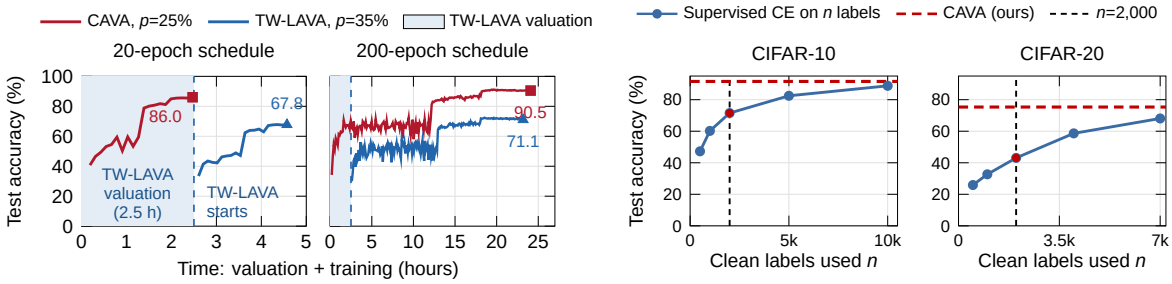

(a) Scaling to one million images (CIFAR-10-1M) (b) Clean labels: validation set vs. training data

Figure 5. **Scale and the value of clean labels.** (a) CIFAR-10-1M ($\eta = 25\%$, SCL-NL): test accuracy after every epoch against time (valuation plus training, measured on one dedicated GPU), each method at its best cut (CAVA $p = 25\%$, TW-LAVA $p = 35\%$), for the 20-epoch and the 200-epoch schedule. Shaded: the valuation of TW-LAVA, so its training starts only after 2.5 hours (dashed line); CAVA needs about 4 minutes. All cuts are in Table 23. (b) Supervised CE trained from scratch on n clean labels (three seeds) against CAVA, which uses 2,000 clean labels only as its validation set ($\eta = 25\%$); CAVA uses the supervised embedder of the main tables, and results on other fixed representations are in Table 21.

noise rates (Table 18 in Appendix D.3 lists all rows). With recycling, CAVA is far above TW-LAVA at every noise rate, with either cut, and the gap is largest at $\eta = 5\%$, near the crossover $1/M$, where TW-LAVA reaches only 8.27%, 3.00%, and 38.95% against 47.33%, 71.93%, and 80.49% for CAVA. Without recycling, CAVA with the fixed cut is above TW-LAVA in every setting, by 0.4 points on RESISC45-20 at $\eta = 25\%$ and by more elsewhere, while its adaptive cut falls below TW-LAVA on FOOD15 and RESISC45-20 at $\eta = 25\%$. At $\eta = 5\%$, the selection of TW-LAVA stays within 4 points of training on the full pool, whereas CAVA gains 14 to 30 points.

Small noise ratios. Real complementary labels are mostly correct: both CLCIFAR datasets have noise rates below $1/M$, the regime in which the noisy-group label cost of TW-LAVA reverses. Figure 6 reproduces this regime on CIFAR-10 with $\eta \in \{2\%, 5\%, 10\%\}$ and trains SCL-NL on the retained pool. All methods use the same feature extractor, the same 2,000 validation images, and the same task pools; each baseline is used as published, keeping the samples it values most (SAVA [17], KNN-Shapley [15], Data-OOB [19]) or the samples with the smallest loss [11]. CAVA peaks at the cut $p = \eta$ at $\eta = 5\%$ and 10% (82.2% and 81.4%) and at $p = 5\%$ at $\eta = 2\%$ (82.6%; 82.3% at $p = 2\%$), against 68.5%, 38.7%, and 7.9% without pruning. TW-LAVA removes valid samples first, so pruning lowers its accuracy at $\eta = 2\%$ and 5% and adds at most 4.6 points at $\eta = 10\%$. Among the baselines, only Data-OOB keeps up with CAVA, and only below the crossover: it reaches 82.6% and 82.0% at $\eta = 2\%$ and 5%, where its natural direction is the correct one (Figure 10 in Appendix C.4), but 11.0% at $\eta = 10\% = 1/M$, where its score is blind. Random pruning, KNN-Shapley,

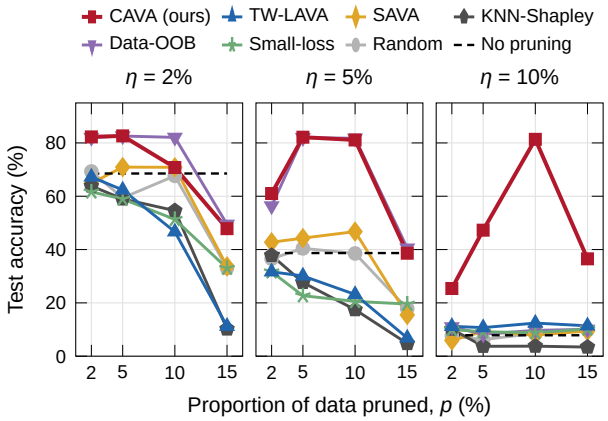

Figure 6. **Small noise ratios.** Final-epoch test accuracy (held-out test images) of SCL-NL on the retained pool against the pruning ratio p, on CIFAR-10 with small diagonal-noise ratios η; CAVA, TW-LAVA, and no pruning are means over three seeds, the other baselines single runs (seed 1126); same validation set for all methods; TW-LAVA keeps its original OT settings here (Appendix C.2). Each baseline is used as published, keeping its highest-valued or lowest-loss samples. The dashed line trains on the full pool.

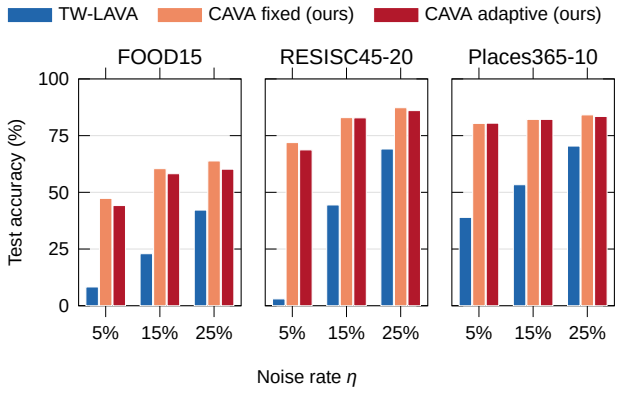

Figure 7. **Beyond CIFAR.** Test accuracy with recycling (CE+SCL-NL) on FOOD15 (Food-101 [1]), RESISC45-20 (NWPU-RESISC45 [2]), and Places365-10 (Places365 [41]) with synthetic diagonal noise; protocol of Table 1, fixed cut $p = \eta$, seed 1126. All rows, including selection only, are in Table 18.

and small-loss selection stay near or below no pruning at every noise rate, SAVA gains at most 8.1 points at $\eta = 5\%$, and at $\eta = 10\%$ no baseline exceeds 12.4%, while CAVA reaches 81.4%.

Selection without a pruning ratio. In practice the noise rate is unknown, so a fixed ratio p is hard to tune. The adaptive cut of Eq. (3.12) needs no ratio: it flags a different fraction in each task, and on MNIST, CIFAR-10, CIFAR-20, and the two CLCIFAR datasets its flagged set reaches 96.0% to 99.98% precision and 87.2% to 99.9% recall (Tables 6 and 14), with a single OT solve per task. A score whose direction depends on η, such as that of TW-LAVA, admits no such threshold. In training, the adaptive cut performs on par with the best fixed cut without choosing a ratio, except on CIFAR-20, and with recycling it is the best setting on five of the seven datasets (Table 1).

Why existing valuation fails, and what fixes it. Figure 10 in Appendix C.4 scores every method on the same pool and tasks, at the small noise rates of real complementary labels. KNN-Shapley [15] credits a sample whose label agrees with its validation neighbors, and a diagonal label agrees by construction, so it ranks the noise as the most valuable data on every dataset. Data-OOB [19] and small-loss selection [11] learn from the given labels themselves, so their direction follows the flip law: they are correct below the crossover $1/M$ (CIFAR-10 at $\eta = 5\%$ and both CLCIFAR datasets), near chance on CIFAR-20 at $\eta = 5\% = 1/M$, and reversed above it (both datasets at $\eta = 15\%$). TW-LAVA follows the same law with the opposite sign: it is reversed on CIFAR-10 at $\eta = 5\%$ and on both human-annotated datasets, near chance on CIFAR-20 at $\eta = 5\%$, and correct only above $1/M$, at $\eta = 15\%$, as Theorem 1 predicts for the synthetic data and consistently with it on the human labels. Every score that reads a complementary label as an ordinary one therefore has a direction that depends on whether η lies above or below $1/M$, which is unknown in practice. CAVA reaches the same detection inside the OT valuation framework, where the cut at zero and the precision of Theorem 3 come with the score at no extra cost. The transport solve itself is not what detects the noise: Theorem 2 holds for any validation potential, and on CLIP features the margin with the potential set to zero is within 0.006 AUC of CAVA on each of the five datasets where we computed it, and the marginal constraint of the transport lowers the AUC when the class marginal of the pool and of the validation set differ (Appendix C.5). The gain over existing valuation comes from the complementary label cost, which removes the flip, and from the margin, which isolates the label; the transport problem is the framework in which both are defined, not a better detector.

Recycling detected diagonal-noise samples further improves both methods on CIFAR-10 and CIFAR-20 and CAVA on all three synthetic datasets (Table 1, block (b) against block (a)); for TW-LAVA with the fixed cut on CIFAR-10, accuracy rises from 76.04% to 88.75%. Detected diagonal-noise samples thus provide useful ordinary-label supervision when they are recycled with sufficient precision.

Real-world complementary labels. CLCIFAR-10 and CLCIFAR-20 contain 3.98% and 2.32% annotation errors, both below $1/M$, which is the regime where the label cost of TW-LAVA reverses (Theorem 1). Figure 4(a) shows the consequence. With the noisy-group label cost, annotation errors and valid labels overlap, and at $p = 3\%$ only 43.5% (CLCIFAR-10) and 9.9% (CLCIFAR-20) of the recycled

samples are true errors (Tables 16 and 17). With the complementary label cost and the label margin, the two groups separate almost completely under the same feature extractor and validation set: the recycled branch at $p = 3\%$ is 100% true errors on CLCIFAR-10 and 75.9% on CLCIFAR-20, where it already contains 98.0% of all errors, and the adaptive cut on CLCIFAR-10 recycles $1{,}990$ samples at 99.7% precision. This selection quality carries over to training (Figure 4(b)). CE+FWD rises from 31.28% to 77.87% on CLCIFAR-10 (82.14% with the adaptive cut) and from 14.89% to 43.67% on CLCIFAR-20, and every objective improves on both datasets, including SCL-NL, which needs no transition matrix and reaches 80.00% on CLCIFAR-10 with the adaptive cut, above known-Q FWD. The gain is largest for the recycling objectives, whose CE branch depends directly on detector precision (Theorem 3). The margin adds no OT solve: it reuses the validation potential of the same transport problem and evaluates M soft-minima per sample. Full results are given in Appendix D.2.

Is the validation set better spent on training? Trained from scratch with supervised CE, the $2{,}000$ clean labels reach 71.52% on CIFAR-10 and 43.08% on CIFAR-20, against 91.57% and 75.34% for CAVA, which uses them only for scoring (Figure 5(b)). CAVA keeps this advantage on a fixed representation that has seen no label of the pool: with CLIP features it reaches 87.91% and 67.37% (Appendix D.8), and the corresponding results on CLIP features for every dataset are in Appendix C.5.

Scalability on the large-scale CIFAR-10-1M benchmark. To evaluate CAVA beyond the regime where full-set OT valuation is feasible, we enlarge CIFAR-10 to one million training samples (CIFAR-10-1M) via label-preserving augmentation under the same diagonal-noise protocol ($\eta = 25\%$). Task-wise decomposition splits the valuation into five tasks of 200,000 samples, and CAVA scores all one million samples in about 4 minutes. CAVA is best at the cut $p = 25\%$ that equals the noise rate and reaches 86.04% with SCL-NL after 20 epochs, against 0.12% without pruning. As Figure 5(a) shows, it finishes this training before the valuation of TW-LAVA does, and TW-LAVA at its best cut $p = 35\%$ reaches 67.76% under the same 20-epoch schedule and 71.05% after 200 epochs, against 90.51% for CAVA. Full results are in Appendix D.9.

5. Conclusion

We have introduced CAVA, an OT-based data valuation framework for CLL. Starting from the task-wise transfer of LAVA, we showed that its noisy-group label cost carries an effective label weight $\lambda_y \alpha(\eta)$ that changes sign at the noise rate $\eta = 1/M$: the direction of its score depends on the noise rate, and the method is blind at the crossover, which is the regime of human-annotated complementary labels. CAVA replaces this cost by a complementary label cost defined between clean classes and scores each sample by a label margin, the part of its transport potential that the given label contributes. The sign of the margin has the same meaning on every dataset, flags diagonal noise with an adaptive cut that needs no pruning ratio, and comes with a closed-form precision that gates the recycling of detected samples as ordinary-label supervision. On synthetic and real-world benchmarks, CAVA with recycling is the best of the compared valuation methods on every dataset, and its gains are largest on real complementary labels, where the transferred method fails. The detection itself is matched by a classifier fit on the validation set, and the transport potential adds nothing to it; what the analysis contributes is the reason transferred valuation fails, a cost under which it does not, and a recycling rule gated by a precision that can be estimated from the validation set.

6. Limitations

Our work has several limitations. First, we evaluate CAVA primarily on relatively small-scale CLL benchmarks; extending OT-based valuation to truly large-scale complementary-label datasets remains an open challenge [28, 38]. Second, our framework assumes access to a clean validation set representative of the downstream task, a common assumption in data valuation research that may not always hold in practice; Appendix D.8 quantifies it. Third, recycling is most directly applicable to diagonal complementary-label noise, where a violated exclusion label reveals ordinary-label information. Addressing scalability, reducing reliance on clean validation sets, and generalizing recycling beyond diagonal noise are important future directions.

References

[1] Lukas Bossard, Matthieu Guillaumin, and Luc Van Gool. Food-101 – mining discriminative components with random forests. In *European Conference on Computer Vision (ECCV)*, pages 446–461, 2014. 5, 8, 9

[2] Gong Cheng, Junwei Han, and Xiaoqiang Lu. Remote sensing image scene classification: Benchmark and state of the art. *Proceedings of the IEEE*, 105(10):1865–1883, 2017. 5, 8, 9

[3] Hongliang Chi, Qiong Wu, Zhengyi Zhou, and Yao Ma. Shapley-guided utility learning for effective graph inference data valuation. In *The Thirteenth International Conference on Learning Representations*, 2025. 2

[4] Yu-Ting Chou, Gang Niu, Hsuan-Tien Lin, and Masashi Sugiyama. Unbiased risk estimators can mislead: A case study of learning with complementary labels. In *International Conference on Machine Learning*, pages 1929–1938, 2020. 1, 3, 6, 9

[5] Marco Cuturi. Sinkhorn distances: Lightspeed computation of optimal transport. *Advances in Neural Information Processing Systems*, 26, 2013. 1

[6] Lei Feng, Takuo Kaneko, Bo Han, Gang Niu, Bo An, and Masashi Sugiyama. Learning with multiple complementary labels. In *International Conference on Machine Learning*, pages 3072–3081, 2020. 9

[7] Lan Feng, Fan Nie, Yuejiang Liu, and Alexandre Alahi. TAROT: Targeted data selection via optimal transport. In *International Conference on Machine Learning*, pages 16837–16852, 2025. 1

[8] Benoît Frénay and Michel Verleysen. Classification in the presence of label noise: a survey. *IEEE transactions on neural networks and learning systems*, 25:845–869, 2013. 23

[9] Aude Genevay, Gabriel Peyré, and Marco Cuturi. Learning generative models with sinkhorn divergences. *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Statistics*, pages 1608–1617, 2018. 1

[10] Amirata Ghorbani and James Zou. Data shapley: Equitable valuation of data for machine learning. *Proceedings of the 36th International Conference on Machine Learning*, pages 2242–2251, 2019. 1, 2, 23

[11] Bo Han, Quanming Yao, Xingrui Yu, Gang Niu, Miao Xu, Weihua Hu, Ivor Tsang, and Masashi Sugiyama. Co-teaching: Robust training of deep neural networks with extremely noisy labels. *Advances in Neural Information Processing Systems*, 31, 2018. 6, 7, 8, 10, 23

[12] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition, 2015. 6

[13] Takashi Ishida, Gang Niu, Weihua Hu, and Masashi Sugiyama. Learning from complementary labels. In *Advances in Neural Information Processing Systems*, page 5639–5649, 2017. 1, 2, 9

[14] Takashi Ishida, Gang Niu, Aditya Menon, and Masashi Sugiyama. Complementary-label learning for arbitrary losses and models. In *International Conference on Machine Learning*, 2019. 3

[15] Ruoxi Jia, David Dao, Boxin Wang, Francesco Hubis, Nick Hynes, Nezihe Merve Gürel, Bo Li, Ce Zhang, and Dawn Song. Efficient task-specific data valuation for nearest neighbor algorithms. *Proceedings of the VLDB Endowment*, 12 (11):1610–1623, 2019. 2, 6, 7, 8, 10

[16] Hoang Anh Just, Feiyang Kang, Jiachen T. Wang, Yi Zeng, Myeongseob Ko, Ming Jin, and Ruoxi Jia. Lava: Data valuation without pre-specified learning algorithms. In *International Conference on Learning Representations*, 2023. 1, 3, 5, 6, 2, 9, 22, 23

[17] Samuel Kessler, Tam Le, and Vu Nguyen. Sava: Scalable learning-agnostic data valuation. In *The Thirteenth International Conference on Learning Representations*, 2025. 1, 3, 5, 6, 7, 2, 9, 22, 23

[18] Alex Krizhevsky. Learning multiple layers of features from tiny images. Computer Science University of Toronto, Canada, 2009. 5, 8

[19] Yongchan Kwon and James Zou. Data-OOB: Out-of-bag estimate as a simple and efficient data value. In *International Conference on Machine Learning*, pages 18135–18152, 2023. 2, 6, 7, 8, 10

[20] Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86:2278–2324, 1998. 5, 8

[21] Junnan Li, Richard Socher, and Steven CH Hoi. Dividemix: Learning with noisy labels as semi-supervised learning. In *International Conference on Learning Representations*, 2020. 23

[22] Wei-I Lin and Hsuan-Tien Lin. Reduction from complementary-label learning to probability estimates. In *Pacific-Asia Conference on Knowledge Discovery and Data Mining*, pages 469–481, 2023. 3, 6, 10

[23] Wei-I Lin, Gang Niu, Hsuan-Tien Lin, and Masashi Sugiyama. Enhancing label sharing efficiency in complementary-label learning with label augmentation, 2023. 1

[24] Xiaoqiang Lin, Xinyi Xu, Zhaoxuan Wu, See-Kiong Ng, and Bryan Kian Hsiang Low. Distributionally robust data valuation. In *Proceedings of the 41st International Conference on Machine Learning*, pages 30362–30391, 2024. 1

[25] Shuqi Liu, Yuzhou Cao, Qiaozhen Zhang, Lei Feng, and Bo An. Consistent complementary-label learning via order-preserving losses. In *International Conference on Artificial Intelligence and Statistics*, pages 8734–8748, 2023. 6

[26] Alexander Long, Wei Yin, Thalaiyasingam Ajanthan, Vu Nguyen, Pulak Purkait, Ravi Garg, Alan Blair, Chunhua Shen, and Anton van den Hengel. Retrieval augmented classification for long-tail visual recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6959–6969, 2022. 25

[27] Tan-Ha Mai and Hsuan-Tien Lin. Intra-cluster mixup: An effective data augmentation technique for complementary-label learning. *Transactions on Machine Learning Research*, 2026. 1, 3

[28] Tan-Ha Mai, Nai-Xuan Ye, Yu-Wei Kuan, Po-Yi Lu, and Hsuan-Tien Lin. The unexplored potential of vision-language models for generating large-scale complementary-label learning data. In *Pacific-Asia Conference on Knowledge Discovery and Data Mining*, pages 90–102, 2025. 1, 9

[29] Tan-Ha Mai, Nai-Xuan Ye, Wei Wang, Gang Niu, Takashi Ishida, Masashi Sugiyama, and Hsuan-Tien Lin. Complementary-label learning: A survey of methods, an open-source toolkit, and fair benchmarking results. *ACM Computing Surveys*, 2026. 1, 6, 9

[30] Mehryar Mohri, Afshin Rostamizadeh, and Ameet Talwalkar. *Foundations of machine learning*. MIT press, 2018. 7

[31] Gabriel Peyré and Marco Cuturi. Computational optimal transport. *Foundations and Trends in Machine Learning*, 11 (5–6):355–607, 2019. 1

[32] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning (ICML)*, 2021. 11

[33] Masashi Sugiyama, Han Bao, Takashi Ishida, Nan Lu, Tomoya Sakai, and Gang Niu. *Machine learning from weak supervision: An empirical risk minimization approach*. MIT Press, 2022. 1

[34] Cédric Villani et al. *Optimal transport: old and new*. Springer, 2008. 1

[35] Hsiu-Hsuan Wang, Mai Tan Ha, Nai-Xuan Ye, Wei-I Lin, and Hsuan-Tien Lin. CLImage: Human-annotated datasets

for complementary-label learning. *Transactions on Machine Learning Research*, 2025. 1, 3, 5, 8, 9, 10

[36] Tianyu Wang, Yifei Sun, Rong Ge, and Ameet Talwalkar. Data valuation via optimal transport. *Proceedings of the 40th International Conference on Machine Learning*, 2023. 2

[37] Wei Wang, Takashi Ishida, Yu-Jie Zhang, Gang Niu, and Masashi Sugiyama. Learning with complementary labels revisited: The selected-completely-at-random setting is more practical. In *International Conference on Machine Learning*, pages 50683–50710, 2024. 6

[38] Nai-Xuan Ye, Tan-Ha Mai, Hsiu-Hsuan Wang, Wei-I Lin, and Hsuan-Tien Lin. libcll: an extendable python toolkit for complementary-label learning, 2024. 9

[39] Xiyu Yu, Tongliang Liu, Mingming Gong, and Dacheng Tao. Learning with biased complementary labels. In *European Conference on Computer Vision*, pages 68–83, 2018. 3, 6, 10

[40] Zilong Zhao, Yiming Li, Tianyi Chen, and Yu-Xiang Wang. Optimal transport for data selection and reweighting. *Advances in Neural Information Processing Systems*, 2024. 2

[41] Bolei Zhou, Agata Lapedriza, Aditya Khosla, Aude Oliva, and Antonio Torralba. Places: A 10 million image database for scene recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(6):1452–1464, 2018. 5, 8, 9, 28

[42] Kaiyang Zhou, Jingkang Yang, Chen Change Loy, and Ziwei Liu. Learning to prompt for vision-language models. *International Journal of Computer Vision*, 130:2337–2348, 2022. 1

[43] Deyao Zhu, Jun Chen, Xiaoqian Shen, Xiang Li, and Mohamed Elhoseiny. Minigpt-4: Enhancing vision-language understanding with advanced large language models. 2023. 1

[44] Jiongli Zhu, Parjanya Prajakta Prashant, Alex Cloninger, and Babak Salimi. KAIROS: Scalable model-agnostic data valuation. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. 2

CAVA: Complementary-Label-Aware Data Valuation

Supplementary Material

A. Background

A.1. Optimal Transport for Data Valuation

OT provides a principled framework for data valuation by quantifying how individual data points contribute to aligning empirical distributions in a task-aware and geometry-aware manner [24, 31, 34]. Let

$$\mu = \sum_{i=1}^{n} a_i \delta_{z_i}, \qquad \nu = \sum_{j=1}^{m} b_j \delta_{z'_j}, \tag{A.1}$$

denote empirical source and target distributions, where $a \in \Delta_n$ and $b \in \Delta_m$ are probability weights on the simplex. Here, z_i and z'_j denote generic source and target samples. In CAVA, these correspond to complementary-labeled training samples and ordinary-labeled validation samples, respectively. Given a ground cost $C(z_i, z'_j)$ that encodes task-relevant dissimilarity, the OT problem is defined as

$$W_C(\mu, \nu) = \min_{\pi \in \Pi(a,b)} \sum_{i=1}^{n} \sum_{j=1}^{m} \pi_{ij} C(z_i, z'_j), \tag{A.2}$$

where the transport polytope is

$$\Pi(a, b) := \Big\{ \pi \in \mathbb{R}_+^{n \times m} \ : \ \pi \mathbf{1} = a, \ \pi^\top \mathbf{1} = b \Big\}. \tag{A.3}$$

Each entry π_{ij} represents how much probability mass from source sample z_i is transported to target sample z'_j, providing a natural mechanism for attributing importance to individual data points.

To improve scalability and numerical stability, entropic regularization is commonly used, yielding the Sinkhorn-regularized OT objective [5]:

$$W_{C,\varepsilon_{\mathrm{OT}}}(\mu, \nu) = \min_{\pi \in \Pi(a,b)} \sum_{i=1}^{n} \sum_{j=1}^{m} \pi_{ij} C(z_i, z'_j) + \varepsilon_{\mathrm{OT}} \sum_{i=1}^{n} \sum_{j=1}^{m} \pi_{ij} \log \pi_{ij}, \tag{A.4}$$

where $\varepsilon_{\mathrm{OT}} > 0$ controls the smoothness of the transport plan. This formulation can be computed efficiently via Sinkhorn iterations and admits dual potentials whose derivatives with respect to the probability weights provide interpretable data values [9].

Let $u^* \in \mathbb{R}^n$ and $w^* \in \mathbb{R}^m$ denote the optimal source-side and target-side dual potentials, respectively. By sensitivity analysis of the OT objective, the derivative of the regularized OT value with respect to the source weights is given by the source-side dual potential:

$$\nabla_a W_{C,\varepsilon_{\mathrm{OT}}}(\mu, \nu) = u^*. \tag{A.5}$$

Thus, u_i^* measures how the OT discrepancy changes when the probability mass of source sample z_i is perturbed.

However, OT dual potentials are unique only up to an additive constant because one of the marginal constraints in the transport polytope is redundant. Following the calibrated OT valuation convention used in LAVA [16], we remove this degree of freedom by using a zero-sum calibrated score. For source sample z_i, the calibrated OT sensitivity score is defined as

$$v_i = u_i^* - \frac{1}{n-1} \sum_{\substack{k=1 \\ k \neq i}}^{n} u_k^*. \tag{A.6}$$

This calibrated score estimates how much source sample z_i contributes to the OT discrepancy between the source and target distributions. A large positive score indicates that increasing the mass of z_i would increase the OT discrepancy, suggesting that the sample is poorly aligned with the target distribution. Conversely, a smaller score indicates better alignment.

In data valuation, these calibrated OT sensitivity scores can be used to identify samples that increase the mismatch between training and validation distributions. Removing or downweighting samples with large positive scores can reduce the OT discrepancy and improve alignment with the validation distribution, as demonstrated in LAVA [16]. CAVA keeps this transport problem and replaces two of its ingredients for complementary-label learning, the label cost and the score (Section 3.3); the

calibrated score v_i is the score of the TW-LAVA baseline (Eq. (3.5)). Specifically, the generic source sample z_i becomes a complementary-labeled training sample $(\mathbf{x}_i, \bar{y}_i)$, the target sample z'_j becomes an ordinary-labeled validation sample $(\mathbf{x}^v_j, y^v_j)$, and the generic cost C is the ground cost of Eq. (3.1).

Despite its effectiveness, Sinkhorn-based OT still requires storing and operating on the pairwise cost matrix, leading to $\mathcal{O}(nm)$ memory complexity, or $\mathcal{O}(N^2)$ when both empirical distributions have size N. Recent works have addressed this limitation through minibatch or hierarchical OT approximations [17, 44]. CAVA instead uses a task-wise decomposition, solving smaller OT problems between each complementary-label task and the clean validation set, which improves computational feasibility while preserving the distributional interpretation needed for data valuation.

Compared to Shapley-value-based data valuation methods [10, 15], OT-based approaches avoid combinatorial subset enumeration and instead leverage distributional geometry to implicitly capture higher-order interactions. This makes OT particularly suitable for large-scale and continuous data settings, where valuation must account for domain shift, covariate imbalance, and sample diversity. Recent work has demonstrated the effectiveness of OT-based data valuation for dataset pruning, reweighting, curriculum learning, and active data acquisition [3, 16, 17, 36, 40], positioning OT as a scalable and principled foundation for data-centric machine learning.

A.2. Diagonal Complementary-Label Noise

Figure 8 shows how diagonal noise appears in the empirical transition matrix $P(\bar{Y} \mid Y)$: valid complementary labels fill the off-diagonal entries, whereas a label that names the true class places mass on the diagonal and asks the learner to suppress the correct class.

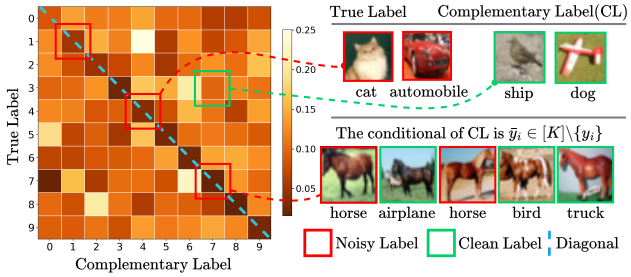

Figure 8. Complementary-label noise in CLL. When the complementary label coincides with the true label, i.e., $\bar{y}_i = y_i$, the empirical transition matrix contains diagonal mass, creating contradictory exclusion constraints.

A.3. Notation

Table 3 collects the main symbols.

Table 3. Notation.

Symbol	Meaning
$\mathcal{X}$, $\mathcal{Y} = \{1, \dots, M\}$	Input space and label space of M classes
$\bar{\mathcal{D}} = \{(\mathbf{x}_i, \bar{y}_i)\}_{i=1}^{N}$, y_i	Pool with given labels, and the unobserved true labels
$\mathcal{N} = \{i : \bar{y}_i = y_i\}$, $\eta = \mathbb{P}(\bar{y} = y)$	Diagonal-noise set and noise rate
$\mathcal{D}_v = \{(\mathbf{x}_j^v, y_j^v)\}_{j=1}^{N_v}$	Clean validation set with ordinary labels
I_t, T	Task t (a subset of the pool) and the number of tasks
ϕ, C_{feat}	Fixed feature extractor and feature cost (Eq. (3.2))
λ_y, $\varepsilon = \varepsilon_{\text{OT}}$	Label weight and entropic regularization
W^{grp}	Noisy-group label cost of TW-LAVA (Eq. (3.3))
W^{cl}, W_D^{cl}	Complementary label cost $(0/1)$ and its class-distance variant
$\alpha(\eta) = (M\eta - 1)/(M - 1)$	Lean of a noisy label group; effective label weight $\lambda_y \alpha(\eta)$ (Theorem 1)
$u^{(t),*}$, $w^{(t),*}$	Optimal pool-side and validation-side potentials of task t
$v_i^{(t)}$	Calibrated score of TW-LAVA (Eq. (3.5))
$f_i(c)$, $p_i(k)$, s_i	Potential under label c, transport posterior, label margin (Eqs. (3.9)–(3.10))
$F = \{i : s_i \le 0\}$, $h_i = \arg\max_k p_i(k)$	Adaptive-cut flagged set and posterior classifier
p	Pruning ratio of the fixed cut
I^{rec}, I^{keep}, κ	Recycled and retained sets, recycling weight (Eq. (3.14))
π_{det}, B, Δ_{gen}	Precision of the recycled set, loss bound, generalization term (Theorem 4)

B. Theory

This section proves the results of Section 3. Appendix B.1 states and proves the flip law, the direction of the label margin, and the precision of the adaptive cut; Appendix B.2 proves the recycling bound that turns this precision into a risk guarantee.

B.1. Analysis of the Complementary Label Cost and the Label Margin

This subsection gives the full statements and the proofs of the three results stated in short form with the method: the flip law (Theorem 1), the decomposition and direction of the margin (Theorem 2), and the detection-to-recycling bound (Theorem 3). Throughout, $d_{ij} = C_{\text{feat}}(\mathbf{x}_i, \mathbf{x}_j^v)$ denotes the feature cost of Eq. (3.2), the validation weights are $b_j = 1/N_v$, the feature weight is $\lambda_{\mathbf{x}} = 1$ (Section 3.1), and the entropic problem of Eq. (3.4) with ground cost C_{ij} has optimal potentials (u, w), written $u^{(t),*}$ and $w^{(t),*}$ in the main text, satisfying

$$u_i = -\varepsilon_{\text{OT}} \log \sum_j b_j \exp\big((w_j - C_{ij})/\varepsilon_{\text{OT}}\big), \tag{B.1}$$

so the counterfactual potentials of Eq. (3.9) satisfy $f_i(\bar{y}_i) = u_i$.

Assumption 1 (Uniform diagonal-noise model)**.** Classes are balanced. Given the true label y, the observed label is $\bar{y} = y$ with probability η and otherwise uniform on the other $M - 1$ labels, independently of $\mathbf{x}$. This is the synthetic protocol of Section 4. Write P_k for the feature distribution of class k, $\bar{P} = \frac{1}{M}\sum_k P_k$, and $\alpha(\eta) = (M\eta - 1)/(M - 1)$.

Lemma 1 (Group composition and shifts)**.** *(a) Under Assumption 1, the feature distribution of the pool samples with given label c is $\bar{\mathcal{D}}_c = (1 - \alpha)\bar{P} + \alpha P_c$ with $\alpha = \alpha(\eta)$. (b) If $C'_{ij} = C_{ij} + r_i + t_j$ for a row term r and a column term t, the entropic problems with costs C and C' have the same optimal plan, and $(u_i + r_i,\ w_j + t_j)$ are optimal potentials for C'. If r is constant, the calibrated score of Eq. (3.5) and the margin of Eq. (3.10) are unchanged.*

Proof. We compute the law of the given label, the posterior of the true label, and assemble the mixture; part (b) is a direct check of the optimality conditions.

Step 1. Marginal of the given label. By Assumption 1 and balanced classes,

$$\mathbb{P}(\bar{y} = c) = \frac{1}{M}\,\eta + \sum_{k \neq c} \frac{1}{M}\,\frac{1 - \eta}{M - 1} = \frac{1}{M}. \tag{B.2}$$

Step 2. Posterior of the true label. By the Bayes rule and Eq. (B.2),

$$\mathbb{P}(y = c \mid \bar{y} = c) = \eta, \qquad \mathbb{P}(y = k \mid \bar{y} = c) = \frac{1 - \eta}{M - 1} \quad (k \neq c). \tag{B.3}$$

Step 3. Mixture form. The label draw is independent of $\mathbf{x}$ given y, so by Eq. (B.3)

$$\bar{\mathcal{D}}_c = \eta\, P_c + \frac{1-\eta}{M-1}\sum_{k\neq c} P_k = \Big[\eta - \frac{1-\eta}{M-1}\Big] P_c + \frac{M(1-\eta)}{M-1}\,\bar{P}.$$

The coefficient of P_c is $(M\eta-1)/(M-1) = \alpha$ and the coefficient of $\bar{P}$ is $1-\alpha$, which proves (a).

Step 4. Same plan under shifts. The optimal plan for C can be rewritten with the shifted potentials:

$$\pi_{ij} = a_i b_j \exp\!\Big(\tfrac{u_i+w_j-C_{ij}}{\varepsilon_{\mathrm{OT}}}\Big) = a_i b_j \exp\!\Big(\tfrac{(u_i+r_i)+(w_j+t_j)-C'_{ij}}{\varepsilon_{\mathrm{OT}}}\Big).$$

It has marginals a and b, and the entropic optimum for C' is the unique plan of this Gibbs form with these marginals, so π is also optimal for C' with potentials $(u+r,\, w+t)$.

Step 5. Scores under shifts. The calibrated score changes by $r_i - \frac{1}{|I_t|-1}\sum_{k\neq i} r_k$, which vanishes when r is constant. For the margin, the shifted potential $w'_j = w_j + t_j$ absorbs the column term, so the counterfactual potentials satisfy $f'_i(c) = f_i(c) + r_i$ for every c, and the difference in Eq. (3.10) is unchanged. This proves (b). □

Full statement of Theorem 1. *Consider population groups under Assumption 1.*

(i) For any label divergence: at $\eta = 1/M$, $\bar{\mathcal{D}}_c = \bar{P}$ for every c, the ground cost no longer involves $\bar{y}_i$, and no score computed from the transport problem can distinguish diagonal noise from valid samples.

(ii) If the label divergence is a squared mean-embedding distance, then

$$W^{\mathrm{grp}}(c,k) = \alpha(\eta)\, W_D^{\mathrm{cl}}(c,k) + r_c + t_k$$

for row and column shifts r_c, t_k, where $W_D^{\mathrm{cl}}(c,k)$ is the same divergence computed between clean classes, the class-distance variant of the complementary label cost of Section 3.3. The TW-LAVA problem therefore has the same optimal plan as the problem under the complementary label cost with label weight $\lambda_y\,\alpha(\eta)$, and with equal class radii the same calibrated score and label margin.

Proof of Theorem 1. We prove part (i) from Lemma 1(a), obtain the shift identity of part (ii) by expanding two squared norms, and conclude with Lemma 1(b).

Step 1. Part (i): the given label leaves the cost. By Lemma 1(a) with $\alpha = 0$, $\bar{\mathcal{D}}_c = \bar{P}$ for every c. Hence

$$W^{\mathrm{grp}}(c,k) = D(\bar{P}, P_k), \qquad C_{ij} = d_{ij} + \lambda_y\, D(\bar{P}, P_{y_j^v}),$$

and $\bar{y}_i$ does not appear in the ground cost. The plan, the potentials, and every score computed from the transport problem are therefore functions of the features and of the validation set only. Under Assumption 1 the indicators $\mathbf{1}[\bar{y}_i = y_i]$ are i.i.d. Bernoulli(η) and independent of the features given the true labels, so every score has the same distribution on diagonal-noise and on valid samples.

Step 2. Mean embedding of the group. The mean embedding is linear in the distribution. Applying it to Lemma 1(a), with $u_k = \mu_{P_k} - \mu_{\bar{P}}$,

$$\mu_{\bar{\mathcal{D}}_c} - \mu_{\bar{P}} = \alpha\,(\mu_{P_c} - \mu_{\bar{P}}) = \alpha\, u_c. \tag{B.4}$$

Step 3. Expanding the squares. Writing $\mu_{\bar{\mathcal{D}}_c} - \mu_{P_k} = (\mu_{\bar{\mathcal{D}}_c} - \mu_{\bar{P}}) - (\mu_{P_k} - \mu_{\bar{P}})$ and using Eq. (B.4),

$$W^{\mathrm{grp}}(c,k) = \|\alpha u_c - u_k\|^2 = \alpha^2\|u_c\|^2 - 2\alpha\langle u_c, u_k\rangle + \|u_k\|^2, \tag{B.5}$$

$$\alpha\, W_D^{\mathrm{cl}}(c,k) = \alpha\|u_c - u_k\|^2 = \alpha\|u_c\|^2 - 2\alpha\langle u_c, u_k\rangle + \alpha\|u_k\|^2. \tag{B.6}$$

Subtracting Eq. (B.6) from Eq. (B.5),

$$W^{\mathrm{grp}}(c,k) - \alpha\, W_D^{\mathrm{cl}}(c,k) = -\alpha(1-\alpha)\|u_c\|^2 + (1-\alpha)\|u_k\|^2,$$

a term depending on the row only plus a term depending on the column only.

Step 4. Plans and scores. In the ground cost the two extra terms become $r_i = -\lambda_y\alpha(1-\alpha)\|u_{\bar{y}_i}\|^2$ and $t_j = \lambda_y(1-\alpha)\|u_{y_j^v}\|^2$. By Lemma 1(b) the TW-LAVA problem has the same optimal plan as the cost $d_{ij} + \lambda_y\alpha\, W_D^{\mathrm{cl}}(\bar{y}_i, y_j^v)$, which is the problem under the complementary label cost with label weight $\lambda_y\alpha(\eta)$. If $\|u_k\| = r$ for all k (equal class radii), then r_i is constant and the second part of Lemma 1(b) gives equality of the calibrated scores and of the margins. □

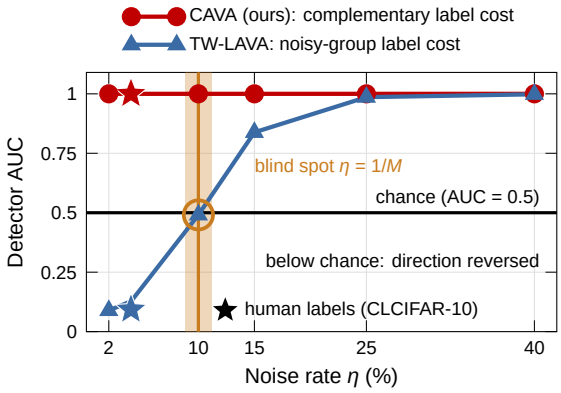

Figure 9. **The direction flip, measured.** Detector AUC against η on CIFAR-10; both scores use the same features, $\varepsilon_{\mathrm{OT}} = 0.01$, $\lambda_y = 1$, and the same $N_v = 2{,}000$ validation images. An AUC of 0.5 is chance. The noisy-group label cost is at chance exactly at the blind spot $\eta = 1/M$ (orange) predicted by Theorem 1, and below chance for smaller η, where it ranks the noise in the wrong direction. The complementary label cost stays at the top for every η. Stars: human-labeled CLCIFAR-10 (3.98% errors).

Figure 9 shows the flip of Theorem 1 as measured on CIFAR-10.

Full statement of Theorem 2. *For the $0/1$ complementary label cost, any validation potential w, any pool, any noise rate and any M, define the OT posterior and the function ψ by*

$$p_i(k) \propto \sum_{j:\, y_j^v = k} b_j \, e^{(w_j - C_{\mathrm{feat}}(\mathbf{x}_i, \mathbf{x}_j^v))/\varepsilon_{\mathrm{OT}}}, \qquad \psi(p) = -\varepsilon_{\mathrm{OT}} \log\big(\rho + (1-\rho)p\big), \quad \rho = e^{-\lambda_y/\varepsilon_{\mathrm{OT}}}.$$

Then:

(i) $f_i(c) = T_i + \psi(p_i(c))$, *where the typicality term T_i does not depend on c.*

(ii) $s_i = \psi(p_i(\bar{y}_i)) - \psi(\max_{c \neq \bar{y}_i} p_i(c))$ *and* $|s_i| \leq \lambda_y$.

(iii) $s_i \leq 0$ *if and only if* $\bar{y}_i \in \arg\max_k p_i(k)$.

(iv) If the label weight is replaced by any real λ', the label part of $f_i(\bar{y}_i)$ is strictly decreasing in $p_i(\bar{y}_i)$ for $\lambda' > 0$, constant at $\lambda' = 0$, and strictly increasing for $\lambda' < 0$.

Proof of Theorem 2. We split the sum that defines $f_i(c)$ by validation class, factor out the normalizer, and read off the four statements. Recall $p_i(k) = \frac{1}{Z_i} \sum_{j: y_j^v = k} b_j e^{(w_j - d_{ij})/\varepsilon_{\mathrm{OT}}}$ with $Z_i = \sum_j b_j e^{(w_j - d_{ij})/\varepsilon_{\mathrm{OT}}}$, $T_i = -\varepsilon_{\mathrm{OT}} \log Z_i$, and $\rho = e^{-\lambda_y/\varepsilon_{\mathrm{OT}}}$.

Step 1. Splitting the sum. With $W^{\mathrm{cl}}(c,k) = \mathbf{1}[c \neq k]$, the validation samples of class c pay no label cost and all others acquire the factor ρ:

$$\sum_j b_j \, e^{(w_j - d_{ij} - \lambda_y \mathbf{1}[c \neq y_j^v])/\varepsilon_{\mathrm{OT}}} = \sum_{j:\, y_j^v = c} b_j \, e^{(w_j - d_{ij})/\varepsilon_{\mathrm{OT}}} + \rho \sum_{j:\, y_j^v \neq c} b_j \, e^{(w_j - d_{ij})/\varepsilon_{\mathrm{OT}}}. \tag{B.7}$$

Step 2. Factoring the normalizer. The first sum in Eq. (B.7) equals $Z_i \, p_i(c)$ and the second equals $Z_i \, (1 - p_i(c))$, so

$$\sum_j b_j \, e^{(w_j - d_{ij} - \lambda_y \mathbf{1}[c \neq y_j^v])/\varepsilon_{\mathrm{OT}}} = Z_i \big[\rho + (1-\rho)\, p_i(c)\big]. \tag{B.8}$$

Taking $-\varepsilon_{\mathrm{OT}}$ log of Eq. (B.8) gives $f_i(c) = T_i + \psi(p_i(c))$, which is (i).

Step 3. The margin. Since ψ is strictly decreasing, the cheapest alternative label is the one with the largest posterior:

$$\min_{c \neq \bar{y}_i} f_i(c) = T_i + \psi\Big(\max_{c \neq \bar{y}_i} p_i(c)\Big).$$

Subtracting this from $f_i(\bar{y}_i) = T_i + \psi(p_i(\bar{y}_i))$ cancels T_i and yields the closed form of (ii). Because $0 \leq \psi \leq \lambda_y$ on $[0,1]$, $|s_i| \leq \lambda_y$.

Step 4. Direction. By monotonicity of ψ,

$$s_i \le 0 \iff \psi\big(p_i(\bar{y}_i)\big) \le \psi\Big(\max_{c\neq\bar{y}_i} p_i(c)\Big) \iff p_i(\bar{y}_i) \ge \max_{c\neq\bar{y}_i} p_i(c),$$

which is (iii). No property of the pool, of η, or of M was used.

Step 5. Sign of the weight. The computation leading to Eq. (B.8) does not use the sign of the weight. With $\rho' = e^{-\lambda'/\varepsilon_{\mathrm{OT}}}$,

$$f_i(\bar{y}_i) = T_i - \varepsilon_{\mathrm{OT}} \log\big(\rho' + (1-\rho')\, p_i(\bar{y}_i)\big).$$

The map $p \mapsto \rho' + (1-\rho')p$ is increasing when $\rho' < 1$, constant when $\rho' = 1$, and decreasing when $\rho' > 1$, and $\rho' < 1$ exactly when $\lambda' > 0$. Composing with $-\varepsilon_{\mathrm{OT}} \log$ reverses the monotonicity, which gives (iv). □

Full statement of Theorem 3. *Let $h_i = \arg\max_k p_i(k)$ be the classifier induced by the OT posterior of Theorem 2 (assumed unique). Let $\mathcal{N} = \{i : \bar{y}_i = y_i\}$ be the diagonal-noise set of the pool of size N, $\hat{\eta} = |\mathcal{N}|/N$, $F = \{i : s_i \le 0\}$, $e_{\mathcal{N}}$ the error rate of h on $\mathcal{N}$, and q_V the fraction of valid samples with $h_i = \bar{y}_i$. With the $0/1$ complementary label cost:*

(i) The recall and the precision of F are

$$\text{recall} = 1 - e_{\mathcal{N}}, \qquad \text{precision} = \frac{\hat{\eta}\,(1-e_{\mathcal{N}})}{\hat{\eta}\,(1-e_{\mathcal{N}}) + (1-\hat{\eta})\, q_V}.$$

(ii) Under the uniform noise model with rate η, if h has overall error e and is conditionally independent of the given label, then $q_V = e/(M-1)$ and the flagged-set purity

$$\pi = \frac{\eta\,(1-e)}{\eta\,(1-e) + (1-\eta)\, e/(M-1)}$$

exceeds the base rate η whenever $e \le 1 - 1/M$.

(iii) Training the CE branch on F satisfies, with probability at least $1-\delta$,

$$R_{\mathrm{true}}(\widehat{g}) - R_{\mathrm{true}}(g^\star) \le 2B\,(1 - \text{precision}) + \Delta_{\mathrm{gen}}(|F|, \delta),$$

by inserting the precision of (i) into Theorem 4.

Proof of Theorem 3. We express the flagging rule through the posterior classifier h, count the flagged samples on the noise set and on the valid set, specialize to uniform complementary labels, and insert the precision into Theorem 4. Let V denote the complement of $\mathcal{N}$.

Step 1. Flagging and the classifier. By Theorem 2(iii) and uniqueness of the maximizer,

$$i \in F \iff h_i = \bar{y}_i. \tag{B.9}$$

Step 2. Count on the noise set. For $i \in \mathcal{N}$ we have $\bar{y}_i = y_i$, so by Eq. (B.9) the sample is flagged exactly when $h_i = y_i$. Hence $|F \cap \mathcal{N}| = |\mathcal{N}|\,(1 - e_{\mathcal{N}})$, and the recall is $1 - e_{\mathcal{N}}$.

Step 3. Count on the valid set. For $i \in V$, Eq. (B.9) reads $h_i = \bar{y}_i$, so $|F \cap V| = |V|\, q_V$. Since $\bar{y}_i \neq y_i$ on V, such a sample is an error of h.

Step 4. Precision. Dividing the two counts by N,

$$\text{precision} = \frac{|F \cap \mathcal{N}|}{|F \cap \mathcal{N}| + |F \cap V|} = \frac{\hat{\eta}(1-e_{\mathcal{N}})}{\hat{\eta}(1-e_{\mathcal{N}}) + (1-\hat{\eta})\, q_V},$$

which completes (i).

Step 5. Uniform complementary labels. Given $(\mathbf{x}_i, y_i)$ and $i \in V$, the label $\bar{y}_i$ is uniform on the $M-1$ labels other than y_i and independent of h_i, so

$$\mathbb{P}\big(h_i = \bar{y}_i \mid i \in V,\, \mathbf{x}_i,\, y_i\big) = \frac{\mathbf{1}[h_i \neq y_i]}{M-1}.$$

Averaging over $(\mathbf{x}_i, y_i)$ gives $q_V = e/(M-1)$. Given $i \in \mathcal{N}$, the Bernoulli draw is independent of h_i, so $\mathbb{P}(i \in F \mid i \in \mathcal{N}) = \mathbb{P}(h_i = y_i) = 1 - e$. The Bayes rule with prior η and these two conditional probabilities gives the purity π of (ii).

Step 6. Enrichment condition. $\pi \ge \eta$ holds if and only if $1 - e \ge \eta(1-e) + (1-\eta)\, e/(M-1)$, that is, $(1-\eta)(1-e) \ge (1-\eta)\, e/(M-1)$. Dividing by $(1-\eta) > 0$ and rearranging gives $e \le 1 - 1/M$, which completes (ii).

Step 7. Recycling. Theorem 4 bounds the excess risk of the empirical minimizer on any recycled set of size m and precision π_{det} by $2B(1-\pi_{\mathrm{det}}) + \Delta_{\mathrm{gen}}(m, \delta)$. Taking the set to be F and inserting the precision of (i) gives (iii). □

Remark 1 (Scope)**.** Theorem 1(ii) covers mean-embedding divergences (squared distance between class means, or squared MMD). The pipeline uses a Sinkhorn divergence for the class-distance variant W_D^{cl}, for which we observe the same sign change with a measured proportionality factor of about $0.32\,\alpha(\eta)$; we state the exact law for mean embeddings and report the Sinkhorn case numerically. Theorems 2 and 3 are exact for the default $0/1$ complementary label cost, which performed within 0.01 AUC of the class-distance variant in every test.

B.2. Recycling Risk Bound

This subsection proves the recycling bound used by Theorem 3(iii). It shows that the benefit of recycling is governed by a single quantity, the *precision* of the flagged set.

Diagonal-noise model and detector. Under the diagonal-noise setting, a training sample is either *clean* ($\bar{y}_i \neq y_i$, a valid exclusion) or *noisy* ($\bar{y}_i = y_i$, a violated exclusion). CAVA flags a recycled set $\widehat{\mathcal{N}} = I^{\mathrm{rec}}$ (Section 3.5) and trains it with cross-entropy treating $\bar{y}$ as an ordinary label. Write $m = |\widehat{\mathcal{N}}|$ and define the detector *precision*

$$\pi_{\mathrm{det}} \;:=\; \frac{|\{i \in \widehat{\mathcal{N}} : \bar{y}_i = y_i\}|}{m}. \tag{B.10}$$

For a correctly flagged (noisy) sample $\bar{y}_i = y_i$, so the CE term is supervised by the *true* label; for a falsely flagged (clean) sample $\bar{y}_i \neq y_i$, the CE term is supervised by a wrong label. Thus the recycled branch is exactly an ordinary-label dataset corrupted at rate $(1 - \pi_{\mathrm{det}})$.

Assumption 2 (Bounded loss and capacity)**.** The per-example cross-entropy is bounded, $\mathcal{L}_{\mathrm{CE}}(y, g(\mathbf{x})) \in [0, B]$ (e.g. via logit clipping, temperature scaling, or label smoothing), and the hypothesis class $\mathcal{G}$ has Rademacher complexity $\mathfrak{R}_m(\mathcal{G})$ on m samples.

Theorem 4 (Recycling excess risk)**.** *Let* $\widehat{g} \in \arg\min_{g\in\mathcal{G}} \widehat{R}_{\mathrm{rec}}(g)$ *minimize the* observed *recycled risk* $\widehat{R}_{\mathrm{rec}}(g) = \frac{1}{m}\sum_{i\in\widehat{\mathcal{N}}} \mathcal{L}_{\mathrm{CE}}\big(\bar{y}_i, g(\mathbf{x}_i)\big)$ *over the recycled set* $\widehat{\mathcal{N}} = I^{\mathrm{rec}}$ *of size* m *and precision* π_{det}*, and let* $g^\star \in \arg\min_{g\in\mathcal{G}} R_{\mathrm{true}}(g)$ *minimize the true-label risk* $R_{\mathrm{true}}(g) = \mathbb{E}_{(\mathbf{x},y)}[\mathcal{L}_{\mathrm{CE}}(y, g(\mathbf{x}))]$*. Under Assumption 2, with probability at least* $1 - \delta$*,*

$$R_{\mathrm{true}}(\widehat{g}) - R_{\mathrm{true}}(g^\star) \leq \underbrace{2(1 - \pi_{\mathrm{det}})\, B}_{\text{noise bias (irreducible)}} + \underbrace{4B\,\mathfrak{R}_m(\mathcal{G}) + 2B\sqrt{\tfrac{\ln(2/\delta)}{2m}}}_{\Delta_{\mathrm{gen}}(m,\delta)\,\to\,0}. \tag{B.11}$$

Proof of Theorem 4. All statements below hold under Assumption 2, with probability at least $1 - \delta$.

Step 1. The observed risk tracks the clean risk. Let $\widehat{R}_{\mathrm{clean}}(g) = \frac{1}{m}\sum_{i\in\widehat{\mathcal{N}}} \mathcal{L}_{\mathrm{CE}}(y_i, g(\mathbf{x}_i))$ denote the empirical risk under the (unobserved) true labels. For a correctly flagged sample, $\bar{y}_i = y_i$, so the corresponding summands of $\widehat{R}_{\mathrm{rec}}$ and $\widehat{R}_{\mathrm{clean}}$ agree. For each of the $(1 - \pi_{\mathrm{det}})m$ falsely flagged samples, the two summands differ by at most B, since both losses lie in $[0, B]$. Averaging over the m samples gives, for all g,

$$\left|\widehat{R}_{\mathrm{rec}}(g) - \widehat{R}_{\mathrm{clean}}(g)\right| \leq (1 - \pi_{\mathrm{det}})B. \tag{B.12}$$

Step 2. Uniform convergence on the clean risk. Standard uniform convergence [30] gives, with probability at least $1 - \delta$,

$$\sup_{g\in\mathcal{G}} \Big(R_{\mathrm{true}}(g) - \widehat{R}_{\mathrm{clean}}(g)\Big) \leq \Delta_{\mathrm{gen}}(m, \delta), \tag{B.13}$$

together with the analogous lower-tail bound.

Step 3. Chaining the two estimates. Combining Steps 1 and 2,

$$\begin{aligned} R_{\mathrm{true}}(\widehat{g}) &\leq \widehat{R}_{\mathrm{clean}}(\widehat{g}) + \Delta_{\mathrm{gen}} \leq \widehat{R}_{\mathrm{rec}}(\widehat{g}) + (1 - \pi_{\mathrm{det}})B + \Delta_{\mathrm{gen}} \leq \widehat{R}_{\mathrm{rec}}(g^\star) + (1 - \pi_{\mathrm{det}})B + \Delta_{\mathrm{gen}} \\ &\leq \widehat{R}_{\mathrm{clean}}(g^\star) + 2(1 - \pi_{\mathrm{det}})B + \Delta_{\mathrm{gen}} \leq R_{\mathrm{true}}(g^\star) + 2(1 - \pi_{\mathrm{det}})B + 2\Delta_{\mathrm{gen}}, \end{aligned}$$

where the first and fifth inequalities use Step 2, the second and fourth use Eq. (B.12), and the third uses that $\widehat{g}$ minimizes $\widehat{R}_{\mathrm{rec}}$. Rearranging completes the proof. □

Theorem 4 isolates the recycled branch by analyzing the minimizer of $\widehat{R}_{\mathrm{rec}}$ alone. The next corollary extends the same argument to the mixed objective actually optimized in Eq. (3.14).

Corollary 1 (Recycling within the mixed objective)**.** Extend Assumption 2 by $\mathcal{L}_{\mathrm{comp}}(\bar{y}, g(\mathbf{x})) \in [0, B]$. Let $\widehat{R}_{\mathrm{keep}}$ denote the empirical complementary risk on the $n_k = |I^{\mathrm{keep}}|$ retained samples, let $\widehat{g} \in \arg\min_{g\in\mathcal{G}} \ \kappa\,\widehat{R}_{\mathrm{rec}}(g) + (1-\kappa)\,\widehat{R}_{\mathrm{keep}}(g)$ minimize the empirical mixed objective, and define the population mixed risk $R_{\mathrm{mix}}(g) = \kappa\,R_{\mathrm{true}}(g) + (1-\kappa)\,R_{\mathrm{comp}}(g)$ with $R_{\mathrm{comp}}(g) = \mathbb{E}_{\mathbf{x}_t\sim\rho_t}[\mathcal{L}_{\mathrm{comp}}(\bar{y}_t, g(\mathbf{x}_t))]$, together with $\Delta_{\mathrm{mix}} = \kappa\,\Delta_{\mathrm{gen}}(m, \delta/2) + (1-\kappa)\,\Delta_{\mathrm{gen}}^{\mathrm{comp}}(n_k, \delta/2)$, where $\Delta_{\mathrm{gen}}^{\mathrm{comp}}$ is the analogous uniform-convergence term for the complementary loss class. Then, with probability at least $1-\delta$, for every $g^\star \in \mathcal{G}$,

$$R_{\mathrm{mix}}(\widehat{g}) - R_{\mathrm{mix}}(g^\star) \ \le\ 2\kappa\,(1-\pi_{\mathrm{det}})\,B \ +\ 2\,\Delta_{\mathrm{mix}}. \tag{B.14}$$

Proof. **Step 1. Bias of the observed mixture.** For every g, Eq. (B.12) gives $\big|\widehat{R}_{\mathrm{mix}}(g) - \kappa\widehat{R}_{\mathrm{clean}}(g) - (1-\kappa)\widehat{R}_{\mathrm{keep}}(g)\big| = \kappa\big|\widehat{R}_{\mathrm{rec}}(g) - \widehat{R}_{\mathrm{clean}}(g)\big| \le \kappa(1-\pi_{\mathrm{det}})B$.

Step 2. Uniform convergence of the unbiased mixture. A union bound over the two branches gives, with probability at least $1-\delta$, $\sup_{g\in\mathcal{G}}\big|\kappa\widehat{R}_{\mathrm{clean}}(g) + (1-\kappa)\widehat{R}_{\mathrm{keep}}(g) - R_{\mathrm{mix}}(g)\big| \le \Delta_{\mathrm{mix}}$.

Step 3. Chaining. Combining Steps 1 and 2 exactly as in the proof of Theorem 4,

$$R_{\mathrm{mix}}(\widehat{g}) \le \widehat{R}_{\mathrm{mix}}(\widehat{g}) + \kappa(1-\pi_{\mathrm{det}})B + \Delta_{\mathrm{mix}} \le \widehat{R}_{\mathrm{mix}}(g^\star) + \kappa(1-\pi_{\mathrm{det}})B + \Delta_{\mathrm{mix}} \le R_{\mathrm{mix}}(g^\star) + 2\kappa(1-\pi_{\mathrm{det}})B + 2\Delta_{\mathrm{mix}},$$

where the middle inequality uses that $\widehat{g}$ minimizes the empirical mixed objective. □

Setting $\kappa = 1$ recovers Theorem 4. In the mixed objective the noise bias enters only through the recycled branch, scaled by the weight κ.

Remark 2 (Data-dependent selection)**.** The recycled set I^{rec} is obtained by thresholding scores computed on the same pool, so its elements are not an i.i.d. sample drawn independently of the selection, and Theorem 4 and Corollary 1 should be read conditionally on the realized selection, with π_{det} its realized precision; this is exactly the quantity we measure post hoc and report (Appendix D.7). An unconditional statement follows from a standard sample-splitting variant: compute the scores and the cuts on one half of the pool and apply recycling only to the other half. Conditionally on the first half, the selection region is fixed, the selected samples of the second half form an i.i.d. draw from it, and both results apply verbatim with m the split size. We verified this device empirically with the TW-LAVA configuration (noisy-group score, fixed cut) on synthetic CIFAR-10 ($\eta = 25\%$, $N_v = 2{,}000$, seed 1126): scoring each half of the pool with the entropic c-transform of the dual potentials fitted on the other half gives detector AUC 0.962 and recycled-branch precision 83.5%, compared with 0.971 and 86.8% for full-pool scoring, and end-to-end training on the split-fitted selection reaches 85.16% against 86.33%, so the unconditional guarantee costs about one accuracy point in practice. Because the per-task rules are one-dimensional thresholds, letting a fitted threshold be chosen on the evaluation half adds only the complexity of threshold classes, a union of families of VC dimension one that inflates Δ_{gen} by a logarithmic factor; the adaptive cut of CAVA fits no threshold, so it adds nothing. Both devices change constants, not the structure of the bounds.

The bound separates two effects: the estimation term Δ_{gen} shrinks as the recycled set grows, whereas false positives leave a bias floor $2(1-\pi_{\mathrm{det}})B$ that more data cannot remove. Recycling therefore helps only when the flagged set is pure, which is why CAVA recycles the adaptive-cut set, whose precision Theorem 3 gives in closed form, and why the selection-only objective is the default when that precision cannot be checked.

C. Experimental Setup

This section gives the datasets, the training and valuation settings, the metrics used to evaluate selections, and the baselines.

C.1. Datasets

Synthetic benchmarks. MNIST [20], CIFAR-10, and CIFAR-20 [18] (CIFAR-100 grouped into its 20 superclasses) receive synthetic diagonal noise: each training label equals the true class with probability η and is otherwise drawn uniformly from the other $M-1$ classes (Assumption 1); $\eta = 25\%$ unless stated otherwise.

Human-annotated benchmarks. CLCIFAR-10 and CLCIFAR-20 [35] pair the CIFAR training images with crowdsourced complementary labels; we use the first of the three labels per image, of which 3.98% ($1{,}990$ images) and 2.32% ($1{,}161$) name the true class. CLMIN-10 and CLMIN-20 from the same collection contain $5{,}000$ and $10{,}000$ ImageNet training images of size 64×64, resized to 32×32, with 4.80% and 3.28% such annotation errors. All four have error rates below $1/M$, the regime in which the label cost of TW-LAVA reverses (Theorem 1).

Further domains. Three benchmarks outside CIFAR and ImageNet receive the same synthetic diagonal noise ($\eta \in \{5, 15, 25\}\%$). FOOD15 uses the first 15 of 20 dishes we take from Food-101 [1] (hamburger, pizza, sushi, ramen, steak, ice cream, french fries, caesar salad, fried rice, chocolate cake, donuts, dumplings, pancakes, spaghetti bolognese, tacos) with the official split (750 training and 250 test images per class); the training images of Food-101 were left uncleaned by its authors and may carry a small amount of label noise of their own, while the test images were cleaned. RESISC45-20 uses 20 of the 45 aerial scene classes of NWPU-RESISC45 [2] (airport, beach, bridge, church, desert, forest, harbor, island, parking lot, stadium, circular farmland, cloud, freeway, golf course, lake, mountain, railway, roundabout, ship, storage tank); the dataset has no official split, so the 700 images per class are split into 500 training and 200 test images with seed 1126. Places365-10 uses 10 scene classes of Places365-Standard [41] (airfield, beach, bedroom, bridge, forest path, kitchen, mountain, skyscraper, football stadium, outdoor swimming pool), with 4,000 training and 1,000 test images per class drawn from the official training set with seed 1126, since the official validation split holds only 100 images per class. All images are center-cropped and resized to 32×32; at 64×64 and 128×128 the accuracy did not improve in a pilot. Each dataset has its own feature extractor of the same architecture, trained on the clean labels of its training split and taken after the final epoch. The validation set holds 50 test images per class (200 for Places365-10), and accuracy is measured on the remaining test images.

Scale. CIFAR-10-1M enlarges the CIFAR-10 training set to one million samples by label-preserving augmentation (Appendix D.9).

Scope. Benchmarks with at most 20 classes are standard in CLL, and current CLL methods are workable only at this scale [4, 6, 13, 29, 35], because a complementary label excludes a single class and the supervision weakens as the label space grows; extending CLL to larger label spaces is beyond the scope of this work.

C.2. Training and Valuation Settings

For all image classification experiments, we use ResNet-18 as the backbone network. Models are trained for 200 epochs with batch size 512 using stochastic gradient descent with momentum 0.9 and weight decay 5×10^{-4}. The initial learning rate is set to 0.1 and decayed to 0.01 and 0.001 at epochs 100 and 150, respectively. We set random seeds to control data loading and training randomness. Table 1 and every table or figure whose caption says so report the mean $\pm$ standard deviation over three independent runs with seeds 1126, 42, and 2023, the same seeds for all compared methods; the other results are single runs with seed 1126. Unless a caption says otherwise, every method in a table uses the same feature extractor, the same 2,000 validation images, the same tasks, and AutoAugment, and accuracy is measured after the final epoch on the held-out test images. The recycling weight is $\kappa = 0.5$ in all experiments. For computationally expensive scalability experiments such as CIFAR-10-1M, we report single-run results and explicitly mark them as such.

OT valuation settings. Following LAVA and SAVA [16, 17], the feature extractor ϕ is a fixed pretrained network shared by every valuation method and baseline, so no method gains an advantage from it. We use the embedder released with LAVA for CIFAR-10 (which CLCIFAR-10 reuses) and networks trained the same way for the other datasets: a supervised PreAct-ResNet-18 trained with the clean ordinary labels of the training split, which is the same image set as the complementary-label pool (the LAVA network reaches 99.87% on the CIFAR-10 training split and 93.07% on its test split; LAVA describes its embedder as trained on the validation set, which this released network is not). CLMIN and the further domains have their own such networks, and CIFAR-10-1M an augmentation-matched one. We use its class logits (the penultimate 512-dimensional features for CIFAR-20 and CLCIFAR-20). In the complementary-label setting such an embedder is an upper bound on feature quality, since it has seen the true label of every pool image, and it is the only component of the pipeline that has: the selection, the thresholds and the training never use a true label. The detection quality of the label margin is bounded by how well ϕ separates the classes (Theorem 3(i)), so weaker, label-free features lower the absolute numbers of every method without changing which method is better; Appendix C.4 reports all detectors on the same features. Representations are L2-normalized and rescaled so that the median feature cost between pool and validation samples equals one. CAVA uses the ground cost of Eq. (3.1) with $\lambda_{\mathbf{x}} = \lambda_y = 1$, the $0/1$ complementary label cost, entropic regularization $\varepsilon_{\mathrm{OT}} = 0.01$, uniform marginals, $T = 5$ random tasks of equal size, and a fixed class-stratified validation set of $N_v = 2{,}000$ test images (drawn once with seed 1126 and shared by all seeds and methods), used only to score the training samples and never for training. On CLMIN-10 and CLMIN-20 the test split holds only 50 images per class, so 20 per class form the validation set (200 and 400 images) and the other 30 per class the evaluation set. TW-LAVA uses the same feature extractor, $\lambda_{\mathbf{x}} = \lambda_y = 1$ and the noisy-group label cost of Eq. (3.3); it is computed at the CAVA settings ($\varepsilon_{\mathrm{OT}} = 0.01$, $N_v = 2{,}000$). The small-noise experiments (Figure 6 and Table 19) are the one exception: there TW-LAVA keeps the OT settings of its original implementation ($\varepsilon_{\mathrm{OT}} = 0.1$, label costs given by the OT distance between label groups), while still using the same 2,000 validation images as every other method. The direct score comparisons of Figures 2 and 9 recompute the noisy-group score at the CAVA settings (same normalized features,

$\varepsilon_{\mathrm{OT}} = 0.01$, $N_v = 2{,}000$), so that only the label cost and the score differ.

Transition matrix for FWD and CPE. Forward correction and CPE take a class transition matrix Q. As in their original papers and in the CLImage benchmark [22, 35, 39], Q is the empirical transition matrix of the training pool (the known-Q protocol); it is the same for every method in a column, so it affects the objective and not the selection. The empirical Q of a noisy pool contains the diagonal noise rates, so these two objectives receive information that SCL-NL and SCL-EXP do not; our main results use SCL-NL, which needs no transition matrix, and CAVA improves it over TW-LAVA by 3.9 to 43.1 points on the CLCIFAR datasets, against 7.3 to 46.6 points for FWD (Table 1, Figure 4).

C.3. Branch-Quality Metrics

Several tables evaluate the two branches of Eq. (3.14) with the true labels, which are available for the synthetic benchmarks and for the annotated CLCIFAR datasets. In the actual CLL training setup, the true label y_i is not observed, and CAVA only uses $(\mathbf{x}_i, \bar{y}_i)$ together with the clean validation set. However, in the synthetic benchmark, the original labels are known to the evaluator because diagonal complementary-label noise is artificially generated. We therefore use y_i only after CAVA has completed selection and recycling, and only to evaluate the quality of the recycled and retained subsets. The true labels are never used during training, OT valuation, threshold estimation, pruning, or recycling, apart from the fixed feature extractor described in Appendix C.2.

The two selected subsets affect different branches of Eq. (3.14). The recycled set I^{rec} is trained with the CE loss by treating the observed complementary label $\bar{y}_i$ as an ordinary label. If a recycled sample satisfies $\bar{y}_i = y_i$, then the CE branch receives correct ordinary-label supervision. If a recycled sample satisfies $\bar{y}_i \neq y_i$, then the CE branch receives an incorrect ordinary label. Similarly, the retained set I^{keep} is trained with SCL-NL as the complementary-label branch, so any diagonal-noise sample remaining in I^{keep} still imposes a contradictory exclusion constraint. Therefore, the quality of I^{rec} affects the CE branch, while the quality of I^{keep} affects the SCL-NL branch.

Let $I^{\mathrm{noise}} = \{i : \bar{y}_i = y_i\}$ be the true diagonal-noise set, which is used only for this post-hoc evaluation. I^{rec} is the recycled set sent to the CE branch (the flagged tail for the fixed cut, or the adaptive-cut set $\{i : s_i \leq 0\}$ of CAVA), and I^{keep} is the retained set sent to the complementary-label branch. Table 4 lists the metrics; the first three describe the CE branch as a detector of diagonal noise, and the last two give the share of wrong supervision in each branch.

Table 4. Branch-quality metrics. Higher is better for precision, recall, and F1; lower is better for the two contamination rates.

Metric	**Definition**	**What it measures**
CE Precision	$\lvert I^{\mathrm{rec}} \cap I^{\mathrm{noise}} \rvert \,/\, \lvert I^{\mathrm{rec}} \rvert$	recycled labels that are correct ordinary labels
CE Recall	$\lvert I^{\mathrm{rec}} \cap I^{\mathrm{noise}} \rvert \,/\, \lvert I^{\mathrm{noise}} \rvert$	diagonal noise recovered into the CE branch
F1	harmonic mean of the two	balance between precision and recall
CE Contam.	$1 -$ CE Precision	wrong ordinary labels in the CE branch
CLL Contam.	$\lvert I^{\mathrm{keep}} \cap I^{\mathrm{noise}} \rvert \,/\, \lvert I^{\mathrm{keep}} \rvert$	diagonal noise left in the complementary-label branch

C.4. Detection Baselines

Table 5 gives the values plotted in Figure 10, which shows the synthetic datasets at $\eta = 5\%$ and $\eta = 15\%$, and adds $\eta = 25\%$. At $\eta = 5\%$, below the crossover $1/M$ of CIFAR-10 and at the crossover of CIFAR-20, TW-LAVA is reversed and at chance, respectively; above both crossovers, Data-OOB and small-loss selection are reversed instead. All methods are evaluated on the same pool and five tasks (seed 1126), all valuation methods use the same feature extractor and the same 2,000 validation images, and each score is evaluated in its natural direction, with valid samples expected to score high. KNN-Shapley [15] uses the exact recursion with the k-nearest-neighbor accuracy ($k = 10$) on the validation set as utility and $\bar{y}_i$ as the training label. Data-OOB [19] bags 50 logistic regressions fit on $(\mathbf{x}_i, \bar{y}_i)$ and scores each sample by the out-of-bag probability of $\bar{y}_i$. Small-loss selection [11] needs no clean data; it warms up a ResNet-18 from scratch with SCL-NL on the full pool for 20 epochs and scores each sample by its negative loss averaged over the last 5 epochs. TW-LAVA is the calibrated score of Eq. (3.5) at the CAVA OT settings. Every score in this table is measured on the first task; over all five tasks, TW-LAVA gives 0.996, 0.945, 0.076, and 0.149 at $\eta = 25\%$ on CIFAR-10 and CIFAR-20 and on CLCIFAR-10 and CLCIFAR-20, the values used in the rest of the appendix.

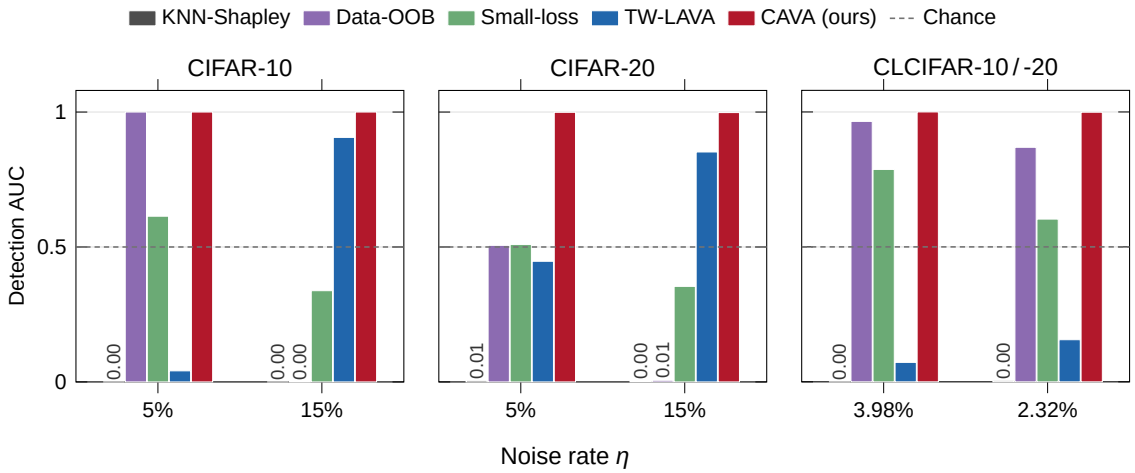

Figure 10. **Detection of diagonal noise.** Detection AUC of each score in its natural direction; the synthetic datasets are shown at $\eta = 5\%$ and 15%, around the crossover $1/M$ (10% for CIFAR-10, 5% for CIFAR-20). The gray dashed line marks chance (0.5): a bar below it ranks the noise above valid data, and bars with an AUC near 0 show their value.

Table 5. Detection AUC of diagonal noise in the natural direction of each score (valid samples high; 0.5 is chance, values near 0 rank the noise as the most valuable data). Same pool and tasks for every method (seed 1126).

	CIFAR-10			CIFAR-20			CLCIFAR-10	CLCIFAR-20
Noise rate η	25%	15%	5%	25%	15%	5%	3.98%	2.32%
KNN-Shapley	0.000	0.000	0.000	0.002	0.003	0.005	0.000	0.004
Data-OOB	0.000	0.001	1.000	0.002	0.006	0.506	0.965	0.869
Small-loss	0.217	0.338	0.614	0.313	0.354	0.509	0.787	0.603
TW-LAVA	0.995	0.906	0.041	0.940	0.852	0.447	0.072	0.156
CAVA (ours)	1.000	1.000	1.000	0.997	0.998	0.999	1.000	0.999

C.5. Target-Label-Free Representations

The supervised embedder of the main tables has seen the true label of every pool image (Appendix C.2), so it bounds the quality any representation can give the valuation. This subsection repeats the detection experiments with two extractors that have seen no label of the pool, under the protocol of Table 1: *CLIP*, the image encoder of CLIP ViT-B/32 [32] applied to the images upsampled to 224×224, and *validation only*, a PreAct-ResNet-18 trained from scratch on the $2{,}000$ validation images (200 or 400 on CLMIN) and nothing else, the protocol LAVA describes for its own embedder, with penultimate features and no model selection. Both are L2-normalized and median-rescaled as before, and all methods in a row share the representation, the validation set and the five tasks. Seeds 42 and 2023 change the AUC of CAVA by at most 0.005 and that of TW-LAVA by at most 0.011 on every dataset of Table 6, so the table reports seed 1126.

What the transport contributes. Theorem 2 holds for any validation potential, including $w = 0$, where p_i is a kernel classifier over the validation set with bandwidth $\varepsilon_{\mathrm{OT}}$. The column "$w = 0$" is therefore the exact ablation of the OT solve. Table 6 shows that it is within 0.006 AUC of CAVA on both representations and every dataset where we computed it: the potentials, which enforce that every validation image receives the same mass, do not improve detection. Two further settings test the one effect the marginal constraint can have. With $N_v \in \{50, 100, 200, 500\}$ on CIFAR-10 (CLIP), CAVA reaches AUC 0.978, 0.984, 0.987 and 0.990, and $w = 0$ is within 0.001 of it at $N_v \geq 100$ and higher at $N_v = 50$ (0.981). When the class marginals of the pool and of the validation set differ, the constraint forces mass onto the wrong classes: thinning five validation classes to 20% lowers CAVA to 0.920 while $w = 0$ stays at 0.985, and thinning five pool classes lowers CAVA to 0.983 against 0.994 for $w = 0$. The transport is thus the framework in which the complementary label cost and the margin are defined, and with mismatched class marginals the margin without the marginal constraint ($w = 0$) is preferable.

The flip law and the two components on label-free features. Table 7 repeats the sweep of Figure 9 on CLIP features for the four combinations of label cost and score. The noisy-group cost flips with either score: with the calibrated potential

Table 6. Detection of diagonal noise on two representations that have never seen a label of the pool (seed 1126, all five tasks). AUC treats valid samples as positive; "adaptive cut" gives the precision and recall of $\{s_i \le 0\}$, and "$w = 0$" is the CAVA margin without the validation potential. The supervised embedder of the main tables is shown for reference; "–": not run.

		CAVA (ours)			
Features	**Dataset**	AUC	adaptive cut P / R	$w = 0$	**TW-LAVA**
CLIP	MNIST	0.990	96.7 / 88.5	–	0.771
	CIFAR-10	0.994	96.6 / 88.1	0.992	0.773
	CIFAR-20	0.971	94.5 / 66.3	0.969	0.711
	CLCIFAR-10	0.996	87.1 / 86.4	0.995	0.758
	CLCIFAR-20	0.981	75.4 / 59.6	0.976	0.621
	CLMIN-10	0.968	73.0 / 70.8	0.965	0.540
	CLMIN-20	0.949	66.7 / 55.5	–	0.514
validation only	MNIST	0.999	99.4 / 98.1	–	0.985
	CIFAR-10	0.940	86.7 / 68.5	0.941	0.751
	CIFAR-20	0.853	81.6 / 37.8	0.855	0.624
	CLCIFAR-10	0.943	52.9 / 66.6	0.942	0.742
	CLCIFAR-20	0.861	28.7 / 32.6	0.857	0.595
	CLMIN-10	0.797	25.0 / 36.7	0.794	0.585
	CLMIN-20	0.754	20.4 / 22.6	–	0.499
supervised (oracle)	CIFAR-10	1.000	99.98 / 99.8	–	0.996
	CIFAR-20	0.997	98.8 / 87.2	–	0.945
	CLCIFAR-10	1.000	99.7 / 99.7	–	0.076
	CLCIFAR-20	0.999	96.0 / 87.9	–	0.149

(TW-LAVA) the AUC rises from 0.34 at $\eta = 2\%$ through chance at $\eta = 1/M$ to 0.77 at 25%, and with the margin from 0.04 to 0.98. The complementary label cost does not flip with either score; the calibrated potential under it stays at 0.97 to 0.98 and the margin at 0.994 to 0.995, the difference being the typicality term T_i that the margin cancels (Theorem 2(i)). The cost therefore removes the flip and the margin adds the last points; CIFAR-20 behaves the same way with a crossover at 5%.

Table 7. Detection AUC of the four combinations of label cost and score on CLIP features (seed 1126, $N_v = 2{,}000$) against the noise rate; the crossover is $\eta = 1/M$. The supervised embedder gives the same pattern (Figure 9).

		Noisy-group cost		Complementary cost	
Dataset	η	calibrated (TW-LAVA)	margin	calibrated	margin (CAVA)
CIFAR-10	2%	0.342	0.041	0.983	0.995
	5%	0.400	0.063	0.981	0.995
	10%	0.503	0.508	0.979	0.994
	15%	0.610	0.952	0.977	0.994
	25%	0.773	0.982	0.974	0.994
CIFAR-20	2.5%	0.457	0.250	0.946	0.974
	5%	0.489	0.488	0.940	0.972
	15%	0.613	0.926	0.939	0.974
	25%	0.711	0.946	0.933	0.971

Estimating the precision without true labels. Theorem 3(ii) needs the error e of the posterior classifier and the noise rate η. On CLIP features we estimate e by leave-one-out prediction of h on the validation set and η from the flagged fraction (Appendix B.1). Table 8 compares the estimate with the precision measured post hoc. On the synthetic benchmarks, which satisfy Assumption 1, the estimate is within 1.5 points. On the human-annotated benchmarks it underestimates the precision by 15 to 57 points, because their valid labels are not uniform and their errors concentrate on ambiguous images, so the uniform model overstates how often a valid sample is flagged; the error is on the safe side for gating recycling. The estimate is not available for the validation-only embedder, which fits the validation images and makes the leave-one-out error vanish.

Table 8. Estimating the precision of the adaptive cut without true labels (CLIP features, seed 1126): leave-one-out error $\hat{e}$ of the posterior classifier, estimated noise rate $\hat{\eta}$, and estimated precision $\hat{\pi}$ (Eq. (3.15)), next to the measured precision and recall.

Dataset	$\hat{e}$	$\hat{\eta}$	η	$\hat{\pi}$ (%)	measured (%)	recall (%)
MNIST	0.093	0.244	0.250	96.6	96.7	88.5
CIFAR-10	0.107	0.245	0.250	96.1	96.6	88.1
CIFAR-20	0.304	0.234	0.250	93.0	94.5	66.3
CLCIFAR-10	0.104	0.032	0.040	71.7	87.1	86.4
CLCIFAR-20	0.285	0.005	0.023	18.6	75.4	59.6
CLMIN-10	0.248	0.026	0.048	42.4	73.0	70.8
CLMIN-20	0.349	0.014	0.033	33.8	66.7	55.5

Training on label-free features. Table 9 reports Table 1 on CLIP features and on the validation-only features, with three seeds each. On both representations and with both objectives, CAVA is above TW-LAVA on all seven datasets; with recycling on the validation-only features it reaches 75.22% on CIFAR-10 and 49.21% on CIFAR-20, against 50.61% and 22.27% for TW-LAVA.

Table 9. Test accuracy (%) of Table 1 on CLIP features and on the validation-only features (mean ± std over seeds 1126, 42, 2023). Best per column, block and representation in **bold**.

			MNIST	CIFAR-10	CIFAR-20	CLCIFAR-10	CLCIFAR-20	CLMIN-10	CLMIN-20
(a) Selection only: SCL-NL on the retained pool									
CLIP	TW-LAVA	Fixed p	26.25±1.31	8.41±1.86	3.10±0.31	42.79±1.79	7.69±0.08	15.22±2.17	4.50±0.17
CLIP	CAVA (ours)	Fixed p	**90.59**±0.21	**74.11**±0.32	**12.05**±1.05	**79.06**±0.67	**10.64**±0.89	**27.11**±1.90	**10.66**±0.76
CLIP	CAVA (ours)	Adaptive cut	84.58±0.35	64.19±3.36	7.01±1.74	77.94±0.43	8.48±0.11	25.11±4.43	7.83±2.92
validation only	TW-LAVA	Fixed p	97.91±0.10	29.90±2.29	2.71±0.12	43.43±0.72	7.68±0.15	16.22±3.02	4.39±0.59
validation only	CAVA (ours)	Fixed p	**98.83**±0.06	**66.25**±0.31	**18.57**±1.40	73.33±0.08	**10.41**±1.77	21.56±2.52	**12.83**±0.67
validation only	CAVA (ours)	Adaptive cut	98.74±0.09	56.61±0.76	4.18±2.27	**73.90**±0.34	8.57±0.34	**22.78**±3.56	9.78±0.51
(b) Selection with recycling: CE+SCL-NL									
CLIP	TW-LAVA	Fixed p	56.42±2.59	46.68±0.58	28.85±0.96	28.18±2.10	5.24±0.29	16.11±1.65	6.45±1.11
CLIP	CAVA (ours)	Fixed p	**94.93**±0.13	**87.91**±0.55	**67.37**±0.24	80.59±0.23	**42.12**±0.77	**50.44**±1.26	**29.55**±0.86
CLIP	CAVA (ours)	Adaptive cut	94.31±0.03	87.24±0.28	62.67±0.34	**81.84**±0.29	41.66±0.90	49.89±0.70	29.39±0.26
validation only	TW-LAVA	Fixed p	97.78±0.47	50.61±1.32	22.27±0.56	34.47±1.22	5.85±0.60	19.11±1.17	4.78±1.77
validation only	CAVA (ours)	Fixed p	99.29±0.05	**75.22**±0.67	**49.21**±0.85	**72.53**±0.14	**35.07**±0.47	40.11±1.07	22.22±1.27
validation only	CAVA (ours)	Adaptive cut	**99.33**±0.05	73.50±0.31	41.81±0.48	72.46±0.13	34.06±0.32	**43.89**±0.70	**22.55**±0.95

Training without the transport solve. Table 10 trains on the selections of Table 6 with CLIP features. The accuracies follow the detection results: with recycling, CAVA without the OT solve ($w = 0$) is within about one point of CAVA on every dataset except CLCIFAR-20 (3.0 points), and both are far above TW-LAVA. Table 11 separates the two components of CAVA in training. With the noisy-group label cost, the margin helps on CIFAR-10, above the crossover, but not on CLCIFAR-10, below it (41.21% and 31.08% against 79.06% and 80.59% for CAVA). The complementary label cost alone removes the reversal on CLCIFAR-10 (71.42% and 73.26%), and the margin adds 7.1 to 15.3 points on top of it. With fewer validation images (Table 12), CAVA reaches 82.92% and 84.36% at $N_v = 50$ and 100, against 87.91% at $N_v = 2{,}000$.

D. Additional Results

D.1. Synthetic Benchmarks: Branch Composition and Complementary-Label Bias

This subsection analyzes the fixed cuts on MNIST, CIFAR-10, and CIFAR-20 with $\eta = 25\%$ (AutoAugment). For every cut, Tables 14, 13, and 15 report the composition of the two branches of Eq. (3.14), measured by CE Contam. and CLL Contam. (Table 4), and the test accuracy with and without recycling. CAVA rows follow the protocol of the main paper; TW-LAVA rows use the original TW-LAVA protocol with 10,000 validation images.

Pruning is a prerequisite for complementary losses. Without selection, SCL-NL collapses to 1.66% on MNIST, 5.10% on CIFAR-10, and 2.25% on CIFAR-20 (Table 1): the 25% diagonal-noise samples teach "not the true class", which a loss

Table 10. Training on the selections of Table 6 with CLIP features (cuts of Table 1; mean ± std over three seeds). "$w = 0$" is the CAVA margin without the validation potential. Best per column and block in **bold**.

Selection		CIFAR-10	CIFAR-20	CLCIFAR-10	CLCIFAR-20	CLMIN-10
(a) Selection only: SCL-NL on the retained pool						
TW-LAVA	Fixed p	8.41±1.86	3.10±0.31	42.79±1.79	7.69±0.08	15.22±2.17
CAVA (ours)	Fixed p	**74.11**±0.32	12.05±1.05	**79.06**±0.67	**10.64**±0.89	27.11±1.90
CAVA (ours)	Adaptive cut	64.19±3.36	7.01±1.74	77.94±0.43	8.48±0.11	25.11±4.43
CAVA, $w = 0$	Fixed p	73.20±0.06	**13.98**±0.81	78.27±0.37	8.75±0.16	25.44±2.14
CAVA, $w = 0$	Adaptive cut	64.29±0.65	8.77±0.51	76.96±0.88	9.39±1.01	**29.00**±1.76
(b) Selection with recycling: CE+SCL-NL						
TW-LAVA	Fixed p	46.68±0.58	28.85±0.96	28.18±2.10	5.24±0.29	16.11±1.65
CAVA (ours)	Fixed p	**87.91**±0.55	**67.37**±0.24	80.59±0.23	**42.12**±0.77	50.44±1.26
CAVA (ours)	Adaptive cut	87.24±0.28	62.67±0.34	**81.84**±0.29	41.66±0.90	49.89±0.70
CAVA, $w = 0$	Fixed p	87.38±0.44	66.71±0.28	80.06±0.20	39.08±0.86	**51.56**±1.17
CAVA, $w = 0$	Adaptive cut	87.01±0.38	63.15±0.61	80.99±0.13	40.58±0.63	50.00±0.88

Table 11. The two components of CAVA in training (CLIP features, fixed cut; mean ± std over three seeds). Table 7 gives the detection AUC of the same four combinations on CIFAR-10. Best per column in **bold**.

		CIFAR-10 ($\eta = 25\%$)		CLCIFAR-10 ($\eta = 3.98\%$)	
Label cost	**Score**	SCL-NL	CE+SCL-NL	SCL-NL	CE+SCL-NL
noisy-group	calibrated (TW-LAVA)	8.41±1.86	46.68±0.58	42.79±1.79	28.18±2.10
noisy-group	margin	67.57±1.31	85.46±0.08	41.21±4.00	31.08±1.89
complementary	calibrated	58.79±2.12	80.81±0.32	71.42±0.16	73.26±0.49
complementary	margin (CAVA)	**74.11**±0.32	**87.91**±0.55	**79.06**±0.67	**80.59**±0.23

Table 12. Training with small validation sets (CLIP features, CIFAR-10, CE+SCL-NL, fixed cut $p = 25\%$, mean ± std over three seeds, final-epoch accuracy on the test images outside the validation set). Best per column in **bold**.

N_v	50	100	200	500	2,000
CAVA (ours)	**82.92**±0.17	**84.36**±0.06	**85.47**±0.10	**86.61**±0.15	**87.91**±0.55
CAVA, $w = 0$	82.76±0.21	83.74±0.11	85.37±0.41	86.08±0.34	87.38±0.44

that treats every label as an exclusion cannot absorb. Every pruned configuration improves on this baseline by an order of magnitude.

CAVA cuts at the noise rate. Because the label margin ranks the noise at the low end of every task (Figure 11), the cut $p = \eta$ gives an almost oracle split on all three datasets: the recycled branch contains 0.70%, 0.54%, and 4.44% valid labels on MNIST, CIFAR-10, and CIFAR-20, the retained pool at most 1.48% diagonal noise, and $p = \eta$ is the best cut for the recycling objective CE+SCL-NL (Table 13). The adaptive cut flags nearly the same set without knowing η (last CAVA row of Tables 14 and 15). TW-LAVA degrades with its detector: as its AUC falls from 0.9994 to 0.9833 and 0.9087, the valid labels in its recycled branch at $p = \eta$ rise from 1.24% to 9.84% and 24.34%, which is why its best cut moves away from η.

Cutting beyond the noise rate biases the complementary labels. A cut larger than η can only move valid labels into the CE branch, close to the floor $(p - \eta)/p$ of 28.6% and 44.4% at $p = 35\%$ and 45%. It also changes the labels that remain: the lowest-ranked valid samples are those whose given label is a plausible class for the image, so over-pruning removes confusable pairs such as $4 \leftrightarrow 9$ and $1 \leftrightarrow 7$ and leaves a transition matrix with near-empty cells ($14/90$ cells at $p = 45\%$ on MNIST versus $0/90$ at $p = 25\%$, and $38/380$ versus $5/380$ on CIFAR-20; Figures 12 and 13). SCL-NL assumes uniform transitions and degrades accordingly (CAVA: $99.35\% \to 79.50\%$ on MNIST and $81.67\% \to 55.13\%$ on CIFAR-10 from $p = 25\%$ to 45%), whereas FWD, which estimates the transition matrix explicitly, is affected less. With TW-LAVA on CIFAR-20 the ordering is reversed, because the contamination left by its weaker detector outweighs this bias, so larger cuts help its selection-only training.

Table 13. Detector quality and branch composition across datasets under the identical synthetic protocol ($\eta = 25\%$, AutoAug, combined CE+SCL-NL loss), for the baseline TW-LAVA and for CAVA with the same feature extractor. The detection AUC is 0.9994, 0.9833, 0.9087 for TW-LAVA and 1.0000, 1.0000, 0.9971 for CAVA on MNIST, CIFAR-10, CIFAR-20. CE Contam. denotes the fraction of the recycled CE branch that consists of valid complementary samples ($\bar{y}_i \neq y_i$), which become false-positive ordinary labels when recycled. CLL Contam. denotes the fraction of the retained complementary-label pool that still consists of diagonal-noise samples ($\bar{y}_i = y_i$). Accuracy denotes the CE+SCL-NL test accuracy: a single run (seed 1126) for TW-LAVA, which uses all 10,000 test images as the validation set, and the mean over seeds 1126, 42, and 2023 for CAVA, which uses a fixed validation set of 2,000 test images the selection statistics are from seed 1126. For $p > \eta$ the CE branch necessarily contains valid samples, so its contamination is at least $1 - \eta/p$ for any detector.

		MNIST			CIFAR-10			CIFAR-20		
Method	**Cut**	**CE Con.**	**CLL Con.**	**Acc.**	**CE Con.**	**CLL Con.**	**Acc.**	**CE Con.**	**CLL Con.**	**Acc.**
TW-LAVA	p=0.25	1.24%	0.41%	99.53	9.84%	3.28%	86.34	24.34%	8.11%	64.10
	p=0.35	28.73%	0.09%	99.15	30.03%	0.78%	85.31	36.81%	4.44%	63.78
	p=0.45	44.54%	0.08%	99.20	44.99%	0.44%	84.13	47.70%	2.66%	61.87
CAVA (ours)	p=0.25	0.70%	0.23%	**99.57**	0.54%	0.18%	**92.45**	4.44%	1.48%	**76.45**
	p=0.35	28.57%	0.00%	98.47	28.57%	0.00%	85.18	28.69%	0.06%	70.44
	p=0.45	44.44%	0.00%	98.60	44.44%	0.00%	82.32	44.47%	0.02%	64.15

Table 14. MNIST ablation under the synthetic diagonal-noise protocol with noise ratio $\eta = 25\%$ and AutoAug. CLL Contam. denotes the fraction of the retained complementary-label pool that still consists of diagonal-noise samples ($\bar{y}_i = y_i$). CE Contam. denotes the fraction of the recycled CE branch that consists of valid complementary samples ($\bar{y}_i \neq y_i$), for which treating $\bar{y}_i$ as an ordinary label is incorrect. FWD uses the transition matrix computed from the synthetic corruption record, which includes the diagonal. The upper block is the baseline TW-LAVA and the lower block is CAVA with the same feature extractor. Static p denotes the fixed cut. For CAVA, the adaptive cut $s_i \leq 0$ needs no ratio (*: seed 1126). TW-LAVA entries are single runs (seed 1126) that use all 10,000 test images as the validation set; CAVA entries are means over seeds 1126, 42, and 2023 with a fixed validation set of 2,000 test images (largest standard deviation 7.40, FWD at $p = 0.35$; at most 1.20 elsewhere except FWD at $p = 0.45$, 4.55), and its selection statistics are from seed 1126; the best accuracy in each column is in **bold**.

	Branch Composition			Combined	CLL-only	
Setting	**CLL Size**	**CLL Contam.**	**CE Contam.**	**CE+SCL-NL**	**SCL-NL**	**FWD**
TW-LAVA (noisy-group label cost)						
Static p=0.25	45,000	0.41%	1.24%	99.53	99.32	99.22
Static p=0.35	39,000	0.09%	28.73%	99.15	96.39	98.14
Static p=0.45	33,000	0.08%	44.54%	99.20	55.90	88.99
CAVA (ours; complementary label cost, label margin)						
Static p=0.25	45,000	0.23%	0.70%	99.57	99.35	99.41
Static p=0.35	39,000	0.00%	28.57%	98.47	97.54	89.36
Static p=0.45	33,000	0.00%	44.44%	98.60	79.50	93.33
Adaptive cut $s_i \leq 0$	45,015	0.05%	0.05%	**99.90***	**99.71***	**99.61***

D.2. Real-world Complementary Labels: CLCIFAR-10 and CLCIFAR-20

CLCIFAR-10 and CLCIFAR-20 replace our synthetic corruption with *human-annotated* complementary labels: each of the 50,000 training images carries one crowdsourced complementary label, of which 3.98% (CLCIFAR-10; 1,990 samples) and 2.32% (CLCIFAR-20; 1,161 samples) are annotation errors, in which the worker selected the *true* class of the image as complementary. Two properties distinguish this regime from the synthetic sections. First, both error rates lie below $1/M$, so under the noisy-group label cost of TW-LAVA the effective label weight $\lambda_y \alpha(\eta)$ is negative (Theorem 1) and annotation errors sit on the high-score side; the TW-LAVA rules therefore flag the high-score tail here, the reverse of the synthetic benchmarks, where diagonal noise concentrates at the validation-consistent low end. Under the label margin of CAVA the errors sit on the low side on every dataset (Theorem 2), and the fixed cut flags the low tail as usual. Second, for TW-LAVA the errors are far harder to detect than synthetic flips (error-detection AUC 0.8790 and 0.7243 versus 0.9087–0.9994 in the synthetic regimes), because human mistakes concentrate on genuinely ambiguous images, whereas the margin separates them almost completely (AUC 1.000 and 0.999, Figure 4); the pruning ratios are matched to the small error rates ($p \in \{3\%, 5\%, 8\%, 10\%\}$).

For the real-world CLImage datasets, we report two CE-branch diagnostics. Err. Prec. is the fraction of the recycled CE branch that consists of true annotation errors ($\bar{y}_i = y_i$); it measures whether recycling is reliable. Err. Rec. is the fraction of all annotation errors captured by the CE branch; it measures how much annotation noise is removed from the retained CLL

Table 15. CIFAR-20 ablation under noise ratio $\eta = 25\%$ and AutoAug. Column semantics follow Table 14. CLL Contam. denotes the fraction of the retained complementary-label pool that still consists of diagonal-noise samples ($\bar{y}_i = y_i$), while CE Contam. denotes the fraction of the recycled CE branch that consists of valid complementary samples ($\bar{y}_i \neq y_i$). The corruption record is used only for post-hoc evaluation of branch composition and is never used by either method during OT valuation, thresholding, pruning, recycling, or training. The upper block is the baseline TW-LAVA and the lower block is CAVA with the same feature extractor. Static p denotes the fixed cut. For CAVA, the adaptive cut $s_i \leq 0$ needs no ratio (*: seed 1126). TW-LAVA entries are single runs (seed 1126) that use all 10,000 test images as the validation set; CAVA entries are means over seeds 1126, 42, and 2023 with a fixed validation set of 2,000 test images (largest standard deviation 2.36), and its selection statistics are from seed 1126; the best accuracy in each column is in **bold**.

	Branch Composition			Combined	CLL-only	
Setting	**CLL Size**	**CLL Contam.**	**CE Contam.**	**CE+SCL-NL**	**SCL-NL**	**FWD**
TW-LAVA (noisy-group label cost)						
Static p=0.25	37,500	8.11%	24.34%	64.10	12.74	15.53
Static p=0.35	32,500	4.44%	36.81%	63.78	16.31	19.34
Static p=0.45	27,500	2.66%	47.70%	61.87	17.77	27.16
CAVA (ours; complementary label cost, label margin)						
Static p=0.25	37,500	1.48%	4.44%	**76.45**	17.23	18.34
Static p=0.35	32,500	0.06%	28.69%	70.44	**28.11**	35.57
Static p=0.45	27,500	0.02%	44.47%	64.15	23.24	**40.17**
Adaptive cut $s_i \leq 0$	38,975	4.12%	1.19%	73.63*	14.70*	15.82*

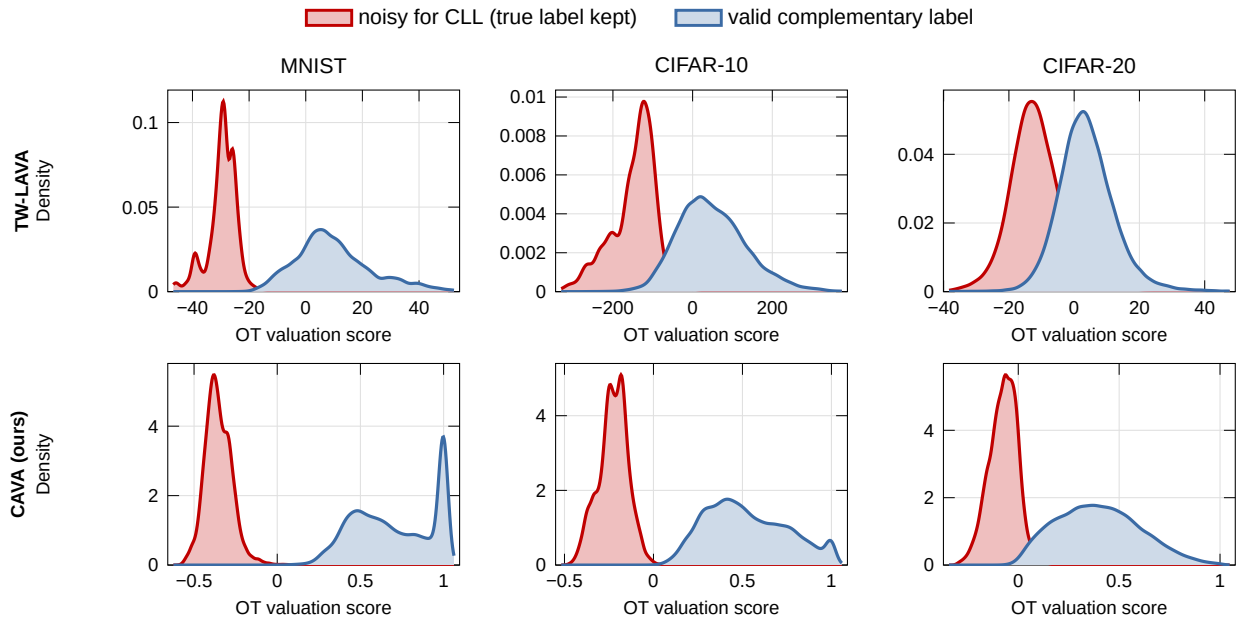

Figure 11. OT valuation score distributions for diagonal-noise samples ($\bar{y}_i = y_i$) and valid complementary samples ($\bar{y}_i \neq y_i$) on MNIST, CIFAR-10, and CIFAR-20 under noise ratio $\eta = 25\%$ and seed 1126, with the same feature extractor in both rows (validation set: 10,000 test images for TW-LAVA, a fixed subset of 2,000 for CAVA). *Top:* TW-LAVA, the baseline with the noisy-group label cost. The two groups become less separable as the dataset becomes harder, with AUC decreasing from 0.9994 on MNIST to 0.9833 on CIFAR-10 and 0.9087 on CIFAR-20, which explains the increasing branch contamination reported in Table 13. *Bottom:* CAVA with the complementary label cost and the label margin. Diagonal noise stays on the low side and the separation is close to complete on all three datasets (AUC 1.000, 1.000, and 0.998).

pool. CLL Contam. is the fraction of annotation errors remaining in the retained complementary-label pool.

Detection quality drives everything downstream. Figure 4(a) in the main text shows the score distributions. Under TW-LAVA, annotation errors ($\bar{y}_i = y_i$) lie on the *high* side, the reverse of the synthetic setting (Figure 11), and overlap strongly with valid complementary labels (detection AUC 0.879 on CLCIFAR-10 and 0.724 on CLCIFAR-20); under CAVA they lie on the low side and are almost disjoint from the valid labels (AUC 1.000 and 0.999). For TW-LAVA, the two datasets place the

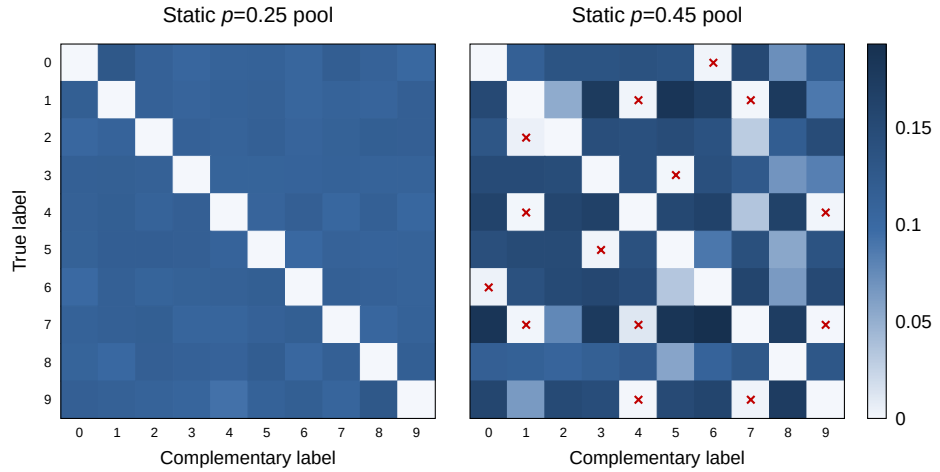

Figure 12. Pruning-induced complementary-label bias (MNIST, $\eta = 25\%$; pools of the original TW-LAVA protocol). Row-normalized transition matrices $P(\bar{Y} \mid Y)$ of the CLL pool at p=0.25 (left, near-uniform) and p=0.45 (right); $\times$ marks near-empty transitions ($< 2\%$ row mass). Aggressive pruning empties the transitions of visually confusable pairs, violating the uniform assumption of SCL-NL while the estimated Q of FWD absorbs the bias.

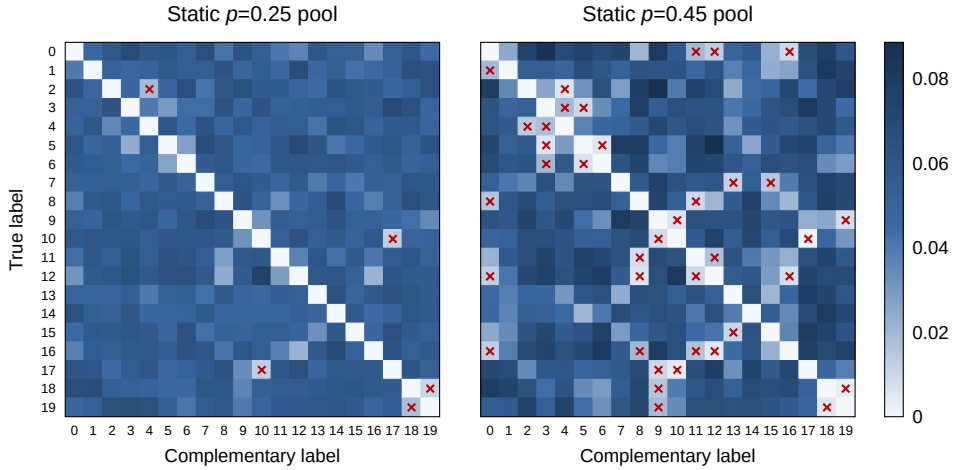

Figure 13. Pruning-induced complementary-label bias on CIFAR-20 ($\eta = 25\%$), as in Figure 12: row-normalized transitions $P(\bar{Y} \mid Y)$ of the CLL pool at p=0.25 (left) and p=0.45 (right); $\times$ marks near-empty transitions ($< 2\%$ row mass; $5/380$ versus $38/380$ cells).

detector at different operating points. On CLCIFAR-10, of the $1{,}500$ highest-scoring samples (the top 3% of the OT ranking, i.e. the CE branch at p=0.03), 481 samples (32.1%) are actual annotation errors. Since errors make up only 3.98% of the dataset, a random cut of the same size would contain $\sim$60; the detector concentrates them $8.1\times$ above the base rate, so even shallow static cuts remove a meaningful fraction of the errors (residual $3.98 \to 2.22\%$ at p=0.10) and every pruned FWD run beats the unpruned baseline ($51.76 \to 59.4$–67.3). On CLCIFAR-20 the detector is far weaker: the same $1{,}500$-sample cut contains only 117 errors (7.8%, a $3.4\times$ enrichment over the 2.32% base rate), the residual barely moves ($2.32 \to 1.73\%$), and static pruning yields only $+1.4$ to $+2.6$ points over the baseline.

Recycling fails when detection precision is low. On synthetic data, the combined objective performs well because the recycled CE branch has sufficiently high precision. On real-world complementary labels, however, the base rate of annotation errors is only 2–4%, so even a detector with useful recall can produce a low-precision recycled branch. Tables 16 and 17 show this clearly. For CLCIFAR-10, TW-LAVA at $p = 10\%$ captures 49.9% of all annotation errors, but only 19.9% of its CE branch consists of true errors. For CLCIFAR-20 at the same cut, it captures 32.9% of the errors with a CE precision of only 7.6%. Most of its recycled samples are thus valid complementary labels falsely treated as ordinary labels. Therefore, with the noisy-group score, real-world complementary labels call for discarding flagged samples rather than recycling them. The

Table 16. CLCIFAR-10 ablation with real-world human complementary labels under AutoAug, for the baseline TW-LAVA and for CAVA with the same feature extractor. Static p denotes the fixed cut; for CAVA, the adaptive cut $s_i \leq 0$ needs no ratio (*: seed 1126). TW-LAVA rows use the original TW-LAVA protocol with 10,000 validation images. Err. Prec. is the fraction of the recycled CE branch that consists of true annotation errors ($\bar{y}_i = y_i$), while Err. Rec. is the fraction of all annotation errors captured by the CE branch. CLL Contam. is the fraction of annotation errors remaining in the retained complementary-label pool. The annotation record is used only for post-hoc evaluation of branch composition and, where indicated, to construct known-Q FWD; neither method uses true labels for selection. TW-LAVA entries are single runs (seed 1126) that use all 10,000 test images as the validation set; CAVA entries are means over seeds 1126, 42, and 2023 with a fixed validation set of 2,000 test images (largest standard deviation 2.04), and its selection statistics are from seed 1126; the best accuracy in each column is in **bold**; "–": not applicable.

Setting	CE branch			CLL pool		Combined		CLL-only	
	Size	**Err. Prec.**	**Err. Rec.**	**Size**	**Contam.**	**CE+SCL-NL**	**CE+FWD**	**SCL-NL**	**FWD**
No pruning ($p = 0$)	–	–	–	50,000	3.98%	–	–	50.68	51.76
TW-LAVA (noisy-group label cost)									
Static $p = 0.03$	1,500	32.1%	24.2%	48,500	3.11%	46.78	46.19	53.25	59.38
Static $p = 0.05$	2,500	26.7%	33.5%	47,500	2.79%	46.48	52.44	53.81	63.81
Static $p = 0.08$	4,000	21.7%	43.7%	46,000	2.44%	43.65	46.78	50.68	67.29
Static $p = 0.10$	5,000	19.9%	49.9%	45,000	2.22%	41.21	45.41	49.41	67.21
CAVA (ours; complementary label cost, label margin)									
Static $p = 0.03$	1,500	100.0%	75.4%	48,500	1.01%	82.42	82.71	74.28	74.05
Static $p = 0.05$	2,500	79.6%	100.0%	47,500	0.00%	83.46	83.20	80.45	**75.99**
Static $p = 0.08$	4,000	49.8%	100.0%	46,000	0.00%	71.16	72.37	75.68	73.57
Static $p = 0.10$	5,000	39.8%	100.0%	45,000	0.00%	63.54	65.46	65.76	67.14
Adaptive cut $s_i \leq 0$	1,990	99.8%	99.8%	48,010	0.01%	**87.11***	**85.25***	**81.45***	75.10*

Table 17. CLCIFAR-20 ablation with real-world human complementary labels under AutoAug, for the baseline TW-LAVA and for CAVA with the same feature extractor. Static p denotes the fixed cut; for CAVA, the adaptive cut $s_i \leq 0$ needs no ratio (*: seed 1126). TW-LAVA rows use the original TW-LAVA protocol with 10,000 validation images. Err. Prec. is the fraction of the recycled CE branch that consists of true annotation errors ($\bar{y}_i = y_i$), while Err. Rec. is the fraction of all annotation errors captured by the CE branch. CLL Contam. is the fraction of annotation errors remaining in the retained complementary-label pool. The annotation record is used only for post-hoc evaluation of branch composition and, where indicated, to construct known-Q FWD; neither method uses true labels for selection. TW-LAVA entries are single runs (seed 1126) that use all 10,000 test images as the validation set; CAVA entries are means over seeds 1126, 42, and 2023 with a fixed validation set of 2,000 test images (largest standard deviation 2.38), and its selection statistics are from seed 1126; the best accuracy in each column is in **bold**; "–": not applicable.

Setting	CE branch			CLL pool		Combined		CLL-only	
	Size	**Err. Prec.**	**Err. Rec.**	**Size**	**Contam.**	**CE+SCL-NL**	**CE+FWD**	**SCL-NL**	**FWD**
No pruning ($p = 0$)	–	–	–	50,000	2.32%	–	–	8.98	23.73
TW-LAVA (noisy-group label cost)									
Static $p = 0.03$	1,500	7.8%	10.1%	48,500	2.15%	9.77	16.75	8.98	25.15
Static $p = 0.05$	2,500	7.2%	15.5%	47,500	2.07%	9.08	14.10	9.08	24.95
Static $p = 0.08$	4,000	7.7%	26.5%	46,000	1.85%	10.30	16.02	9.38	25.44
Static $p = 0.10$	5,000	7.6%	32.9%	45,000	1.73%	11.07	15.62	9.65	26.29
CAVA (ours; complementary label cost, label margin)									
Static $p = 0.03$	1,500	75.9%	98.0%	48,500	0.05%	51.28	**53.74**	12.11	32.88
Static $p = 0.05$	2,500	46.3%	99.7%	47,500	0.01%	48.09	52.65	13.41	34.92
Static $p = 0.08$	4,000	29.0%	99.9%	46,000	0.00%	44.50	48.88	14.95	37.71
Static $p = 0.10$	5,000	23.2%	99.9%	45,000	0.00%	42.25	46.29	**15.79**	**41.78**
Adaptive cut $s_i \leq 0$	1,062	96.0%	87.9%	48,938	0.29%	**51.66***	53.56*	11.72*	30.54*

label margin changes this picture: at $p = 3\%$ its recycled branch consists of 100% true errors on CLCIFAR-10 and 75.9% on CLCIFAR-20 (Figure 4), which is why the recycling objectives gain the most in Table 1.

SCL-NL degrades with annotator bias. The uniform-transition assumption of SCL-NL is violated here: the human complementary-label marginal is non-uniform: per-class counts range over $4{,}097$–$6{,}445$ (uniform $5{,}000$) on CLCIFAR-10 but

1,676–3,538 (uniform 2,500) on CLCIFAR-20. The milder CLCIFAR-10 bias leaves SCL-NL functional (50.7% unpruned, 53.8% at the best static cut) while the stronger CLCIFAR-20 bias breaks it outright (∼9% at chance 5%, in every configuration). FWD, whose Q models the biased transition explicitly, is unaffected by this violation and is the only complementary loss that remains usable across both datasets.

D.3. Further Domains

Table 18 lists the results plotted in Figure 7; the datasets are described in Appendix C.1.

Table 18. Test accuracy (%) on three further domains with synthetic diagonal noise (Figure 7; datasets in Appendix C.1). The validation set holds 50 test images per class (200 for Places365-10), and the fixed cut is $p = \eta$. Seed 1126; best per column and block in **bold**.

		FOOD15			RESISC45-20			Places365-10		
Method	**Selection**	5%	15%	25%	5%	15%	25%	5%	15%	25%
(a) Selection only: SCL-NL on the retained pool										
No pruning	–	8.33	3.17	1.90	2.77	0.70	0.17	44.73	3.09	0.84
TW-LAVA	Fixed p	7.60	7.20	13.00	3.23	13.57	24.33	48.39	46.04	56.94
CAVA (ours)	Fixed p	**22.50**	**21.10**	**18.90**	23.17	23.43	**24.70**	**75.08**	**74.03**	**73.90**
CAVA (ours)	Adaptive cut	18.67	17.57	11.60	**25.17**	**26.23**	22.93	71.60	73.86	71.22
(b) Selection with recycling: CE+SCL-NL										
TW-LAVA	Fixed p	8.27	22.93	42.17	3.00	44.43	69.10	38.95	53.35	70.40
CAVA (ours)	Fixed p	**47.33**	**60.43**	**63.83**	**71.93**	**82.93**	**87.30**	80.39	**82.16**	**84.14**
CAVA (ours)	Adaptive cut	44.17	58.20	60.20	68.70	82.83	86.03	**80.49**	**82.16**	83.45

D.4. Other Noise Rates

The flip law makes the noise rate the decisive variable for TW-LAVA, so we vary it on CIFAR-10. Figure 14 reports CIFAR-10 results under $\eta \in \{0.15, 0.35, 0.45\}$ using AutoAugment, against the fixed cut p. Comparing the best cut of each method, CAVA is higher for both objectives at every noise ratio, by 2.5 to 10.2 points. At $\eta = 35\%$ and 45% the best cut of both methods is $p = \eta$; at $\eta = 15\%$ it is $p = 25\%$, larger than η. Table 19 reports the complementary regime $\eta \leq 10\%$, where TW-LAVA is reversed or blind and CAVA is best at every noise ratio.

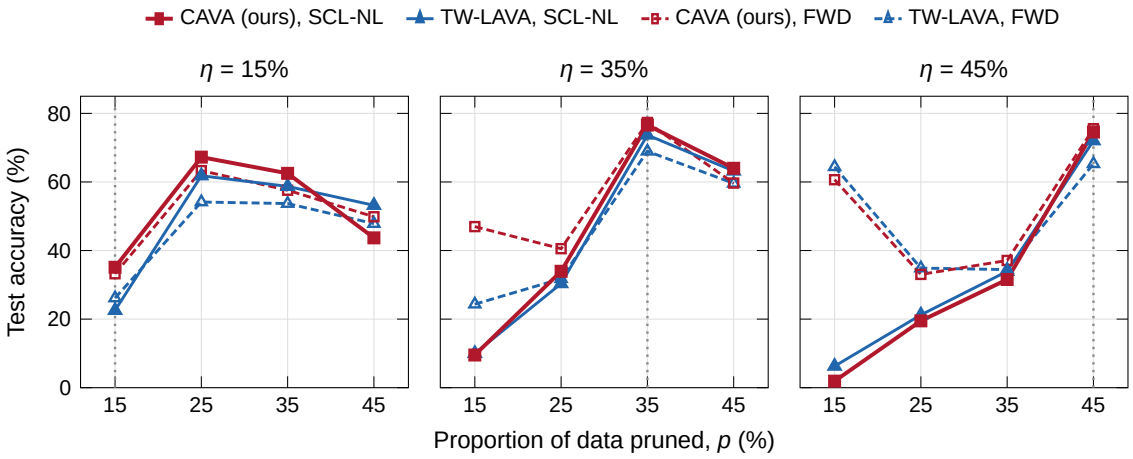

Figure 14. **Other noise ratios.** CIFAR-10 test accuracy (%) of TW-LAVA and CAVA against the fixed cut $p \in \{15\%, 25\%, 35\%, 45\%\}$ at diagonal-noise ratios $\eta \in \{15\%, 35\%, 45\%\}$, above the crossover $1/M$, with the selection-only objective (SCL-NL solid, FWD dashed). Same feature extractor and validation set for both methods; AutoAugment; single runs (seed 1126) under the protocol of Table 1, accuracy after the final epoch on the held-out test images. The dotted line marks $p = \eta$.

Table 19. CIFAR-10 test accuracy (%) at small noise rates, $\eta \leq 1/M = 10\%$, where the label cost of TW-LAVA is reversed or blind (Theorem 1). Both methods use the same fixed cuts; TW-LAVA keeps its original OT settings here (Appendix C.2). SCL-NL and No pruning: mean±std over three seeds; other objectives: seed 1126. Best per noise rate and objective in **bold**; "–": no recycled branch without pruning.

		TW-LAVA (noisy-group label cost)				CAVA (ours, complementary label cost)			
η	Cut	CE+SCL-NL	CE+FWD	SCL-NL	FWD	CE+SCL-NL	CE+FWD	SCL-NL	FWD
2%	none	–	–	68.53±2.28	69.30	–	–	68.53±2.28	69.30
	$p = 2\%$	48.38	29.57	67.22±3.19	68.36	**84.51**	**84.41**	82.26±0.25	**82.16**
	$p = 5\%$	24.57	22.64	62.38±1.80	66.03	77.31	75.89	**82.64±0.11**	81.58
	$p = 10\%$	9.95	16.25	46.66±1.27	64.78	56.74	53.64	70.78±1.37	67.61
	$p = 15\%$	7.70	14.18	11.25±1.30	31.19	42.59	40.69	47.84±9.72	44.61
5%	none	–	–	38.69±0.85	38.94	–	–	38.69±0.85	38.94
	$p = 2\%$	18.71	2.84	31.65±1.67	33.51	66.12	56.27	61.02±1.02	55.23
	$p = 5\%$	6.33	1.86	30.10±1.38	31.91	**86.46**	**86.25**	**82.15±0.23**	**82.05**
	$p = 10\%$	2.95	2.51	23.16±3.82	37.92	72.53	70.53	81.11±0.19	71.50
	$p = 15\%$	3.41	2.59	6.84±0.72	18.15	58.26	54.84	38.63±5.02	54.52
10%	none	–	–	7.85±1.66	10.00	–	–	7.85±1.66	10.00
	$p = 2\%$	18.32	18.85	11.20±2.04	10.00	30.52	31.11	25.39±0.62	27.10
	$p = 5\%$	21.65	18.36	10.72±0.74	10.00	59.31	57.55	47.27±1.63	44.46
	$p = 10\%$	14.89	14.28	12.42±4.39	18.11	**89.04**	**88.42**	**81.35±0.48**	**80.85**
	$p = 15\%$	12.51	13.32	11.40±1.88	10.00	77.31	76.42	36.52±2.18	35.25

D.5. Other CLL Objectives

We examine whether the benefit of CAVA is tied to a specific CLL objective or carries over as a plug-in data-selection module. To isolate selection, recycling is disabled: each objective is trained only on the retained complementary-labeled pool (Eq. (3.13)). We use two settings that cover both regimes of the flip law. On synthetic CIFAR-10 with $\eta = 25\%$, above the crossover $1/M$, CAVA uses the adaptive cut $s_i \leq 0$, which needs no hyperparameter, and TW-LAVA the fixed cut $p = 25\%$. On human-annotated CLCIFAR-10 with $\eta = 3.98\%$, below the crossover, both TW-LAVA and CAVA use the fixed cut $p = 5\%$, the smallest grid cut that covers the error rate. Both methods are scored with the same feature extractor and the same 2,000 validation images, and No selection trains on the full pool.

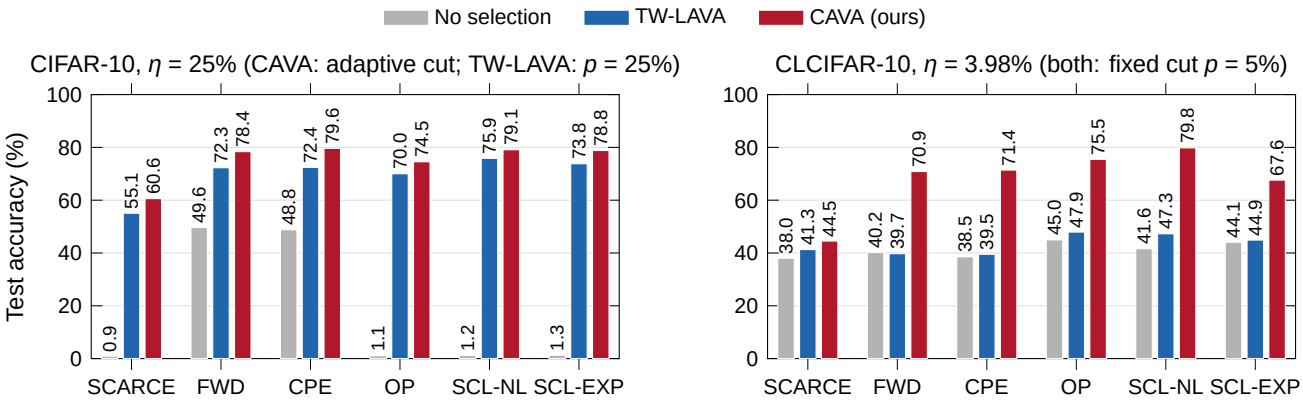

Figure 15. **CAVA as a plug-in selection module across CLL objectives.** Test accuracy (%) of six objectives without selection, with TW-LAVA, and with CAVA (selection only, no recycling; AutoAugment; seed 1126; final-epoch accuracy on the held-out test images, protocol of Table 1). All objectives use SGD as in the main experiments except OP, which uses Adam with learning rate 10^{-4} as in its official implementation. On synthetic CIFAR-10 ($\eta = 25\%$), CAVA uses the adaptive cut and TW-LAVA, which has no adaptive cut, the fixed cut $p = 25\%$; on human-annotated CLCIFAR-10 ($\eta = 3.98\%$), both methods use the fixed cut $p = 5\%$. TW-LAVA and CAVA use the same 2,000 validation images.

Figure 15 reports the test accuracy of each objective without selection, after TW-LAVA selection, and after CAVA selection. On CIFAR-10, every objective needs selection: without it, SCL-NL, SCL-EXP, OP, and SCARCE collapse (1.18%, 1.26%,

1.06%, and 0.90%), and FWD and CPE reach only 49.64% and 48.81%. With the adaptive cut, CAVA lifts SCL-NL, SCL-EXP, OP, and SCARCE to 79.09%, 78.81%, 74.54%, and 60.58%, and FWD and CPE to 78.56% and 79.59%. TW-LAVA also helps here, where η lies above the crossover and its score has the correct direction, but CAVA is higher for all six objectives by 3.2 to 7.2 points, without choosing a ratio. On CLCIFAR-10, below the crossover, TW-LAVA stays within 6 points of No selection, whereas CAVA improves FWD, CPE, OP, SCL-NL, and SCL-EXP by 23.5 to 38.2 points and SCARCE by 6.5 points. Because recycling depends on detector precision, this analysis uses selection-only training; Appendix D.7 studies when removed samples can be reused as ordinary-labeled data.

D.6. Data Augmentation

AutoAugment is the default augmentation in all experiments. Figure 16 repeats the CIFAR-10 setting ($\eta = 25\%$) with Cutout and RandAugment for TW-LAVA and for CAVA under the protocol of Table 1. The conclusions do not depend on the augmentation policy: among the fixed cuts both methods are best at $p = \eta$ for every objective and policy, and the adaptive cut of CAVA is within one point of that cut or above it; with recycling, CAVA reaches 91.4 to 92.3% under the three policies (best setting per policy), 3.3 to 4.2 points above the best TW-LAVA setting under the same policy and objective.

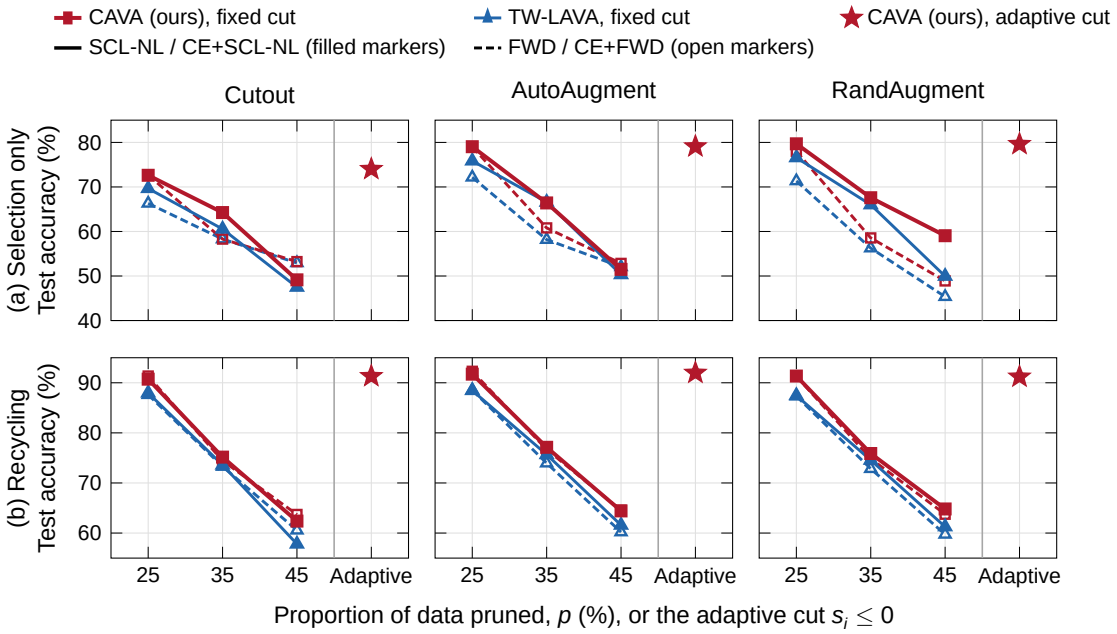

Figure 16. **Data augmentation.** CIFAR-10 test accuracy (%) at $\eta = 25\%$ of TW-LAVA and CAVA against the fixed cut $p \in \{25\%, 35\%, 45\%\}$, and of CAVA with the adaptive cut $s_i \leq 0$ (stars right of the gray line, with SCL-NL and CE+SCL-NL), under Cutout, AutoAugment, and RandAugment. Row (a) uses the selection-only objective in Eq. (3.13) with SCL-NL (solid, filled markers) and FWD (dashed, open markers); row (b) the recycling objective in Eq. (3.14) with CE+SCL-NL and CE+FWD. Same feature extractor, validation set and OT settings for both methods; single runs (seed 1126), accuracy after the final epoch on the held-out test images, protocol of Table 1.

D.7. Recycling and Branch Quality

Recycling helps only when the flagged set is pure (Theorem 4). This subsection measures that purity with the metrics of Table 4.

CAVA. Wherever the cut does not exceed the noise rate, the recycled branch of CAVA is nearly pure: it contains 0.54% valid labels on CIFAR-10 at $p = \eta$ (Table 13), and the adaptive cut reaches 99.7% and 96.0% precision on CLCIFAR-10 and CLCIFAR-20 (Tables 16 and 17). Recycling then turns the detected noise into supervision: on CIFAR-10 it raises SCL-NL from 79.26% to 91.57%, and on CLCIFAR-20 from 10.78% to 47.71% with the fixed cut (Table 1).

TW-LAVA. The baseline shows the failure mode that the bound predicts. On CIFAR-10 (Table 20), fixed cuts trade precision for recall: at $p = 25\%$ the recycled branch is 95.4% pure, and larger cuts recover more noise but recycle more valid labels,

so accuracy falls from 88.46% to 61.58%. On human labels its recycled branches contain only 9.3 to 43.5% true errors, and recycling lowers its accuracy on CLCIFAR-10 by 6 to 13 points; on CLCIFAR-20, where SCL-NL alone is near chance, it adds 3 to 5 points (Appendix D.2).

Pruning and recycling tolerate errors differently. An extreme score is not proof of noise: atypical valid images, distribution shift between pool and validation set, or limited coverage of rare modes can also produce one. The two uses of a flagged sample are asymmetric. For pruning, a false positive costs one weak exclusion out of tens of thousands, whereas a retained diagonal label teaches a wrong exclusion, so pruning remains safe under imperfect detection. For recycling, a false positive injects a wrong ordinary label. Recycling is therefore an option gated by the measured precision, and selection only is the default.

Table 20. Branch quality of recycling with TW-LAVA on synthetic CIFAR-10 ($\eta = 25\%$, CE+SCL-NL, seed 1126). The metrics of Table 4 use the true labels for evaluation only; the last column is the test accuracy.

Setting	CE Precision	CE Recall	F1	$\lvert\mathbf{I}^{\text{rec}}\rvert$	CE Contam.	CLL Contam.	Test Accuracy
Fixed cut $p = 25\%$	95.37%	95.37%	95.37%	12,500	4.63%	1.54%	88.46%
Fixed cut $p = 35\%$	71.12%	99.57%	82.97%	17,500	28.88%	0.17%	75.67%
Fixed cut $p = 45\%$	55.46%	99.82%	71.30%	22,500	44.54%	0.08%	61.58%

D.8. The Validation Set

This subsection examines the clean validation set from two sides: how much its labels are worth when spent on direct supervision instead, and how small it can be.

Valuation versus direct supervision.

CAVA follows the standard protocol of OT-based data valuation [16, 17]: a small clean validation set defines the target distribution against which training samples are valued. Every valuation baseline in our comparison (LAVA, SAVA) assumes the same validation set, so this assumption grants no advantage specific to CAVA. Moreover, the validation set enters only the OT scores; it is never used to train the classifier. Its cost is also one-time and independent of the pool size: the CIFAR-10-1M study in Appendix D.9 keeps the same 2,000 validation examples while the training pool grows twentyfold.

A natural question remains: would the 2,000 clean labels be more valuable if they were spent directly on supervised training rather than on valuation? To quantify this, we train a classifier with the standard CE loss on n clean labels, stratified per class, for budgets $n \in \{500, 1{,}000, 2{,}000, 5{,}000, 10{,}000\}$ on CIFAR-10 and $n \in \{500, 1{,}000, 2{,}000, 4{,}000, 7{,}000\}$ on CIFAR-20. The protocol matches the main experiments exactly (ResNet-18, AutoAugment, SGD with the same schedule and weight decay, 200 epochs, accuracy after the final epoch on the held-out test images, three runs per budget with seeds 1126, 42, and 2023), except that the batch size is $\min(128, n)$ so that small budgets receive a comparable number of updates. This baseline faces no label noise at all, whereas CAVA operates under 25% complementary-label noise.

Table 21. What 2,000 clean labels are worth ($\eta = 25\%$; mean $\pm$ std over three seeds). CE from scratch trains on these labels alone. CAVA uses them only as its validation set, together with the 50,000 noisy complementary labels (CE+SCL-NL, $p = 25\%$), on three fixed representations.

Representation	Method with 2,000 clean labels	CIFAR-10	CIFAR-20
none	CE from scratch	71.52±0.66	43.08±2.24
supervised embedder	CAVA (ours), recycling	91.57±0.14	75.34±0.17
CLIP ViT-B/32	CAVA (ours), recycling	87.91±0.55	67.37±0.24
validation images only	CAVA (ours), recycling	75.22±0.67	49.21±0.85

Table 21 and Figure 5(b) give the answer. Against training from scratch, the validation set is worth more as an anchor: with 2,000 labels direct supervision reaches 71.52% on CIFAR-10 and 43.08% on CIFAR-20, CAVA 91.57% and 75.34%, and direct supervision does not reach this level even with 10,000 and 7,000 labels (88.79% and 68.09%). The table also checks that CAVA does not depend on one particular fixed representation. The supervised embedder of the main tables has itself been trained on the 50,000 clean training labels (Appendix C.2); with CLIP features, which have seen no label of the pool, CAVA reaches 87.91% and 67.37%, still far above training from scratch, and with the network trained on the 2,000 validation images only, the weakest of the three representations, it reaches 75.22% and 49.21%, still above training from scratch on

the same labels (Appendix C.5). Table 22 adds that, with the supervised embedder, CAVA keeps its level with as few as 200 validation images; on label-free features the detection AUC on CIFAR-10 falls from 0.994 at $N_v = 2{,}000$ to 0.978 at $N_v = 50$ (Appendix C.5).

Scope and fairness of the comparison. CAVA is a valuation and denoising method for CLL: the question it answers is which samples of a given noisy complementary-label pool should be kept, recycled, or removed. Every compared method answers this same question under the same validation set, so the comparison is information-equal across the table. Holding out a small clean set for scores, thresholds, and hyperparameter selection, without using it for training, is standard practice throughout the learning-with-noisy-labels literature [8, 11, 21] and is inherited from the data valuation line [10, 16, 17]; every baseline with tunable hyperparameters implicitly relies on such a set. Methods that instead consume the validation set as training data, whether by direct supervision or by semi-supervised learning over the pool images, solve a different problem and discard the complementary labels that are the object of study; the budget study above therefore measures a fixed clean-annotation budget rather than information equivalence, and Table 21 bounds what direct supervised use achieves at the same budget. Such methods are pipeline alternatives rather than competing valuation methods: any consumer of the validation set can be composed with the CAVA output, which supplies a cleaned complementary-label pool and a recycled ordinary-label set on top, and the validation set itself can also be added to the CE branch of the recycling objective in Eq. (3.14).

Validation-set size. We shrink the validation set on both synthetic benchmarks. For each $N_v \in \{200, 500, 1{,}000, 1{,}500, 2{,}000\}$ we draw a class-stratified subset of the test split ($N_v/10$ per class on CIFAR-10, $N_v/20$ per class on CIFAR-20, seed 1126), recompute the task-wise OT scores against it, and rerun the full fixed-cut pipeline ($p = 0.25$, combined CE+SCL-NL with $\kappa = 0.5$, AutoAugment; three full-pipeline seeds per setting). Table 22 reports the detector AUC, the precision of the recycled branch, and the final test accuracy for TW-LAVA and CAVA. The two methods react differently. TW-LAVA degrades as the validation set shrinks: from $N_v = 2{,}000$ to 200 the precision of its recycled branch falls by 4 points on CIFAR-10 and 6 points on CIFAR-20, and its accuracy by 3.8 and 6.7 points. CAVA is almost insensitive to N_v: its precision stays above 99% and 93%, and its accuracy within 0.2 points on CIFAR-10 and 2.1 points on CIFAR-20 of the $N_v = 2{,}000$ result. With only 200 validation images, CAVA reaches 91.47% on CIFAR-10 and 73.21% on CIFAR-20, above TW-LAVA with ten times more validation data (88.75% and 59.03%). These numbers use the supervised embedder of the main tables; on CLIP features the detection AUC of CAVA falls from 0.994 at $N_v = 2{,}000$ to 0.978 at $N_v = 50$ (Appendix C.5). The label margin needs only enough validation images to place each class, not to estimate the feature distribution of every noisy label group.

Table 22. Effect of the validation-set size N_v on CIFAR-10 and CIFAR-20 ($\eta = 25\%$, fixed cut $p = 25\%$, CE+SCL-NL). Prec. is the share of true diagonal noise in the recycled branch. Mean ± std over three seeds; best per row in **bold**, ties broken by the unrounded value.

Dataset	Method	N_v	200	500	1,000	1,500	2,000
CIFAR-10	TW-LAVA	Detector AUC	0.98±0.00	0.99±0.00	0.99±0.00	1.00±0.00	**1.00**±0.00
		CE-branch Prec. (%)	91.52±0.58	92.99±0.29	94.21±0.23	95.27±0.18	**95.46**±0.09
		Accuracy (%)	84.94±0.06	86.29±0.31	87.82±0.25	88.64±0.07	**88.75**±0.40
	CAVA (ours)	Detector AUC	1.00±0.00	1.00±0.00	1.00±0.00	1.00±0.00	**1.00**±0.00
		CE-branch Prec. (%)	99.23±0.11	99.29±0.14	99.32±0.13	**99.34**±0.14	99.33±0.15
		Accuracy (%)	91.47±0.16	**91.69**±0.17	91.51±0.17	91.69±0.46	91.57±0.14
CIFAR-20	TW-LAVA	Detector AUC	0.90±0.00	0.92±0.00	0.93±0.00	0.94±0.00	**0.94**±0.00
		CE-branch Prec. (%)	72.32±0.31	75.04±0.18	76.77±0.22	78.00±0.04	**78.54**±0.18
		Accuracy (%)	52.29±0.56	55.67±1.07	57.49±0.75	58.18±0.79	**59.03**±0.77
	CAVA (ours)	Detector AUC	0.99±0.00	1.00±0.00	1.00±0.00	1.00±0.00	**1.00**±0.00
		CE-branch Prec. (%)	93.11±0.04	94.30±0.17	94.62±0.14	95.31±0.09	**95.52**±0.15
		Accuracy (%)	73.21±0.49	74.13±0.56	74.32±0.54	74.79±0.34	**75.34**±0.17

D.9. Scaling to One Million Samples

To further evaluate the scalability of CAVA with respect to the number of training samples, we conduct an additional large-sample ablation study on CIFAR-10 under 25% diagonal complementary-label noise ($\eta = 25\%$). Specifically, we expand the original CIFAR-10 training set from 50,000 images to 1,000,000 training samples using label-preserving data augmentation, denoted as CIFAR-10-1M.[1] After augmentation, complementary labels are generated following the same diagonal-noise

[1] CIFAR-10-1M is an augmented benchmark generated from CIFAR-10 using label-preserving augmentation, increasing the training set from 50K to 1M samples. Therefore, it evaluates scalability with respect to training-set size rather than naturally collected semantic diversity.

protocol as in the main experiments, with noise ratio $\eta = 25\%$. The original CIFAR-10 test set is kept unchanged for evaluation. This design allows us to assess whether CAVA remains effective under a substantially enlarged noisy complementary-label training regime without introducing test-set leakage or changing the evaluation protocol.

Figure 5(a) in the main text visualizes the accuracy–efficiency trade-off on CIFAR-10-1M, and Table 23 reports the corresponding numerical results. Without pruning, SCL-NL fails on the enlarged noisy pool and reaches only 0.12%. TW-LAVA peaks at the cut $p = 35\%$, above the noise rate, with 71.05% after 200 epochs (68.26% at $p = 25\%$ and 70.28% at $p = 45\%$) and 67.76% under the 20-epoch schedule of CAVA (61.01% and 55.95% at $p = 25\%$ and 45%), so the gap to CAVA does not come from the training budget. CAVA ranks the noise at the low end of each task (detection AUC 0.996) and peaks at the cut $p = 25\%$ that equals the noise rate, where it reaches 86.04% with SCL-NL after only 20 epochs; deeper cuts remove valid samples and lower accuracy (78.05% and 67.30% at $p = 35\%$ and 45% after 20 epochs, 84.95% and 73.25% after 200 epochs, against 90.51% at $p = 25\%$). Cuts below the noise rate leave diagonal noise in the retained pool and fail for both methods: at $p = 15\%$ TW-LAVA reaches 27.79% and 42.90% after 20 and 200 epochs and CAVA 24.99% after 20 epochs, while its 200-epoch run collapses to a single class (10.00%) from epoch 38 on and does not recover.

The computational results show that the task-wise decomposition keeps both methods feasible at this scale, and that the label margin makes valuation nearly free. The valuation is decomposed into five tasks of 200,000 samples each: TW-LAVA needs about 30 minutes per task, about 2.5 hours in total, whereas CAVA scores all one million samples in about 4 minutes, mainly because the complementary label cost needs no label-to-label transport between noisy groups. Training is the dominant cost, and it depends only on the size of the retained pool: on one dedicated GPU an epoch takes 566 seconds without pruning and 436 and 432 seconds at $p = 25\%$ for TW-LAVA and CAVA (Table 23). These are the times of the second epoch, without the one-time start-up of about 20 seconds; the two methods agree within 4% at the same cut, and the CAVA times at $p = 15\%$ come from a linear fit in the pool size whose residuals at the measured cuts are below 2%. TW-LAVA at its best cut ($p = 35\%$) needs 20.7 hours for 200 epochs, whereas CAVA reaches a higher accuracy, 86.04%, with the 20-epoch schedule in 2.4 hours, a tenfold smaller training budget; with 200 epochs CAVA reaches 90.51% (24.0 hours).

Table 23. Scaling to CIFAR-10-1M ($\eta = 25\%$, SCL-NL, single run). Training time is the measured time of one epoch on a dedicated NVIDIA RTX A6000, multiplied by 20 or 200; it depends only on the size of the retained pool. Valuation adds about 2.5 hours for TW-LAVA and 4 minutes for CAVA. The CAVA times at $p = 15\%$ are estimated by a linear fit in the pool size. The best-epoch columns are for reference only; the text uses the final epoch. Best per column in **bold**; "–": not run.

		Training time		Acc. (%), final epoch		Acc. (%), best epoch	
Method	p (%)	20 ep.	200 ep.	20 ep.	200 ep.	20 ep.	200 ep.
No pruning	0	3.1 h	31.5 h	0.12	–	2.52	–
TW-LAVA	15	2.7 h	27.4 h	27.79	42.90	27.96	43.66
	25	2.4 h	24.2 h	61.01	68.26	61.01	69.64
	35	2.1 h	20.7 h	67.76	71.05	67.80	72.40
	45	1.8 h	17.7 h	55.95	70.28	56.00	71.14
CAVA (ours)	15	2.7 h	27.2 h	24.99	10.00	24.99	21.02
	25	2.4 h	24.0 h	**86.04**	**90.51**	**86.04**	**91.19**
	35	2.0 h	20.0 h	78.05	84.95	78.09	86.70
	45	1.7 h	17.2 h	67.30	73.25	67.30	73.42

Overall, these results indicate that CAVA scales favorably with increased training-set size through task-wise decomposition. However, because CIFAR-10-1M is generated through label-preserving augmentation rather than collected as a naturally large and semantically diverse dataset, we view evaluation on naturally large-scale CLL datasets as an important direction for future work.

D.10. Computational Efficiency of Task-Wise Decomposition

The Sinkhorn algorithm can become computationally expensive when OT is computed over the full training and validation sets. To examine whether the proposed task-wise decomposition improves computational feasibility, we conduct an additional efficiency comparison among LAVA, SAVA, and CAVA on CIFAR-10. Specifically, we use 10K training samples and 10K validation samples and report the wall-clock runtime required for one OT-based data-selection step.

As shown in Table 24, LAVA and task-wise selection with one solve per task require a comparable runtime in the 10K-sample setting because both solve an OT problem of similar size. CAVA solves a problem of the same size in 0.2 seconds on

Table 24. Runtime of one OT-based selection step on CIFAR-10 with 10K training and 10K validation samples. The CAVA time is the mean of three runs on one GPU, with features computed beforehand.

Method	Training Size	Validation Size	Data-Selection Runtime	Remark
LAVA	10K	10K	45 seconds	Full-set OT computation
SAVA	10K	10K	300 seconds	Batch-wise OT computation
TW-LAVA (task-wise)	10K	10K	45 seconds	One task-wise OT selection step
CAVA (ours)	10K	10K	0.2 seconds	OT solve and label margin on precomputed features

precomputed features; its label cost is a fixed $0/1$ table and needs no OT between label groups. However, the key difference lies in their scalability. LAVA performs OT over the full data selection problem, and therefore its computational and memory costs increase rapidly as the training set becomes larger. In contrast, CAVA keeps the OT problem size fixed by decomposing the enlarged training set into task-level subproblems. We further compare the peak memory usage of different methods as the training-set size increases. The results are reported in Table 25.

Table 25. Peak GPU memory of OT-based selection as the training set grows. LAVA, SAVA, and TW-LAVA run their original pipelines with 10,000 validation images (TW-LAVA with tasks of 10K samples). For CAVA we report its valuation step, without feature extraction, with the 2,000 validation images of the paper and with 10,000 for a direct comparison. OOM: out of memory; "–": not measured.

				CAVA (ours)	
Training set	**LAVA**	**SAVA**	**TW-LAVA**	$N_v = 2{,}000$	$N_v = 10{,}000$
10K	21.65 GB	3.00 GB	21.65 GB	**0.10 GB**	0.35 GB
20K	43.00 GB	3.00 GB	21.65 GB	**0.17 GB**	0.65 GB
30K	OOM	3.00 GB	21.65 GB	**0.23 GB**	0.95 GB
40K	OOM	3.00 GB	21.65 GB	**0.30 GB**	1.25 GB
50K	OOM	3.00 GB	21.65 GB	**0.35 GB**	1.55 GB
1M	OOM	–	–	**0.42 GB**	1.86 GB

Table 25 shows that LAVA becomes increasingly memory-intensive as the training-set size grows. In particular, its memory usage reaches 43.00 GB at 20K training samples and runs out of memory from 30K samples onward in our environment. This behavior is expected because full-set OT requires constructing and optimizing over a large transport problem whose size grows with the number of training and validation samples.

SAVA, by contrast, maintains low memory usage of approximately 3.00 GB because it performs OT at the batch level. However, this memory advantage comes with substantially longer runtime in our experiment, as shown in Table 24. Moreover, batch-wise selection may restrict the comparison scope of OT to small local batches.

TW-LAVA, the task-wise transfer of LAVA, avoids the out-of-memory failure by fixing the task size, but its original pipeline still needs about 21.65 GB. CAVA needs far less: its complementary label cost is a lookup table, so the valuation step stores only the task-by-validation cost matrix and a few vectors. At the paper protocol of 2,000 validation images it uses 0.10 to 0.35 GB for 10K to 50K training samples, and 0.42 GB for one million samples, where pools larger than 12,000 samples are scored in chunks; even with 10,000 validation images it stays below 2 GB. Memory therefore grows with the validation set and the chunk size, and with the training set only until the pools are scored in chunks, and the task-wise subproblems remain independent, so they can be run in parallel across devices without changing the method.

E. Visualizing the Label Margin

Figures 17 and 18 show how CAVA decides on individual training samples, in the style of the retrieval visualization of Long et al. [26]. Each figure shows eight samples drawn at random from the first task, half of them annotation errors ($\bar{y}_i = y_i$) and half valid complementary labels, with the OT solution of the main experiments. Each row gives the image (red border: annotation error), the labels of the six validation images with the largest transport weight from this sample, with these weights (green: the true class), the label potentials $f_i(c)$ relative to their minimum for the six cheapest classes (the given label $\bar{y}_i$ in red), the margins of all samples in the task with this sample marked and the threshold $s_i = 0$ of the adaptive cut drawn as a dashed line, and the decision of the adaptive cut.

The rows make the mechanism of Theorem 2 visible. The validation images that receive the transport mass of a sample almost always carry its true class. For an annotation error, the given label is that class, so $f_i(\bar{y}_i)$ is the smallest potential, the

margin $s_i = f_i(\bar{y}_i) - \min_{c \neq \bar{y}_i} f_i(c)$ is negative, and the sample is flagged. For a valid complementary label, the given label names another class, its potential lies above the minimum, and the sample is kept. The examples are not selected for success: on Places365-10 (Figure 18), one annotation error labeled "bridge" shows mostly a forest stream, its transport mass goes to forest-path images, and the adaptive cut keeps it. Such failures are rare, as the margin histograms show: the two groups overlap only in a thin band around the threshold $s_i = 0$.

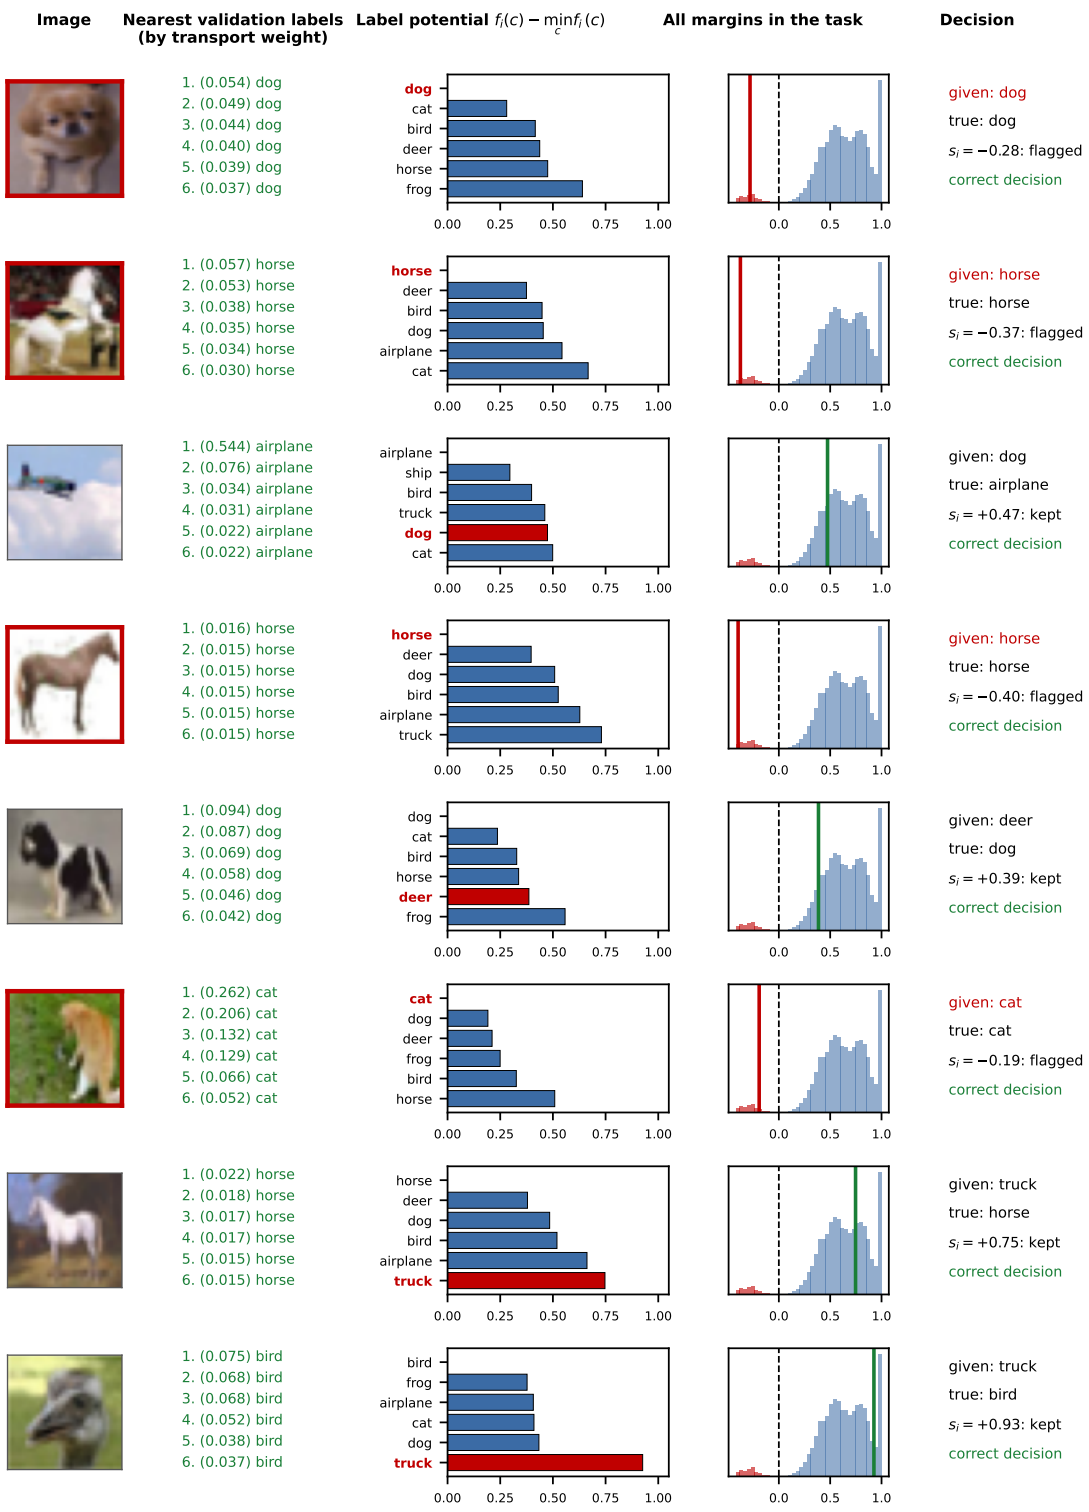

Figure 17. Label-margin visualization for randomly drawn samples from CLCIFAR-10 (human complementary labels, 3.98% annotation errors; first task, 2,000 validation images). Columns: image (red border: annotation error), labels of the validation images with the largest transport weight, with the weights (green: true class), label potentials $f_i(c) - \min_c f_i(c)$ with the given label in red, margins of all samples in the task (blue: valid, red: annotation error; dashed: adaptive-cut threshold $s_i = 0$) with this sample marked, and the decision of the adaptive cut.

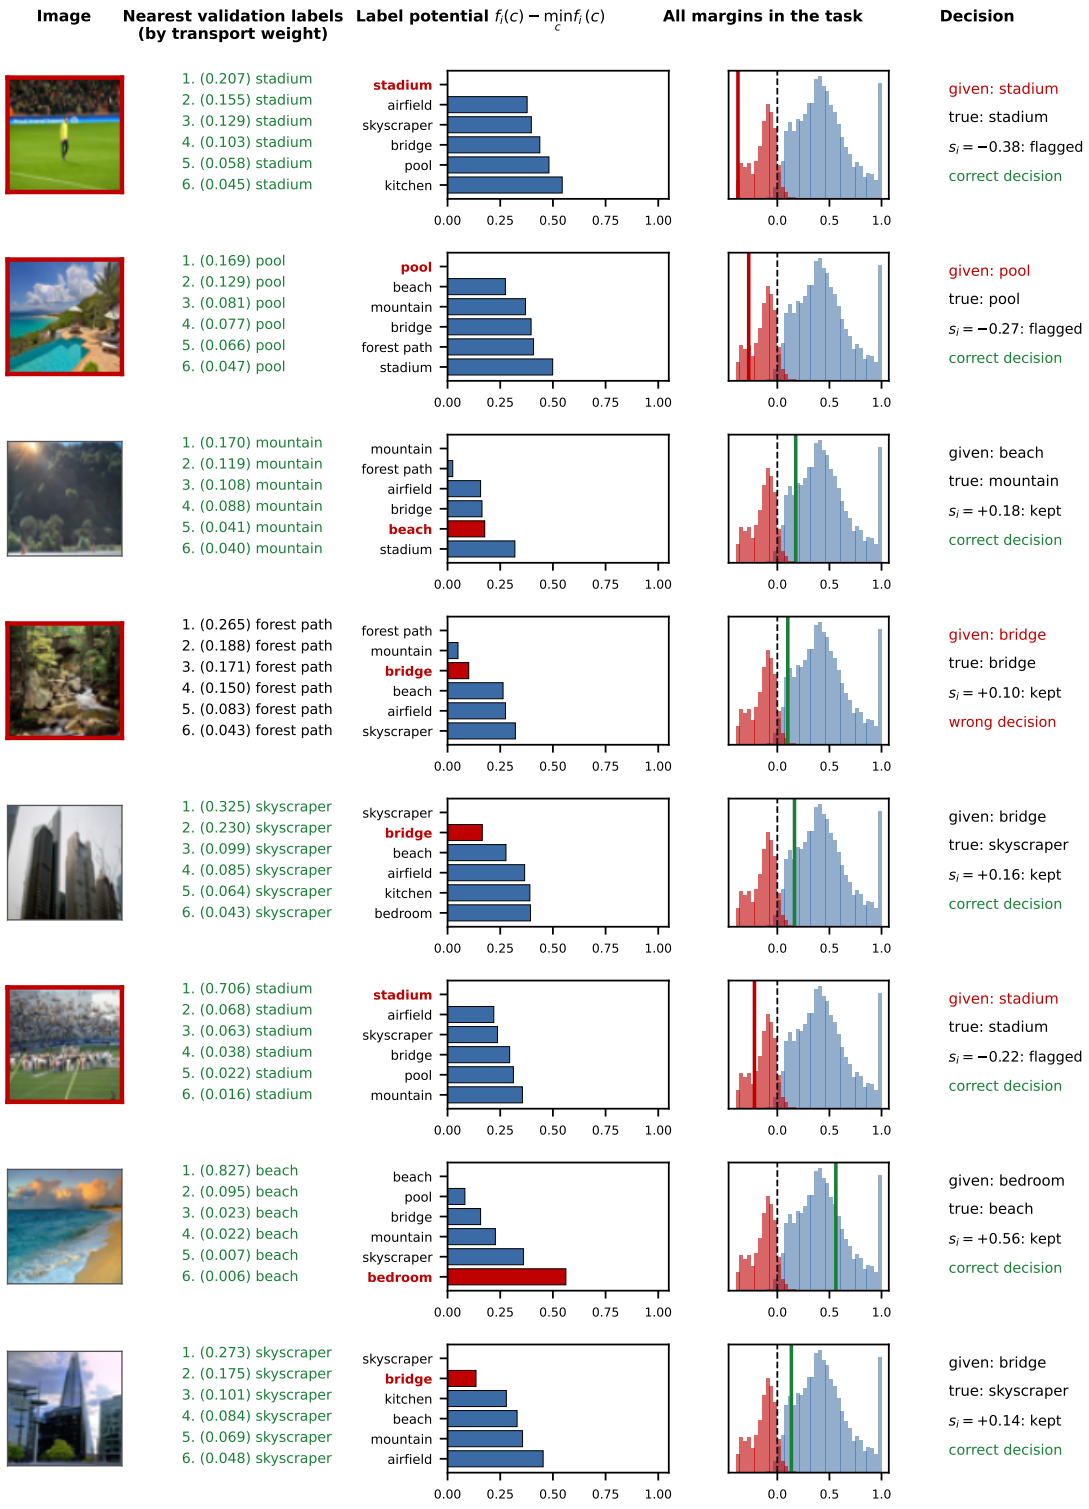

Figure 18. Label-margin visualization for randomly drawn samples from Places365-10, ten scene classes of Places365 [41] with synthetic diagonal noise ($\eta = 25\%$; first task, 2,000 validation images). Columns as in Figure 17. The fourth row is a failure: an annotation error whose image resembles a forest path is kept.