arXiv:2610.29495v1 [cs.CV] 21 Oct 2026

Balancing Foundation Models and Trait-Specific Learning for Forehead-Crease Recognition

Geetanjali Sharma[1], Aryan Jain[1], Dev Nagaich[1], Aditya Nigam[1], and Raghavendra Ramachandra[2]

[1] Indian Institute of Technology Mandi, India
[2] SAFE Center, Norwegian University of Science and Technology (NTNU), Gjøvik, Norway

Abstract. Biometric recognition often degrades under partial facial visibility, where conventional facial cues are occluded. The forehead region remains consistently visible and contains distinctive crease patterns that can serve as a complementary biometric trait. We present a unified evaluation of forehead-crease recognition across explainable zero-shot reasoning, domain-adaptive learning, and trait-specific modeling. FIRE enables interpretable zero-shot verification using multimodal large language models, while pretrained face-recognition models are adapted to the forehead domain through margin-based learning. Under open-set evaluation, the best zero-shot model GPT-4o+FIRE achieves 18.22% EER, while Edge-Forehead achieves 8.30% EER and 54.01% TMR@0.1% FMR. The results show that although MLLMs offer interpretable zero-shot verification, trait-specific learning remains more effective for robust forehead-crease biometric recognition.

Keywords: Trait-Specific · Forehead creases · Recognition.

1 Introduction

Biometric verification systems [2,6] have achieved remarkable progress in recent years, largely driven by deep learning, supervised approaches. These methods leverage large-scale annotated datasets and powerful representation learning techniques [14,15], enabling high recognition accuracy under controlled conditions. As a result, such systems are widely deployed in applications including access control, surveillance, and identity verification. However, their performance degrades significantly in unconstrained environments, particularly when faces are partially occluded due to masks, scarves, pose variations, or low-resolution imaging. In these scenarios, a substantial portion of the facial region becomes unavailable, limiting the effectiveness of holistic face-based representations. To address this limitation, recent studies have explored partial facial regions and soft biometric traits that remain visible under occlusion [2,15]. Among these, the forehead is particularly promising due to its consistent visibility and rich structural characteristics. Forehead creases provide discriminative identity cues,

with prior studies demonstrating verification across images acquired in two sessions; however, their temporal stability remains insufficiently quantified [2,15]. Furthermore, fine-grained skin texture and wrinkle patterns retain identity-discriminative information even in the absence of global facial cues.

Despite these advantages, existing forehead-crease recognition methods face several limitations. Most approaches rely on task specific architectures and do not fully exploit transferable representations learned by modern face recognition models. Although foundation models(FMs) and multimodal large language models (MLLMs) have demonstrated strong Zero-Shot(ZS) capabilities, they are not optimized for subtle biometric cues such as crease topology, spacing, and texture continuity. In addition, prior studies often compare ZS and supervised approaches under different protocols, limiting fair evaluation. Finally, explainability in forehead-based biometric recognition remains largely underexplored despite its importance for transparency and real-world deployment. In parallel, modern face recognition models have been shown to learn hierarchical representations that capture not only global facial structure but also mid-level and fine-grained texture patterns [14]. This observation motivates an important research question: *can pretrained face recognition models be effectively adapted to operate on spatially constrained regions such as the forehead?* Recent advances in margin-based metric learning, including AdaFace [7], CosFace [19], and related formulations, provide a promising direction by enforcing strong intra-class compactness and inter-class separability, even under subtle appearance variations. These properties make them particularly suitable for modelling fine-scale biometric traits such as forehead creases. This work systematically compares FMs with domain-adaptive and trait-specific approaches for trustworthy and interpretable biometric verification.

Motivated by these insights, we propose a unified framework for explainable and domain-adaptive forehead-crease biometric recognition that integrates three complementary paradigms within a single system. Unlike prior work that treats these directions independently, our approach combines: **(i)** explainable zero-shot reasoning using multimodal models, **(ii)** domain adaptation of pretrained face recognition networks using margin-based metric learning, and **(iii)** a trait-specific model designed to explicitly capture fine-grained crease structures. This unified formulation enables a comprehensive understanding of the trade-off between interpretability, generalisation, and discriminative performance. We evaluate the proposed framework under open-set verification protocols on real-world and synthetic forehead datasets. To ensure a fair and rigorous comparison, we benchmark multiple paradigms, including multimodal large language models, vision foundation models, domain-adapted face recognition systems, and specialised forehead-crease architectures. In addition, we conduct detailed ablation studies to analyse the impact of structured prompting, margin-based supervision, and training strategies on overall performance. From an application perspective, the proposed system is particularly relevant in scenarios where full-face information is unavailable or unreliable, such as masked authentication, surveillance in unconstrained environments, and low-resolution or partial-face recognition. In

such cases, forehead-crease biometrics can act as a complementary modality, improving robustness without requiring additional sensors or intrusive acquisition processes.

Contributions. The main contributions of this work are: (i) a unified framework integrating explainable zero-shot reasoning, domain-adaptive learning, and trait-specific modeling for forehead-crease recognition; (ii) adaptation of pretrained face-recognition models to forehead biometrics through margin-based metric learning; (iii) FIRE, a structured framework for interpretable multimodal verification; (iv) a balanced open-set evaluation of zero-shot, domain-adaptive, and trait-specific approaches; and (v) an assessment of forehead-crease biometrics for partial-face recognition scenarios. The remainder of this paper is structured as follows. Section 2 reviews related work. Section 3 presents the proposed methodology. Section 4 describes the datasets and evaluation protocols. Section 5 presents the results and discussion, followed by additional ablation studies and prompt analyses are provided in the supplementary material. Finally, Section 6 concludes the paper.

2 Related work

In this section, we review existing literature on forehead-based biometrics, face recognition and domain adaptation, and explainable verification using multimodal models. We also highlight key limitations and position our contribution within these research directions.

2.1 Forehead-Based Biometric Recognition

Forehead-based biometrics has emerged as a promising complementary modality for identity verification, particularly in scenarios where the lower facial region is occluded. Early approaches relied on handcrafted features, such as texture descriptors and edge-based representations, to model wrinkle patterns and skin structure. While these methods demonstrated that forehead creases contain identity-discriminative information, they were limited in handling variations in illumination, facial expression, and image resolution.

With the advancement of deep learning, recent works have shifted towards data-driven approaches for modelling forehead characteristics. Methods such as FHCVS [2] and FH-SSTNet [15] introduced specialised architectures designed to capture the spatial arrangement and structural properties of forehead creases. These models explicitly learn representations based on line patterns, curvature, and local texture variations, achieving improved performance over handcrafted pipelines. Similarly, convolutional neural network–based approaches trained directly on forehead regions have demonstrated the ability to extract discriminative features under controlled settings [2,15]. To address the limited availability of annotated datasets, synthetic data generation has also been explored. Synthetic forehead datasets attempt to model variations in crease patterns, shape, and illumination, enabling controlled experimentation and data augmentation.

However, models trained on synthetic data often suffer from domain shift when applied to real-world images, resulting in degraded performance.

Despite these advancements, several limitations remain. Most existing approaches are highly task-specific and do not leverage transferable representations from large-scale face recognition models. In addition, evaluations are often conducted on relatively small datasets, limiting insights into generalisation. Furthermore, comparisons across methods are typically restricted, making it difficult to establish a comprehensive understanding of performance across different paradigms. These challenges motivate the need for a more unified and scalable approach that combines domain-specific modelling with transferable feature learning.

2.2 Face Recognition and Domain Adaptation

Face recognition has achieved significant success through deep metric learning frameworks trained on large-scale datasets [14]. Models such as FaceNet learn embedding spaces where identity-related features are well separated, enabling robust verification across varying conditions. A key strength of these models lies in their ability to learn hierarchical representations that capture both global facial structure and fine-grained texture information. This property suggests that pretrained face recognition models may be transferable to localised regions such as the forehead. Recent studies have explored domain adaptation strategies, including fine-tuning pretrained networks and employing margin-based loss functions such as AdaFace [7] and CosFace [19]. These approaches improve feature discrimination by enforcing intra-class compactness and inter-class separability, particularly under subtle appearance variations. However, forehead-crease biometrics remain largely underexplored. Existing methods rarely investigate the transferability of pretrained face-recognition representations to this fine-grained and spatially constrained trait, leaving the effectiveness of domain adaptation for forehead-crease verification an open question.

2.3 Multimodal Large Language Models

Recent advances in multimodal learning have led to the development of powerful vision-language models that jointly process visual and textual information. These models combine visual encoders with large language models, enabling them to interpret images and generate descriptive or reasoning-based outputs. They have demonstrated strong performance across tasks such as visual question answering, image captioning, and visual reasoning. A key advantage of these models is their ability to perform zero-shot inference, allowing them to generalise to new tasks without task-specific training. This makes them particularly appealing for biometric applications where labelled data is limited, such as forehead-crease recognition. Moreover, multimodal models offer a unique advantage in terms of explainability. Unlike traditional biometric systems that provide only similarity scores, these models can generate human-interpretable justifications based on visual evidence, enhancing transparency and trust. Despite these benefits, multimodal models face several limitations in fine-grained biometric tasks. Forehead-

crease recognition requires precise analysis of subtle texture patterns and structural variations, which general-purpose models are not explicitly designed to capture. Their performance is sensitive to variations in illumination, resolution, and expression, and they may produce inconsistent predictions or hallucinated reasoning when differences between samples are subtle. FMs and MLLMs are pretrained on large-scale multimodal corpora and exhibit strong zero-shot and few-shot generalization capabilities. These properties make them attractive for biometric applications, where acquiring large annotated datasets is often difficult. Unlike conventional biometric systems that require task-specific training, MLLMs can perform verification through prompt-driven reasoning and provide human-interpretable explanations. However, biometric recognition operates under stringent security constraints, requiring extremely low EER and high TMR at low FMR operating points. Generic foundation models are optimized for broad visual-semantic understanding rather than fine-grained biometric discrimination. Furthermore, large-scale deployment introduces practical challenges related to computational cost, inference latency, and scalability. These limitations motivate a systematic comparison between generic foundation models, domain-adaptive learning, and trait-specific biometric models.

2.4 Limitations of Existing Work

Despite recent progress, forehead-crease biometrics remain underexplored, with limited use of pretrained face-recognition models, inconsistent evaluation protocols, and insufficient study of explainability and domain shift. To address these gaps, we systematically evaluate foundation, domain-adaptive face-recognition, and trait-specific models under a unified protocol, comparing their interpretability, transferability, and verification performance.

3 Proposed Method

In this section, we present the proposed unified framework for forehead-crease–based biometric recognition. The goal is to design a reliable verification system that operates solely on the forehead region, particularly in scenarios where full-face information is unavailable due to occlusion or masking. Unlike conventional approaches that rely on a single paradigm, our framework integrates three complementary components: **(i)** an explainable **Zero-Shot** verification module using multimodal reasoning, **(ii)** a **Domain-adaptive** learning module that transfers knowledge from pretrained face recognition models, and **(iii)** a **Trait-Specific** model designed to capture fine-grained crease structures. The overall pipeline is illustrated in Figure 1.

3.1 Problem Formulation

Given a pair of forehead images $(\mathbf{x}_i, \mathbf{x}_j)$, the objective is to determine whether they belong to the same identity. This is formulated as a binary verification task:

$$y_{ij} = \begin{cases} 1, & \text{if } \mathbf{x}_i \text{ and } \mathbf{x}_j \text{ belong to the same subject,} \\ 0, & \text{otherwise.} \end{cases} \tag{1}$$

The system produces a similarity score $s_{ij} \in [0, 1]$. For explainable models, it additionally outputs a structured textual explanation describing the reasoning behind the decision.

3.2 Forehead Region Extraction and Preprocessing

Face images are first processed to extract the forehead region using a fixed cropping strategy, ensuring spatial consistency across samples. As forehead creases are predominantly structural and textural cues, the cropped regions are converted to grayscale to reduce appearance variability. The resulting images are replicated across three channels and resized to 224×224 for compatibility with standard CNN architectures. During training, horizontal flipping, small rotations, and brightness perturbations are applied to improve generalization while preserving discriminative crease patterns.

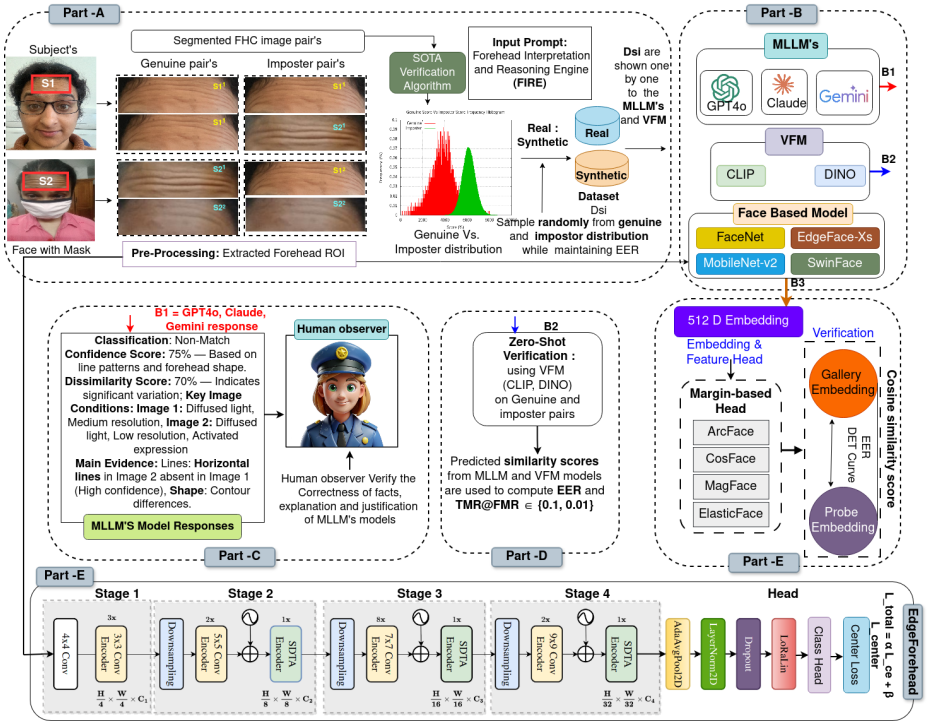

Fig. 1. Forehead-crease verification pipeline using real and synthetic image pairs evaluated with MLLMs, FMs and trait-specific models.

3.3 Explainable Zero-Shot Verification via FIRE

To enable interpretable biometric verification, we introduce the **Forehead Interpretation and Reasoning Engine (FIRE**(detailed prompt in supplemen-

tary material), a structured prompting framework that guides multimodal large language models (MLLMs) such as Claude [1], GPT-4o [10], and Gemini [18]. Unlike supervised models, these systems are not explicitly trained for biometric verification. Therefore, we formulate the task as a structured reasoning problem using Chain-of-Thought (CoT) prompting. Each input consists of a pair of forehead images along with a carefully designed prompt that encourages step-by-step comparison. The **FIRE** framework decomposes the verification process into multiple stages, including: **(i)** assessment of image conditions (lighting, occlusion, resolution), **(ii)** extraction of structural features (crease patterns, spacing, orientation), **(iii)** comparison of primary and secondary features, **(iv)** decision-making with confidence and dissimilarity estimation. This structured reasoning reduces hallucinations and improves consistency compared to generic prompts. The model outputs both a classification (match/non-match) and a human-readable explanation, enabling transparent and interpretable verification. The complete FIRE (detailed prompt provided in the supplementary).

3.4 Zero-Shot Verification using Vision Foundation Models

To further analyse zero-shot generalisation, we evaluate vision foundation models (VFMs), including DINOv2 [11] and CLIP [13], without task-specific fine-tuning. Given an input forehead image $\mathbf{x}_i$, an embedding $\mathbf{z}_i$ is extracted using a pretrained model. Verification is performed using cosine similarity:

$$s_{ij} = \frac{\mathbf{z}_i \cdot \mathbf{z}_j}{\|\mathbf{z}_i\|\|\mathbf{z}_j\|} \tag{2}$$

DINOv2 captures fine-grained structural features due to self-supervised training, while CLIP benefits from semantic alignment learned from image-text pairs. These models serve as strong zero-shot baselines for evaluating generalisation beyond traditional face recognition cues. However, since these models are not trained specifically for forehead biometrics, their performance is sensitive to variations in illumination, expression, and domain shift, motivating the need for domain adaptation.

3.5 Domain-Adaptive Learning using Pretrained Models

To address the limitations of zero-shot models, we adopt a domain-adaptive learning strategy that transfers knowledge from pretrained face recognition models to the forehead domain.

Feature Representation Let $f(\cdot)$ denote a pretrained face recognition network. For an input image $\mathbf{x}_i$, the model produces an embedding: $\mathbf{z}_i = f(\mathbf{x}_i), \quad \mathbf{z}_i \in \mathbb{R}^{512}$ These embeddings capture identity-related information, which is further refined for forehead-specific features through fine-tuning. We employ multiple backbone architectures, including MobileFaceNet-V2 [4], EdgeFace-Xs [5], SwinFace [12], and FaceNet [14].

Margin-Based Metric Learning To enhance discriminability, we employ margin-based loss functions that enforce compact intra-class distributions and well-separated inter-class boundaries.

We further evaluate several margin-based objectives, including **ElasticFace** (stochastic margin regularization) [3], **AdaFace** [7] (quality-adaptive margin), **MagFace** [9] (magnitude-aware representation learning) , and **CosFace** [19] (additive cosine margin) . These objectives enhance intra-class compactness and inter-class separability, enabling more discriminative modeling of fine-grained forehead-crease patterns for robust verification.

Forehead Interpretation and Reasoning Engine (FIRE)

Input: Two forehead images. **Output:** Match/Non-Match, Confidence Score, Dissimilarity Score, and Explanation. **Step 1:** Assess image quality (lighting, occlusion, resolution, expression). **Step 2:** Extract and compare forehead-crease characteristics, including topology, spacing, orientation, continuity, forehead shape, and permanent marks. **Step 3:** Normalize illumination, scale, occlusion, and expression variations. **Step 4:** Evaluate feature consistency, assign confidence weights, and estimate similarity. **Step 5:** Generate the final decision with confidence, dissimilarity score, and reasoning. **Constraint:** Use only observable forehead features and explicitly report uncertainty when evidence is insufficient.

3.6 Trait-Specific Model: Edge-Forehead

While domain-adaptive models improve performance, they are still derived from general face representations. To further enhance fine-grained discrimination, we introduce a trait-specific model, **Edge-Forehead**, tailored for forehead-crease recognition. We fine-tune the EdgeFace model [5] on forehead images using a gradual unfreezing strategy, allowing the network to retain general facial features while progressively specialising in crease-specific patterns. To improve embedding quality, we employ a dual-loss objective:

$$\mathcal{L}_{\text{total}} = \alpha \mathcal{L}_{\text{ce}} + \beta \mathcal{L}_{\text{center}} \tag{3}$$

where $\mathcal{L}_{\text{ce}}$ enforces inter-class separability and $\mathcal{L}_{\text{center}}$ reduces intra-class variance. The weighting parameters α and β balance classification accuracy and feature compactness. By focusing exclusively on forehead regions, the model(Figure1, Part E) learns fine-grained structural characteristics such as crease continuity, spacing, curvature, and texture patterns. This leads to improved robustness under variations in illumination, pose, and expression. Finally, all models are evaluated under a unified verification protocol, enabling a fair comparison between zero-shot, domain-adaptive, and trait-specific approaches.

4 Dataset & Evaluation Protocol

In this section, we describe the datasets and evaluation protocol used to assess the proposed framework across zero-shot, domain-adaptive, and trait-specific models. Forehead crease patterns are considered as the primary biometric trait. Experiments are conducted on both real and synthetic datasets to ensure comprehensive evaluation under diverse conditions. We design a unified protocol to

enable fair comparison across MLLMs, VFMs, and supervised forehead-crease recognition systems.

4.1 Datasets for Zero-Shot Evaluation (LLMs & VFMs)

We evaluate zero-shot verification on two benchmark datasets subsets: a real dataset (FH-V2) and a synthetic dataset [16]. The synthetic dataset was obtained from [16] and used to evaluate forehead verification under synthetic image conditions. Masked face images are first processed to extract the forehead region, as shown in Figure 1. To construct a reliable benchmark, a pretrained SOTA verification model [16] is used to compute similarity scores and select representative genuine and impostor pairs. Specifically, 50 genuine pairs from 50 unique subjects and 450 impostor pairs from 250 distinct subjects are sampled from both datasets. This high negative-to-positive ratio follows standard biometric evaluation protocols, ensuring robust score separation and stable performance estimation. The selected pairs capture variations in illumination, expression, and image quality, enabling a realistic assessment of MLLMs and VFMs for forehead-crease-based verification.

4.2 Datasets for Domain-Adaptive & Trait-Specific Models

For supervised training, we use the FH-V1 dataset [2], obtained by extracting forehead regions from face images using a fixed YOLOv6[8] based cropping strategy. The dataset comprises 247 subjects with 2,462 gallery images and 2,502 probe images, totaling 4,964 images and variations in illumination, pose, expression, and image quality. Subjects are treated as unique identity classes, and a subject-disjoint train-validation split is employed to prevent identity leakage and support model selection. For open-set evaluation, all models are tested on FH-V2 [16], comprising 295 unseen subjects and 29,500 genuine and approximately 8.67 million impostor verification pairs.

4.3 Training Protocols

Domain-adaptive and trait-specific models are trained on FH-V1 using a stratified 90%/10% train-validation split. Pretrained face-recognition backbones are fine-tuned using margin-based objectives (CosFace, AdaFace, MagFace, and ElasticFace). Edge-Forehead is initialized with EdgeFace-Base weights and trained in three stages: (i) cross-entropy with label smoothing ($\epsilon = 0.1$) and progressive unfreezing (70%,→,0%), (ii) joint cross-entropy and center loss ($\lambda_{center} = 0.3$, $\alpha = 0.5$) with augmentation (flip, $\pm 9^\circ$ rotation, 0.8–1.2× scaling, grayscale), and (iii) semi-hard triplet mining with a dynamic margin (0.4–1.5). Stages 1–2 use SGD (lr=10^{-3}, momentum=0.9, weight decay=5×10^{-4}, batch size=32), while Stage 3 uses Adam (lr=10^{-5}, batch size=2). Early stopping (patience=3) is applied throughout training. Images are resized to 112×112. Verification uses Euclidean distance between 512-D embeddings, with thresholds set at EER. We use EdgeFace-Base (18.23M parameters) with dropout ($p = 0.5$) and a 247-class

FC head. Training uses three stages: progressive unfreezing (70%, 50%, 30%, 10%, and 0% frozen; up to 15 epochs per setting), Stage 2 (up to 30 epochs), and Stage 3 (until convergence). The center-update rate and center-loss weight are $\alpha = 0.5$ and $\beta = 0.3$, respectively.

4.4 Evaluation Protocol

To comprehensively evaluate the proposed framework, we consider three categories of models: (i) Multimodal LLMs: GPT-4o, Claude Sonnet 4, and Gemini 2.5 Pro. (ii) Vision foundation models: DINOv2 and CLIP (iii) Supervised models: Edge-Forehead (proposed), and FHCVS. All models are evaluated under a unified open-set verification protocol on both real (FH-V2) and synthetic datasets. For **zero-shot models** (MLLMs and VFMs), each image pair undergoes three independent stateless inference runs to reduce output stochasticity. Final decisions are obtained via majority voting, with similarity/confidence scores averaged across runs. When refusals or conflicting predictions occur, the highest-confidence prediction is retained. For **domain-adaptive** and **trait-specific** models, verification is performed using cosine similarity in the learned embedding space. Scores are computed and aggregated uniformly across models to ensure fair comparison. Performance is evaluated using standard biometric metrics, including Equal Error Rate (EER) and True Match Rate (TMR) at fixed False Match Rate (FMR) operating points of 0.1 and 0.01. These metrics provide a consistent framework for comparing zero-shot, domain-adaptive, and trait-specific approaches under varying operating conditions.

Reproducibility Details: We use a subject-disjoint 90%/10% train validation split(seed=42). Verification pairs follow Section 4.1. Images are resized to 112×112, normalized, and augmented. Edge-Forehead is initialized from pretrained EdgeFace-Base and trained using the three-stage schedule in Section 4.3. FIRE prompts, model versions, and three-run stateless inference settings are fixed. Verification thresholds are selected at EER using Euclidean distances between 512-D embeddings. It has approximately 18.23M parameters and requires 1,398.83 MFLOPs per forward pass.

5 Results & Discussion

This section evaluates whether generic foundation models (GPT-4o, Gemini 2.5 Pro , Claude Sonnet 4) can provide competitive verification performance to domain-adaptive and trait-specific biometric systems. While foundation models offer zero-shot generalization and explainability, biometric verification demands highly discriminative representations and strict operating characteristics. The experiments analyse this trade-off across all three paradigms under a unified protocol.

5.1 FIRE-based LLM and VFM Results

We first evaluate the performance of FIRE-based multimodal LLMs and vision FMs on both real and synthetic datasets for zero-shot verification. The

Table 1. Identity verification results on the sampled FH-V2 verification subset (50 genuine and 450 impostor pairs) and synthetic subset. [16].

Model	EER (%) ↓	TMR@FMR =0.1 (%) ↑	TMR@FMR =0.01 (%) ↑
FH-V2(Real			
Edge-Forehead(Our)	**8.63**	**70.00**	**70.00**
FH-SSTNet[15]	14.22	38.00	38.00
FHCVS [16]	21.56	32.00	32.00
CLIP [13]	24.11	14.00	6.00
DINO-V2 [11]	24.00	40.00	38.00
Claude Sonnet 4[1]	27.11	24.00	24.00
Gemini 2.5 Pro [18]	24.89	26.00	26.00
GPT-4o [10]	18.22	52.00	52.00
Synthetic			
Edge-Forehead(Our)	**0.25**	**94.00**	**94.00**
FH-SSTNet[15]	29.74	5.88	3.92
FHCVS [16]	24.22	22.00	16.00
CLIP [13]	31.56	10.00	4.00
DINO-V2 [11]	28.00	12.00	12.00
Claude Sonnet 4[1]	31.89	18.00	18.00
Gemini 2.5 Pro[18]	30.89	22.00	10.00
GPT-4o [10]	25.44	20.00	12.00

quantitative results are summarised in Table 1 and explainability examples in supplementary material Figure 3.

Quantitative Analysis For reference, the trait-specific Edge-Forehead model achieves the best performance on the sampled FH-V2 evaluation subset (EER = 8.63%), while its full open-set evaluation on FH-V2 is reported separately in Section 5.2. Results on the complete FH-V2 open-set evaluation are reported separately in Table 3, where Edge-Forehead attains an EER of 8.30% with a True Match Rate (TMR) of 70.00% at both operating points. This represents a significant improvement over existing supervised approaches such as FHCVS (EER: 21.56%) and FH-SSTNet (EER: 14.22%). Among zero-shot methods, GPT-4o achieves the strongest performance (EER: 18.22%, TMR: 52.00%), indicating that structured reasoning enables effective capture of identity-related cues even without task-specific training. DINOv2 also demonstrates competitive performance (TMR: 40.00%, 38.00%), suggesting that self-supervised representations retain sensitivity to fine-grained structural patterns. In contrast, CLIP exhibits relatively weaker performance, likely due to its emphasis on semantic alignment rather than fine texture discrimination. On synthetic data, most zero-shot models degrade under domain shift, whereas Edge-Forehead maintains strong performance (EER: 0.25%; TMR: 94.00%). GPT-4o performs best among zero-shot models (EER: 25.44%) but struggles with synthetic variations. Overall, trait-specific learning improves verification accuracy and robustness, while zero-shot models offer interpretability.

Qualitative Analysis Table 2 presents a three-round response analysis to evaluate the consistency and reliability of FIRE-based verification. Several key be-

havioural patterns can be observed: (i) **Progressive Refinement**, where predictions often improve across rounds; (ii) **High-Confidence Stability**, where clear samples yield consistent decisions; (iii) **Recovery from Refusal**, where initial refusals may be resolved through subsequent prompting; and (iv) **Response Variability**, where challenging cases can produce conflicting predictions due to model sensitivity and stochasticity. Additional prompt-ablation and stability analyses are provided in the supplementary material.

Scenario	Round 1	Round 2	Round 3
Improved classification results and confidence	CLS: Non-Match Conf: 55% Dis Score: 65%	CLS: Match Conf: **70%** Dis Score: 15%	CLS: Match Conf: **85%** Dis Score: 20%
Improved match consistency with lower dissimilarity scores	CLS: Match Conf: 85% Dis Score: **15%**	CLS: Match Conf: 85% Dis Score: **15%**	CLS: Match Conf: 85% Dis Score: **10%**
Improved forehead verification after iterative prompting	Unable to perform forensic analysis	CLS: NM Conf: 60% Dis Score: 75%	CLS: NM Conf: 75% Dis Score: 80%
Ambiguity due to conflicting or missing classifications	CLS: Match Conf: 75% Dis Score: 25%	Unable to perform forensic analysis or classify	CLS: NM Conf: 70% Dis Score: 60%

Table 2. Round-wise response comparison across three inference passes. CLS: classification; Conf: confidence; Dis: dissimilarity; M: match; NM: non-match.

Explainability Analysis make it crisp but keep important information A key advantage of FIRE framework is its ability to provide human-interpretable explanations alongside verification decisions. Representative examples of genuine and impostor pairs are presented in the supplementary material with model-generated reasoning. For impostor pairs, FIRE correctly identifies differences in crease alignment and spacing despite illumination and resolution variations, achieving up to 85% confidence. For genuine pairs, it recognizes consistent crease patterns, producing high-confidence matches (95%) with low dissimilarity. These results demonstrate that FIRE enhances interpretability while generally aligning with human reasoning, although occasional errors remain under challenging visual conditions.

5.2 Domain-Adaptive & Trait specific Forehead-Crease Verification

Quantitative Results Table 3 reports open-set verification performance for domain-adaptive models trained on FH-V1 and evaluated on FH-V2. The results reveal consistent trends across backbone architectures and margin-based loss functions. For the FaceNet backbone, margin-based supervision significantly improves performance compared to standard formulations. ElasticFace achieves the best results (EER: 11.17%), benefiting from stochastic margin sampling that improves generalisation under limited identity evidence. MobileFaceNet-V2 demonstrates the strongest overall performance across all architectures. In particular, MobileFaceNet-V2 combined with CosFace achieves the lowest EER of 9.31% and the highest TMR values. This can be attributed to its lightweight convolutional design, which effectively captures local texture patterns critical for forehead-based recognition. SwinFace exhibits comparatively lower performance, with AdaFace achieving the best results within this group (EER: 13.44%). The

Table 3. Open-set verification performance of foundation, domain-adaptive, face-recognition, and trait-specific models on FH-V2 after fine-tuning on FH-V1.

Model	EER (%)	TMR (%) @ FMR (%) =	
		0.1 (%)	0.01 (%)
FaceNet[14] + Arcface	20.65	11.72	2.49
FaceNet[14] + AdaFace	15.38	26.13	8.07
FaceNet[14] + CosFace	17.35	17.51	4.14
FaceNet[14] + MagFace	26.73	7.33	1.56
FaceNet[14] + ElasticFace	**11.17**	**45.48**	**27.87**
MFN-V2[4] + Arcface	10.73	55.76	41.04
MFN-V2[4] + AdaFace	13.68	46.35	32.32
MFN-V2[4] + CosFace	**9.31**	**60.26**	**44.59**
MFN-V2[4] + MagFace	11.51	51.00	32.22
MFN-V2[4] + ElasticFace	15.37	41.91	26.33
SwinFace[12] + Arcface	17.72	28.61	**10.12**
SwinFace[12] + AdaFace	**13.44**	**31.64**	9.43
SwinFace[12] + CosFace	21.33	21.50	7.55
SwinFace[12] + MagFace	22.13	6.80	1.93
SwinFace[12] + ElasticFace	22.13	14.80	4.82
EF-Xs[5] + Arcface	21.54	24.14	7.73
EF-Xs[5] + AdaFace	12.72	49.01	33.83
EF-Xs[5] + CosFace	19.46	27.47	11.52
EF-Xs[5] + MagFace	**11.01**	**52.14**	**37.35**
EF-Xs[5] + ElasticFace	22.81	22.12	6.43
DINO-V2 (ViT-b)[11]	8.83	**57.16**	**40.39**
CLIP(ViT-b)[13]	10.72	44.25	25.29
FHCVS[17]	12.37	47.61	35.93
ResNet 18 + CAM[2]	14.61	47.61	35.93
EdgeFace-Base[5]	9.84	51.95	35.12
Edgeface-Base[5](inference)	26.89	8.50	3.00
Edgeface-Base[5] + CosFace	18.15	38.97	25.19
Edgeface-Base[5] + Adaface	20.95	36.75	22.80
EdgeForehead(Our)	**8.30**	54.01	37.08

reduced performance suggests that global attention mechanisms may be less effective for modelling highly localised crease structures. For EdgeFace-Xs, MagFace achieves the best performance (EER: 11.01%), highlighting the importance of quality-aware feature learning in handling variations in image quality and crease visibility. These findings indicate that both backbone architecture and margin-based supervision play a crucial role in capturing fine-grained biometric features.

Trait-Specific Forehead-Crease Verification: Table 3 compares Edge-Forehead with existing forehead-recognition approaches. Edge-Forehead achieves the lowest EER (8.30%) and competitive TMRs, surpassing Vision Foundation Models (DINOv2, CLIP). The results demonstrate that explicit modeling of fine-grained crease structures yields superior discriminability and robustness, underscoring the effectiveness of trait-specific learning for forehead-crease verification. EdgeFace-Base is fine-tuned on FH-V1, while EdgeFace-Base (inference) uses the original pretrained model without fine-tuning. Figure 2 shows the **DET curves** for all models. EdgeForehead achieves the best overall trade-off, attaining the lowest EER (8.30%), while MFN-V2+CosFace delivers the strongest low-FMR performance. SwinFace variants exhibit higher FNMRs, suggesting weaker mod-

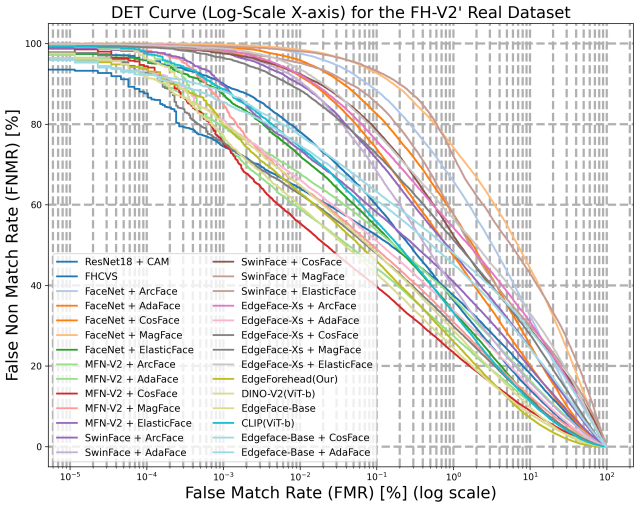

Fig. 2. DET curves for forehead-crease verification on the open-set FH-V2 dataset.

eling of fine-grained crease patterns. These results highlight the effectiveness of locality-preserving architectures and trait-specific learning for robust forehead-crease verification.

6 Conclusion

This work presents a systematic comparison of MLLMs (GPT-4o, Claude, Gemini), domain-adaptive learning, and trait-specific models for forehead-crease biometric verification. While MLLMs enable interpretable zero-shot verification, their performance remains below specialized biometric systems. Domain-adaptive and trait-specific models consistently achieve lower EER and higher TMR, highlighting the importance of discriminative biometric learning. These findings indicate that FMs offer interpretable verification and valuable zero-shot capabilities, but security-critical biometric applications still require trait-specific learning for reliable deployment.

References

1. Anthropic: Claude 4 Opus and claude 4 sonnet. `https://claude.ai/` (2025), latest release May 22, 2025
2. Bharadwaj, R., Jaswal, G., Nigam, A., Tiwari, K.: Mobile based human identification using forehead creases: Application and assessment under covid-19 masked face scenarios. In: Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision. pp. 3693–3701 (2022)
3. Boutros, F., Damer, N., Kirchbuchner, F., Kuijper, A.: Elasticface: Elastic margin loss for deep face recognition. In: Proceedings of the IEEE/CVF conference on computer vision and pattern recognition. pp. 1578–1587 (2022)

4. Chen, S., Liu, Y., Gao, X., Han, Z.: Mobilefacenets: Efficient cnns for accurate real-time face verification on mobile devices. In: Chinese conference on biometric recognition. pp. 428–438. Springer (2018)
5. George, A., Ecabert, C., Shahreza, H.O., Kotwal, K., Marcel, S.: Edgeface: Efficient face recognition model for edge devices. IEEE Transactions on Biometrics, Behavior, and Identity Science **6**(2), 158–168 (2024). https://doi.org/10.1109/TBIOM.2024.3352164
6. Jain, A.K., Ross, A., Prabhakar, S.: An introduction to biometric recognition. IEEE Transactions on circuits and systems for video technology **14**(1), 4–20 (2004)
7. Kim, M., Jain, A.K., Liu, X.: Adaface: Quality adaptive margin for face recognition. In: Proceedings of the IEEE/CVF conference on computer vision and pattern recognition. pp. 18750–18759 (2022)
8. Li, C., Li, L., Jiang, H., Weng, K., Geng, Y., Li, L., Ke, Z., Li, Q., Cheng, M., Nie, W., et al.: Yolov6: A single-stage object detection framework for industrial applications. arXiv preprint arXiv:2209.02976 (2022)
9. Meng, Q., Zhao, S., Huang, Z., Zhou, F.: Magface: A universal representation for face recognition and quality assessment. In: Proceedings of the IEEE/CVF conference on computer vision and pattern recognition. pp. 14225–14234 (2021)
10. OpenAI: Gpt-4o. https://openai.com/index/hello-gpt-4o/ (2024), accessed: 2025-06-23
11. Oquab, M., Darcet, T., Moutakanni, T., Vo, H., Szafraniec, M., Khalidov, V., Fernandez, P., Haziza, D., Massa, F., El-Nouby, A., et al.: Dinov2: Learning robust visual features without supervision. arXiv preprint arXiv:2304.07193 (2023)
12. Qin, L., Wang, M., Deng, C., Wang, K., Chen, X., Hu, J., Deng, W.: Swinface: a multi-task transformer for face recognition, expression recognition, age estimation and attribute estimation. IEEE Transactions on Circuits and Systems for Video Technology **34**(4), 2223–2234 (2023)
13. Radford, A., Kim, J.W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., et al.: Learning transferable visual models from natural language supervision. In: International conference on machine learning. pp. 8748–8763. PmLR (2021)
14. Schroff, F., Kalenichenko, D., Philbin, J.: Facenet: A unified embedding for face recognition and clustering. In: Proceedings of the IEEE conference on computer vision and pattern recognition. pp. 815–823 (2015)
15. Sharma, G., Jaswal, G., Nigam, A., Ramachandra, R.: Fh-sstnet: Forehead creases based user verification using spatio-spatial temporal network. In: 2024 12th International Workshop on Biometrics and Forensics (IWBF). pp. 1–6. IEEE (2024)
16. Tandon, A., Sharma, G., Jaswal, G., Nigam, A., Ramachandra, R.: Synthetic forehead-creases biometric generation for reliable user verification. In: 2024 IEEE International Joint Conference on Biometrics (IJCB). pp. 1–9. IEEE (2024)
17. Tandon, A., Sharma, G., Jaswal, G., Nigam, A., Ramachandra, R.: Generating realistic forehead-creases for user verification via conditioned piecewise polynomial curves. arXiv preprint arXiv:2501.13889 (2025)
18. Team, G., Anil, R., Borgeaud, S., Alayrac, J.B., Yu, J., Soricut, R., Schalkwyk, J., Dai, A.M., Hauth, A., Millican, K., et al.: Gemini: a family of highly capable multimodal models. arXiv preprint arXiv:2312.11805 (2023)
19. Wang, H., Wang, Y., Zhou, Z., Ji, X., Gong, D., Zhou, J., Li, Z., Liu, W.: Cosface: Large margin cosine loss for deep face recognition. In: Proceedings of the IEEE conference on computer vision and pattern recognition. pp. 5265–5274 (2018)