arXiv:2610.31065v1 [cs.CV] 22 Oct 2026

A General-Purpose Dual-Path Interaction Framework for Multi-View Medical Image Segmentation: Alleviating Data Volume Dependency in Small-Sample Regimes

Shuhuan Wang[a,1], Yuanyuan Cui[b,1], Pengju Wang[a], Huimin Guo[a], Ling Yue[b], Yin Gu[a], Yanchao Zhang[b], Wei Feng[b], Yunhe Liu[b], Wei Qian[b], Mingjun Xu[d,*], Shuang Liu[b,*], He Ma[a,c,*]

[a] College of Medicine and Biological Information Engineering, Northeastern University, Shenyang, Liaoning, 110169, China
[b] Department of Ultrasound, The Fourth Affiliated Hospital of China Medical University, Shenyang, Liaoning, 110000, China
[c] The Key Laboratory of Intelligent Computing in Medical Image, Ministry of Education, Northeastern University, Shenyang, Liaoning, 110819, China
[d] Department of Cardiology, Qilu Hospital of Shandong University
[1] These authors contributed equally

October 23, 2026

Abstract

The performance of medical image segmentation depends on a large amount of annotated data, and most segmentation models experience a significant drop in performance in low-sample-size scenarios. Existing semi-supervised and weakly supervised methods rely on specific network architectures, while traditional multi-view-assisted segmentation methods require strictly paired samples, which limits their clinical applications. To address these issues, we propose DPIF, a general-purpose plug-and-play dual-path interaction framework compatible with arbitrary encoder-decoder segmentation networks, which leverages cross-view anatomical complementarity to reduce annotation reliance. DPIF first performs Sinkhorn optimal transport-based cross-view distribution alignment to establish global semantic correspondence without sample-level pairing. It then adopts conditional mutual information-guided feature decoupling to enable bidirectional knowledge transfer only within the shared semantic space, suppressing view-specific noise. Finally, persistent homology-based topological regularization encodes anatomical invariants as unsupervised priors to mitigate overfitting. Experiments were conducted using 24 baseline models on a private multicenter echocardiography dataset and two public datasets. The results show that DPIF-enhanced models outperform single-view baseline models when trained on comparable amounts of data. Notably, most dual-path variants trained on 40% - 60% of the data achieve performance comparable to that of single-view models trained on the full dataset. This demonstrates that DPIF provides a general-purpose and scalable solution for data-scarce multi-view medical image segmentation problems.

Keywords: Medical Image Segmentation, Multi-view Analysis, Few-shot Learning, Dual-path Interaction

1 Introduction

Accurate medical image segmentation is the foundation for clinical quantitative analysis, disease diagnosis, and surgical planning. Although deep learning has advanced this field, its performance still depends on large-scale datasets with pixel-level annotationsGan et al. [2025]. Unlike natural images, annotating medical images requires specialized clinical knowledge, making the process both time-consuming and costly. For multi-view imaging protocols, each view must be annotated separately, leading to a significant increase in workload and costs. Therefore, the scarcity of annotated data remains a major bottleneck for clinical translation, particularly for rare diseasesDissanayake et al. [2025].

In standard clinical practice, images of the same anatomical structure are typically acquired from multiple angles to enable a comprehensive assessment of lesionsBerahmand et al. [2025]. For example, a standard echocardiogram simultaneously acquires both the long-axis and short-axis views of the left ventricle. Although these views contain highly complementary anatomical information, they differ significantly in terms of geometric structure and imaging characteristicsDu et al. [2026]. Current mainstream segmentation methods still process each view independently, failing to fully leverage the inherent structural and semantic relationships between views. Most existing multi-view fusion methods employ simple strategiesLiu et al. [2024b], Majerník et al. [2026]. These methods overlook the substantial spatial structural differences and variations in semantic distribution across views, often introducing redundant noise and yielding only limited performance improvements. More critically, most multi-view frameworks are tightly coupled with specific backbone architectures, which hinders their rapid adaptation to new tasks and poses challenges for cross-scenario deploymentZhou et al. [2026].

To address these fundamental limitations, we propose a general-purpose dual-path interaction framework. Its core idea is to leverage the complementarity of cross-view anatomical structures as an alternative to expanding the annotated dataset. The framework features a universal, plug-and-play design that is compatible with various segmentation backbone models and achieves high-precision segmentation with only a small number of annotated samples. The main contributions of this study are summarized as follows:

We propose a general-purpose dual-path interaction framework designed to reduce reliance on segmented annotated data. This framework can be integrated into any encoder-decoder Segmentation network. We propose a cross-view alignment mechanism based on optimal transport theory. It establishes semantic correspondences by minimizing the Wasserstein distance, thereby enabling distribution-level alignment of unpaired data. We propose a two-layer decoupling strategy for cross-view knowledge transfer. Under conditional mutual information constraints, features are decomposed into shared and view-specific components, and knowledge transfer and enhancement are performed only in the shared space to suppress noise propagation. We propose a regularization method based on cross-view topological consistency. This method uses anatomical topological invariants as shared prior information to impose constraints, thereby narrowing the hypothesis space and mitigating the problem of overfitting in small-sample settings.

2 Related WORK

2.1 Multi-view Medical Image Segmentation

Existing multi-view medical image segmentation methods can be broadly categorized into three paradigms: input-level, feature-level, and decision-level fusion. Input-level fusion stacks multi-view images as multi-channel inputs but ignores spatial differences between views. Feature-level

fusion directly concatenates or superimposes intermediate features within the network, but carries the risk of introducing view-specific noise and artifactsWang and Voiculescu [2022]. Decision-level fusion, while capable of merging independent prediction outputs, discards complementary semantic information present in the hidden layersLi et al. [2025]. Although improvements based on attention mechanisms have shown some promise, most fusion schemes remain constrained by their network architectures, making it difficult to transfer them across different segmentation pipelinesZhan et al. [2023]. Furthermore, most of these schemes are designed for fully annotated datasets and perform poorly when data is limited. Crucially, nearly all existing methods require strictly paired multi-view patient data, which imposes severe constraints on real-world clinical inputs.

2.2 Small-Sample Medical Image Segmentation

To address the issue of scarce annotations, existing research primarily follows three approaches: semi-supervised learning, weakly supervised learning, and few-shot learning. Semi-supervised methods utilize large amounts of unlabeled data to aid training, but still rely on a certain scale of unlabeled data, thus imposing requirements on the data volumeCao et al. [2025]. Weakly supervised methods employ weak labels, such as point annotations and bounding box annotations, to reduce annotation costs. However, due to the limitations imposed by the simplicity of these labels, there is an upper bound on accuracy.Chen et al. [2025a] Few-shot learning methods aim to achieve generalization with small datasets through cross-domain knowledge transfer. However, due to significant variations across medical imaging domains, their practical effectiveness remains uncertainDissanayake et al. [2025]. It is worth noting that existing methods rarely exploit the inherent correlations among multi-view images of the same object, nor have they established a universal interaction scheme that can be reused across models. Meanwhile, the cost of acquiring multi-view images in clinical settings is far lower than the cost of annotating data required for a single view, and this potential has yet to be fully utilized.

3 Method

3.1 Overview

As shown in Fig 1, our proposed Dual-Path Interaction Framework (DPIF) adopts a symmetric dual-branch architecture, in which the backbone networks share the same structure but are initialized independently. Each branch receives a single-view image for segmenting the target anatomical structure. The training process is divided into three stages. In the feature encoding stage, a dual-branch encoder extracts multiscale features. Subsequently, the ODI module is invoked at designated encoder-decoder layers to perform feature decoupling, unpaired distribution alignment, and adaptive bidirectional feature enhancement, after which the enhanced features are fed back to the decoder. Meanwhile, a topological regularization loss serves as an anatomical prior to ensure topological consistency across different views. In the final decoding and joint training stage, each decoder upsamples and aggregates the features to generate segmentation outputs and soft probability maps. The entire framework is optimized via a composite loss function. During the inference stage, the interaction module is disabled, and the two branches operate independently.

3.2 Optimal Transmission Decoupling Interaction Module

The ODI module is the core functional unit of DPIF, enabling cross-view feature interaction. The specific workflow of this module consists of three consecutive steps: two-layer feature decoupling,

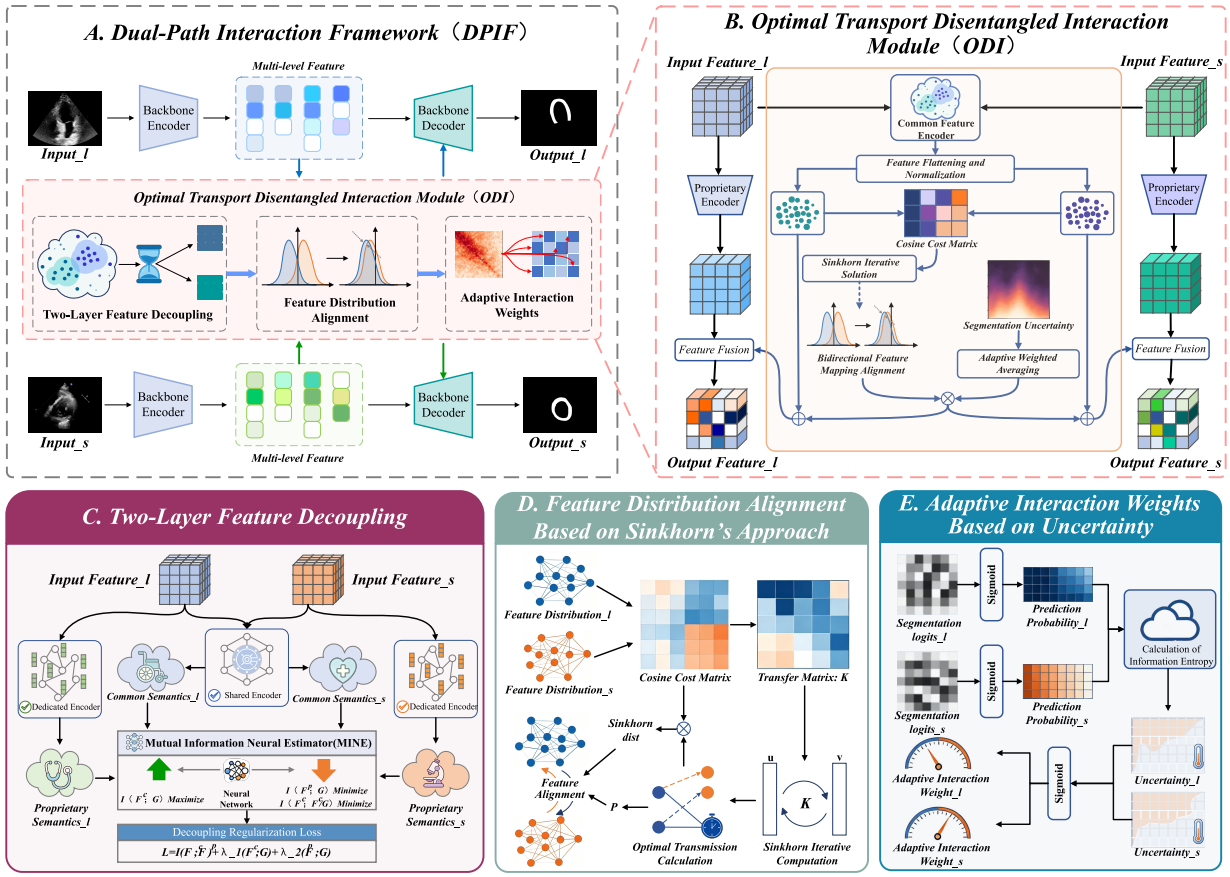

Figure 1: DPIF Framework and Functional Modules. (A) Schematic diagram of the dual-branch interaction framework deployment. (B) Core interaction module, used to connect the two segmentation pipelines. (C) Two-layer feature decoupling method, which distinguishes between shared and proprietary features. (D) Feature distribution alignment algorithm, which enables bidirectional semantic interaction and enhancement. (E) Segmentation uncertainty-guided weight interaction method, which enables adaptive adjustment of interaction intensity.

unpaired distribution alignment, and adaptive bidirectional feature enhancement.

3.2.1 Two-Layer Feature Decoupling

Multi-view medical images contain both shared anatomical semantics and view-specific noise, intensity shifts, and artifacts. Directly propagating raw features across branches reduces the quality of feature interaction, especially when data is limited. To address this issue, our proposed two-layer feature decoupling method decomposes the input features of each view into two mutually exclusive components: a shared semantic component for cross-view information exchange, and a view-specific component that is excluded from inter-branch interaction. Each branch employs two parallel 1×1 convolutional encoders. One encoder, with shared weights, extracts shared semantic features F^c to construct a unified semantic space, while the view-independent encoder generates view-specific features F^p:

$$F^c = \mathrm{Enc}_c(F), \quad F^p = \mathrm{Enc}_p(F) \tag{1}$$

To guarantee effective decoupling, we introduce a conditional mutual information constraint as a regularization term:

$$\mathcal{L}_{\text{decouple}} = I(F^c; F^p|G) - \lambda_1 I(F^c; G) - \lambda_2 I(F^p; G) \tag{2}$$

where $I(\cdot;\cdot)$ denotes mutual information and G is the segmentation label. This loss maximizes the dependence between shared components and segmentation labels, and minimizes conditional mutual-information between F^c and F^p given labels, guaranteeing their statistical independence and preventing noise from polluting the shared semantic space.

3.2.2 Sinkhorn-Based Unpaired Distribution Alignment

Given the scarcity of multi-view data for sample-level pairing, which limits the scope of application, we reformulate cross-view semantic matching as an sinkhorn-Based optimal transport problem to achieve global alignment at the distribution levelWatanabe and Isobe [2024]. First, the shared semantic features from both branches are flattened into sets of feature points. For the long-axis branch, the feature set is $X = \{x_i\}_{i=1}^{N_l}$ with $x_i \in \mathbb{R}^{C_c}$, corresponding to the discrete distribution $\mu = \frac{1}{N_l}\sum_{i=1}^{N_l} \delta_{x_i}$. For the short-axis branch, the feature set is $Y = \{y_j\}_{j=1}^{N_s}$ with $y_j \in \mathbb{R}^{C_c}$, corresponding to the discrete distribution $\nu = \frac{1}{N_s}\sum_{j=1}^{N_s} \delta_{y_j}$. Here, $N_l = B_l \cdot H_l \cdot W_l$ and $N_s = B_s \cdot H_s \cdot W_s$ denote the total number of feature points in the two branches, respectively. The two values can differ, making the formulation compatible with unpaired settings. We define the cost matrix $C \in \mathbb{R}^{N_l \times N_s}$, where each element $C_{ij} = \|x_i - y_j\|_2^2$ represents the squared Euclidean distance between two feature points, i.e., the transport cost. The optimal transport problem aims to find a transport plan $P \in \Pi(\mu, \nu)$ that satisfies the marginal distribution constraints while minimizing the total transport cost: $\min_{P\in\Pi(\mu,\nu)} \langle P, C\rangle_F$. To improve computational efficiency and ensure differentiability, we introduce an entropy regularization term, converting the problem into a Sinkhorn distance optimization:

$$d_{\text{Sinkhorn}}(\mu, \nu) = \min_{P\in\Pi(\mu,\nu)} \langle P, C\rangle_F - \varepsilon H(P) \tag{3}$$

where $H(P) = -\sum_{i,j} P_{ij} \log P_{ij}$ is the entropy of the transport plan and ε is the regularization coefficient. This problem can be efficiently solved via the Sinkhorn–Knopp iterative algorithm to yield the optimal transport plan P^*. Based on the optimal transport plan, we achieve bidirectional

feature mapping and enhancement. Short-axis features are mapped to the long-axis feature space to generate complementary features for the long-axis view: $F^c_{l \leftarrow s} = P^* \cdot F^c_s$. Similarly, complementary features for the short-axis view are obtained as $F^c_{s \leftarrow l} = (P^*)^T \cdot F^c_l$. This process establishes semantic correspondence between the two views at the global distribution level, enabling knowledge transfer without requiring sample-level pairing.

3.2.3 Adaptive Interaction Weights Based on Segmentation Uncertainty

In clinical practice, there are inherent differences in imaging quality and segmentation difficulty across different views. Fixed, symmetric interaction cannot adapt to such differences, often resulting in underutilization of high-quality information while allowing low-quality information to cause excessive interference. To address this issue, we designed an energy function based on segmentation uncertainty to dynamically adjust the intensity of bidirectional interaction, thereby enabling the directed transmission of high-confidence features to the low-confidence branch.

First, we calculate the segmentation uncertainty for each view based on its predicted map:

$$E = -\frac{1}{N} \sum_{i=1}^{N} \left(p_i \log p_i + (1 - p_i) \log(1 - p_i) \right) \tag{4}$$

Here, p_i is the predicted probability that pixel i belongs to the target region, and $E \in [0, 1]$. A higher value indicates greater segmentation uncertainty and lower confidence.

Based on the difference in uncertainty between the two branches, adaptive interaction weights are generated using the sigmoid function:

$$\alpha_l = \alpha_0 \cdot \sigma \left(\frac{E_l - E_s}{\tau} \right), \quad \alpha_s = \alpha_0 \cdot \sigma \left(\frac{E_s - E_l}{\tau} \right) \tag{5}$$

where α_0 is the base interaction strength, and τ is the temperature coefficient that controls the smoothness of weight adjustments. When the confidence in the long-axis segmentation is significantly lower than that of the short-axis segmentation, α_l increases, causing the long-axis branch to receive more guidance from the short-axis branch, the opposite is also true.

3.3 Cross-View Topological Consistency Regularization Loss

In small-sample scenarios, limited training coverage may induce spurious pixel-level correlations and produce unreasonable segmentation results, while pixel-based loss functions struggle to mitigate this phenomenon. We consider that anatomical structures possess inherent topological invariance, and that key topological statistics across multi-view projections exhibit high consistency, which can serve as unsupervised priors to narrow the model's hypothesis space and alleviate overfitting. Therefore, drawing inspiration from persistent homology-based topological data analysis methods, we propose a cross-view topological consistency regularization lossKapulkin and Kershaw [2026]. Starting from the segmentation probability map of each view, we obtain a persistent diagram (PD) through sub-level set filtering, where the points encode topological features characterized by thresholds and corresponding persistent values. Higher persistence values indicate more stable and anatomically plausible structures. Since the exact persistence homology calculation is non-differentiable, we employ a differentiable approximation method: applying multiscale Gaussian smoothing to the probability graph to obtain an approximate persistence graph. The second-order Wasserstein distance between the persistence graphs of the two views forms our topological loss function:

$$\mathcal{L}_{\text{topo}} = W_2(\text{PD}(P_l), \text{PD}(P_s)) \tag{6}$$

where $\text{PD}(\cdot)$ denotes the persistent diagram from which the segmentation results are extracted, and $W_2(\cdot)$ is the second-order Wasserstein distance. This loss ensures that the segmentation results from the two viewpoints remain consistent in their topological statistical distributions, guiding the model to learn segmentation patterns that conform to anatomical priors.

3.4 Joint Loss Function

The total loss of DPIF integrates four terms to balance segmentation performance, feature-decoupling quality, topological consistency and cross-view knowledge complementarity for end-to-end joint optimization:

$$\mathcal{L}\text{total} = \mathcal{L}\text{seg} + \lambda_d \mathcal{L}\text{decouple} + \lambda_t \mathcal{L}\text{topo} + \lambda_k \mathcal{L}_{\text{kd}} \tag{7}$$

where λ_d, λ_t, λ_k denote loss-weight hyper-parameters. The primary segmentation loss $\mathcal{L}_{\text{seg}}$ combines Dice loss and binary cross-entropy (BCE) to handle medical class imbalance, computed and averaged independently for each view.The decoupling loss enforces mutual-information constraints: it minimizes conditional mutual information between shared/view-specific features given labels, maximizes relevance between shared features and annotations, and suppresses task-irrelevant view-specific signals:

$$\mathcal{L}_{\text{decouple}} = I(F^c; F^p \mid G) - \lambda_1 I(F^c; G) - \lambda_2 I(F^p; G) \tag{8}$$

Here $I(\cdot;\cdot)$ and $I(\cdot;\cdot \mid \cdot)$ represent mutual and conditional mutual information. F^c, F^p are shared and view-specific features, G is the segmentation label, and λ_1, λ_2 are weighting coefficients. The topological consistency loss regularizes cross-view topological statistics via persistent-homology priors to shrink the model hypothesis space:

$$\mathcal{L}\text{topo} = W_2\big(\text{PD}(P_l), \text{PD}(P_s)\big) \tag{9}$$

$\text{PD}(\cdot)$ extracts persistence diagrams from predicted probability maps. $W_2(\cdot,\cdot)$ is the second-order Wasserstein distance measuring topological-distribution discrepancy. For discrete persistence point-sets $X = \{x_i\}_{i=1}^m$, $Y = \{y_j\}_{j=1}^n$:

$$W_2(X,Y) = \left(\min \gamma \in \Pi(X,Y) \sum_{i=1}^{m} \sum_{j=1}^{n} \gamma_{ij} |x_i - y_j|_2^2 \right)^{\frac{1}{2}} \tag{10}$$

where $\Pi(X,Y)$ denotes feasible coupling matrices and γ_{ij} stands for coupling weight entries.Bidirectional KL-divergence knowledge distillation is applied to segmentation probability maps to mine implicit complementary information and align cross-view semantics:

$$\begin{aligned} \mathcal{L}_{\text{kd}} = \frac{T^2}{2} \Big(& \text{KL}\left(\text{softmax}(P_l/T) \parallel \text{softmax}(P_s/T)\right) \\ & + \text{KL}\left(\text{softmax}(P_s/T) \parallel \text{softmax}(P_l/T)\right) \Big) \end{aligned} \tag{11}$$

where temperature T smooths soft-label outputs while retaining fine structural details.

4 Experiment

4.1 Datasets

We conducted our evaluation using three multi-view medical imaging datasets, including a private multicenter echocardiography dataset (referred to as PME), the publicly available FLARE2023 liver CT dataset, and the MSD hippocampal MRI dataset. The specific datasets used are shown in Fig 2.

PME Dataset: This dataset compiles data from 429 patients across five clinical centers, providing images of the apical four-chamber long-axis view and the papillary muscle-level short-axis view, along with annotations of the left ventricular myocardium. One image at the end of systole and one at the end of diastole are selected for each cardiac cycle, resulting in 858 images per view.

FLARE2023 DatasetMa et al. [2024]: 2D slices with liver annotations from two views were extracted from the 3D abdominal CT volumes. After filtering based on liver size and the presence of tumors, 581 images were retained for each view for 581 patients.

MSD DatasetSimpson et al. [2019]: We extracted 2D sections with dual views labeled for the anterior hippocampus. After image quality screening, 172 images were retained for each view for 172 patients.

To distinguish between the two views, the results are uniformly labeled as the "long-axis view" and the "short-axis view." All datasets were split at the patient level to prevent data leaks. First, 20% of the data was selected as an independent test set (not used for training), and the remaining data was subjected to five-fold cross-validation. All datasets underwent a standardized preprocessing workflow. For each dataset, images were cropped to a size of 512×512 pixels. During training, data augmentation was applied only to the training set to improve the model's generalization ability. Neither the validation set nor the test set underwent data augmentation to ensure the objectivity and consistency of the evaluation.

4.2 Evaluation Metrics

We adopted four commonly used evaluation metrics in medical image segmentation: The Dice similarity coefficient(DSC), which measures the degree of overlap between the segmentation results and the gold standard. The Jaccard index (IoU), which is the ratio of the intersection to the union. The 95% Hausdorff distance (HD95), which measures the degree of alignment of the segmentation boundaries. The average surface distance (ASD), which is the average distance between points on the segmentation boundaries. Using 5-fold cross-validation, we identified five optimal models, tested them on the independent test set, and reported the average values of the aforementioned metrics.

4.3 Experimental Setup

We implemented all networks using the PyTorch 1.12 framework. The Adam optimizer was used with an initial learning rate of 2e-4, which was reduced to a minimum learning rate of 1e-5 using a cosine annealing strategy. The batch size was set to 2, with a total of 100 training epochs. The backbone network was frozen for the first 30 epochs. The similarity threshold was set to $\tau = 0.2$, the distillation temperature to $T = 2.0$, and the loss weights to $\alpha = 0.1, \beta_s = 0.05, \beta_l = 0.05$. Except for specific parameters, the settings for the comparative experiments were identical to those of our method. All experiments were conducted on two NVIDIA RTX 4090 GPUs.

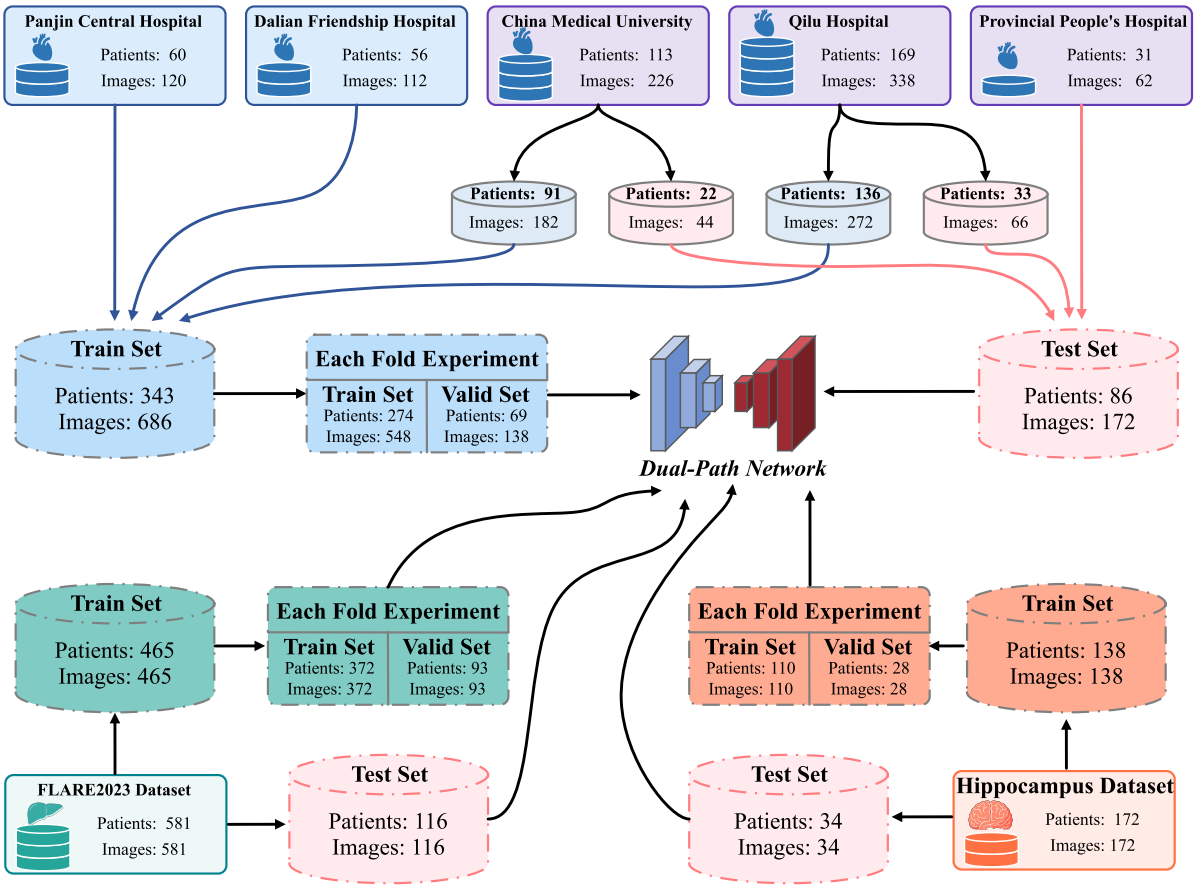

Figure 2: Details on the use of the three datasets. For each dataset, 20% is randomly allocated as an independent test set, and the remaining data is used for training and validation using five-fold cross-validation.

Table 1: Results of the single-path and dual-path models on the PME dataset from the long-axis perspective. Here, PI represents the percentage improvement in the DSC metric achieved by the DP method relative to the SP method. The best result is shown in red, the second-best in purple, and the third-best in blue.

Method	Long Axis									Short Axis								
	SP				DP				PI(+%)	SP				DP				PI(+%)
	DSC	HD95	ASD	IOU	DSC	HD95	ASD	IOU		DSC	HD95	ASD	IOU	DSC	HD95	ASD	IOU	
KiUNet Valanarasu et al. [2020]	0.5232	22.1285	8.4941	0.3636	0.5866	20.6775	5.5894	0.4149	12.12	0.4827	22.8059	9.0319	0.3182	0.6394	16.7441	5.1721	0.4706	32.46
RollingUNet Liu et al. [2024d]	0.5959	20.8334	7.5748	0.4339	0.7841	12.5132	3.8435	0.6449	31.58	0.3449	31.2219	11.9033	0.2084	0.3866	26.6497	10.0282	0.2398	12.09
UTNet Gao et al. [2021]	0.6303	18.0414	6.0789	0.4767	0.7218	16.0721	4.9658	0.5647	14.52	0.6085	20.2116	5.2387	0.4375	0.6467	17.7436	5.547	0.4774	6.28
DSCNet Qi et al. [2023]	0.6667	17.5818	6.0677	0.5119	0.7095	16.1644	4.8706	0.5498	6.42	0.4348	24.793	9.7012	0.2778	0.5694	21.5876	7.4392	0.3982	30.96
CGNet Wu et al. [2020]	0.6746	17.2862	5.7327	0.5244	0.7411	14.7987	4.2023	0.5887	9.86	0.7115	16.278	4.3075	0.5526	0.7557	14.6831	4.2626	0.6064	6.21
PSPNet Zhao et al. [2017]	0.6943	16.2117	4.4803	0.5425	0.7674	14.1164	3.9533	0.6226	10.53	0.6265	19.3502	5.1265	0.4560	0.7244	16.0846	4.1613	0.5687	15.63
SegNet Badrinarayanan et al. [2017]	0.7008	16.1720	4.4165	0.5566	0.7542	14.7153	4.2537	0.6054	7.62	0.6513	17.4458	5.8651	0.4826	0.6647	16.6071	5.0469	0.4973	2.06
SUNet Li et al. [2023]	0.7009	16.9809	4.4306	0.5531	0.7852	13.2798	3.9445	0.6464	12.03	0.6145	20.1446	5.1867	0.4436	0.6494	17.7968	5.9124	0.4804	5.68
TinyUNet Chen et al. [2024]	0.7013	16.1626	4.3507	0.5545	0.7423	14.8933	4.2606	0.5902	5.85	0.4982	23.1913	9.5385	0.3317	0.6955	16.1829	4.8591	0.5335	39.60
TransUNet Chen et al. [2021]	0.7065	16.1293	4.3538	0.5605	0.7119	16.1154	4.8751	0.5527	0.76	0.4708	24.1321	9.5772	0.3077	0.6055	20.5665	7.3924	0.4348	28.61
TCSegNet Lal [2024]	0.7075	16.1478	5.1605	0.5474	0.7088	16.1949	4.4496	0.5653	0.18	0.5884	20.9776	7.6207	0.4165	0.7397	14.7148	4.2574	0.5868	25.71
UNext Valanarasu and Patel [2022]	0.7086	16.1927	4.4988	0.5633	0.7505	14.8914	4.2439	0.6006	5.91	0.5856	20.6931	5.6859	0.4137	0.6918	16.3731	4.4644	0.5283	18.14
ICNet Zhao et al. [2018]	0.7101	16.1261	4.4558	0.5619	0.7895	13.2946	3.8133	0.6522	11.18	0.6656	16.6114	5.0518	0.4985	0.7009	16.1627	4.5369	0.5396	5.30
CFANet Zhang and Yan [2023]	0.7104	16.2234	4.3477	0.5651	0.8042	12.7379	3.3653	0.6725	13.20	0.4974	23.3267	9.2401	0.3302	0.5914	20.7906	7.6695	0.4195	18.90
LETNet Xu et al. [2023]	0.7124	16.1467	4.2671	0.5705	0.7703	13.3076	4.0313	0.6264	8.13	0.6851	16.2453	4.9756	0.5192	0.7052	16.2683	4.5233	0.5445	2.93
MAUNet Chen et al. [2025b]	0.7174	16.1472	4.3334	0.5753	0.7245	15.6569	4.1872	0.5681	0.99	0.6267	18.3737	6.1602	0.4562	0.7424	14.8253	4.2448	0.5903	18.46
SCSEUNet Liu et al. [2024c]	0.7178	16.2116	4.8531	0.5599	0.7336	14.3787	4.0558	0.5929	2.20	0.5818	20.8129	5.5795	0.4106	0.5843	21.4814	7.6181	0.4127	0.43
DUCKNet Dumitru and Peteleaza [2023]	0.7197	15.8825	4.8644	0.5621	0.7391	14.351	4.0117	0.5997	2.70	0.7267	22.5628	4.7218	0.5708	0.7507	15.9823	4.6254	0.6013	3.30
DeepLabv3 Chen et al. [2018]	0.7301	14.4663	4.1456	0.5874	0.7625	14.2943	4.0832	0.6162	4.44	0.6457	16.7308	5.1075	0.4765	0.7491	13.5302	4.0513	0.5988	16.01
I2UNet Dai et al. [2024]	0.7323	14.3905	4.1328	0.5924	0.7829	12.8860	3.8676	0.6433	6.91	0.7265	16.2223	4.7267	0.5709	0.7371	14.3615	3.9148	0.5836	1.46
CPFNet Feng et al. [2020]	0.7409	14.3676	3.9011	0.6023	0.8097	12.6148	3.3271	0.6802	9.29	0.7596	14.3028	4.1601	0.6128	0.7759	13.5073	3.2019	0.6344	2.15
UNet Ronneberger et al. [2015]	0.7417	14.3168	4.0564	0.6038	0.8125	12.5713	3.3209	0.6842	9.55	0.7639	13.3085	3.7427	0.6182	0.7794	13.2097	3.4591	0.6386	2.03
EANet Wang et al. [2022]	0.7439	14.6818	4.0547	0.5922	0.7517	13.5500	3.9219	0.6019	1.05	0.6615	17.6969	6.1624	0.4936	0.7126	16.2989	4.2979	0.5537	7.72
UMNet Du et al. [2025]	0.8058	12.6179	3.2484	0.6748	0.8519	11.1098	3.1702	0.7421	5.72	0.7984	12.5134	3.7446	0.6646	0.8495	11.0153	3.1246	0.7384	6.40

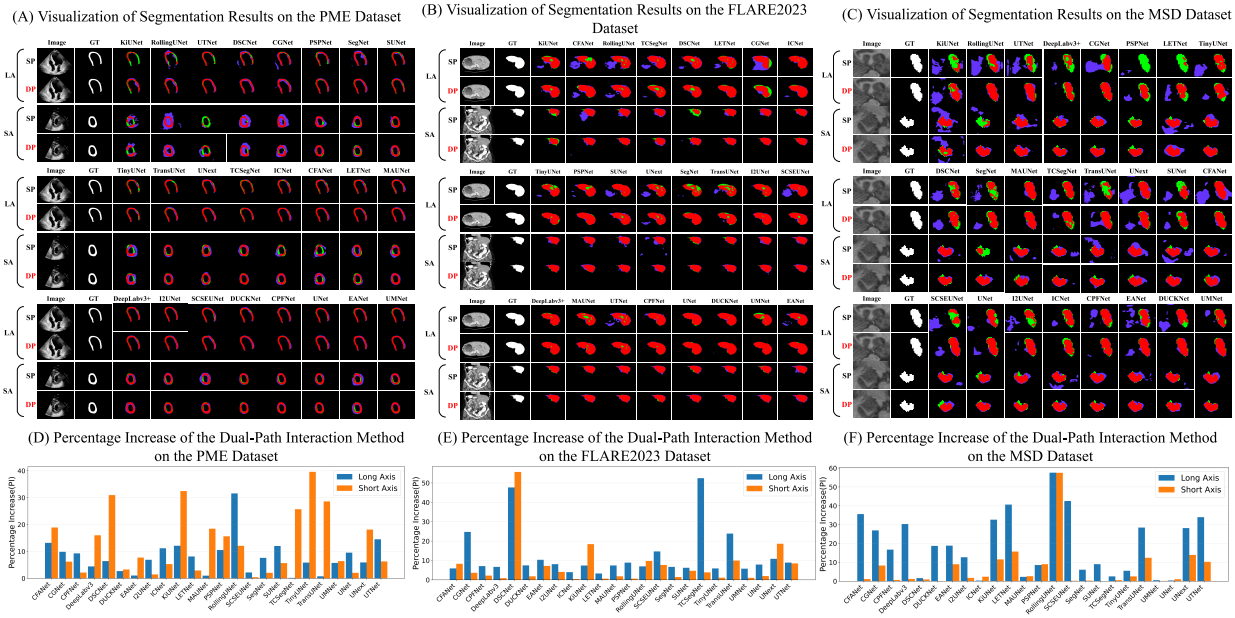

Figure 3: Visualization of the segmentation results and the bar chart showing the percentage increase in DSC. (A)–(C) show the validation results of the segmentation results against the mask, where red indicates the overlapping regions, purple indicates over-segmented regions, and green indicates under-segmented regions. (D)–(F) show the percentage improvements in the DSC metric achieved by the dual-path method.

5 Results

5.1 Comparison with Single-Path Segmentation Models

To validate the DPIF's generality and its advantages in few-shot scenarios, we selected 24 representative encoder-decoder networks as single-path (SP) baseline models and constructed corresponding dual-path (DP) variants by integrating the DPIF framework. All models were trained and evaluated on three multi-view datasets. The complete quantitative results are shown in TABLE 1 to TABLE 3 and Fig 3. Experimental results show that all dual-path variants achieved performance improvements on the three datasets. In terms of the DSC metric, the dual-path variants achieved an average improvement of 13.33% on the long-axis views and 8.97% on the short-axis views across the three datasets. Notably, the hippocampus long-axis dataset showed the largest average DSC improvement, reaching 19.2%. Among these, RollingUNet demonstrated the most significant improvement, with a 57.57% increase in DSC on the hippocampus long-axis dataset. This confirms that DPIF can effectively leverage the complementary anatomical information inherent in both views without introducing directional performance bias. Regarding boundary accuracy metrics, the average HD95 decreased by 4.27 in the long-axis view and by 2.38 in the short-axis view across the three datasets. It is worth noting that a few models (such as SCSEUNet) showed a slight decrease in HD95. Nevertheless, the overall trend confirms that DPIF effectively enhances the boundary continuity and anatomical fidelity of segmentation results for most architectures, delivering stable and consistent performance improvements across a variety of segmentation network architectures.

Table 2: Results of the single-path and dual-path models on the FLARE2023 dataset.

Method	Long Axis									Short Axis								
	SP				DP				PI(+%)	SP				DP				PI(+%)
	DSC	HD95	ASD	IOU	DSC	HD95	ASD	IOU		DSC	HD95	ASD	IOU	DSC	HD95	ASD	IOU	
TCSegNet	0.6198	18.3918	5.8360	0.4491	0.9449	4.5801	2.0367	0.8956	52.45	0.8589	7.2738	2.7845	0.7527	0.8920	6.5329	2.5299	0.8056	3.85
DSCNet	0.6424	17.4416	5.7550	0.4732	0.9486	4.5296	2.0461	0.9022	47.67	0.5731	20.5699	5.7285	0.4016	0.8921	6.4391	2.3964	0.8053	55.66
KiUNet	0.7584	14.5091	4.1135	0.6108	0.8145	12.3562	3.6908	0.6871	7.40	0.7152	16.2628	4.8605	0.5567	0.8470	7.4423	2.8995	0.7345	18.43
CGNet	0.7669	14.3293	4.1465	0.6213	0.9565	3.6063	1.7078	0.917	24.72	0.9001	6.3828	2.3235	0.8203	0.9328	5.5925	1.9791	0.8767	3.63
TransUNet	0.7861	12.7978	3.8189	0.6475	0.9739	2.4119	1.6209	0.9498	23.89	0.8476	7.3489	2.8246	0.7353	0.9319	5.4060	2.2814	0.8727	9.95
CFANet	0.8179	12.3989	3.6128	0.6918	0.8661	7.2256	2.6655	0.7637	5.89	0.7964	12.5354	3.7402	0.6620	0.8625	7.1838	2.8289	0.7582	8.30
SCSEUNet	0.8516	7.2741	2.8618	0.7416	0.9765	2.1880	1.6218	0.9547	14.67	0.8723	7.1075	2.4796	0.7735	0.9395	4.7404	2.2647	0.8864	7.70
RollingUNet	0.8545	7.3239	2.8240	0.7460	0.9140	6.3252	2.1794	0.8417	6.96	0.8565	7.2412	2.7991	0.7491	0.9401	4.4841	1.9321	0.8879	9.76
UNext	0.8751	7.1859	2.4934	0.7780	0.9701	2.4416	1.6413	0.9424	10.86	0.7834	12.4592	3.9104	0.6439	0.9294	5.7157	2.3337	0.8679	18.64
PSPNet	0.8815	6.7152	2.4384	0.7881	0.9598	3.9024	1.7795	0.9223	8.88	0.9142	5.8518	2.1452	0.8422	0.9197	6.0335	2.3273	0.8516	0.60
EANet	0.8915	6.5431	2.6729	0.8042	0.9837	2.1538	1.5255	0.9683	10.34	0.8848	6.7879	2.6847	0.7935	0.9485	4.3918	1.9111	0.9021	7.20
UTNet	0.8975	6.6537	2.5574	0.8141	0.9774	2.4198	1.5849	0.9562	8.90	0.8597	7.2433	2.8360	0.7539	0.9327	5.4161	1.9518	0.8740	8.49
I2UNet	0.9013	6.5120	2.3801	0.8203	0.9742	2.3025	1.6284	0.9493	8.09	0.9041	6.4323	2.3329	0.8249	0.9413	4.5895	2.1356	0.8901	4.11
TinyUNet	0.9047	6.4935	2.3857	0.8260	0.9581	3.8255	1.7483	0.9201	5.90	0.9004	6.4887	2.4627	0.8190	0.9105	6.0412	2.2603	0.8358	1.12
UNet	0.9076	6.5079	2.3974	0.8308	0.9793	2.1446	1.5695	0.9594	7.90	0.9235	5.6229	2.1299	0.8601	0.9417	4.4993	1.9436	0.8901	1.97
SUNet	0.9078	6.5068	2.3987	0.8312	0.9644	3.2756	1.8659	0.9318	6.23	0.8788	7.1364	2.4687	0.7838	0.9198	5.9112	2.1908	0.8518	4.67
MAUNet	0.9097	6.4789	2.3726	0.8344	0.9772	2.4285	1.5712	0.9558	7.42	0.9228	5.6887	2.1799	0.8567	0.9399	4.5761	2.0137	0.8876	1.85
SegNet	0.9106	6.4370	2.3451	0.8361	0.9716	2.4916	1.6149	0.9445	6.70	0.9187	5.8774	2.1985	0.8495	0.9321	5.3141	2.1446	0.8732	1.46
DUCKNet	0.9114	6.1049	2.3078	0.8372	0.9794	2.1459	1.5851	0.9602	7.46	0.9172	5.8638	2.1904	0.8471	0.9340	5.5109	2.1776	0.8761	1.83
CPFNet	0.9139	6.2375	2.3126	0.8415	0.9788	2.3454	1.6046	0.9588	7.10	0.9177	5.8527	2.1849	0.8479	0.9379	5.3627	2.1594	0.8835	2.20
DeepLabv3	0.9158	6.1744	2.2782	0.8447	0.9771	2.6096	1.7213	0.955	6.69	0.9387	5.1275	1.9552	0.8845	0.9468	4.3287	1.9462	0.8997	0.86
LETNet	0.9186	6.1337	2.1355	0.8495	0.9488	4.4889	1.9556	0.9025	3.29	0.9285	5.3223	2.1837	0.8665	0.9341	5.6013	2.1619	0.8786	0.60
ICNet	0.9216	5.8365	2.1246	0.8545	0.9581	3.8212	1.7641	0.9197	3.96	0.9254	5.6478	2.1225	0.8616	0.9293	5.4848	2.0835	0.8698	0.42
UMNet	0.9288	5.5999	2.0612	0.8672	0.9824	2.1293	1.5425	0.9659	5.77	0.9404	4.4648	2.0467	0.8875	0.9503	4.1747	2.0689	0.9060	1.05

Table 3: Results of the single-path and dual-path models on the MSD dataset.

Method	Long Axis									Short Axis								
	SP				DP				PI(+%)	SP				DP				PI(+%)
	DSC	HD95	ASD	IOU	DSC	HD95	ASD	IOU		DSC	HD95	ASD	IOU	DSC	HD95	ASD	IOU	
KiUNet	0.5395	22.0747	7.0769	0.3693	0.7155	15.9292	4.9417	0.5573	32.62	0.5304	22.5674	7.1577	0.3611	0.5916	18.4389	6.2965	0.4191	11.54
LETNet	0.5526	20.3445	6.6149	0.3817	0.7772	12.9451	3.9491	0.6355	40.64	0.8073	11.8559	3.6003	0.6768	0.9341	5.6013	1.9176	0.8764	15.71
RollingUNet	0.5779	19.0086	6.6645	0.4064	0.9106	6.2819	2.1666	0.8359	57.57	0.5806	19.5479	6.7108	0.4089	0.9146	5.8424	2.0725	0.8426	57.53
SCSEUNet	0.6243	18.1700	5.8097	0.4538	0.8898	7.0346	2.8565	0.8019	42.53	0.8954	6.4212	2.4971	0.8105	0.8974	6.5153	2.4608	0.8138	0.22
UTNet	0.6470	18.2934	6.1232	0.4779	0.8668	7.1968	2.7462	0.7658	33.97	0.7711	12.8753	4.0784	0.6286	0.8502	7.3667	2.8764	0.7391	10.26
CFANet	0.6501	17.7222	7.2335	0.4818	0.8815	7.0559	2.4861	0.7887	35.6	0.8625	6.8438	2.5889	0.7591	0.8721	6.6657	2.6299	0.7726	1.11
TransUNet	0.6620	17.5615	5.8473	0.4945	0.8503	7.2486	2.7199	0.7392	28.44	0.7639	14.2961	4.1338	0.6178	0.8587	7.2445	2.7498	0.7527	12.41
UNext	0.6727	16.5286	6.5304	0.5063	0.8622	6.8691	2.8336	0.7582	28.17	0.7834	12.4592	3.9107	0.6442	0.8925	6.5413	2.4716	0.8061	13.93
DeepLabv3	0.6924	16.5496	5.3898	0.5294	0.9024	6.4139	2.2435	0.8219	30.33	0.9011	6.4812	2.4076	0.8208	0.9114	5.7941	2.0826	0.8379	1.14
CGNet	0.7147	16.2685	4.9114	0.5562	0.9071	6.4182	2.2781	0.8309	26.92	0.8614	6.7453	2.7558	0.7571	0.9328	5.5925	1.9791	0.8747	8.29
PSPNet	0.7453	14.8329	4.1227	0.5938	0.8091	12.0723	3.5719	0.6795	8.56	0.8342	7.4021	2.7154	0.7159	0.9092	6.4922	2.3012	0.8337	8.99
TinyUNet	0.7498	14.7967	4.2405	0.5995	0.7912	12.4828	3.8143	0.6542	5.52	0.8781	6.6749	2.6857	0.7834	0.9004	6.4887	2.4627	0.8185	2.54
EANet	0.7669	14.2132	4.1877	0.6222	0.9114	6.2923	2.1254	0.8373	18.84	0.8388	7.7428	3.1776	0.7224	0.9141	6.1077	2.1490	0.8416	8.98
DUCKNet	0.7684	14.2412	4.1801	0.6238	0.9122	6.3269	2.1127	0.8386	18.71	0.9115	6.3722	2.2152	0.8367	0.9123	6.1891	2.1926	0.8388	0.09
CPFNet	0.7801	12.6025	4.0218	0.6399	0.9106	6.4693	2.1234	0.8359	16.73	0.9120	6.1170	2.1381	0.8391	0.9177	5.8527	2.0849	0.8474	0.62
I2UNet	0.7963	12.7034	3.5625	0.6616	0.8974	6.5074	2.6047	0.8135	12.70	0.8939	6.4095	2.4645	0.8082	0.9094	6.4713	2.2896	0.8345	1.73
SUNet	0.8017	12.0986	3.6487	0.6693	0.874	7.0476	2.7537	0.7759	9.02	0.9102	6.2244	2.2494	0.8359	0.9104	6.4164	2.2819	0.8354	0.02
DSCNet	0.8345	11.2599	3.1635	0.7163	0.848	7.3059	2.8134	0.7361	1.62	0.8696	6.6449	2.6017	0.7693	0.8774	6.9324	2.5916	0.7815	0.90
SegNet	0.8394	7.6623	3.1598	0.7236	0.8903	6.4261	2.5607	0.8026	6.06	0.8834	6.5274	2.6071	0.7906	0.8861	7.0214	2.7643	0.7955	0.31
MAUNet	0.8467	7.3213	2.7552	0.7341	0.8661	7.1427	2.7153	0.7640	2.29	0.8991	6.4084	2.4267	0.8167	0.9228	5.7887	2.0799	0.8576	2.64
TCSegNet	0.8501	7.2451	2.6984	0.7389	0.8716	7.0703	2.5782	0.7722	2.53	0.8545	7.2355	2.7102	0.7459	0.8589	7.1738	2.6478	0.7534	0.51
UNet	0.8926	6.7326	2.6387	0.8063	0.8957	6.5928	2.438	0.8113	0.35	0.8901	6.464	2.5096	0.8026	0.8993	6.2387	2.4641	0.8169	1.03
ICNet	0.9025	6.6902	2.2249	0.8226	0.9054	6.4716	2.2952	0.8268	0.32	0.9077	6.5199	2.4805	0.8317	0.9293	5.6848	2.0835	0.8688	2.38
UMNet	0.9077	6.4167	2.2079	0.8313	0.9124	6.3704	2.1709	0.8395	0.52	0.9172	6.0197	2.1543	0.8477	0.9185	5.9608	2.1767	0.8487	0.14

5.2 Comparison with Multi-View Segmentation Methods

We selected five representative multi-view medical image segmentation methods for comparison. To ensure fairness in the experiments, all comparison methods used the same UMNet as the backbone network, were trained on exactly the same data split. The quantitative results are shown in TABLE 4. Fig 4 visualizes the percentage improvement for all methods. It can be seen that the DPIF framework demonstrates a consistent performance advantage across all views on the three datasets. In particular, on the long-axis view of the liver dataset, DP-UMNet achieves a 5.77% improvement over the single-path baseline, significantly outperforming all comparison methods. In contrast, on all three test datasets, all other methods underperform DP-UMNet. Most of these methods even performed worse than the single-path UMNet baseline model, with only a few achieving a slight improvement on the liver dataset. Among them, MSMVNet and CoBioNet exhibited the most significant performance declines. Overall, compared to existing multi-view segmentation methods, the DPIF framework achieves stable and consistent performance improvements in multi-view tasks spanning different imaging modalities and anatomical structures.

Table 4: Comparison with Other Multi-View Segmentation Methods.

Method	LV dataset									
	Long Axis					Short Axis				
	DSC	HD95	ASD	IOU	PI(%)	DSC	HD95	ASD	IOU	PI(%)
SP-UMNet	0.8058	12.6179	3.2484	0.6748	–	0.7984	12.5134	3.7446	0.6646	–
AMDANetSu et al. [2025]	0.7241	15.8557	4.9166	0.5680	−10.14	0.7289	15.8725	4.9742	0.5734	−8.70
CoBioNetPeiris et al. [2023]	0.6948	16.2034	4.1474	0.5323	−13.78	0.5041	20.4393	6.8256	0.3371	−36.86
DSMZhu et al. [2025]	0.7351	15.7973	4.5805	0.5812	−8.77	0.7075	16.1543	4.9275	0.5474	−11.39
MSMVNetLiu et al. [2024a]	0.6901	16.5968	5.3385	0.5267	−14.36	0.4905	20.7283	7.1322	0.3249	−38.56
PGRNetHuang et al. [2024]	0.7284	15.8456	4.9452	0.5728	−9.61	0.6754	17.2599	5.5299	0.5097	−15.41
DP-UMNet	0.8519	11.1098	3.1702	0.7421	+5.72	0.8495	11.0153	3.1246	0.7384	+6.40
Method	FL dataset									
	Long Axis					Short Axis				
	DSC	HD95	ASD	IOU	PI(%)	DSC	HD95	ASD	IOU	PI(%)
SP-UMNet	0.9288	5.5999	2.0612	0.8672	–	0.9404	4.4648	2.0467	0.8875	–
AMDANet	0.9519	3.5387	1.8056	0.9098	+2.49	0.9191	6.3722	2.1324	0.8504	−2.26
CoBioNet	0.9669	3.1477	1.7196	0.9359	+4.10	0.9225	5.6113	2.0414	0.8569	−1.90
DSM	0.9818	2.3618	1.4945	0.9648	+5.71	0.9378	5.3893	1.9318	0.8826	−0.28
MSMVNet	0.9222	5.7246	2.0569	0.8562	−0.71	0.8921	6.4842	2.4385	0.8053	−5.14
PGRNet	0.9774	2.4743	1.5172	0.9555	+5.23	0.9438	4.5985	2.0680	0.8935	+0.36
DP-UMNet	0.9824	2.2193	1.4525	0.9659	+5.77	0.9503	3.7747	1.6089	0.9060	+1.05
Method	Hip dataset									
	Long Axis					Short Axis				
	DSC	HD95	ASD	IOU	PI(%)	DSC	HD95	ASD	IOU	PI(%)
SP-UMNet	0.9077	6.4167	2.2079	0.8313	–	0.9172	6.0197	2.1543	0.8477	–
AMDANet	0.8863	6.9053	2.7587	0.7962	−2.36	0.8733	7.1605	2.4956	0.7748	−4.79
CoBioNet	0.7336	15.5838	4.3405	0.5795	−19.18	0.8617	7.2136	2.5765	0.7573	−6.05
DSM	0.9029	6.5943	2.4251	0.8232	−0.53	0.9131	6.2486	2.1561	0.8406	−0.45
MSMVNet-UMNet	0.6300	17.6188	5.2609	0.4615	−30.59	0.6819	16.0451	5.4792	0.5188	−25.65
PGRNet	0.9004	6.6239	2.4178	0.82187	−0.80	0.9031	7.0377	2.3682	0.8235	−1.54
DP-UMNet	0.9124	6.3704	2.1709	0.8395	+0.52	0.9182	1.9608	2.1767	0.8487	+0.11

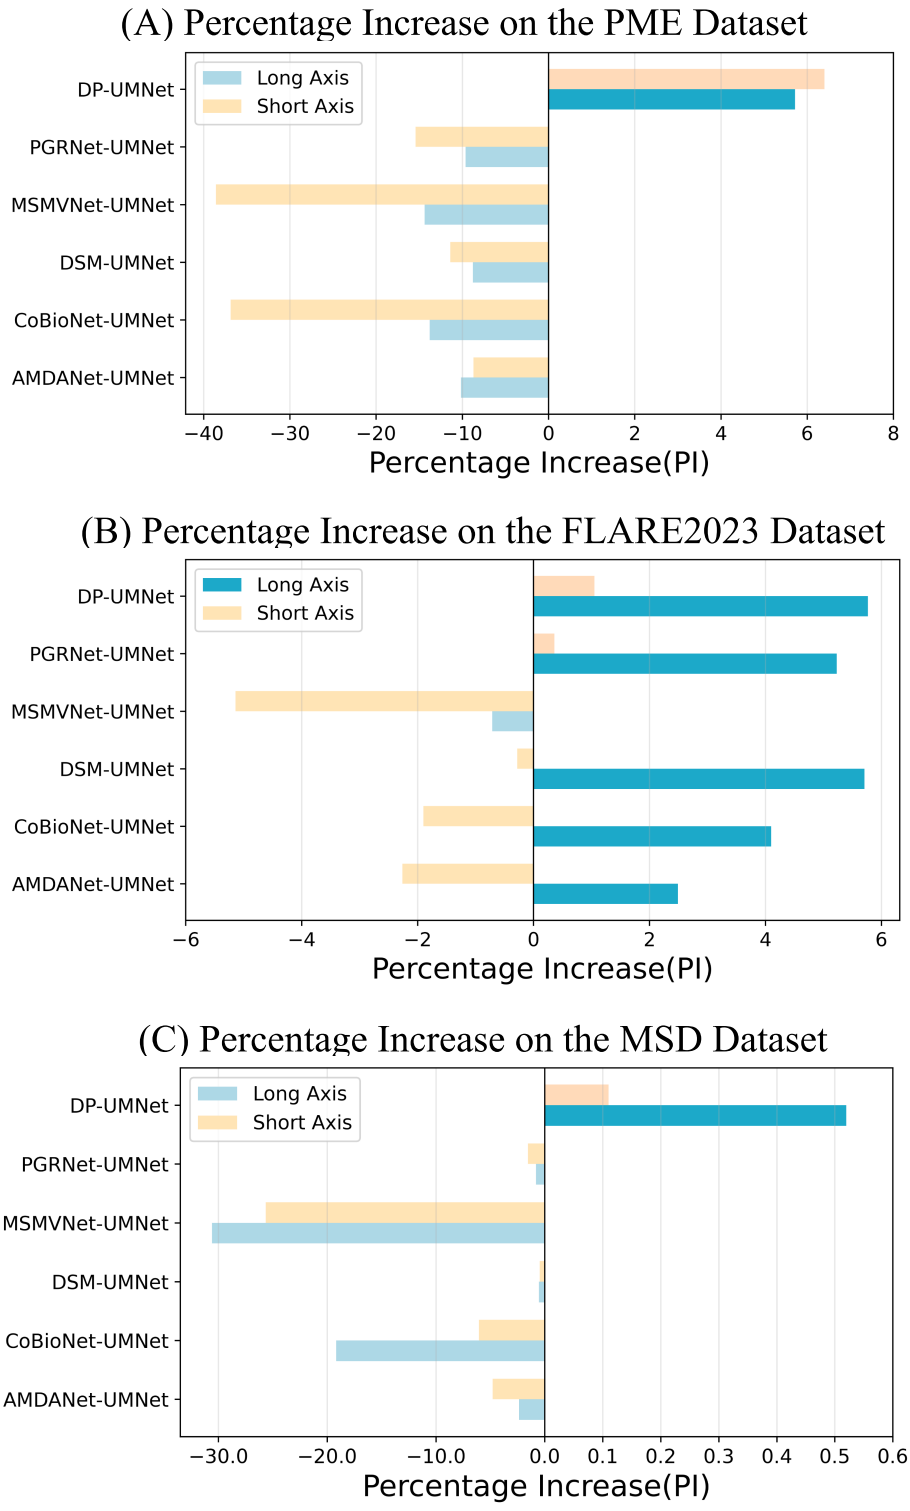

Figure 4: Comparison of Performance Improvements Achieved by the Multi-View Approach Across Three Datasets.

Table 5: Ablation study of the main components in the DPIF framework. The first row indicates the use of only the single-path baseline model.

Method	A	B	C	D	Long Axis					Short Axis				
					DSC	HD95	ASD	IOU	PI(+%)	DSC	HD95	ASD	IOU	PI(+%)
UMNet	×	×	×	×	0.8058	12.6179	3.2484	0.6748	–	0.7984	12.5134	3.7446	0.6646	–
	✓	×	×	×	0.8082	12.4836	3.2157	0.6785	0.30	0.8011	12.3952	3.7024	0.6689	0.34
	×	✓	×	×	0.8245	10.7124	3.0842	0.7018	2.32	0.8192	10.5287	3.4519	0.6953	2.61
	×	×	✓	×	0.8073	12.5471	3.2268	0.6772	0.19	0.8003	12.4506	3.7185	0.6676	0.24
	×	×	×	✓	0.8152	11.2587	3.1395	0.6883	1.17	0.8097	11.1843	3.5426	0.6801	1.42
	×	✓	✓	✓	0.8406	8.2875	2.8053	0.7251	4.32	0.8372	8.3104	2.8479	0.7206	4.86
	✓	×	✓	✓	0.8337	9.0982	2.8704	0.7151	3.46	0.8294	9.1257	2.9318	0.7097	3.88
	✓	✓	×	✓	0.8475	7.6391	2.7468	0.7350	5.18	0.8448	7.6823	2.7695	0.7312	5.81
	✓	✓	✓	×	0.8452	8.3957	2.8164	0.7317	4.89	0.8419	8.4236	2.8351	0.7275	5.45
	✓	✓	✓	✓	0.8519	7.2098	2.7102	0.7415	5.72	0.8495	7.2153	2.7246	0.7379	6.40

5.3 Ablation Studies and Analyses

To quantitatively verify individual contributions and synergies among DPIF components, we perform progressive ablation experiments on the multicenter PME echocardiography dataset with UMNet as the backbone. Taking single-path UMNet as baseline, nine ablation variants are built to evaluate four modules: two-layer feature decoupling (A), Sinkhorn unpaired distribution alignment (B), adaptive interaction weights (C), and cross-view topological consistency regularization (D). Quantitative results are reported in TABLE 5. The results show that decoupling features across two layers yields only a slight improvement, as it achieves feature separation and noise suppression within branches but does not facilitate cross-view information exchange. Sinkhorn-based distribution alignment is the primary driver of performance improvement, increasing DSC by 2% for the long-axis view and 2.24% for the short-axis view, respectively. This validates that distribution-level semantic matching is the primary channel for cross-view knowledge transfer. When used alone, adaptive interaction weights yielded a negligible improvement, indicating that dynamic weighting cannot function effectively without the support of distribution alignment. Cross-view topological regularization independently optimized segmentation boundaries, increasing the long-axis DSC by 1.14% and reducing HD95 by approximately 0.39.

5.4 Analysis of Different Training Data Usage

To investigate the performance and few-shot adaptability of DPIF under different levels of data usage, we conducted gradient data experiments on three multi-view datasets using three representative backbone networks (RollingUNet, UNet and UMNet). Among these, RollingUNet represents a class of networks with poor baseline performance, UNet represents a class of classic segmentation networks, and UMNet represents a class of networks with superior baseline performance. We employed a stratified inclusive sampling method, stratified by multicenter origin and lesion proportion, to construct training subsets of 20%, 40%, 60%, and 80%. Smaller subsets were fully nested within larger ones to isolate data volume as the sole variable while preserving the statistical distribution of the full dataset. All experiments used a fixed test set and five-fold cross-validation. The corresponding performance curves are illustrated in Fig 5. The results show that the performance of the dual-path model increases monotonically as the amount of training data increases. The performance improvement is most significant in the range from 20% to 40%. In most cases, the DPIF-enhanced model can achieve the same accuracy as the full-data single-path model even when the annotation rate is below 60For weaker backbone networks, DPIF delivers greater performance improvements. DP-RollingUNet, trained on only 20% of the labeled data, outperformed the full-

data single-path baseline model on the PME short-axis view. Cross-modal analysis indicates that the benefits of few-shot learning are positively correlated with task difficulty. In the PME echocardiography task, DP-UMNet outperforms full-data single-path models with a 60% annotation rate. On the high-contrast FLARE2023 liver CT dataset, it requires only 40% annotated data. On a high-resolution hippocampal MRI dataset, DP-UMNet requires a 60% annotation rate to approach the performance of the full-data single-path model.

5.5 Blind Review of Segmentation Results by Clinicians

To verify whether the improvements achieved through dual-pathway interaction could translate into clinically usable segmentation results, we conducted a blinded review study across three datasets. Three ultrasound specialists and three radiology specialists evaluated the segmentation results using a 1–5-point Likert scale for 30 randomly selected cases from each view. For each case, the original image and the segmentation results from both DP-UMNet and SP-UMNet were displayed simultaneously, with the positions of the results randomized to avoid subjective bias. The results are shown in Table6. The advantages of DP-UMNet were most pronounced in the sections with the weakest baseline performance: in the PME short-axis section, DP-UMNet raised the average score to 3.49 $\pm$ 0.60, and the proportion of cases with a score of $\geq$4 increased from 24.4% to 52.2%. In cases where single-path baseline models had already achieved high clinical scores (e.g., FLARE2023 short-axis and MSD short-axis), the dual-path enhancement did not yield further significant improvements. The inter-reader consistency (ICC) across views ranged from moderate to high. These results confirm that the introduction of DPIF not only improves the agreement between segmentation results and ground-truth labels but also demonstrates that DPIF achieves consistent clinical validation.

Table 6: Comparison of subjective scores and inter-rater reliability for DP and SP across datasets and views.

Dataset	View	Mean±SD(DP)	Mean±SD(SP)	Wilcoxon p (n=30)	Cliff's δ	Mixed-effects β (p)	DP-ICC(2,3)	SP-ICC(2,3)
PME	Long	3.90±0.85	3.79±0.97	0.647	0.17	0.11 (0.341)	0.857	0.830
	Short	3.49±0.60	3.07±0.67	0.004	0.57	0.42 (<0.001)	0.767	0.687
FLARE	Long	4.40±0.82	4.00±0.89	0.035	0.40	0.40 (<0.001)	0.837	0.717
	Short	4.23±0.85	4.21±0.79	0.917	0.10	0.02 (0.819)	0.815	0.813
MSD	Long	4.23±0.81	3.96±0.70	0.049	0.47	0.28 (0.001)	0.811	0.806
	Short	4.21±0.74	4.14±0.77	0.519	0.00	0.07 (0.465)	0.674	0.632

6 Discussion

In the comparative experiment, traditional multi-view methods cannot guarantee consistent performance improvements and sometimes even underperform compared to single-view baseline models, particularly on noisier, unpaired data. Thanks to its multiple functional modules—particularly feature alignment and two-layer decoupling, DPIF mitigates the effects of noise and differences in interaction targets. As a result, it continues to demonstrate performance improvements even on challenging datasets such as ultrasound and MRI.

Gradient data experiments confirm that DPIF improves data utilization efficiency. The magnitude of performance improvement increases as task difficulty and baseline model performance decrease. For low-performance backbone networks such as RollingUNet, DPIF delivers significant improvements, indicating its immense potential in the field of rare diseases where annotations are scarce. Notably, performance continues to improve steadily even when the data usage is 100%.

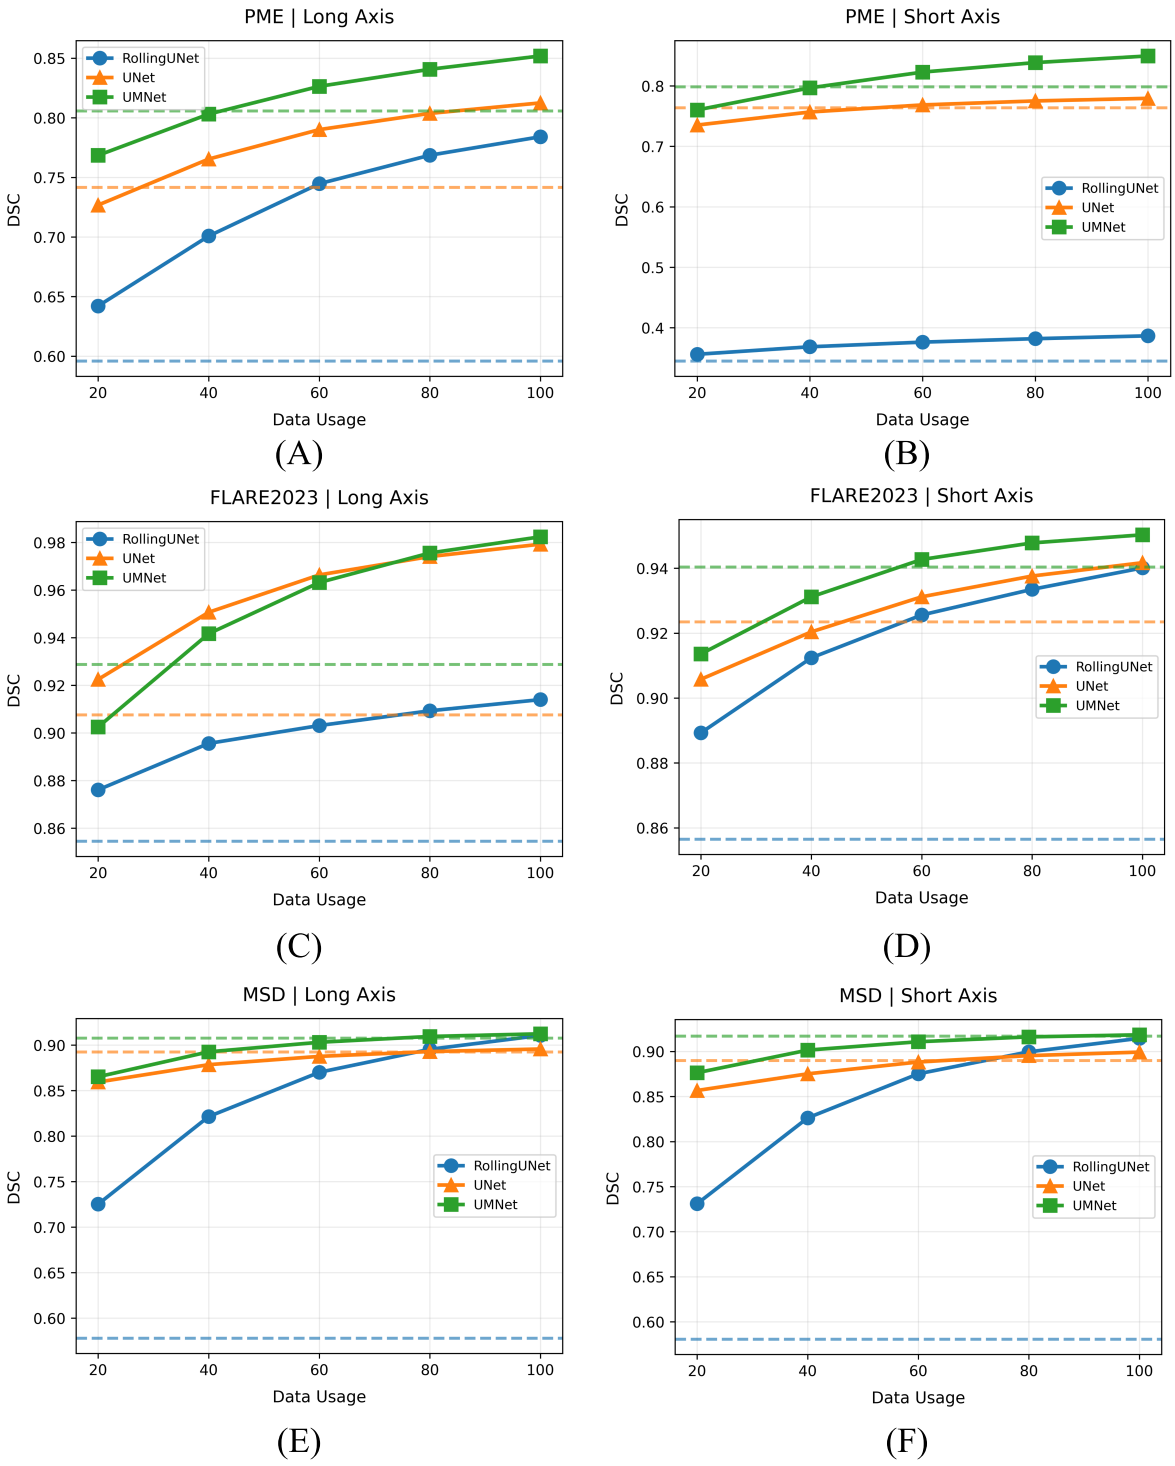

Figure 5: Ablation analysis regarding the amount of training data. The dashed line represents the model's DSC for training on the entire dataset in a single-path pipeline.

Therefore, DPIF serves not only as a solution for few-shot learning but also as a general-purpose enhancement module for fully supervised multi-view clinical pipelines.

Nevertheless, the framework presented in this paper has some limitations. The current implementation supports only two-view interactions, extending it to three or more views would facilitate the volumetric segmentation of 3D data. Topological regularization employs a differentiable persistent homology approximation, which may overlook subtle topological features. Therefore, a more precise differentiable topological solver is needed.

7 Conclusion

This paper proposes a general-purpose, plug-and-play dual-path interaction framework called DPIF for multi-view medical image segmentation, which improved data utilization when data is limited. DPIF achieves consistent performance improvements across 24 encoder-decoder backbone networks and outperforms existing multi-view segmentation methods. It can match or exceed the performance of single-path models trained on the full dataset using only a subset of data of a certain size. This demonstrates that DPIF provides a practical and scalable solution for multi-view medical image segmentation under data-scarce conditions.

References

Vijay Badrinarayanan, Alex Kendall, and Roberto Cipolla. Segnet: A deep convolutional encoder-decoder architecture for image segmentation. *IEEE transactions on pattern analysis and machine intelligence*, 39(12):2481–2495, 2017.

Kamal Berahmand, Fatemeh Daneshfar, Maryam Rahmaninia, Maryam Haghighat, and Mahdi Jalili. A comprehensive survey on multi-view classification: Methods, applications, and challenges. *ACM Transactions on Intelligent Systems and Technology*, 16(6):1–34, 2025.

Xunhang Cao, Xiaoxin Guo, Guangqi Yang, Qi Chen, and Hongliang Dong. Multi-view cross-consistency and multi-scale cross-layer contrastive learning for semi-supervised medical image segmentation. *Expert Systems with Applications*, 277:127223, 2025.

Jieneng Chen, Yongyi Lu, Qihang Yu, Xiangde Luo, Ehsan Adeli, Yan Wang, Le Lu, Alan L Yuille, and Yuyin Zhou. Transunet: Transformers make strong encoders for medical image segmentation. *arXiv preprint arXiv:2102.04306*, 2021.

Jingkun Chen, Haoran Duan, Xiao Zhang, Boyan Gao, Vicente Grau, and Jungong Han. From gaze to insight: Bridging human visual attention and vision language model explanation for weakly-supervised medical image segmentation. *IEEE Transactions on Medical Imaging*, 2025a.

Junren Chen, Rui Chen, Wei Wang, Junlong Cheng, Lei Zhang, and Liangyin Chen. Tinyu-net: Lighter yet better u-net with cascaded multi-receptive fields. In *International Conference on Medical Image Computing and Computer-Assisted Intervention*, pages 626–635. Springer, 2024.

Liang-Chieh Chen, Yukun Zhu, George Papandreou, Florian Schroff, and Hartwig Adam. Encoder-decoder with atrous separable convolution for semantic image segmentation. In *Proceedings of the European conference on computer vision (ECCV)*, pages 801–818, 2018.

Wenna Chen, Chuanqi Cai, Xinghua Tan, Rongfu Lv, Jincan Zhang, and Ganqin Du. Maunet: a mixed attention u-net with spatial multi-dimensional convolution and contextual feature calibration for 3d brain tumor segmentation in multimodal mri. *Frontiers in Neuroscience*, 19:1682603, 2025b.

Duwei Dai, Caixia Dong, Qingsen Yan, Yongheng Sun, Chunyan Zhang, Zongfang Li, and Songhua Xu. I2u-net: A dual-path u-net with rich information interaction for medical image segmentation. *Medical Image Analysis*, 97:103241, 2024.

Theekshana Dissanayake, Yasmeen George, Dwarikanath Mahapatra, Shridha Sridharan, Clinton Fookes, and Zongyuan Ge. Few-shot learning for medical image segmentation: A review and comparative study. *ACM Computing Surveys*, 58(1):1–36, 2025.

Feixiang Du, Xiang Wang, Shengkun Wu, Zhongliang Wang, and Joel Than. Recent advances in medical image segmentation by using deep learning algorithms: a brief review. *Discover Applied Sciences*, 2026.

Xiuquan Du, Xuebin Xu, Jiajia Chen, Xuejun Zhang, Lei Li, Heng Liu, and Shuo Li. Um-net: Rethinking icgnet for polyp segmentation with uncertainty modeling. *Medical Image Analysis*, 99:103347, 2025.

Razvan-Gabriel Dumitru and Darius Peteleaza. Using duck-net for polyp image segmentation. *Scientific Reports*, 13, 06 2023. doi: 10.1038/s41598-023-36940-5.

Shuanglang Feng, Heming Zhao, Fei Shi, Xuena Cheng, Meng Wang, Yuhui Ma, Dehui Xiang, Weifang Zhu, and Xinjian Chen. Cpfnet: Context pyramid fusion network for medical image segmentation. *IEEE transactions on medical imaging*, 39(10):3008–3018, 2020.

Hong-Seng Gan, Muhammad Hanif Ramlee, Zimu Wang, and Akinobu Shimizu. A review on medical image segmentation: Datasets, technical models, challenges and solutions. *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery*, 15(1):e1574, 2025.

Yunhe Gao, Mu Zhou, and Dimitris N Metaxas. Utnet: a hybrid transformer architecture for medical image segmentation. In *International conference on medical image computing and computer-assisted intervention*, pages 61–71. Springer, 2021.

Wendong Huang, Jinwu Hu, Junhao Xiao, Yang Wei, Xiuli Bi, and Bin Xiao. Prototype-guided graph reasoning network for few-shot medical image segmentation. *IEEE Transactions on Medical Imaging*, 44(2):761–773, 2024.

Chris Kapulkin and Nathan Kershaw. Redzed: Computing persistent homology by reduction to zero differentials. *arXiv preprint arXiv:2606.06310*, 2026.

Shyam Lal. Tc-segnet: Robust deep learning network for fully automatic two-chamber segmentation of two-dimensional echocardiography. *Multimedia Tools and Applications*, 83(2):6093–6111, 2024.

Heng Li, Mingyang Ou, Haojin Li, Zhongxi Qiu, Ke Niu, Huazhu Fu, and Jiang Liu. Multi-view test-time adaptation for semantic segmentation in clinical cataract surgery. *IEEE Transactions on Medical Imaging*, 44(5):2307–2318, 2025.

Xiaosen Li, Xiao Qin, Chengliang Huang, Yuer Lu, Jinyan Cheng, Liansheng Wang, Ou Liu, Jianwei Shuai, and Chang-an Yuan. Sunet: A multi-organ segmentation network based on multiple attention. *Computers in Biology and Medicine*, 167:107596, 2023.

Caiqi Liu, Han Liu, Xuehui Zhang, Jierui Guo, and Pengju Lv. Multi-scale and multi-view network for lung tumor segmentation. *Computers in Biology and Medicine*, 172:108250, 2024a.

Huabing Liu, Jiawei Huang, Dengqiang Jia, Qian Wang, Jun Xu, and Dinggang Shen. Transferring adult-like phase images for robust multi-view isointense infant brain segmentation. *IEEE Transactions on Medical Imaging*, 44(1):69–78, 2024b.

Xiuli Liu, Xiangqiong Wu, Yang Chen, and Peng Wang. Scse-unet: A spatial channel squeeze-and-excitation unet for medical image segmentation. In *2024 12th International Conference on Information Systems and Computing Technology (ISCTech)*, pages 1–7. IEEE, 2024c.

Yutong Liu, Haijiang Zhu, Mengting Liu, Huaiyuan Yu, Zihan Chen, and Jie Gao. Rolling-unet: Revitalizing mlp's ability to efficiently extract long-distance dependencies for medical image segmentation. In *Proceedings of the AAAI conference on artificial intelligence*, volume 38, pages 3819–3827, 2024d.

Jun Ma, Yao Zhang, Song Gu, Cheng Ge, Ershuai Wang, Qin Zhou, Ziyan Huang, Pengju Lyu, Jian He, and Bo Wang. Automatic organ and pan-cancer segmentation in abdomen ct: the flare 2023 challenge. *arXiv preprint arXiv:2408.12534*, 2024.

Tomáš Majerník, Dávid Truhlář, Tomáš Brček, Martin Komák, Martin Bobák, and Fedor Lehocki. Deep learning approaches for brain tumor segmentation: A comparative study of multi-class pituitary region and multi-view binary segmentation. In *2026 Information Technology Scientific Research Conference (IIT. SRC)*, pages 1–6. IEEE, 2026.

Himashi Peiris, Munawar Hayat, Zhaolin Chen, Gary Egan, and Mehrtash Harandi. Uncertainty-guided dual-views for semi-supervised volumetric medical image segmentation. *Nature Machine Intelligence*, 5(7):724–738, 2023.

Yaolei Qi, Yuting He, Xiaoming Qi, Yuan Zhang, and Guanyu Yang. Dynamic snake convolution based on topological geometric constraints for tubular structure segmentation. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 6047–6056. IEEE, 2023.

Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical image computing and computer-assisted intervention*, pages 234–241. Springer, 2015.

Amber L Simpson, Michela Antonelli, Spyridon Bakas, Michel Bilello, Keyvan Farahani, Bram Van Ginneken, Annette Kopp-Schneider, Bennett A Landman, Geert Litjens, Bjoern Menze, et al. A large annotated medical image dataset for the development and evaluation of segmentation algorithms. *arXiv preprint arXiv:1902.09063*, 2019.

Yuting Su, Peng Ma, Weiming Wang, Shaochu Wang, Yuting Wu, Yun Li, and Peiguang Jing. Amdanet: Augmented multiscale difference aggregation network for image change detection. *IEEE Transactions on Geoscience and Remote Sensing*, 63:1–12, 2025.

Jeya Maria Jose Valanarasu and Vishal M Patel. Unext: Mlp-based rapid medical image segmentation network. In *International conference on medical image computing and computer-assisted intervention*, pages 23–33. Springer, 2022.

Jeya Maria Jose Valanarasu, Vishwanath A Sindagi, Ilker Hacihaliloglu, and Vishal M Patel. Kiu-net: Towards accurate segmentation of biomedical images using over-complete representations. In *International conference on medical image computing and computer-assisted intervention*, pages 363–373. Springer, 2020.

Kun Wang, Xiaohong Zhang, Xiangbo Zhang, Yuting Lu, Sheng Huang, and Dan Yang. Eanet: Iterative edge attention network for medical image segmentation. *Pattern Recognition*, 127: 108636, 2022.

Ziyang Wang and Irina Voiculescu. Triple-view feature learning for medical image segmentation. In *MICCAI workshop on resource-efficient medical image analysis*, pages 42–54. Springer, 2022.

Kazuki Watanabe and Noboru Isobe. Sinkhorn algorithm for sequentially composed optimal transports. *arXiv preprint arXiv:2412.03120*, 2024.

Tianyi Wu, Sheng Tang, Rui Zhang, Juan Cao, and Yongdong Zhang. Cgnet: A light-weight context guided network for semantic segmentation. *IEEE Transactions on Image Processing*, 30: 1169–1179, 2020.

Guoan Xu, Juncheng Li, Guangwei Gao, Huimin Lu, Jian Yang, and Dong Yue. Lightweight real-time semantic segmentation network with efficient transformer and cnn. *IEEE Transactions on Intelligent Transportation Systems*, 24(12):15897–15906, 2023.

Bangcheng Zhan, Enmin Song, Hong Liu, Zhenyu Gong, Guangzhi Ma, and Chih-Cheng Hung. Cfnet: A medical image segmentation method using the multi-view attention mechanism and adaptive fusion strategy. *Biomedical Signal Processing and Control*, 79:104112, 2023.

Qing Zhang and Weiqi Yan. Cfanet: A cross-layer feature aggregation network for camouflaged object detection. In *2023 IEEE International Conference on Multimedia and Expo (ICME)*, pages 2441–2446. IEEE, 2023.

Hengshuang Zhao, Jianping Shi, Xiaojuan Qi, Xiaogang Wang, and Jiaya Jia. Pyramid scene parsing network. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 2881–2890, 2017.

Hengshuang Zhao, Xiaojuan Qi, Xiaoyong Shen, Jianping Shi, and Jiaya Jia. Icnet for real-time semantic segmentation on high-resolution images. In *Proceedings of the European conference on computer vision (ECCV)*, pages 405–420, 2018.

Zhengchao Zhou, Pingping Wang, Xinggui Ji, Wanbo Xu, Zhongyi Han, and Benzheng Wei. From contrast-driven segmentation to central lumbar spinal stenosis grading: a comprehensive multi-view spinal mri image analysis. *IEEE Transactions on Medical Imaging*, 2026.

Yazhou Zhu, Shidong Wang, Tao Zhou, Zechao Li, Haofeng Zhang, and Ling Shao. Cross-domain few-shot medical image segmentation via dynamic semantic matching. *IEEE Transactions on Image Processing*, 2025.