arXiv:2610.34662v1 [cs.CV] 26 Oct 2026

StreamVibe: Proactive Video Agents for Evolving User Intent

Chengzhi Liu[1,2,*], Yuxiang Lai[1,6], Yuzhe Yang[2], Yanqing Liu[3], Jiheng Hou[2], Yue Cao[2], Haonan Wang[7], Ruixi Fan[1], Lan Wang[1], Sheng Liu[4], Songyou Peng[1], Yale Song[1], Kevin Qinghong Lin[5], Xin Eric Wang[2], Jui-Hsin Lai[1] and Youbao Tang[1]

[1] Google, [2] University of California, Santa Barbara, [3] University of California, Santa Cruz, [4] Stanford University, [5] University of Oxford, [6] Emory University, [7] Columbia University

Vision-language models are shifting from offline passive understanding toward online proactive interaction. However, prevailing methods largely rely on explicit requests while miss the implicit user intent (*e.g.,*"where should I click next?" while a user navigate a website), meanwhile fail to update the state simultaneously as user requests might update continuously. These limitations make them less suited to intent-aware interaction in practical. To address these challenges, we introduce StreamVibe-Bench, a new benchmark serving as testbed for user intent interaction, with 4.1K evaluation samples and 85K training samples spanning explicit, implicit, and dynamic intents. Building on this benchmark, we develop StreamVibe-VL, a streaming proactive model that infers implicit user intents from continuous observations and dynamically tailors its interactions as these intents evolve. We further introduce StreamVibe-Harness, which maintains user intents as persistent states with a fleixble lifecycles via a hierarchical memory to track the latest user intent and environmental state. Our experiments reveal three key findings: *(i)* StreamVibe-VL preserves strong general video understanding while improving response timing and answer quality by 10.7% over other baselines; *(ii)* existing proactive models remain limited in understanding user intent and deciding when to respond, especially under implicit and dynamic intents; and *(iii)* StreamVibe-Harness further improves response consistency and adaptation over sustained interaction, resulting in a 12.3% gain in proactive interaction.

stream-vibe.github.io

Figure 1 | **Overview of StreamVibe.** (a) StreamVibe studies sustained proactive interaction under explicit, implicit, and dynamically evolving user intents, where the agent continuously observes the environment and decides when and how to assist. (b) Existing paradigms rely on explicit intents, with limited tracking of interaction state and evolving intents. StreamVibe addresses these limitations through proactive streaming policy learning and persistent interaction state modeling that tracks task progress and updates active intents.

* This work was done while Chengzhi Liu was a Student Researcher at Google.
Corresponding author(s): chengzhi@ucsb.edu

1. Introduction

"Mind is activity, a range of dynamic processes of interaction between organism and world."

— Stanford Encyclopedia of Philosophy, John Dewey

Vision-language models (Lu et al., 2026; Wang et al., 2026) are advancing from offline content understanding (Wang et al., 2025c; Zhu et al., 2025) toward continuous proactive assistance in real-world streaming scenarios (Cui et al., 2026; Ding et al., 2025; Orantqing et al., 2026; Wang et al., 2025d). However, existing proactive interaction approaches are largely designed for explicit queries (e.g. "Notify me when a goal is scored."), and isolated, single-turn requests (Kundu et al., 2026; Li et al., 2026a), as illustrated in Figure 1(i).

However, these paradigm is less suited to real-world open-ended settings across three challenges, where ***(i)*** *implicit user intents* (e.g. "Help me finish this task whenever I get stuck.") must be inferred from ongoing observations and interactions. Beyond identifying user intent, effective assistance requires ***(ii)*** *task-relevant state tracking*. As shown in Figure 1(ii), prior interactions can change the environment and task progress, making a response that was appropriate earlier unnecessary or incorrect later. Although existing streaming models can retain long interaction histories (Wang et al., 2025a; Yao et al., 2026a), they still struggle to maintain an up-to-date state, resulting in repeated, outdated, or missed responses. Lastly, as shown in Figure 1(iii), ***(iii)*** *user intents is evolving* as the interaction unfolds (Tack et al., 2026; Tao et al., 2026; Yang et al., 2026d). Later feedback may introduce new goals, revise existing constraints, or shift the current focus, while existing methods often fail to update which intents remain relevant and instead continue to rely on earlier requests. These challenges make proactive interaction more than a problem of real-time long-context modeling: it requires the agent to continuously infer and update user intent while tracking the evolving interaction state. This motivates three key research questions forward "**Intent-aware interaction**":

How can we construct a testbed for proactive interaction with implicit user intent?
How can such intent-aware interaction be internalized as a native model capability?
How to flexible manage context to track environmental and user intent changes?

In this work, we address these three questions through benchmarking, modeling, and agentic design. Firstly, we introduce StreamVibe-Bench, a unified intent-aware benchmark for sustained proactive interaction with 4.1K evaluation samples and 85K training samples spanning explicit, implicit, and dynamic user intents. Building on this benchmark, we further develop StreamVibe-VL, a streaming proactive model that infers implicit user intents and adapts its interactions as these intents evolve. Streaming supervised fine-tuning (SFT) jointly learns when and how to assist, while localized trajectory reinforcement learning (RL) focuses optimization on interaction-critical moments to improve response timing and quality. Lastly, we further develop StreamVibe-Harness, which maintains user intents as persistent, updatable states with flexible lifecycles, while hierarchical memory retrieves relevant history to keep subsequent decisions aligned with the current interaction context and user intent. Our main contributions are as follows:

- To the best of our knowledge, we are the first work to study user intents in proactive streaming interaction. To this end, we construct StreamVibe-Bench, a unified benchmark with 4.1K evaluation and 85K training samples spanning explicit, implicit, and dynamic intents.
- We introduce StreamVibe-VL, a proactive streaming video model that jointly learns when and how to assist under implicit and evolving user intents.
- We develop the StreamVibe-Harness, which maintains updatable intent states with flexible lifecycles and hierarchical memory for consistent behavior over sustained interaction.

Table 1 | **Comparison with other proactive interaction benchmarks.** StreamVibe-Bench covers Explicit, Implicit, and Dynamic Intents across both physical and digital environments, and supports both Pointwise and Trajectory evaluation. ✓, ▲, and ✗denote full, partial, and no support, respectively. Dur. denotes the average sample duration in minutes; Evid. and Traj. denote the numbers of evidence points and evaluation trajectories.

Benchmark	User Intent Type			Dur. (min)	Evid.	Environment		Eval Setting		Traj.
	Explicit	Implicit	Dynamic			Physical	Digital	Pointwise	Trajectory	
ProactiveVideoQA (Wang et al., 2025e)	✓	✗	✗	2.1	3.5K	✓	✗	✗	✓	1.4K
EgoProactive (Kundu et al., 2026)	▲	✓	✗	3.5	5.4K	✓	✗	✓	✗	700
EgoPro-Bench (Ran et al., 2026)	✓	✓	✗	3.0	2.4K	✓	✗	✗	✓	2.4K
IPIBench (Li et al., 2026b)	✓	✗	✓	1.0	3.7K	✓	✗	✓	▲	1.8K
EgoServe (Sitong et al., 2026)	▲	✓	▲	6.2	3.4K	✓	✗	✗	✓	293
GuideMe (Liu et al., 2026)	✓	✓	✗	5.5	47.8K	✓	✗	✓	▲	2.5K
GUIDE (Yang et al., 2026c)	✓	✓	▲	33.8	0.7K	✗	✓	✓	✗	120
StreamVibe-Bench (Ours)	✓	✓	✓	**8.3**	**20.5K**	✓	✓	✓	✓	**4.1K**

2. Related Work

Streaming Video Benchmark. Early video benchmarks focus on offline understanding (Fu et al., 2025a; Zhao et al., 2025), while online benchmarks evaluate sequential perception and response before video completion (Huang et al., 2025; Li et al., 2025b). Recent proactive benchmarks further require models to decide when to respond, often using annotated trigger points and extending to egocentric or task-guidance settings (Kundu et al., 2026; Ran et al., 2026; Wang et al., 2025e). However, most focus on a single intent type or environment and evaluate responses independently. StreamVibe-Bench instead unifies explicit, implicit, and dynamic intents across physical and digital tasks, with both pointwise and trajectory evaluation.

Proactive Interaction Models. Online video-language models support continuous perception and low-latency responses (Cui et al., 2026; Zeng et al., 2025), while real-time and task-guidance systems further learn when to speak or provide assistance during ongoing interaction (Liu et al., 2026; Lu et al., 2026; Sitong et al., 2026). However, most existing approaches center on explicitly specified requests or task objectives, with limited support for tracking interaction state and adapting to evolving user intent over time (Tack et al., 2026; Tao et al., 2026). StreamVibe-VL targets open-ended proactive interaction, where intents must be inferred from user goals, environment changes, and task progress, and updated throughout the interaction.

Memory and Agent Harness. Agentic systems increasingly incorporate persistent memory (Xu et al., 2025), planning (Xiong et al., 2025), reflection (Yuan et al., 2025), and reusable skills (Zheng et al., 2025) to preserve and reuse historical context over long-horizon tasks. Recent memory systems further study how to retrieve, integrate, and abstract historical information into reusable experience for long-term interaction (Chhikara et al., 2025; Yu et al., 2026). However, these methods mainly organize and access past information without explicitly tracking evolving user intents and task progress, which can cause proactive responses to drift from the current interaction state. StreamVibe-Harness addresses this gap by maintaining user intents and task progress as persistent, updatable interaction states and retrieving relevant history before each decision.

3. StreamVibe-Bench: Intent-aware Proactive Benchmark

3.1. Task Formulation

We formulate proactive assistance as a causal streaming decision process. The agent observes a video stream $\mathcal{V} = \{v_t\}_{t=1}^{T}$ while the user may express one or more intents i_j over time, explicitly or implicitly,

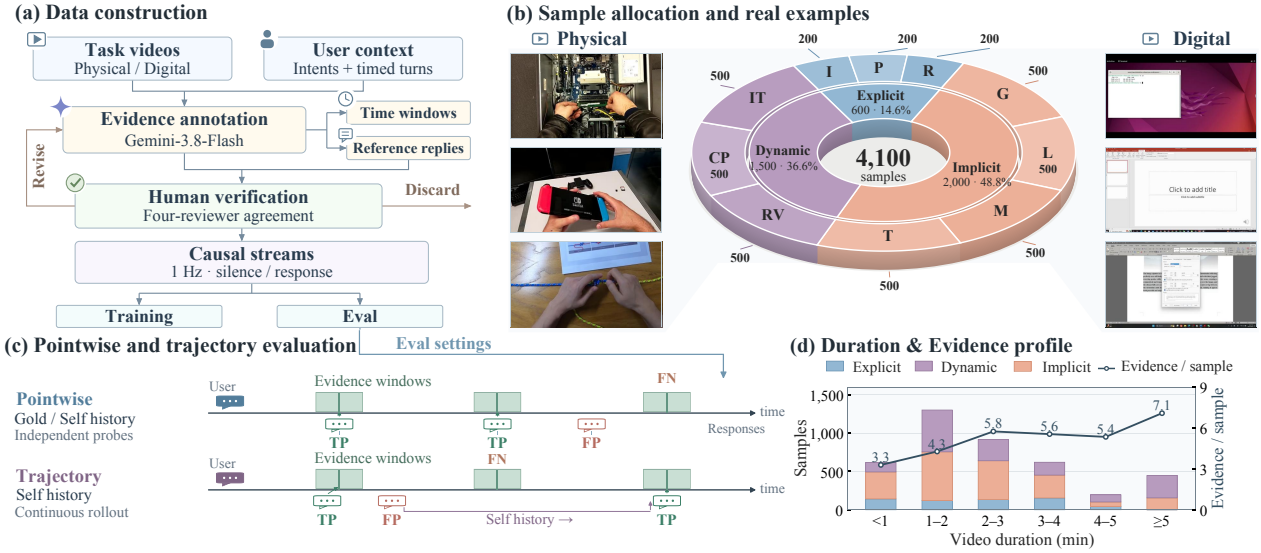

Figure 2 | **Construction and statistics of StreamVibe-Bench.** (a) Data construction from task videos and user context to evidence annotation, human verification, and causal stream generation. (b) Sample allocation across three requirement types, with examples from physical and digital environments. (c) Pointwise and trajectory evaluation settings for local response decisions and continuous interaction behavior. TP: true positives; FP: false positives; FN: false negatives. (d) Joint distribution of video duration, sample count, and average number of evidence points per sample.

possibly before earlier intents are fulfilled. Each intent becomes available to the agent at time τ_j. Let $\mathcal{I}_{\leq t}$ denote all intents available by time t, and $\mathcal{A}_t \subseteq \mathcal{I}_{\leq t}$ the currently active ones. At each timestep, the agent accesses only $c_t = \langle \mathcal{V}_{\leq t}, \mathcal{I}_{\leq t}, \mathcal{H}_{<t} \rangle$ and predicts $a_t = \pi(c_t) \in \{\textsc{silence}, \textsc{response}\}$. If responding, it generates y_t for some $i_j \in \mathcal{A}_t$, without access to future video or user inputs.

Unlike retrospective VideoQA, relevant evidence may emerge only after an intent is expressed. Each intent i_j can correspond to one or more evidence points ξ_{jk}, each with a valid response window $[s_{jk}, e_{jk}]$ and reference response $\hat{y}_{jk}$. An evidence point is covered if the agent responds within its window with a semantically correct response to the corresponding intent. An intent remains active until all its evidence points are covered. The agent must therefore determine both when to respond and what to respond, while avoiding premature, missed, repeated, or irrelevant responses.

3.2. Hierarchical Difficulty

StreamVibe-Bench characterizes proactive interaction along two dimensions: *intent explicitness* and *intent dynamics*. Existing benchmarks primarily focus on **(i) Level 1: Explicit Intents**, where the user specifies a condition, event, or state to monitor, and the agent responds when it occurs, e.g. Let me know when the water starts boiling.

We further introduce **(ii) Level 2: Implicit Intents**, where the user provides only a high-level goal, requiring the agent to infer when and how to assist from ongoing observations and task progress, e.g. Help me keep an eye on this dish.

Finally, **(iii) Level 3: Dynamic Intents** considers intents that evolve throughout the interaction. As the task unfolds, the user may provide multiple new inputs, requiring the agent to continually update the currently valid task state.

We consider three common forms: *Revision (RV)*, which modifies an existing intent; *Composition (CP)*, which adds a new intent or constraint while preserving existing ones; and *Interruption (IT)*, which introduces an immediate request while the ongoing task remains active. These changes may occur repeatedly and in combination within a single interaction.

3.3. Benchmark Construction

Task Collection. We collect task videos from both physical and digital environments to cover two complementary forms of proactive interaction. In (i) *physical* environments, assistance mainly relies on understanding continuous visual processes, including user activities, task progress, and emerging intents. In contrast, user intent in (ii) *digital* environments is often implicit in interaction trajectories such as navigation, content editing, and tool use, requiring the agent to infer user goals, constraints, and task state from the current interface and interaction history. The two settings thus capture user intent through physical behavior and digital interaction, respectively, covering proactive assistance across different interaction environments. Specifically, physical tasks mainly include egocentric daily activities and task execution, while digital tasks focus on GUI and computer-use scenarios.

Annotation and Verification. For each video-request pair, we use Gemini-3.8-Flash (Google DeepMind, 2026) to clean the data and identify candidate help points with response windows, preferred times, and reference answers. For dynamic intents, we further construct timestamped multi-turn interactions in which intents may be revised, extended, or interrupted as the task unfolds. Four reviewers then independently verify the user requests and response timing, correcting or removing samples without sufficient evidence and retaining only those approved by all reviewers. The validated annotations are aligned to 1 FPS streams and converted into causal interaction sequences for training and evaluation. Details are provided in Appendix B.3.

3.4. Data Statistics

Categories and Scale. Building on the three proactive interaction settings in Section 3.2, we organize StreamVibe-Bench into finer-grained task subtypes, as illustrated in Figure 2(b). The 600 *Explicit Intents* samples span Single Event (I), Transition (P), and Recurring Event (R). For *Implicit Intents*, 2,000 samples cover Guidance (G), Learning (L), Monitoring (M), and Tool (T). The remaining 1,500 *Dynamic Intents* samples consist of Revision (RV), Composition (CP), and Interruption (IT). Subtypes are evenly distributed within each category. In total, StreamVibe-Bench contains 4,100 evaluation samples, 20,546 help points, and 183.4 hours of video. Additional statistics and detailed explanations are provided in the Appendix B.

Training Data Curations. We construct the StreamVibe-85K training set, consisting of 70K samples for supervised fine-tuning and 15K tasks for reinforcement learning. During supervised fine-tuning, we further incorporate 768K general streaming-video samples from sources including TimeChat-Online (Yao et al., 2025) and VideoChat-Flash (Li et al., 2025a) to improve general streaming-video modeling. Detailed components are provided in Appendix B.4.

3.5. Evaluation Protocol

Settings. As illustrated in Figure 2(c), two complementary evaluation settings are adopted to distinguish local proactive decision-making from sustained interaction behavior.

(i) Pointwise Evaluation samples positive and negative timestamps around annotated help points and independently predicts whether the model should respond. Both *Gold History* and *Self History* are considered, where the former isolates local decision ability and the latter measures the effect of model-generated history. Independent evaluation avoids error accumulation across time.

(ii) Trajectory Evaluation runs the model continuously over the full video stream, with each decision conditioned on prior observations and self-generated interaction history. This setting captures error propagation, requirement tracking, and repeated or missed responses. Together, the two settings evaluate local decisions under a given history and sustained interaction under self-generated history.

Metrics. We evaluate models along two dimensions: response timing and response quality. Timing is measured by Precision, Recall, F1, and False Positive Rate (FPR), capturing response accuracy, coverage, and false activations. Quality is evaluated using Answer Score for overall correctness and relevance, and Content F1 for key-content alignment with reference answers.

4. StreamVibe: Grounding Evidence, Tracking Intent

Overview. As shown in Figure 3, the formulation in Section 3.1 highlights two key challenges in proactive interaction. **(i) Evidence Grounding.** The policy π must ground each active intent i_j in emerging visual evidence ξ_{jk} and determine when and how to respond. We address this with a dedicated training framework that directly improves the streaming video policy. **(ii) State Tracking.** As intents evolve, maintaining an up-to-date active set $\mathcal{A}_t$ and interaction history $\mathcal{H}_{<t}$ becomes challenging for the policy alone. We therefore introduce the StreamVibe-Harness ϕ, which explicitly maintains these interaction states and provides π with continuously updated context, $\tilde{c}_t = \phi(c_t, \mathcal{A}_t)$. This separates video-grounded response decisions from persistent state tracking: the Harness tracks dynamic intents and task progress, while response decisions and generation remain with π. Details of the framework and workflow are provided in Appendix C.

4.1. StreamVibe-Harness Framework

Interaction State. The Harness maintains an explicit interaction state $\mathcal{S}_t = \langle \mathcal{A}_t, \mathcal{P}_t, \mathcal{R}_t \rangle$, where $\mathcal{A}_t$ denotes the currently active intents, $\mathcal{P}_t$ captures task progress and environment state, and $\mathcal{R}_t$ records previously delivered responses. As new user inputs and observations arrive, the Harness incrementally updates $\mathcal{S}_{t-1}$ to $\mathcal{S}_t$ by adding, revising, or terminating intents and updating task progress. State changes are applied only when supported by sufficient interaction or perceptual evidence, avoiding premature updates from uncertain observations.

Intent Lifecycle. Each intent i_j is maintained as an independent stateful instance, allowing multiple intents to remain active and evolve concurrently. The Harness tracks its current status, relevant evidence, and completion state, and updates the lifecycle as the interaction unfolds. Persistent intents may remain active across multiple events, whereas one-time intents are closed once satisfied. Importantly, the lifecycle only determines which intents remain valid; when and how to respond are still decided by the agent policy π.

Hierarchical Memory. The Harness organizes past observations and interactions into hierarchical memory $\mathcal{M}_{<t}$, preserving relevant history without continuously expanding the agent context. Recent information is retained in fine-grained form, while older content is progressively compressed to preserve task-relevant entities, actions, constraints, and state changes. Given the active intents $\mathcal{A}_t$ and task state $\mathcal{P}_t$, the Harness retrieves relevant history as $m_t = \text{Retrieve}(\mathcal{M}_{<t}, \mathcal{A}_t, \mathcal{P}_t)$, prioritizing recent and task-relevant information. The retrieved memory serves as additional evidence for the current decision without directly modifying the persistent interaction state.

State Integration. Before each decision, the Harness integrates the persistent interaction state and retrieved memory with the causal context to obtain $\tilde{c}_t = \phi(c_t, \mathcal{S}_t, m_t)$, which is then passed to the agent policy π for response decision and generation. After a response is delivered, the resulting interaction is written back to the state and memory for subsequent decisions.

4.2. Evidence-Grounded Policy Learning for Streaming Interaction

***Stage 1*: Adaptive Streaming SFT.** We introduce adaptive streaming supervision, which adjusts decision-token weights to the current training distribution while separately optimizing response

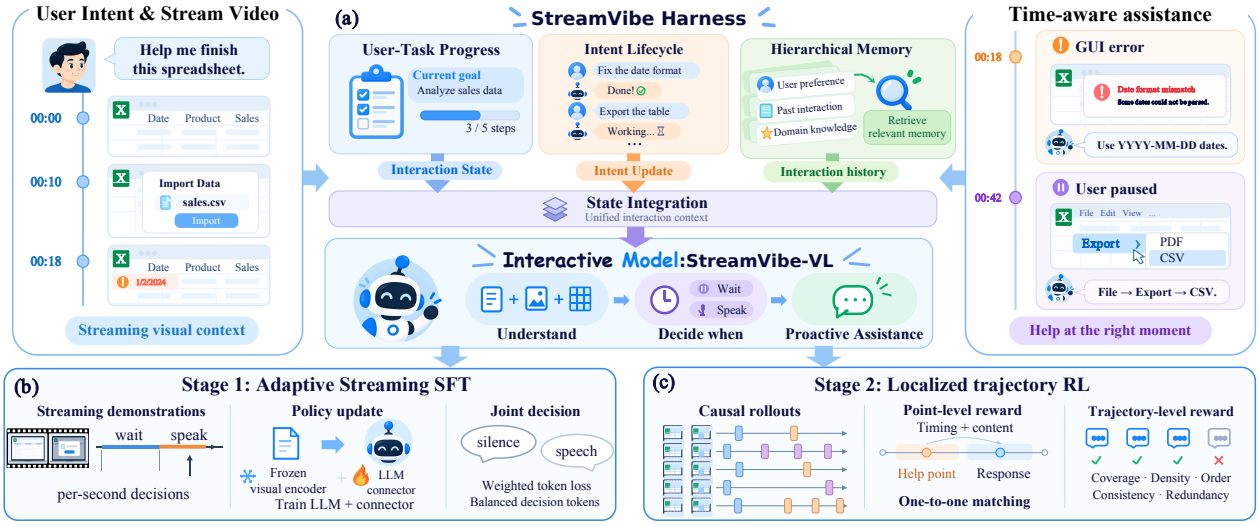

Figure 3 | **Overview of StreamVibe.** (a) During inference, StreamVibe-Harness maintains persistent interaction state, tracks each intent through an independent lifecycle, and retrieves relevant history from hierarchical memory, providing StreamVibe-VL with a unified context for response decisions. (b) Adaptive streaming SFT balances wait/speak decisions and jointly supervises timing and content. (c) Localized trajectory RL samples causal rollouts around help points, combining point-level timing and content rewards with trajectory-level rewards on coverage, density, order, consistency, and redundancy.

generation. This helps balance varying silence/response ratios across training sources and reduces bias toward excessive silence or over-response. Our SFT objective is

$$\mathcal{L}_{\text{SFT}} = -\frac{1}{|\mathcal{Y}|}\sum_{j\in\mathcal{Y}} \omega_j \log \pi_\theta\left(z_j \mid c_{t(j)}, z_{<j}\right), \qquad \omega_j = \begin{cases} \gamma_{B,c(j)}, & j \in \mathcal{A}, \\ \lambda_{\text{text}}, & j \in \mathcal{R}, \end{cases} \tag{1}$$

where $\mathcal{Y}$ denotes supervised assistant tokens, $\mathcal{A}$ and $\mathcal{R}$ denote decision and response tokens, and $c(j) \in \{\text{response}, \text{silence}\}$ is the decision class. The adaptive weight $\gamma_{c(j)}^{(t)}$ follows the current training distribution, while λ_{text} preserves response-generation learning.

***Stage 2*: Localized Trajectory RL.** Building on SFT, we introduce localized trajectory reinforcement learning. Full-trajectory optimization mixes correct and erroneous behaviors, leading to diffuse credit assignment and reward interference. We therefore construct short causal rollouts around annotated help points and optimize them with GRPO (Shao et al., 2024) using complementary point-level and trajectory-level rewards.

(i) Point-Level Reward. For each evidence point g and predicted response p, the temporal offset is defined as $\Delta_{g,p} = t_p - t_g$. The timing score $\tau(\Delta_{g,p}) \in [0, 1]$ measures alignment within an asymmetric response window, while $\rho(g, p) \in [0, 1]$ measures LLM-judged content quality. Their product $\tau(\Delta_{g,p})\rho(g, p)$ serves as the matching score, and one-to-one matching yields the set $\mathcal{M}$, where each evidence point and response can be matched at most once. The matched scores are then aggregated while accounting for missed evidence points and excessive responses.

$$R_{\text{align}} = \frac{(1+\alpha)\sum_{(g,p)\in\mathcal{M}} \tau(\Delta_{g,p})\rho(g,p)}{N_{\text{pred}} + \alpha N_{\text{gt}}} \tag{2}$$

Here, N_{pred} and N_{gt} denote the numbers of predicted responses and evidence points, respectively, while α controls the trade-off between missed evidence and over-response.

(ii) Trajectory-Level reward. We further introduce a trajectory-level reward R_{glob} to evaluate each rollout as a whole. Since point-level reword alone does not ensure coherent interaction, R_{glob} captures the overall consistency and naturalness of the response sequence, discouraging redundant or conflicting responses and favoring coherent streaming behavior. The final reward is defined as $R = \lambda R_{\text{align}} + (1 - \lambda) R_{\text{glob}}$, with $\lambda = 0.7$ to place greater emphasis on response timing and content quality while preserving trajectory-level coherence. Since response formatting is already well learned during SFT, the RL stage optimizes only these two objectives without introducing an additional format reward.

5. Experiments

In this section, we conduct extensive experiments to answer the following research questions:

- ♠ **RQ1**: How top-performing baseline behave and can StreamVibe-VL outperform them in intent-aware interaction setting?
- ♠ **RQ2**: What limitations do current proactive agents exhibit in sustained proactive interaction?
- ♠ **RQ3**: What gains does the StreamVibe-Harness provide beyond the end-to-end policy?

5.1. Setup

Benchmarks. We evaluate StreamVibe across two settings: ***(i) Proactive Interaction*** on StreamVibe-Bench, General Timing (Yao et al., 2026b), EgoProactive (Kundu et al., 2026), and IPIBench (Li et al., 2026b); and ***(ii) General Video Understanding*** on MVBench (Li et al., 2024), Video-MME (Fu et al., 2025a), OVO-Bench (Li et al., 2025b), and OVBench (Huang et al., 2025). External suites retain their official protocols. Further details are provided in Appendix A.

Baselines. We benchmark StreamVibe against representative models from the three categories used in Table 2: *(i) Proprietary Models*, including Gemini-3.1-Pro and Gemini-3.8-Flash (Google DeepMind, 2026), and GPT-Realtime-2.1 (OpenAI, 2026); *(ii) Open-Source Offline Models*, including Qwen3-VL-8B (Bai et al., 2025), Qwen3.5-9B (Qwen Team, 2026), and JoyAI-VL-Interaction (Yao et al., 2026b); and *(iii) Open-Source Streaming Models*, including MMDuet2 (Wang et al., 2025d), Mage-VL (Yang et al., 2026d), LiveStar (Yang et al., 2025), ViSpeak (Fu et al., 2025b), and VideoLLM-Online (Chen et al., 2024).

Evaluation Settings. For StreamVibe-Bench, we use the protocol and metrics defined in Section 3.5, with Claude Sonnet 4.6 (Anthropic, 2026) as a shared judge and response quality scored on a 0–4 scale. For external benchmarks, we strictly follow their official evaluation protocols and metrics. See Appendix B.5 for evaluation details.

5.2. Analysis

♠ **RQ1:** ***StreamVibe-VL achieves stronger proactive interaction while preserving general video understanding.*** Table 2 shows that StreamVibe-VL consistently outperforms open-source baselines on StreamVibe-Bench. Under Pointwise Evaluation, it achieves the highest Content F1 while remaining competitive with proprietary models in response quality. RL further improves over SFT, with larger gains on dynamic intents. This trend persists under Trajectory Evaluation, where RL yields higher Content F1 and response quality, indicating stronger sustained interaction when decisions depend on prior history. The gains also transfer to external benchmarks. Figure 4(a-b) shows competitive results on EgoProactive and consistent gains over other models on IPIBench. Detailed results and analysis are provided in Appendix A.

Table 2 | Main results on StreamVibe-Bench (±2 s tolerance). *Pointwise evaluation* scores each decision given the gold interaction history, whereas *trajectory evaluation* runs the model over the full stream with its self-generated history. Streaming models do not support pointwise evaluation and are thus reported only under trajectory evaluation. **bold** and underline mark the best and second-best results within each block.

Model	Explicit						Implicit						Dynamic					
	P↑	R↑	F1↑	FPR↓	ANS↑	C-F1↑	P↑	R↑	F1↑	FPR↓	ANS↑	C-F1↑	P↑	R↑	F1↑	FPR↓	ANS↑	C-F1↑
Pointwise Evaluation *(gold interaction history)*																		
Proprietary Models																		
Gemini-3.1-Pro	**62.34**	68.52	65.28	**41.40**	45.22	29.52	70.92	85.58	77.56	43.57	67.46	52.32	62.87	84.59	72.13	58.81	51.15	36.89
GPT-Realtime-2.1	57.78	74.09	64.93	64.87	51.10	33.18	69.67	70.95	70.30	29.83	55.92	39.31	**62.99**	84.47	**72.17**	55.37	45.64	32.94
Gemini-3.8-Flash	60.34	77.40	**67.81**	50.69	51.37	**34.84**	**81.55**	80.45	**81.00**	**27.09**	82.63	**66.93**	62.96	84.52	72.16	60.25	**60.56**	**43.70**
Open-Source offline Models																		
Qwen3-VL-8B	51.25	82.28	63.16	89.05	40.91	25.84	56.38	57.85	57.11	83.71	80.51	45.98	54.37	87.56	67.08	93.88	45.01	30.19
Qwen3.5-9B	52.95	**91.36**	67.04	81.45	44.59	29.89	57.42	61.23	59.26	49.17	78.37	46.45	56.74	**90.93**	69.88	78.89	47.50	33.19
JoyAI-VL-Interaction	52.50	90.10	66.34	81.88	40.06	26.58	58.38	79.21	67.22	67.32	81.63	54.87	57.48	72.20	64.00	62.60	45.23	28.95
Ours																		
StreamVibe-VL (SFT)	57.14	62.90	59.88	47.90	49.89	29.88	59.71	85.64	70.36	71.34	**85.92**	60.46	58.56	58.85	58.70	**49.13**	56.68	33.27
StreamVibe-VL (RL)	57.33	64.89	60.88	46.19	**55.32**	33.68	61.59	**86.66**	72.01	67.91	85.11	61.28	58.94	75.86	66.34	62.93	59.25	39.31
Trajectory Evaluation *(self-generated interaction history)*																		
Open-Source offline Models																		
Qwen3-VL-8B	3.56	40.05	6.54	19.30	36.78	2.40	3.99	35.54	7.17	26.46	40.34	2.89	2.88	48.89	5.44	18.64	25.31	1.38
Qwen3.5-9B	9.37	47.62	15.66	19.30	16.63	2.60	7.94	52.48	13.79	14.56	44.98	6.20	4.87	51.65	8.90	20.95	27.67	2.46
JoyAI-VL-Interaction	4.67	**95.44**	8.90	30.87	20.21	1.80	4.39	37.96	7.87	18.87	46.83	3.69	3.15	**76.70**	6.05	64.52	21.64	1.31
Open-Source Streaming Models																		
MMDuet2	19.45	38.65	25.88	7.90	22.71	5.88	15.98	36.50	22.23	4.13	**74.26**	16.51	13.06	35.51	19.10	5.24	39.10	7.47
Mage-VL	20.70	50.25	29.32	6.11	10.23	3.00	8.56	51.94	14.70	5.50	8.19	1.20	11.97	52.88	19.52	6.20	6.10	1.19
LiveStar	14.96	19.95	17.10	4.81	3.60	0.62	6.84	24.95	10.74	6.27	25.43	2.73	7.29	12.92	9.32	2.86	8.05	0.75
ViSpeak	20.66	74.71	**32.37**	12.75	10.19	3.30	8.41	**74.03**	15.10	12.67	62.98	9.51	13.51	72.32	**22.77**	14.19	16.85	3.84
VideoLLM-Online	12.29	5.04	7.15	**1.44**	17.89	1.28	6.33	17.66	9.32	5.32	31.88	2.97	**17.86**	14.75	16.16	**1.28**	8.69	1.40
Ours																		
StreamVibe-VL (SFT)	18.20	29.39	22.48	6.33	38.93	8.75	18.49	28.51	22.43	3.80	61.79	13.86	10.01	43.20	16.25	12.25	38.98	6.34
StreamVibe-VL (RL)	**21.53**	34.48	26.51	2.02	**40.86**	**10.83**	**24.57**	28.20	**26.26**	**1.49**	72.83	**19.13**	14.80	30.65	19.96	5.68	**44.62**	**8.91**

Figure 4(c) further shows that our model outperforms its base model across all general video benchmarks, indicating that proactive policy learning preserves general video understanding.

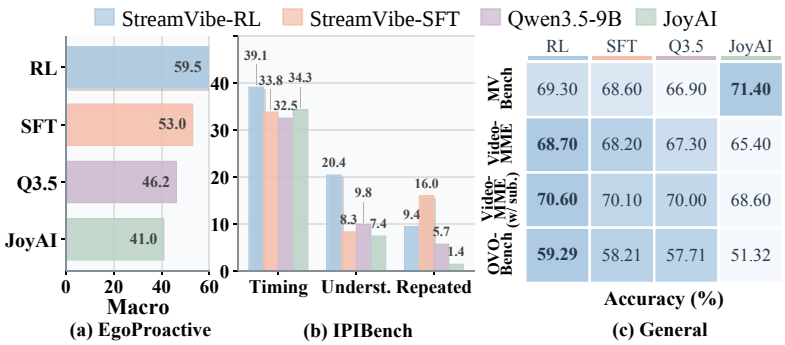

Figure 4 | Generalization to external benchmarks. (a) Macro score on EgoProactive; (b) submetric scores on IPIBench; (c) accuracy on general video benchmarks, with darker cells indicating higher scores within each row.

♠ **RQ2:** ***Current proactive agents struggle to sustain reliable behavior over continuous interaction.*** Figure 5 and Figure 6 reveal three main findings. *(i) Intent understanding remains difficult.* Implicit and Dynamic Intents are markedly harder than Explicit Intents, indicating that agents struggle to determine what help is needed when intents are unstated or keep changing. *(ii) Response selectivity is insufficient.* Higher Recall often comes with more false triggers, suggesting that current methods cannot yet balance coverage against unnecessary responses. *(iii) Sustained interaction amplifies these problems.* Timing F1 declines steadily with interaction length; Revision shows a sharper late-stage Recall drop, and Composition exhibits weaker content quality, indicating difficulty in continuously updating intents while maintaining multiple constraints. Correct individual judgments are thus insufficient for stable proactive behavior throughout interaction.

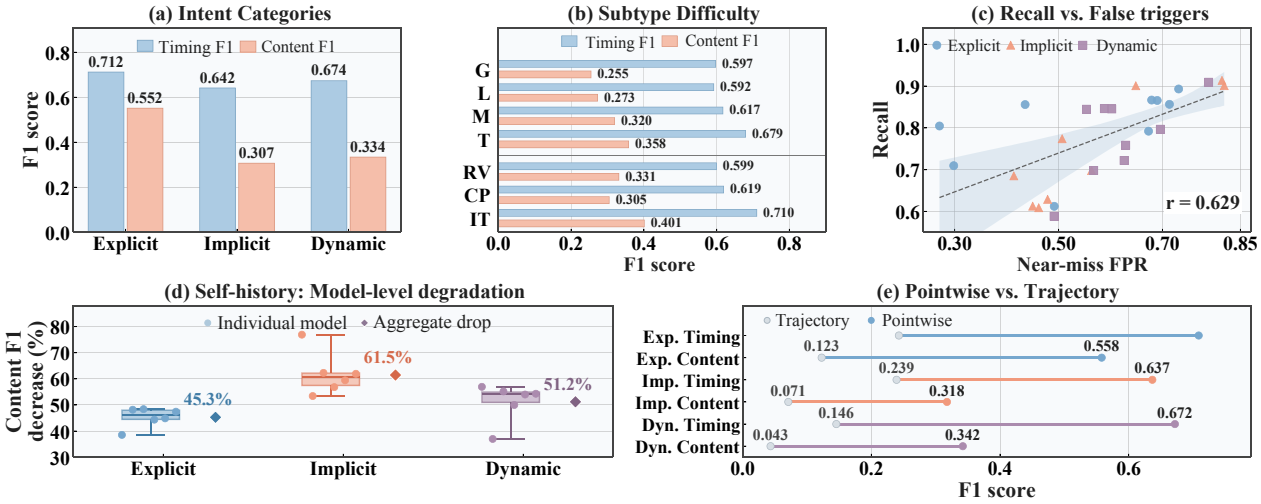

Figure 5 | **Diagnostic analysis of proactive agents.** (a)-(b) Timing and Content F1 across intent categories and subtypes; (c) Recall versus near-miss FPR across models; (d) Content F1 degradation under accumulated self-generated history; (e) pointwise versus trajectory F1 per category and metric.

♠ **RQ3:** ***Policy learning improves proactive decision-making, while state tracking enhances sustained interaction stability.*** As shown in Figure 6(2), SFT makes the model more sensitive to potential assistance opportunities, leading to higher Recall but also more unnecessary triggers. RL further improves decision selectivity, substantially reducing FPR while also improving answer quality. Figure 6(3) further shows that, with the policy fixed, adding the Harness continues to suppress unnecessary responses and improves content quality in most settings, with the clearest gains on Dynamic Intents. These results suggest that policy learning mainly improves local proactive decisions, whereas state tracking helps preserve valid intents and task progress over long-horizon interactions, thereby reducing outdated, redundant, or state-inconsistent outputs.

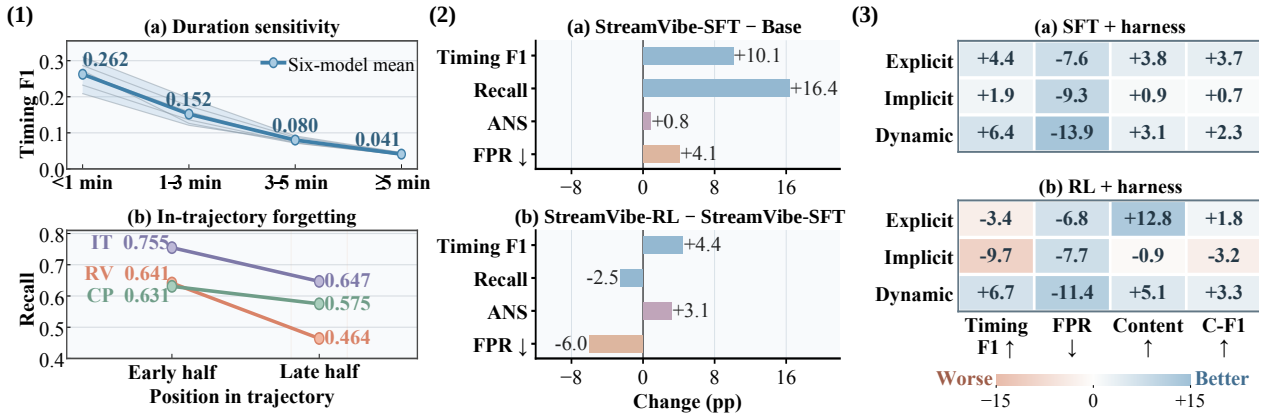

Figure 6 | **Effects of interaction length, policy learning and the harness.** (1) Timing F1 versus video duration, and Recall in the early versus late half of the trajectory for each Dynamic Intent subtype. (2) Absolute metric changes of SFT over the base model and of RL over SFT. (3) Gains from adding the Harness to the SFT and RL policies across intent categories. Tolerance: $\tau = 2$ s.

Qualitative Analysis. Figure 7 illustrates how StreamVibe-VL assists users in representative physical and digital interactions. Across multiple assistance points within the same task, it grounds responses in the current visual evidence and task state, providing assistance at appropriate moments. In contrast, streaming baselines often respond before sufficient evidence appears or fail to update later responses as the task evolves, yielding premature, inconsistent, or redundant assistance. These cases show

that our agent keeps both response timing and content consistent throughout sustained interaction. Additional qualitative examples are provided in Appendix D.2.

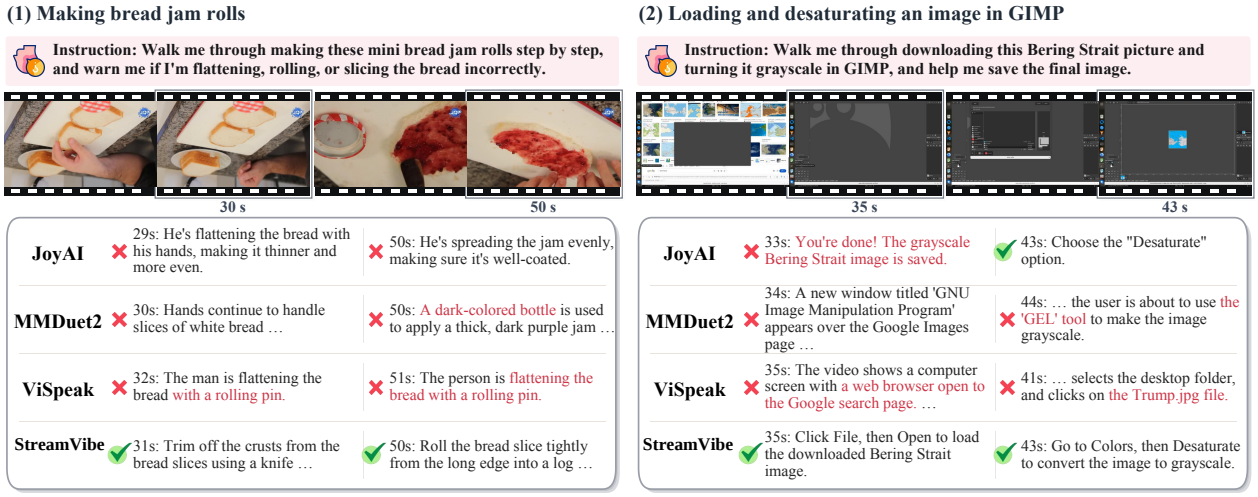

Figure 7 | **Case studies of implicit proactive interaction.** Across physical and digital tasks, StreamVibe-VL better aligns response timing and content with the evolving visual evidence and task state, while streaming baselines more often produce premature, missing, or inconsistent assistance.

6. Conclusion

We study proactive interaction over streaming video, where user intents may be implicit, evolve over time, and must be tracked alongside a changing environment. We introduce StreamVibe-Bench, a unified benchmark spanning explicit, implicit, and dynamic intents with pointwise and trajectory evaluation; StreamVibe-VL, a proactive streaming policy trained with adaptive streaming SFT and localized trajectory RL; and StreamVibe-Harness, which maintains user intents and task progress as persistent interaction states. Experiments show that our model improves response timing and quality while preserving general video understanding, while explicit state tracking further reduces outdated and redundant responses. More broadly, real-world proactive interaction may involve richer intent evolution and tighter coupling between user needs and environmental changes than captured in our current setting. New observations, task outcomes, and prior assistance may jointly reshape both user needs and the environment, leading to new goals and constraints. Future work could further study such co-evolving interactions, particularly the complex feedback between intents and environmental states over longer horizons.

References

Anthropic. Introducing claude sonnet 4.6. Anthropic model announcement, 2026. URL https://www.anthropic.com/news/claude-sonnet-4-6.

S. Bai et al. Qwen3-vl technical report. 2025. URL https://arxiv.org/abs/2511.21631.

M. Bain, A. Nagrani, G. Varol, and A. Zisserman. Frozen in time: A joint video and image encoder for end-to-end retrieval. 2022. URL https://arxiv.org/abs/2104.00650.

F. Caba Heilbron, V. Escorcia, B. Ghanem, and J. C. Niebles. ActivityNet: A large-scale video benchmark for human activity understanding. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2015. URL https://openaccess.thecvf.com/content_cvpr_2015/html/Heilbron_ActivityNet_A_Large-Scale_2015_CVPR_paper.html.

D. Chen, Y. Huang, S. Wu, J. Tang, L. Chen, Y. Bai, Z. He, C. Wang, H. Zhou, Y. Li, T. Zhou, Y. Yu, C. Gao, Q. Zhang, Y. Gui, Z. Li, Y. Wan, P. Zhou, J. Gao, and L. Sun. Gui-world: A video benchmark and dataset for multimodal gui-oriented understanding. 2025a. URL https://arxiv.org/abs/2406.10819.

J. Chen, Z. Lv, S. Wu, K. Q. Lin, C. Song, D. Gao, J.-W. Liu, Z. Gao, D. Mao, and M. Z. Shou. Videollm-online: Online video large language model for streaming video. 2024. URL https://arxiv.org/abs/2406.11816.

J. Chen, Z. Zeng, Y. Lin, W. Li, Z. Ma, and M. Z. Shou. Livecc: Learning video llm with streaming speech transcription at scale. 2025b. URL https://arxiv.org/abs/2504.16030.

P. Chhikara, D. Khant, S. Aryan, T. Singh, and D. Yadav. Mem0: Building production-ready ai agents with scalable long-term memory. 2025. URL https://arxiv.org/abs/2504.19413.

J. Cohen. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46, 1960. doi: 10.1177/001316446002000104.

J. Cui, B. Xu, C. Wang, T. Yu, W. Sun, Y. Xu, T. Wang, Z. He, W. Ma, T. Cai, J. Gui, L. Zhang, X. Sun, F. Huang, M. Chen, Z. Lin, H. Liu, Q. Gui, Q. Han, Y. Wen, H. Liu, R. Wang, Y. Zhang, H. Wei, C. Chen, Y. Li, K. Fang, J. Zhou, Y. Li, G. Zeng, C. Xiao, Y. Lin, X. Han, M. Sun, Z. Liu, and Y. Yao. Minicpm-o 4.5: Towards real-time full-duplex omni-modal interaction. 2026. URL https://arxiv.org/abs/2604.27393.

X. Ding, H. Wu, Y. Yang, S. Jiang, D. Bai, Z. Chen, and T. Cao. Streammind: Unlocking full frame rate streaming video dialogue through event-gated cognition, 2025. URL https://arxiv.org/abs/2503.06220.

J. L. Fleiss. Measuring nominal scale agreement among many raters. *Psychological Bulletin*, 76(5): 378–382, 1971. doi: 10.1037/h0031619.

C. Fu, Y. Dai, Y. Luo, L. Li, S. Ren, R. Zhang, Z. Wang, C. Zhou, Y. Shen, M. Zhang, P. Chen, Y. Li, S. Lin, S. Zhao, K. Li, T. Xu, X. Zheng, E. Chen, C. Shan, R. He, and X. Sun. Video-mme: The first-ever comprehensive evaluation benchmark of multi-modal llms in video analysis. 2025a. URL https://arxiv.org/abs/2405.21075.

S. Fu, Q. Yang, Y.-M. Li, Y.-X. Peng, K.-Y. Lin, X. Wei, J.-F. Hu, X. Xie, and W.-S. Zheng. Vispeak: Visual instruction feedback in streaming videos. 2025b. URL https://arxiv.org/abs/2503.12769.

Google DeepMind. Gemini 3 model documentation. Google AI for Developers, 2026. URL https://ai.google.dev/gemini-api/docs/models.

K. Grauman, A. Westbury, L. Torresani, K. Kitani, J. Malik, et al. Ego-Exo4D: Understanding Skilled Human Activity from First- and Third-Person Perspectives. 2024. URL https://arxiv.org/abs/2311.18259.

M. Han, L. Yang, X. Chang, L. Yao, and H. Wang. Shot2story: A new benchmark for comprehensive understanding of multi-shot videos. 2025. URL https://arxiv.org/abs/2312.10300.

K. Hu, P. Wu, F. Pu, W. Xiao, Y. Zhang, X. Yue, B. Li, and Z. Liu. Video-mmmu: Evaluating knowledge acquisition from multi-discipline professional videos. 2025. URL https://arxiv.org/abs/2501.13826.

Z. Huang, X. Li, J. Li, J. Wang, X. Zeng, C. Liang, T. Wu, X. Chen, L. Li, and L. Wang. Online video understanding: Ovbench and videochat-online. 2025. URL https://arxiv.org/abs/2501.00584.

X. Jian, S. Nayak, K. Q. Lin, A. Feizi, K. Li, P. Bechard, S. Gella, and S. Rajeswar. Cua-suite: Massive human-annotated video demonstrations for computer-use agents. 2026. URL https://arxiv.org/abs/2603.24440.

P. Knab, O. Xhelili, I. Buzi, D. A. G. Nilo, M. S. Khan, L. Kolb, M. Scherzer, K. Yildirir, C. Bartelt, and P. J. Schubert. Barista: A multi-task egocentric benchmark for compositional visual understanding. 2026. URL https://arxiv.org/abs/2605.12074.

K. Kundu, R. Shrivastava, M. Arap, N. Wang, X. Zhu, Q. Fettes, G. Tiwari, P. Suresh, T. Moutakanni, A. C. Munoz, A. Bolourchi, P. Fung, P. Donmez, B. Damavandi, A. Kumar, and S. Moon. Plan, watch, recover: A benchmark and architectures for proactive procedural assistance. 2026. URL https://arxiv.org/abs/2606.04970.

J. R. Landis and G. G. Koch. The measurement of observer agreement for categorical data. *Biometrics*, 33(1):159–174, 1977. doi: 10.2307/2529310.

A. Li, Z. Xiao, Z. Yue, B. Xu, L. Yao, J. Li, P. Fu, J. Ju, J. Luan, and Q. Jin. Streampro: From reactive perception to proactive decision-making in streaming video. *arXiv preprint arXiv:2605.16381*, 2026a.

J. Li, Y. Chen, W. Song, Y. Lei, Y. Zhang, H. Yan, P. Pan, and M. Liu. Ipibench: Evaluating interactive proactive intelligence of mllms under continuous streams. 2026b. URL https://arxiv.org/abs/2605.27074.

K. Li, Y. Wang, Y. He, Y. Li, Y. Wang, Y. Liu, Z. Wang, J. Xu, G. Chen, P. Luo, L. Wang, and Y. Qiao. Mvbench: A comprehensive multi-modal video understanding benchmark. 2024. URL https://arxiv.org/abs/2311.17005.

X. Li, Y. Wang, J. Yu, X. Zeng, Y. Zhu, H. Huang, J. Gao, K. Li, Y. He, C. Wang, Y. Qiao, Y. Wang, and L. Wang. Videochat-flash: Hierarchical compression for long-context video modeling. 2025a. URL https://arxiv.org/abs/2501.00574.

Y. Li, J. Niu, Z. Miao, C. Ge, Y. Zhou, Q. He, X. Dong, H. Duan, S. Ding, R. Qian, P. Zhang, Y. Zang, Y. Cao, C. He, and J. Wang. Ovo-bench: How far is your video-llms from real-world online video understanding? 2025b. URL https://arxiv.org/abs/2501.05510.

F. Liu, J. Chen, K. Xu, Y. Liu, H. Guan, X. Lu, B. Yang, G. Hancke, R. Liu, and R. W. H. Lau. Guideme: Multi-domain task guidance and intervention in streaming video. 2026. URL https://arxiv.org/abs/2607.02991.

X. Lu, Y. Bo, J. Chen, S. Li, X. Guo, H. Guan, F. Liu, D. Xu, P. Sun, H. Sun, R. Liu, and H. Li. Aura: Always-on understanding and real-time assistance via video streams. 2026. URL https://arxiv.org/abs/2604.04184.

M. Maaz, H. Rasheed, S. Khan, and F. Khan. Videogpt+: Integrating image and video encoders for enhanced video understanding. 2024. URL https://arxiv.org/abs/2406.09418.

K. Mangalam, R. Akshulakov, and J. Malik. Egoschema: A diagnostic benchmark for very long-form video language understanding. 2023. URL https://arxiv.org/abs/2308.09126.

J. Mu, C. Zhang, C. Ni, L. Wang, B. Qiao, K. Mathur, Q. Wu, Y. Xie, X. Ma, M. Zhou, S. Qin, L. Li, Y. Kang, M. Ma, Q. Lin, S. Rajmohan, and D. Zhang. Gui-360°: A comprehensive dataset and benchmark for computer-using agents. 2025. URL https://arxiv.org/abs/2511.04307.

OpenAI. Gpt-realtime-2.1 model documentation. OpenAI API documentation, 2026. URL https://developers.openai.com/api/docs/models/gpt-realtime-2.1.

Orantqing, S. Ji, J. Tong, J. Zuo, D. Fu, D. Cao, Y. Li, S. Wu, Franz, Evan, T. Veyra, C. Pan, J. Lu, D. Yang, Z. Xie, Y. Tan, X. Shen, X. Yang, W. Wang, T. Sun, S. Yves, and Z. Zhao. Multimodal duplex interaction agent. 2026. URL https://arxiv.org/abs/2609.08977.

V. Pătrăucean, L. Smaira, A. Gupta, A. R. Continente, L. Markeeva, D. Banarse, S. Koppula, J. Heyward, M. Malinowski, Y. Yang, C. Doersch, T. Matejovicova, Y. Sulsky, A. Miech, A. Frechette, H. Klimczak, R. Koster, J. Zhang, S. Winkler, Y. Aytar, S. Osindero, D. Damen, A. Zisserman, and J. Carreira. Perception test: A diagnostic benchmark for multimodal video models. 2023. URL https://arxiv.org/abs/2305.13786.

Qwen Team. Qwen3.5-9b. Hugging Face model card, 2026. URL https://huggingface.co/Qwen/Qwen3.5-9B.

D. Ran, L. Ou, X. Li, W. Tong, C. Guo, H. Guo, K. Wang, and L. Lu. Egopro-bench: Benchmarking personalized proactive interaction in egocentric video streams. 2026. URL https://arxiv.org/

abs/2605.07299.

F. Sener, D. Chatterjee, D. Shelepov, K. He, D. Singhania, R. Wang, and A. Yao. Assembly101: A large-scale multi-view video dataset for understanding procedural activities. 2022. URL https://arxiv.org/abs/2203.14712.

Z. Shao, P. Wang, Q. Zhu, R. Xu, J. Song, X. Bi, H. Zhang, M. Zhang, Y. K. Li, Y. Wu, and D. Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. 2024. URL https://arxiv.org/abs/2402.03300.

G. A. Sigurdsson, G. Varol, X. Wang, A. Farhadi, I. Laptev, and A. Gupta. Hollywood in homes: Crowdsourcing data collection for activity understanding. 2016. URL https://arxiv.org/abs/1604.01753.

G. A. Sigurdsson, A. Gupta, C. Schmid, A. Farhadi, and K. Alahari. Charades-ego: A large-scale dataset of paired third and first person videos. 2018. URL https://arxiv.org/abs/1804.09626.

G. Sitong, T. Yan, C. Kang, B. Zheng, X. Ruan, H. Lu, K. Zhang, Y. Sato, and Y. Huang. Vinci2: Providing proactive assistance in continuous egocentric videos. 2026. URL https://arxiv.org/abs/2607.11523.

E. Song, W. Chai, W. Xu, J. Xie, Y. Liu, and G. Wang. Video-mmlu: A massive multi-discipline lecture understanding benchmark. 2025. URL https://arxiv.org/abs/2504.14693.

J. Tack, P. Laban, and J. Neville. Llms get lost in evolving user intent. 2026. URL https://arxiv.org/abs/2607.20734.

M. Tao, Y. Tao, and P. Wang. Intent-driven situation tracking for user-centric multi-turn agents. 2026. URL https://arxiv.org/abs/2608.15755.

H. Wang, B. Feng, Z. Lai, M. Xu, S. Li, W. Ge, A. Dehghan, M. Cao, and P. Huang. Streambridge: Turning your offline video large language model into a proactive streaming assistant. 2025a. URL https://arxiv.org/abs/2505.05467.

P. Wang, C. Tan, S. Zhou, Q. Zhou, Y. Chen, X. He, H. Zeng, J. Cheng, C. Wang, X. Qian, P. Wang, Z. Huang, S. Gao, W. Huang, L. Cao, W. Ran, J. Liu, C. Zhu, H. Wang, Y. Tian, C. Liu, Z. Ye, X. Wang, B. Jiang, G. Feng, Z. Fei, R. Li, M. Chen, Y. Gao, Q. Cheng, S. Li, and X. Qiu. Moss-vl technical report, 2026. URL https://arxiv.org/abs/2608.15045.

X. Wang, T. Kwon, M. Rad, B. Pan, I. Chakraborty, S. Andrist, D. Bohus, A. Feniello, B. Tekin, F. V. Frujeri, N. Joshi, and M. Pollefeys. Holoassist: an egocentric human interaction dataset for interactive ai assistants in the real world. 2023. URL https://arxiv.org/abs/2309.17024.

X. Wang, B. Wang, D. Lu, J. Yang, T. Xie, J. Wang, J. Deng, X. Guo, Y. Xu, C. H. Wu, Z. Shen, Z. Li, R. Li, X. Li, J. Chen, B. Zheng, P. Li, F. Lei, R. Cao, Y. Fu, D. Shin, M. Shin, J. Hu, Y. Wang, J. Chen, Y. Ye, D. Zhang, D. Du, H. Hu, H. Chen, Z. Zhou, H. Yao, Z. Chen, Q. Gu, Y. Wang, H. Wang, D. Yang, V. Zhong, F. Sung, Y. Charles, Z. Yang, and T. Yu. Opencua: Open foundations for computer-use agents. 2025b. URL https://arxiv.org/abs/2508.09123.

Y. Wang, Y. He, Y. Li, K. Li, J. Yu, X. Ma, X. Li, G. Chen, X. Chen, Y. Wang, C. He, P. Luo, Z. Liu, Y. Wang, L. Wang, and Y. Qiao. Internvid: A large-scale video-text dataset for multimodal understanding and generation. 2024. URL https://arxiv.org/abs/2307.06942.

Y. Wang, X. Li, Z. Yan, Y. He, J. Yu, X. Zeng, C. Wang, C. Ma, H. Huang, J. Gao, M. Dou, K. Chen, W. Wang, Y. Qiao, Y. Wang, and L. Wang. Internvideo2.5: Empowering video mllms with long and rich context modeling. 2025c. URL https://arxiv.org/abs/2501.12386.

Y. Wang, S. Liu, D. Wang, N. Xu, G. Wan, H. Zhang, and D. Zhao. Mmduet2: Enhancing proactive interaction of video mllms with multi-turn reinforcement learning. 2025d. URL https://arxiv.org/abs/2512.06810.

Y. Wang, X. Meng, Y. Wang, H. Zhang, and D. Zhao. Proactivevideoqa: A comprehensive benchmark evaluating proactive interactions in video large language models. 2025e. URL https://arxiv.org/abs/2507.09313.

J. Xia, P. Chen, M. Zhang, X. Sun, and K. Zhou. Streaming video instruction tuning. 2026. URL https://arxiv.org/abs/2512.21334.

W. Xiong, Y. Song, Q. Dong, B. Zhao, F. Song, X. Wang, and S. Li. Mpo: Boosting llm agents with meta plan optimization. 2025. URL https://arxiv.org/abs/2503.02682.

W. Xu, Z. Liang, K. Mei, H. Gao, J. Tan, and Y. Zhang. A-mem: Agentic memory for llm agents. 2025. URL https://arxiv.org/abs/2502.12110.

J. Yang, S. Liu, H. Guo, Y. Dong, X. Zhang, S. Zhang, P. Wang, Z. Zhou, B. Xie, Z. Wang, B. Ouyang, Z. Lin, M. Cominelli, Z. Cai, Y. Zhang, P. Zhang, F. Hong, J. Widmer, F. Gringoli, L. Yang, B. Li, and Z. Liu. Egolife: Towards egocentric life assistant. 2026a. URL https://arxiv.org/abs/2503.03803.

Q. Yang, H. Li, H. Zhao, X. Yan, J. Ding, F. Xu, and Y. Li. Fingertip 20k: A benchmark for proactive and personalized mobile llm agents. 2026b. URL https://arxiv.org/abs/2507.21071.

S. Yang, J. Yu, Y.-H. Peng, K. Q. Lin, J. W. Cho, Y. Song, and J. Kim. Guide: A benchmark for understanding and assisting users in open-ended gui tasks. 2026c. URL https://arxiv.org/abs/2603.25864.

S. Yang, K. Zhang, Z. Jia, J. Guo, Y. Shen, X. Zhang, X. Zhang, H. Wang, X. Li, P. Zhang, X. An, Y. Xie, Z. Liu, X. Guo, J. Li, S. Zheng, J. Wang, Z. Guo, W. Xie, Z. Zheng, Y. Luo, B. Li, and Y. Lu. Mage-vl: An efficient codec-native streaming multimodal foundation model. 2026d. URL https://arxiv.org/abs/2607.24904.

Z. Yang, K. Zhang, Y. Hu, B. Wang, S. Qian, B. Wen, F. Yang, T. Gao, W. Dong, and C. Xu. Livestar: Live streaming assistant for real-world online video understanding. 2025. URL https://arxiv.org/abs/2511.05299.

D. Yao, S. Gu, Q. Si, J. Zhou, C. Yang, C. Qin, N. Gu, Z. Lin, W. Wang, N. Duan, and J. Wang. Harnessing streaming video in the wild, 2026a. URL https://arxiv.org/abs/2606.08615.

D. Yao, J. Zhou, C. Yang, C. Qin, X. Zeng, Y. Li, H. Hou, Z. Liang, C. Wang, K. Tuo, J. Zhang, Y. Zhu, Y. Cao, S. Ye, S. Xie, S. Gu, H. Huang, Q. Si, N. Duan, and J. Wang. Joyai-vl-interaction: Real-time vision-language interaction intelligence. 2026b. URL https://arxiv.org/abs/2606.14777.

L. Yao, Y. Li, Y. Wei, L. Li, S. Ren, Y. Liu, K. Ouyang, L. Wang, S. Li, S. Li, L. Kong, Q. Liu, Y. Zhang, and X. Sun. Timechat-online: 80% visual tokens are naturally redundant in streaming videos. 2025. URL https://arxiv.org/abs/2504.17343.

Y. Yu, L. Yao, Y. Xie, Q. Tan, J. Feng, Y. Li, and L. Wu. Agentic memory: Learning unified long-term and short-term memory management for large language model agents. 2026. URL https://arxiv.org/abs/2601.01885.

S. Yuan, Z. Chen, Z. Xi, J. Ye, Z. Du, and J. Chen. Agent-r: Training language model agents to reflect via iterative self-training. 2025. URL https://arxiv.org/abs/2501.11425.

X. Zeng, K. Qiu, Q. Zhang, X. Li, J. Wang, J. Li, Z. Yan, K. Tian, M. Tian, X. Zhao, Y. Wang, and L. Wang. Streamforest: Efficient online video understanding with persistent event memory. 2025. URL https://arxiv.org/abs/2509.24871.

J. Zhang, T. Wang, Y. Ge, Y. Ge, X. Li, Y. Shan, and L. Wang. Timelens: Rethinking video temporal grounding with multimodal llms. 2026. URL https://arxiv.org/abs/2512.14698.

Y. Zhao, L. Xie, H. Zhang, G. Gan, Y. Long, Z. Hu, T. Hu, W. Chen, C. Li, J. Song, Z. Xu, C. Wang, W. Pan, Z. Shangguan, X. Tang, Z. Liang, Y. Liu, C. Zhao, and A. Cohan. Mmvu: Measuring expert-level multi-discipline video understanding. 2025. URL https://arxiv.org/abs/2501.12380.

B. Zheng, M. Y. Fatemi, X. Jin, Z. Z. Wang, A. Gandhi, Y. Song, Y. Gu, J. Srinivasa, G. Liu, G. Neubig, and Y. Su. Skillweaver: Web agents can self-improve by discovering and honing skills. 2025. URL https://arxiv.org/abs/2504.07079.

J. Zhou, Y. Shu, B. Zhao, B. Wu, Z. Liang, S. Xiao, M. Qin, X. Yang, Y. Xiong, B. Zhang, T. Huang, and Z. Liu. Mlvu: Benchmarking multi-task long video understanding. 2025. URL https:

//arxiv.org/abs/2406.04264.

L. Zhou, C. Xu, and J. J. Corso. Towards automatic learning of procedures from web instructional videos. 2017. URL https://arxiv.org/abs/1703.09788.

J. Zhu, W. Wang, Z. Chen, Z. Liu, S. Ye, L. Gu, H. Tian, Y. Duan, W. Su, J. Shao, Z. Gao, E. Cui, X. Wang, Y. Cao, Y. Liu, X. Wei, H. Zhang, H. Wang, W. Xu, H. Li, J. Wang, N. Deng, S. Li, Y. He, T. Jiang, J. Luo, Y. Wang, C. He, B. Shi, X. Zhang, W. Shao, J. He, Y. Xiong, W. Qu, P. Sun, P. Jiao, H. Lv, L. Wu, K. Zhang, H. Deng, J. Ge, K. Chen, L. Wang, M. Dou, L. Lu, X. Zhu, T. Lu, D. Lin, Y. Qiao, J. Dai, and W. Wang. Internvl3: Exploring advanced training and test-time recipes for open-source multimodal models. 2025. URL https://arxiv.org/abs/2504.10479.

A. Additional Evaluation Results

We report Pointwise evaluation with gold or model-generated interaction history and Trajectory evaluation with model-generated history. The matching tolerances are ±2, ±1, and ±0 s; tighter tolerances rescore the same predictions. The following model-comparison tables report StreamVibe-VL without the Harness, followed by separate Harness comparisons. **Bold** and underline denote the best and second-best eligible results in each column; FPR is lower-is-better and all other metrics are higher-is-better.

A.1. Results Under the Three Evaluation Protocols

Tables 3 and 4 report Pointwise evaluation with gold interaction history, Tables 5 and 6 report Pointwise evaluation with self-generated interaction history, and Tables 7 and 8 report Trajectory evaluation with self-generated interaction history.

Model	Explicit						Implicit						Dynamic					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models																		
Gemini-3.1-Pro	70.92	85.58	77.56	43.57	67.46	52.32	**62.34**	68.52	65.28	41.40	45.22	29.52	62.87	84.59	72.13	58.81	51.15	36.89
GPT-Realtime-2.1	69.67	70.95	70.30	29.83	55.92	39.31	57.78	90.09	**70.41**	64.87	51.10	**35.98**	**62.99**	84.47	**72.17**	55.37	45.64	32.94
Gemini-3.8-Flash	**81.55**	80.45	**81.00**	27.09	82.63	**66.93**	60.34	77.40	67.81	50.69	51.37	34.83	62.96	84.52	72.16	60.25	**60.56**	**43.70**
Open-Source Offline Models																		
Qwen3.5-9B	57.42	61.23	59.26	49.17	78.37	46.44	52.95	91.36	67.04	81.45	44.59	29.89	56.74	90.93	69.88	78.89	47.50	33.19
Qwen3-VL-8B	56.38	**90.87**	69.59	83.71	80.51	56.03	51.25	**95.28**	66.65	91.05	47.91	31.93	54.37	**97.56**	69.83	93.88	50.91	35.55
JoyAI-VL-Interaction	58.38	79.21	67.22	67.32	81.63	54.87	52.50	90.10	66.34	81.88	40.06	26.58	57.48	72.20	64.00	62.60	45.23	28.95
Open-Source Streaming Models																		
MMDuet2	15.98	36.50	22.23	**4.13**	74.26	16.51	19.45	38.65	25.88	7.90	22.71	5.88	13.06	35.51	19.10	5.24	39.10	7.47
Mage-VL	8.56	51.94	14.70	5.50	8.19	1.20	20.70	50.25	29.32	6.11	10.23	3.00	11.97	52.88	19.52	6.20	6.10	1.19
LiveStar	6.84	24.95	10.74	6.27	25.43	2.73	14.96	19.95	17.10	4.81	3.60	0.62	7.29	12.92	9.32	2.86	8.05	0.75
ViSpeak	8.41	74.03	15.10	12.67	62.98	9.51	20.66	74.71	32.37	12.75	10.19	3.30	13.51	72.32	22.77	14.19	16.85	3.84
VideoLLM-Online	6.33	17.66	9.32	5.32	31.88	2.97	12.29	5.04	7.15	**1.44**	17.89	1.28	17.86	14.75	16.16	**1.28**	8.69	1.40
Ours																		
StreamVibe-VL (SFT)	60.19	89.31	71.91	73.11	85.02	61.14	55.74	69.79	61.98	56.35	53.65	33.25	57.56	79.68	66.84	69.58	57.67	38.55
StreamVibe-VL (RL)	61.59	86.66	72.01	67.91	**85.11**	61.29	57.33	60.89	59.06	46.19	**55.32**	32.67	58.94	75.86	66.34	62.93	59.25	39.31

Table 3 | ***Pointwise evaluation with gold interaction history*** on StreamVibe-Bench without the Harness (±2 s tolerance).

Model	Explicit						Implicit						Dynamic					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models																		
Gemini-3.1-Pro	67.47	81.43	73.80	51.13	68.92	50.86	53.46	58.76	55.98	52.96	46.73	26.16	58.03	78.07	66.57	68.27	52.94	35.24
GPT-Realtime-2.1	63.10	64.25	63.67	42.00	57.79	36.79	50.76	79.14	61.85	77.83	51.78	32.03	55.68	74.67	63.79	69.60	46.60	29.73
Gemini-3.8-Flash	**79.04**	77.97	**78.50**	31.60	83.41	**65.48**	53.48	68.61	60.11	61.10	52.45	31.53	**59.01**	79.21	67.63	67.94	**61.25**	**41.42**
Open-Source Offline Models																		
Qwen3.5-9B	54.78	58.42	56.54	54.27	79.23	44.80	51.40	88.69	65.08	84.61	44.91	29.23	55.39	88.77	68.22	82.03	47.61	32.48
Qwen3-VL-8B	55.38	**89.25**	68.35	86.65	81.16	55.47	50.52	**93.93**	**65.70**	92.66	47.82	31.42	54.18	**97.22**	**69.58**	94.37	50.97	35.46
JoyAI-VL-Interaction	57.78	78.40	66.53	68.79	81.96	54.53	51.82	88.94	65.49	83.25	40.27	26.37	56.31	70.74	62.71	64.73	45.40	28.47
Open-Source Streaming Models																		
MMDuet2	4.77	10.91	6.64	7.13	75.87	5.04	5.49	10.90	7.30	11.10	24.07	1.76	7.12	19.35	10.41	6.95	38.66	4.02
Mage-VL	1.97	11.93	3.38	10.19	9.84	0.33	7.58	18.40	10.74	9.78	9.65	1.04	3.05	13.48	4.97	10.38	7.07	0.35
LiveStar	2.95	10.75	4.63	7.93	27.01	1.25	7.31	9.75	8.36	5.98	3.85	0.32	2.61	4.63	3.34	3.73	10.36	0.35
ViSpeak	2.16	19.01	3.88	19.12	64.28	2.49	5.52	19.96	8.65	19.06	10.99	0.95	4.86	26.02	8.19	19.10	16.76	1.37
VideoLLM-Online	2.53	7.07	3.73	**6.56**	29.77	1.11	6.66	2.73	3.87	**1.71**	19.82	0.77	15.24	12.59	13.79	**1.51**	5.52	0.76
Ours																		
StreamVibe-VL (SFT)	59.83	88.77	71.48	74.09	85.14	60.86	53.46	66.95	59.45	59.72	54.21	**32.23**	56.36	78.01	65.44	72.00	58.14	38.05
StreamVibe-VL (RL)	61.01	85.85	71.33	69.38	**85.24**	60.80	**54.45**	57.83	56.09	49.81	**55.90**	31.35	57.45	73.94	64.66	65.71	59.70	38.60

Table 4 | ***Pointwise evaluation with gold interaction history*** on StreamVibe-Bench without the Harness (±0 s tolerance).

Model	Explicit						Implicit						Dynamic					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models																		
Gemini-3.1-Pro	66.28	**82.45**	73.49	33.17	40.95	30.09	68.23	70.21	69.21	29.31	29.20	20.21	73.03	80.05	76.38	28.13	36.70	28.03
GPT-Realtime-2.1	65.42	72.84	68.93	34.54	54.10	37.29	61.47	85.03	71.36	50.95	**48.94**	34.92	66.86	79.83	72.77	42.50	41.99	30.56
Gemini-3.8-Flash	**89.93**	75.22	**81.92**	5.79	71.88	**58.88**	79.24	79.49	79.36	13.77	46.90	**37.22**	**85.98**	**82.75**	**84.33**	7.27	**56.27**	**47.45**
Open-Source Offline Models																		
Qwen3.5-9B	66.55	29.86	41.22	8.24	68.04	28.05	56.04	89.04	68.79	67.93	22.90	15.75	62.64	80.52	70.46	51.20	32.39	22.82
JoyAI-VL-Interaction	64.78	67.33	66.03	29.64	68.08	44.95	78.54	**91.23**	**84.41**	14.37	21.65	18.27	83.76	59.29	69.43	4.54	33.53	23.28
Open-Source Streaming Models																		
MMDuet2	15.98	36.50	22.23	**4.13**	**74.26**	16.51	19.45	38.65	25.88	7.90	22.71	5.88	13.06	35.51	19.10	5.24	39.10	7.47
Mage-VL	8.56	51.94	14.70	5.50	8.19	1.20	20.70	50.25	29.32	6.11	10.23	3.00	11.97	52.88	19.52	6.20	6.10	1.19
LiveStar	6.84	24.95	10.74	6.27	25.43	2.73	14.96	19.95	17.10	4.81	3.60	0.62	7.29	12.92	9.32	2.86	8.05	0.75
ViSpeak	8.41	74.03	15.10	12.67	62.98	9.51	20.66	74.71	32.37	12.75	10.19	3.30	13.51	72.32	22.77	14.19	16.85	3.84
VideoLLM-Online	6.33	17.66	9.32	5.32	31.88	2.97	12.29	5.04	7.15	**1.44**	17.89	1.28	17.86	14.75	16.16	**1.28**	8.69	1.40
Ours																		
StreamVibe-VL (SFT)	69.05	76.03	72.37	27.09	70.42	50.96	77.54	62.79	69.39	12.51	42.46	29.46	76.98	69.37	72.98	16.38	45.94	33.53
StreamVibe-VL (RL)	73.58	76.67	75.09	17.66	72.43	54.39	**84.28**	65.75	73.87	5.00	45.71	33.77	82.70	68.45	74.90	8.05	48.53	36.35

Table 5 | ***Pointwise evaluation with self-generated interaction history*** on StreamVibe-Bench without the Harness (±2 s tolerance).

Model	Explicit						Implicit						Dynamic					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models																		
Gemini-3.1-Pro	54.69	**68.03**	60.63	59.37	41.25	25.01	45.91	47.24	46.57	56.49	27.20	12.67	50.21	55.03	52.51	64.44	33.65	17.67
GPT-Realtime-2.1	57.03	63.50	60.09	51.52	55.53	33.37	**49.81**	68.90	57.82	70.04	49.80	**28.79**	**54.90**	65.55	**59.75**	63.23	43.11	**25.76**
Gemini-3.8-Flash	**72.18**	60.37	**65.75**	32.78	73.90	**48.59**	35.61	35.72	35.66	65.59	49.80	17.76	43.63	42.00	42.80	66.43	**58.46**	25.02
Open-Source Offline Models																		
Qwen3.5-9B	53.43	23.97	33.09	18.94	69.09	22.86	48.11	**76.44**	**59.05**	82.85	23.04	13.61	52.34	**67.28**	58.88	70.42	31.84	18.75
JoyAI-VL-Interaction	47.90	49.78	48.82	61.53	68.28	33.33	27.04	31.40	29.06	85.19	21.03	6.11	37.12	26.27	30.77	52.47	34.70	10.68
Open-Source Streaming Models																		
MMDuet2	4.77	10.91	6.64	7.13	**75.87**	5.04	5.49	10.90	7.30	11.10	24.07	1.76	7.12	19.35	10.41	6.95	38.66	4.02
Mage-VL	1.97	11.93	3.38	10.19	9.84	0.33	7.58	18.40	10.74	9.78	9.65	1.04	3.05	13.48	4.97	10.38	7.07	0.35
LiveStar	2.95	10.75	4.63	7.93	27.01	1.25	7.31	9.75	8.36	5.98	3.85	0.32	2.61	4.63	3.34	3.73	10.36	0.35
ViSpeak	2.16	19.01	3.88	19.12	64.28	2.49	5.52	19.96	8.65	19.06	10.99	0.95	4.86	26.02	8.19	19.10	16.76	1.37
VideoLLM-Online	2.53	7.07	3.73	**6.56**	29.77	1.11	6.66	2.73	3.87	**1.71**	19.82	0.77	15.24	12.59	13.79	**1.51**	5.52	0.76
Ours																		
StreamVibe-VL (SFT)	47.87	52.70	50.17	69.48	67.37	33.80	33.89	27.44	30.33	54.36	46.02	13.96	40.67	36.65	38.56	63.87	45.61	17.59
StreamVibe-VL (RL)	47.36	49.35	48.33	67.32	69.37	33.53	31.78	28.02	29.78	46.10	**50.35**	14.99	39.12	37.49	38.29	57.34	49.03	18.77

Table 6 | ***Pointwise evaluation with self-generated interaction history*** on StreamVibe-Bench without the Harness (±0 s tolerance).

Model	Explicit						Implicit						Dynamic					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models																		
Gemini-3.1-Pro	32.74	**79.21**	46.33	4.60	85.60	39.66	17.73	58.78	27.24	12.80	52.12	14.20	14.23	57.82	22.84	8.92	64.95	14.83
GPT-Realtime-2.1	9.70	55.56	16.52	11.63	66.84	11.04	14.91	56.26	23.57	15.71	48.21	11.36	9.69	52.67	16.37	14.49	46.86	7.67
Gemini-3.8-Flash	**47.69**	79.10	**59.50**	**2.37**	**87.49**	**52.06**	**21.03**	71.82	**32.53**	12.94	**59.96**	**19.50**	16.00	66.17	**25.77**	8.42	**72.48**	**18.68**
Open-Source Offline Models																		
Qwen3.5-9B	7.94	52.48	13.79	14.56	44.98	6.20	9.37	47.62	15.66	19.30	16.63	2.60	4.87	51.65	8.90	20.95	27.67	2.46
JoyAI-VL-Interaction	4.39	37.96	7.87	18.87	46.83	3.69	4.67	**95.44**	8.90	90.87	10.21	0.91	3.15	**76.70**	6.05	64.52	21.64	1.31
Open-Source Streaming Models																		
MMDuet2	15.98	36.50	22.23	4.13	74.26	16.51	19.45	38.65	25.88	7.90	22.71	5.88	13.06	35.51	19.10	5.24	39.10	7.47
Mage-VL	8.56	51.94	14.70	5.50	8.19	1.20	20.70	50.25	29.32	6.11	10.23	3.00	11.97	52.88	19.52	6.20	6.10	1.19
LiveStar	6.84	24.95	10.74	6.27	25.43	2.73	14.96	19.95	17.10	4.81	3.60	0.62	7.29	12.92	9.32	2.86	8.05	0.75
ViSpeak	8.41	74.03	15.10	12.67	62.98	9.51	20.66	74.71	32.37	12.75	10.19	3.30	13.51	72.32	22.77	14.19	16.85	3.84
VideoLLM-Online	6.33	17.66	9.32	5.32	31.88	2.97	12.29	5.04	7.15	**1.44**	17.89	1.28	**17.86**	14.75	16.16	**1.28**	8.69	1.40
Ours																		
StreamVibe-VL (SFT)	11.55	40.66	17.99	11.40	57.97	10.43	13.13	47.91	20.61	15.59	38.06	7.84	5.38	56.76	9.83	26.13	35.88	3.53
StreamVibe-VL (RL)	15.80	41.68	22.91	8.27	60.07	13.76	18.02	43.64	25.51	9.68	41.74	10.65	7.60	52.58	13.28	17.11	39.55	5.25

Table 7 | ***Trajectory evaluation with self-generated interaction history*** on StreamVibe-Bench without the Harness (±2 s tolerance).

Model	Explicit						Implicit						Dynamic					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models																		
Gemini-3.1-Pro	19.21	**46.49**	27.19	8.43	85.54	23.26	9.01	29.88	13.85	16.13	55.51	7.69	7.74	31.46	12.42	11.71	69.94	8.69
GPT-Realtime-2.1	4.84	27.70	8.24	14.90	71.20	5.87	6.62	24.96	10.46	19.32	52.70	5.51	4.34	23.61	7.33	17.57	51.95	3.81
Gemini-3.8-Flash	**27.51**	45.63	**34.33**	**6.29**	**88.34**	**30.33**	**10.60**	36.20	**16.40**	17.04	**64.64**	**10.60**	8.52	35.23	13.72	11.70	**76.23**	**10.46**
Open-Source Offline Models																		
Qwen3.5-9B	3.69	24.41	6.41	17.85	44.26	2.84	4.75	24.15	7.94	22.00	17.36	1.38	2.34	24.81	4.28	23.79	29.24	1.25
JoyAI-VL-Interaction	3.23	27.92	5.79	20.05	41.05	2.38	4.53	**92.53**	8.64	91.20	10.04	0.87	2.92	**70.92**	5.61	65.11	21.13	1.19
Open-Source Streaming Models																		
MMDuet2	4.77	10.91	6.64	7.13	75.87	5.04	5.49	10.90	7.30	11.10	24.07	1.76	7.12	19.35	10.41	6.95	38.66	4.02
Mage-VL	1.97	11.93	3.38	10.19	9.84	0.33	7.58	18.40	10.74	9.78	9.65	1.04	3.05	13.48	4.97	10.38	7.07	0.35
LiveStar	2.95	10.75	4.63	7.93	27.01	1.25	7.31	9.75	8.36	5.98	3.85	0.32	2.61	4.63	3.34	3.73	10.36	0.35
ViSpeak	2.16	19.01	3.88	19.12	64.28	2.49	5.52	19.96	8.65	19.06	10.99	0.95	4.86	26.02	8.19	19.10	16.76	1.37
VideoLLM-Online	2.53	7.07	3.73	6.56	29.77	1.11	6.66	2.73	3.87	**1.71**	19.82	0.77	**15.24**	12.59	**13.79**	**1.51**	5.52	0.76
Ours																		
StreamVibe-VL (SFT)	6.47	22.79	10.08	13.49	52.78	5.32	7.62	27.78	11.96	17.91	38.43	4.60	3.92	41.33	7.16	27.76	35.18	2.52
StreamVibe-VL (RL)	7.90	20.84	11.46	10.71	56.67	6.49	9.01	21.81	12.75	12.20	42.85	5.46	4.98	34.48	8.70	19.03	41.00	3.57

Table 8 | ***Trajectory evaluation with model-generated interaction history*** on StreamVibe-Bench without the Harness (±0 s tolerance).

A.2. General Timing and External Benchmarks

The general-timing results use Pointwise or Trajectory evaluation as indicated. EgoProactive, IPIBench, and OVO-Bench retain their benchmark-specific evaluation protocols.

Model	$\tau = 2$ s						$\tau = 0$ s					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models												
Gemini-3.1-Pro	73.78	78.71	76.17	34.17	47.57	36.23	65.06	69.40	67.16	47.97	49.82	33.46
GPT-Realtime-2.1	65.29	82.16	72.76	48.30	40.64	29.57	56.15	70.66	62.57	65.34	41.09	25.71
Gemini-3.8-Flash	**75.14**	79.58	**77.30**	35.71	47.59	36.79	**68.79**	72.85	70.76	45.70	49.06	34.71
Open-Source Offline Models												
Qwen3.5-9B	52.64	83.96	64.71	88.56	61.52	39.81	51.17	81.61	62.90	92.05	61.44	38.65
JoyAI-VL-Interaction	69.34	86.15	76.84	53.41	67.11	51.57	67.44	83.80	**74.74**	56.90	67.67	50.58
Open-Source Streaming Models												
MMDuet2	37.06	45.76	40.95	9.17	49.16	20.13	14.98	18.50	16.55	14.40	45.18	7.48
Mage-VL	35.96	37.99	36.95	4.84	22.73	8.40	9.64	10.18	9.90	10.18	22.58	2.24
LiveStar	23.57	40.34	29.75	12.18	28.80	8.57	12.79	21.89	16.15	15.72	34.38	5.55
ViSpeak	28.21	84.84	42.34	22.66	33.06	14.00	12.19	36.67	18.30	31.91	32.39	5.93
VideoLLM-Online	45.06	17.46	25.17	**3.40**	12.54	3.16	30.65	11.88	17.12	**4.47**	14.40	2.47
Ours												
StreamVibe-VL (SFT)	59.39	**96.93**	73.65	94.32	**71.32**	52.53	59.22	96.66	73.44	94.72	**71.35**	52.40
StreamVibe-VL (RL)	60.37	96.60	74.30	85.47	71.28	**52.96**	59.90	**97.91**	74.33	86.53	71.30	**53.00**

Table 9 | ***Pointwise evaluation with gold interaction history*** on general timing without the Harness.

Model	$\tau = 2$ s						$\tau = 0$ s					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Proprietary Models												
Gemini-3.1-Pro	29.75	91.35	44.88	26.46	68.26	30.64	20.62	63.33	31.11	31.85	68.00	21.15
GPT-Realtime-2.1	30.17	68.86	41.96	17.84	60.90	25.55	17.84	40.72	24.81	23.25	61.91	15.36
Gemini-3.8-Flash	**48.33**	85.50	**61.75**	12.53	**72.78**	**44.94**	30.38	53.75	**38.82**	18.63	**74.77**	**29.03**
Open-Source Offline Models												
Qwen3.5-9B	24.21	31.53	27.39	5.36	34.18	9.36	9.79	12.75	11.08	8.99	32.45	3.60
JoyAI-VL-Interaction	30.00	70.61	42.11	26.05	54.85	23.10	20.74	48.82	29.11	30.27	51.49	14.99
Open-Source Streaming Models												
MMDuet2	37.06	45.76	40.95	9.17	49.16	20.13	14.98	18.50	16.55	14.40	45.18	7.48
Mage-VL	35.96	37.99	36.95	4.84	22.73	8.40	9.64	10.18	9.90	10.18	22.58	2.24
LiveStar	23.57	40.34	29.75	12.18	28.80	8.57	12.79	21.89	16.15	15.72	34.38	5.55
ViSpeak	28.21	84.84	42.34	22.66	33.06	14.00	12.19	36.67	18.30	31.91	32.39	5.93
VideoLLM-Online	45.06	17.46	25.17	**3.40**	12.54	3.16	**30.65**	11.88	17.12	**4.47**	14.40	2.47
Ours												
StreamVibe-VL (SFT)	22.95	86.92	36.31	49.13	47.87	17.38	19.51	73.89	30.87	51.65	43.09	13.30
StreamVibe-VL (RL)	25.44	**92.28**	39.88	44.94	53.80	21.46	20.42	**74.06**	32.01	48.47	48.06	15.38

Table 10 | ***Trajectory evaluation with self-generated interaction history*** on general timing without the Harness.

Model	EgoProactive				IPIBench			
	Macro	G-mean	Interrupt	Silent	Timing	Underst.	Repeated	Macro
Proprietary Models								
Gemini-3.1-Pro	58.75	58.75	59.42	58.09	53.40	33.12	7.55	31.36
GPT-Realtime-2.1	45.28	45.28	67.09	23.46	41.52	18.25	8.02	22.60
Gemini-3.8-Flash	**64.96**	**64.96**	66.19	63.72	**58.77**	**37.82**	12.74	**36.44**
Open-Source Offline Models								
Qwen3.5-9B	46.20	43.20	62.60	29.80	32.50	9.80	5.70	16.00
JoyAI-VL-Interaction	41.00	28.60	**70.40**	11.60	34.30	7.40	1.40	14.40
Open-Source Streaming Models								
MMDuet2	38.21	30.96	15.82	60.61	30.88	6.49	5.66	14.34
Mage-VL	40.52	35.48	60.10	20.95	26.65	11.47	5.19	14.44
LiveStar	53.62	53.16	46.58	60.67	14.65	3.40	0.47	6.17
ViSpeak	45.21	43.84	34.17	56.26	19.89	0.43	0.00	6.77
VideoLLM-Online	44.71	39.14	23.10	**66.32**	7.67	0.38	2.36	3.47
Ours								
StreamVibe-VL (SFT)	53.04	52.20	62.45	43.63	33.80	8.30	**16.00**	19.40
StreamVibe-VL (RL)	59.50	59.50	58.90	60.00	39.10	20.40	9.40	23.00

Table 11 | Results on EgoProactive and IPIBench without the Harness.

Model	OVO-Bench			
	BT	RT	FAR	Overall
Open-Source Offline Models				
Qwen3.5-9B	53.94	63.94	55.24	57.71
JoyAI-VL-Interaction	47.09	67.99	38.87	51.32
Open-Source Streaming Models				
MMDuet2	47.02	61.59	37.91	48.84
Mage-VL	46.76	**81.12**	43.50	**63.94**
LiveStar	42.25	55.12	4.29	33.88
ViSpeak	**57.00**	65.80	53.61	60.55
VideoLLM-Online	31.99	29.58	5.69	22.42
Ours				
StreamVibe-VL (SFT)	54.12	64.12	60.10	59.44
StreamVibe-VL (RL)	53.42	64.55	**62.00**	59.29

Table 12 | Results on OVO-Bench without the Harness.

A.3. StreamVibe-Harness

We compare StreamVibe-VL with and without the Harness under Trajectory evaluation at ±2 s, using the corresponding base model as a reference. The external-benchmark comparison also reports the matched SFT pair evaluated with and without the Harness.

Model	Explicit						Implicit						Dynamic					
	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1	P	R	F1	FPR	ANS	C-F1
Ours + Harness																		
StreamVibe-VL (SFT) + Harness	18.49	28.51	22.43	3.80	61.79	13.86	18.20	29.39	22.48	6.33	38.93	8.75	10.01	43.20	16.25	12.25	38.98	6.33
StreamVibe-VL (RL) + Harness	**24.57**	16.20	19.53	**1.49**	**72.83**	**14.22**	**21.53**	29.48	24.89	**2.02**	40.86	10.17	**14.80**	30.65	**19.96**	**5.68**	**44.62**	**8.91**
Ours																		
StreamVibe-VL (SFT)	11.55	40.66	17.99	11.40	57.97	10.43	13.13	**47.91**	20.61	15.59	38.06	7.84	5.38	**56.76**	9.83	26.13	35.88	3.53
StreamVibe-VL (RL)	15.80	41.68	**22.91**	8.27	60.07	13.76	18.02	43.64	**25.51**	9.68	**41.74**	**10.65**	7.60	52.58	13.28	17.11	39.55	5.25
Open-Source Offline Models																		
Qwen3.5-9B	7.94	**52.48**	13.79	14.56	44.98	6.20	9.37	47.62	15.66	19.30	16.63	2.60	4.87	51.65	8.90	20.95	27.67	2.46

Table 13 | ***Trajectory evaluation*** of StreamVibe-VL with and without the Harness on StreamVibe-Bench (±2 s tolerance).

Model	General timing						IPIBench				EgoProactive			
	P	R	F1	FPR	ANS	C-F1	Timing	Underst.	Repeated	Macro	Macro	G-mean	Interrupt	Silent
Ours + Harness														
StreamVibe-VL (SFT) + Harness	48.39	68.25	56.63	12.57	61.44	34.79	44.13	15.67	38.68	32.83	46.71	43.48	**63.75**	29.66
StreamVibe-VL (RL) + Harness	**53.23**	67.22	**59.41**	6.84	**62.88**	**37.36**	**45.34**	16.88	**40.05**	**35.67**	48.79	44.0	3.87	29.80
Ours														
StreamVibe-VL (SFT)	22.95	86.92	36.31	49.13	47.87	17.38	34.85	13.40	28.77	25.67	35.74	23.33	62.82	8.66
StreamVibe-VL (RL)	25.44	**92.28**	39.88	44.94	53.80	21.46	39.10	**20.40**	29.40	23.00	**53.04**	**52.20**	62.45	**43.63**
Open-Source Offline Models														
Qwen3.5-9B	24.21	31.53	27.39	**5.36**	34.18	9.36	32.50	9.80	5.70	16.00	46.20	43.20	62.60	29.80

Table 14 | StreamVibe-Harness comparison on general timing (Trajectory, ±2 s), IPIBench, and EgoProactive.

A.4. Category Difficulty Across Intent Types

Building on the overall experimental results above, we further analyze performance across different subtypes to obtain a more fine-grained understanding of proactive interaction capabilities. Specifically, we examine *(i) the relative difficulty of different subtypes for existing systems*, *(ii) how model performance changes across evaluation settings*, and *(iii) how RL shifts response behavior and performance trade-offs across subtypes*. All analyses use a matching tolerance of ±2 s. Figures 8–10 report subtype-level results under the Pointwise and Trajectory evaluation settings used in Tables 3 and 7. Figure 8 summarizes the average performance of all available models across subtypes under each evaluation protocol to characterize the overall difficulty of different task types. Pointwise includes seven models, while Trajectory additionally includes five native streaming baselines.

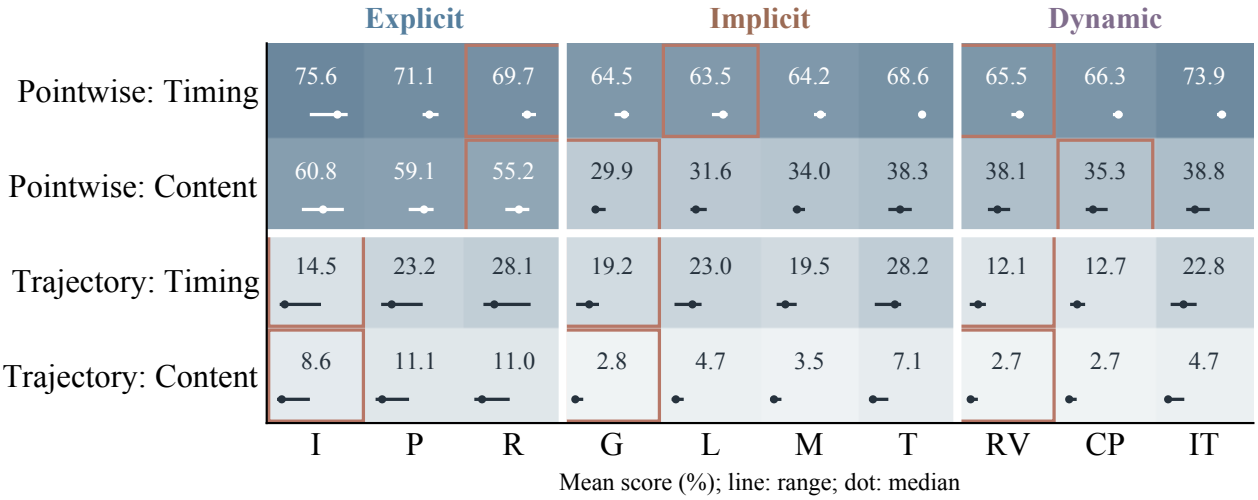

Figure 8 | **Category difficulty across StreamVibe-Bench.** Each cell gives the mean score across all models available in that setting; the line spans the per-model range and the dot marks the median. Outlines identify the lowest mean within each intent type and metric. Pointwise uses gold interaction history (n=7 models), whereas Trajectory uses model-generated interaction history and also includes native streaming baselines (n=12 models). All results use the ±2 s tolerance.

Trajectory evaluation reveals substantially greater difficulty under sustained interaction. As shown in Figure 8, performance drops markedly from Pointwise to Trajectory evaluation across all subtypes. Under Pointwise evaluation, Timing F1 remains relatively high at 63.5–75.6, indicating that models can often recognize locally appropriate response moments when evaluated at isolated timestamps. In contrast, Trajectory Timing F1 falls to 12.1–28.2 and Content-F1 to 2.7–11.1, revealing a substantial gap between local perception and sustained interaction. This suggests that maintaining reliable assistance over a full trajectory requires more than recognizing individual events: models must continuously update active intents, track task state, and regulate responses based on their evolving interaction history.

Figure 8 further shows that subtype difficulty changes under sustained interaction. For Explicit intents, I achieves the highest Pointwise Timing F1 at 75.6 but the lowest Trajectory Timing F1 at 14.5, whereas R reaches 28.1. Among Implicit intents, Tool/GUI remains relatively strong, while Guidance becomes more challenging. For Dynamic intents, Revision and Composition underperform Interruption, highlighting the difficulty of continuously revising and composing active requirements. Overall, Pointwise evaluation primarily reflects local recognition and timing, whereas Trajectory evaluation exposes the additional challenges of maintaining coherent and reliable behavior throughout continuous, evolving interaction.

Different subtypes exhibit distinct performance patterns. As shown in Figure 9, different subtypes exhibit clear differences in both response timing and content quality. StreamVibe-VL performs best on Tool/GUI, achieving a Timing F1 of 37.9 and a Content-F1 of 13.0, while Monitoring reaches only 17.2 and 4.1, respectively, suggesting that tasks with explicit operational goals are easier to handle reliably than continuous monitoring. For Explicit intents, P and R achieve Timing F1 scores of 26.0 and 25.1, substantially higher than I at 8.9, indicating that single-event cases are more prone to missed or mistimed responses during sustained interaction. For Dynamic intents, Revision and Composition both achieve Timing F1 scores around 10, whereas Interruption reaches 19.0 and also attains higher Content-F1, suggesting that continuously revising or composing existing requirements places greater demands on maintaining the currently valid intent and task state.

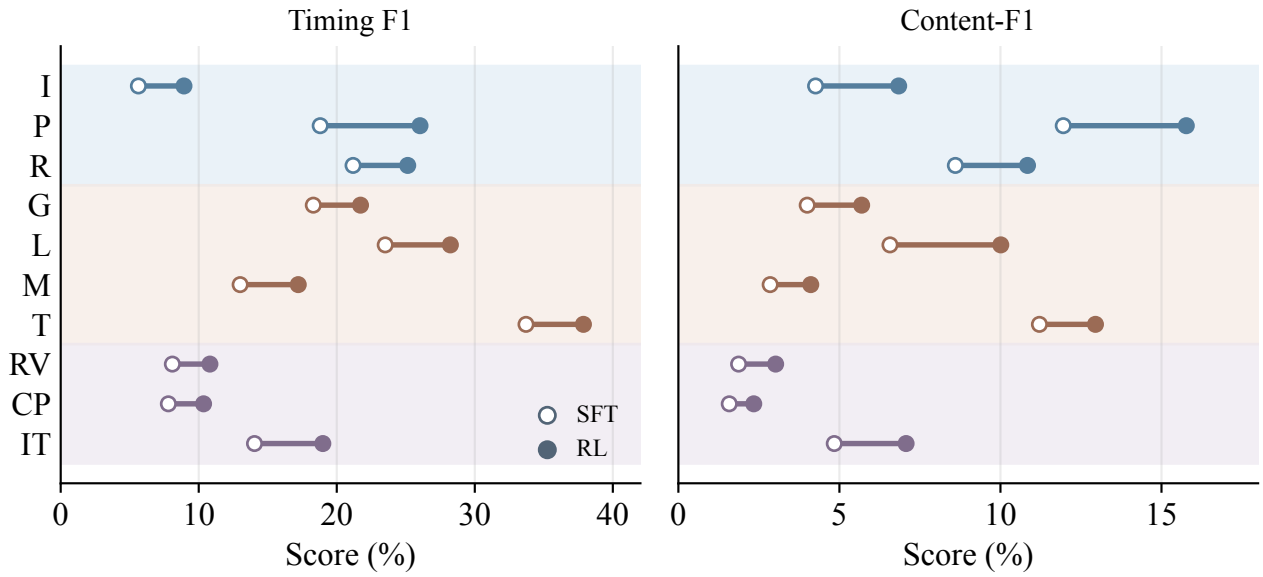

Figure 9 | **Absolute subtype profiles of StreamVibe-VL (SFT) and StreamVibe-VL (RL) under Trajectory evaluation.** Hollow and filled endpoints show SFT and RL, respectively; the connecting segment exposes both direction and magnitude. Background bands encode the three intent types. Scores are percentage points at ±2 s.

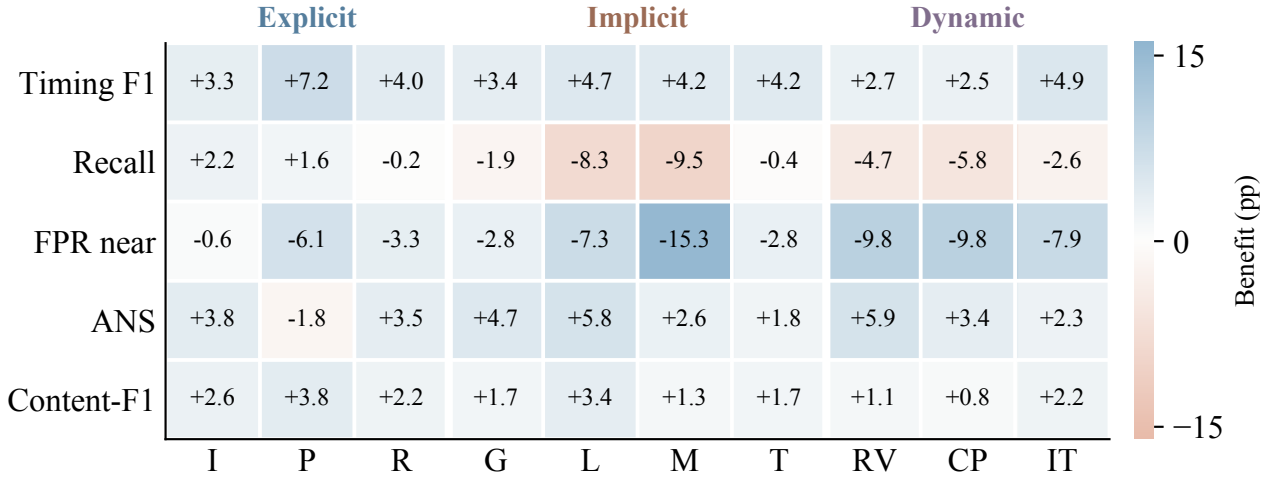

Figure 10 | **How RL changes the operating point by subtype.** Cells report StreamVibe-VL (RL) minus StreamVibe-VL (SFT) in percentage points under Trajectory evaluation. Numbers retain the metric's natural sign; color encodes benefit, so a negative FPR_{near} change is blue. The visualization jointly exposes timing, coverage, false interventions, answer quality, and Content-F1.

RL shifts the policy toward more selective intervention. As shown in Figure 10, RL consistently improves Timing F1 and Content-F1 while reducing FPR_{near} across all subtypes, indicating that its primary effect is to suppress unnecessary triggers and concentrate responses around more appropriate moments. The gains vary across subtypes: P shows the largest Timing F1 improvement of +7.2, while Monitoring exhibits the largest FPR reduction of −15.3, suggesting that RL is particularly effective in settings prone to over-triggering. RL makes proactive behavior more selective and better aligned with appropriate intervention opportunities across subtypes.

B. Dataset Construction and Evaluation Details

This section consolidates the taxonomy, source provenance, training mixture, evaluation cohorts, and metrics summarized in the main paper.

B.1. Overview of Evaluation Datasets

Our evaluation covers proactive interaction and general video understanding. The former tests when to intervene and whether the response supports the ongoing task; the latter checks whether proactive policy learning preserves video perception and temporal reasoning. We describe the datasets below; protocols and metrics are detailed in Appendix B.5.

StreamVibe-Bench. Our primary benchmark evaluates sustained proactive interaction under Explicit, Implicit, and Dynamic Intents. Its 4,100 evaluation samples cover ten subtypes, ranging from responding to specified events to inferring useful assistance and tracking revised or interleaved requests. Pointwise evaluation isolates response decisions at individual timestamps, while Trajectory evaluation tests continuous interaction with model-generated history. The following subsections describe its taxonomy, sources, and construction.

General-Timing. We construct General-Timing as a separate evaluation subset drawn from the publicly released training data of JoyAI-VL-Interaction (Yao et al., 2026b). The selected samples are held out from our training: they are used for neither supervised fine-tuning nor reinforcement learning, and are reserved exclusively for evaluation. This subset tests response timing and content quality on additional streaming-video interactions beyond StreamVibe-Bench. We apply our Pointwise and Trajectory protocols and report results in Tables 9 and 10.

EgoProactive. EgoProactive (Kundu et al., 2026) evaluates proactive procedural assistance from wearable, egocentric videos. Its annotations include out-of-plan deviations and recovery guidance, requiring an assistant to recognize task progress, decide when an interruption is useful, and provide appropriate corrective assistance. We use its official evaluation protocol to assess proactive assistance in physical task execution.

IPIBench. IPIBench (Li et al., 2026b) evaluates interactive proactive intelligence under continuous video streams. It covers proactive monitoring, proactive task management, and interleaved reactive–proactive requests, including multi-turn interactions in which users add, modify, or cancel requests. It complements StreamVibe-Bench by testing the coordination of response timing and evolving interaction requirements under an external protocol. We follow its official evaluation setup; results for IPIBench and EgoProactive are reported in Table 11.

MVBench. MVBench (Li et al., 2024) evaluates video understanding across diverse temporal tasks that require reasoning beyond a single frame. We use it to assess whether proactive policy learning preserves general temporal perception and reasoning, following its official evaluation protocol.

Video-MME. Video-MME (Fu et al., 2025a) assesses video understanding across varied content domains and video durations. It complements MVBench by testing broader video comprehension over different temporal scales. We follow its official evaluation protocol and report the corresponding benchmark score.

OVO-Bench. OVO-Bench (Li et al., 2025b) evaluates timestamp-aware online video understanding through backward tracing, real-time perception, and forward active responding. These settings test whether a model can recall relevant past events, interpret the current observation, and respond as a stream unfolds. We retain its official task definitions and evaluation metrics.

OVBench. OVBench (Huang et al., 2025) evaluates online video understanding as visual observations arrive over time. We use it to assess comprehension of ongoing video streams alongside the proactive interaction benchmarks, following its official evaluation protocol and reported metrics.

B.2. Task Taxonomy and Evaluation Sources

Each example contains a video stream, an initial user instruction, optional timestamped user turns, and one or more verified evidence points. An evidence point stores an acceptable intervention window, a preferred time, and a gold response. The ten subtypes in Table 15 describe what must be tracked rather than merely the visual domain.

Category	Subtype	Operational definition	Samples	Evidence	Hours
Explicit	I	One requested event or answer; verified Explicit Intents pool	200	200	7.5
	P	State transition or irreversible point; verified Explicit Intents pool	200	723	8.2
	R	Repeated occurrences of a requested event; verified Explicit Intents pool	200	794	6.4
Implicit	G	Guidance for the next task action	500	3,702	26.7
	L	Learning and explanation	500	3,287	18.4
	M	State monitoring for task progress and changes	500	2,701	18.9
	T	Tool use in software interfaces	500	2,861	15.3
Dynamic	RV	Revision replaces an active intent; Dynamic Intents construction pool	500	1,885	25.2
	CP	Composition adds a intent while preserving prior ones; Dynamic Intents construction pool	500	2,340	31.2
	IT	Interruption inserts an immediate request while the proactive task remains active; Dynamic Intents construction pool	500	2,053	25.6
		Target balanced evaluation allocation	**4,100**	**20,546**	**183.4**

Table 15 | **StreamVibe-Bench target allocation and source coverage.** "Evidence" counts proactive gold evidence points; reactive questions are excluded. Implicit and Dynamic statistics are measured on their balanced manifests. Explicit evidence-point and duration statistics are proportional estimates from the qualified I/P/R pool at 200 samples per subtype.

Source coverage. Physical-task videos include BARISTA (Knab et al., 2026), Assembly101 (Sener et al., 2022), HoloAssist (Wang et al., 2023), YouCook2 (Zhou et al., 2017), and Ego-Exo4D (Grauman et al., 2024). Learning sources additionally include ActivityNet (Caba Heilbron et al., 2015), Video-MMLU (Song et al., 2025), and Video-MMMU (Hu et al., 2025); monitoring sources include EgoSchema (Mangalam et al., 2023), Perception Test (Pătrăucean et al., 2023), Charades (Sigurdsson et al., 2016), and MLVU (Zhou et al., 2025). GUI sources include AgentNet (Wang et al., 2025b), VideoCUA (Jian et al., 2026), FingerTip (Yang et al., 2026b), GUI-World (Chen et al.,

2025a), and GUI-360 (Mu et al., 2025). Source videos are reused with our intent-specific requests and evidence-point annotations; their original benchmark scores are not used as StreamVibe-Bench labels.

Annotation and provenance. Candidate evidence points are produced from the source video and task record, then checked by four reviewers for request fidelity, visible evidence, timing, and assistance content. Accepted videos are serialized at 1 FPS; at time t, the record exposes only frames and user turns available by t. The normalized metadata retains the source collection and video identifier, duration, subtype, initial and updated requests, evidence-point windows, gold responses, review record, and negative-probe type. Dynamic examples additionally retain the active intent transition, making superseded and newly valid constraints distinguishable.

B.3. Human Review Criteria and Inter-Annotator Agreement

Review unit and acceptance rule. The review unit is a complete candidate sample: the source video, initial instruction, timestamped user turns, evidence windows, preferred response times, and reference answers. Four human reviewers independently assess the same sample against Table 16. A reviewer approves a sample only when all applicable criteria are satisfied; a required correction counts as non-approval of the current version. A sample is retained only if all four reviewers approve the same final version. Thus, three approvals are insufficient. Samples with any objection must be corrected and reviewed again by all four reviewers, or discarded; changing an annotation invalidates approval of its earlier version. Unresolvable ambiguity or insufficient visual evidence requires exclusion rather than a guessed reference answer.

Table 16 | **Human review rubric.** Each reviewer checks every applicable criterion. Acceptance requires four independent approvals of the same sample version, with no unresolved issue.

Criterion	Standard for approval
Intent fidelity	The instruction is coherent and compatible with the task. The assigned Explicit, Implicit, or Dynamic category matches the interaction; inferred intents are supported by the user's context, not invented.
Evidence and causality	Every annotated intervention has identifiable visual or interaction evidence available by its response time. Future frames and later user turns are not needed to justify the response.
Coverage and non-redundancy	Distinct, necessary intervention opportunities are covered, without duplicate evidence points or repeated assistance after the relevant intent has already been satisfied.
Temporal validity	Each response window and preferred timestamp correspond to a moment when the evidence is sufficient and assistance is still useful; the preferred time lies within the window and neither anticipates evidence nor misses the actionable moment.
Reference-answer quality	The reference is correct, relevant, and useful for the active intent. It avoids unsupported details, contradictory advice, and unnecessary instructions, while remaining consistent with the visible state.
Dynamic-intent consistency	For Dynamic samples, timestamped revisions, compositions, and interruptions are coherent. References follow the latest valid intent and preserve unaffected constraints without acting on superseded ones; this criterion is not applicable to other categories.
Record integrity	The video, source identifier, user-turn timestamps, evidence windows, and reference answers are complete and aligned. Each decision refers to an identifiable sample version.

Independent agreement audit. Unanimous acceptance is a quality-control rule, not itself evidence of high inter-annotator agreement. The agreement audit must use the original, independent judgments on a fixed set of candidate samples *before* discussion, correction, or acceptance filtering, including candidates that are not retained. Each reviewer supplies a binary decision, approve or require revision/reject, under the same rubric. Repeated versions of a candidate must not be counted as independent samples. The audit should record the candidate count and decision distribution for each intent type.

Human–human agreement. Fleiss' κ (Fleiss, 1971) measures chance-corrected agreement among the four reviewers. For N candidates, let n_{ic} be the number of reviewers assigning candidate i to decision $c \in \{0, 1\}$, with $\sum_c n_{ic} = 4$. Then

$$\bar{P} = \frac{1}{N} \sum_{i=1}^{N} \frac{\sum_c n_{ic}(n_{ic} - 1)}{4 \cdot 3}, \qquad p_c = \frac{\sum_i n_{ic}}{4N}, \qquad \kappa_{\mathrm{F}} = \frac{\bar{P} - \sum_c p_c^2}{1 - \sum_c p_c^2}.$$

The audit should additionally report the raw pairwise agreement $\bar{P}$ and the fraction of candidates with four identical decisions. Computing κ only on retained, unanimously approved samples is invalid as evidence of general reviewer reliability: filtering enforces agreement, and a single decision class makes the chance-correction denominator zero.

Gemini–human agreement. For this comparison, Gemini-3.8-Flash must independently assign the same binary review decision to the same frozen candidates, using the rubric without access to human judgments or post-review corrections. Its original candidate annotations alone are not binary audit decisions. Compute Cohen's κ (Cohen, 1960) separately between Gemini and each reviewer, $\kappa_{\mathrm{C}}^{(r)} = (p_o^{(r)} - p_e^{(r)})/(1 - p_e^{(r)})$, where $p_o^{(r)}$ is observed agreement and $p_e^{(r)}$ is agreement expected from their respective decision marginals. Report the mean of the four pairwise coefficients and their range, rather than treating the four human judgments as one independent rater.

Table 17 | **Inter-annotator agreement.** Human–human agreement is measured by Fleiss' κ; Gemini–human agreement is the mean of four pairwise Cohen's κ coefficients, with ranges in parentheses.

Audit subset	Candidates	Human–Human Fleiss' κ ↑	Gemini-3.8-Flash–Human Cohen's κ ↑ (mean of four)
Explicit	872	0.87	0.79 (0.76–0.82)
Implicit	2,946	0.84	0.75 (0.72–0.78)
Dynamic	2,187	0.82	0.73 (0.70–0.77)
All candidates	6,005	0.84	0.76 (0.70–0.82)

Analysis. As shown in Table 17, human–human Fleiss' κ ranges from 0.82 to 0.87 across all categories, falling within the "almost perfect agreement" band (>0.81) (Landis and Koch, 1977), indicating that the four reviewers maintained highly consistent judgments under the shared rubric. Explicit achieves the highest agreement (0.87), while Dynamic is slightly lower (0.82), consistent with the inherently greater complexity of annotating evolving intents. Gemini–human Cohen's κ ranges from 0.73 to 0.79, within the "substantial agreement" band (0.61–0.80), with narrow per-pair ranges, indicating that Gemini's review decisions are well aligned with human judgments. Together, these results confirm the reliability and quality of the StreamVibe-Bench annotation pipeline.

B.4. Training Mixture and Evaluation Cohorts

Table 18 reports the SFT data manifest containing 70K task-aligned examples and 768K general streaming examples. Task-aligned Explicit and Implicit examples preserve their original task structure, while the general streaming mixture broadens coverage across visual domains, narration styles, temporal scales, and GUI scenarios. Localized RL starts from a separate balanced 15K seed pool, consisting of 5K tasks each from Implicit Intents, Explicit Intents, and general streaming data.

Training component	Recorded source families	Examples	Use
Task-aligned Explicit	Explicit Intents event families	35,000	Streaming SFT
Task-aligned Implicit	Guidance, learning, monitoring, and tool/GUI	35,000	Streaming SFT
General streaming core	JoyAI mixture: Live-WhisperX, Shot2Story, TimeLens, GUI-World, VideoGPT+, Charades-Ego, InternVid, ActivityNet, WebVid, and smaller sources	763,607	Streaming SFT
General streaming auxiliary	TimeChat-Online, VideoChat-Flash, TimeLens, EgoIT, and StreamO-Instruct	4,393	Streaming SFT
	Full SFT manifest	**838,000**	**1 FPS streams**
Localized RL seed	Implicit / Explicit / general streaming, 5K each	15,000	Local rollout windows

Table 18 | **Training-data provenance.** The SFT mixture contains 70K task-aligned examples and 768K general streaming examples. Localized RL uses a separate 15K seed pool, yielding 85K task-oriented training examples in total when combined with task-aligned SFT. Dataset names indicate source families and do not imply a shared annotation protocol.

General-streaming sources. The inherited mixture includes Live-WhisperX (Chen et al., 2025b), Shot2Story (Han et al., 2025), TimeLens (Zhang et al., 2026), VideoGPT+ (Maaz et al., 2024), Charades-Ego (Sigurdsson et al., 2018), InternVid (Wang et al., 2024), WebVid (Bain et al., 2022), and the ActivityNet and GUI-World sources cited above. Auxiliary data comes from TimeChat-Online (Yao et al., 2025), VideoChat-Flash (Li et al., 2025a), TimeLens, EgoIT-99K (released with EgoLife) (Yang et al., 2026a), and Streamo-Instruct (Xia et al., 2026).

B.5. Evaluation Protocols and Metrics

Temporal and content scores. At tolerance τ, predictions and evidence points within $\pm\tau$ are matched chronologically and one-to-one. Thus $P = \mathrm{TP}/N_{\mathrm{pred}}$, $R = \mathrm{TP}/N_{\mathrm{gold}}$, and F1 is their harmonic mean; $\mathrm{FPR}_{\mathrm{near}}$ measures responses on designated nearby negative probes. For each temporally matched response, Claude Sonnet 4.6 (Anthropic, 2026) assigns a 0–4 content score r that is normalized to $s_{g,p} = r/4 \in [0, 1]$. ANS is its mean over matched responses. Content-F1 is $2\sum_{(g,p)\in\mathcal{M}} s_{g,p}/(N_{\mathrm{pred}} + N_{\mathrm{gold}})$, so missed intents and extra interventions remain penalized instead of disappearing from a conditional answer average. For Dynamic Intents, ANS and Content-F1 score proactive responses only; reactive answers remain in the interaction history.

Evaluation	What is scored	Reported measures	Judge / scorer
Pointwise	Independent probes; supplied history	Precision, recall, F1, FPR_{near}, ANS, Content-F1	Temporal matcher; Claude Sonnet 4.6 (Anthropic, 2026)
Trajectory	Complete 1 FPS rollout; generated history	Same six scores at ±2, ±1, ±0 s	Same matcher and judge
General Timing	Streaming transfer cohort	Same timing and content scores	Same protocol
EgoProactive / IPIBench	Session-level interaction	Macro, G-mean, Interrupt, Silent; Timing, Understanding, Repeated	Official protocols
Video suites	MVBench, Video-MME, OVO-Bench, OVBench	Accuracy or official aggregate	Official scorers

Table 19 | **Metrics and judges used in the paper.** External benchmarks retain their official evaluation protocols and are not rescored by the StreamVibe-Bench content judge.

C. StreamVibe-Harness Implementation Details

This section provides implementation details for the four functional components introduced in Section 4.1. StreamVibe-Harness maintains persistent interaction state and organizes causal evidence for the interaction policy. These components jointly support a single interaction process and should not be interpreted as four independent models. We use Qwen3-VL-4B-Instruct (Bai et al., 2025) as an auxiliary visual-memory model for memory summarization and consolidation with greedy decoding, while the StreamVibe-VL checkpoint under evaluation serves as the interaction policy. The auxiliary model is used off-the-shelf and receives no additional training.

Table 20 | **Harness configuration.** Temporal parameters correspond to the chunked-streaming setting; benchmark-specific native adapters are described separately below. Retrieval limits denote maximum capacities rather than fixed numbers of items injected into the policy context.

Component	Setting
Memory model	Qwen3-VL-4B-Instruct; separate inference process
Memory input	At most 64 frames per call; 262,144 pixels per frame
Chunk and submission schedule	T_s = 32 steps; submission lookahead = 8 steps
Consolidation and retention	M = 5 mid-term summaries; L = 8 long-term blocks
Evidence selection	Up to 3 facts and 3 ordered transitions
Compact state evidence	First rendered evidence line, capped at 180 characters
Skill recall	10 fixed Skills; at most 1 selected per context construction
Procedural memory	Up to 8 completed steps and 12 temporal anchors

C.1. Interaction State

At time t, the interaction state is represented as $\mathcal{S}_t = \langle \mathcal{A}_t, \mathcal{P}_t, \mathcal{R}_t \rangle$, where $\mathcal{A}_t$ stores active intents, $\mathcal{P}_t$ records observed task progress, and $\mathcal{R}_t$ tracks responses that have actually been delivered. State changes are represented as typed events with an identifier, timestamp, source, and payload. A deterministic update function applies valid events to the corresponding state fields and ignores duplicate updates. User instructions may introduce, revise, or retire requests; visual observations may establish task-relevant events; and observed progress or explicit user confirmation may update procedural state. A delivery record is created only after a response has been emitted to the user.

We enforce source-specific update constraints so that information retains its semantic role across the interaction. In particular, a request specifies what should be monitored rather than what has already occurred; a reference plan describes possible future steps rather than completed actions; and an assistant suggestion is not treated as evidence of user progress. Fields without sufficient support remain unresolved. Historical evidence may be exposed to the policy as context, but cannot by itself create, revise, or retire an intent.

Procedural memory retains at most eight completed steps and twelve evidence intervals that serve as temporal retrieval anchors. Progress is updated only when supported by observation or explicit confirmation. For event-monitoring requests, the auxiliary model additionally extracts and normalizes evidence from causal frames within a two-second window. This event parser updates evidence and event state only; it does not generate user-facing responses and therefore does not constitute a second action policy.

C.2. Intent Lifecycle

Each active intent is assigned a stable identifier and maintains an independent record of its criterion, associated event, evidence reference, and delivery status. A newly registered criterion begins in an evidence-waiting state. Once sufficient evidence supports a match, the criterion is linked to a concrete event instance. If the policy subsequently responds, the resulting delivery record identifies both the addressed event and the associated intents. This design allows multiple intents to remain active concurrently without treating a response to one intent as completion of the others.

The lifecycle explicitly separates *request validity* from *event coverage*. Answering a one-time question closes that pending question. By contrast, reporting an occurrence under a standing request marks only that event instance as addressed while keeping the request active for future evidence. Closing an event returns the corresponding instance to an evidence-waiting state, whereas explicit user cancellation retires the request itself. These transitions update interaction bookkeeping rather than policy behavior: neither an active intent nor a matched event deterministically triggers a response.

Skill library and selection. The Harness maintains ten fixed, read-only Skill records covering visual question answering, multiple-choice answering, event detection, explanation, hazard monitoring, milestone reporting, mistake monitoring, narration, next-step guidance, and stage summarization. Each Skill specifies request cues, evidence to inspect, response and waiting conditions, interaction cadence, and answer form. Skills are reusable behavioral specifications rather than neural models. Each active criterion maintains its own execution instance, including phase, evidence reference, event association, and delivery status.

Skill selection first restricts candidates to those compatible with the current interaction type and requires non-empty observed memory, procedural context, or event context. A pending question prioritizes the question-answering Skill, followed by direct matches to active criterion types. If neither applies, lexical overlap between the request and observed context is used as a fallback, with at least two shared terms required. At most one Skill is selected for each context construction, and selecting no Skill is a valid outcome. For example, next-step guidance is conditioned on observed task progress, whereas event detection tracks a requested transition and whether the corresponding occurrence has already been reported. Importantly, selecting a Skill does not imply that its response condition has been satisfied.

C.3. Hierarchical Memory

Storage and consolidation. The memory system contains three levels: recent raw visual context, timestamped mid-term summaries, and older consolidated text blocks. Under the chunked-streaming configuration, we use $T_s = 32$ steps per chunk, consolidate every $M = 5$ mid-term summaries, and retain up to $L = 8$ long-term blocks. At one step per second, each chunk corresponds to 32 seconds of video. The auxiliary model processes frames from the short-term buffer in chronological order, uniformly subsampling when the 64-frame input limit is exceeded, and conditions each summary on the current request. Consolidation of older summaries is performed using text-only input.

Mid-term summaries preserve task-relevant object states, quantities, on-screen text, and temporal transitions. Intervals without meaningful change are collapsed, unsupported attributes remain unknown, and the most recently observed state is retained when visually established. Each ordered transition records its temporal interval, actor, action, object, prior and resulting states, temporal relation, and confidence. The mid-term summarization prompt is capped at 220 words, while procedural fields distinguish observed current and completed steps from request-supported references to possible next steps.

Every M consecutive mid-term summaries are consolidated into one long-term block using a 300-word prompt budget. Consolidation removes redundant static descriptions while preserving timestamps, distinct actions, temporal order, and supported procedural information. Each mid-term summary belongs to exactly one non-overlapping consolidation group. Once the number of long-term blocks reaches the capacity L, older blocks are evicted in first-in, first-out order.

To preserve causality under asynchronous execution, summarization is submitted before the chunk boundary using only frames already observed at submission time. Frames arriving afterward remain in the next short-term buffer, preventing gaps in temporal coverage. Mid-term inputs remain available while long-term consolidation is pending. If consolidation fails or produces an empty output, their concatenation is retained instead, preserving the available evidence without assuming that generated summaries are lossless.

Indexing and retrieval. Summary facts are indexed using normalized text together with stable identifiers, source intervals, and lexical terms. Repeated observations of the same fact update its observation interval rather than creating duplicate entries. Ordered transitions are indexed separately to preserve event sequence. Request descriptions and reference-plan fields are excluded from the factual index, and response records remain separate from visual evidence.

At retrieval time, the Harness constructs a focus string from the pending question, current request, and relevant procedural state. Candidate entries are ranked primarily by lexical overlap with this focus and secondarily by recency, with additional relevance assigned to entries overlapping task-relevant temporal anchors. Facts and ordered transitions are selected independently, with up to three candidates from each category and at most six entries before context projection. If no candidate matches the focus, the most recent entries are retained as a fallback. Selected transitions are reordered chronologically before projection. For native event-monitoring adapters, target retrieval is further restricted to at most three matching evidence lines, excluding negative or uncertain statements. No embedding model or learned retriever is used.

The retrieval capacity is intentionally decoupled from the amount of evidence exposed to the policy. In the compact lifecycle projection, only the first rendered evidence line is retained and truncated to 180 characters; this line may contain a joined excerpt of multiple selected facts. For a pending question, the bounded fact and transition blocks may instead be exposed directly. Retrieved items are always treated as timestamped historical context, and current observations take precedence whenever

they supersede an earlier state.

C.4. State Integration

Before each policy call, the Harness compiles the relevant state fields and retrieved evidence into a compact textual context and combines it with the causal video input,

$$\tilde{c}_t = \phi(c_t, \mathcal{S}_t, m_t),$$

where c_t denotes the causal interaction context, $\mathcal{S}_t$ the persistent interaction state, and m_t the retrieved memory. The projected context explicitly distinguishes pending questions, observed progress, event status, delivery records, uncertainty, and historical evidence. Inactive or empty fields are omitted, and the full memory store is not appended by default. Within a context epoch, the header remains fixed while subsequent observations and actions are appended, allowing the policy to reuse the shared prompt prefix.

Skill Projection for Policy Inference. Although each Skill stores a complete behavioral specification, the policy receives only the subset required for the current interaction. For single-criterion guidance, monitoring, and explanation, the projected context includes the selected Skill's execution phase, evidence reference, response basis, and a compact cadence description. Question-answering and multi-criterion contexts expose only the answer-form field, while single-criterion narration retains a compact symbolic representation. The complete Skill specification is therefore not copied into the prompt, and Skill metadata does not replace policy generation. Request changes update the interaction type causally; after an intervening question is answered, the context of the standing request is restored.

The interaction policy π remains the sole component that decides whether to remain silent or respond and, when responding, generates the user-facing text. Once a non-empty response is delivered, its timestamp and intent/event associations are committed to $\mathcal{R}_t$, while its exact text is preserved in the interaction history. Subsequent context projections may retain the most recent response verbatim and compress older responses into action markers. Assistant outputs are excluded from visual summarization. Response verification, cadence gating, and post-generation rewriting are disabled. Consequently, this feedback path records delivered assistance for future state construction without introducing new visual evidence or modifying the policy output.

D. Qualitative Examples and Case Studies

This section presents annotated examples from all ten task subtypes, followed by additional comparisons of model responses during sustained interaction.

D.1. Representative Samples from Every Subtype

Figures 11–30 show two verified samples per subtype. Each full-width sample contains the original instruction, a chronological six-frame filmstrip, and timestamped reference assistance. Three annotated interventions are shown per sample, except for single-event I, whose definition permits only one gold evidence point. Each pair of adjacent frames shows the second before an intervention and its annotated time; gaps between pairs are indicated by the timestamps. The outlined frame marks the intervention. Dynamic examples include the user updates applicable to the displayed timeline.

I | One requested event — Example A

User instruction: Notify me as soon as the man throws the book onto the armchair.

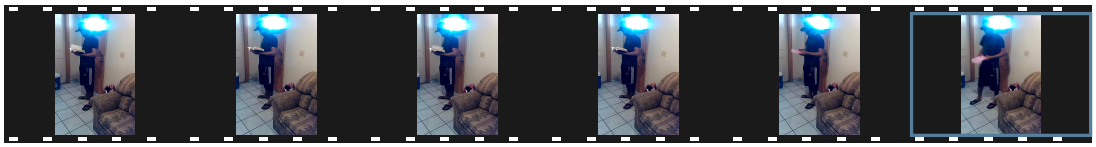

Evidence 1 (6s): The man is throwing the book onto the armchair.

Figure 11 | **Explicit intent, single event (I), example A.**

I | One requested event — Example B

User instruction: Notify immediately when a woman holding a beverage cup appears in the office.

Evidence 1 (24s): A woman holding a beverage cup has entered the office area.

Figure 12 | **Explicit intent, single event (I), example B.**

P | State transition — Example A

User instruction: Monitor the state of the playing card deck and report whenever it transitions between being a single stack and being split into two stacks.

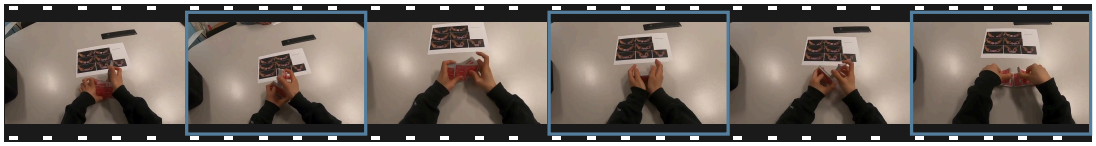

Evidence 1 (4s): The deck has transitioned from a single stack to two stacks.

Evidence 2 (10s): The deck has transitioned from two stacks to a single stack.

Evidence 3 (14s): The deck has transitioned from a single stack to two stacks.

Figure 13 | **Explicit intent, state transition (P), example A.**

P | State transition — Example B

User instruction: Monitor the raccoon and identify transitions in its eating state (eating vs not eating).

Evidence 1 (14s): The raccoon has finished the kibble and stopped eating.

Evidence 2 (38s): The raccoon has reached the bowl on the counter and started eating.

Evidence 3 (46s): The raccoon was pulled away from the food bowl and stopped eating.

Figure 14 | **Explicit intent, state transition (P), example B.**

R | Repeated occurrences — Example A

User instruction: Notify me every time the man takes a sip from his mug.

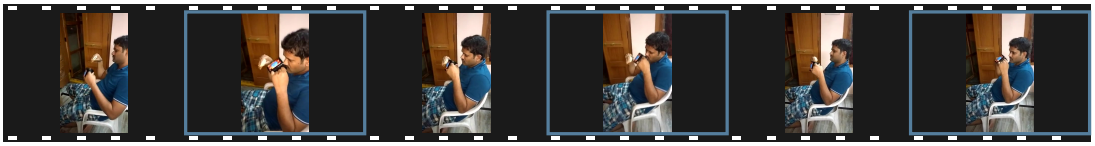

Evidence 1 (8s): The man has started taking his first sip from the mug.

Evidence 2 (17s): The man has started taking a second sip from the mug.

Evidence 3 (28s): The man has started taking a third sip from the mug.

Figure 15 | **Explicit intent, repeated event (R), example A.**

R | Repeated occurrences — Example B

User instruction: Notify me each time the mechanic spins the bicycle's front wheel.

Evidence 1 (44s): The mechanic spins the front wheel for the first test.

Evidence 2 (59s): The mechanic spins the front wheel for a second inspection.

Evidence 3 (88s): The mechanic spins the front wheel for a third check after tool adjustment.

Figure 16 | **Explicit intent, repeated event (R), example B.**

G | Hands-on guidance — Example A

User instruction: Walk me through using this capsule espresso machine step by step as I prepare my drink, and let me know if I do something incorrectly.

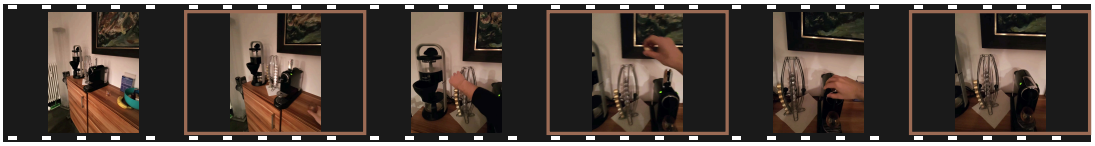

Evidence 1 (3s): Now go ahead and pick a coffee capsule from the holder, then place it into the top slot of the machine.

Evidence 2 (8s): Now push the top lever all the way down to lock the capsule firmly into place.

Evidence 3 (13s): Now press one of the top brew buttons to start extracting your espresso shot.

Figure 17 | **Implicit intent, guidance (G), example A.**

G | Hands-on guidance — Example B

User instruction: Guide me step by step through assembling this metal plant stand tray table, and stop me if I make any assembly errors.

Evidence 1 (120s): Cross the two leg frames into an X shape and line up the center screw holes where they intersect.

Evidence 2 (136s): Cross the two frame pieces at the central notch to form the X-base, and make sure the middle screw holes line up.

Evidence 3 (165s): Before inserting the central screw, double-check that the holes on both interlocked frame legs line up straight.

Figure 18 | **Implicit intent, guidance (G), example B.**

L | Learning — Example A

User instruction: Walk me through this hedge trimming tutorial, explaining the core trimming techniques, and leaf inspection methods.

Evidence 1 (3s): Before turning on the trimmer, taking stock of your garden layout helps you plan your cutting path and ensure your extension cord stays safely out of the way.

Evidence 2 (59s): Hold the trimmer firmly with both hands for full control, keeping the blade horizontal and parallel to the surface so you can guide it in smooth, sweeping motions.

Evidence 3 (120s): Checking the freshly trimmed surface by hand lets you feel for hidden uneven spots and ensure the blades made clean cuts rather than tearing the leaves.

Figure 19 | **Implicit intent, learning (L), example A.**

L | Learning — Example B

User instruction: Walk me through this visual sum of cubes proof, and explain how the tile geometry translates into the algebraic summation formula.

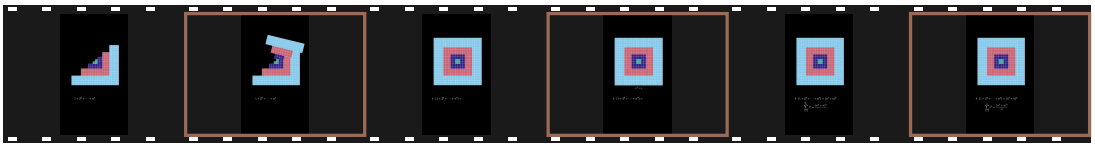

Evidence 1 (18s): Combining four identical tile sets fills in the open gaps to create a complete, symmetrical square, making it much easier to express the total area algebraically as four times the sum of ...

Evidence 2 (24s): That $n^2 + n$ side length comes from combining the dimensions of all n tile layers across the interlocked square.

Evidence 3 (39s): Since four identical tile stacks were combined to form that large square, dividing its total area by four isolates the formula for a single sum of cubes.

Figure 20 | **Implicit intent, learning (L), example B.**

M | Monitoring — Example A

User instruction: Keep an eye on this living room camera feed, noting human movements coming through and any pets getting onto furniture or causing potential household disruptions.

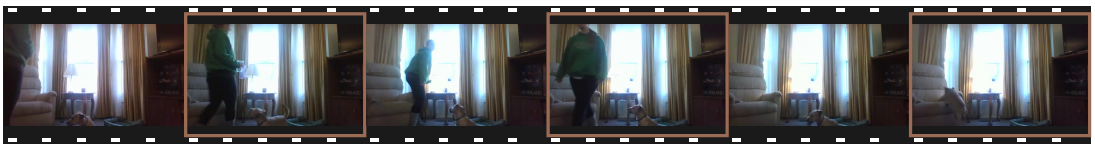

Evidence 1 (2s): A person holding a spray bottle just entered the living room.

Evidence 2 (10s): Watch out — wet cleaning spray was left on the window glass, and the dog is now unattended in the room.

Evidence 3 (30s): The dog has climbed up onto the recliner armchair.

Figure 21 | **Implicit intent, monitoring (M), example A.**

M | Monitoring — Example B

User instruction: Monitor this doorway safety test feed for items encroaching on the egress pathway or hazards blocking the partition door from closing smoothly.

Evidence 1 (4s): A purple ball was rolled into the doorway, directly obstructing the partition door path.

Evidence 2 (18s): A blue neck pillow was thrown into the doorway near the box, adding to pathway hazards.

Evidence 3 (25s): The green tinsel on the floor is getting caught directly beneath the closing partition door.

Figure 22 | **Implicit intent, monitoring (M), example B.**

T | Tool / GUI operation — Example A

User instruction: Walk me through cloning a repository in GitHub Desktop and managing my local folders, and steer me through switching between active projects.

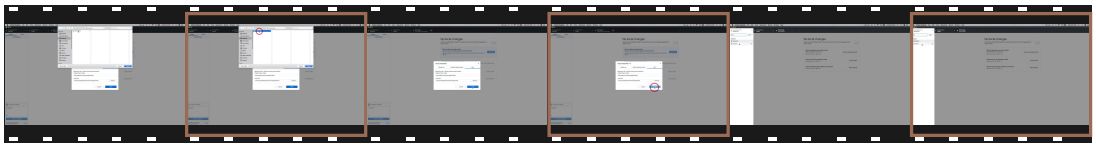

Evidence 1 (9s): Double-check that 'TestProject' is the intended target folder before confirming your selection in the file dialog.

Evidence 2 (19s): The cloning process is complete, and 'githubTest' has been automatically set as your active repository in the top bar.

Evidence 3 (39s): You've switched back to LiveALife as your active repository, so your original project workspace is restored and ready.

Figure 23 | **Implicit intent, tool / GUI (T), example A.**

T | Tool / GUI operation — Example B

User instruction: Walk me through creating a personalized workload view in Asana, and guide me if I misstep setting up board filters or tab options.

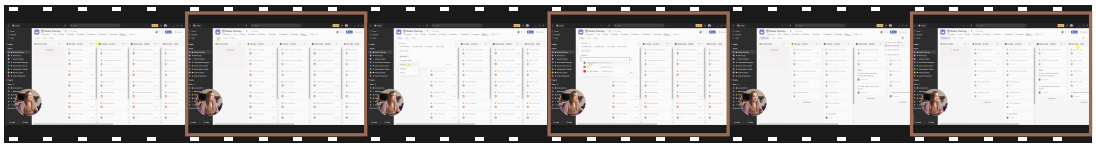

Evidence 1 (21s): Select 'Just my tasks' from the quick filter menu to show only items assigned to you.

Evidence 2 (24s): Just double-check that 'Just my tasks' is selected for Natalia Kalinska under Assigned to before moving on.

Evidence 3 (30s): Your board view is now filtered to show just your assigned tasks, and the layout has been saved.

Figure 24 | **Implicit intent, tool / GUI (T), example B.**

RV | Revision — Example A

User instruction: Please provide progress updates as I disassemble and rebuild the toy bulldozer.

User update (182s): Stop giving progress updates. Instead, check the quality of my assembly to ensure parts are aligned and properly secured.

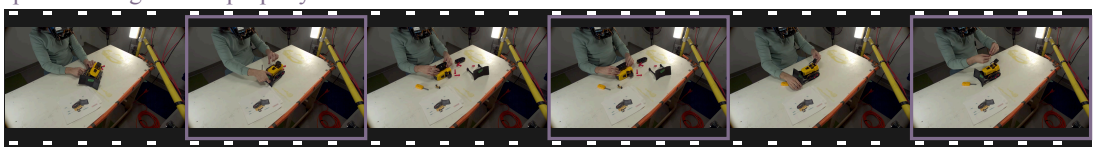

Evidence 1 (36s): The front shovel blade has been completely detached from the bulldozer chassis.

Evidence 2 (259s): Before driving the screws in, double-check that the side track is sitting completely flush against the chassis frame.

Evidence 3 (329s): Looks like everything is back together. Give it a quick push across the desk to make sure the tracks roll smoothly and nothing is catching.

Figure 25 | **Dynamic intent, revision (RV), example A.**

RV | Revision — Example B

User instruction: Please provide next-step guidance as I make this lemon garlic aioli.

User update (79s): Stop with the next steps. Now, verify the quality of my aioli mixture.

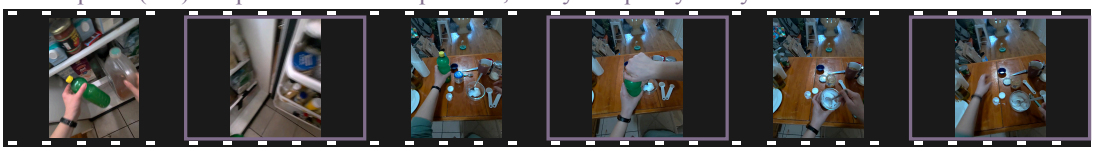

Evidence 1 (6s): Now that you have the mayo and lemon juice, grab your garlic and a small mixing bowl to start building the base.

Evidence 2 (78s): Next, measure out about a teaspoon of the lemon juice and pour it into the mayonnaise.

Evidence 3 (152s): That looks perfectly smooth and thoroughly combined.

Figure 26 | **Dynamic intent, revision (RV), example B.**

CP | Composition — Example A

User instruction: Let me know if I make a mistake while tying this knot.

User update (10s): Also, notify me whenever I complete a key step or finish the knot.

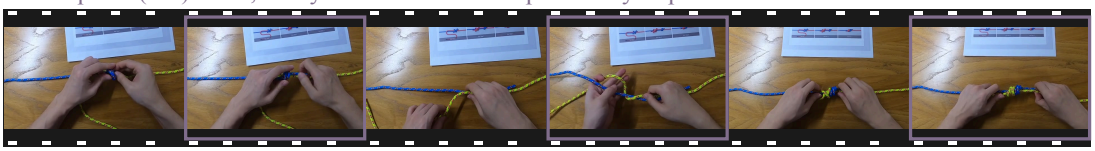

Evidence 1 (24s): You have completed the first knot.

Evidence 2 (36s): [...] wrap the yellow cord twice around both strands, moving back towards the blue knot, before threading the end through.

Evidence 3 (51s): You've finished tying the double fisherman's bend.

Figure 27 | **Dynamic intent, composition (CP), example A.**

CP | Composition — Example B

User instruction: Could you guide me on the next steps while I replace the RAM in my desktop computer tower?

User update (99s): Also, please prompt me to do a quality check whenever I complete an installation step.

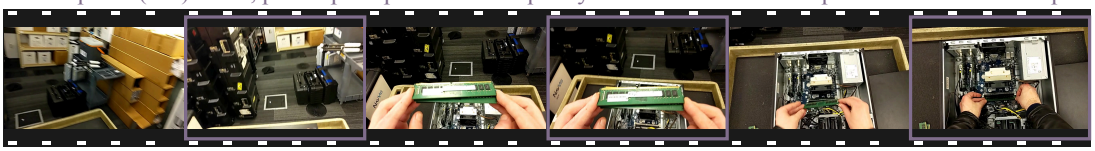

Evidence 1 (76s): You've finished placing the desktop tower flat. Next, remove the side panel.

Evidence 2 (136s): You've finished inspecting the RAM stick. Next, align its notch with the ridge inside the motherboard slot.

Evidence 3 (160s): You've finished placing the RAM stick into the slot channels. Next, press straight down firmly on both ends.

Figure 28 | **Dynamic intent, composition (CP), example B.**

IT | Interruption — Example A

User instruction: I'm assembling this Nintendo Switch setup for the first time. Let me know if I attach a controller wrong or insert something incorrectly.

User update (31s): Are the Joy-Cons attached to the correct sides of the console?

User update (76s): Did I put the Joy-Cons on the correct sides of the grip?

30s 31s 51s 52s 95s 96s

Immediate answer 1 (31s): Yes, the blue Joy-Con is on the left and the red one is on the right.

Evidence 2 (52s): You attached the wrist strap upside down; align the plus and minus signs properly.

Evidence 3 (96s): You're inserting the microSD card into the game card slot; put it in the slot under the kickstand instead.

Figure 29 | **Dynamic intent, interruption (IT), example A.**

IT | Interruption — Example B

User instruction: I'm putting together this desk riser stand. Let me know if I attach any frame pieces upside down or miss a screw.

User update (48s): How many screws are left on the table?

User update (112s): What tool am I using to tighten the screws?

User update (174s): How many legs have I attached so far?

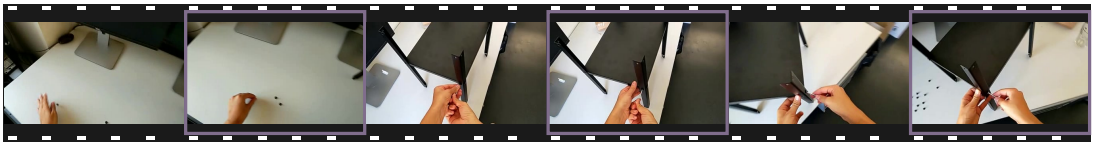

47s 48s 133s 134s 187s 188s

Immediate answer 1 (48s): There are 3 screws left on the table.

Evidence 2 (134s): You're attaching that leg bracket backwards. Turn it around so the lip faces inward.

Evidence 3 (188s): You're attaching that leg piece upside down. Flip it so the mounting holes align correctly.

Figure 30 | **Dynamic intent, interruption (IT), example B.**

D.2. Additional Qualitative Case Studies

Figures 31–38 compare responses during sustained interaction. StreamVibe uses task progress, interface state, and user goals to give timely, concrete guidance. Baselines more often respond too early or too late, miss key evidence, repeat earlier advice, or produce guidance inconsistent with the current interface state.

Case 01. Drawing and saving shapes in GIMP

Instruction: Walk me through drawing and coloring some basic shapes in GIMP, and make sure I save the project correctly when I'm done.

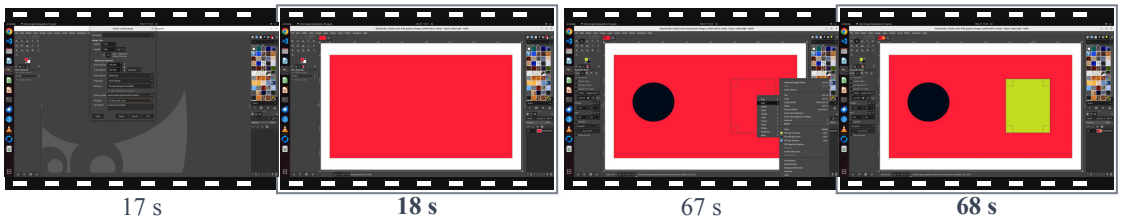

	18 s	68 s
JoyAI	✗ 20s: Great job! Now, let's add some color. Click the Paintbrush tool again and choose a different color from the color picker.	✗ 68s: Now, click the Paintbrush tool again to draw your shape. You can adjust the brush size in the tool options if needed.
MMDuet2	✗ No response within ±2 s	✗ No response within ±2 s
ViSpeak	✗ 19s: The video shows a computer screen with a red background. The video is shot from a first-person perspective, with the cursor …	✗ 68s: The video shows a person drawing and coloring basic shapes in GIMP. The person first draws a circle and fills it with black, then draws a …
StreamVibe (ours)	✓ 18s: Your canvas is set up. Now you can select the shape tool from the left toolbar to start drawing.	✓ 69s: Both shapes are drawn and colored. Now you can save the project by going to File and selecting Save.

Figure 31 | **Additional case study 01.** Drawing and saving shapes in GIMP. Model responses at two annotated evidence points, matched within ±2 s.

Case 02. Editing and exporting a video in Filmora

Instruction: Walk me through editing this green smoothie video in Filmora, exporting it, and setting up the sharing permissions in Google Drive.

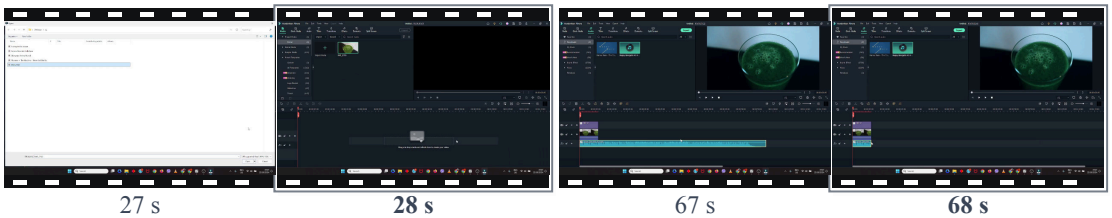

	27 s	67 s
JoyAI	✗ 28s: The user returns to the Filmora editing interface, where the video is still in the timeline.	✗ 68s: The user continues editing the video, adjusting the timeline and previewing the footage.
MMDuet2	✓ 29s: … text instructing to drag and drop media and effects to create a video.	✓ 66s: … the user is about to click the export button to save the project.
ViSpeak	✗ 27s: The video shows a computer screen with a Google Drive folder open. The user clicks on the "Share" button and selects "Invite People" …	✓ 66s: The video shows a person editing a green smoothie video in Filmora. The person then exports the video and sets up the sharing permissions in …
StreamVibe (ours)	✓ 30s: Drag the green smoothie clip from the media library onto the timeline to begin editing.	✓ 70s: Click the 'Export' button in the export settings window to save your video.

Figure 32 | **Additional case study 02.** Editing and exporting a video in Filmora. Model responses at two annotated evidence points, matched within ±2 s.

Case 03. Setting a French-learning reminder

Instruction: Guide me through switching to the French course in Duolingo and creating a daily study task in Microsoft To Do on this tablet.

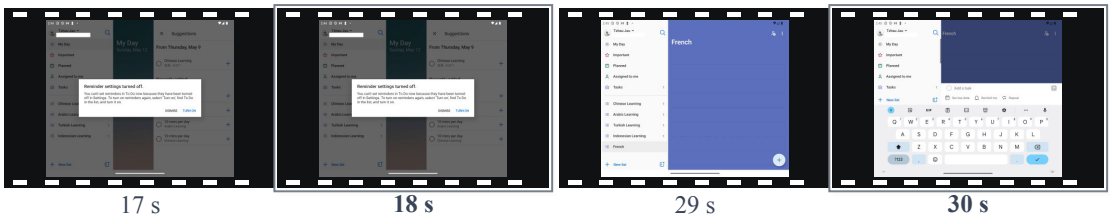

	17 s	29 s
JoyAI	✗ 17s: Now switch to the Microsoft To Do app. Open it from the home screen.	✗ 30s: You're back in Duolingo. Complete the lesson by selecting the correct image for "le chat."
MMDuet2	✗ No response within ±2 s	✗ 29s: After typing 'French' in the 'Enter list title' field, the 'CREATE LIST' button is clicked, resulting in a new 'French' list appearing on the right …
ViSpeak	✗ 19s: To switch to the French course in Duolingo, tap the French flag on the screen. Then tap the "Add" button on the screen. Then tap the "Add task" …	✗ 29s: The video shows a person using a tablet to switch to the French course in Duolingo and create a daily study task in Microsoft To Do on the tablet.
StreamVibe (ours)	✓ 18s: Tap 'TURN ON' to enable reminders for your daily study task.	✓ 30s: Type '10 min per day' into the task field.

Figure 33 | **Additional case study 03.** Setting a French-learning reminder. Model responses at two annotated evidence points, matched within ±2 s.

Case 04. Finding and selecting jeans in a shopping app

Instruction: Walk me through finding a pair of jeans on this shopping app, picking my size and color, and adding them to my cart.

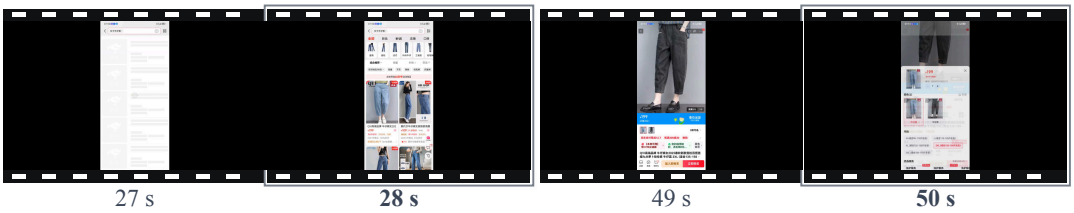

JoyAI	✗ 28s: The user taps the "Add to Cart" button, and the interface switches to the shopping cart page.	✗ 50s: The user taps the "Place Order" button, and the order confirmation page loads.
MMDuet2	✗ 29s: The search results page displays various styles of women's jeans including straight-leg, loose, and knitted, with options for sorting by …	✓ 51s: … options to select colors (including blue and black) and sizes (M, L, XL, 2XL) with weight ranges for each size; there is also a section …
ViSpeak	✗ 28s: The video shows a person using a shopping app to search for jeans. The person types in the search bar and selects the type of jeans they want …	✗ 49s: The person is browsing the shopping app and selecting a pair of jeans.
StreamVibe (ours)	✓ 28s: Now that the search results are loaded, scroll through the list to find a pair of jeans you like.	✓ 50s: Select your preferred size and color from the options, then tap the red 'Add to Cart' button at the bottom.

Figure 34 | **Additional case study 04.** Finding and selecting jeans in a shopping app. Model responses at two annotated evidence points, matched within ±2 s.

Case 05. Making capsule coffee

Instruction: Walk me step by step through making coffee with this capsule espresso machine, and jump in if I make any mistakes.

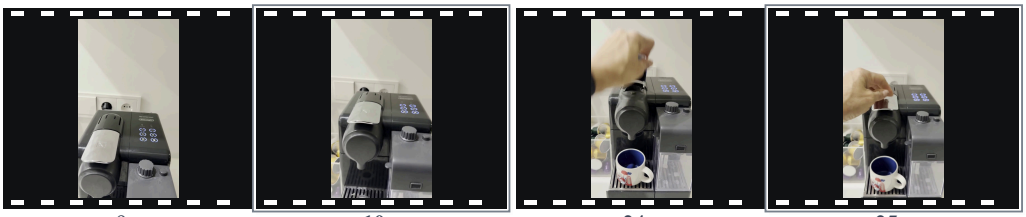

JoyAI	✗ 10s: Now pull out the capsule holder, and you can see the capsule slot inside.	✗ 25s: You can see the capsule is being pushed into the machine.
MMDuet2	✗ No response within ±2 s	✓ 23s: Press the button on the machine, which dispenses coffee into the mug, and after dispensing, lift the coffee capsule out of the port.
ViSpeak	✗ 8s: The first step is to press the button on the machine.	✗ 26s: The person puts a cup under the coffee machine, presses the button, and then puts a coffee capsule into the machine.
StreamVibe (ours)	✓ 9s: The machine is now powered on. Next, place your cup under the spout and insert a coffee capsule.	✓ 25s: The capsule is locked in. Press the coffee button on the control panel to start brewing.

Figure 35 | **Additional case study 05.** Making capsule coffee. Model responses at two annotated evidence points, matched within ±2 s.

Case 06. Assembling a metal shelf

Instruction: Walk me through assembling this flat-pack metal shelf unit, and check that I am aligning the frame legs and securing every screw joint properly.

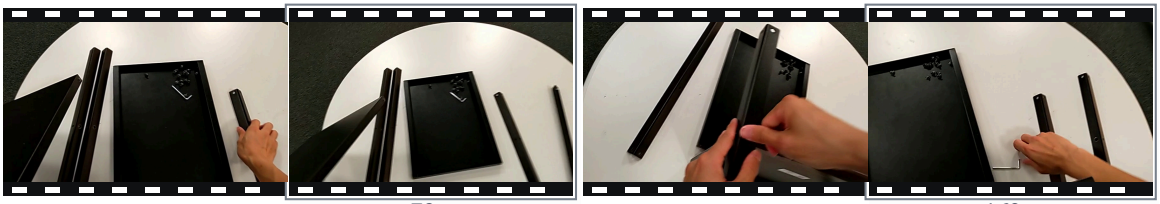

JoyAI	✓ 73s: Great, now attach the top to the frame. Make sure it's aligned properly.	✗ 162s: Great, now attach the first shelf to the frame. Make sure it's aligned properly.
MMDuet2	✗ No response within ±2 s	✗ No response within ±2 s
ViSpeak	✓ 73s: The person is assembling a flat-pack metal shelf unit by placing the frame legs on the shelf and securing them with screws.	✗ No response within ±2 s
StreamVibe (ours)	✓ 73s: Align the side leg with the mounting holes on the base frame and insert the screws to secure it.	✓ 162s: Tighten the screw on the side leg using the Allen key to secure the joint.

Figure 36 | **Additional case study 06.** Assembling a metal shelf. Model responses at two annotated evidence points, matched within ±2 s.

Case 07. Browning butter without scorching

Instruction: Walk me through making brown butter in this skillet, and warn me if the heat is wrong or if it starts to burn.

JoyAI	✗ 40s: The butter is turning a deeper golden color, and the aroma is getting stronger.	✗ 80s: The warning text remains the same, with a red overlay.
MMDuet2	✗ 39s: Chunks of butter are melting in a stainless-steel bowl on a stovetop burner, and a person in a dark chef's coat stands nearby observing the …	✗ 79s: Chunks of butter are melting in a stainless-steel bowl on a stovetop burner, and the melted butter has a frothy surface with numerous small …
ViSpeak	✓ 39s: … The person is adding butter to the pan and stirring it with a spatula.	✓ 81s: The heat is too high, and the butter is burning.
StreamVibe (ours)	✓ 38s: Keep the heat at medium and stir the butter continuously with a spatula to prevent the milk solids from scorching.	✓ 78s: Immediately remove the pan from the heat to stop cooking and prevent the butter from burning.

Figure 37 | **Additional case study 07.** Browning butter without scorching. Model responses at two annotated evidence points, matched within ±2 s.

Case 08. Making successive creases for a paper butterfly

Instruction: Walk me through folding this paper butterfly step by step, and tell me exactly how to align and crease the paper for each fold.

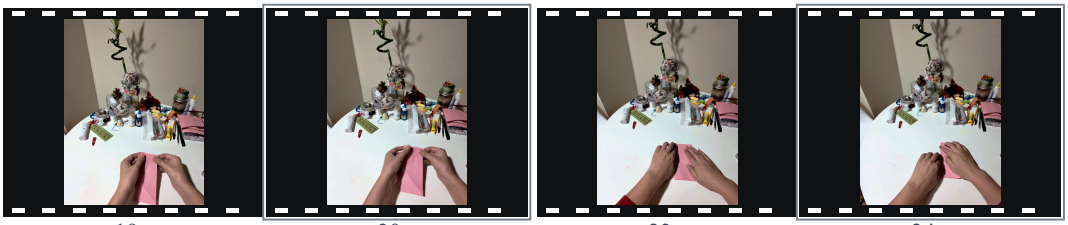

JoyAI	✗ 20s: The hand folds the left corner down to the center crease, creating a new triangular shape.	✗ 24s: Then the bottom corner is folded up to meet the top, making the paper narrower.
MMDuet2	✗ No response within ±2 s	✗ No response within ±2 s
ViSpeak	✗ No response within ±2 s	✗ 26s: Fold the paper in half lengthwise, then unfold it. Fold the top corners of the paper to the center crease, then unfold it. Fold the bottom corners …
StreamVibe (ours)	✓ 19s: Now that you've opened the paper back up, fold the top edge down to meet the bottom edge to create a horizontal crease.	✓ 23s: Now that you've opened the paper back up, fold the top edge down to meet the bottom edge to create a horizontal crease.

Figure 38 | **Additional case study 08.** Making successive creases for a paper butterfly. Model responses at two annotated evidence points, matched within ±2 s.