arXiv:2610.39147v1 [cs.CV] 29 Oct 2026

Kinematic Causal Attribution: Black-Box Explanations of Motion versus Geometry for 4D Point-Cloud Video Models

Zerui Wang

Abstract

4D perception models, which classify 3D point clouds that evolve over time, are being built into robots, assistive monitoring and AR, yet before deploying one a practitioner cannot check *which cue* it relies on: how a body part moves, or what shape it has. Existing explanations do not fit this setting. 3D attribution is single-frame, video attribution is confined to RGB pixels, gradient maps need white-box access, attention is absent from state-space backbones, and point-level perturbations assume a point correspondence across frames that real depth and LiDAR data lack. We present Kinematic Causal Attribution (KCA), a black-box, interventional attribution method that queries only the model's outputs. Its FREEZE operator removes a region's motion while keeping its shape, and a shape-free proxy removes the shape while keeping the motion. Together they span a motion$\times$geometry cube for every tracked region, whose order-symmetric Shapley decomposition is complete and whose interaction is measured rather than assumed away. The operators rewrite per-frame point *sets*, so they need no correspondences, and a context-averaged estimator keeps the scores reliable when a structure spans several regions. We validate KCA on the official PSTNet, P4Transformer and Mamba4D trained on MSR-Action3D. The decisive test is an audit: two binary tasks are built from real clips whose classes differ only in the motion or only in the shape of the same arm, and the motion share of the arm's importance follows the defining cue on all three backbones (0.83 vs. 0.60 on P4Transformer, 0.86 vs. 0.30 on PSTNet, 0.92 vs. 0.61 on Mamba4D), a distinction no single attribution map can make. The split costs no faithfulness: the necessity maps have the lowest deletion AUC on P4Transformer and PSTNet, tie with the best baseline on Mamba4D, and lose their correlation with the trained model's once the weights are randomised. Skeleton ground truth, re-sampling stability and a cost study complete the validation.

1 Introduction

Point-cloud video is the stream of 3D point sets produced by depth and LiDAR sensors, and it underpins action recognition, robotic manipulation, assistive monitoring and augmented reality. Dedicated *4D* backbones now model it directly: point spatio-temporal convolutions [Fan et al., 2021b], spatio-temporal transformers [Fan et al., 2021a] and attention-free state-space models such as Mamba4D [Liu et al., 2025]. They are accurate but opaque. Before such a model is deployed, a practitioner should be able to check *what* it bases a decision on. Did it label a clip "high arm wave" because of *how* the arm moved, or because of *what shape* the body had? A model that relies on shape where the task is defined by motion has learned a shortcut, as RGB action models do with static appearance [Li et al., 2023], and that shortcut can fail under new subjects, clothing or sensor placement.

What a deployable 4D explanation needs. We argue that an explanation method for deployed 4D models has to meet four requirements, and that no existing method meets them all. *(R1) Black-box access.* Deployed models are often compiled, quantised or served behind an interface, and the backbones differ in kind. Gradients need white-box access. Attention is a contested proxy for importance [Jain and Wallace, 2019,

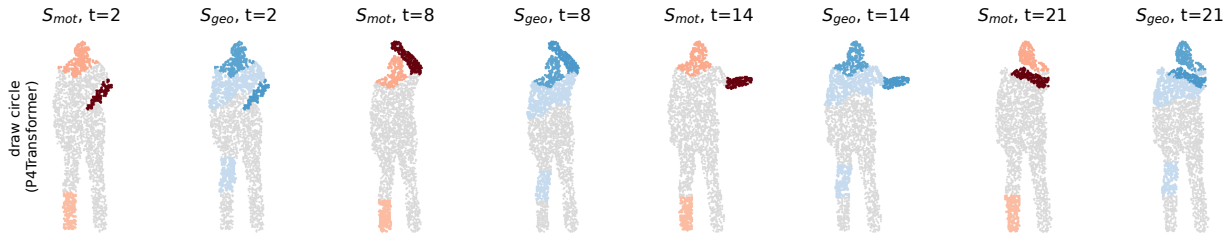

Figure 1: **Kinematic Causal Attribution (KCA)** attributes a 4D model's decision to *how* each spacetime region moves (S_{mot}, red) and to *what shape* it has (S_{geo}, blue). It intervenes on the input and reads only the model's output. The clip is a real MSR-Action3D test clip (*draw circle*) at four instants, explained for the official P4Transformer. The motion term is strongest on the drawing hand (weaker on the head and one foot), while geometry importance covers the upper body. Fig. 6 in the appendix shows two more actions on all three backbones, including one where the motion term falls on the torso. No attention, gradient or point correspondence is used, and the same code explains convolutional and attention-free state-space backbones (Sec. 4.3).

Abnar and Zuidema, 2020], and Mamba4D has no attention to read. An explanation should therefore need only forward queries. *(R2) Real sensor data.* Sensors re-sample the surface in every frame, so "the same point at two times" does not exist. Any operator that moves or freezes indexed points silently breaks on real data. *(R3) The right question.* A prediction can depend on the *shape* of a structure or on its *motion*. In pixel space the two are entangled in one saliency value. In 4D they are physically distinct, yet 3D attribution is single-frame [Zheng et al., 2019, Schinagl et al., 2022], video attribution works on RGB pixels [Price and Damen, 2020, Lee et al., 2025], and no method attributes a decision to motion versus geometry. *(R4) Verifiable and affordable.* An explanation that will be acted on must be shown to be faithful to the model, must depend on the learned weights [Adebayo et al., 2018], should be checked against ground truth where it exists, and must run at a practical cost.

Our approach. Kinematic Causal Attribution (KCA) treats explanation as intervention on tracked spacetime regions. Its key operator, FREEZE, holds a region static in its anchor-frame configuration: the region loses its motion and keeps its shape. A shape-free proxy that follows the region's centroid does the converse. The four resulting inputs form a motion$\times$geometry cube for each region. Averaging the cube's two paths yields an order-symmetric geometry term S_{geo} and motion term S_{mot} that sum exactly to the region's necessity $\mathcal{N}$, and the cube's mixed difference measures their interaction. All operators rewrite only a region's per-frame point *sets*, so they need no correspondences. A context-averaged estimator, a Banzhaf-type average over random interventions on the other regions, keeps the scores reliable when a structure spans several regions. KCA only queries the model, so it explains convolutional, attention and state-space backbones with the same code (Fig. 1).

Contributions.

- **KCA, a black-box explanation that separates motion from geometry** (R1, R3). To our knowledge it is the first method that attributes a 4D model's decision to the *motion* versus the *geometry* of each spacetime region. Its decomposition is complete and order-symmetric (Shapley), and it reports the interaction instead of assuming it away (Sec. 3).
- **A correspondence-free instantiation with guarantees** (R2). Set-based operators, a local and a context-averaged estimator, and the conditions under which each score is a causal effect.

- **A validation protocol for deployment** (R4). On three unmodified official backbones we test faithfulness against 13 baselines, localisation against skeleton ground truth, dependence on the learned weights, stability under sensor re-sampling, and cost. We also give a *crossover audit*: two tasks built from real clips whose classes differ only in the motion or only in the shape of one arm, which checks whether the explanation reports the cue a model was trained to use. The crossover is our primary evidence that the split is correct; the other tests show that it costs no faithfulness (Sec. 4).

2 Related Work

Point-cloud video models. Static point-cloud networks [Qi et al., 2017a,b, Wang et al., 2019, Wu et al., 2024] were extended to point-cloud video by MeteorNet [Liu et al., 2019], point spatio-temporal convolutions (PSTNet [Fan et al., 2021b]), point 4D convolution with self-attention (P4Transformer [Fan et al., 2021a]; PST-Transformer [Fan et al., 2023]), primitive-level modelling [Wen et al., 2022], temporal point patches [Ben-Shabat et al., 2024] and the attention-free state-space Mamba4D [Liu et al., 2025]. None of them explains its decisions. The "motion saliency" of recent pre-training [Wang et al., 2026] selects tokens to mask, and Mamba4D's "disentangled" and "causal" refer to efficiency and scan order, not to attribution.

Attribution for images, video and point clouds. Surrogates [Ribeiro et al., 2016, Lundberg and Lee, 2017] gave way to occlusion and mask optimisation [Petsiuk et al., 2018, Fong et al., 2019], gradients [Selvaraju et al., 2017, Sundararajan et al., 2017] and transformer-specific methods [Abnar and Zuidema, 2020, Chefer et al., 2021]; attention as explanation is contested [Jain and Wallace, 2019] and unavailable for state-space models. Video methods attribute frames [Price and Damen, 2020], read spatio-temporal attention [Wang and Liu, 2025], learn temporal concept bottlenecks [Knab et al., 2025, Lee et al., 2025] or test concept sensitivity [Saha et al., 2024]. The closest in spirit, DANCE [Lee et al., 2025], separates motion, object and scene *concepts*, but it is ante hoc, works in a learned concept space and is correlational; KCA intervenes post hoc on physical 3D regions of any trained model. All of these work on 2D RGB. Point-cloud saliency [Zheng et al., 2019], gradient maps [Tayyub et al., 2022], occlusion maps for LiDAR detectors [Schinagl et al., 2022] and interpretable classifiers [De Vries et al., 2025] explain a *single* cloud with no time axis. Interpretability for dynamic scene *representations*—language-embedded 4D Gaussians [Fiebelman et al., 2025, Li et al., 2025], static/dynamic separation for editing [Kwon et al., 2025], importance for pruning [Zhang et al., 2024], prototypes for static splats [Galus et al., 2026]—serves grounding, editing or compression, not the attribution of a recognition decision.

Causal attribution and trustworthy evaluation. LENS [Watson et al., 2021] operationalises Pearl's necessity and sufficiency for tabular and text data, CaCE [Goyal et al., 2019] contrasts interventional effects with concept sensitivity, and LD-ViCE [Varshney et al., 2025] generates RGB video counterfactuals; "motion attribution" has been used for *training-data* attribution in video generation (Motive [Wu et al., 2026]). None of them reaches 3D geometry and motion, and none separates the two. Uncertainty methods for 4D occupancy [Wang et al., 2025, Kälble et al., 2025, Heidrich et al., 2025] ask how reliable a prediction is, not what caused it. Whether an explanation can be trusted is itself an empirical question. Deletion-based faithfulness [Petsiuk et al., 2018, DeYoung et al., 2020] tests whether highly ranked inputs matter to the model, and model randomisation [Adebayo et al., 2018] tests whether the explanation depends on what the model learned. Point-cloud XAI still lacks an agreed evaluation standard [Chakraborty et al., 2026]. We therefore validate KCA with all three kinds of evidence: faithfulness, sanity checks and ground truth.

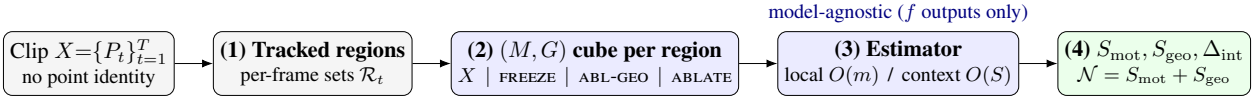

Figure 2: **Method overview.** Regions are tracked through time as per-frame point *sets*, so no cross-frame point correspondence is needed. Each region carries two factors—its geometry G and its motion M—and four interventions realize the corners of the (M, G) cube; every edge of the cube changes exactly one factor. A local or a context-averaged estimator turns the cube's outputs into order-symmetric motion and geometry importance whose interaction is measured. Only outputs of f are used, so the same pipeline explains convolutional, attention and state-space backbones.

3 Method

KCA explains a trained 4D model by *intervening* on spacetime regions of its input and measuring the causal effect on the prediction. After fixing the causal setting (Sec. 3.1) and notation (Sec. 3.2), we define correspondence-free interventions (Sec. 3.3) and the causal scores they induce (Sec. 3.4), derive the motion-vs-geometry decomposition that is our central contribution (Sec. 3.5), establish its properties (Sec. 3.6), and give two estimators (Sec. 3.7). Fig. 2 gives an overview.

3.1 Preliminaries: causal attribution

We adopt the interventional view of explanation [Watson et al., 2021, Janzing et al., 2020, Pearl, 2009]: the importance of an input component is the change it *causes* in the output under an intervention $\mathrm{do}(\cdot)$, rather than a gradient or attention read-out. For a target φ and a manipulable factor Z, the *average causal effect* (ACE) of setting $Z{=}z$ against a reference z_0 is $\mathrm{ACE}_Z = \mathbb{E}[\varphi \mid \mathrm{do}(Z{=}z)] - \mathbb{E}[\varphi \mid \mathrm{do}(Z{=}z_0)]$. Pearl's probability of *necessity* (PN) asks whether removing a cause flips the outcome, and of *sufficiency* (PS) whether the cause alone produces it [Pearl, 2009]; Watson et al. [2021] operationalize local PN/PS for feature subsets. We (i) instantiate these for *spacetime* regions of a 4D model, without assuming point correspondences, and (ii) factor each region's effect into *motion* and *geometry*—a factorization unique to dynamic-3D data. Appendix A summarises the notation and gives pseudocode (Alg. 1).

3.2 Problem formulation

A point-cloud video is a sequence $X = \{P_t\}_{t=1}^{T}$, $P_t = \{\mathbf{p}_{t,i}\}_{i=1}^{N_t} \subset \mathbb{R}^3$. Depth and LiDAR sensors re-sample the surface in every frame, so there is *no point identity* across frames: index i in frame t and index i in frame t' are unrelated points. A trained model f maps X to logits; we explain $\varphi(X) = f(X)_c$ for the predicted class c (any scalar read-out is admissible, including non-differentiable ones). The spacetime support is partitioned into m *tracked regions*: region $\mathcal{R}$ is a sequence of per-frame point sets $\{\mathcal{R}_t \subseteq P_t\}_t$ covering the same physical structure (tracked k-means supervoxels; Appendix C). KCA assigns each region $S_{\text{geo}}(\mathcal{R})$ and $S_{\text{mot}}(\mathcal{R})$, aggregated into per-point fields.

3.3 Correspondence-free spacetime interventions

A region has two factors: its *geometry* G (the shape it has in each frame) and its *motion* M (how it moves through time). We realize every combination with operators that act only on the per-frame sets $\mathcal{R}_t$—never on indexed points:

- $\mathrm{do}(\textsc{freeze}, \mathcal{R})$ $(M{=}0, G{=}1)$: every $\mathcal{R}_t$ is replaced by the region's anchor-frame set $\mathcal{R}_{t_0}$ (re-sampled to $|\mathcal{R}_t|$ points): the region stays static in its t_0 configuration, with its real shape. This removes *all* of its motion—translation, rotation and articulation alike—without any correspondence.

- $\mathrm{do}(\textsc{abl-geo}, \mathcal{R})$ (M=1, G=0): $\mathcal{R}_t$ is replaced by a *shape-free proxy*—a point mass—that follows the region's centroid trajectory $c_t(\mathcal{R})$: the region keeps its motion but loses its shape.
- $\mathrm{do}(\textsc{ablate}, \mathcal{R})$ (M=0, G=0): the same proxy, held static at $c_{t_0}(\mathcal{R})$.

The four corners X, $\mathrm{do}(\textsc{freeze}, \mathcal{R})X$, $\mathrm{do}(\textsc{abl-geo}, \mathcal{R})X$, $\mathrm{do}(\textsc{ablate}, \mathcal{R})X$ form the region's (M, G) cube, and *every edge changes exactly one factor*. This is the property that makes the decomposition below meaningful: had ABLATE *deleted* the region while ABL-GEO kept a proxy, the edge ABL-GEO $\rightarrow$ ABLATE would change presence as well as motion, and a static region's presence would be booked as "motion" (we measure this failure mode in Sec. 4.5). Keeping the proxy at the region's own location also keeps every corner on the scene's support. A third operator, $\mathrm{do}(\textsc{displace}, \mathcal{R})$, replays the region's sets in reversed or permuted order ($\mathcal{R}_t \leftarrow \mathcal{R}_{\pi(t)}$): a sign-sensitive probe of motion *direction*. When point identity does exist (synthetic data), all operators reduce exactly to their index-aligned forms, *e.g.* $\mathbf{p}_{t,i} \leftarrow \mathbf{p}_{t_0,i}$ for FREEZE.

3.4 Causal attribution: necessity and sufficiency

Necessity measures the drop when a region's spatio-temporal content—shape and motion—is replaced by the static proxy,

$$\mathcal{N}(\mathcal{R}) \;=\; \varphi(X) \;-\; \varphi\big(\mathrm{do}(\textsc{ablate}, \mathcal{R})\, X\big), \tag{1}$$

and *sufficiency* how much the region alone preserves, $\mathcal{S}(\mathcal{R}) = \varphi(\mathrm{do}(\textsc{ablate}, \Omega\backslash\mathcal{R})X) - \varphi(\mathrm{do}(\textsc{ablate}, \Omega)X)$. High $\mathcal{N}$ marks regions the decision *depends on*. Unlike attention or gradients, both are defined for *any* f, including attention-free state-space models.

3.5 Disentangling motion from geometry

Necessity conflates a region's motion and its geometry; the cube separates them. Writing $\varphi \equiv \varphi(X)$ and $\varphi_{\mathrm{op}} \equiv \varphi(\mathrm{do}(\mathrm{op}, \mathcal{R})X)$, the two paths from X to $\mathrm{do}(\textsc{ablate}, \mathcal{R})X$ remove motion first or geometry first; averaging them yields the *order-symmetric* scores

$$\begin{aligned} S_{\mathrm{mot}}(\mathcal{R}) &= \tfrac{1}{2}\big[(\varphi - \varphi_{\textsc{freeze}}) + (\varphi_{\textsc{abl-geo}} - \varphi_{\textsc{ablate}})\big], \\ S_{\mathrm{geo}}(\mathcal{R}) &= \tfrac{1}{2}\big[(\varphi - \varphi_{\textsc{abl-geo}}) + (\varphi_{\textsc{freeze}} - \varphi_{\textsc{ablate}})\big], \end{aligned} \tag{2}$$

each an average of two single-factor contrasts. S_{mot} asks *does the decision depend on how this structure moves?*, S_{geo} *does it depend on what shape it has?*—a question no prior attribution can pose. Their gap to the single-order scores is the interaction

$$\Delta_{\mathrm{int}}(\mathcal{R}) = \varphi - \varphi_{\textsc{freeze}} - \varphi_{\textsc{abl-geo}} + \varphi_{\textsc{ablate}}, \tag{3}$$

which we *measure* rather than assume away.

3.6 Theoretical properties

KCA queries f only through outputs of intervened inputs, so its guarantees hold for every architecture. The headline results require only interventional validity **(I)**; separability **(S)** is *measured*, not assumed (Tab. 5; proofs in Appendix B).

Assumption 1 (Interventional validity **(I)** & separability **(S))**. A region's effect on φ is mediated by a geometry factor $G_{\mathcal{R}}$ and a motion factor $M_{\mathcal{R}}$. **(I)** FREEZE, ABL-GEO and ABLATE set the targeted factor(s) without altering off-region content (they only rewrite $\mathcal{R}_t$). **(S)** *(diagnostic)* φ is additively separable in $(M_{\mathcal{R}}, G_{\mathcal{R}})$, *i.e.* the discrete mixed difference vanishes.

Proposition 1 (Completeness and the order-2 Shapley decomposition)**.** *For any f, $S_{\text{mot}}(\mathcal{R}) + S_{\text{geo}}(\mathcal{R}) = \mathcal{N}(\mathcal{R})$. $(S_{\text{mot}}, S_{\text{geo}})$ is the Shapley value of the two-player game $\{M_{\mathcal{R}}, G_{\mathcal{R}}\}$ with worth $v(\{\cdot\}) = \varphi$ at the corresponding corner; Δ_{int} is its Shapley interaction index [Grabisch and Roubens, 1999]; each single-order score deviates from the symmetric one by exactly $\pm\Delta_{\text{int}}/2$; and $\Delta_{\text{int}}{=}0$ iff **(S)** holds for $\mathcal{R}$.*

Proposition 2 (Causal identification of motion)**.** *Under **(I)**, $\varphi - \varphi_{\textsc{freeze}}$ is the interventional effect of region $\mathcal{R}$'s motion at its real geometry and $\varphi_{\textsc{abl-geo}} - \varphi_{\textsc{ablate}}$ the same effect for a shape-free body (CaCE-type contrasts [Goyal et al., 2019]); S_{mot} is their average. Under **(S)** the two coincide and equal the marginal motion ACE; symmetrically for S_{geo}.*

Proposition 3 (Necessity correspondence)**.** *With the decision event $E{=}\mathbf{1}[\varphi{\geq}\kappa]$, $\mathcal{N}(\mathcal{R})$ is the margin relaxation of the empirical probability of necessity $\text{PN}(\mathcal{R}) = \Pr[E(\text{do}(\textsc{ablate}, \mathcal{R})X){=}0 \mid E(X){=}1]$ [Pearl, 2009, Watson et al., 2021]; under monotonicity and exogeneity **(M)** it is identified by the excess-risk ratio. $\mathcal{S}$ is the dual* PS *score.*

3.7 Estimation

Local estimator. Evaluating the cube of every region at the observed input costs $3m{+}1$ forward passes, which we batch. The local scores answer "what happens if this region alone is intervened on"—the classic occlusion question.

Context-averaged estimator. Local contrasts under-credit *redundant* structure: if a decision-relevant part spans several regions, removing any one of them leaves enough of the others, and every local necessity is small. We therefore also average the same contrasts over random *contexts* in which the other regions are themselves intervened. We draw S joint states with every factor $M_{\mathcal{R}}, G_{\mathcal{R}} \sim \text{Bernoulli}(q)$ independently, evaluate φ once per state (all regions intervened jointly), and fit

$$\varphi \approx \beta_0 + \textstyle\sum_{\mathcal{R}} \left(\mu_{\mathcal{R}} M_{\mathcal{R}} + \gamma_{\mathcal{R}} G_{\mathcal{R}} + \iota_{\mathcal{R}} M_{\mathcal{R}} G_{\mathcal{R}}\right) \tag{4}$$

by least squares, then set $S_{\text{mot}}{=}\mu + \iota/2$, $S_{\text{geo}}{=}\gamma + \iota/2$, $\Delta_{\text{int}}{=}\iota$, $\mathcal{N}{=}\mu{+}\gamma{+}\iota$.

Proposition 4 (Context-averaged decomposition)**.** *Eq. (4) satisfies completeness $S_{\text{mot}} + S_{\text{geo}} = \mathcal{N}$ exactly. If φ has no cross-region interactions, its population solution equals Eq. (2) (for any context); otherwise it is the context-averaged (Banzhaf-type [Banzhaf III, 1965]) version of Eq. (2) under the sampling distribution.*

The cost is S passes, independent of m. Both estimators use the same interventions; the local one is exact for a single context and cheaper, the context one is robust to redundancy.

3.8 Model-agnostic instantiation

We explain three architecturally distinct official backbones with the *same* code: the convolutional PSTNet [Fan et al., 2021b], the attention-based P4Transformer [Fan et al., 2021a], and the attention-free state-space Mamba4D [Liu et al., 2025]. The interventions and scores (Eqs. (1)–(4)) touch the model only through its outputs; the only backbone-specific code is the input contract (a fixed number of points per frame, so an intervened frame is re-sampled from its remaining points, as the official data loaders do). Attention is never used: KCA is equally defined on models that have none.

4 Experiments

We ask four questions. **Q1** Is KCA's necessity map faithful? **Q2** Does one code path explain convolutional, attention and state-space models? **Q3** Is the motion/geometry split *correct*? **Q4** What do its components contribute and cost?

4.1 Setup

Data and ground truth. We explain action recognition on MSR-Action3D [Li et al., 2010]: 567 depth videos of 20 actions by 10 subjects, with the standard split (subjects 1–5 for training, 270 videos; 6–10 for testing, 297 videos). Depth maps become point-cloud videos with the MeteorNet protocol [Liu et al., 2019] used by all three official backbones; we use centred 24-frame clips with 2048 points per frame, re-sampled independently in every frame (no point identity). MSR-Action3D also ships 20-joint skeletons. We lift them into the same coordinates, where a median 95% of the points lie within about 12 cm of a bone, and label every point with its nearest bone. Some videos store no skeleton for a run of frames; every ground-truth-based number leaves out the 12 benchmark clips concerned. Points of bones that travel at least half as far as the clip's most active bone form the *moving* ground truth (on average 13% of a clip's points); points of bones that travel at most a fifth as far form the *static* ground truth (65%; Appendix C). The main benchmark explains every test video that is at least 24 frames long (287 clips; 275 with a complete skeleton). The cross-backbone study, the ablations and the stability study use the first two test clips of every action (40 clips).

Backbones. We explain the *unmodified official code* of the convolutional PSTNet [Fan et al., 2021b], the attention-based P4Transformer [Fan et al., 2021a] and the attention-free state-space Mamba4D [Liu et al., 2025]. No pretrained weights exist, so we train each from scratch at its official width on one A100 GPU, replacing the CUDA-only operators by numerically equivalent PyTorch versions (Appendix C). Test accuracies (video level) are 89.2/88.9/88.2 % for P4Transformer/PSTNet/Mamba4D (published: 90.9/91.2/92.7). A second, laptop-scale training (reduced width, 1024 points) supports the same conclusions; Appendix F compares the two.

Baselines. *Floors:* random; geometry-edge (local point spacing); motion prior (model-blind nearest-neighbour flow between frames). *Gradient:* Grad×Input, SmoothGrad [Smilkov et al., 2017], Integrated Gradients [Sundararajan et al., 2017], and 3D Grad-CAM [Selvaraju et al., 2017] on the tube-embedding tokens. *Perturbation:* region occlusion (deleting each tracked region, as in point-cloud saliency [Zheng et al., 2019]), OccAM [Schinagl et al., 2022], and extremal perturbations [Fong et al., 2019]. *Attention (P4Transformer only):* rollout [Abnar and Zuidema, 2020], transformer attribution [Chefer et al., 2021], and raw attention: the attention each token receives in the last layer, used *as* the explanation.

KCA. $m{=}12$ tracked regions; local estimator (37 passes) and context estimator ($S{=}256$, $q{=}\frac{1}{2}$). An attribution map gives each point its region's $\mathcal{N}$; localisation uses S_{mot} (moving GT) and S_{geo} (static GT).

Metrics. *Faithfulness* (no GT): deletion/insertion AUC of the target-class probability [Petsiuk et al., 2018], comprehensiveness/sufficiency (AOPC over $\{1, 5, 10, 20, 50\}\%$) [DeYoung et al., 2020], and Sensitivity-n [Ancona et al., 2018]; a deleted point is removed and its frame re-sampled. *Localisation* (skeleton GT): Pointing Game [Zhang et al., 2018] and energy-based EBPG [Wang et al., 2020]. *Sanity:* cascading model randomisation [Adebayo et al., 2018]. *Statistics.* Every score is computed per clip and reported as the mean over clips. A p-value compares two methods on the same clips with a two-sided paired Wilcoxon signed-rank test [Wilcoxon, 1945]; p-values in the text are uncorrected, and a * in a table marks KCA better than the strongest baseline after Bonferroni correction over that table's metrics [Dunn, 1961]. Intervals are 95% percentile bootstrap intervals over clips (2000 resamples) [Efron, 1979]. Appendix E lists the data, models and code behind every table and figure.

Method	Del.↓	Comp.↑	Ins.↑†	Suff.↓†
Random (floor)	0.916±0.14	0.003±0.05	0.920±0.14	0.041±0.12
Geometry-edge (floor)	0.890±0.19	0.011±0.11	0.903±0.16	0.103±0.21
Motion-prior (floor)	0.505±0.38	0.277±0.33	0.860±0.16	0.365±0.22
Grad×Input	0.655±0.28	0.086±0.18	0.722±0.25	0.332±0.28
SmoothGrad	0.772±0.24	0.031±0.12	<u>0.915±0.14</u>	<u>0.083±0.16</u>
Integrated Grad.	0.282±0.26	0.347±0.24	0.862±0.17	0.129±0.19
3D Grad-CAM	0.713±0.26	0.051±0.12	0.554±0.30	0.659±0.30
Region occlusion	<u>0.233±0.23</u>	<u>0.416±0.23</u>	0.784±0.17	0.531±0.24
OccAM	0.853±0.24	0.059±0.13	**0.953±0.07**	**−0.004±0.11**
Extremal Pert.	0.634±0.22	0.052±0.10	0.309±0.26	0.830±0.25
Attn. Rollout	0.649±0.27	0.073±0.15	0.769±0.24	0.431±0.26
Transformer Attr.	0.601±0.28	0.100±0.17	0.766±0.24	0.424±0.28
Raw attention (last layer)	0.542±0.28	0.127±0.16	0.688±0.25	0.551±0.29
KCA (local)	0.256±0.25	0.403±0.24	0.793±0.17	0.511±0.23
KCA (context)	**0.179±0.19***	**0.453±0.21***	0.761±0.18	0.546±0.24

Table 1: **Faithfulness** on P4Transformer, MSR-Action3D test set (all 287 test clips; mean±std over clips). Del.: deletion AUC of the target-class probability; Comp.: comprehensiveness. Both remove the top-ranked points. Ins. and Suff. (†) *keep* the top-ranked points, which rewards spatially scattered rankings on point clouds (a random subset preserves the whole body; see the random floor). **Bold**: best, <u>underlined</u>: second, floors excluded. *: KCA significantly better than the strongest baseline after Bonferroni correction.

4.2 Faithfulness (Q1)

Tab. 1 and Fig. 3 compare the necessity maps. KCA's context estimator attains the lowest deletion AUC (0.179; local estimator 0.256) and the highest comprehensiveness. Its deletion AUC is significantly below that of region occlusion (0.233; paired Wilcoxon $p<10^{-3}$) and of Integrated Gradients (0.282; $p<10^{-3}$), and far ahead of attention rollout (0.649), transformer attribution (0.601) and raw attention (0.542): on the very backbone that has attention, reading it out is a poor explanation. Two findings deserve comment. *(i) Region occlusion is KCA's necessity with a deletion corner.* Replacing the static proxy of Sec. 3.3 by deletion makes the local necessity map *identical* to region occlusion, which the deletion metric (itself deletion-based) favours over the local estimator (p=0.011); the context estimator more than recovers the gap, and the proxy is what makes the motion/geometry split meaningful (Sec. 4.4, Tab. 11). *(ii) Point-cloud models shrug off random point dropout.* Deleting points in random order leaves the prediction almost intact (Fig. 3), so point-level random masking (OccAM, 0.853) and soft-mask optimisation (Extremal, 0.634) carry little signal at this budget, and *keep*-based metrics reward scattered rankings: the random floor scores an insertion AUC of 0.920. We therefore read faithfulness mainly from the deletion-based columns.

4.3 Across backbones (Q2)

Tab. 2 runs the same code on the three official architectures. KCA is defined on all of them—the attention methods are not. On every backbone its context estimator has the lowest or second-lowest deletion AUC, statistically tied with region occlusion on PSTNet and Mamba4D (paired Wilcoxon p=0.38, p=0.83) and below Integrated Gradients, significantly on PSTNet (p=0.007) but not on Mamba4D (p=0.064); its motion term localises the moving parts on par with Integrated Gradients, the best single map (p=0.63, p=0.73), and at least as well as region occlusion. The method ranking is consistent across backbones (Kendall τ, caption), and the state-space Mamba4D, which has no attention to read, is explained exactly like the others. On PSTNet and Mamba4D every informative map puts less mass on the moving parts than on P4Transformer: these

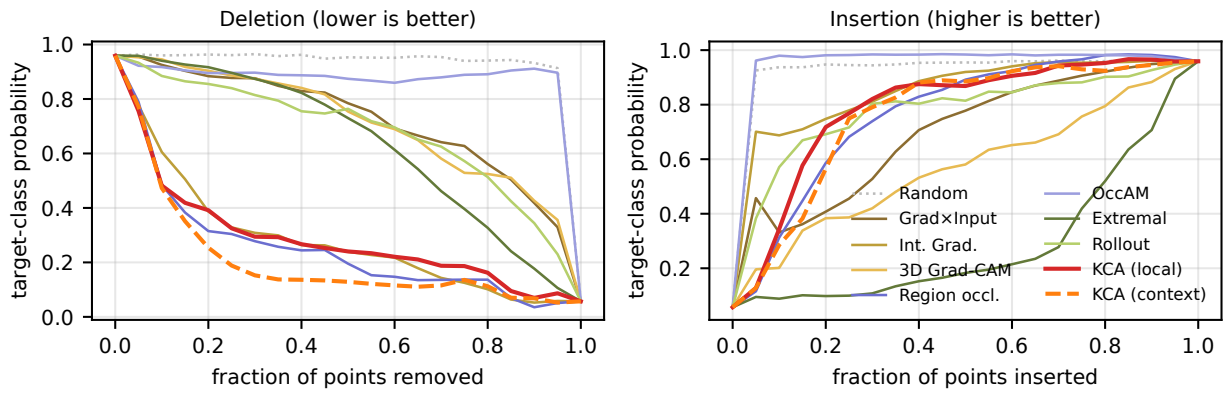

Figure 3: **Deletion and insertion curves** on P4Transformer (mean over the 287 clips of Tab. 1): target-class probability vs. fraction of points removed (left) or inserted (right), in 20 steps.

	PSTNet		P4Transformer		Mamba4D	
Method	Del.↓	EBPG↑	Del.↓	EBPG↑	Del.↓	EBPG↑
Random (floor)	0.950	0.098	0.929	0.098	0.902	0.098
Integrated Grad.	0.303	**0.390**	0.230	**0.567**	0.329	**0.456**
Region occlusion	0.252	0.329	0.230	0.394	**0.246**	0.351
OccAM	0.845	0.098	0.865	0.098	0.778	0.098
Extremal Pert.	0.932	0.098	0.621	0.076	0.577	0.085
Attn. Rollout	*n/a*	*n/a*	0.667	0.141	*n/a*	*n/a*
Raw attention (last layer)	*n/a*	*n/a*	0.591	0.201	*n/a*	*n/a*
KCA (local)	0.262	0.326	0.272	0.504	0.314	0.392
KCA (context)	**0.217**	0.361	**0.219**	0.506	0.269	0.433

Table 2: **Across backbones** (the same 40 test clips, two per action, for every backbone; means over clips): deletion AUC and localisation of the moving ground truth (EBPG, on the clips with a complete skeleton; S_{mot} for KCA). Attention methods are undefined on PSTNet and Mamba4D (n/a). Kendall's τ between the rankings by Del. of the 7 methods defined on all three backbones: PSTNet–P4Transf. 0.71, PSTNet–Mamba4D 0.81, P4Transf.–Mamba4D 0.71.

models spread their evidence more widely (Sec. 4.6).

4.4 Motion versus geometry (Q3)

Localisation. S_{mot} places 0.50 of its positive mass on the skeleton's moving parts (EBPG, Tab. 3), above region occlusion (0.40) and far above the model-blind motion prior (0.27), whose mass is spread over sensor jitter, but below the best single map, Integrated Gradients (0.55; $p<10^{-3}$). The single most-salient point of Integrated Gradients hits the moving parts more often (PG), since region-constant maps have no ranking *within* a region. On moving parts, 0.79 of KCA's positive importance is motion, against 0.69 over the whole body. Localisation, however, is confounded: moving parts are also geometrically distinctive, so a map that highlights the arm for *any* reason scores well.

Method	→ moving GT PG↑	→ moving GT EBPG↑	→ static PG↑	crossover: cue = motion	crossover: cue = shape
Grad×Input	0.46	0.35	–	–	–
Integrated Grad.	**0.92**	**0.55**	–	–	–
3D Grad-CAM	0.19	0.10	–	–	–
Region occlusion	0.40	0.40	–	–	–
OccAM	0.17	0.13	–	–	–
Attn. Rollout	0.19	0.17	–	–	–
Motion-prior (floor)	0.79	0.27	–	–	–
KCA (local)	0.47	0.51	0.41	0.83	0.60
KCA (context)	0.48	0.50	0.38	0.80	0.60

Table 3: **Motion vs. geometry** (P4Transformer). Left: localisation against the skeleton ground truth on the 275 clips of Tab. 1 with a complete skeleton (PG: Pointing Game; S_{mot} against moving, S_{geo} against static parts). Right: *crossover* (Sec. 4.4, 40 test clips per task), motion share $S^{+}_{\text{mot}}/(S^{+}_{\text{mot}} + S^{+}_{\text{geo}})$ of the positive importance on the arm; single maps have no split (–).

Crossover: the decisive test. We therefore build two binary tasks from real clips in which the classes differ *only* in the right arm: in the *motion* task one class replays the arm at 25% speed (its shape is real), in the *shape* task one class thickens the arm 3× about its bones (its motion is untouched). We keep the videos in which the right arm carries at least 30% of the skeleton's motion (163 training and 200 test videos, one centred clip each) and train one P4Transformer per task; its accuracy over both classes of the test clips is 80.5 / 83.5 % (motion / shape task). In each task we explain the class whose defining property is present and removable by the interventions—normal-speed motion (removed by FREEZE) and the thick shape (removed by the proxy). We explain the first 40 test clips and pool the positive S_{mot} and S_{geo} over their arm points. For the *same* arm, the motion share of KCA's importance is 0.83 for the motion model (95% bootstrap CI [0.77, 0.88]; context estimator 0.80) and 0.60 for the shape model ([0.53, 0.67]; context 0.60)—a gap of [0.14, 0.32]—whereas the model-blind motion prior marks the arm identically in both tasks (EBPG 0.28 vs. 0.28). Across 6 training seeds of the two task models, the motion-task model leaves chance in only 3; in all 3 of these the arm's motion share is higher in the motion task, with an interval that excludes zero in 2 (Tab. 10). The split is not all-or-nothing on real models—the shape model still reacts to freezing the arm, whose frozen pose changes its per-frame configuration—but it moves with the cue, which no single map can show. Repeating the test with the convolutional PSTNet (80.5 / 90.2 % accuracy) gives a cleaner split, 0.86 vs. 0.30 (gap [0.48, 0.65]), and the state-space Mamba4D (77.5 / 89.2 %) one similar to P4Transformer, 0.92 vs. 0.61 (gap [0.22, 0.41]). On these two backbones the shape model still draws more than half of the arm's importance from motion, in line with their motion-heavy maps on MSR (whole-body motion share 0.75 for Mamba4D, 0.69 for P4Transformer): the audit shows which cue a model responds to, not that it uses that cue alone. Fig. 4 summarises all three next to the fully controlled setting of Appendix D (1.00 vs. 0.00).

4.5 Ablation, sanity and cost (Q4)

Components (Tab. 11 in the appendix). *Redundancy:* with a fine partition (m=24) a body part spans several regions and the local necessity collapses (deletion AUC 0.334), while the context estimator keeps it (0.212)—the failure mode it was designed for. *The* $(0, 0)$ *corner:* deleting the region instead of using the static proxy aligns the necessity with the deletion metric (local 0.272→0.235; context 0.219→0.126, the best of all variants) but weakens the motion term's localisation (EBPG 0.54→0.49); in the controlled study, where the answer is known, the same change makes a pure-shape model look 49% motion-driven

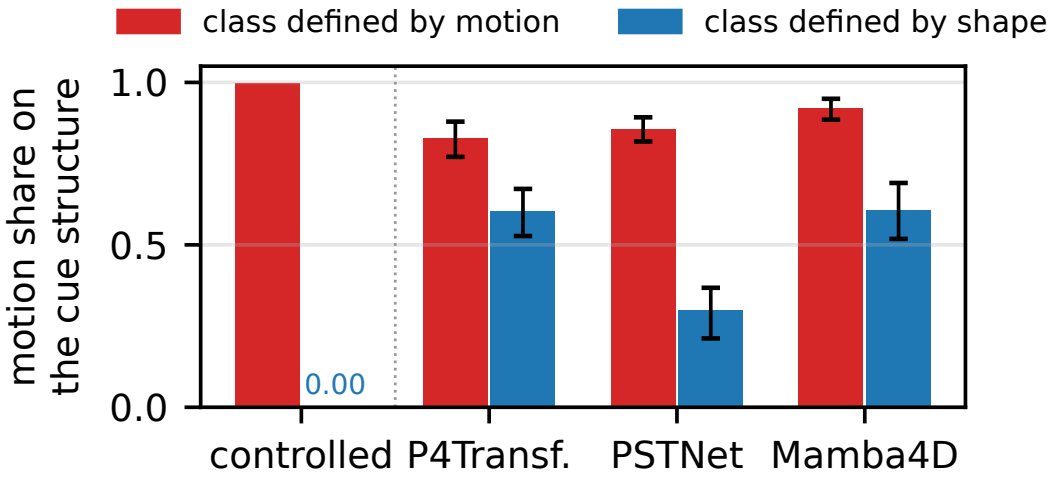

Figure 4: **Double dissociation:** motion share on the cue structure when the class is defined by its motion (red) or by its shape (blue), local estimator. Controlled task: pooled over 3 seeds $\times$ 20 clips. Real crossover: 40 test clips per task and backbone, with 95% bootstrap intervals.

(Appendix D). *The order:* the motion-first scores absorb the interaction into S_{mot} (motion share 0.88 vs. 0.73 symmetric). S=64 samples already give 0.231 (vs. 0.219 at S=256). Anchoring FREEZE and the proxy at the middle instead of the first frame barely changes the split (motion share 0.68 vs. 0.73) and does not hurt faithfulness (Tab. 11). The interaction is *not* negligible on real models ($|\Delta_{\text{int}}|/|\mathcal{N}|$ in the table, against 0.02 for the controlled network), which is exactly why we report order-symmetric scores and the residual rather than assume separability.

Sanity. Under the cascading randomisation of Adebayo et al. [2018], applied to P4Transformer from the head down on 20 clips (Fig. 5 in the appendix), the rank correlation of KCA's map with the trained model's drops to 0.12 as soon as the head is re-initialised and stays within ± 0.2 (-0.10 with every block randomised): the explanation depends on the learned weights. Grad$\times$Input keeps a correlation of 0.81 even then—much of it reflects the input rather than the model. On the region with the largest S_{mot} (40 clips), freezing it changes the prediction in 42% of the clips and reversing its trajectory (DISPLACE) in 22%: most motion evidence is about *whether* a part moves, less about *which way*.

Stability under re-sampling. A sensor never returns the same points twice, so we re-sample each of the 40 clips 5 times, explain every sampling, map each explanation onto the points of the first one and rank-correlate the two (Tab. 8 in the appendix). The prediction stays the same in 99% of the samplings. KCA's maps are the less stable ones point by point: 0.62 for the context estimator against 0.92 for Integrated Gradients and 0.71 for region occlusion, because the regions are re-fitted to every sampling. At the level of body parts (each map averaged over the 19 skeleton bones) the gap narrows (0.82 vs. 0.94), and the whole-body motion share varies little between samplings (standard deviation 0.04). Averaging the maps of two samplings, which costs only forward passes, raises the context estimator to 0.76 (points) and 0.89 (parts).

Cost. On one A100 GPU (mean over the 287 clips), the local estimator takes 0.5 s per clip and the context estimator 2.3 s (S=256 random states, batched), comparable to OccAM (1.5 s) and Integrated Gradients (1.9 s). All passes are independent forward queries, so the cost falls with batch throughput and needs no backward pass.

4.6 Qualitative results

Fig. 1 and Fig. 6 in the appendix overlay S_{mot} and S_{geo} on real test clips. On *draw circle* all three backbones put their strongest motion importance on the drawing hand, and on *forward kick* P4Transformer and PSTNet put it on the kicking leg. *High arm wave* is different: PSTNet marks the raised hand, but P4Transformer and Mamba4D put their motion importance on the torso rather than on the waving arm. Geometry importance spreads more widely over the body on every backbone. Per-action statistics are in Tab. 14.

5 Conclusion

KCA explains 4D point-cloud video models with forward queries alone and needs no point correspondences. It asks a question that no prior 4D explanation can pose: does a decision depend on *how* a structure moves or on *what shape* it has? Three findings support it. **(1) The split is correct.** Where the answer is known by construction, the motion share on the cue structure is 1.00 for a motion-defined and 0.00 for a shape-defined class (Appendix D). On real clips, when two models see the same arm but are trained on its motion or on its shape, the motion share of the arm's importance follows the cue on all three backbones (0.83 vs. 0.60 on P4Transformer, 0.86 vs. 0.30 on PSTNet, 0.92 vs. 0.61 on Mamba4D; every gap's 95% interval excludes zero), whereas a model-blind motion map marks the arm identically in both tasks and a single attribution map has no split at all. **(2) The split costs no faithfulness.** The necessity map has the lowest deletion AUC on P4Transformer (significantly so on the full test set) and on PSTNet, ties with the best baseline on Mamba4D, and loses its correlation with the trained model's map once the weights are randomised (-0.10, against 0.81 for Grad$\times$Input). **(3) It is deployable.** The same code explains convolutional, attention-based and attention-free backbones, while attention read-outs exist for only one of them, at 0.5–2.3 s per clip, comparable to Integrated Gradients. KCA is not the best map on every axis: white-box Integrated Gradients localises the moving body parts better, and its point-level maps are more stable under sensor re-sampling. What KCA adds is the motion/geometry split, which no gradient, attention or occlusion map provides.

Practical guidance. For auditing a 4D model before deployment we recommend the following. (i) Use the context-averaged estimator for the final report and the local one for quick inspection: the local one needs $3m{+}1$ passes (0.5 s per clip on one A100 GPU) and the context one S passes (2.3 s). (ii) Report Δ_{int} alongside S_{mot} and S_{geo}. On real models it is large, so a single-order split would be arbitrary. (iii) Average the maps of a few re-samplings of the same clip when the explanation must be stable; it costs only forward passes. (iv) When the intended cue is known, test it directly with a crossover-style pair of tasks. A model whose arm importance stays motion-dominated on a shape-defined task, as P4Transformer's and Mamba4D's do here, also responds to a cue that does not define its classes.

Limitations. (i) No pretrained weights exist, so we trained the backbones from scratch. P4Transformer comes close to its published accuracy, while PSTNet and Mamba4D fall short of theirs. A laptop-scale training gives the same picture (Appendix F). At full scale the P4Transformer motion-task model trains unreliably: it leaves chance in only 3 of 6 training seeds (Tab. 10). Further datasets (NTU RGB+D, HOI4D) remain to be tested. (ii) Skeleton-derived ground truth is approximate. It marks where motion *is*, not where the model *should* look, which is why our primary evidence for the decomposition is the crossover rather than

localisation. (iii) The decomposition's semantics require the proxy corner. With deletion at that corner the necessity map aligns better with removal-based metrics, but presence leaks into the motion term (Tab. 11). (iv) Regions come from a k-means tracker that can drift under fast or occluded motion, and the scores are only as meaningful as the regions they are computed on; because the regions are re-fitted to every sampling, KCA's point-level maps are less stable under sensor re-sampling than gradient maps.

Future work. Attributing 4D occupancy *forecasts* [Tian et al., 2023, Ma et al., 2024], so far served only by uncertainty estimates [Wang et al., 2025, Kälble et al., 2025, Heidrich et al., 2025], to the past motion that caused them, and interventional probes of adversarial behaviour on 4D models [Tan and Kotthaus, 2023, Li et al., 2021].

References

Samira Abnar and Willem Zuidema. Quantifying attention flow in transformers. In *ACL*, 2020.

Julius Adebayo, Justin Gilmer, Michael Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. Sanity checks for saliency maps. In *NeurIPS*, 2018.

Marco Ancona, Enea Ceolini, Cengiz Öztireli, and Markus Gross. Towards better understanding of gradient-based attribution methods for deep neural networks. In *ICLR*, 2018.

John F Banzhaf III. Weighted voting doesn't work: A mathematical analysis. *Rutgers Law Review*, 19: 317–343, 1965.

Yizhak Ben-Shabat, Oren Shrout, and Stephen Gould. 3DInAction: Understanding human actions in 3d point clouds. In *CVPR*, pages 19978–19987, 2024.

Asmita Chakraborty, Gizem Karagoz, and Nirvana Meratnia. A cross-domain benchmark of intrinsic and post hoc explainability for 3D deep learning models. *Journal of Imaging*, 12(2):63, 2026.

Hila Chefer, Shir Gur, and Lior Wolf. Transformer interpretability beyond attention visualization. In *CVPR*, 2021.

Matt De Vries, Reed Naidoo, Olga Fourkioti, Lucas G Dent, Nathan Curry, Chris Dunsby, and Chris Bakal. Interpretable point cloud classification using multiple instance learning. In *ICCV*, pages 22209–22220, 2025.

Jay DeYoung, Sarthak Jain, Nazneen Fatema Rajani, Eric Lehman, Caiming Xiong, Richard Socher, and Byron C Wallace. ERASER: A benchmark to evaluate rationalized NLP models. In *ACL*, 2020.

Olive Jean Dunn. Multiple comparisons among means. *Journal of the American Statistical Association*, 56 (293):52–64, 1961.

Bradley Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979.

Hehe Fan, Yi Yang, and Mohan Kankanhalli. Point 4d transformer networks for spatio-temporal modeling in point cloud videos. In *CVPR*, 2021a.

Hehe Fan, Yi Yang, and Mohan Kankanhalli. PSTNet: Point spatio-temporal convolution on point cloud sequences. In *ICLR*, 2021b.

Hehe Fan, Yi Yang, and Mohan Kankanhalli. Point spatio-temporal transformer networks for point cloud video modeling. *IEEE TPAMI*, 2023.

Gal Fiebelman, Tamir Cohen, Ayellet Morgenstern, Peter Hedman, and Hadar Averbuch-Elor. 4-LEGS: 4d language embedded gaussian splatting. *Computer Graphics Forum*, 44(2):e70085, 2025.

Ruth Fong, Mandela Patrick, and Andrea Vedaldi. Understanding deep networks via extremal perturbations and smooth masks. In *ICCV*, 2019.

Dominik Galus, Julia Farganus, Tymoteusz Zapala, Mikołaj Czachorowski, Piotr Borycki, Przemysław Spurek, and Piotr Syga. XSPLAIN: XAI-enabling splat-based prototype learning for attribute-aware interpretability. *arXiv preprint arXiv:2602.10239*, 2026.

Yash Goyal, Amir Feder, Uri Shalit, and Been Kim. Explaining classifiers with causal concept effect (CaCE). *arXiv preprint arXiv:1907.07165*, 2019.

Michel Grabisch and Marc Roubens. An axiomatic approach to the concept of interaction among players in cooperative games. *International Journal of Game Theory*, 28(4):547–565, 1999.

Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. In *COLM*, 2024.

Severin Heidrich, Till Beemelmanns, Alexey Nekrasov, Bastian Leibe, and Lutz Eckstein. OCCUQ: Exploring efficient uncertainty quantification for 3d occupancy prediction. In *ICRA*, 2025. arXiv:2503.10605.

Sarthak Jain and Byron C Wallace. Attention is not explanation. In *NAACL*, 2019.

Dominik Janzing, Lenon Minorics, and Patrick Blöbaum. Feature relevance quantification in explainable AI: A causal problem. In *AISTATS*, 2020.

Jonas Kälble, Sascha Wirges, Maxim Tatarchenko, and Eddy Ilg. EvOcc: Accurate semantic occupancy for automated driving using evidence theory. In *CVPR*, 2025.

Patrick Knab, Sascha Marton, Philipp J Schubert, Drago Guggiana, and Christian Bartelt. Concepts in motion: Temporal concept bottleneck model for interpretable video classification. *arXiv preprint arXiv:2509.20899*, 2025.

Joohyun Kwon, Hanbyel Cho, and Junmo Kim. Efficient dynamic scene editing via 4d gaussian-based static-dynamic separation. In *CVPR*, 2025. arXiv:2502.02091.

Jongseo Lee, Wooil Lee, Gyeong-Moon Park, Seong Tae Kim, and Jinwoo Choi. Disentangled concepts speak louder than words: Explainable video action recognition. In *NeurIPS*, 2025.

Haoxin Li, Yuan Liu, Hanwang Zhang, and Boyang Li. Mitigating and evaluating static bias of action representations in the background and the foreground. In *ICCV*, pages 19911–19923, 2023.

Wanhua Li, Renping Zhou, Jiawei Zhou, Yingwei Song, Johannes Herter, Minghan Qin, Gao Huang, and Hanspeter Pfister. 4D LangSplat: 4d language gaussian splatting via multimodal large language models. In *CVPR*, 2025. arXiv:2503.10437.

Wanqing Li, Zhengyou Zhang, and Zicheng Liu. Action recognition based on a bag of 3d points. In *CVPR Workshops*, 2010.

Yiming Li, Congcong Wen, Felix Juefei-Xu, and Chen Feng. Fooling LiDAR perception via adversarial trajectory perturbation. In *ICCV*, 2021.

Jiuming Liu, Jinru Han, Lihao Liu, Angelica I Aviles-Rivero, Chaokang Jiang, Zhe Liu, and Hesheng Wang. Mamba4D: Efficient 4d point cloud video understanding with disentangled spatial-temporal state space models. In *CVPR*, 2025. arXiv:2405.14338.

Xingyu Liu, Mengyuan Yan, and Jeannette Bohg. MeteorNet: Deep learning on dynamic 3d point cloud sequences. In *ICCV*, 2019.

Scott M Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *NeurIPS*, 2017.

Junyi Ma, Xieyuanli Chen, Jiawei Huang, Jingyi Xu, Zhen Luo, Jintao Xu, Weihao Gu, Rui Ai, and Hesheng Wang. Cam4DOcc: Benchmark for camera-only 4d occupancy forecasting in autonomous driving applications. In *CVPR*, pages 21486–21495, 2024.

Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An imperative style, high-performance deep learning library. In *NeurIPS*, 2019.

Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2 edition, 2009.

Vitali Petsiuk, Abir Das, and Kate Saenko. RISE: Randomized input sampling for explanation of black-box models. In *BMVC*, 2018.

Will Price and Dima Damen. Play fair: Frame contributions in video models. In *ACCV*, 2020.

Charles R Qi, Hao Su, Kaichun Mo, and Leonidas J Guibas. PointNet: Deep learning on point sets for 3d classification and segmentation. In *CVPR*, 2017a.

Charles R Qi, Li Yi, Hao Su, and Leonidas J Guibas. PointNet++: Deep hierarchical feature learning on point sets in a metric space. In *NeurIPS*, 2017b.

Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. "why should i trust you?": Explaining the predictions of any classifier. In *KDD*, 2016.

Avinab Saha, Shashank Gupta, Sravan Kumar Ankireddy, Karl Chahine, and Joydeep Ghosh. Exploring explainability in video action recognition. In *CVPR Workshops*, pages 8176–8181, 2024.

David Schinagl, Georg Krispel, Horst Possegger, Peter M Roth, and Horst Bischof. OccAM's Laser: Occlusion-based attribution maps for 3d object detectors on LiDAR data. In *CVPR*, 2022.

Ramprasaath R Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra. Grad-CAM: Visual explanations from deep networks via gradient-based localization. In *ICCV*, 2017.

Daniel Smilkov, Nikhil Thorat, Been Kim, Fernanda Viégas, and Martin Wattenberg. SmoothGrad: removing noise by adding noise. *arXiv preprint arXiv:1706.03825*, 2017.

Mukund Sundararajan, Ankur Taly, and Qiqi Yan. Axiomatic attribution for deep networks. In *ICML*, 2017.

Hanxiao Tan and Helena Kotthaus. Explainability-aware one point attack for point cloud neural networks. In *WACV*, 2023.

Jawad Tayyub, Muhammad Sarmad, and Nicolas Schönborn. Explaining deep neural networks for point clouds using gradient-based visualizations. In *ACCV*, 2022.

Xiaoyu Tian, Tao Jiang, Longfei Yun, Yucheng Mao, Huitong Yang, Yue Wang, Yilun Wang, and Hang Zhao. Occ3D: A large-scale 3d occupancy prediction benchmark for autonomous driving. In *NeurIPS*, 2023.

Payal Varshney, Adriano Lucieri, Christoph Balada, Sheraz Ahmed, and Andreas Dengel. LD-ViCE: Latent diffusion model for video counterfactual explanations. *arXiv preprint arXiv:2509.08422*, 2025.

Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, et al. SciPy 1.0: Fundamental algorithms for scientific computing in Python. *Nature Methods*, 17(3):261–272, 2020.

Haofan Wang, Zifan Wang, Mengnan Du, Fan Yang, Zijian Zhang, Sirui Ding, Piotr Mardziel, and Xia Hu. Score-CAM: Score-weighted visual explanations for convolutional neural networks. In *CVPR Workshops*, 2020.

Song Wang, Zhongdao Wang, Jiawei Yu, Wentong Li, Bailan Feng, Junbo Chen, and Jianke Zhu. Reliable and calibrated semantic occupancy prediction by hybrid uncertainty learning. In *IJCAI*, pages 1973–1981, 2025.

Wei Wang, Yiding Sun, Yuyan Wang, Zhuoyue Zhang, Zhengqiao Li, Dongfu Yin, and Chen Li. Motion-saliency complementary masked modeling for point cloud video understanding. *arXiv preprint arXiv:2608.30279*, 2026.

Yue Wang, Yongbin Sun, Ziwei Liu, Sanjay E Sarma, Michael M Bronstein, and Justin M Solomon. Dynamic graph CNN for learning on point clouds. *ACM Transactions on Graphics (TOG)*, 2019.

Zerui Wang and Yan Liu. STAA: Spatio-temporal attention attribution for real-time interpreting transformer-based AI video models. *IEEE Access*, 13:101647–101661, 2025.

David S Watson, Limor Gultchin, Ankur Taly, and Luciano Floridi. Local explanations via necessity and sufficiency: Unifying theory and practice. In *UAI*, 2021.

Hao Wen, Yunze Liu, Jingwei Huang, Bo Duan, and Li Yi. Point primitive transformer for long-term 4d point cloud video understanding. In *ECCV*, 2022.

Frank Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945.

Xiaoyang Wu, Li Jiang, Peng-Shuai Wang, Zhijian Liu, Xihui Liu, Yu Qiao, Wanli Ouyang, Tong He, and Hengshuang Zhao. Point transformer V3: Simpler, faster, stronger. In *CVPR*, 2024.

Xindi Wu, Despoina Paschalidou, Jun Gao, Antonio Torralba, Laura Leal-Taixé, Olga Russakovsky, Sanja Fidler, and Jonathan Lorraine. Motion attribution for video generation. In *ICML*, 2026.

Jianming Zhang, Sarah Adel Bargal, Zhe Lin, Jonathan Brandt, Xiaohui Shen, and Stan Sclaroff. Top-down neural attention by excitation backprop. *International Journal of Computer Vision (IJCV)*, 2018.

Zhaoliang Zhang, Tianchen Song, Yongjae Lee, Li Yang, Cheng Peng, Rama Chellappa, and Deliang Fan. LP-3DGS: Learning to prune 3d gaussian splatting. In *NeurIPS*, 2024. arXiv:2405.18784.

Tianhang Zheng, Changyou Chen, Junsong Yuan, Bo Li, and Kui Ren. PointCloud saliency maps. In *ICCV*, 2019.

A Notation and pseudocode

Tab. 4 lists the notation, Tab. 5 the assumptions behind each result, and Alg. 1 the procedure.

$X = \{P_t\}_{t=1}^{T}$	point-cloud video; $P_t \subset \mathbb{R}^3$, sampled per frame
$f,\ \varphi(X)=f(X)_c$	model; target logit of class c
$\mathcal{R} = \{\mathcal{R}_t\}_t$	a tracked region: its point set in each frame
$c_t(\mathcal{R})$	centroid of $\mathcal{R}_t$
$\mathrm{do}(\textsc{freeze}, \mathcal{R})$	freeze: $\mathcal{R}$ static at its t_0 shape (M=0, G=1)
$\mathrm{do}(\textsc{abl-geo}, \mathcal{R})$	shape-free proxy moving along c_t (M=1, G=0)
$\mathrm{do}(\textsc{ablate}, \mathcal{R})$	shape-free proxy static at c_{t_0} (M=0, G=0)
$\mathcal{N},\ S_{\mathrm{mot}}, S_{\mathrm{geo}},\ \Delta_{\mathrm{int}}$	necessity; motion, geometry importance; interaction

Table 4: Notation.

	Assumption	Needed by
I	interventional validity (operators rewrite only $\mathcal{R}_t$)	Props. 2, 3, 4
S	separability *(measured by Δ_{int}, not assumed)*	Prop. 2 (marginal ACE only)
M	monotonicity & exogeneity	Prop. 3 (identification)

Table 5: Assumptions and the results that require them.

Algorithm 1 KCA: Kinematic Causal Attribution

Require: clip X, model f, class c, tracked regions $\{\mathcal{R}^j\}_{j=1}^{m}$
Ensure: motion field S_{mot}, geometry field S_{geo}, interaction Δ_{int}

precompute for all regions: frozen sets $\mathcal{R}^j_{t_0}$, proxy paths $c_t(\mathcal{R}^j)$
if local **then**
    evaluate φ at X and 3 corners per region ▷ $3m{+}1$ passes
    Eq. (2), Eq. (3) per region
else
    evaluate φ on S random joint states ▷ S passes
    solve Eq. (4); $S_{\mathrm{mot}}{=}\mu{+}\iota/2$, $S_{\mathrm{geo}}{=}\gamma{+}\iota/2$
end if
return per-point fields: each point receives its region's scores

B Full proofs

Notation: $\varphi(X) = f(X)_c$; for a region $\mathcal{R}$ the four corners of its (M, G) cube are X $(M{=}1, G{=}1)$, $\mathrm{do}(\textsc{freeze}, \mathcal{R})X$ $(0, 1)$, $\mathrm{do}(\textsc{abl-geo}, \mathcal{R})X$ $(1, 0)$ and $\mathrm{do}(\textsc{ablate}, \mathcal{R})X$ $(0, 0)$, with outputs φ, $\varphi_{\textsc{freeze}}$, $\varphi_{\textsc{abl-geo}}$, $\varphi_{\textsc{ablate}}$.

Proposition 1 (completeness, Shapley, interaction). Consider the two-player game on $\{M, G\}$ whose worth of a coalition K is $v(K) = \varphi$ evaluated at the corner where exactly the factors in K are *kept*: $v(\{M, G\}) = \varphi$, $v(\{G\}) = \varphi_{\text{FREEZE}}$, $v(\{M\}) = \varphi_{\text{ABL-GEO}}$, $v(\emptyset) = \varphi_{\text{ABLATE}}$. The Shapley value of M averages its marginal contribution over the two orders: $\frac{1}{2}[(v(\{M, G\}) - v(\{G\})) + (v(\{M\}) - v(\emptyset))] = \frac{1}{2}[(\varphi - \varphi_{\text{FREEZE}}) + (\varphi_{\text{ABL-GEO}} - \varphi_{\text{ABLATE}})] = S_{\text{mot}}$, and likewise for G, giving S_{geo}. Efficiency of the Shapley value gives $S_{\text{mot}} + S_{\text{geo}} = v(\{M, G\}) - v(\emptyset) = \varphi - \varphi_{\text{ABLATE}} = \mathcal{N}$, for every f. The Shapley interaction index of the pair [Grabisch and Roubens, 1999] is $v(\{M, G\}) - v(\{M\}) - v(\{G\}) + v(\emptyset) = \Delta_{\text{int}}$. The motion-first score $\varphi - \varphi_{\text{FREEZE}}$ equals $S_{\text{mot}} + \Delta_{\text{int}}/2$ and the geometry-first one $\varphi_{\text{ABL-GEO}} - \varphi_{\text{ABLATE}}$ equals $S_{\text{mot}} - \Delta_{\text{int}}/2$ (direct substitution), so both deviate from S_{mot} by $\pm\Delta_{\text{int}}/2$. Finally, φ restricted to the cube is additively separable, $\varphi(M, G) = a(M) + b(G)$, iff its mixed second difference vanishes, *i.e.* iff $\Delta_{\text{int}} = 0$. □

Proposition 2 (motion effects). By **(I)** the operators rewrite only $\mathcal{R}$'s per-frame sets, so $\text{do}(\text{FREEZE}, \mathcal{R})$ sets $M_{\mathcal{R}}$=0 while holding $G_{\mathcal{R}}$ (the real shape, at its t_0 configuration) and all off-region content fixed: $\varphi - \varphi_{\text{FREEZE}}$ is the single-factor interventional contrast of $\mathcal{R}$'s motion at its real geometry, a CaCE-type effect [Goyal et al., 2019]; $\varphi_{\text{ABL-GEO}} - \varphi_{\text{ABLATE}}$ is the same contrast for the shape-free proxy (same centroid path vs. static). Under **(S)**, $\varphi(M, G) = a(M) + b(G)$ on the cube and both contrasts equal $a(1) - a(0)$, the marginal ACE of motion; S_{mot}, their average, equals it too. The geometry statement is symmetric. □

Proposition 3 (necessity). Let $E = \mathbf{1}[\varphi \geq \kappa]$. $E(X)$=1 is Pearl's factual event and $E(\text{do}(\text{ABLATE}, \mathcal{R})X)$=0 the counterfactual outcome under the do-operation that replaces $\mathcal{R}$'s content (valid by **(I)**); substituting into the definition of PN [Pearl, 2009] gives $\text{PN}(\mathcal{R}) = \Pr[E(\text{do}(\text{ABLATE}, \mathcal{R})X)=0 \mid E(X)=1]$. Under monotonicity and exogeneity **(M)**, Pearl's identification gives the excess-risk ratio $[P(E=1 \mid \text{intact}) - P(E=1 \mid \text{replaced})]/P(E=1 \mid \text{intact})$, whose numerator is the expectation of the hard-thresholded $\mathcal{N}$; the real-valued $\mathcal{N}$ is its margin relaxation, matching the local PN of [Watson et al., 2021]. The sufficiency statement is dual. □

Proposition 4 (context-averaged decomposition). Completeness holds by construction: $S_{\text{mot}} + S_{\text{geo}} = \mu + \gamma + \iota = \mathcal{N}$ for every region. For the population statement, let $Z = (Z_1, \ldots, Z_m)$ with $Z_j = (M_j, G_j)$ be the independent Bernoulli(q) factors and $Y = \varphi(X_Z)$. Under independence Y has the ANOVA (Hoeffding) decomposition $Y = \mu_0 + \sum_j f_j(Z_j) + \sum_{j<k} f_{jk}(Z_j, Z_k) + \ldots$, where every higher-order term has zero conditional mean given any single block Z_j, hence is L^2-orthogonal to every function of a single block. The regressors of Eq. (4), $\{1\} \cup \{M_j, G_j, M_j G_j\}_j$, span exactly the functions $c + \sum_j g_j(Z_j)$ (a function of two binary variables is spanned by $1, M, G, MG$). The population least-squares solution is therefore the L^2 projection $\mu_0 + \sum_j f_j(Z_j)$, with $f_j(z) = \mathbb{E}[Y \mid Z_j{=}z] - \mathbb{E}[Y]$. The fitted coefficients of block j are thus the cube contrasts of $\bar{\varphi}_j(M, G) = \mathbb{E}[Y \mid M_j{=}M, G_j{=}G]$—the region's four corners averaged over random contexts of all other regions—and $S_{\text{mot}} = \mu + \iota/2$, $S_{\text{geo}} = \gamma + \iota/2$, $\Delta_{\text{int}} = \iota$ are exactly Eqs. (2)–(3) applied to $\bar{\varphi}_j$ (for $q = \frac{1}{2}$: Banzhaf-type [Banzhaf III, 1965] main effects of the two factors, split symmetrically). If φ has no cross-region interactions, every context gives the same contrasts, so $\bar{\varphi}_j$'s contrasts coincide with the local ones. □

C Implementation details

Data and skeleton ground truth. MSR-Action3D [Li et al., 2010] (567 depth videos of 20 actions by 10 subjects, 320×240) is converted to point-cloud videos exactly as in MeteorNet [Liu et al., 2019] (focal 280, principal point at the image centre, z = depth; the input format of all three official backbones), then

sampled to 24-frame centred clips of 2048 points per frame (1024 at laptop scale; independently per frame) and divided by 300. The dataset's 20-joint skeletons store (u, v, d) per joint in depth-image pixels; we lift each joint with the same projection, taking its depth from the body surface around (u, v) (median of the non-zero 7×7 window; joints outside the silhouette use a per-frame unit ratio fitted on the others). Over the 275 benchmark clips with a complete skeleton, a median 95% of the points (10th percentile: 92%) lie within 25 units of a bone. One unit is about 5 mm: the median head-to-foot joint distance is 318 units, which we take to be about 1.55 m, so 25 units are about 12 cm. Each point is assigned to its nearest of 19 bone segments; a bone is *moving* if its midpoint path length over the clip is at least $0.5\times$ the clip's maximum and *static* if at most $0.2\times$; moving-GT / static-GT are the points of moving / static bones (on average 13% / 65% of the points of a benchmark clip). 12 benchmark clips, all from *bend* and *pick up & throw*, store all-zero joints for a run of frames; their nearest-bone labels are meaningless there, so they are left out of every ground-truth-based number (they still count for faithfulness, which needs no ground truth).

Tracked regions and operators. Regions are tracked k-means supervoxels (m=12 unless stated): frame 0 uses farthest-point initialisation from the lowest point (so the partition does not depend on point order), each later frame re-runs k-means from the previous frame's centres (4 iterations). Operators act on the per-frame index sets: FREEZE fills the slots of $\mathcal{R}_t$ with the points of $\mathcal{R}_{t_0}$ (paired cyclically; t_0 = first frame); ABL-GEO writes the centroid $c_t(\mathcal{R})$ into every slot of $\mathcal{R}_t$ (a point mass on the centroid path); ABLATE writes $c_{t_0}(\mathcal{R})$; DISPLACE fills $\mathcal{R}_t$ from $\mathcal{R}_{\pi(t)}$. Because each operator's new values depend only on the region's own points, the joint state of all regions is a per-slot selection among three precomputed views, which lets both estimators evaluate hundreds of intervened clips per batch. A test in the released code checks that shuffling every frame's point order leaves the scores unchanged ($< 10^{-6}$).

Backbones. No pretrained MSR weights exist for any of the three models, so we train the *unmodified official model code* from scratch. Their compiled CUDA ops (FPS, ball query, grouping; `mamba_ssm`; `knn_cuda`) are replaced by numerically equivalent pure-PyTorch versions (FPS seeded at index 0, ball query returning the first in-radius indices, the reference selective scan), so the models also run on CPU/Apple GPUs. On CUDA, two of these loops are replaced by kernels that are checked against them on random inputs before use: a single-kernel FPS with the same selection rule, and the compiled selective scan of `mamba_ssm`. *Full scale* (main text): the official configuration of every model (44.1 M / 8.3 M / 109.8 M parameters for P4Transformer / PSTNet / Mamba4D), 2048 points per frame, one A100 GPU; AdamW (3×10^{-4}, weight decay 0.05), 5 warm-up epochs then cosine decay, the official batch sizes (14 / 16 / 16), 50 / 35 / 50 epochs, 17 random 24-frame clips per training video per epoch (about the number the official loaders enumerate), the official augmentation (random per-axis scaling in $[0.9, 1.1]$), TF32 matrix products during training and FP32 for explanation. *Laptop scale* (Appendix F): P4Transformer dim 256, depth 4, 4 heads of 64 (spatial stride 16 for 1024 points, *i.e.* the official 64 anchors per frame), Mamba4D dim 256 with 2 inter- and 1 intra-frame Mamba layers, PSTNet official; 1024 points, batch 8, 50 / 30 / 20 epochs, 5 / 3 / 3 clips per video per epoch, on an Apple M4. Both scales report the last epoch; no model is selected on the test set. Two details were necessary for from-scratch training to leave its initial plateau: P4Transformer and Mamba4D embed *absolute* (x, y, z, t) with $t = 1, \ldots, L$, so at initialisation the time channel dominates every token ($\|\text{pos}\| \approx 34$ vs. $\|\text{feature}\| \approx 7$) and the token max-pool selects tokens by frame index rather than content; we therefore scale the *initial* weights of that channel by $1/L$ (the architecture is unchanged) and subtract each clip's centroid (a preprocessing step applied identically at explanation time). With the official SGD recipe the reduced models did not leave the plateau within the $\approx10^3$ iterations a laptop affords, so both scales use AdamW.

Baselines and metrics. Integrated Gradients and Extremal Perturbations use the clip centroid as their reference (a far-away constant is off-manifold for real point networks); 3D Grad-CAM uses the gradient-

weighted tube-embedding tokens of P4Transformer/Mamba4D, scattered to points through their nearest anchors (an input-space proxy for PSTNet); attention rollout and transformer attribution (generic attention explainability [Chefer et al., 2021]) use all P4Transformer layers; raw attention is the attention each token receives in the last layer, scattered to points; region occlusion deletes each tracked region; OccAM averages the output over 128 random 50% point masks; Extremal optimises a per-point mask for 30 Adam steps at 10% area. Deletion/insertion remove/insert points in importance order in 20 steps; a removed point is deleted and its frame re-sampled from the remaining points (the official loaders' rule for short frames); all curves score the target-class *probability*. Comprehensiveness/Sufficiency use AOPC over $\{1, 5, 10, 20, 50\}\%$; Sensitivity-n uses 30 random 10% subsets.

D Controlled experiment: ground truth by construction

To test the decomposition where the answer is known *by construction*, every clip has three point clusters with the same layout in all classes; one "odd" cluster differs and the class is *which* cluster is odd. In the **motion** task the odd cluster translates (all shapes identical: only motion is discriminative); in the **shape** task it is larger and nothing moves (only geometry is discriminative). An attention-based 4D classifier (per-point MLP $\rightarrow$ per-frame voxel tokens $\rightarrow$ Transformer) is trained per task and seed (3 seeds, 20 held-out clips each) and explained with tracked regions (m=3). Within the odd cluster, the motion share of KCA's positive importance is 1.00 for the motion model and 0.00 for the shape model (local estimator; context: 1.00 vs. 0.00), *i.e.* a complete double dissociation. With the earlier *deletion* corner instead of the static proxy, a static region's presence leaks into S_{mot} (Tab. 6). On a third task whose class is the motion *direction* and whose two directions share the same spatial envelope, reversing the actor's trajectory (DISPLACE) flips the prediction to the opposite class in 90% of the clips, freezing it in 0%: DISPLACE is sign-sensitive, FREEZE is not.

$(0, 0)$ corner, order	mot.	shape	Del.↓
static proxy (default), symmetric	1.00	0.00	0.428
static proxy, motion-first	1.00	0.00	0.428
deletion, symmetric	0.92	0.49	0.419
deletion, motion-first	0.85	0.00	0.419
donor fill, symmetric	0.92	0.49	0.419
donor fill, motion-first	0.85	0.00	0.419

Table 6: Controlled ablation (1 seed): the $(0, 0)$ corner and the order. Motion share within the odd cluster on the motion (mot.) and shape (shape) models; the decomposition is correct iff mot. ≈ 1 and shape ≈ 0.

E Data, models and code behind each result

Data. MSR-Action3D [Li et al., 2010]: raw depth maps and 20-joint skeletons from the archive distributed with MeteorNet [Liu et al., 2019], converted with MeteorNet's depth-to-point procedure (Appendix C). Every real-data result uses test subjects 6–10 only; models are trained on subjects 1–5. Each clip is the centred 24-frame window of a video, and its points are sampled with a fixed seed per video, so every method explains exactly the same inputs.

Models and code. The unmodified official repositories of P4Transformer [Fan et al., 2021a] (`github.com/hehefan/P4Transformer`, commit `461b83e`), PSTNet [Fan et al., 2021b] (`github.com/hehefan/Point-Spatio-Temporal-Convolution`, `04f9b47`) and Mamba4D [Liu et al., 2025] (commit `8d5bbe2`).

Their compiled operators are replaced as described in Appendix C; on CUDA the selective scan is the kernel of `mamba_ssm` 2.3.2 [Gu and Dao, 2024]. Full scale: PyTorch [Paszke et al., 2019] 2.11 with CUDA 13.0 on one A100 (40 GB). Laptop scale: PyTorch 2.14 on an Apple M4. Statistical tests use SciPy [Virtanen et al., 2020].

From records to numbers. Every explained clip is written to a per-clip record (prediction, attribution scores, every metric). Scripts compute every table body and every measured number in the text from these records; no measured number is typed by hand. Tab. 7 lists the record behind each result.

Result	Data	Model(s)	Record
Tab. 1, 3 (left), 12, 13, 14; Fig. 3, 1, 6	all 287 test clips, 2048 points per frame; ground-truth metrics on the 275 with a complete skeleton	P4Transformer (Fig. 6: all three)	benchmark, `p4transformer/raw.jsonl`
Tab. 2	the first two test clips per action (40)	all three	benchmark, one file per backbone
Tab. 3 (right), Fig. 4, Appendix G	arm-moving videos: 163 train, 200 test; 40 test clips explained per task	one task model per task and backbone	`crossover*/crossover.json`
Tab. 11, Fig. 5; displace probe	40 clips (sanity: the first 20)	P4Transformer	`ablation/ablation.json`
Tab. 8	the same 40 clips, 5 point samplings each	P4Transformer	`stability/stability.jsonl`
Tab. 10	as the crossover, 6 training seeds	P4Transformer task models	`crossover*/crossover.json`
Tab. 6; controlled bar of Fig. 4	synthetic clips, 3 seeds $\times$ 20 test clips	small attention model	`controlled/*.json`
Tab. 9	as above, laptop-scale models	all three	laptop-scale records

Table 7: The data, models and per-clip records behind each table and figure. All real-data results use the MSR-Action3D test subjects 6–10 and, except Tab. 9, the full-scale models.

F Two training scales

The main text uses backbones trained at their official width with 2048 points per frame on one A100 GPU. We also trained all three at laptop scale (reduced width for P4Transformer and Mamba4D, 1024 points, an Apple M4; Appendix C) and ran every real-data experiment again. Tab. 9 sets the headline numbers side by side. The full-scale models are closer to their published accuracies, most of all Mamba4D. The conclusions agree, with three differences in significance. On P4Transformer, KCA's deletion advantage over region occlusion is significant at laptop scale on 100 clips and at full scale on all 287 clips ($p<10^{-3}$), but not at full scale on the same 100 clips ($p=0.10$): the full test set is what makes it significant. On Mamba4D its advantage over Integrated Gradients is significant only at laptop scale. Integrated Gradients localises the moving parts better than S_{mot} on P4Transformer at full scale, where the laptop-scale run had them tied. In the crossover the shape-task models of P4Transformer and Mamba4D draw more of the arm's importance from motion at full scale, but every motion–shape gap still excludes zero.

	one sampling		average of two samplings	
Method	points	body parts	points	body parts
Random (floor)	0.00	−0.01	0.00	0.01
Motion-prior (floor)	0.66	0.99	0.73	0.99
Grad×Input	0.82	0.86	0.81	0.87
Integrated Grad.	0.92	0.94	0.91	0.93
Region occlusion	0.71	0.85	0.81	0.91
OccAM	0.00	0.04	0.00	0.04
Attn. Rollout	0.68	0.70	0.75	0.76
KCA (local)	0.43	0.68	0.57	0.80
KCA (context)	0.62	0.82	0.76	0.89

Table 8: **Stability under sensor re-sampling** (P4Transformer, 40 test clips, 5 independent point samplings each). Spearman rank correlation between the map of sampling 0 and the map of another sampling transferred onto the points of sampling 0 (nearest neighbour per frame), averaged over samplings and clips. Body parts: each map averaged over the points of each of the 19 skeleton bones (one clip without a complete skeleton is left out). Average of two samplings: maps of samplings 0 and 1 averaged and compared with the average of samplings 2 and 3. Method randomness is fixed per clip, so only the point sampling changes.

Quantity	laptop	full scale
Test accuracy, P4Transformer (%)	90.6	89.2
Test accuracy, PSTNet (%)	87.5	88.9
Test accuracy, Mamba4D (%)	83.6	88.2
P4T Del.: KCA (context)	0.178	0.179
P4T Del.: region occlusion	0.219	0.233
P4T Del.: Integrated Gradients	0.268	0.282
P4T Del.: context vs. occlusion	p=0.0043	$p<10^{-3}$
P4T Del.: context vs. IG	$p<10^{-3}$	$p<10^{-3}$
PSTNet Del.: context / occlusion	0.191 / 0.222	0.217 / 0.252
Mamba4D Del.: context / occlusion	0.245 / 0.240	0.269 / 0.246
Mamba4D Del.: context vs. IG	$p<10^{-3}$	p=0.064
P4T S_{mot} EBPG: context / IG	0.54 / 0.58	0.50 / 0.55
Crossover P4T: motion / shape task	0.81 / 0.39	0.83 / 0.60
Crossover PSTNet: motion / shape task	0.85 / 0.23	0.86 / 0.30
Crossover Mamba4D: motion / shape task	0.89 / 0.54	0.92 / 0.61
Randomised model: KCA / Grad×Input	−0.01 / 0.72	−0.10 / 0.81

Table 9: Headline numbers at the two training scales (laptop: 100 P4Transformer clips; full scale: all 287; the other rows use the same clips at both scales. Del.: deletion AUC; crossover: motion share of the arm's importance; randomised model: rank correlation with the trained model's map once every block is re-initialised).

G Additional real-data results

Crossover design. We explain, in each task, the class whose defining property is *present* in the clip and removable by the interventions: the original class on original clips in the motion task (its normal-speed motion is removed by FREEZE) and the thick class on thickened clips in the shape task (its shape is removed by

the proxy). We report the variants we tried on the way. Explaining the *original* class in the shape task—whose evidence is then the *absence* of the cue—gave shares that depend on the sign convention (positive parts 0.88, signed sums 0.49): freezing a normally shaped arm moves that logit strongly in both directions from clip to clip. Thinning the arm instead of thickening it was not learnable by P4Transformer (test accuracy 50%, i.e. chance), so that variant gives no information. With the final design both conventions agree (motion task: 0.83 positive / 0.84 signed; shape task: 0.60 / 0.60). 95% bootstrap intervals over clips (positive parts, local estimator): motion task [0.77, 0.88], shape task [0.53, 0.67], gap [0.14, 0.32]. We retrained both task models with further seeds (Tab. 10). At full scale the motion-task model often fails to leave chance: it trains in 3 of 6 seeds, and the shares of a model at chance carry no information, so we do not use them. In all 3 seeds whose models trained, the arm's motion share is higher in the motion task; the interval of the gap excludes zero in 2 of them, and the smallest gap (seed 3) comes with the weakest shape model. Three of the crossover videos (one training, two test, none of them among the explained clips) have frames without a skeleton; their arm masks are wrong in those frames, and we keep them.

Seed	acc. % (mot. / shape)	share (mot. / shape)	gap 95% CI
0	80.5 / 83.5	0.83 / 0.60	[0.14, 0.32]
1 *(not trained)*	50.0 / 86.5	0.65 / 0.55	[−0.02, 0.21]
2	67.8 / 88.8	0.73 / 0.57	[0.06, 0.26]
3	78.8 / 77.5	0.85 / 0.78	[−0.02, 0.19]
4 *(not trained)*	50.0 / 72.2	0.60 / 0.81	[−0.32, −0.08]
5 *(not trained)*	50.0 / 80.5	0.73 / 0.69	[−0.05, 0.17]

Table 10: P4Transformer crossover over training seeds (local estimator, 40 explained test clips per task). A seed counts only if both task models leave chance (accuracy $\geq 60\%$ on the balanced binary test set); seed 0 is the one reported in the main text.

PSTNet (same tasks and recipe): 0.86 [0.82, 0.89] and 0.30 [0.21, 0.37] (context estimator 0.88 / 0.26; motion prior on the arm 0.28 / 0.28 EBPG in both tasks); the signed shares (1.08, 0.16) agree in direction and fall just outside $[0, 1]$ because the other term's sum over the arm is slightly negative. *Mamba4D*: 0.92 [0.89, 0.95] and 0.61 [0.52, 0.69] (context 0.90 / 0.58; signed 1.03 / 0.61).

Variant	Del.↓	S_{mot} EBPG↑	mot. share	$\frac{\lvert\Delta_{\text{int}}\rvert}{\lvert\mathcal{N}\rvert}$	passes	s
local, m=6	0.235	0.47	0.65	0.46	19	0.3
local, m=12 (default)	0.272	0.54	0.73	0.63	37	0.5
local, m=24	0.334	0.54	0.78	0.70	73	0.8
context, m=6	0.192	0.44	0.65	0.48	257	2.3
context, m=12 (default)	0.219	0.54	0.75	0.58	257	2.3
context, m=24	0.212	0.52	0.82	0.66	257	2.4
local, removal instead of proxy	0.235	0.49	0.70	0.75	37	0.5
context, removal instead of proxy	0.126	0.51	0.72	0.58	257	3.0
local, donor fill instead of proxy	0.208	0.50	0.70	0.72	37	2.1
local, motion-first order	0.272	0.53	0.88	–	25	1.5
context, S=64	0.231	0.49	0.75	0.77	65	0.7
context, S=128	0.204	0.53	0.76	0.64	129	1.2
context, S=512	0.211	0.55	0.76	0.52	513	4.5
context, q=0.75	0.233	0.55	0.74	0.63	257	2.3
local, anchor t_0 = middle frame	0.245	0.56	0.68	0.65	37	0.5
context, anchor t_0 = middle frame	0.173	0.57	0.69	0.58	257	2.3

Table 11: **Ablation and cost** (P4Transformer, 40 clips). EBPG and motion share on the moving ground truth; $|\Delta_{\text{int}}|/|\mathcal{N}|$: mean interaction relative to necessity; forward passes and seconds per clip on one A100 GPU.

Method	Del.↓	Comp.↑	Ins.↑†	Suff.↓†	Sens-n ↑
Random (floor)	0.916±0.14	0.003±0.05	0.920±0.14	0.041±0.12	0.004±0.17
Geometry-edge (floor)	0.890±0.19	0.011±0.11	0.903±0.16	0.103±0.21	0.000±0.17
Motion-prior (floor)	0.505±0.38	0.277±0.33	0.860±0.16	0.365±0.22	0.000±0.18
Grad×Input	0.655±0.28	0.086±0.18	0.722±0.25	0.332±0.28	0.028±0.23
SmoothGrad	0.772±0.24	0.031±0.12	0.915±0.14	0.083±0.16	0.006±0.18
Integrated Grad.	0.282±0.26	0.347±0.24	0.862±0.17	0.129±0.19	**0.046±0.22**
3D Grad-CAM	0.713±0.26	0.051±0.12	0.554±0.30	0.659±0.30	0.000±0.15
Region occlusion	0.233±0.23	0.416±0.23	0.784±0.17	0.531±0.24	0.021±0.16
OccAM	0.853±0.24	0.059±0.13	**0.953±0.07**	**−0.004±0.11**	0.003±0.17
Extremal Pert.	0.634±0.22	0.052±0.10	0.309±0.26	0.830±0.25	0.005±0.17
Attn. Rollout	0.649±0.27	0.073±0.15	0.769±0.24	0.431±0.26	−0.009±0.17
Transformer Attr.	0.601±0.28	0.100±0.17	0.766±0.24	0.424±0.28	0.003±0.17
Raw attention (last layer)	0.542±0.28	0.127±0.16	0.688±0.25	0.551±0.29	0.006±0.16
KCA (local)	0.256±0.25	0.403±0.24	0.793±0.17	0.511±0.23	0.018±0.17
KCA (context)	**0.179±0.19***	**0.453±0.21***	0.761±0.18	0.546±0.24	0.011±0.17

Table 12: Full faithfulness table on P4Transformer, including Sensitivity-n (Pearson correlation between summed attribution and output drop over random 10% subsets; near zero for all methods, because random subsets barely change the output).

Method	s / clip
Random (floor)	0.00
Geometry-edge (floor)	0.14
Motion-prior (floor)	0.16
Grad×Input	0.06
SmoothGrad	1.53
Integrated Grad.	1.93
3D Grad-CAM	0.07
Region occlusion	0.57
OccAM	1.46
Extremal Pert.	2.02
Attn. Rollout	0.05
Transformer Attr.	0.07
Raw attention (last layer)	0.05
KCA (local)	0.47
KCA (context)	2.33

Table 13: Explanation cost per clip on one A100 GPU (P4Transformer, 24 frames × 2048 points; seconds, excluding metric evaluation).

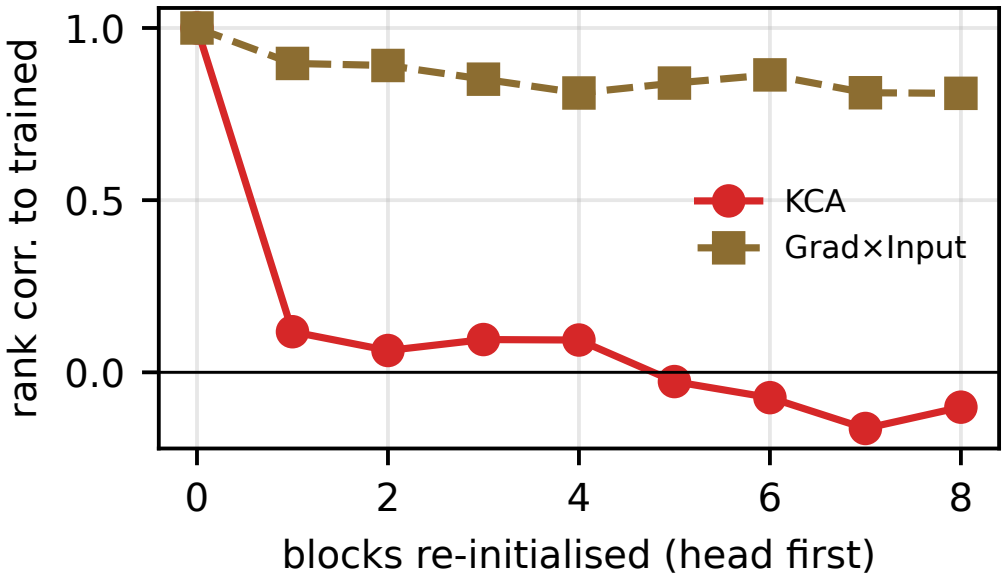

Figure 5: Cascading model randomisation [Adebayo et al., 2018] of P4Transformer, from the classification head down: rank correlation of the attribution map with the trained model's.

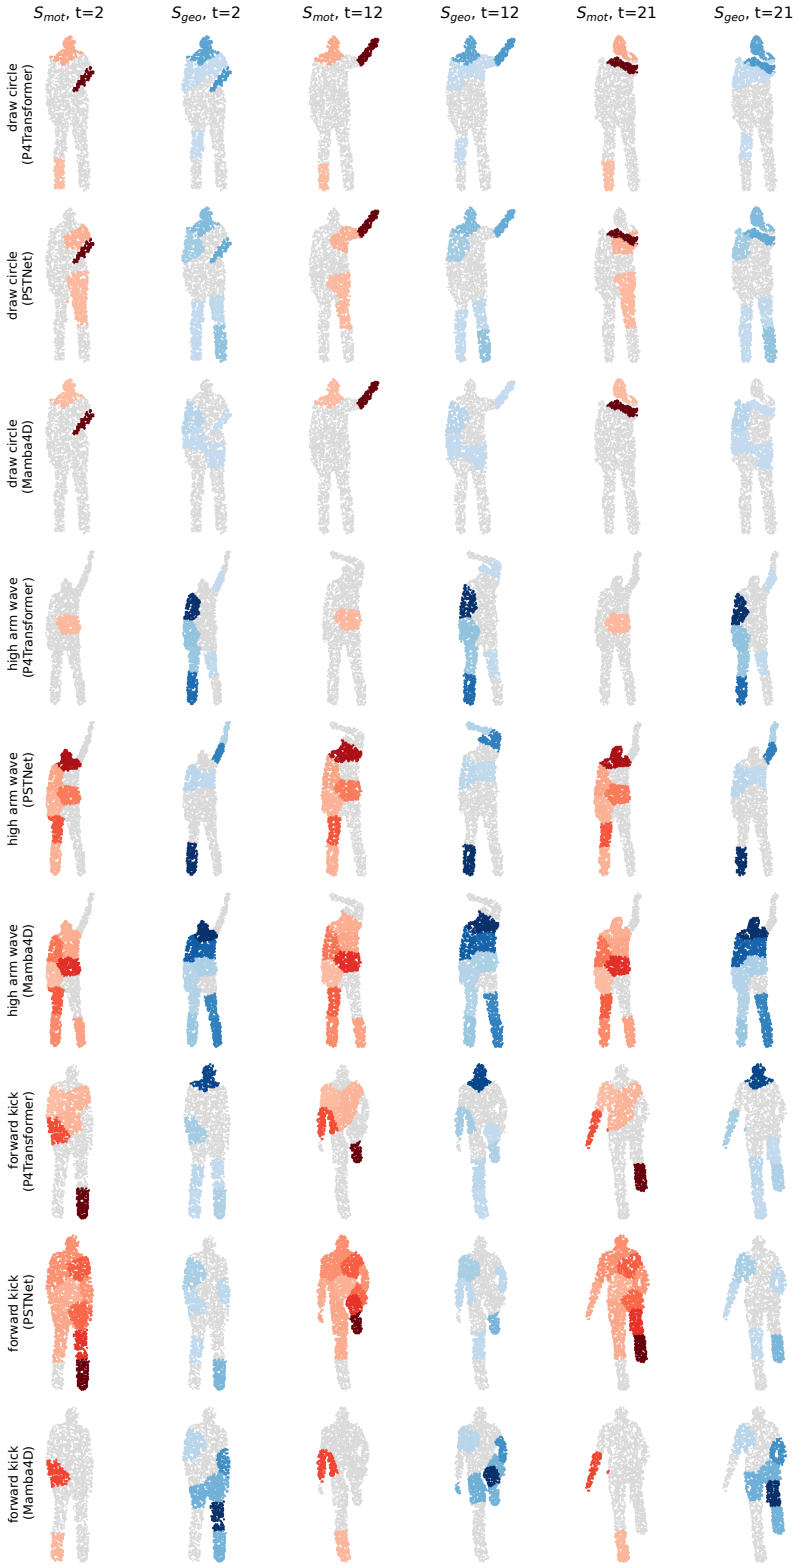

Figure 6: S_{mot} (red) and S_{geo} (blue) on real MSR-Action3D test clips at three instants; rows: three actions (*draw circle*, *high arm wave*, *forward kick*) $\times$ the three backbones, same code.

Action	mot. share	Del. KCA	Del. occl.	acc. %
high arm wave	0.61	0.151	0.265	100
horizontal arm wave	0.61	0.062	0.090	100
hammer	0.63	0.074	0.075	7
hand catch	0.68	0.244	0.359	80
forward punch	0.63	0.193	0.232	92
high throw	0.65	0.103	0.101	86
draw x	0.67	0.060	0.061	77
draw tick	0.72	0.085	0.144	93
draw circle	0.64	0.042	0.045	80
hand clap	0.73	0.071	0.084	100
two hand wave	0.66	0.080	0.078	100
side-boxing	0.69	0.280	0.310	71
bend	0.73	0.325	0.291	93
forward kick	0.80	0.494	0.698	100
side kick	0.71	0.209	0.437	100
jogging	0.54	0.422	0.626	100
tennis swing	0.74	0.078	0.097	73
tennis serve	0.75	0.163	0.176	87
golf swing	0.61	0.261	0.223	100
pick up & throw	0.69	0.231	0.331	100

Table 14: Per action (P4Transformer, all test clips of each action): motion share $S^+_{\mathrm{mot}}/(S^+_{\mathrm{mot}} + S^+_{\mathrm{geo}})$ of KCA (context) over the whole body, deletion AUC of KCA (context) and of region occlusion, and the model's accuracy on these clips.