arXiv:2611.02327v1 [cs.CV] 2 Nov 2026

Findings-aware Language-Image Pretraining for CT Volume Interpretation

Takuro Shimaya, Junya Inoh, Masahiro Saiko
Biometrics Research Labs, NEC corporation, Kawasaki, Japan
takuro-shimaya@nec.com

Abstract

Computed tomography (CT) provides diverse patient information and various clinical tasks, for which numerous deep learning models have been investigated. Recent years have seen the development of vision-language foundation models that can comprehensively extract clinical features, driven by representation learning that combines CT volumes with radiology reports. Existing approaches typically adopt Contrastive Language-Image Pre-training (CLIP)-based representation learning, where pairs from different cases are, in principle, treated as negative pairs such that the volume and report embeddings are pushed apart. In the biomedical imaging, however, treating different cases as negative pairs is not always appropriate, particularly when clinical similarity exists. In this study, we propose a supervised representation learning framework grounded in clinical knowledge from language foundation models, rather than the weak supervision used in CLIP. Specifically, we utilized a large language model to extract abnormal finding phrases from reports. We then explicitly constructed the target features that the image encoder should output, as the semantic representation of the extracted findings using a biomedical text embedding model. Our Findings-aware Language-Image Pre-training (FiLIP) outperformed weakly supervised approaches in multi-label zero-shot finding classification in both in-domain and external-domain evaluations, and also proved effective when fine-tuned for the downstream task of radiology report generation.

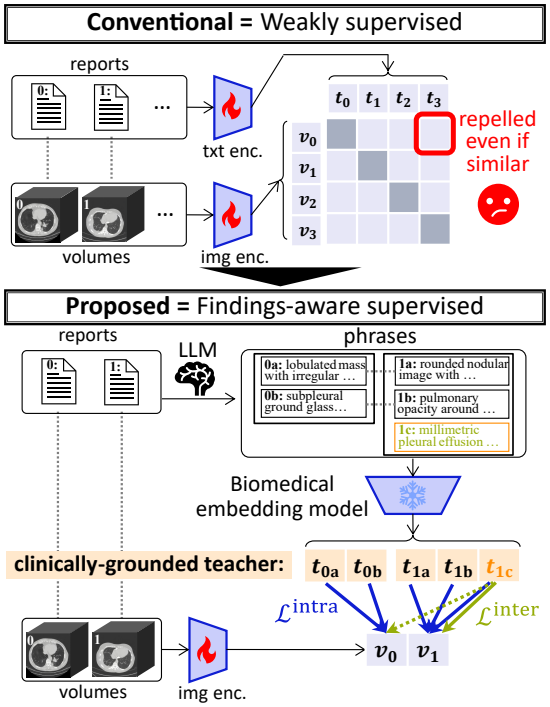

Figure 1. Schematic comparison of conventional methods and the proposed FiLIP. While contrastive learning may treat similar cases as negative pairs, our method directly supervises the model using text embeddings of extracted finding phrases from reports. We leverage a pretrained biomedical embedding model and an LLM as external clinical knowledge sources.

1. Introduction

Computed tomography (CT) is one of the most widely used medical imaging modalities, providing diverse clinical information across multiple organs. CT imaging contributes to a broad range of clinical tasks, *e.g.*, by supplying critical information for disease diagnosis, treatment planning, and digital twin construction [31, 40]. Accordingly, numerous deep learning models have been developed, and translation to clinical practice has become increasingly common, particularly in computer-aided diagnosis [13–15].

Conventional deep learning has relied on time-consuming manual annotations provided by medical experts. To reduce this burden, there has been growing interest in developing foundation models using large-scale medical imaging datasets. Unlike task-specific models, foundation models learn broad clinical representations and can support applications such as zero-shot diagnosis without manual annotation [18, 26]. Accordingly, large-scale pretraining has been explored with datasets across multiple modalities such as RadImageNet [20] and ROCOv2 [29]. In parallel, modality-specific vision-language foundation models

have been developed for chest X-ray [9] and histopathology [4]. This progress has been largely driven by Contrastive Language-Image Pretraining (CLIP) [27, 39], which has achieved substantial success in computer vision.

This trend also applies to CT. Three-dimensional volumes contain various abnormalities within a small fraction of slices, making manual annotation extremely burdensome. Recently, the release of CT-RATE [12], an open dataset that pairs CT volumes with corresponding radiology reports, has accelerated the development of vision-language foundation models for CT [5, 7, 12, 30, 33–35]. Conventional foundation models designed for two-dimensional slices have limited applicability because clinical practice often requires three-dimensional interpretation. To address this problem, CT-CLIP [12], which applies CLIP directly to CT volumes, achieved zero-shot classification of various findings and case retrieval. Further improvement has been achieved by fVLM [33], which performs contrastive learning on an organ-wise basis; SimCroP [35], which learns representations at the sentence level in reports; and ViSD-Boost [8] and CORIPLI [34], which combine vision-language alignment with additional loss functions such as autoencoding losses. While the above methods involve minor modifications, they are still trained primarily using CLIP-style vision-language contrastive learning.

While vision-language alignment has achieved substantial success, concerns remain when applying such approaches to medical imaging. In CLIP-style contrastive learning, different cases are typically treated as negative pairs, meaning their feature representations are pushed apart (Fig. 1). Although this assumption is basically unproblematic for general-purpose image datasets spanning broad domains, domain-specific datasets often include semantically similar samples, whose representations should not be forced apart. This issue is particularly relevant in medical imaging datasets, which often contain many clinically similar cases. The fVLM addresses this issue by treating normal cases as equivalent and assigning them a shared feature representation. The fVLM also estimates target similarities among abnormal cases via co-teaching, thereby drawing the representations of clinically similar cases closer together. However, it is not straightforward to accurately estimate similarity between cases sharing similar findings because the fVLM relies solely on volume-report correspondences without additional clinical knowledge.

This study proposes a vision-language alignment learning framework that avoids treating clinically similar cases as negative pairs. Rather than relying on a weakly supervised framework based on volume-report correspondences, we leverage the clinical knowledge embedded in pretrained language models to explicitly construct target feature representations for the image encoder. This novel supervised framework enables the model to learn clinically appropriate representations without treating clinically similar cases as negative pairs. Our Findings-aware Language-Image Pretraining (FiLIP) makes the following contributions:

1. We propose a method to define teacher representations for the image encoder by utilizing knowledge from a large language model (LLM) and a biomedical text embedding model.
2. We designed an intra-image focal similarity loss and a finding-conditioned inter-image contrastive loss, enabling phrase-level supervision without treating other cases as negatives solely based on case identity.
3. FiLIP outperformed conventional weak supervision in zero-shot finding classification on both in- and external-domain evaluations, and also exhibited superior performance on downstream radiology report generation.

2. Methods

2.1. FiLIP

The core idea is to prepare teacher representations for the image encoder, realizing novel supervised learning grounded in clinical knowledge from language foundation models. Specifically, we extracted phrases describing abnormal findings and instances, such as lesion and medical equipments, from radiology reports using an LLM, and obtained feature representations of extracted phrases using a biomedical embedding model (Fig. 1 and Fig. 2a).

Here, the finding phrases were extracted for each predefined organ as noun phrases with modifiers, as illustrated in Fig. 1 and Fig. 3. During extraction, we removed non-visual descriptions such as temporal changes and patients' clinical background, retaining only visualized instances that the image encoder can recognize. Besides, expressions such as "Nodules are not observed" were still treated as a confirmed-absent (negative) phrase "nodules." Because vision-language models have difficulty learning negation [2], the extracted phrases do not include explicit negation tokens such as "no" and "not". In addition, disease-related phrases such as lung cancer may also be described with expressions of uncertainty, *e.g.*, "suspected", when the radiologist identifies visual characteristics suggestive of the disease before a definitive diagnosis is established, for example by biopsy. This reflects the fact that visual features alone are not sufficient to establish a definitive diagnosis. In this study, phrases described with such expressions of suspicion were all treated as present (positive). The obtained phrases thereby can include diverse modifier sequences as shown in the right bottom of Fig. 3. Each phrase was vectorized using a frozen biomedical text embedding model and used as a teacher representation for the image encoder.

The set of embedded phrases for each predefined organ $o \in \{\mathrm{lung},\ \mathrm{heart},\ \mathrm{esophagus},\ \mathrm{aorta}\}$ in case i is denoted by $\{\boldsymbol{t}_{i,l}^{o,a}\}_{l=0}^{N_i^{o,a}-1}$, where a indicates the appearance (present:

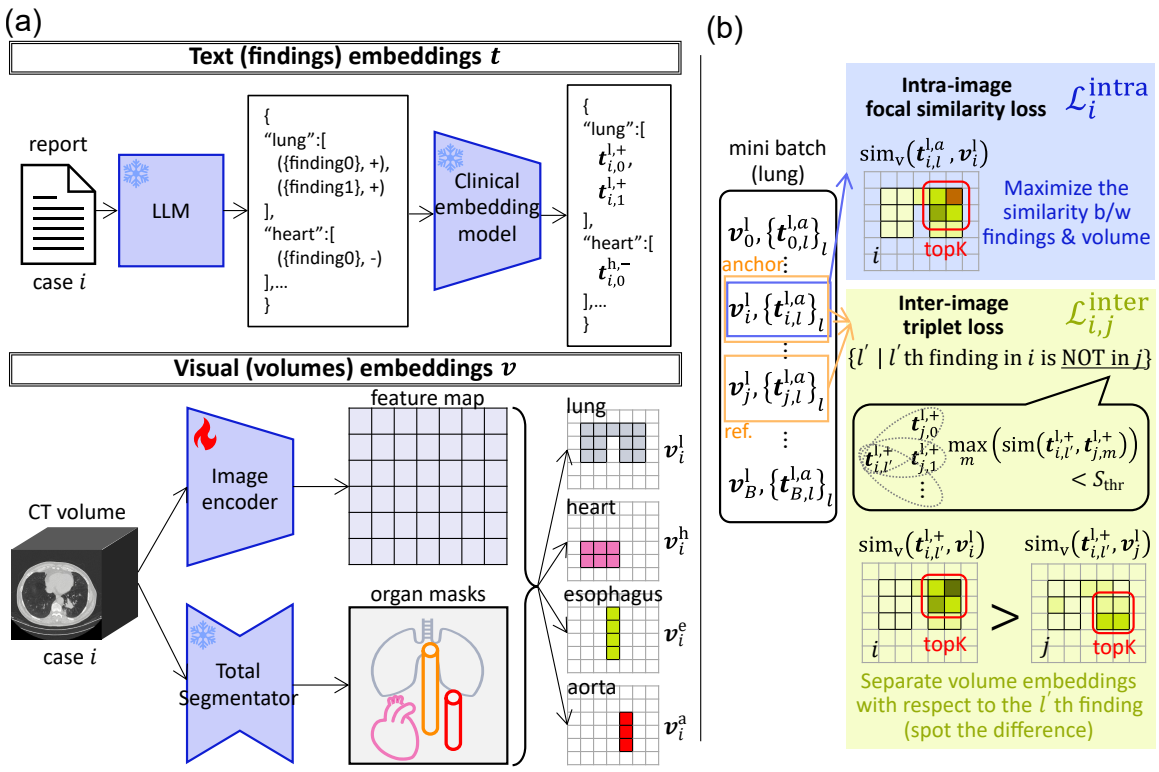

Figure 2. Overview of the proposed strategy. (a) Feature extraction processes. See also Fig. 3. (b) Loss function calculations. See Methods for details.

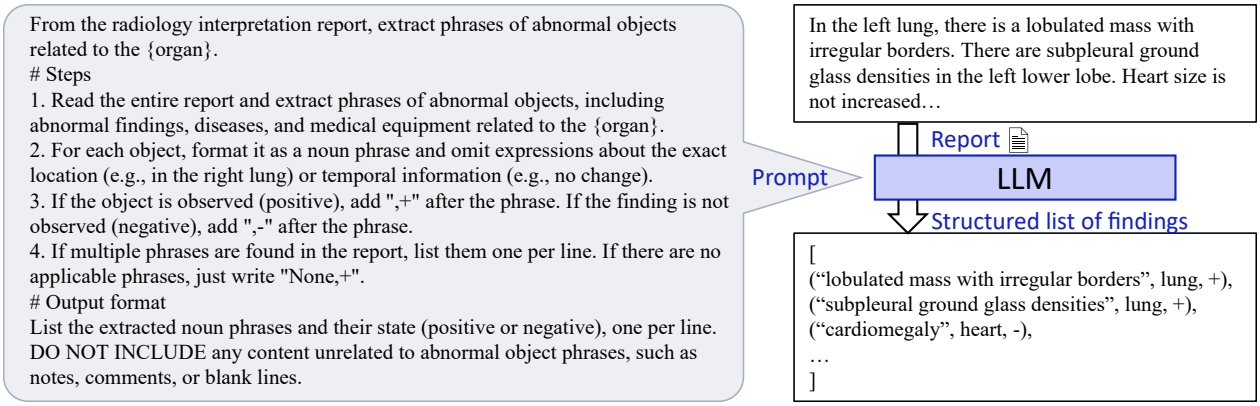

Figure 3. Finding phrases extraction via an LLM.

$+$ or absent: $-$), and $N_i^{o,a}$ is the number of present/absent findings, indexed by l. To push apart features corresponding to background regions from those of abnormal findings, we explicitly define the background embedding shared across all cases and organs:

$$\boldsymbol{t}_{\mathrm{bg}} = -\mathrm{Normal}\left(\sum_{(i,o,a,l)\in \mathrm{Train}} \boldsymbol{t}_{i,l}^{o,a}\right), \qquad (1)$$

i.e., the vector farthest from most abnormal finding representations where $\mathrm{Normal}(\boldsymbol{v}) = \boldsymbol{v}/|\boldsymbol{v}|$. $\boldsymbol{t}_{\mathrm{bg}}$ was then appended to all $\{\boldsymbol{t}_{i,l}^{o,+}\}_{l=0}^{N_i^{o,+}-1}$ in the training data.

In parallel, we computed organ-wise visual features from the three-dimensional CT volume. This process is identical to that proposed by fVLM. We cropped the feature map obtained from the three-dimensional image encoder using the mask of organ o generated by TotalSegmentator [36], yielding an organ-specific volume feature map $\boldsymbol{v}_i^o$ for case i.

We now design novel vision-language alignment losses: an intra-image focal similarity loss and an inter-image triplet loss (Fig. 2b). We hypothesize that a clinically grounded image encoder can be obtained solely by training it to reproduce embeddings from a frozen language model. The simplest way is to design a loss such that a volume's visual feature representation reproduces the list of finding embeddings $\{\boldsymbol{t}_{i,l}^{o,a}\}_{l=0}^{N_i^{o,a}-1}$ extracted from the corresponding report. For organ o in case i, we define the following intra-image focal similarity loss:

$$\mathcal{L}_{i,o}^{\text{intra}} = -\frac{1}{N_i^{o,+}+N_i^{o,-}}\left(\sum_{l\in+}^{N_i^{o,+}}(1-p_{i,l}^{o,+})^{\gamma}\ln p_{i,l}^{o,+} + \sum_{l\in-}^{N_i^{o,-}}(p_{i,l}^{o,-})^{\gamma}\ln(1-p_{i,l}^{o,-})\right), \tag{2}$$

where $p_{i,l}^{o,a}$ is the prediction score based on the similarity between the text embedding and the volume feature. A higher $p_{i,l}^{o,a}$ indicates that the volume feature likely contains the l-th finding, which is formulated as

$$p_{i,l}^{o,a} = \sigma\left(\text{TopK-AvgPool}(\text{sim}_{\text{v}}(\boldsymbol{t}_{i,l}^{o,a}, \boldsymbol{v}_i^{o}), K)\right), \tag{3}$$

$$\sigma(x) = (1+\exp(-\alpha x+\beta))^{-1}. \tag{4}$$

The score $p_{i,l}^{o,a}$ reflects the relevance of the l-th finding to the volume, which is computed via the sigmoid function $\sigma(\cdot)$ and voxel-wise cosine similarities $\text{sim}_{\text{v}}(\cdot)$ with learnable parameters α and β, following ref. [37]. The voxel-wise cosine similarity between the finding embedding and each voxel-wise feature is aggregated via the TopK-AvgPool function, where the top $K\times 100$ % of elements are selected per each organ and averaged. This TopK average pooling is motivated because voxels relevant to a target finding are localized, and global average pooling may not be suitable. The loss $\mathcal{L}_{i,o}^{\text{intra}}$ encourages volume features $\boldsymbol{v}_i^o$ to be close to the present finding embedding $\boldsymbol{t}_{i,l}^{o,+}$, while pushing them away from the absent findings $\boldsymbol{t}_{i,l}^{o,-}$. In addition, because the target representation $\boldsymbol{t}_{i,l}^{o,a}$ may be non-uniformly distributed, we regard this as "class imbalance" and thus introduce the coefficients from the focal loss controlled by γ [19].

We additionally introduce an inter-image triplet loss that trains the image encoder in a "spot-the-difference" manner (Fig. 2b). The model trained solely with the intra-image loss may confuse visual features of coexisting findings in each single volume. To address this concern, FiLIP provides the additional opportunities for the image encoder to distinguish between findings that frequently co-occur by comparing two volumes. To guide spot-the-difference training, we construct a pair consisting of an anchor volume i and a reference volume j, and define the set of unique findings

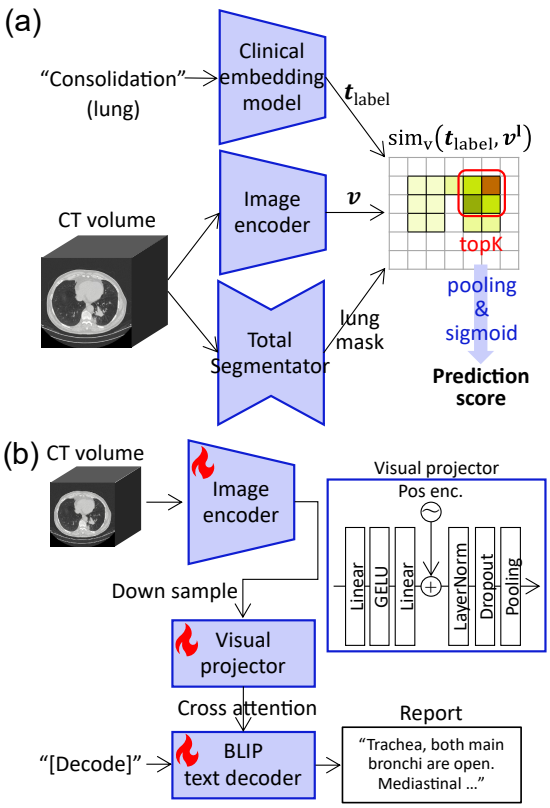

Figure 4. Performance evaluation methods. (a) Zero-shot classification of finding presence. (b) Fine-tuning for radiology report generation as a downstream task. See Methods for details.

present only in the anchor, $U_{i,j}^o$, as:

$$U_{i,j}^o = \left\{ l \,\middle|\, \max_m \left(\text{sim}(\boldsymbol{t}_{i,l}^{o,+}, \boldsymbol{t}_{j,m}^{o,+})\right) < S_{\text{thr}} \right\}, \tag{5}$$

where $\text{sim}(\cdot)$ denotes scalar cosine similarity between embeddings, and S_{thr} is a hyperparameter that serves as a threshold controlling how similar two finding phrases are to be regarded as identical. For example, in Fig. 1 bottom, "effusion" is included in U. We then define the following binary InfoNCE-style triplet loss such that the anchor volume feature is pulled toward the embeddings of findings in $U_{i,j}^o$, while the reference volume feature is pushed away:

$$\mathcal{L}_{i,j,o}^{\text{inter}} = -\frac{1}{|U_{i,j}^o|} \sum_{l\in U_{i,j}^o} \ln\left(\frac{\exp(s(\boldsymbol{t}_{i,l}^{o,+}, \boldsymbol{v}_i^o)/\tau)}{\exp(s(\boldsymbol{t}_{i,l}^{o,+}, \boldsymbol{v}_i^o)/\tau)+\exp(s(\boldsymbol{t}_{i,l}^{o,+}, \boldsymbol{v}_j^o)/\tau)}\right), \tag{6}$$

$$s(\boldsymbol{t}, \boldsymbol{v}) = \text{TopK-AvgPool}(\text{sim}_{\text{v}}(\boldsymbol{t}, \boldsymbol{v})), \tag{7}$$

with a learnable temperature τ. This inter-image loss emphasizes semantic differences between volumes and, in par-

ticular, promotes the discrimination between findings that coexist within a single image.

We finally use the combination of intra- and inter-image losses as the total loss for a mini-batch B:

$$\mathcal{L} = \sum_o \left(\frac{1}{|B|} \sum_{i \in B} \mathcal{L}_{i,o}^{\text{intra}} + \frac{1}{|B|(|B|-1)} \sum_{i \in B} \sum_{j \in B, j \neq i} \mathcal{L}_{i,j,o}^{\text{inter}} \right). \quad (8)$$

taking the sum over all specified organs. Note that we train only the image encoder; the language models are kept frozen throughout this study.

As noted earlier, the objective of FiLIP is to leverage the external medical knowledge encoded in language models for training the image encoder (Fig. 1). Conventional weakly supervised learning based on report-volume correspondence, as in CLIP, relies on the strong assumption that different cases are clinically distinct. However, this assumption can be substantially violated in domain-specific datasets such as medical images. Our intra-image focal similarity loss aligns local features with predefined teacher representations aiming to train the vision-language model in a supervised rather than weakly supervised manner. The inter-image triplet loss further separates cases only along findings on which they differ, thereby replacing case-identity-based learning with clinically grounded supervision. The time required for inter-image triplet loss calculation and the number of phrases considered per anchor image, *i.e.*, the effective number of negative examples per anchor, are provided in the supplementary text.

2.2. Datasets, implementations, and experiments

Following fVLM [33] and related studies, we trained the image encoder and performed in-domain validation using an publicly available chest CT dataset, CT-RATE [12]. In accordance with the official data split, we used 24,114 unique cases for training and 1,557 for validation. For external-domain testing, we used 3,630 cases from Rad-ChestCT [10]. Preprocessing, including resolution downsampling, also followed fVLM.

We adopted a Vision Transformer (ViT) [1] as the image encoder, following fVLM. We used GPT-4.1 [22] as the LLM for structuring reports, and obtained 25,315 unique free-form finding phrases exclusively from the training reports. We used BioLORD-2023 [28] as the biomedical text embedding model, yielding 768-dims embeddings. We expected BioLORD-2023 to provide broader specialized biomedical knowledge than both the general-purpose embedding model text-embedding-3-large [21] and CXR-BERT [6], a chest X-ray-specific model frequently used as pretrained weights for medical vision-language models; this expectation was subsequently verified through ablation studies. For training, we adopted the hyperparameters from fVLM official implementation [38], such as the learning rate and batch size. The FiLIP-specific learnable parameters α, β, and τ were initialized to 4, 2, and 0.07, respectively. The hyperparameters K, γ, and S_{thr} were set to 0.75, 2, and 0.6, respectively, and we evaluated performance dependence on them via ablation studies.

We evaluated the performance of the proposed method by a zero-shot classification of finding presence in each CT volume. Specifically, for each of the 16 findings defined in the fVLM evaluation protocol [33], we calculated the area under the ROC curve (AUC). During inference, we computed embeddings of each label term using BioLORD-2023 and performed classification by taking the similarity with the volume feature map of the organ where the finding occurs; the score p in Eq. (3) was used as the probability of the finding being present (see Fig. 4a for the case of the consolidation in lungs). To demonstrate the advantage of FiLIP over existing loss functions for vision-language alignment, we compared the zero-shot finding classification performance against recently proposed CT-CLIP [12], fVLM [33], and SimCroP [35], which have similar model architectures and are therefore suitable for comparison. The baseline methods were re-evaluated using their published weights trained on CT-RATE, with confidence intervals (CI) computed for fair comparison. For SimCroP, note that because a zero-shot inference method is not officially provided, the cosine similarity between the globally averaged volume feature and the text embedding of the prompt "{finding} is present." was leveraged as the score, following the training strategy reported in ref. [35]. The classification performance was evaluated by averaging the AUC across all finding classes. The 95% CI of the mean AUC were determined via a non-parametric bootstrap procedure with 1,000 resamples, and statistical significance was assessed at a 5% significance level using two-sided DeLong tests. Besides, we also compared the AUCs by CT-CLIP variants fine-tuned with class labels (VocabFine and ClassFine) and Merlin [5] evaluated in the same setup and reported in ref.[8]. Note that these models are not directly comparable because they were trained with additional data beyond radiology reports.

To verify that the FiLIP-trained model comprehensively captures clinical features beyond finding presence, we also evaluated its performance on a downstream task. Specifically, we trained an end-to-end model for radiology report generation using CT-RATE. The pretrained image encoder was connected to the BLIP text decoder [16] through cross-attention mediated by a visual projector, and the entire network was fine-tuned (Fig. 4b, see the supplementary text for details). Performance was evaluated using lexical metrics (BLEU [25], METEOR [3], and ROUGE [17])

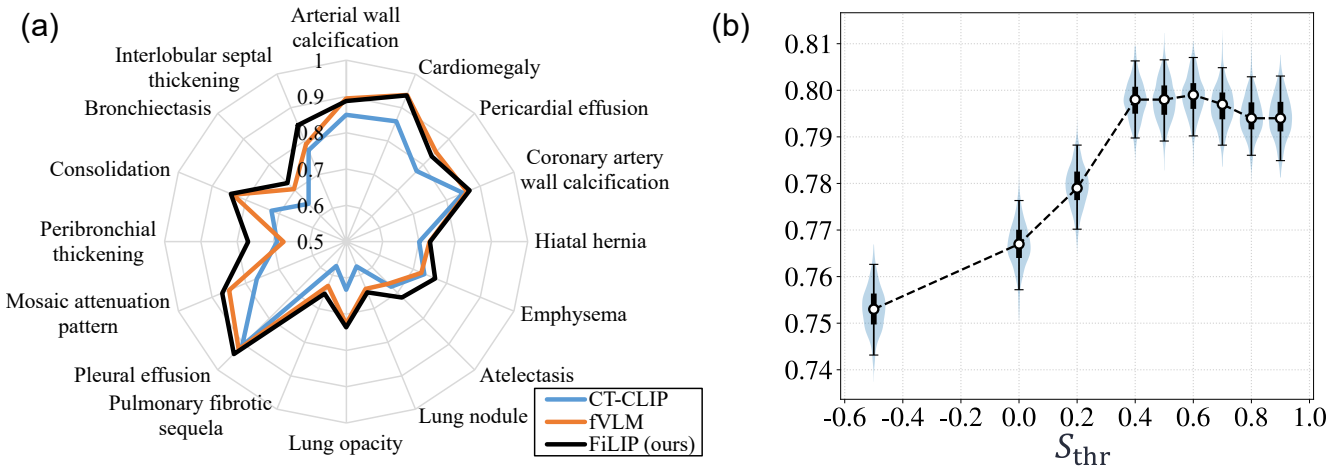

Figure 5. Zero-shot finding classification performance for CT-RATE. (a) AUC for each finding. See the supplementary text for correspondence between labels and organs. (b) Dependency of the mean AUC on the similarity threshold S_{thr} for computing $\mathcal{L}^{\text{inter}}$. Violin plots show the AUC distributions, with bootstrap-estimated 95% confidence intervals and interquartile ranges overlaid.

Table 1. The mean AUC for zero-shot finding classification. 95% confidence intervals are shown in the brackets. † indicates models requiring additional data beyond radiology reports for training, and values are from ref.[8]. FiLIP significantly outperformed fVLM on both evaluation datasets (see the main and supplementary text for details).

Model	Evaluation dataset	
	CT-RATE	Rad-ChestCT
CT-CLIP [12]	0.733 [0.727-0.740]	0.632 [0.625-0.640]
(VocabFine)†	0.760	0.657
(ClassFine)†	0.761	0.647
Merlin† [5]	0.728	0.644
fVLM [33]	0.778 [0.770-0.786]	0.680 [0.674-0.686]
SimCroP [35]	0.620 [0.613-0.633]	0.509 [0.505-0.521]
FiLIP (ours)	**0.799 [0.791-0.808]**	**0.690 [0.683-0.697]**

Table 2. Ablation studies on losses evaluated using CT-RATE.

$\mathcal{L}^{\text{intra}}$	t_{bg}	TopK	Focal	$\mathcal{L}^{\text{inter}}$	AUC
✓					0.761
✓	✓				0.764
✓	✓	✓			0.768
✓	✓		✓		0.774
✓	✓	✓	✓		0.779
				✓	0.788
✓	✓			✓	0.797
✓	✓	✓	✓	✓	**0.799**

and the GREEN score [24], a LLM-based metric that accounts for clinical semantic similarity between texts. In addition to overall score, the GREEN analyzes the following error/match counts: false reporting of a finding in the candidate report, omission of a finding present in the reference report, misidentification of a finding's anatomic location or position, misassessment of a finding's severity, inclusion of a comparison not described in the reference report, omission of a comparison describing interval change from a prior study, and matched findings. Because the report generation model in this study analyzes a single CT volume, evaluation categories related to longitudinal comparison were not considered. For each metric, uncertainty was estimated using the standard error. For the BLEU-4 and the GREEN including subcategories, statistical significance of differences between FiLIP and prior methods was assessed at a 5% significance level using two-sided Wilcoxon signed-rank tests. All p-values in this study were adjusted for multiple comparisons according to the Holm correction.

Table 3. Ablation studies on foundation models evaluated using CT-RATE. # 4 is the default setting in this study. Statistical significance was observed between # 1 vs # 4, and # 2 vs # 4, while no significant difference was found between # 3 vs # 4 (see the main and supplementary text for details).

	LLM	embedding model	AUC
# 1	GPT-4.1	CXR-BERT	0.744 [0.734-0.754]
# 2	GPT-4.1	text-embedding-3-large	0.787 [0.778-0.795]
# 3	gpt-oss-120B	BioLORD-2023	**0.796 [0.787-0.804]**
# 4	GPT-4.1	BioLORD-2023	**0.799 [0.791-0.808]**

3. Results

The mean AUC of zero-shot finding classification is reported in Table 1. Our model outperformed existing models trained with conventional CLIP-style methods on both the in-domain validation and the external-domain test for zero-shot finding classification. While CIs are overlapped for the external test, the improvement was shown to be statistically significant for both datasets by DeLong tests (see the sup-

Table 4. Results on the downstream task of report generation using CT-RATE. Standard errors are shown in the brackets. For FiLIP scores, *1, *2, and *3 indicate statistically significant differences from ImageNet, fVLM, and SimCroP, respectively. *0 indicates no significant difference from any model. Metrics without an significance marker were not statistically tested. See p-values for the supplementary text.

(a) Evaluation using lexical metrics.

Model	BLEU-1	BLEU-2	BLEU-3	BLEU-4	METEOR	ROUGE-1	ROUGE-2	ROUGE-L
ImageNet	0.3794(38)	0.2852(38)	0.2258(37)	0.1771(34)	0.3903(34)	0.5816(38)	0.3832(57)	0.4483(53)
fVLM [33]	0.3765(38)	0.2824(39)	0.2239(37)	0.1759(35)	0.3889(35)	0.5797(40)	0.3823(58)	**0.4485(54)**
SimCroP [35]	0.3915(38)	0.2925(39)	**0.2307(39)**	**0.1810(37)**	**0.3975(37)**	0.5886(39)	0.3815(58)	0.4444(54)
FiLIP (ours)	**0.3938(34)**	**0.2934(36)**	0.2306(36)	0.1798(33)^{*1,2}	0.3974(33)	**0.5893(37)**	**0.3837(56)**	0.4472(52)

(b) Evaluation using the clinical metric, the GREEN score. The subcategories represent the counts of following subjects: g-1, false report of a finding in the candidate; g-2, missing a finding present in the reference; g-3, misidentification of a finding's anatomic location/position; g-4, misassessment of the severity of a finding; g-5, matched findings.

		subcategory				
Model	overall	g-1 (↓)	g-2 (↓)	g-3 (↓)	g-4 (↓)	g-5 (↑)
ImageNet	0.5416(67)	0.982(31)	**3.376(60)**	0.1019(85)	0.0679(68)	6.13(11)
fVLM [33]	0.5426(68)	0.969(29)	3.381(60)	0.0936(81)	0.0724(69)	6.23(11)
SimCroP [35]	0.5464(67)	0.883(31)	3.383(58)	0.1091(97)	0.0842(75)	6.19(11)
FiLIP (ours)	**0.5533(68)**^{*1,2}	**0.747(27)**^{*1,2,3}	3.387(57)^{*0}	**0.0641(65)**^{*1,2,3}	**0.0615(64)**^{*0}	**6.30(11)**^{*0}

plementary text for p-values). A per-finding comparison is also shown in Fig. 5a. Classification performance improved particularly for lung findings, and we did not observe substantial degradation in AUC for any specific finding. As noted in the original paper [35], SimCroP does not execute instance-level alignment for the entire volume and therefore showed relatively lower classification performance, while it demonstrated strong performance on a downstream task as described later.

The validity of the proposed method was further underpinned by ablation studies (Table 2 and Fig. 5b). We observed performance drop when removing $\boldsymbol{t}_{\mathrm{bg}}$, replacing the focal loss with a cross-entropy loss, or replacing TopK average pooling with global average pooling in $\mathcal{L}^{\mathrm{intra}}$. Remarkably, using $\mathcal{L}^{\mathrm{intra}}$ alone, i.e., using only correspondences within the same case, still achieved performance comparable to fVLM. We also found that training with only the inter-image triplet loss $\mathcal{L}^{\mathrm{inter}}$ outperformed fVLM. The best performance was realized by combining $\mathcal{L}^{\mathrm{intra}}$ and $\mathcal{L}^{\mathrm{inter}}$. When varying the similarity threshold S_{thr} to determine the set of unique cases U, we observed peak performance around $S_{\mathrm{thr}} \approx 0.6$. This is consistent with the following intuition: If S_{thr} is set too high or too low, common findings between anchor and reference volumes cannot be identified reliably, which in turn hampers the spot-the-difference training. The plateau near the peak further suggests that our method is not highly sensitive to the choice of the hyperparameter S_{thr}. Moreover, we observed that the performance was significantly affected by the choice of the text embedding model (Table 3), indicating that FiLIP successfully utilized the clinical knowledge embedded in BioLORD. We also confirm that the performance was not substantially decayed when replacing GPT-4.1 with gpt-oss-120B [23], showing that models can be trained in private environments as well (see p-values for the supplementary text).

Table 4 summarizes the fine-tuning results for radiology report generation with pretrained image encoders. We compared a ViT with attention weights pretrained on ImageNet with those trained with fVLM, SimCroP, and FiLIP. While FiLIP exhibited no significant difference to SimCroP regarding conventional lexical metrics, significant difference was observed regarding the GREEN, which measures the clinical semantic similarity beyond word-level overlap and has been reported to correlate well with radiologists' assessments [24, 32] (see p-values for the supplementary text). Across the subcategories, FiLIP significantly reduced false reporting of findings (g-1) but also improved the localization of findings (g-3) compared to SimCroP. The overall GREEN score and BLEU-4 was significantly higher for FiLIP compared to fVLM. These results suggest that the FiLIP-trained encoder learned to capture not only features required for classification of finding presence but also diverse clinical information to reproduce radiology reports.

To visually assess the improvement in the g-3 score, we observed the zero-shot capability of finding localization by the FiLIP-trained model. For models trained with fVLM and FiLIP, we visualized pseudo-color maps of the cosine similarity between voxel-wise features and the each finding embedding, normalized using the minimum and maximum values over the entire volume (Fig. 6). Consolidation and pleural effusion were recognized relatively well by both models (Fig. 5a). In contrast, for emphysema, the fVLM-

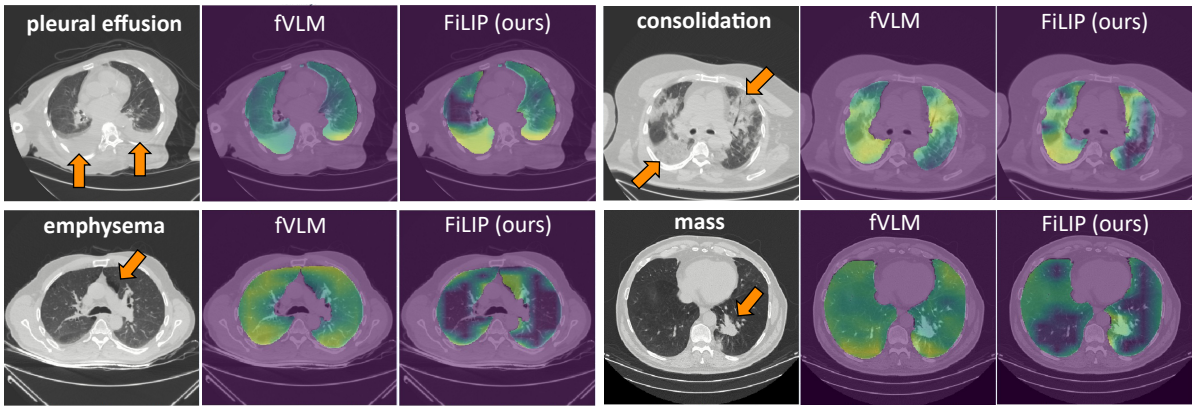

Figure 6. Visualization of voxel-wise similarity map between target finding embeddings and volume features. The arrows indicate the position of each finding. The brighter color corresponds to the higher similarity.

trained model produced an almost uniform similarity map, whereas FiLIP yielded an anatomically plausible map. Our model also exhibited a spatially coherent map correlated with mass, which was not included in the zero-shot finding classification evaluation. Note that quantitative evaluation is required in future studies to prove superior performance in zero-shot finding localization.

4. Discussions and Conclusions

From a clinical perspective, the following limitations remain. Although we demonstrated improved zero-shot finding classification, more comprehensive zero-shot prediction remains to be investigated, regarding information beyond findings that can be omitted during LLM-based extraction. Even for finding classification, the performance gain for lung nodules, one of the most critical findings relevant to cancer, is relatively modest. Consistently, the similarity map for the cancer-related mass (Fig. 6) exhibits a false positive near the slightly brighter region in the right lung (left in the image). To address this issue, since cancer-related findings are often highly localized, the optimal K may also be smaller for such findings. It may be possible to adaptively select visual tokens for each finding, as in ref. [11]. Regarding the downstream report generation, no significant improvement from SimCroP was observed in the overall GREEN and several subcategories. Severity assessment may require more expertised clinical knowledge. In addition, the reference reports used for evaluation include information missed in the chest CT volume, such as head, neck, and abdomen findings and patient background, which may make score differences less detectable. Future work should consider preprocessing the reference reports and obtaining more precise evaluation by radiologists.

In conclusion, we propose supervised vision-language alignment, motivating a rethinking of the common strategy for multimodal medical representation learning. By leveraging the knowledge encoded in the LLM and the biomedical text embedding model to define teacher representations for the image encoder, our approach outperforms prior vision-language alignment methods for zero-shot finding classification and downstream radiology report generation. Experiments on downstream report generation suggest that the FiLIP-trained encoder captures features related not only to the presence of findings but also to their location. In other words, FiLIP has the potential to learn comprehensive clinical features contained in CT volumes. Regarding the hyperparameters, since the value $\gamma = 2$ employed in the experiments is the standard choice for focal loss, only two hyperparameters, K and S_{thr}, need to be tuned in practice. The ablation study showed that the performance variations associated with K, *i.e.*, the TopK pooling parameter, and S_{thr} were limited, suggesting that the burden of hyperparameter search is not substantial.

The performance of the FiLIP-trained model may further improve as language models continue to advance. As the biomedical embedding models improve, image encoders may learn to capture increasingly fine-grained clinical feature representations. For report structuring, improvements in phrase extraction are also expected as LLMs evolve. Further data augmentation is also possible, *e.g.*, by using external medical knowledge to supplement missing relationships among phrases. In addition, FiLIP can be combined with any network architecture and auxiliary losses, such as image reconstruction and report generation objectives, as explored for ViSD-Boost [8], SimCroP [35], and CORIPLI [34]. Questions remain regarding how to construct the list of teacher representations, including how to define $\boldsymbol{t}_{\text{bg}}$ (Eq. (1)) and incorporate negative phrases not mentioned in the report into training. FiLIP has the potential for scalability and transferability to other organs and modalities, as performance decay is modest across hyperparameter choices, possibly encouraging broader investigation and adoption in domain-specific vision-language alignment beyond the medical field.

References

[1] Alexey Dosovitskiy and Lucas Beyer and Alexander Kolesnikov and Dirk Weissenborn and Xiaohua Zhai and Thomas Unterthiner and Mostafa Dehghani and Matthias Minderer and Georg Heigold and Sylvain Gelly and Jakob Uszkoreit and Neil Houlsby. An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. In *ICLR*, 2021. 5

[2] Kumail Alhamoud, Shaden Alshammari, Yonglong Tian, Guohao Li, Philip H.S. Torr, Yoon Kim, and Marzyeh Ghassemi. Vision-Language Models Do Not Understand Negation. In *CVPR*, pages 29612–29622, 2025. 2

[3] Satanjeev Banerjee and Alon Lavie. METEOR: An Automatic Metric for MT Evaluation with Improved Correlation with Human Judgments. In *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, pages 65–72, 2005. 5

[4] Mohsin Bilal, Aadam, Manahil Raza, Youssef Altherwy, Anas Alsuhaibani, Abdulrahman Abduljabbar, Fahdah Almarshad, Paul Golding, and Nasir Rajpoot. Foundation Models in Computational Pathology: A Review of Challenges, Opportunities, and Impact. *arXiv preprint, arXiv:2502.08333*, 2025. 2

[5] Louis Blankemeier, Ashwin Kumar, Joseph Paul Cohen, Jiaming Liu, Longchao Liu, Dave Van Veen, Syed Jamal Safdar Gardezi, Hongkun Yu, Magdalini Paschali, Zhihong Chen, Jean-Benoit Delbrouck, Eduardo Reis, Robbie Holland, Cesar Truyts, Christian Bluethgen, Yufu Wu, Long Lian, Malte Engmann Kjeldskov Jensen, Sophie Ostmeier, Maya Varma, Jeya Maria Jose Valanarasu, Zhongnan Fang, Zepeng Huo, Zaid Nabulsi, Diego Ardila, Wei-Hung Weng, Edson Amaro Junior, Neera Ahuja, Jason Fries, Nigam H. Shah, Greg Zaharchuk, Marc Willis, Adam Yala, Andrew Johnston, Robert D. Boutin, Andrew Wentland, Curtis P. Langlotz, Jason Hom, Sergios Gatidis, and Akshay S. Chaudhari. Merlin: a computed tomography vision–language foundation model and dataset. *Nature*, 652: 1318–1328, 2024. 2, 5, 6

[6] Benedikt Boecking, Naoto Usuyama, Shruthi Bannur, Daniel C. Castro, Anton Schwaighofer, Stephanie Hyland, Maria Wetscherek, Tristan Naumann, Aditya Nori, Javier Alvarez-Valle, Hoifung Poon, and Ozan Oktay. Making the Most of Text Semantics to Improve Biomedical Vision–Language Processing. In *ECCV*, 2022. 5

[7] Weiwei Cao, Jianpeng Zhang, Yingda Xia, Tony C. W. Mok, Zi Li, Xianghua Ye, Le Lu, Jian Zheng, Yuxing Tang, and Ling Zhang. Bootstrapping Chest CT Image Understanding by Distilling Knowledge from X-Ray Expert Models. In *CVPR*, pages 11238–11247, 2024. 2

[8] Weiwei Cao, Jianpeng Zhang, Zhongyi Shui, Sinuo Wang, Zeli Chen, Xi Li, Le Lu, Xianghua Ye, Tingbo Liang, Qi Zhang, and Ling Zhang. Boosting Vision Semantic Density with Anatomy Normality Modeling for Medical Vision-language Pre-training. In *ICCV*, pages 23041–23050, 2025. 2, 5, 6, 8

[9] Zhihong Chen, Maya Varma, Jean-Benoit Delbrouck, Magdalini Paschali, Louis Blankemeier, Dave Van Veen, Jeya Maria Jose Valanarasu, Alaa Youssef, Joseph Paul Cohen, Eduardo Pontes Reis, Emily Tsai, Andrew Johnston, Cameron Olsen, Tanishq Mathew Abraham, Sergios Gatidis, Akshay S Chaudhari, and Curtis Langlotz. CheXagent: Towards a Foundation Model for Chest X-Ray Interpretation. In *AAAI*, 2024. 2

[10] Rachel Lea Draelos, David Dov, Maciej A. Mazurowski, Joseph Y. Lo, Ricardo Henao, Geoffrey D. Rubin, and Lawrence Carin. Machine-learning-based multiple abnormality prediction with large-scale chest computed tomography volumes. *Med. Image Anal.*, 67:101857, 2021. 5

[11] Satya Krishna Gorti, Noel Vouitsis, Junwei Ma, Keyvan Golestan, Maksims Volkovs, Animesh Garg, and Guangwei Yu. X-Pool: Cross-Modal Language-Video Attention for Text-Video Retrieval. In *CVPR*, pages 4996–5005, 2022. 8

[12] Ibrahim E. Hamamci, Sezgin Er, Chenyu Wang, Furkan Almas, Ayse G. Simsek, Sevval Nil Esirgun, Irem Dogan, Omer Faruk Durugol, Benjamin Hou, Suprosanna Shit, Weicheng Dai, Murong Xu, Hadrien Reynaud, Muhammed F. Dasdelen, Bastian Wittmann, Tamaz Amiranashvili, Enis Simsar, Mehmet Simsar, Emine B. Erdemir, Abdullah Alanbay, Anjany Sekuboyina, Berkan Lafci, Ahmet Kaplan, Zhiyong Lu, Malgorzata Polacin, Bernhard Kainz, Christian Bluethgen, Kayhan Batmanghelich, Mehmet K. Ozdemir, and Bjoern Menze. Generalist foundation models from a multimodal dataset for 3D computed tomography. *Nat. Biomed. Eng.*, 2026. 2, 5, 6

[13] Cecile E.J. Kleber, Ramez Karius, Lucas E. Naessens, Coen O. Van Toledo, Jochen A. C. van Osch, Martijn F. Boomsma, Jan W.T. Heemskerk, and Aart J. van der Molen. Advancements in supervised deep learning for metal artifact reduction in computed tomography: A systematic review. *Eur. J. Radiol.*, 181:111732, 2024. 1

[14] Lennart R. Koetzier, Domenico Mastrodicasa, Timothy P. Szczykutowicz, Niels R. van der Werf, Adam S. Wang, Veit Sandfort, Aart J. van der Molen, Dominik Fleischmann, and Martin J. Willemink. Deep Learning Image Reconstruction for CT: Technical Principles and Clinical Prospects. *Radiology*, 306(3):e221257, 2023.

[15] Lennart R. Koetzier, Domenico Mastrodicasa, Timothy P. Szczykutowicz, Niels R. van der Werf, Adam S. Wang, Veit Sandfort, Aart J. van der Molen, Dominik Fleischmann, and Martin J. Willemink. Deep learning in pulmonary nodule detection and segmentation: a systematic review. *Eur. Radiol.*, 35:255–266, 2025. 1

[16] Junnan Li, Dongxu Li, Caiming Xiong, and Steven Hoi. BLIP: Bootstrapping Language-Image Pre-training for Unified Vision-Language Understanding and Generation. In *Proc. Mach. Learn. Res.*, pages 12888–12900, 2022. 5

[17] Chin-Yew Lin. ROUGE: A Package for Automatic Evaluation of Summaries. In *Text Summarization Branches Out*, pages 74–81, 2004. 5

[18] Haoneng Lin, Cheng Xu, and Jing Qin. Taming Vision-Language Models for Medical Image Analysis: A Comprehensive Review. *arXiv preprint, arXiv:2506.18378*, 2025. 1

[19] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and

Piotr Dollár. Focal Loss for Dense Object Detection. *IEEE Trans. Pattern Anal. Mach. Intell.*, 42(2):318–327, 2020. 4

[20] Xueyan Mei, Zelong Liu, Philip M. Robson, Brett Marinelli, Mingqian Huang, Amish Doshi, Adam Jacobi, Chendi Cao, Katherine E. Link, Thomas Yang, Ying Wang, Hayit Greenspan, Timothy Deyer, Zahi A. Fayad, and Yang Yang. RadImageNet: An Open Radiologic Deep Learning Research Dataset for Effective Transfer Learning. *Radiol. Artif. Intell.*, 4:e210315, 2022. 1

[21] OpenAI. New embedding models and API updates, 2024. https://openai.com/index/new-embedding-models-and-api-updates/ [Accessed: 2026-08-21]. 5

[22] OpenAI. Introducing GPT-4.1 in the API, 2025. https://openai.com/index/gpt-4-1/ [Accessed: 2026-06-16]. 5

[23] OpenAI. Introducing gpt-oss, 2025. https://openai.com/index/introducing-gpt-oss/ [Accessed: 2026-08-27]. 7

[24] Sophie Ostmeier, Justin Xu, Zhihong Chen, Maya Varma, Louis Blankemeier, Christian Bluethgen, Arne Edward Michalson, Michael Moseley, Curtis Langlotz Akshay S Chaudhari, and Jean-Benoit Delbrouck. GREEN: Generative Radiology Report Evaluation and Error Notation. In *EMNLP Findings*, pages 374–390, 2024. 6, 7

[25] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a Method for Automatic Evaluation of Machine Translation. In *ACL*, pages 311–318, 2002. 5

[26] Magdalini Paschali, Zhihong Chen, Louis Blankemeier, Maya Varma, Alaa Youssef, Christian Bluethgen, Curtis Langlotz, Sergios Gatidis, and Akshay Chaudhari. Foundation Models in Radiology: What, How, Why, and Why Not. *Radiology*, 314:e240597, 2025. 1

[27] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning Transferable Visual Models From Natural Language Supervision. In *ICML*, pages 8748–8763, 2021. 2

[28] François Remy, Kris Demuynck, and Thomas Demeester. BioLORD-2023: semantic textual representations fusing large language models and clinical knowledge graph insights. *Journal of the American Medical Informatics Association*, 31(9):1844–1855, 2024. 5

[29] Johannes Rückert, Louise Bloch, Raphael Brüngel, Ahmad Idrissi-Yaghir, Henning Schäfer, Cynthia S. Schmidt, Sven Koitka, Obioma Pelka, Asma Ben Abacha, Alba G. Seco de Herrera, Henning Müller, Peter A. Horn, Felix Nensa, and Christoph M. Friedrich. ROCOv2: Radiology Objects in COntext Version 2, an Updated Multimodal Image Dataset. *Sci. Data*, 11:1–10, 2024. 1

[30] Ji Seung Ryu, Hyunyoung Kang, Yuseong Chu, and Sejung Yang. Vision-language foundation models for medical imaging: a review of current practices and innovations. *Biomed. Eng. Lett.*, 15:809–830, 2025. 2

[31] Euclid Seeram. *Computed Tomography: Physical Principles, Patient Care, Clinical Applications, and Quality Control*. Elsevier, 5 edition, 2022. 1

[32] Vanshali Sharma, Andrea Mia Bejar, Gorkem Durak, and Ulas Bagci. CTest-Metric: A Unified Framework to Assess Clinical Validity of Metrics for CT Report Generation. In *ISBI*, 2026. 7

[33] Zhongyi Shui, Jianpeng Zhang, Weiwei Cao, Sinuo Wang, Ruizhe Guo, Le Lu, Lin Yang, Xianghua Ye, Tingbo Liang, Qi Zhang, and Ling Zhang. Large-scale and Fine-grained Vision-language Pre-training for Enhanced CT Image Understanding. In *ICLR*, 2025. 2, 5, 6, 7

[34] Tassilo Wald, Ibrahim E. Hamamci, Yuan Gao, Sam Bond-Taylor, Harshita Sharma, Maximilian Ilse, Cynthia Lo, Olesya Melnichenko, Anton Schwaighofer, Noel C. F. Codella, Maria T. Wetscherek, Klaus H. Maier-Hein, Panagiotis Korfiatis, Valentina Salvatelli, Javier Alvarez-Valle, and Fernando Pérez-García. Comprehensive language-image pre-training for 3D medical image understanding. *arXiv preprint, arXiv:2510.15042*, 2025. 2, 8

[35] Rongsheng Wang, Fenghe Tang, Qingsong Yao, Rui Yan, Xu Zhang, Zhen Huang, Haoran Lai, Zhiyang He, Xiaodong Tao, Zihang Jiang, and S. Kevin Zhou. SimCroP: Radiograph Representation Learning with Similarity-driven Cross-granularity Pre-training. In *MICCAI*, pages 563 – 573, 2025. 2, 5, 6, 7, 8

[36] Jakob Wasserthal, Hanns-Christian Breit, Manfred T. Meyer, Maurice Pradella, Daniel Hinck, Alexander W. Sauter, Tobias Heye, Daniel T. Boll, Joshy Cyriac, Shan Yang, Michael Bach, and Martin Segeroth. TotalSegmentator: Robust Segmentation of 104 Anatomic Structures in CT Images. *Radiol. Artif. Intell.*, 5(5):e230024, 2023. 3

[37] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid Loss for Language Image Pre-Training. In *ICCV*, pages 11975–11986, 2023. 4

[38] Jianpeng Zhang. GitHub repository. https://github.com/alibaba-damo-academy/fvlm [Accessed: 2026-06-16], 2026. 5

[39] Jingyi Zhang, Jiaxing Huang, Sheng Jin, and Shijian Lu. Vision-Language Models for Vision Tasks: A Survey. *IEEE Trans. Pattern Anal. Mach. Intell.*, 46(8):5625–5644, 2024. 2

[40] Feng Zhao, Yizhou Wu, Mingzhe Hu, Chih-Wei Chang, Ruirui Liu, Richard Qiu, and Xiaofeng Yang. Current progress of digital twin construction using medical imaging. *J. Appl. Clin. Med. Phys.*, 26(9):e70226, 2025. 1

Supplementary text for "Findings-aware Language-Image Pretraining for CT Volume Interpretation"

Supplementary tables

Table S1. P-values for the comparison of different models on various metrics BEFORE Holm correction.

	Report gen. metrics		GREEN subcategory					Zero-shot AUC	
models	BLEU-4	GREEN	g-1	g-2	g-3	g-4	g-5	CT-RATE	RadChestCT
FiLIP vs ImageNet	8.5×10^{-4}	7.9×10^{-4}	$< 10^{-4}$	1.1×10^{-1}	$< 10^{-4}$	4.1×10^{-1}	4.5×10^{-2}	-	-
FiLIP vs CT-CLIP	-	-	-	-	-	-	-	$< 10^{-4}$	$< 10^{-4}$
FiLIP vs fVLM	$< 10^{-4}$	1.9×10^{-3}	$< 10^{-4}$	3.0×10^{-1}	1.3×10^{-3}	1.5×10^{-1}	7.6×10^{-1}	$< 10^{-4}$	1.1×10^{-3}
FiLIP vs SimCroP	1.8×10^{-1}	6.5×10^{-2}	$< 10^{-4}$	2.6×10^{-1}	$< 10^{-4}$	8.9×10^{-3}	1.7×10^{-1}	$< 10^{-4}$	$< 10^{-4}$
FiLIP (w/ BioLORD) vs FiLIP (w/ CXR-BERT)	-	-	-	-	-	-	-	$< 10^{-4}$	-
FiLIP (w/ BioLORD) vs FiLIP (w/ text-embedding-3-large)	-	-	-	-	-	-	-	$< 10^{-4}$	-
FiLIP (w/ GPT-4.1) vs FiLIP (w/ gpt-oss-120b)	-	-	-	-	-	-	-	4.4×10^{-1}	-

Table S2. P-values for the comparison of different models on various metrics AFTER Holm correction. The **bold and underlined values** indicate statistical significance at a 5% significance level.

	Report gen. metrics		GREEN subcategory					Zero-shot AUC	
models	BLEU-4	GREEN	g-1	g-2	g-3	g-4	g-5	CT-RATE	RadChestCT
FiLIP vs ImageNet	$\underline{\mathbf{1.4 \times 10^{-2}}}$	$\underline{\mathbf{1.4 \times 10^{-2}}}$	$\underline{\mathbf{< 10^{-4}}}$	> 0.99	$\underline{\mathbf{6.5 \times 10^{-4}}}$	> 0.99	4.9×10^{-1}	-	-
FiLIP vs CT-CLIP	-	-	-	-	-	-	-	$\underline{\mathbf{< 10^{-4}}}$	$\underline{\mathbf{< 10^{-4}}}$
FiLIP vs fVLM	$\underline{\mathbf{< 10^{-4}}}$	$\underline{\mathbf{2.4 \times 10^{-2}}}$	$\underline{\mathbf{< 10^{-4}}}$	> 0.99	$\underline{\mathbf{1.8 \times 10^{-2}}}$	> 0.99	7.6×10^{-1}	$\underline{\mathbf{< 10^{-4}}}$	$\underline{\mathbf{1.7 \times 10^{-2}}}$
FiLIP vs SimCroP	> 0.99	6.5×10^{-1}	$\underline{\mathbf{1.7 \times 10^{-4}}}$	> 0.99	$\underline{\mathbf{< 10^{-4}}}$	1.1×10^{-1}	> 0.99	$\underline{\mathbf{< 10^{-4}}}$	$\underline{\mathbf{< 10^{-4}}}$
FiLIP (w/ BioLORD) vs FiLIP (w/ CXR-BERT)	-	-	-	-	-	-	-	$\underline{\mathbf{< 10^{-4}}}$	-
FiLIP (w/ BioLORD) vs FiLIP (w/ text-embedding-3-large)	-	-	-	-	-	-	-	$\underline{\mathbf{1.4 \times 10^{-2}}}$	-
FiLIP (w/ GPT-4.1) vs FiLIP (w/ gpt-oss-120B)	-	-	-	-	-	-	-	8.9×10^{-1}	-

Table S3. Correspondence between finding labels and organs, which is equivalent to that adopted in fVLM.

Finding label	Organ
Arterial wall calcification	aorta
Cardiomegaly	heart
Pericardial effusion	
Coronary artery wall calcification	
Hiatal hernia	esophagus
Emphysema	lung
Atelectasis	
Lung nodule	
Lung opacity	
Pulmonary fibrotic sequela	
Pleural effusion	
Mosaic attenuation pattern	
Peribronchial thickening	
Consolidation	
Bronchiectasis	
Interlobular septal thickening	

Supplementary note for the calculation of the inter-image triplet loss

The average and std of unique phrases present only in the anchor, i.e., the number of element in $U^o_{i,j}$ (Eq. (5)) is 4.6 ± 0.4 /anchor when $S_{\mathrm{thr}} = 0.6$. We set the batch size to 56, and the number of unique phrases in a batch is $4.6 \times 55 \approx 254$ on average. This is the effective number of negative samples for the inter-image triplet loss.

Regarding the time required for training, the average and std time per epoch was $1,743 \pm 22$ sec on NVIDIA A100 (80GB) * 8 GPUs, while 440 ± 39 sec without the inter image triplet loss and 597 ± 34 sec for fVLM without co-teaching. The inter-image triplet loss requires $O(|B|^2)$ computational cost as shown in Eq. (8), and is thus more computationally expensive than the intra-image focal similarity loss. However, considering that fVLM with co-teaching (whose implementation is not publicly available) requires training two models simultaneously, including the text encoder, the computational cost for training FiLIP is not substantially different from that of fVLM.

Supplementary note for the architecture for radiology report generation tasks

We trained an end-to-end model for radiology report generation using CT-RATE. The pretrained image encoder was connected to the BLIP text decoder through cross-attention mediated by a visual projector, and the entire network was fine-tuned (Fig. 4b). The visual projector sequentially applies a linear layer, a GELU activation, a second linear layer, positional encoding, layer normalization, and 10% dropout, followed by 3D adaptive average pooling, while preserving the embedding dimension of the image features. We then obtained a feature map with 768 channels and spatial dimensions of $4 \times 4 \times 8$ along the right-to-left, anterior-to-posterior, and craniocaudal axes, respectively. During training, we used nearly the same hyperparameter settings as for image-encoder pretraining, except that the learning rate was reduced from 1e-4 to 1e-5. For the loss function, we used the same formulation as that used for the text decoder in BLIP.