arXiv:2611.04092v1 [cs.CV] 3 Nov 2026

CHROMA-MILL: Joint Low-Light Enhancement and White Balance under Colored Illumination

Pol Riubrogent-Comas[1,2] David Serrano-Lozano[1,2] Maria Pilligua[1,2]
Michael S. Brown[3] Ramon Baldrich[1,2] Javier Vazquez-Corral[1,2]
[1]Computer Vision Center [2]Universitat Autònoma de Barcelona [3]York University

https://color.cvc.uab.cat/chromamill/

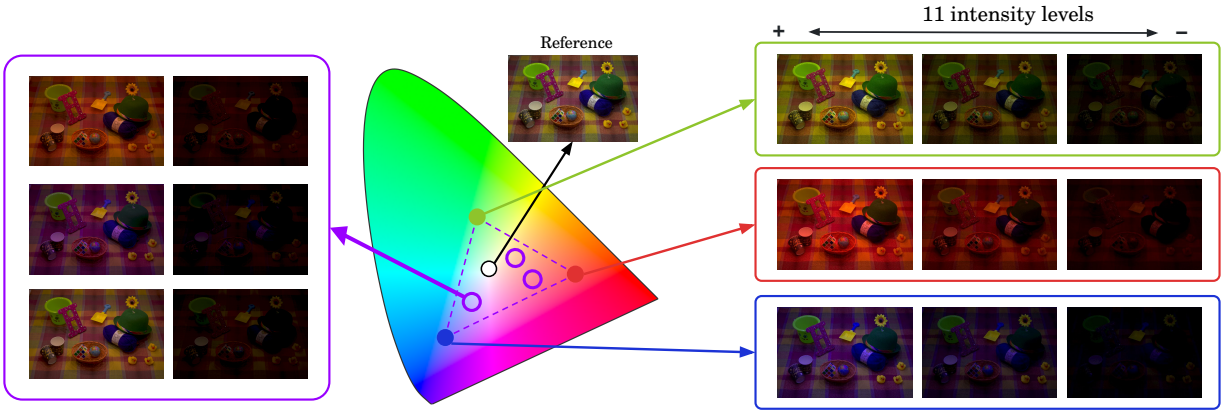

Figure 1. Overview of CHROMA-MILL. Scenes are captured under jointly varying illumination and intensity. The CIE 1931 diagram shows the four captured illuminants (three colored and one white); linear-RAW blending places the effective illuminant anywhere in their convex hull (dashed triangle). Each mixture is captured over 11 intensity levels, from well-lit to near-dark. The well-lit white-illuminant capture (Level 11, Reference) is the ground truth.

Abstract

Capturing an image in a low-light environment produces a weak, noisy sensor signal. Low-light image enhancement (LLIE) aims to restore it. Auto white balance (AWB) is also required for every image, and the same noisy signal degrades its ability to infer the scene illuminant. LLIE is typically performed first, followed by AWB, but this strategy remains largely unexamined because existing low-light datasets span a narrow range of illuminants and lack illuminant ground truth for evaluating AWB. We address this with CHROMA-MILL, a RAW dataset of 50 scenes captured with a DSLR and smartphone under four calibrated illuminants at eleven intensity levels. Because light is additive, we blend these captures in the linear RAW domain to synthesize physically valid illuminant mixtures with exact chromaticity ground truth. Benchmarking on this data, we find joint correction is best, and introduce CTI-Retinexformer, which supervises an illumination map, a chromaticity prediction, and an illumination-invariant scene embedding, outperforming the best cascade by 0.7 *dB PSNR*$_L$ *and* 1.1 ΔE_{00}.

1. Introduction

Capturing an image in a low-light environment produces a weak sensor signal. Reduced scene radiance lowers the signal-to-noise ratio, leaving captures that are dark, noisy, and stripped of detail [8, 28]. Low-light image enhancement (LLIE) aims to restore this signal, and deep models are now effective at brightening and denoising [7, 11, 12, 36, 39].

Restoring the signal is not the only correction such an image needs. Auto white balance (AWB) is required for every image a camera captures, and the same noisy signal degrades it. An AWB algorithm infers the scene illuminant from cues in the image itself, such as bright or near-neutral surfaces, and every estimator—statistical or learned—needs those cues to be well exposed [38]. In low light they are buried in the same noise that corrupts the scene content, so the darker the capture, the less reliable the illuminant estimate. The problem compounds because low-light capture happens indoors and at night, where artificial sources impose color casts that are both strong and varied: the illuminants hardest to estimate are the ones most likely present

when the signal is weakest.

LLIE is typically performed first, followed by AWB, but this strategy remains largely unexamined. LLIE methods [7, 11, 12, 39] are scored on fidelity to a ground truth carrying the input's own illuminant, so a residual color cast costs them nothing; color constancy methods [1, 4, 5, 15, 33] are scored on well-exposed inputs, and their behavior under severe underexposure is almost unexplored [21]. Neither evaluation measures the pipeline the two are deployed in.

The main issue is the lack of appropriate datasets. Low-light datasets such as LOL [36, 40], SID [8], and MILL [28] span a narrow range of illuminants, having been captured under approximately neutral lighting, and provide no illuminant ground truth against which AWB could be evaluated. As a result, neither separate enhancement error from color error nor present the challenging illuminants that make AWB hard in practice.

We address this with **CHROMA-MILL**. Building on the capture protocol of MILL [28], we capture 50 scenes with a DSLR and smartphone in RAW with fixed camera parameters, under four calibrated illuminants—three colored and one white, with measured CIE xy chromaticities—at eleven calibrated intensity levels. Each input is paired with two intermediate ground truths isolating the individual tasks and a final ground truth combining them, so the two error sources can be attributed separately. Because light is additive and the captures are linear RAW and aligned, we can linearly combine a scene under different illuminants—as in one-light-at-a-time (OLAT) relighting [10, 27]—to synthesize it under any convex mixture of those lights, with exact chromaticity ground truth and no additional capture cost. This densifies our coverage of the chromaticity plane and brings it closer to the variety of light colors found in real indoor scenes, each blend acting as a new illuminant of intermediate chromaticity.

Benchmarking cascades on this data shows that neither ordering is satisfactory on average: whichever module runs first, hands the second an input outside its training distribution, and the errors compound at the interface, most visibly in dark inputs. Joint correction is better overall, and we introduce **CTI-Retinexformer** as a joint baseline to perform it. It builds on TI-Retinexformer [28], which disentangles illumination intensity from scene content in the latent space, and extends that disentanglement to color: a new loss supervises an explicit chromaticity prediction under a uniform-illuminant prior, and the content loss is redesigned so the scene embedding is invariant to illuminant color as well as intensity.

Our contributions are:

- CHROMA-MILL, the first dataset supporting joint evaluation of low-light enhancement and white balance: four calibrated illuminants crossed with eleven calibrated intensity levels over 50 scenes and two cameras, extended with physically valid illuminant mixtures with exact chromaticity ground truth.
- The first systematic benchmark of enhancement and white-balance cascades, in both orderings, quantifying how the two corrections interfere.
- CTI-Retinexformer, a joint architecture and objective disentangling intensity, chromaticity, and scene content in the latent space, which achieves the best overall performance across both cameras.

2. Related Work

2.1. Low-Light Image Enhancement

Early paired LLIE datasets include LOL [36] and its extension LOLv2 [40], captured by varying camera settings, and SID [8], which introduced RAW-domain capture. Later datasets scaled size and diversity [23, 26] or exploited multi-exposure sequences [6]. All of them capture each scene under a single low-light condition or simulate brightness change through camera parameters. MILL [28] addressed this limitation by capturing each scene at eleven calibrated intensity levels with fixed camera settings, revealing that LLIE performance varies substantially across input intensity. However, all these datasets are captured under approximately neutral illumination; none provides color-cast ground truth.

On the method side, classical approaches build on Retinex theory [19], estimating an illumination map to recover reflectance [13]. RetinexNet [36] introduced learned decomposition, refined by KinD [48] and KinD++ [49]. Reference-free formulations include Zero-DCE [12], SCI [25], RUAS [24], and EnlightenGAN [16]. Transformer architectures [44] now dominate: Retinexformer [7] couples a one-stage Retinex framework with illumination-guided attention, and subsequent work explores frequency-domain modeling [34, 47], state-space models [37], color-space design [39], geometry-guided normalization [31], brightness-mismatch losses [22], and diffusion priors [41, 50]. Recently, TI-Retinexformer [28] modified Retinexformer to handle multiple low-light intensity levels. However, like other LLIE methods, it brightens and denoises without estimating or correcting the chromaticity of the scene illuminant.

2.2. Color Constancy and White Balance

Statistical illuminant estimation descends from the gray-world assumption [5] and the white-patch heuristic [19], unified in the gray-edge framework [33]. Learned estimators improved accuracy substantially: FC4 [15] pools patch-wise estimates with learned confidence weights, and FFCC [4] solves estimation as localization on a torus in the frequency domain. Deep White-Balance Editing [1] reframes correction as direct sRGB re-rendering, and [2]

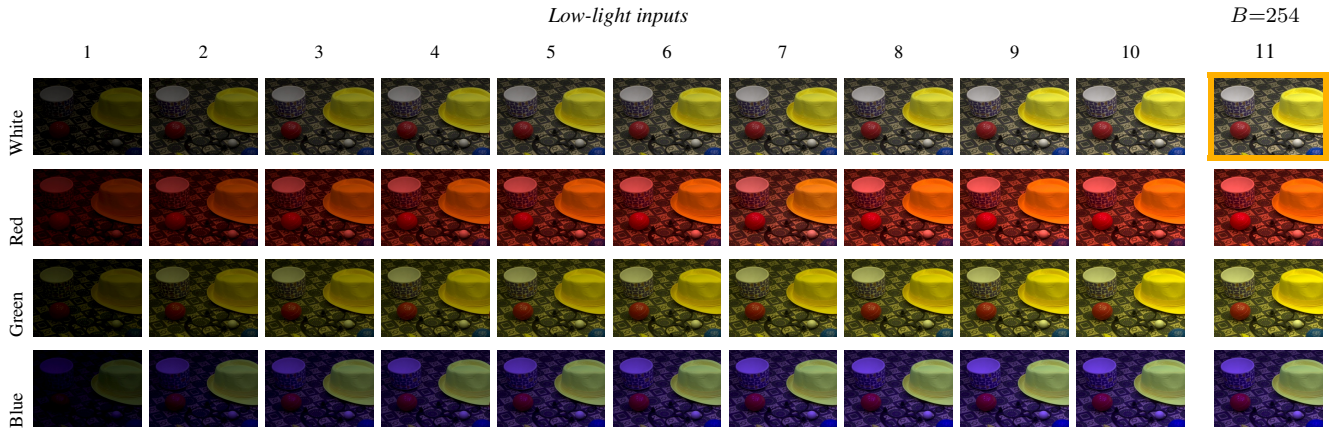

Figure 2. One scene of CHROMA-MILL under the four illuminants (white and three colors), shown at the ten low-light levels (Level 1 is the darkest) and at Level 11, each illuminant's maximum-power capture. The white Level 11 capture (amber frame) is the ground truth and is shared by all rows. All cells use the same linear-to-sRGB transform.

shows that errors baked in by the camera's nonlinear rendering are difficult to undo afterwards. Beyond the single-illuminant assumption, real scenes often have mixed lighting conditions [3]: LSMI [17] provides large-scale captures with per-pixel mixture ratios, AID [18] decomposes illumination with slot attention, and fusion-based correction [30] renders multi-illuminant training data at scale. All of these methods presume adequately exposed input, and recent analysis shows deep color constancy models exhibit notable sensitivity to brightness perturbations [38]. Estimation under severe underexposure remains almost unexplored, with only recent work on nighttime scenes [21]. Thus, our benchmark addresses a timely and largely unexplored gap.

2.3. Joint Restoration

Joint handling of co-occurring degradations has precedent: LEDNet [51] and DarkIR [11] couple low-light enhancement with deblurring, and all-in-one restoration frameworks argue that cascading independently trained modules accumulates errors [45]. For enhancement and color, JCR-Net [42] balances brightness and color in a multi-stage design, and a zero-shot diffusion prior also performs white balance without task-specific training [9]. Neither trains nor evaluates against measured illuminant ground truth, since, to the authors' knowledge, no dataset providing it currently exists.

In this work, we introduce (i) CHROMA-MILL, a new dataset, and (ii) CTI-Retinexformer, a new model, to overcome the current limitations in the literature. CHROMA-MILL, thanks to its intermediate ground truths (white-illuminant capture at the same brightness, maximum brightness capture with the same color cast as the input) and its final ground truth (white-illuminant capture at maximum brightness), allows training cascaded modules for low-light enhancement and color constancy; however, as will be shown in this work, this cascade training falls short at solving the problem, as cascading models leads to error propagation. This is solved by our new model, CTI-Retinexformer, which introduces a color-constancy loss in the latent space of a state-of-the-art low-light model, enabling joint correction of both tasks.

3. The CHROMA-MILL Dataset

Our goal is to capture a calibrated dataset for training and evaluating joint low-light enhancement at different intensity levels and white-balance correction. We build on the capture protocol of MILL [28] and extend it along the color axis, collecting over 4,400 RAW captures. All images are captured in a controlled indoor environment, using a platform with LED programmable lights, a Nikon D5200 DSLR, and a Samsung Galaxy S7 smartphone [28]. Camera parameters remain fixed throughout: ISO 100, f/9 and 1/5 s for the DSLR, f/1.7 and 1/10 s for the smartphone, following the 95%-saturation criterion of [28]. All images are captured in RAW format and processed in the linear domain. Although LED illuminants can exhibit spectral peaks and do not perfectly reproduce continuous Planckian spectra, they are the preferred technology for fine-grained control of both intensity and color [28].

Colored illuminants and intensity levels. We capture each scene under four calibrated illuminants: three colored and one white. The colored illuminants are selected as the maximum red, green, and blue with x, y values of $(0.5235, 0.3297)$, $(0.2821, 0.4727)$, and $(0.1968, 0.1623)$ respectively. The white illuminant mea-

Table 1. Quantitative comparison on the DSLR test set of CHROMA-MILL. Results are averaged over all test images. Every two-stage pipeline is evaluated in both cascade orderings and we report the better one, marking low-light enhancement-first rows with †. Best, second best, and third best results are highlighted.

	Venue	PSNR$_L$ ↑	PSNR$_C$ ↑	SSIM ↑	LPIPS ↓	MAE$_\angle$ ↓	ΔE_{76} ↓	ΔE_{00} ↓
Unprocessed	–	18.764	17.676	0.552	0.253	23.160	22.843	16.152
TI-Retinexformer [28] + DWB [1]	CVPRF'26 / CVPR'20	24.776	22.677	0.759	0.172	14.260	14.206	10.178
TI-Retinexformer [28] + FC4 [15]	CVPRF'26 / CVPR'17	23.784	19.546	0.653	0.280	26.035	24.476	16.652
TI-Retinexformer [28] + FFCC [4]†	CVPRF'26 / CVPR'17	23.723	19.638	0.677	0.253	26.506	24.429	16.861
TI-Retinexformer [28] + RL-AWB [21]†	CVPRF'26 / ECCV'26	22.915	22.034	0.685	0.183	13.618	14.603	10.792
Retinexformer [7] + DWB [1]†	ICCV'23 / CVPR'20	22.525	19.757	0.671	0.274	18.756	18.894	13.079
Retinexformer [7] + FC4 [15]	ICCV'23 / CVPR'17	23.599	19.875	0.649	0.277	26.750	23.853	16.555
Retinexformer [7] + FFCC [4]†	ICCV'23 / CVPR'17	23.397	19.818	0.673	0.252	27.539	24.137	16.922
Retinexformer [7] + RL-AWB [21]	ICCV'23 / ECCV'26	25.666	22.885	0.702	0.195	19.149	17.665	12.041
DarkIR [11] + DWB [1]	CVPR'25 / CVPR'20	26.087	23.676	0.783	0.165	14.738	13.787	9.838
DarkIR [11] + FC4 [15]†	CVPR'25 / CVPR'17	20.548	19.239	0.622	0.272	25.926	22.302	16.755
DarkIR [11] + FFCC [4]†	CVPR'25 / CVPR'17	23.985	20.160	0.694	0.256	27.660	24.912	17.292
DarkIR [11] + RL-AWB [21]†	CVPR'25 / ECCV'26	23.486	22.470	0.706	0.178	13.259	14.133	10.629
MIRNet [43] + DWB [1]	ECCV'20 / CVPR'20	26.350	23.663	0.787	0.165	14.847	13.510	9.697
MIRNet [43] + FC4 [15]	ECCV'20 / CVPR'17	23.716	19.959	0.660	0.273	26.720	23.663	16.385
MIRNet [43] + FFCC [4]†	ECCV'20 / CVPR'17	23.426	19.821	0.685	0.245	27.689	24.129	16.851
MIRNet [43] + RL-AWB [21]	ECCV'20 / ECCV'26	25.429	22.604	0.712	0.192	19.636	17.935	12.175
PromptNorm [31] + DWB [1]†	CVPRW'25 / CVPR'20	22.235	20.855	0.773	0.157	13.495	15.173	11.221
PromptNorm [31] + FC4 [15]†	CVPRW'25 / CVPR'17	22.886	19.909	0.678	0.251	25.027	21.959	16.531
PromptNorm [31] + FFCC [4]†	CVPRW'25 / CVPR'17	25.578	19.813	0.732	0.240	26.498	25.973	17.606
PromptNorm [31] + RL-AWB [21]†	CVPRW'25 / ECCV'26	24.864	23.429	0.753	0.153	11.812	12.618	9.714
Ours (CTI-Retinexformer)	–	27.034	24.914	0.809	0.134	11.538	11.144	8.562

sures $(0.3127, 0.3290)$. All four chromaticities were measured and captured with a Macbeth ColorChecker.

For each illuminant, we capture eleven brightness levels, from minimum to maximum light power linearly based on lux meter readings. The maximum level (B=254) serves as reference; the remaining ten levels serve as low-light inputs.

Thus, each low-light input is paired with two "middle" ground truths: a white-illuminant capture at the same brightness level, which isolates the white-balance task, and a maximum brightness capture with the same color cast, which isolates the low-light enhancement task. The final ground truth is the white-illuminant capture at maximum brightness combining both tasks. Fig. 2 shows the original captured images for a given scene. We capture 50 scenes with both cameras, split into 30 training, 12 validation, and 8 test scenes with no object overlap, following [28].

Extending chromatic coverage via superposition. The four calibrated illuminants sample the chromaticity plane coarsely. We densify this sampling at no additional capture cost. Because light is additive and our images were captured in the linear RAW domain with static parameters, aligned captures of the same scene under different lighting can be linearly combined to synthesize the same scene under the superposition of those lights. Concretely, given aligned captures $\{I_c\}$ of a scene at a fixed brightness level, one per calibrated illuminant, we combine them with convex weights

Figure 3. Examples of illuminant mixtures: three blends of the blue and red illuminants and one three-illuminant mixture. Same linear-to-sRGB transform as Fig. 2.

in the linear domain.

$$I_{\text{mix}} = \sum_c \alpha_c I_c, \qquad \alpha_c \geq 0, \qquad \sum_c \alpha_c = 1. \quad (1)$$

Each weight is a global scalar, so the synthesized image is again lit by a single spatially uniform illuminant, now of intermediate chromaticity.

The chromaticity of the resulting illuminant follows the same additivity: we recover the CIE XYZ values of each calibrated illuminant from its measured chromaticity and its luminance at full power, combine them with the same weights, and normalize. The result always lies within the convex hull of the calibrated chromaticities. Fig. 3 shows four examples.

The final dataset contains 14,520 images in training, 5,808 validation images, and 3,872 test images, separated into 30, 12, and 8 scenes respectively at 11 intensity levels.

Table 2. Quantitative comparison on the phone test set of CHROMA-MILL. Each enhancer is shown with its best-ranking white-balance module; the full grid is in the supplementary. Orderings and highlighting follow Table 1.

	$\text{PSNR}_L \uparrow$	$\text{PSNR}_C \uparrow$	SSIM ↑	LPIPS ↓	$\text{MAE}_\angle \downarrow$	$\Delta\text{E}_{76} \downarrow$	$\Delta\text{E}_{00} \downarrow$
Unprocessed	16.634	15.761	0.439	0.284	23.062	27.275	19.841
TI-Retinexformer + DWB†	22.636	21.360	0.638	0.163	11.231	13.379	10.162
Retinexformer + DWB†	24.421	22.170	0.647	0.166	11.440	12.463	9.327
DarkIR + DWB†	23.597	21.545	0.627	0.167	12.249	13.395	9.970
MIRNet + DWB†	23.330	21.362	0.632	0.181	12.103	13.626	10.092
PromptNorm + DWB†	21.834	20.121	0.622	0.181	12.256	15.146	11.344
Ours (CTI-Retinexformer)	24.352	22.776	0.649	0.136	10.505	11.387	8.881

Table 3. Performance across input intensity levels on the DSLR test set, from the darkest (Level 1) to the brightest (Level 10). We report PSNR_L, PSNR_C, $\text{MAE}_\angle$ and ΔE_{00}. Orderings and highlighting follow Table 1.

	All levels				Level 1				Level 5				Level 10			
	$\text{PSNR}_L \uparrow$	$\text{PSNR}_C \uparrow$	$\text{MAE}_\angle \downarrow$	$\Delta\text{E}_{00} \downarrow$	$\text{PSNR}_L \uparrow$	$\text{PSNR}_C \uparrow$	$\text{MAE}_\angle \downarrow$	$\Delta\text{E}_{00} \downarrow$	$\text{PSNR}_L \uparrow$	$\text{PSNR}_C \uparrow$	$\text{MAE}_\angle \downarrow$	$\Delta\text{E}_{00} \downarrow$	$\text{PSNR}_L \uparrow$	$\text{PSNR}_C \uparrow$	$\text{MAE}_\angle \downarrow$	$\Delta\text{E}_{00} \downarrow$
Unprocessed	18.764	17.676	23.160	16.152	11.984	11.842	23.318	21.523	17.726	17.317	23.478	15.750	26.631	22.830	20.891	13.061
TI-Retinexformer + DWB	24.776	22.677	14.260	10.178	18.493	17.223	25.359	16.058	25.520	23.384	12.411	9.284	26.125	24.160	10.413	8.387
Retinexformer + RL-AWB	25.666	22.885	19.149	12.041	16.088	15.013	28.937	17.756	26.420	22.960	20.335	12.669	30.048	27.509	9.600	7.181
DarkIR + DWB	26.087	23.676	14.738	9.838	19.880	18.152	30.087	18.462	27.337	24.743	12.349	8.421	27.014	24.960	10.328	7.670
MIRNet + DWB	26.350	23.663	14.847	9.697	20.938	18.176	30.293	17.160	26.990	24.442	12.306	8.639	27.131	24.977	10.358	7.718
PromptNorm + DWB†	22.235	20.855	13.495	11.221	22.428	19.834	23.328	15.707	22.402	21.139	12.208	10.514	22.190	21.105	10.660	10.062
Ours (CTI-Retinexformer)	27.034	24.914	11.538	8.562	22.432	20.117	25.771	16.023	28.216	25.924	9.476	7.362	27.326	25.600	7.906	6.943

4. Method

We propose CTI-Retinexformer, a method that disentangles intensity, color, and scene content in the latent space. The backbone architecture is Retinexformer [7], an architecture comprising an illumination estimator and a U-shaped denoiser built from illumination-guided attention blocks, with base width $C = 40$ and channels doubling at each of two downsampling stages, giving $4C = 160$ channels in the latent space.

In TI-Retinexformer [28], the authors partitioned these latent features of Retinexformer into an intensity channel and a scene embedding. Our model goes one step further by extending the latent disentanglement strategy from intensity to color: we retain the intensity channel Z_I (1 channel), modify the scene embedding Z_s to allocate for color variations (157 channels), and propose a new chromaticity prediction region Z_{ch} (2 channels). Each component is supervised by a dedicated loss. In what follows, we explain the different disentanglement losses.

Intensity prediction loss. Following [28], the first channel predicts the normalized intensity $i_{in} \in [0, 1]$ of the input, given by its light-power level, $i_{in} = B/254$. We penalize the squared error of its spatial mean:

$$\mathcal{L}_i = \left(\overline{Z_I} - i_{in}\right)^2, \tag{2}$$

where $\bar{\cdot}$ denotes the spatial mean.

Chromaticity losses. We propose new chromaticity losses to focus on the color constancy task. In detail, our goal is that the second and third channels of the latent space predict the CIE xy chromaticity c_{in} of the scene illuminant:

$$\mathcal{L}_c = \left\|\overline{Z_{ch}} - c_{in}\right\|_2^2. \tag{3}$$

Moreover, since each training image is lit by a spatially uniform (single or uniformly blended) illuminant, we additionally penalize the spatial variance of the prediction to force the encoding to output a uniform-illuminant prior.

$$\mathcal{L}_u = \text{Var}(Z_{ch}). \tag{4}$$

Scene content loss. In [28], they proposed that the remaining channels should encode the content of the scene independently of illumination. In our case, we extend this scene content loss to color. For that, we consider a triplet margin loss [29] in which the positive is sampled from the same scene as the anchor under a different illumination condition —a different brightness level, a different illuminant color, or both— and the negative from a different scene. Anchor and positive share identical crop coordinates to enforce spatial correspondence.

$$\mathcal{L}_{tr} = \max(\|Z_s^q - Z_s^p\|_2 - \|Z_s^q - Z_s^n\|_2 + m,\ 0). \tag{5}$$

where Z_s^q, Z_s^p, and $Z_s^n \in \mathbb{R}^{\frac{H}{4} \times \frac{W}{4} \times (4C-3)}$ correspond to the anchor, positive and negative latent features after the bottleneck attention block, and m=1; distances are taken along each feature-map element and later averaged.

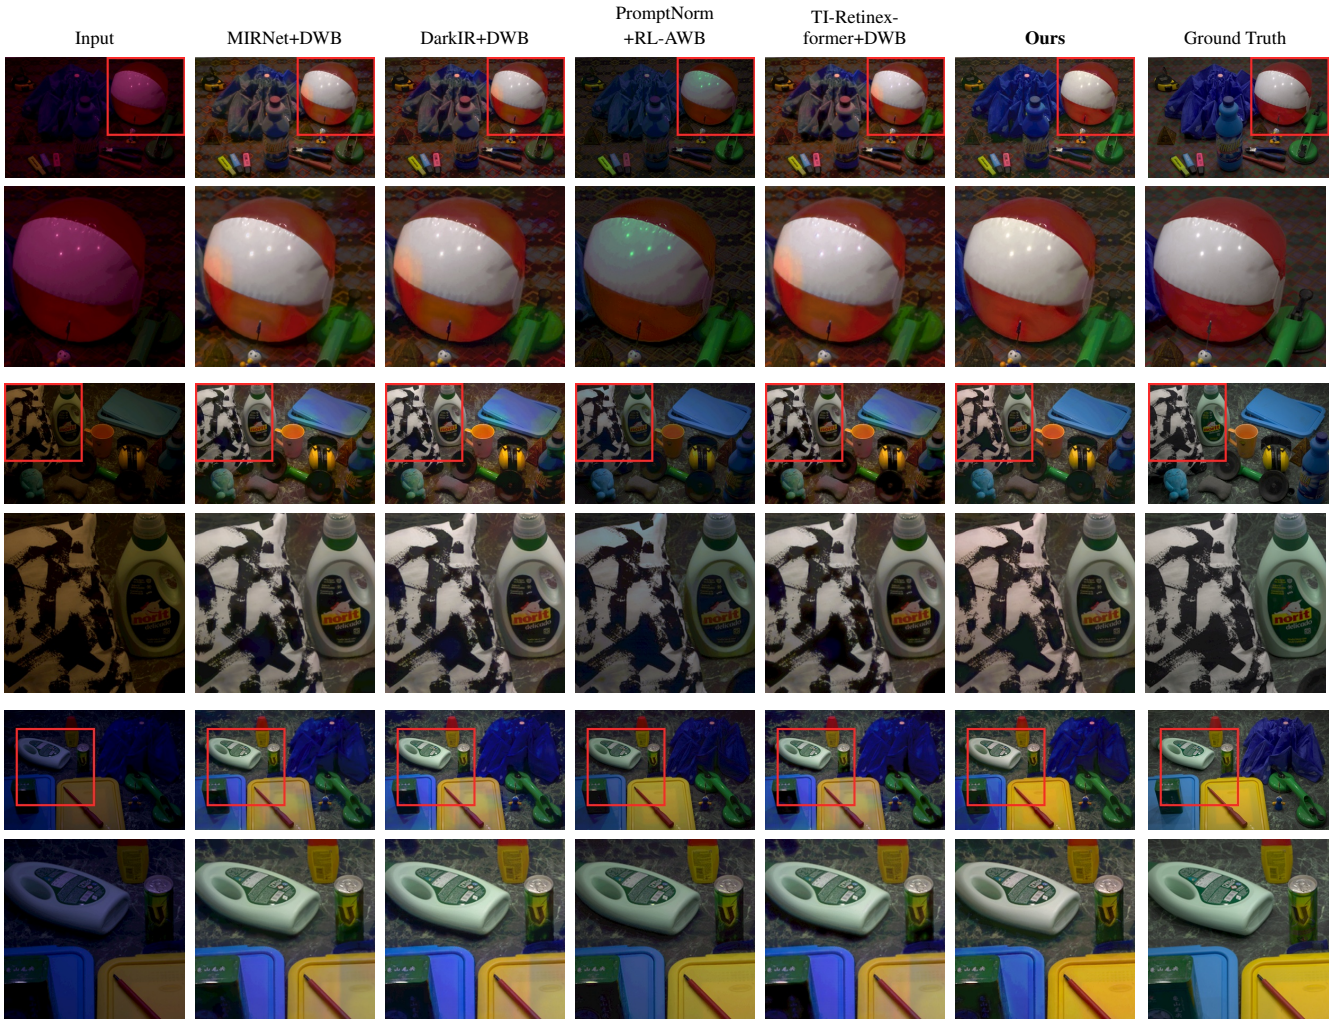

Figure 4. Visual comparison on CHROMA-MILL. From left to right: input, MIRNet+DWB, DarkIR+DWB, PromptNorm+RL-AWB, TI-Retinexformer+DWB, ours, and ground truth. We show three examples under different illuminant colors, with a zoomed-in region below each. The baselines are the four best-performing pipelines of Table 1. Same linear-to-sRGB transform as Fig. 2.

Reconstruction loss. We also add a reconstruction loss $\mathcal{L}_{re}$ defined as the mean L1 distance between the network output and the ground truth image.

Final loss. The complete objective loss is the unweighted sum of all the aforementioned ones:

$$\mathcal{L} = \mathcal{L}_{re} + \mathcal{L}_i + \mathcal{L}_c + \mathcal{L}_u + \mathcal{L}_{tr}. \tag{6}$$

5. Experiments

5.1. Benchmark on CHROMA-MILL

We compare our approach against two-stage pipelines pairing a low-light enhancer with a white-balance module, and against the unprocessed input. The enhancers are Retinexformer [7], TI-Retinexformer [28], DarkIR [11], MIRNet [43], and PromptNorm [31]; the white-balance modules are Deep White-Balance Editing [1], FC4 [15], FFCC [4], and RL-AWB [21]. All models are fine-tuned on the respective training set of CHROMA-MILL, using the intermediate ground truths. We report PSNR on luminance (PSNR_L) and RGB (PSNR_C) [28], SSIM [35], LPIPS [46], mean recovery angular error [14], ΔE_{76}, and ΔE_{2000} [32]. Every two-stage pipeline is evaluated in both cascade orderings —white balance + enhance or enhance + white balance. Table 1 shows the results for the Nikon camera, reporting the ordering that ranks higher across the seven metrics; the complete evaluation with both orderings listed separately is in the supplementary.

Our model leads on every metric in Table 1, and the margin is widest on color. Against MIRNet+DWB, the strongest cascade on fidelity, it gains 0.68 dB PSNR_L and lowers ΔE_{00} by 1.14 and ΔE_{76} by 2.37; against

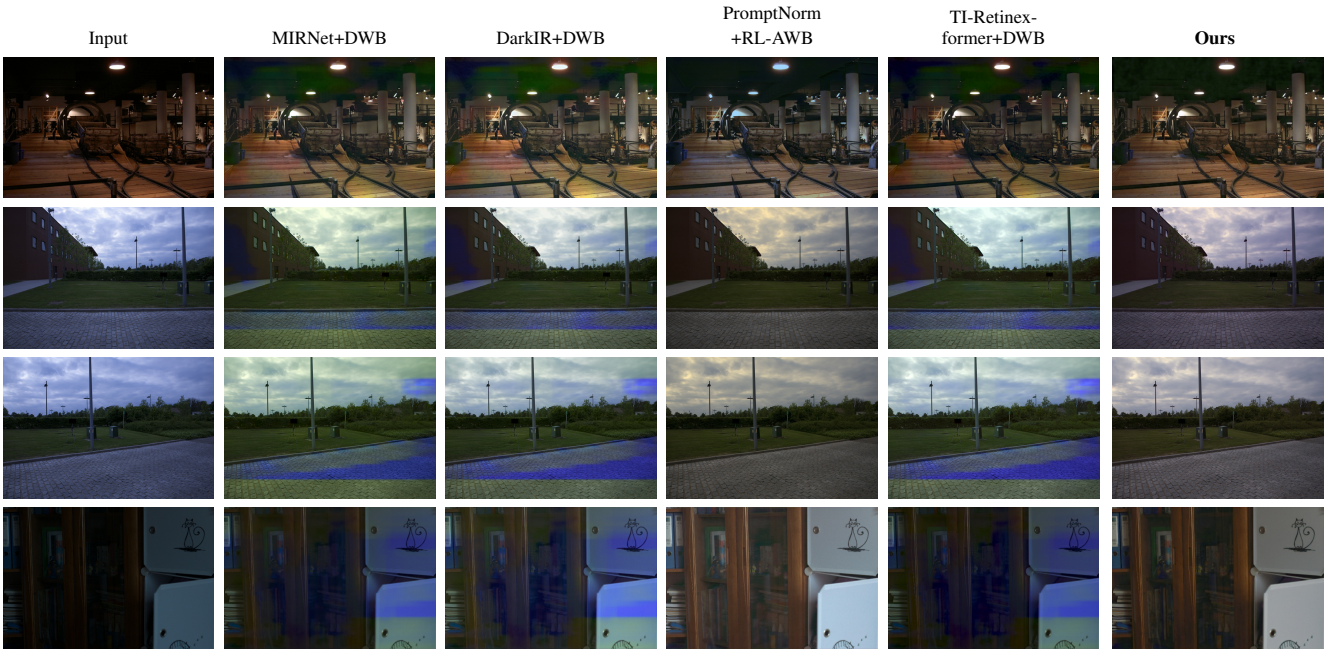

Figure 5. Generalization to other datasets. Rows, top to bottom: one image from DICM [20], two from the Rendered WB dataset [2], and one from LOLv1 [36]. Columns follow Fig. 4: input, the four best-performing pipelines of Table 1, and ours. No ground truth is available for these datasets, so no metrics are reported. All models are trained on CHROMA-MILL only, with no fine-tuning. Same linear-to-sRGB transform as Fig. 2.

PromptNorm+RL-AWB, the strongest on color, it gains 2.17 dB PSNR$_L$ and lowers ΔE_{00} by 1.15 and ΔE_{76} by 1.47. Fidelity and color improve together, and neither is traded for the other; Section 5.3 examines the contribution of each auxiliary loss term.

Two-stage pipelines suffer from error compounding, and no cascade ordering avoids it. A stronger enhancer reduces the fidelity gap but not the color gap: DarkIR [11] and MIRNet [43], both cascaded with Deep White-Balance Editing [1], come within 1 dB of our PSNR$_L$ while remaining more than one ΔE_{00} unit behind. PromptNorm cascaded with RL-AWB is the strongest competitor on color metrics, but it still trails on all metrics and gives up 2.2 dB of PSNR$_L$. This behavior is expected, since the residuals of the first stage are out-of-distribution inputs for the second. In contrast, a single network sharing features across both tasks has no such interface. The cascade ordering does not compensate for the decomposition on average. Which ordering performs better depends on the pipeline: white-balance-first for the DWB cascades built on DarkIR and MIRNet, and enhancement-first for all the PromptNorm pipelines. The difference between the orderings can exceed 3 dB of PSNR$_L$, so it must be tuned for each pair of methods. Even so, averaged over all levels, the better of the two orderings never reaches our jointly trained model.

Table 2 provides the results for the Smartphone test set, with all the models retrained on the phone training split. Only the best combination for each low-light enhancer is shown; the full table is in the supplementary. Our method obtains the best results on every metric except PSNR$_L$ where Retinexformer cascaded with Deep White-Balance Editing is 0.07 dB higher while being worse on the six remaining metrics. The margin on color is even larger than on the Nikon split. The absolute values are lower for all the methods due to the noisier sensor.

Table 3 reports performance across input intensity levels for the Nikon camera. Only the best results from Table 1 are presented. Our method obtains the best results in 12 out of 16 cases, being second in 3 other ones. The advantage is not uniform across levels: at the brightest inputs (Levels 7–10), Retinexformer+RL-AWB leads on PSNR$_L$ by up to 2.7 dB, while ours still leads on ΔE_{00}. At high SNR, enhancement is close to the identity and the white-balance module receives a well-exposed input, so little error compounds at the interface.

5.2. Qualitative Results

Figure 4 shows three examples at Levels 2 and 3, under two synthesized blends and the pure blue primary, against the four strongest baselines of Table 1. In each case, CTI-Retinexformer leads all four on PSNR$_L$, PSNR$_C$, and ΔE_{00}. The failure modes depend more on the illuminant than on the enhancer. The three DWB cascades *overcorrect* the two blends, carrying whites past neutral, and *undercorrect* the blue primary, leaving a deep blue cast.

Table 4. Ablation of the loss terms on the DSLR test set. Each row enables a different subset of the auxiliary losses; the reconstruction loss $\mathcal{L}_{re}$ is always active and every row is trained jointly to the final ground truth, from scratch, with the same training configuration.

Model	$\mathcal{L}_i$	$\mathcal{L}_{tr}$	$\mathcal{L}_c$	$\mathcal{L}_u$	$\text{PSNR}_L \uparrow$	$\text{PSNR}_C \uparrow$	$\text{MAE}_\angle \downarrow$	$\Delta\text{E}_{76} \downarrow$
Rec. only					24.885	23.522	**11.255**	11.329
I	✓				26.408	24.460	11.641	11.309
S		✓			25.632	23.735	12.306	12.345
TI	✓	✓			26.208	24.374	11.913	11.547
CTI w/o $\mathcal{L}_u$	✓	✓	✓		26.843	24.515	11.683	11.339
CTI (ours)	✓	✓	✓	✓	**27.034**	**24.914**	11.538	**11.144**

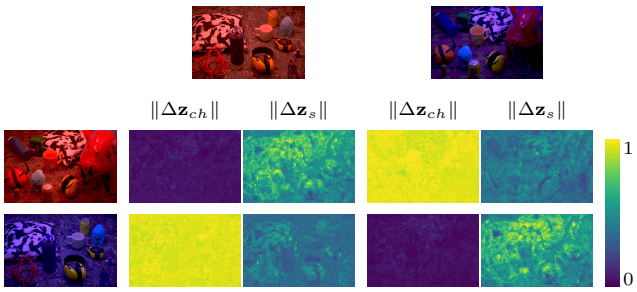

Figure 6. Response of the bottleneck to a change of a single factor. Each anchor (left column) is compared against each capture (top row); every pair differs in exactly one factor, so the main diagonal changes the scene at a fixed illuminant and the anti-diagonal changes the illuminant at a fixed scene.

PromptNorm+RL-AWB shifts the cast rather than removing it and reconstructs the scene far more darkly, more than 5 dB below the DWB cascades on PSNR_L in both blends. Both failure modes are consistent with the error compounding of Sec. 5.1. Our joint model stays closest to the ground truth, though the blue primary remains hard for it too, with a residual cast and visible grain.

Generalization to other datasets. Figure 5 applies the models of Table 1 to DICM [20], Rendered WB [2] and LOLv1 [36] images without fine-tuning. On the Rendered WB scenes the cascades reproduce their failure modes under domain shift: MIRNet+DWB overshoots into a yellow-green cast, while DarkIR+DWB and TI-Retinexformer+DWB leave much of the input's blue tint. CTI-Retinexformer removes the cast most completely while preserving the warm tungsten character of the DICM. On the last example PromptNorm+RL-AWB gives a brighter exposure, but ours still avoids the strong blue cast of the other cascades.

5.3. Ablations

We now isolate the contribution of each auxiliary supervision in Eq. (6). Table 4 reports PSNR_L, PSNR_C, $\text{MAE}_\angle$ and ΔE_{76} on the Nikon test set. Removing $\mathcal{L}_u$ degrades all four metrics, most notably PSNR_C (-0.40 dB). Further removing $\mathcal{L}_c$, which leaves TI-Retinexformer's losses (TI), costs another 0.64 dB PSNR_L, and the triplet then becomes harmful: keeping only $\mathcal{L}_i$ (I) is better on all four metrics. Since our triplet positives vary the illuminant color, $\mathcal{L}_{tr}$ pushes color out of Z_s, with nowhere else to go without a chromaticity channel. Removing all auxiliary terms (Reconstruction loss only) costs 2.15 dB PSNR_L but slightly lowers $\text{MAE}_\angle$, which ignores intensity. Even so, this jointly trained model lowers ΔE_{76} from 14.21 to 11.33 relative to the TI-Retinexformer+DWB cascade: joint training brings most of the color gain, the auxiliary losses most of the fidelity gain.

5.4. Feature Disentanglement Analysis

We extend the latent-space analysis of TI-Retinexformer [28], which separates an intensity map from a scene embedding, to the three components of our bottleneck by adding the chromaticity loss. Figure 6 reports the response of $\mathbf{z}_{ch}$ and $\mathbf{z}_s$ to pairs of images that differ in exactly one factor. Each code moves with the factor it is supervised for and stays comparatively still under the other. The separation is sharpest for chromaticity. Changing the illuminant displaces $\mathbf{z}_{ch}$ more than $\mathbf{z}_s$ does for a change of scene, and the response is close to uniform over the frame. $\mathbf{z}_s$ is complementary, responding most to a change of scene, with a spatial structure that follows scene content, while still retaining a bit of response to the illuminant change.

6. Conclusion

We studied the interaction between low-light enhancement and auto white balance, both challenged by the weak, noisy sensor signal in low light. To evaluate them jointly, we introduced CHROMA-MILL, a RAW dataset of 50 scenes captured with a DSLR and a smartphone under four calibrated illuminants at eleven intensity levels, extended with physically valid illuminant mixtures with exact chromaticity ground truth.

Our benchmark shows that cascading low-light enhancement and AWB is insufficient. CTI-Retinexformer, which jointly supervises an illumination map, a chromaticity prediction, and a scene embedding, outperforms the best two-stage alternative by 0.7 dB PSNR_L and 1.1 ΔE_{00}, showing the importance of addressing both tasks jointly.

Limitations. Every CHROMA-MILL image is lit by a spatially uniform illuminant, which $\mathcal{L}_u$ builds into the model as a prior; spatially varying lighting, synthesizable with per-pixel weights in Eq. (1), is left to future work. Our LEDs span CIE F2 and the Planckian locus above about 3400 K, but not warmer sources such as CIE A. Finally, the dataset has 50 indoor scenes (8 for testing), models are trained per camera, and generalization beyond CHROMA-MILL is shown only qualitatively (Fig. 5).

Acknowledgements

This work was supported by Grant PID2024-162555OB-I00 funded by MCIN/AEI/10.13039/501100011033 and by ERDF, EU, AIA2025-163919-C52 funded by MCIN/AEI/10.13039/501100011033, the grant Catedra ENIA UAB-Cruïlla (TSI-100929-2023-2) from the Ministry of Economic Affairs and Digital Transformation of Spain, and by the Generalitat de Catalunya CERCA Program. JVC also acknowledges the 2025 Leonardo Grant for Scientific Research and Cultural Creation from the BBVA Foundation. The BBVA Foundation accepts no responsibility for the opinions, statements and contents included in the project and/or the results thereof, which are entirely the responsibility of the authors. DSL also acknowledges the FPI grant from Spanish Ministry of Science and Innovation (PRE2022-101525).

References

[1] Mahmoud Afifi and Michael S. Brown. Deep white-balance editing. In *CVPR*, 2020. 2, 4, 6, 7

[2] Mahmoud Afifi, Brian Price, Scott Cohen, and Michael S. Brown. When color constancy goes wrong: Correcting improperly white-balanced images. In *CVPR*, 2019. 2, 7, 8

[3] Mahmoud Afifi, Marcus A. Brubaker, and Michael S. Brown. Auto white-balance correction for mixed-illuminant scenes. In *WACV*, 2022. 3

[4] Jonathan T. Barron and Yun-Ta Tsai. Fast fourier color constancy. In *CVPR*, 2017. 2, 4, 6

[5] Gershon Buchsbaum. A spatial processor model for object colour perception. *Journal of the Franklin Institute*, 310(1): 1–26, 1980. 2

[6] Jianrui Cai, Shuhang Gu, and Lei Zhang. Learning a deep single image contrast enhancer from multi-exposure images. *IEEE Transactions on Image Processing*, 27(4):2049–2062, 2018. 2

[7] Yuanhao Cai, Hao Bian, Jing Lin, Haoqian Wang, Radu Timofte, and Yulun Zhang. Retinexformer: One-stage retinex-based transformer for low-light image enhancement. In *ICCV*, 2023. 1, 2, 4, 5, 6

[8] Chen Chen, Qifeng Chen, Jia Xu, and Vladlen Koltun. Learning to see in the dark. In *CVPR*, 2018. 1, 2

[9] Joshua Cho, Sara Aghajanzadeh, Zhen Zhu, and D. A. Forsyth. Mr. illuminate: Zero-shot low light image enhancement with diffusion prior. In *CVPR*, 2026. 3

[10] Paul Debevec, Tim Hawkins, Chris Tchou, Haarm-Pieter Duiker, Westley Sarokin, and Mark Sagar. Acquiring the reflectance field of a human face. In *SIGGRAPH*, 2000. 2

[11] Daniel Feijoo, Juan C. Benito, Alvaro Garcia, and Marcos V. Conde. DarkIR: Robust low-light image restoration. In *CVPR*, 2025. 1, 2, 3, 4, 6, 7

[12] Chunle Guo, Chongyi Li, Jichang Guo, Chen Change Loy, Junhui Hou, Sam Kwong, and Runmin Cong. Zero-reference deep curve estimation for low-light image enhancement. In *CVPR*, 2020. 1, 2

[13] Xiaojie Guo, Yu Li, and Haibin Ling. LIME: Low-light image enhancement via illumination map estimation. *IEEE Transactions on Image Processing*, 26(2):982–993, 2017. 2

[14] Steven D. Hordley and Graham D. Finlayson. Reevaluation of color constancy algorithm performance. *JOSA A*, 23(5): 1008–1020, 2006. 6

[15] Yuanming Hu, Baoyuan Wang, and Stephen Lin. FC4: Fully convolutional color constancy with confidence-weighted pooling. In *CVPR*, 2017. 2, 4, 6

[16] Yifan Jiang, Xinyu Gong, Ding Liu, Yu Cheng, Chen Fang, Xiaohui Shen, Jianchao Yang, Pan Zhou, and Zhangyang Wang. EnlightenGAN: Deep light enhancement without paired supervision. *IEEE Transactions on Image Processing*, 30:2340–2349, 2021. 2

[17] Dongyoung Kim, Jinwoo Kim, Seonghyeon Nam, Dongwoo Lee, Yeonkyung Lee, Nahyup Kang, Hyong-Euk Lee, ByungIn Yoo, Jae-Joon Han, and Seon Joo Kim. Large scale multi-illuminant (LSMI) dataset for developing white balance algorithm under mixed illumination. In *ICCV*, 2021. 3

[18] Dongyoung Kim, Jinwoo Kim, Junsang Yu, and Seon Joo Kim. Attentive illumination decomposition model for multi-illuminant white balancing. In *CVPR*, 2024. 3

[19] Edwin H. Land and John J. McCann. Lightness and retinex theory. *Journal of the Optical Society of America*, 61(1):1–11, 1971. 2

[20] Chulwoo Lee, Chul Lee, and Chang-Su Kim. Contrast enhancement based on layered difference representation. In *ICIP*, 2012. 7, 8

[21] Yuan-Kang Lee, Kuan-Lin Chen, Chia-Che Chang, and Yu-Lun Liu. Rl-awb: Deep reinforcement learning for auto white balance correction in low-light nighttime scenes. In *European Conference on Computer Vision*, pages 77–99. Springer, 2026. 2, 3, 4, 6

[22] Jingxi Liao, Shijie Hao, Richang Hong, and Meng Wang. GT-mean loss: A simple yet effective solution for brightness mismatch in low-light image enhancement. In *ICCV*, 2025. 2

[23] Jiaying Liu, Dejia Xu, Wenhan Yang, Minhao Fan, and Haofeng Huang. Benchmarking low-light image enhancement and beyond. *International Journal of Computer Vision*, 129:1153–1184, 2021. 2

[24] Risheng Liu, Long Ma, Jiaao Zhang, Xin Fan, and Zhongxuan Luo. Retinex-inspired unrolling with cooperative prior architecture search for low-light image enhancement. In *CVPR*, 2021. 2

[25] Long Ma, Tengyu Ma, Risheng Liu, Xin Fan, and Zhongxuan Luo. Toward fast, flexible, and robust low-light image enhancement. In *CVPR*, 2022. 2

[26] Alexandra Malyugina, Nantheera Anantrasirichai, and David Bull. A topological loss function for image denoising on a new BVI-lowlight dataset. *Signal Processing*, 211:109081, 2023. 2

[27] Jeffry S. Nimeroff, Eero Simoncelli, and Julie Dorsey. Efficient re-rendering of naturally illuminated environments. In *Eurographics Workshop on Rendering*, 1994. 2

[28] Maria Pilligua, David Serrano-Lozano, Pai Peng, Ramon Baldrich, Michael S. Brown, and Javier Vazquez-Corral. MILL: Evaluating low-light image enhancement across multiple intensity levels. In *CVPR Findings*, 2026. 1, 2, 3, 4, 5, 6, 8, 11

[29] Florian Schroff, Dmitry Kalenichenko, and James Philbin. FaceNet: A unified embedding for face recognition and clustering. In *CVPR*, 2015. 5

[30] David Serrano-Lozano, Aditya Arora, Luis Herranz, Konstantinos G. Derpanis, Michael S. Brown, and Javier Vazquez-Corral. Revisiting image fusion for multi-illuminant white-balance correction. In *ICCV*, 2025. 3

[31] David Serrano-Lozano, Francisco A. Molina-Bakhos, Danna Xue, Yixiong Yang, Maria Pilligua, Ramon Baldrich, Maria Vanrell, and Javier Vazquez-Corral. PromptNorm: Image geometry guides ambient light normalization. In *CVPR Workshops*, 2025. 2, 4, 6

[32] Gaurav Sharma, Wencheng Wu, and Edul N. Dalal. The CIEDE2000 color-difference formula: Implementation notes, supplementary test data, and mathematical observations. *Color Research & Application*, 30(1):21–30, 2005. 6

[33] Joost van de Weijer, Theo Gevers, and Arjan Gijsenij. Edge-based color constancy. *IEEE Transactions on Image Processing*, 16(9):2207–2214, 2007. 2

[34] Chenxi Wang, Hongjun Wu, and Zhi Jin. FourLLIE: Boosting low-light image enhancement by fourier frequency information. In *ACM International Conference on Multimedia*, 2023. 2

[35] Zhou Wang, Alan C. Bovik, Hamid R. Sheikh, and Eero P. Simoncelli. Image quality assessment: From error visibility to structural similarity. *IEEE Transactions on Image Processing*, 13(4):600–612, 2004. 6

[36] Chen Wei, Wenjing Wang, Wenhan Yang, and Jiaying Liu. Deep retinex decomposition for low-light enhancement. In *BMVC*, 2018. 1, 2, 7, 8

[37] Jiangwei Weng, Zhiqiang Yan, Ying Tai, Jianjun Qian, Jian Yang, and Jun Li. MambaLLIE: Implicit retinex-aware low light enhancement with global-then-local state space. In *NeurIPS*, 2024. 2

[38] Mengda Xie, Chengzhi Zhong, Yiling He, Zhan Qin, and Meie Fang. Boosting illuminant estimation in deep color constancy through brightness robustness enhancement. *Pattern Recognition*, 171:112153, 2026. 1, 3

[39] Qingsen Yan, Yixu Feng, Cheng Zhang, Guansong Pang, Kangbiao Shi, Peng Wu, Wei Dong, Jinqiu Sun, and Yanning Zhang. HVI: A new color space for low-light image enhancement. In *CVPR*, 2025. 1, 2

[40] Wenhan Yang, Wenjing Wang, Haofeng Huang, Shiqi Wang, and Jiaying Liu. Sparse gradient regularized deep retinex network for robust low-light image enhancement. *IEEE Transactions on Image Processing*, 30:2072–2086, 2021. 2

[41] Xunpeng Yi, Han Xu, Hao Zhang, Linfeng Tang, and Jiayi Ma. Diff-retinex: Rethinking low-light image enhancement with a generative diffusion model. In *ICCV*, 2023. 2

[42] Nana Yu, Hong Shi, and Yahong Han. Joint correcting and refinement for balanced low-light image enhancement. *IEEE Transactions on Multimedia*, 26:6310–6324, 2024. 3

[43] Syed Waqas Zamir, Aditya Arora, Salman Khan, Munawar Hayat, Fahad Shahbaz Khan, Ming-Hsuan Yang, and Ling Shao. Learning enriched features for real image restoration and enhancement. In *ECCV*, 2020. 4, 6, 7

[44] Syed Waqas Zamir, Aditya Arora, Salman Khan, Munawar Hayat, Fahad Shahbaz Khan, and Ming-Hsuan Yang. Restormer: Efficient transformer for high-resolution image restoration. In *CVPR*, 2022. 2

[45] Huimin Zeng, Jiacheng Li, Ziqiang Zheng, and Zhiwei Xiong. All-in-one image compression and restoration. In *WACV*, 2025. 3

[46] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *CVPR*, 2018. 6

[47] Tongshun Zhang, Pingping Liu, Yubing Lu, Mengen Cai, Zijian Zhang, Zhe Zhang, and Qiuzhan Zhou. CWNet: Causal wavelet network for low-light image enhancement. In *ICCV*, 2025. 2

[48] Yonghua Zhang, Jiawan Zhang, and Xiaojie Guo. Kindling the darkness: A practical low-light image enhancer. In *ACM International Conference on Multimedia*, 2019. 2

[49] Yonghua Zhang, Xiaojie Guo, Jiayi Ma, Wei Liu, and Jiawan Zhang. Beyond brightening low-light images. *International Journal of Computer Vision*, 129:1013–1037, 2021. 2

[50] Dewei Zhou, Zongxin Yang, and Yi Yang. Pyramid diffusion models for low-light image enhancement. In *IJCAI*, 2023. 2

[51] Shangchen Zhou, Chongyi Li, and Chen Change Loy. LEDNet: Joint low-light enhancement and deblurring in the dark. In *ECCV*, 2022. 3

Supplementary Material

We provide additional material to supplement our main submission, covering:

1. Extended method details.
2. Additional quantitative results.

A. Extended Method Details

Figure 7 depicts an overview of CTI-Retinexformer. Given a low-light input image x, the encoder maps the image into a latent representation that is explicitly disentangled into three components: an intensity representation Z_I, a chromaticity representation Z_{ch}, and a scene-content representation Z_s. Following TI-Retinexformer [28], Z_I consists of a single channel and encodes the normalized illumination intensity of the input. We extend this decomposition with a two-channel chromaticity representation Z_{ch}, which predicts the CIE xy chromaticity of the scene illuminant, while the remaining 157 channels form Z_s and encode scene content independently of illumination. The three representations are jointly processed by the decoder to reconstruct the enhanced and white-balanced output $\hat{y}$. Each latent component is explicitly supervised during training. The intensity loss $\mathcal{L}_i$ encourages the spatial mean of Z_I to match the normalized light-power level of the input, $i_{in} = B/254$. The chromaticity loss $\mathcal{L}_c$ supervises Z_{ch} using the ground-truth illuminant chromaticity, while the uniformity loss $\mathcal{L}_u$ minimizes its spatial variance to enforce the spatially uniform illuminant prior. Finally, the scene-content loss $\mathcal{L}_{tr}$ uses a triplet objective to encourage Z_s to remain invariant to changes in illumination intensity and color while preserving scene identity. The reconstruction loss $\mathcal{L}_{re}$ supervises the final output as a mean L1 distance with the same weight for every input level. The complete training objective is the sum of the reconstruction and four auxiliary losses: $\mathcal{L} = \mathcal{L}_{re} + \mathcal{L}_i + \mathcal{L}_c + \mathcal{L}_u + \mathcal{L}_{tr}$.

B. Additional Quantitative Results

Tables 5 and 6 report the complete evaluation, with both cascade orderings listed separately for every two-stage pipeline for the DSLR and the smartphone camera, respectively. Table 7 then breaks the DSLR evaluation down by input intensity level, reporting every method at each of the ten levels.

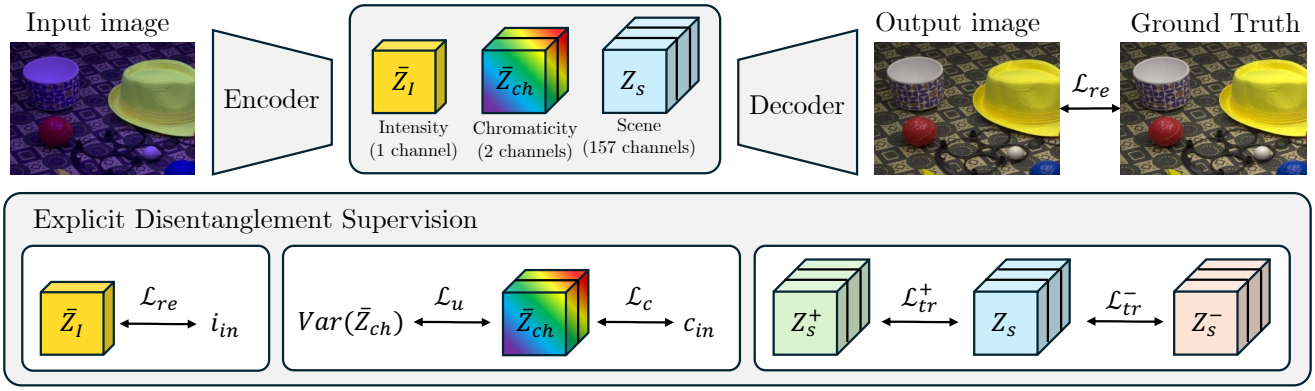

Figure 7. **Overview of CTI-Retinexformer.**

Table 5. Complete evaluation on the Augmented (blended-illuminant) Nikon test set, with both cascade orderings of every two-stage pipeline. White-balance-first applies the white-balance module to the raw low-light input; enhancement-first applies it to the enhancer's output. The main paper reports the better of the two per pipeline. Best results are highlighted in red, second in orange, and third in yellow.

	Order	PSNR$_L$ ↑	PSNR$_C$ ↑	SSIM ↑	LPIPS ↓	MAE$_\angle$ ↓	ΔE_{76} ↓	ΔE_{00} ↓
Unprocessed	–	18.764	17.676	0.552	0.253	23.160	22.843	16.152
TI-Retinexformer + DWB	WB-first	24.776	22.677	0.759	0.172	14.260	14.206	10.178
	LLIE-first	21.092	18.821	0.643	0.289	18.933	20.162	13.871
TI-Retinexformer + FC4	WB-first	23.784	19.546	0.653	0.280	26.035	24.476	16.652
	LLIE-first	20.264	18.958	0.602	0.275	26.055	22.999	16.907
TI-Retinexformer + FFCC	WB-first	24.235	19.280	0.660	0.262	27.544	25.939	17.328
	LLIE-first	23.723	19.638	0.677	0.253	26.506	24.429	16.861
TI-Retinexformer + RL-AWB	WB-first	24.972	22.129	0.696	0.200	19.126	18.723	12.568
	LLIE-first	22.915	22.034	0.685	0.183	13.618	14.603	10.792
Retinexformer + DWB	WB-first	23.226	20.059	0.626	0.286	21.015	20.133	13.381
	LLIE-first	22.525	19.757	0.671	0.274	18.756	18.894	13.079
Retinexformer + FC4	WB-first	23.599	19.875	0.649	0.277	26.750	23.853	16.555
	LLIE-first	20.036	18.873	0.603	0.273	26.109	22.481	16.901
Retinexformer + FFCC	WB-first	24.566	19.700	0.662	0.257	27.464	24.672	16.751
	LLIE-first	23.397	19.818	0.673	0.252	27.539	24.137	16.922
Retinexformer + RL-AWB	WB-first	25.666	22.885	0.702	0.195	19.149	17.665	12.041
	LLIE-first	22.962	22.085	0.688	0.173	12.239	13.824	10.259
DarkIR + DWB	WB-first	26.087	23.676	0.783	0.165	14.738	13.787	9.838
	LLIE-first	22.740	19.804	0.677	0.289	18.852	18.694	13.110
DarkIR + FC4	WB-first	23.787	20.102	0.664	0.283	26.685	23.934	16.646
	LLIE-first	20.548	19.239	0.622	0.272	25.926	22.302	16.755
DarkIR + FFCC	WB-first	24.882	20.019	0.680	0.260	27.671	24.858	16.763
	LLIE-first	23.985	20.160	0.694	0.256	27.660	24.912	17.292
DarkIR + RL-AWB	WB-first	25.578	22.515	0.718	0.198	19.399	18.224	12.282
	LLIE-first	23.486	22.470	0.706	0.178	13.259	14.133	10.629
MIRNet + DWB	WB-first	26.350	23.663	0.787	0.165	14.847	13.510	9.697
	LLIE-first	25.295	22.704	0.739	0.200	18.058	15.698	11.545
MIRNet + FC4	WB-first	23.716	19.959	0.660	0.273	26.720	23.663	16.385
	LLIE-first	20.543	18.996	0.616	0.261	26.177	22.016	16.541
MIRNet + FFCC	WB-first	24.525	19.759	0.673	0.253	27.878	24.744	16.755
	LLIE-first	23.426	19.821	0.685	0.245	27.689	24.129	16.851
MIRNet + RL-AWB	WB-first	25.429	22.604	0.712	0.192	19.636	17.935	12.175
	LLIE-first	22.761	21.861	0.687	0.168	12.924	14.040	10.473
PromptNorm + DWB	WB-first	21.922	20.625	0.774	0.164	13.934	16.024	11.674
	LLIE-first	22.235	20.855	0.773	0.157	13.495	15.173	11.221
PromptNorm + FC4	WB-first	24.828	19.087	0.698	0.274	26.244	26.641	17.875
	LLIE-first	22.886	19.909	0.678	0.251	25.027	21.959	16.531
PromptNorm + FFCC	WB-first	25.049	18.972	0.713	0.257	27.389	27.734	18.119
	LLIE-first	25.578	19.813	0.732	0.240	26.498	25.973	17.606
PromptNorm + RL-AWB	WB-first	23.435	20.469	0.736	0.191	18.688	20.577	13.724
	LLIE-first	24.864	23.429	0.753	0.153	11.812	12.618	9.714
Ours (CTI-Retinexformer)	–	27.034	24.914	0.809	0.134	11.538	11.144	8.562

Table 6. Complete evaluation on the Augmented (blended-illuminant) phone test set, with both cascade orderings of every two-stage pipeline. White-balance-first applies the white-balance module to the raw low-light input; enhancement-first applies it to the enhancer's output. The main paper reports the better of the two per pipeline. Best results are highlighted in red, second in orange, and third in yellow.

	Order	$PSNR_L \uparrow$	$PSNR_C \uparrow$	SSIM ↑	LPIPS ↓	$MAE_\angle \downarrow$	$\Delta E_{76} \downarrow$	$\Delta E_{00} \downarrow$
Unprocessed	–	16.634	15.761	0.439	0.284	23.062	27.275	19.841
Retinexformer + DWB	WB-first	24.090	21.931	0.626	0.178	12.295	13.301	9.804
	LLIE-first	24.421	22.170	0.647	0.166	11.440	12.463	9.327
Retinexformer + FC4	WB-first	23.662	20.602	0.601	0.232	18.286	18.444	14.303
	LLIE-first	19.845	18.500	0.561	0.232	19.372	20.579	16.181
Retinexformer + FFCC	WB-first	22.048	19.512	0.584	0.243	22.342	25.213	17.334
	LLIE-first	20.552	18.845	0.584	0.244	22.978	25.388	18.214
Retinexformer + RL-AWB	WB-first	23.457	21.774	0.611	0.184	12.924	14.797	11.068
	LLIE-first	19.046	18.626	0.577	0.185	12.027	18.182	13.798
DarkIR + DWB	WB-first	23.235	21.304	0.613	0.172	12.908	14.048	10.247
	LLIE-first	23.597	21.545	0.627	0.167	12.249	13.395	9.970
DarkIR + FC4	WB-first	23.470	19.943	0.578	0.237	20.197	20.572	15.449
	LLIE-first	19.275	17.895	0.525	0.236	20.314	21.670	16.649
DarkIR + FFCC	WB-first	21.979	19.091	0.564	0.244	23.499	26.478	17.752
	LLIE-first	20.391	18.613	0.562	0.243	24.207	26.351	18.418
DarkIR + RL-AWB	WB-first	23.652	21.267	0.593	0.186	14.558	16.426	11.776
	LLIE-first	18.859	18.407	0.551	0.183	13.072	18.509	13.985
MIRNet + DWB	WB-first	23.247	21.276	0.627	0.190	12.658	14.035	10.305
	LLIE-first	23.330	21.362	0.632	0.181	12.103	13.626	10.092
MIRNet + FC4	WB-first	21.679	19.336	0.563	0.250	20.912	21.083	16.048
	LLIE-first	18.439	17.380	0.517	0.249	20.222	22.329	17.111
MIRNet + FFCC	WB-first	20.781	18.659	0.555	0.251	23.952	26.775	18.243
	LLIE-first	19.169	17.967	0.550	0.257	24.477	26.726	19.035
MIRNet + RL-AWB	WB-first	22.709	21.017	0.591	0.199	15.218	17.125	12.425
	LLIE-first	17.981	17.702	0.535	0.203	13.285	19.774	14.902
PromptNorm + DWB	WB-first	21.187	19.713	0.613	0.185	12.357	15.827	11.715
	LLIE-first	21.834	20.121	0.622	0.181	12.256	15.146	11.344
PromptNorm + FC4	WB-first	22.504	19.196	0.582	0.247	19.363	21.610	16.139
	LLIE-first	20.157	18.194	0.547	0.247	20.177	21.585	16.584
PromptNorm + FFCC	WB-first	21.649	18.523	0.571	0.256	22.909	27.766	18.443
	LLIE-first	21.340	18.886	0.577	0.257	23.691	26.671	18.516
PromptNorm + RL-AWB	WB-first	22.250	20.238	0.592	0.199	13.978	17.439	12.742
	LLIE-first	19.614	18.753	0.569	0.198	13.239	18.142	13.770
Ours (CTI-Retinexformer)	–	24.352	22.776	0.649	0.136	10.505	11.387	8.881

Table 7. Performance across all ten input intensity levels on the DSLR test set. Level 1 is the darkest input and level 10 the brightest; the reference is captured at maximum brightness. Every method of the main comparison is listed at its selected cascade ordering, rows marked † enhancement-first and the rest white-balance-first. Within each level, best, second best and third best results are highlighted per metric. Level-averaged results are in the main paper.

	Level 1				Level 2				Level 3			
DSLR	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$
Unprocessed	11.984	11.842	23.318	21.523	13.524	13.435	25.303	19.416	14.863	14.740	25.497	17.980
TI-Retinexformer + DWB	18.493	17.223	25.359	16.058	23.470	20.873	20.340	12.876	24.990	22.543	16.114	10.694
TI-Retinexformer + FC4	16.048	14.839	28.792	18.541	20.430	17.811	29.190	17.604	22.616	19.043	27.831	17.018
TI-Retinexformer + FFCC†	17.858	16.776	23.990	16.571	21.007	18.381	28.715	17.538	22.852	19.338	28.432	17.258
TI-Retinexformer + RL-AWB†	16.761	16.245	24.437	17.079	20.097	19.451	18.692	13.630	21.557	20.830	16.686	12.199
Retinexformer + DWB†	19.591	17.108	26.730	16.392	22.274	19.464	22.248	13.828	22.475	19.718	20.063	13.230
Retinexformer + FC4	15.853	14.657	32.633	18.989	20.477	17.966	30.889	17.763	22.476	19.190	28.615	17.087
Retinexformer + FFCC†	18.122	16.817	31.245	18.362	20.354	18.324	30.325	17.874	21.995	19.233	28.831	17.336
Retinexformer + RL-AWB	16.088	15.013	28.937	17.756	21.347	18.673	28.445	16.177	23.546	20.134	25.933	15.330
DarkIR + DWB	19.880	18.152	30.087	18.462	23.870	21.058	20.968	13.236	26.341	23.486	16.057	10.153
DarkIR + FC4†	16.179	15.912	37.630	22.442	18.884	18.258	27.533	17.502	19.561	18.767	25.074	16.437
DarkIR + FFCC†	18.551	17.452	32.709	22.091	21.554	18.575	31.198	18.860	22.405	19.571	29.061	17.349
DarkIR + RL-AWB†	17.664	17.233	27.645	19.048	20.625	19.952	19.302	13.567	21.683	20.957	15.027	11.534
MIRNet + DWB	20.938	18.176	30.293	17.160	24.741	21.377	21.340	12.527	26.883	23.413	16.451	10.188
MIRNet + FC4	16.566	15.158	33.910	18.918	19.825	17.702	31.553	17.713	22.137	19.145	28.617	16.800
MIRNet + FFCC†	18.242	16.859	33.236	18.396	20.191	18.132	31.128	17.694	21.830	19.191	28.830	17.061
MIRNet + RL-AWB	17.393	15.779	32.955	17.616	21.159	18.421	29.618	16.436	23.401	19.952	26.277	15.449
PromptNorm + DWB†	22.428	19.834	23.328	15.707	21.451	20.052	16.532	12.777	21.901	20.560	14.358	11.645
PromptNorm + FC4†	18.391	17.028	33.871	21.282	20.933	18.871	27.857	18.272	22.053	19.555	25.464	16.931
PromptNorm + FFCC†	22.173	18.749	30.536	20.008	24.287	19.294	27.975	18.457	25.137	19.585	27.076	17.838
PromptNorm + RL-AWB†	20.046	18.696	23.850	16.751	22.562	21.230	16.714	12.259	23.666	22.367	13.578	10.649
Ours (CTI-Retinexformer)	22.432	20.117	25.771	16.023	25.617	23.457	15.781	10.588	27.389	25.049	12.425	8.737

	Level 4				Level 5				Level 6			
DSLR	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$
Unprocessed	16.259	16.042	24.507	16.801	17.726	17.317	23.478	15.750	19.107	18.435	22.842	15.031
TI-Retinexformer + DWB	25.119	22.921	13.669	9.849	25.520	23.384	12.411	9.284	25.867	23.746	11.622	8.870
TI-Retinexformer + FC4	23.895	19.811	26.555	16.582	24.708	20.184	25.743	16.389	25.235	20.478	25.116	16.231
TI-Retinexformer + FFCC†	23.512	19.650	27.597	17.141	24.170	19.930	26.956	16.929	24.640	20.126	26.537	16.822
TI-Retinexformer + RL-AWB†	22.615	21.736	14.622	11.146	23.404	22.477	12.678	10.227	23.924	22.951	11.552	9.696
Retinexformer + DWB†	22.766	19.966	18.626	12.825	22.906	20.167	17.720	12.587	22.924	20.190	17.204	12.512
Retinexformer + FC4	23.620	20.030	26.982	16.516	24.430	20.544	25.973	16.211	25.021	20.880	25.197	16.011
Retinexformer + FFCC†	23.143	19.850	27.608	16.910	23.852	20.168	27.055	16.669	24.179	20.334	26.743	16.600
Retinexformer + RL-AWB	25.091	21.645	23.039	14.000	26.420	22.960	20.335	12.669	27.180	23.805	18.459	11.762
DarkIR + DWB	27.072	24.274	13.689	9.018	27.337	24.743	12.349	8.421	27.407	24.975	11.502	8.048
DarkIR + FC4†	20.386	19.361	24.111	15.906	21.109	19.733	23.951	15.762	21.280	19.782	24.048	15.845
DarkIR + FFCC†	23.782	20.257	27.724	16.829	24.738	20.605	26.914	16.581	24.930	20.764	26.436	16.419
DarkIR + RL-AWB†	23.068	22.070	13.006	10.522	24.055	22.988	11.323	9.592	24.570	23.428	10.371	9.119
MIRNet + DWB	26.969	24.163	13.779	9.141	26.990	24.442	12.306	8.639	27.415	24.906	11.486	8.117
MIRNet + FC4	23.639	20.170	26.620	16.170	24.639	20.695	25.518	15.920	25.297	21.051	24.809	15.778
MIRNet + FFCC†	23.261	19.909	27.539	16.814	23.814	20.192	26.989	16.674	24.282	20.448	26.515	16.531
MIRNet + RL-AWB	24.915	21.470	23.135	14.133	26.312	22.775	20.318	12.730	26.816	23.513	18.323	11.859
PromptNorm + DWB†	22.275	20.947	13.086	10.940	22.402	21.139	12.208	10.514	22.458	21.235	11.628	10.260
PromptNorm + FC4†	22.957	20.013	24.160	16.191	23.481	20.279	23.535	15.746	23.737	20.391	23.250	15.555
PromptNorm + FFCC†	25.691	19.842	26.425	17.486	25.935	19.965	26.078	17.276	26.066	20.025	25.831	17.161
PromptNorm + RL-AWB†	24.718	23.324	11.572	9.545	25.365	23.920	10.213	8.872	25.737	24.287	9.291	8.381
Ours (CTI-Retinexformer)	28.057	25.667	10.607	7.870	28.216	25.924	9.476	7.362	28.156	26.003	8.778	7.093

Table 7 – continued from the previous page

DSLR	Level 7				Level 8				Level 9			
	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$
Unprocessed	20.650	19.565	22.388	14.436	22.837	21.019	21.602	13.716	24.396	21.777	21.688	13.652
TI-Retinexformer + DWB	26.144	24.023	11.095	8.587	26.219	24.091	10.800	8.513	25.883	23.876	10.555	8.567
TI-Retinexformer + FC4	25.663	20.664	24.722	16.142	26.251	20.898	24.055	15.915	26.457	20.729	24.457	16.268
TI-Retinexformer + FFCC†	25.073	20.279	26.231	16.738	25.880	20.618	25.564	16.477	25.972	20.465	25.919	16.835
TI-Retinexformer + RL-AWB†	24.421	23.453	10.145	8.949	25.215	24.197	9.305	8.445	25.387	24.336	9.131	8.425
Retinexformer + DWB†	23.055	20.226	16.821	12.454	23.140	20.295	16.164	12.282	23.182	20.285	15.879	12.251
Retinexformer + FC4	25.470	21.116	24.778	15.894	26.040	21.446	24.040	15.600	26.257	21.350	24.447	15.936
Retinexformer + FFCC†	24.706	20.582	26.354	16.482	25.530	20.901	25.761	16.265	25.886	20.838	26.055	16.587
Retinexformer + RL-AWB	28.103	25.280	15.106	9.983	29.455	26.962	11.024	7.807	29.619	27.128	10.068	7.466
DarkIR + DWB	27.431	25.092	10.982	7.809	27.391	25.114	10.700	7.718	27.197	24.986	10.465	7.721
DarkIR + FC4†	21.483	19.871	24.177	15.941	22.035	20.260	23.943	15.737	22.073	20.086	24.603	16.146
DarkIR + FFCC†	25.092	20.880	26.107	16.324	25.941	21.184	25.440	16.048	26.130	21.007	25.865	16.407
DarkIR + RL-AWB†	25.084	23.855	9.568	8.635	25.888	24.600	8.939	8.168	25.906	24.635	8.754	8.154
MIRNet + DWB	27.588	25.126	10.972	7.828	27.463	25.065	10.723	7.793	27.443	25.069	10.506	7.742
MIRNet + FC4	25.765	21.254	24.441	15.714	26.287	21.501	23.775	15.509	26.373	21.354	24.222	15.849
MIRNet + FFCC†	24.672	20.630	26.027	16.432	25.662	20.956	25.528	16.223	25.948	20.815	25.914	16.572
MIRNet + RL-AWB	27.681	24.790	14.989	10.177	28.665	26.273	10.836	8.020	29.040	26.564	9.903	7.653
PromptNorm + DWB†	22.481	21.280	11.213	10.094	22.405	21.224	11.006	10.079	22.371	21.190	10.772	10.060
PromptNorm + FC4†	23.939	20.520	23.044	15.431	24.474	20.888	22.734	15.146	24.330	20.674	23.377	15.520
PromptNorm + FFCC†	26.197	20.068	25.623	17.081	26.692	20.223	25.075	16.830	26.762	20.085	25.570	17.178
PromptNorm + RL-AWB†	25.978	24.519	8.788	8.077	26.768	25.243	8.138	7.592	26.768	25.243	8.023	7.583
Ours (CTI-Retinexformer)	27.969	25.953	8.340	6.958	27.779	25.867	8.023	6.900	27.421	25.540	8.062	7.056

DSLR	Level 10			
	$PSNR_L \uparrow$	$PSNR_C \uparrow$	$MAE_\angle \downarrow$	$\Delta E_{00} \downarrow$
Unprocessed	26.631	22.830	20.891	13.061
TI-Retinexformer + DWB	26.125	24.160	10.413	8.387
TI-Retinexformer + FC4	26.694	21.072	23.793	15.810
TI-Retinexformer + FFCC†	26.396	20.865	25.080	16.297
TI-Retinexformer + RL-AWB†	25.919	24.801	8.664	7.985
Retinexformer + DWB†	22.975	20.188	15.935	12.383
Retinexformer + FC4	26.506	21.658	23.808	15.507
Retinexformer + FFCC†	26.347	21.196	25.327	16.112
Retinexformer + RL-AWB	30.048	27.509	9.600	7.181
DarkIR + DWB	27.014	24.960	10.328	7.670
DarkIR + FC4†	22.587	20.415	24.112	15.793
DarkIR + FFCC†	26.859	21.354	25.040	15.961
DarkIR + RL-AWB†	26.462	25.110	8.389	7.801
MIRNet + DWB	27.131	24.977	10.358	7.718
MIRNet + FC4	26.795	21.641	23.581	15.444
MIRNet + FFCC†	26.512	21.136	25.082	16.098
MIRNet + RL-AWB	29.124	26.734	9.423	7.411
PromptNorm + DWB†	22.190	21.105	10.660	10.062
PromptNorm + FC4†	24.658	20.917	22.882	15.173
PromptNorm + FFCC†	26.912	20.310	24.738	16.713
PromptNorm + RL-AWB†	27.147	25.571	7.724	7.297
Ours (CTI-Retinexformer)	27.326	25.600	7.906	6.943