arXiv:2611.05917v1 [cs.CV] 4 Nov 2026

DecoupAttn: Decoupled Sparse Attention for Training-Free Acceleration of Video Diffusion

Yanjin Li[1] Bingbing Qiu[1] Qiang Huang[2] Rui Qian[2]
Yanfeng Li[3] Jian Wang[2] Xiaoying Tang[4] Mingda Zhang[4,*]

[1]University of Illinois Urbana-Champaign [2]Fudan University [3]Macao Polytechnic University
[4]The Chinese University of Hong Kong, Shenzhen [*]Corresponding author: mingdazhang@ieee.org

Abstract

Sparse attention accelerates few-step video diffusion transformers by computing a few key blocks exactly and compensating for the rest. However, existing sparse attention still faces key challenges, including identification overhead *from finding critical blocks at every denoising step,* suboptimal budget allocation *that sets sparsity budgets uniformly or by local metrics, and* cross-step error accumulation *as reused masks or clusters go stale. Observing that the errors left on skipped blocks cancel each other once every key is compensated, we propose DecoupAttn, a new paradigm of decoupled sparse attention: semantic clusters compensate every key, and exact attention corrects their error on a few blocks fixed once offline by position. To support this paradigm, we introduce Distortion Budget, a rate–distortion objective for budget allocation that places exact computation across timesteps, layers, and heads where it removes the most distortion that compensation leaves in the final latent. Built on this budget, we further propose Reassign Cache, in which each step reassigns the drifting keys to cached centroids and refreshes compensation while the budgeted blocks stay fixed. Moreover, DecoupAttn is training-free: a few calibration clips compile one plan offline, and inference runs standard block-sparse kernels with one fused compensation kernel. Our code will be released after submission.*

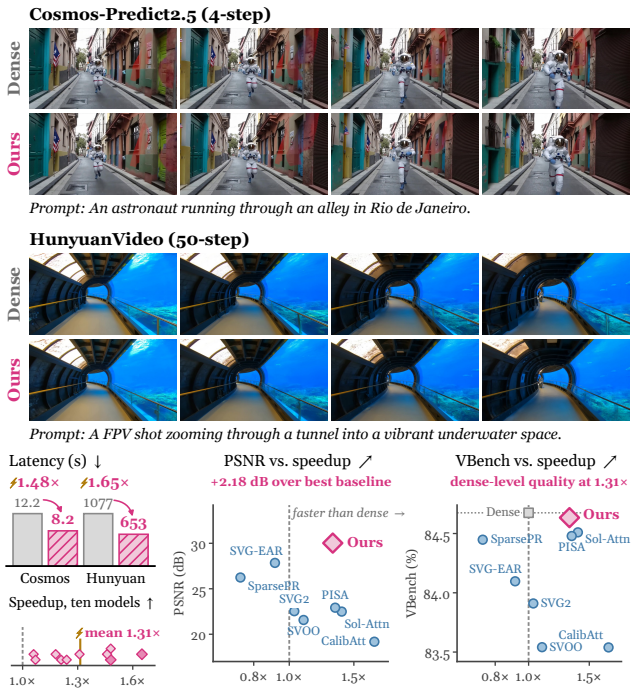

Figure 1. **DecoupAttn keeps dense-attention quality at a lower latency.** Frames from dense attention and DecoupAttn on Cosmos-Predict2.5 (4 steps) and HunyuanVideo (50 steps), with latency, speedup on all ten models (filled: these two), and PSNR and VBench against speedup for DecoupAttn and the seven baselines of Table 1, averaged over the ten models.

1. Introduction

In recent years, video diffusion transformers have been applied to a wide range of generation tasks, from text-to-video synthesis [34] to controlled generation [24] and neural rendering [15], gradually moving from many-step samplers toward few-step distilled generators [6, 44]. Sparse attention makes these generators fast without retraining: since attention grows quadratically with video length, it computes a few key blocks exactly and compensates for the rest, keeping near-dense fidelity (Figure 1). However, existing methods tie compensation to block selection: the groups that pick the exact blocks also summarize the skipped keys, so compensation is costly at every step of a few-step sampler and impossible once the selection is fixed offline.

Sparse attention for video generation has followed three main lines of work, as shown in Figure 2. First, online routing (dynamic sparse attention) identifies critical blocks at every step from the current queries and keys, by scoring pooled blocks [39, 45] or clustering tokens by semantics [33, 41], treating sparsity as a property of each input. Second, offline calibration (static sparse attention) fixes sparse patterns before deployment, measured on a few prompts [43] or derived from how attention decays with distance [13], treating sparsity as a property of the model. Third, compensation approximates the skipped blocks instead of dropping them, from cluster centroids [53], blockwise Taylor expansions [12], or fitted residuals [32], recov-

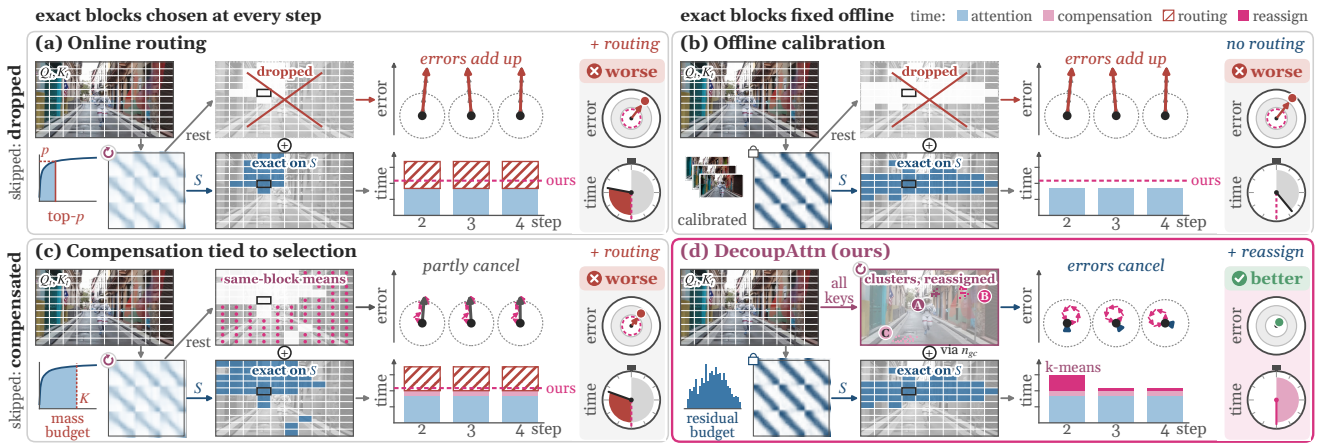

Figure 2. **Three lines of sparse attention and DecoupAttn.** (a) Online routing re-picks exact blocks each step and drops the rest, (b) offline calibration fixes them but drops the rest, and (c) compensation reuses them and reroutes every step. (d) DecoupAttn fixes them offline and compensates every key by semantic clusters. Right: schematic error and time against DecoupAttn (dashed), excess in red.

ering much of the attention that selection leaves out.

However, these methods still face three challenges. **(i) Identification overhead.** Online routing and the compensation built on its groups identify critical blocks at every denoising step [11, 12, 17, 32, 33, 37, 39–41, 45, 48, 49, 53], an overhead that can offset the gains of sparsity in few-step samplers, while static patterns fixed in advance [13, 43, 46] discard the remaining blocks and incur information loss. **(ii) Suboptimal budget allocation.** Existing methods set sparsity budgets across heads, layers, and timesteps uniformly or by local metrics such as retained attention mass [20, 41, 43, 48, 53], or by the effect of dropping blocks on the output [38, 51, 52]. Retained attention mass, however, does not determine the output error [32], which depends on the residual left after compensation and on how unevenly later layers and steps amplify it, so exact computation is spent where it barely changes the final video. **(iii) Cross-step error accumulation.** To avoid redundant recomputation, recent methods reuse masks, clusters, or earlier centroids across steps, or predict masks ahead [2, 10, 16, 18, 20, 29, 40, 41], but attention patterns evolve across denoising steps, so the reused groups go stale and their errors accumulate, while refreshing them repeats the identification.

Our key observation is that once every key is compensated, the residuals of the skipped blocks point in many directions and cancel each other like dithered quantization errors [27], so blocks fixed once offline match blocks reselected at every step. Based on this observation, we propose DecoupAttn, a new paradigm of decoupled sparse attention: semantic clusters compensate every key, and exact attention corrects their residual on a few blocks fixed once offline by position (Figure 2d). Selecting blocks costs nothing at inference, and the residual of every fixed block can be priced offline. Few-step samplers amplify this residual unevenly, so we introduce Distortion Budget, a rate–distortion objective for budget allocation that prices every timestep, layer, and head by the distortion the residual leaves in the final latent. Built on this budget, we further propose Reassign Cache, which limits the growth of these residuals by reassigning the drifting keys to cached centroids at every step, while the blocks and the budget stay fixed. DecoupAttn thus moves block identification and budget allocation into one plan compiled offline with no training, and inference only refreshes the clusters, running standard block-sparse kernels and one fused compensation kernel.

We evaluate DecoupAttn on ten video generation, controlled generation, and neural rendering models, where it outperforms online-routing, offline-calibrated, and compensation-based methods on nine of them and delivers a substantial speedup over dense attention.

2. Related Work

Summarizing skipped attention. Transformers summarize attention by clusters, near- and far-field splits, or sparse-plus-low-rank estimators [1, 14, 21, 25, 35, 50], and video sparse attention compensates the skipped keys by cluster or block means [17, 53], Taylor terms [12, 48], or fitted or cached residuals [2, 32, 40]. Masks frozen early or predicted ahead [16, 18] drop the keys outside them, and routers that reuse centroids [10, 20, 41] still let the clusters pick the blocks. DecoupAttn summarizes the skipped keys by semantic clusters that never pick blocks.

Allocating computation by sensitivity. Coders and quantizers allocate bits by rate–distortion cost or layer sensitivity [3, 30, 31], and sparse attention sets density by retained attention mass, profiling, sampler curvature, or block error [20, 41, 43, 48, 53], or by each unit's effect on the

output [38, 51, 52]. Distortion Budget instead prices the residual left after compensation, not the error of dropping a block, by its effect on the final latent.

3. Preliminaries

Definition 1: Block-sparse attention call. A self-attention call at denoising step t, layer l, and head h maps the queries q_i, keys k_j, and values v_j of its L tokens to outputs y_i, with the tokens partitioned by position into query tiles g and key blocks b, the units of block-sparse kernels [42]:

$$y_i = \frac{N_i}{Z_i} = \frac{\sum_b N_{ib}}{\sum_b Z_{ib}}, \quad N_{ib} = \sum_{j \in b} e^{s_{ij}} v_j, \quad Z_{ib} = \sum_{j \in b} e^{s_{ij}}, \tag{1}$$

where $s_{ij} = q_i^\top k_j / \sqrt{d}$ and d is the head dimension. A sparse call computes N_{ib} and Z_{ib} exactly on a block set $\mathcal{S}_g$ and summarizes the rest. Existing compensation writes $\hat{N}_i = \sum_{b \in \mathcal{S}_g} N_{ib} + \sum_{b \notin \mathcal{S}_g} \tilde{N}_{ib}$ with $\tilde{N}_{ib}$ from the statistics of block b itself, and every change of $\mathcal{S}_g$ rebuilds the compensation terms [12, 53].

Definition 2: Sparse attention configuration. A sampler runs T denoising steps. Each head of each call is a unit $u = (t, l, h)$ of the unit set $\mathcal{U}$, and the heads of a call (t, l) share one compensation estimator. A configuration, or plan, $\pi = (\{o_u\}_{u \in \mathcal{U}}, \{\eta_{tl}\})$ assigns every unit an option and every call an estimator, which determines the final latent and the density:

$$\pi \;\Rightarrow\; \hat{x}_0(\pi), \qquad X(\pi) = \frac{\sum_u w_u X_u(o_u)}{\sum_u w_u}, \tag{2}$$

where $o_u \in \mathcal{O}$ is the unit's option (Eq. (8)), $\eta_{tl} \in \{\textsc{Full}, \textsc{Centroid}\}$ the call's estimator (Eq. (4)), X_u the density of an option, the fraction of blocks it executes with every block padded to the kernel block size and the DENSE option counted as one, and w_u the number of sampler steps the unit represents. The first step and the first layer always stay dense and lie outside $\mathcal{U}$, and the estimator adds a fixed overhead to X (Appendix A.9).

Problem statement. Given a pretrained video diffusion transformer with frozen weights and a few calibration clips, we seek a plan that is fixed before deployment, that keeps the density within a target ρ, and that keeps the final latent closest to the dense one:

$$\pi^\star = \arg\min_{\pi} \; \mathbb{E}_x \big\| \hat{x}_0(\pi) - x_0 \big\|_2^2 \;\text{ s.t. }\; X(\pi) \le \rho, \tag{3}$$

where x is the input (prompt, conditions, and seed) and x_0 the dense final latent for the same input.

4. Method

As Figure 3 shows, DecoupAttn centers on one quantity, the residual that compensation leaves on each key block: the decoupled estimator makes this residual the only error that reaches the output and keeps it small (Section 4.1), Distortion Budget spends exact computation where the residual most distorts the final latent (Section 4.2), and Reassign Cache limits its growth as keys drift (Section 4.3).

4.1. Decoupled Sparse Attention

To keep compensation when the selection is fixed offline, decoupled sparse attention follows the difference estimators of survey sampling [26]: a cheap algorithm predicts the contribution of every key, and exact attention on a few blocks replaces the prediction there (Figure 3c).

Centroid compensation with two corrections. At the first sparse step, the keys of every head are clustered by k-means into clusters c, and every later call summarizes each cluster by its centroid with two second-order corrections computed from the cluster's current members:

$$\omega_{ic} = \phi_{ic}\,(1 + \gamma_{ic}), \qquad a_{ic} = \omega_{ic}\,\bar{v}_c + \phi_{ic}\,\frac{H q_i}{\sqrt{d}}, \tag{4}$$

where $\phi_{ic} = e^{q_i^\top \bar{k}_c / \sqrt{d}}$ is the centroid weight, and $\bar{k}_c$, $\bar{v}_c$, and σ_c^2 are the cluster's mean key, mean value, and per-dimension key variance. The *variance correction* $\gamma_{ic} = \frac{1}{2d} \sum_p q_{ip}^2 \sigma_{cp}^2$ is the diagonal part of the second-order term of $e^{q_i^\top k / \sqrt{d}}$ around $\bar{k}_c$ and restores this part of the mass that a centroid misses on a spread-out cluster (Lemmas A.3 and A.4). The within-cluster cross-covariance $H = \frac{1}{L}\big(\sum_j v_j k_j^\top - \sum_c |c|\, \bar{v}_c \bar{k}_c^\top\big)$, pooled over the clusters of a head, gives their shared *first-order value correction* (Lemma A.5). The CENTROID estimator drops both corrections.

Decoupled estimator. Each query tile g computes its exact blocks $\mathcal{S}_g$ in full and covers every other key with its cluster centroid, and both parts share one softmax denominator:

$$\hat{y}_i = \frac{\hat{N}_i}{\hat{Z}_i} = \frac{\sum_{b \in \mathcal{S}_g} N_{ib} + \sum_c n_{gc}\, a_{ic}}{\sum_{b \in \mathcal{S}_g} Z_{ib} + \sum_c n_{gc}\, \omega_{ic}}, \tag{5}$$

where a_{ic} and ω_{ic} are what one member of cluster c contributes to the numerator and the denominator (Eq. (4)). The count $n_{gc} = |c| - \sum_{b \in \mathcal{S}_g} n_{bc}$, with $n_{bc} = |b \cap c|$, keeps only the members of c that the exact blocks do not already cover. The exact blocks are fixed offline from position and calibration statistics, the clusters come from the keys alone, and the counts n_{gc} are the only place the two meet, so no block identification runs at inference.

Compensate every key, correct some blocks. Substituting n_{gc} into Eq. (5) splits the estimate into two steps. First, the centroids compensate all keys as if no block were exact. Then, on each exact block, the exact value replaces the

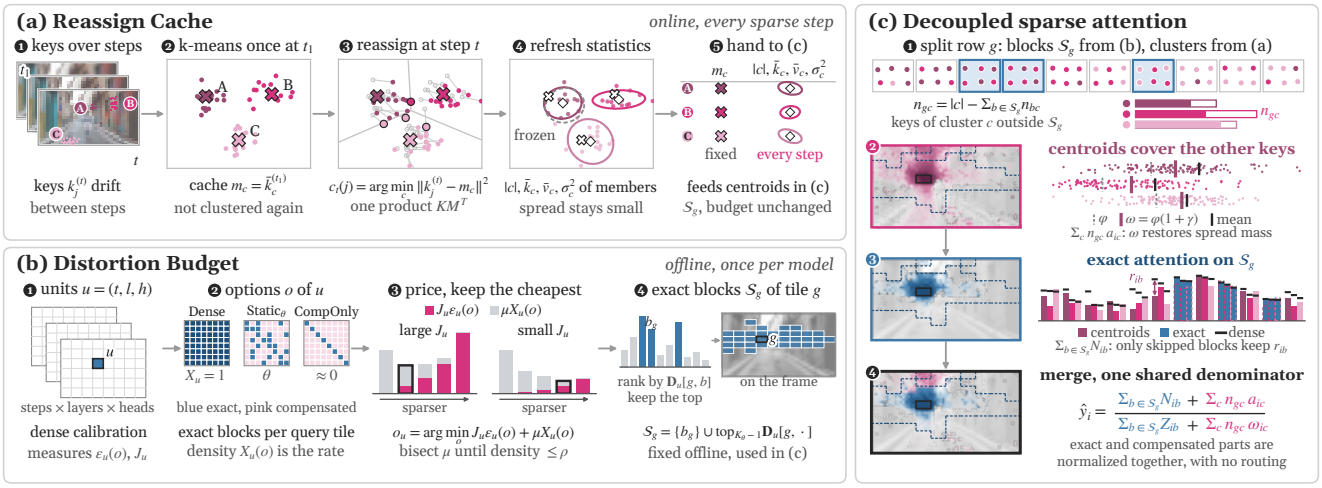

Figure 3. **How DecoupAttn computes one attention call** (Cosmos-Predict2.5). (a) Reassign Cache clusters the keys once, caches the centroids, and at every later step moves each key to its nearest cached centroid and refreshes the cluster statistics, (b) Distortion Budget prices the options of every unit offline and keeps the cheapest, which fixes the exact blocks $\mathcal{S}_g$, and (c) exact attention on $\mathcal{S}_g$ and centroid compensation of every other key meet only through the counts n_{gc} and share one softmax denominator (pink: approximation, blue: exact).

centroids' prediction, as in Scatterbrain and EVA [1, 50]:

$$\hat{N}_i = \sum_c |c|\, a_{ic} + \sum_{b \in \mathcal{S}_g} \Delta N_{ib},$$
$$\hat{y}_i - y_i = -\frac{1}{\hat{Z}_i} \sum_{b \notin \mathcal{S}_g} \underbrace{\left(\Delta N_{ib} - y_i\, \Delta Z_{ib}\right)}_{r_{ib}}, \tag{6}$$

where $\Delta N_{ib} = N_{ib} - \sum_c n_{bc} a_{ic}$ and $\Delta Z_{ib} = Z_{ib} - \sum_c n_{bc} \omega_{ic}$ are the gaps between exact value and centroid prediction on block b, the residual r_{ib} is the error block b passes to the output, and skipped blocks leave error.

Why fixed blocks suffice. By Eq. (6), the error is a sum of block residuals, so its size depends on the estimated mass, on how large the residuals are, and on how they align:

$$\|\hat{y}_i - y_i\|_2^2 = \frac{1}{\hat{Z}_i^2} \Big\| \sum_{b \notin \mathcal{S}_g} r_{ib} \Big\|_2^2 = \frac{\kappa_i R_i}{\hat{Z}_i^2}, \tag{7}$$

where $R_i = \sum_{b \notin \mathcal{S}_g} \|r_{ib}\|_2^2$ is the residual energy of the skipped blocks, and the coherence ratio $\kappa_i = \|\sum_{b \notin \mathcal{S}_g} r_{ib}\|_2^2 / R_i$, set to one when $R_i = 0$, is one for mutually orthogonal residuals and grows as they align. Dropped blocks each remove their whole attention and push the output the same way, whereas a compensated residual comes from its members' deviations from their centroids, which sum to zero over the cluster, and the residuals of the skipped blocks cancel each other like dithered quantization errors [27] (Theorem A.10). With $\hat{Z}_i \approx Z_i$ (Corollary A.2) and κ_i kept small, the residual energy R_i is what Sections 4.2 and 4.3 then control.

Proposition 1. *For count-based compensation, decoupled sparse attention equals compensating every key and correcting the residual of each exact block, and its error is determined by the summed residual of the other blocks for any blocks and clusters, while for blocks fixed offline the deviation parts of these residuals add like random errors and nearly cancel each other.* (Proof in Appendices A.2 and A.4, evidence in Section 5.3.)

4.2. Distortion Budget

Section 4.1 leaves the residual energy R_i to control, and by Proposition 1 a block's residual does not depend on which other blocks are exact (Appendix A.2). Distortion Budget measures each residual once on the calibration set and solves Eq. (3) as a rate–distortion bit allocation [30]: units act as quantizers, options as their resolutions, density as the rate, and final-latent error as the distortion (Figure 3b).

Options. Each unit $u = (t, l, h)$ (Definition 2) picks one of three kinds of option, which differ in how many key blocks it computes exactly. DENSE computes every block, as dense attention does. STATIC$_\theta$ computes the fraction θ of key blocks that rank highest for the query tile and compensates the rest. COMPONLY computes only the diagonal block and compensates everything else, the analog of the skip mode in rate–distortion mode decision [31]. The option set and the exact blocks of STATIC$_\theta$ are

$$\mathcal{O} = \{\text{DENSE}, \text{COMPONLY}\} \cup \{\text{STATIC}_\theta\}_{\theta \in \Theta},$$
$$\mathcal{S}_g(\text{STATIC}_\theta) = \{b_g\} \cup \text{top}_{K_\theta - 1}\big(\mathbf{D}_u[g, \cdot]\big), \tag{8}$$

where Θ is the set of density levels (Appendix B), $K_\theta = \max(1, \lfloor \theta N_B \rceil)$ for N_B key blocks, $\mathbf{D}_u[g, \cdot]$ the ranking of the other key blocks for query tile g (Eq. (9)), and b_g its diagonal block, which every option computes. The ranking orders these blocks by how badly compensation handles

them, measured by their residual in Eq. (6) on the calibration set with the model run densely:

$$\mathbf{D}_u[g,b] = \mathbb{E}_x\,\mathbb{E}_{i\in g}\,\frac{\|r_{ib}\|_2^2}{Z_i^2\,\|y_i\|_2^2}, \tag{9}$$

where r_{ib} is taken under full compensation, and dividing by $Z_i^2\|y_i\|_2^2$ puts large and small outputs on one scale.

Distortion of an option. Distortion Budget measures, unit by unit on dense calibration runs, the projected error of the head under every option o and estimator η:

$$\begin{aligned}\varepsilon_u(o,\eta) &= \mathbb{E}_x\big\langle M_h,\ (\hat{O}_u - O_u)^\top(\hat{O}_u - O_u)\big\rangle,\\ M_h &= W_{o,h}^\top W_{o,h},\end{aligned} \tag{10}$$

where O_u and $\hat{O}_u$ are the head's dense output and its output under o and η (Eq. (5)), the expectation runs over calibration inputs x, the inner product is the squared error after projection, $\|(\hat{O}_u - O_u)W_{o,h}^\top\|_F^2$, and $W_{o,h}$ is the head's slice of the output projection into the residual stream.

Sensitivity. The distortion of each unit is further scaled by a sensitivity J_u, the squared gain from the head to the final latent through the remaining layers and sampler steps (Appendix A.10), which plays the role of Hessian sensitivity in mixed-precision quantization [3]. To reduce the number of measurements, units are measured in groups $\mathcal{G}$, each covering a range of steps and of layers (Appendix B): the units of $\mathcal{G}$ run sparse at a reference level θ_0 with full compensation while the rest stays dense, so the injected errors are the real sparse errors, and the extra final-latent error is divided by the group's summed distortion:

$$J_u = \frac{\mathrm{MSE}_{\mathcal{G}} - \mathrm{MSE}_0}{\sum_{u'\in\mathcal{G}} w_{u'}\,\varepsilon_{u'}(\textsc{Static}_{\theta_0}, \textsc{Full})}, \tag{11}$$

where $\mathrm{MSE}_{\mathcal{G}}$ is the final-latent error of this run against the dense output, and the subtracted MSE_0 is the run-to-run error of the same dense attention backend. This reference gain prices every option of the unit (Lemma A.17).

Allocation. With a price $J_u\varepsilon_u$ for every unit and option, choosing the plan becomes a budgeting problem under the density target ρ:

$$\begin{aligned}&\min_{\pi}\sum_u w_u J_u\,\varepsilon_u(o_u,\eta_{tl}) \quad \text{s.t.} \quad X(\pi)\le\rho,\\ &\mathcal{L}_\mu(\pi) = \sum_u w_u\big(J_u\,\varepsilon_u(o_u,\eta_{tl}) + \mu X_u(o_u)\big),\end{aligned} \tag{12}$$

where the first line approximates Eq. (3) (Proposition 2). The second line is the Lagrangian relaxation of Everett [4], which splits the search over plans into one choice per call: under each estimator every head takes the option with the lowest $J_u\varepsilon_u + \mu X_u$, and the call keeps the lower-total estimator. Bisection on μ finds the densest plan that still fits ρ, and a final greedy pass spends the density left over.

Proposition 2. *The summed priced distortion approximates the final-latent error, and for every nonnegative multiplier, the plan that minimizes $\mathcal{L}_\mu$ has the lowest priced distortion among all plans that are no denser, so bisection returns the best plan at its own density.* (Proof in Appendix A.7, evidence in Section 5.3.)

4.3. Reassign Cache

The plan prices the residuals of clusters built at the first sparse step, and a cluster's residual is of second order in the spread of its members around the centroid (Proposition A.6). As keys drift between denoising steps, the spread grows and the residual outgrows what the plan priced. Reassign Cache bounds this spread at every step at little cost (Figure 3a).

Cached centroids. We treat the centroids of the first sparse step as a fixed codebook and, at every later step, run the assignment step of Lloyd's algorithm [19]: each key moves to the cluster of its nearest cached centroid, and the statistics of Eq. (4) are recomputed from the current members:

$$c_t(j) = \arg\min_c \big\|k_j^{(t)} - m_c\big\|_2^2, \qquad m_c = \bar{k}_c^{(t_1)}, \tag{13}$$

where t_1 is the first sparse step, $k_j^{(t)}$ is key j at step t, and m_c is the cached centroid of cluster c. One call costs one key–centroid product and an argmin (Appendix A.9).

Assignment rule. One measurement on the calibration set decides whether reassignment pays:

$$\text{reassign} = \mathbf{1}\big[\delta_{\text{stale}} > \tau_s\big], \tag{14}$$

where $\mathbf{1}[\cdot]$ is the indicator function, δ_{stale} is the PSNR lost when the cluster assignments of the first sparse step stay frozen, and τ_s is a threshold (Appendix B).

Proposition 3. *Reassigning the keys at a step gives the lowest total squared distance to the cached centroids of any assignment, and it changes only the compensation: the exact blocks and every unit's budget stay as planned.* (Proof in Appendix A.8, evidence in Section 5.3.)

5. Experiments

We ask whether DecoupAttn is more accurate than prior sparse attention at equal density yet faster than dense attention (RQ1), whether one procedure fits all families and samplers (RQ2), what each component adds (RQ3), the quality–speed trade-off (RQ4), and what the plan costs (RQ5).

5.1. Experimental Setup

Models. We test ten models in three families (Table 5). Video generation covers Cosmos-Predict1 (CP1) [23], Cosmos-Predict2.5 (CP2.5) [24], Wan2.1-14B at 480p (Wan2.1) and Wan2.2-A14B (Wan2.2) [34], HunyuanVideo (HY) [9], and image-to-video Cosmos3-Nano (Nano) [22].

Model	Floor		CalibAtt		SVOO		SVG2		SVG-EAR		Sol-Attn		PISA		SparsePR		**Ours**		Δ	
(steps)	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS	PSNR SSIM	× LPIPS
CP1	32.66 ±7.17	1.00×	19.35 ±2.46	1.73×	21.27 ±3.70	1.35×	22.03 ±3.51	1.34×	24.32 ±4.27	1.18×	21.87 ±3.63	1.47×	22.27 ±4.05	1.43×	24.51 ±4.31	1.06×	**27.77** ±4.60	1.48×	+5.54	+0.11×
(35)	0.897 ±0.100	0.072 ±0.071	0.668 ±0.127	0.397 ±0.074	0.717 ±0.136	0.297 ±0.097	0.732 ±0.136	0.266 ±0.099	0.765 ±0.130	0.181 ±0.080	0.732 ±0.133	0.252 ±0.092	0.724 ±0.143	0.260 ±0.099	0.766 ±0.131	0.183 ±0.089	**0.835** ±0.090	**0.121** ±0.056	+0.106	−0.141
CP2.5	27.78 ±3.90	1.00×	12.35 ±1.48	2.44×	17.78 ±2.57	1.08×	19.74 ±2.85	0.98×	21.85 ±3.32	0.90×	17.95 ±2.65	1.50×	17.68 ±2.48	1.40×	20.97 ±3.00	0.60×	**24.70** ±3.78	1.48×	+6.37	+0.21×
(4)	0.921 ±0.053	0.054 ±0.030	0.400 ±0.107	0.584 ±0.071	0.594 ±0.127	0.381 ±0.099	0.693 ±0.104	0.225 ±0.065	0.781 ±0.096	0.157 ±0.066	0.600 ±0.128	0.314 ±0.067	0.615 ±0.116	0.313 ±0.070	0.744 ±0.092	0.186 ±0.057	**0.855** ±0.077	**0.102** ±0.046	+0.223	−0.207
Wan2.1	24.93 ±3.03	1.00×	10.56 ±1.31	1.54×	19.27 ±1.56	0.81×	18.56 ±1.56	0.72×	22.95 ±2.16	0.67×	18.32 ±1.63	1.31×	17.87 ±1.50	1.16×	20.29 ±1.81	0.40×	**24.70** ±2.29	1.06×	+6.44	+0.12×
(4)	0.902 ±0.049	0.050 ±0.028	0.312 ±0.074	0.589 ±0.067	0.726 ±0.069	0.150 ±0.036	0.696 ±0.073	0.171 ±0.043	0.858 ±0.054	0.074 ±0.027	0.685 ±0.070	0.175 ±0.038	0.669 ±0.071	0.191 ±0.037	0.773 ±0.066	0.122 ±0.034	**0.889** ±0.044	**0.057** ±0.022	+0.215	−0.153
Wan2.2	21.10 ±3.91	1.00×	11.73 ±2.20	1.39×	16.53 ±2.29	0.81×	16.09 ±2.12	0.71×	18.12 ±2.64	0.66×	15.03 ±1.86	1.31×	15.38 ±2.03	1.26×	17.09 ±2.35	0.40×	**18.43** ±2.66	1.07×	+2.72	+0.14×
(4)	0.749 ±0.135	0.132 ±0.073	0.304 ±0.118	0.576 ±0.090	0.521 ±0.139	0.296 ±0.086	0.496 ±0.136	0.314 ±0.095	0.614 ±0.142	0.218 ±0.085	0.452 ±0.127	0.361 ±0.073	0.468 ±0.131	0.340 ±0.080	0.559 ±0.134	0.260 ±0.079	**0.634** ±0.130	**0.202** ±0.071	+0.146	−0.136
HY	37.13 ±11.63	1.00×	17.26 ±1.82	1.90×	20.14 ±3.28	1.66×	19.98 ±3.54	1.60×	22.06 ±5.05	1.37×	17.44 ±2.66	1.71×	17.24 ±2.37	1.54×	23.68 ±5.11	1.36×	**27.47** ±5.84	1.65×	+7.78	+0.06×
(50)	0.929 ±0.142	0.073 ±0.168	0.664 ±0.109	0.428 ±0.111	0.755 ±0.123	0.287 ±0.141	0.742 ±0.118	0.279 ±0.158	0.820 ±0.110	0.222 ±0.137	0.698 ±0.129	0.356 ±0.147	0.685 ±0.129	0.375 ±0.133	0.820 ±0.124	0.183 ±0.122	**0.881** ±0.102	**0.130** ±0.101	+0.140	−0.174
Nano	25.34 ±6.20	1.00×	15.46 ±2.53	1.54×	14.58 ±1.68	1.16×	15.51 ±2.28	1.12×	17.20 ±3.19	0.97×	16.61 ±2.56	1.28×	17.42 ±4.11	1.28×	18.21 ±3.45	0.94×	**21.56** ±5.04	1.21×	+5.13	+0.03×
(35)	0.832 ±0.117	0.091 ±0.075	0.494 ±0.128	0.398 ±0.115	0.423 ±0.110	0.498 ±0.078	0.473 ±0.121	0.427 ±0.097	0.535 ±0.144	0.356 ±0.087	0.534 ±0.122	0.334 ±0.099	0.590 ±0.172	0.291 ±0.146	0.609 ±0.166	0.271 ±0.123	**0.731** ±0.146	**0.154** ±0.093	+0.208	−0.214
T2.5	41.93 ±4.94	1.00×	23.27 ±1.29	1.72×	25.20 ±1.83	1.08×	26.85 ±2.05	0.99×	32.34 ±2.46	0.92×	26.99 ±1.93	1.49×	25.63 ±1.58	1.40×	30.94 ±3.37	0.60×	**33.81** ±2.87	1.46×	+6.49	+0.29×
(4)	0.988 ±0.010	0.014 ±0.013	0.808 ±0.048	0.281 ±0.057	0.823 ±0.049	0.261 ±0.077	0.860 ±0.042	0.174 ±0.058	0.942 ±0.025	0.067 ±0.034	0.874 ±0.037	0.161 ±0.052	0.858 ±0.042	0.208 ±0.056	0.929 ±0.028	0.082 ±0.039	**0.958** ±0.020	**0.050** ±0.025	+0.087	−0.126
DR	54.70 ±1.61	1.00×	31.95 ±7.85	1.68×	32.91 ±3.43	0.73×	36.23 ±3.61	0.72×	48.79 ±2.42	0.71×	38.58 ±2.69	1.21×	43.11 ±3.16	1.23×	45.40 ±2.85	0.43×	**50.87** ±1.73	1.18×	+11.30	+0.22×
(4)	0.998 ±0.001	0.001 ±0.001	0.968 ±0.055	0.033 ±0.047	0.973 ±0.019	0.038 ±0.022	0.984 ±0.011	0.029 ±0.020	0.995 ±0.001	0.003 ±0.001	0.989 ±0.008	0.011 ±0.004	0.994 ±0.002	0.005 ±0.003	0.995 ±0.002	0.004 ±0.002	**0.997** ±0.001	**0.002** ±0.001	+0.012	−0.016
Alaya	39.13 ±3.56	1.00×	18.87 ±2.89	1.57×	19.24 ±2.69	1.05×	20.11 ±2.99	0.98×	26.36 ±4.61	0.77×	18.64 ±2.97	1.29×	18.67 ±2.87	1.31×	20.45 ±4.08	0.67×	**28.15** ±3.68	1.24×	+7.82	+0.15×
(50)	0.969 ±0.014	0.012 ±0.007	0.636 ±0.126	0.433 ±0.120	0.673 ±0.115	0.369 ±0.117	0.733 ±0.106	0.262 ±0.115	**0.871** ±0.058	0.096 ±0.052	0.705 ±0.102	0.302 ±0.093	0.688 ±0.102	0.327 ±0.091	0.728 ±0.161	0.274 ±0.173	0.868 ±0.070	**0.091** ±0.055	+0.149	−0.204
UniRel.	55.49 ±1.38	1.00×	30.99 ±4.89	1.55×	28.76 ±3.52	1.21×	30.25 ±3.89	1.16×	**44.45** ±4.46	0.99×	33.47 ±5.44	1.33×	33.93 ±5.31	1.33×	40.84 ±4.39	0.90×	42.75 ±4.21	1.31×	+8.08	+0.10×
(35)	0.998 ±0.001	0.001 ±0.000	0.956 ±0.023	0.069 ±0.036	0.954 ±0.064	0.055 ±0.023	0.961 ±0.074	0.041 ±0.040	0.993 ±0.010	0.008 ±0.024	0.967 ±0.056	0.042 ±0.071	0.973 ±0.027	0.036 ±0.037	0.992 ±0.006	0.008 ±0.008	**0.994** ±0.002	**0.006** ±0.004	+0.023	−0.031
Avg.	36.02	1.00×	19.18	1.71×	21.57	1.09×	22.54	1.03×	27.84	0.91×	22.49	1.39×	22.92	1.33×	26.24	0.74×	**30.02**	1.31×	+6.77	+0.14×
	0.918	0.050	0.621	0.379	0.716	0.263	0.737	0.219	0.817	0.138	0.724	0.231	0.726	0.235	0.792	0.157	**0.864**	**0.092**	+0.131	−0.140

Table 1. **Quality and speed at density 0.20.** Each model shows PSNR (dB) and speedup over the fastest dense attention backend (top) and SSIM and LPIPS (bottom) against same-seed dense outputs, with test-prompt standard deviations below. Floor compares two dense attention backends, bold marks the best sparse method, and Δ is DecoupAttn minus the baseline mean (Section 5.1).

Component	Variant	CP1	CP2.5	Wan2.1	Wan2.2	HY	Nano	T2.5	DR	Alaya	UniRel.	Avg.	Supp.
	Full method	23.42	24.07	25.31	17.21	25.84	20.92	33.34	50.20	27.97	41.54	28.98	
	ms	171,006	8,212	12,940	12,820	654,610	188,278	9,620	3,905	46,057	108,377	121,583	
Decoupling (§4.1)	Tied: blocks as groups	−1.95	−1.63	−3.72	−1.48	−1.61	−1.71	−2.21	−2.75	−4.07	−7.39	−2.85	10/10
	ms	−10,326	−725	−1,181	−1,147	−27,949	−10,911	−804	−396	−3,673	−6,113	−6,323	
	Tied†: same, our plan	−2.45	−1.90	−4.19	−1.66	−0.85	−2.17	−2.67	−5.29	−5.43	−8.59	−3.52	10/10
	ms	−9,592	−707	−1,119	−1,146	−29,257	−10,610	−802	−423	−3,482	−6,308	−6,345	
	Centroid: no corrections	−1.60	−1.03	−1.24	−0.53	−1.06	−0.60	−1.20	−2.48	−1.43	−3.77	−1.49	10/10
	ms	−2,621	−95	−176	−148	−23,626	−6,768	−75	+121	−1,154	−1,885	−3,643	
Distortion Budget (§4.2)	U: every $J_u{=}1$	−1.93	−2.30	−2.14	−1.06	−4.10	−2.05	−0.55	−0.77	−1.06	−0.85	−1.68	9/10
	ms	+1,814	+71	+38	−10	+13,275	+1,473	+51	−15	−129	+1,367	+1,794	
	D: uncompensated ε_u	−0.34	−0.96	−1.65	−0.65	−0.87	−1.06	−1.31	−1.71	−0.86	−2.41	−1.18	9/10
	ms	−1,738	+21	+19	+24	+2,311	−171	+3	−6	+37	+284	+78	
Reassign Cache (§4.3)	Frozen: no reassignment	−0.95	−0.38	−1.14	−0.36	−0.38	−0.12	−0.69	–	−2.59	−4.45	−1.23	8/9
	ms	−7,484	−150	−218	−255	−26,463	−8,050	−167	–	−2,850	−4,780	−5,602	

Table 2. **Component ablations (RQ3).** Variant minus the full method (first row) on eight validation clips per model, in PSNR (dB) and latency (ms rows). Supp. counts the models where the variant loses at least 0.3 dB on average and on at least 6 of 8 clips (Table 6).

Controlled generation uses edge-conditioned Cosmos-Transfer2.5 (T2.5) [24], and neural rendering uses DiffusionRenderer (DR) [15], AlayaRenderer (Alaya) [8], and UniRelight (UniRel.) [5]. Five run four-step distilled samplers and five their original many-step samplers. Each model has disjoint calibration (at most 16 clips), validation (8), and test (8–39) splits, all with seed 1000.

Baselines. We compare with seven methods from three lines of work, all run from their official code: offline calibration (CalibAtt [43]), online routing by clustering (SVOO [20], SVG2 [41]) or by block scoring (Sol-Attn [11]), and compensation-based sparse attention (SVG-EAR [53], PISA [12], SparsePR [32]). For a fair comparison, every method keeps the first step and layer dense,

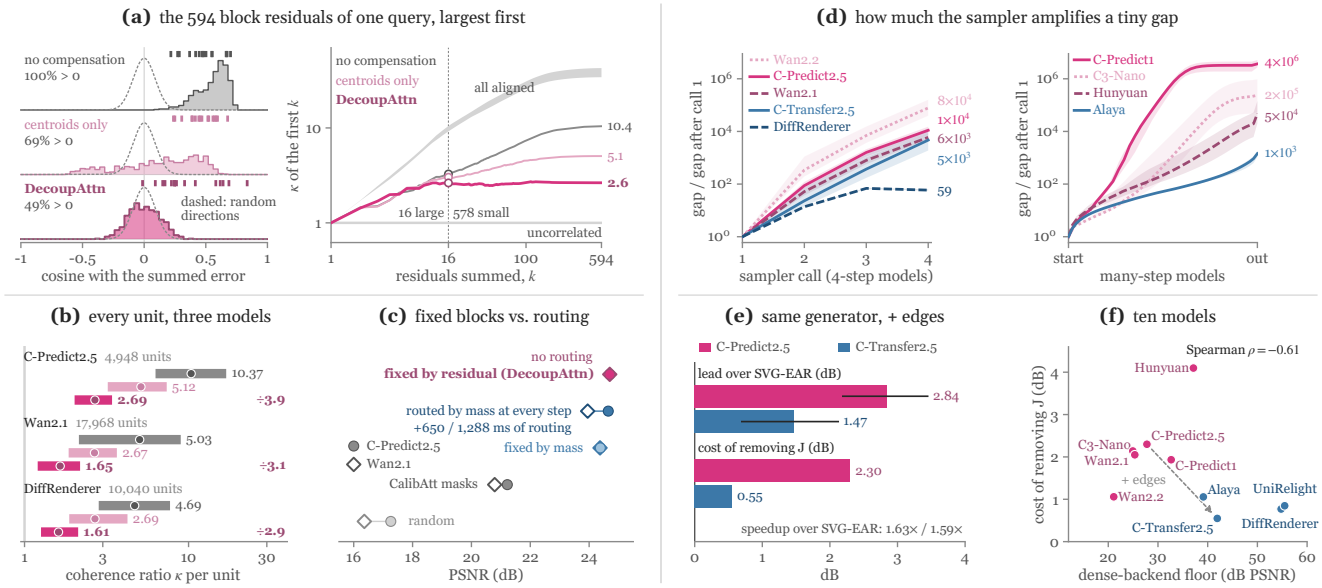

Figure 4. **Compensated block residuals cancel like dithered quantization errors, and conditioning damps error amplification.** (a) For one Cosmos-Predict2.5 query, the small residuals align with the summed error without compensation but point randomly with DecoupAttn, so the coherence ratio κ of Eq. (7) stops growing after the 16 largest (ticks: these 16, 90% of the energy). (b) Per-unit κ drops 2.9–3.9-fold on three models (bars: quartiles, dots: medians, ÷: drop in median), and (c) under one compensator, blocks fixed by their residual match per-step routing by attention mass at no routing cost during inference. (d) Two dense attention backends diverge along the sampler, less under per-pixel conditions (blue) than under prompts (pink). (e) Adding edges to Cosmos-Predict2.5 halves our lead over SVG-EAR and cuts the value of the sensitivity fourfold (lines: 95% bootstrap CI), and (f) on ten models, it falls as the dense-backend floor rises.

runs the other calls at density 0.20 or its official minimum (Appendix C, 0.23–0.42 with padding in Table 14), and is scored and timed against the same dense outputs in one process.

Metrics and implementation. Quality is the PSNR of each video against the dense output with the same seed, together with SSIM [36] and LPIPS [47] and six VBench dimensions [7] (Appendix C), and speed is the latency of the denoising loop (DiT) and its speedup over dense attention. DecoupAttn sparsifies the spatiotemporal self-attention of every model and keeps cross-attention dense. Exact blocks run on FlashInfer's variable-block FlashAttention-3 kernel [28, 42], and one kernel merges the cluster terms (Appendix B).

5.2. Main Results

DecoupAttn is accurate and fast at once (RQ1). As shown in Table 1, DecoupAttn has the highest PSNR on nine of the ten models, 2.18 dB above the strongest baseline on average, while running 1.31× faster than dense attention and up to 1.65×. Among the methods faster than dense attention it is the most accurate on all ten and leads every baseline in pairwise PSNR (Figure 5f), with the lowest LPIPS on all ten, VBench within 0.007 of dense attention (Tables 7 and 28), and dense-level error to ground truth (Table 35, cases in Appendix E). The three lines of prior work show where this comes from: dropping the skipped keys is fastest but least accurate, compensating with cheap position blocks lands in between, and compensation that re-clusters at every step falls below dense speed on all four-step models (Figure 5a, Table 27). Decoupling removes this trade-off: blocks and clusters meet only through the counts n_{gc}, so it fixes the blocks offline, clusters once, and compensates accurately at fixed-mask speed.

One procedure serves every family and sampler (RQ2). With the same procedure and constants, DecoupAttn leads the best baseline in all three families and both sampler types, by up to 3.79 dB, with each model's sampler setting its own J_u (Eq. (11)).

5.3. Component Analysis (RQ3)

We validate each component by removing it from the full method and re-solving the plan at the same density (Table 2). Every removal lowers PSNR in all 59 cells.

Decoupling matters most. Grouping compensation by the selection blocks (Tied), as in Definition 1, costs 2.85 dB on average at almost unchanged speed, and 3.52 dB under our plan (Tied†). Position blocks group keys by place, not content, so their members spread far from the centroid and leave large residuals: under one exact mask, semantic clusters raise PSNR from 18.36 to 23.47 dB (Figure 5e, Figure 6). A centroid underestimates the mass of a spread-out cluster, so dropping both corrections (Centroid) costs 1.49 dB (Appendix D.3).

(a) Time of one video (s) and the plan's cost

Method	k-means	per step	Attn.	DiT	E2E
CP2.5	plan: 5.9 GPU-h, repaid after 5.3k videos				
Dense	–	–	9.35	12.15	19.68
SVG-EAR	4.73	0.92	10.69	13.35	20.96
Ours	**0.62**	**0.57**	**5.48**	**8.20**	**15.71**
Wan2.1	plan: 5.3 GPU-h, repaid after 23.7k videos				
Dense	–	–	4.84	13.65	16.40
SVG-EAR	7.14	1.33	11.91	20.52	23.27
Ours	**0.92**	**0.91**	**4.27**	**12.91**	**15.66**
UniRel.	plan: 19.6 GPU-h, repaid after 2.1k videos				
Dense	–	–	68.42	141.10	144.40
SVG-EAR	18.44	20.88	77.65	144.76	147.92
Ours	**1.22**	**15.27**	**39.69**	**108.43**	**111.72**

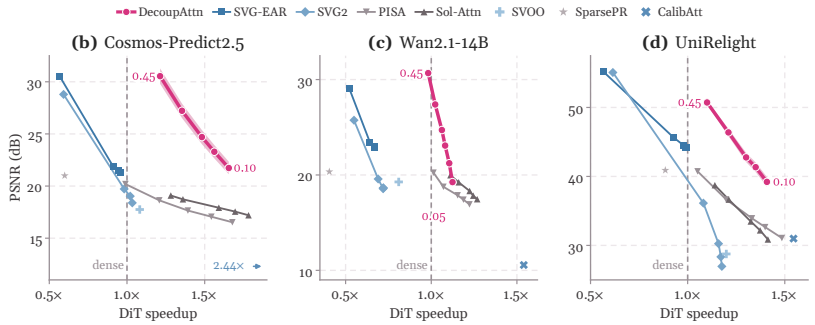

(e) Skipped keys and exact blocks, changed one at a time

		Cosmos-Predict2.5					Wan2.1-14B				
Grouping	Corrections	Attn.	Latent	PSNR	LPIPS	DiT (s)	Attn.	Latent	PSNR	LPIPS	DiT (s)
–	none	415	214	17.57±2.33	0.385	7.18	577	186	17.08±1.10	0.300	11.29
position	centroid	38.6	174	18.36±2.72	0.297	7.39	109	88.6	19.69±1.46	0.155	11.44
position	both	23.1	69.0	22.33±3.26	0.152	7.47	73.2	68.5	20.66±1.62	0.123	11.55
semantic	centroid	10.8	54.1	23.47±3.68	0.124	8.13	14.7	36.4	23.41±2.05	0.078	12.76
semantic	both (ours)	6.23	41.2	24.70±3.78	0.102	8.21	8.80	26.0	24.70±2.29	0.057	12.90
Selection	Ranking	Routing (ms)		PSNR	LPIPS	DiT (s)	Routing (ms)		PSNR	LPIPS	DiT (s)
offline	attention mass	0		24.38±3.66	0.107	8.22	0		24.37±2.24	0.061	12.91
offline	residual (ours)	0		24.70±3.78	0.102	8.21	0		24.70±2.29	0.057	12.90
online	attention mass	650		24.65±3.46	0.098	8.86	1,288		23.95±2.08	0.066	14.19
offline	random	0		17.26±2.61	0.415	8.22	0		16.36±1.42	0.289	12.87
CalibAtt masks		0		21.22±3.39	0.183	8.21	0		20.79±1.73	0.116	12.88

(f) Pairwise PSNR differences

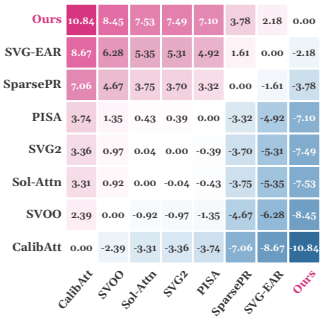

Figure 5. **Where the speed and accuracy come from, and how they hold.** (a) Medians of interleaved runs; per step covers assignment, statistics, counts, and merging. (b–d) DecoupAttn sweeps its density (labels; shaded: ±1 standard error). (e) Top: our exact mask, bottom: our compensator (test prompts; MSE $\times 10^3$; small ±: standard deviation). (f) Row minus column PSNR (dB), mean over ten models.

With compensation on, fixed blocks are as good as routing. Under the same compensator, routing by attention mass at every call matches our fixed blocks on Cosmos-Predict2.5 at 650 ms of routing and falls behind on Wan2.1-14B (Figure 5e, Table 16). The cancellation of Section 4.1 explains this: compensated residuals cancel each other, cutting the median coherence ratio on Cosmos-Predict2.5 from 10.37 to 2.69 (Figure 4a–c, Table 11). Ranking blocks offline by this residual also beats ranking by mass or CalibAtt's masks (Figure 5e).

Pricing the residual by its effect on the output matters next. Setting every sensitivity to one (U) costs 1.68 dB on average, and pricing options by their error before compensation (D), which buys exact blocks where compensation already works, costs 1.18 dB, both at unchanged speed (Table 8). The price orders 21 candidate plans by their held-out error as Proposition 2 expects, and the deployed plan is the second best (Figure 8, Table 20). The sensitivity pays most where the sampler amplifies errors: removing it costs 2.30 dB on Cosmos-Predict2.5 but 0.55 dB with edges (Figure 4e,f, Appendix D.4).

Reassignment keeps the clusters current at little cost. Drifting keys spread frozen clusters beyond what the plan priced, so freezing them (Frozen) costs 1.23 dB on average over the nine models that reassign (Table 26). Reassignment never increases this spread (Proposition 3, Figure 11), so per-step k-means (Rebuild) changes PSNR by at most 0.95 dB for 11–43% more latency (Tables 9, 8, and 25).

5.4. Trade-off and Offline Plan (RQ4, RQ5)

The lead holds across densities (RQ4). Sweeping its density from 0.10 (0.05 on Wan2.1-14B) to 0.45, every baseline lies below its curve within its speed range (Figure 5b–d), since each point is its density's lowest-priced plan (Proposition 2), and at the latency of the strongest baseline it is 5.74–8.40 dB more accurate (Figure 12, Table 29).

The offline plan is cheap and holds on new content (RQ5). Calibration takes 5.3–19.6 GPU-hours per model, which the faster inference repays after 2.1k–23.7k videos (Figure 5a, Table 30). As the frozen plan fixes only where exact attention runs and each video clusters its own keys, it beats SVG-EAR on all held-out content (Table 32), and Appendix D.8 varies calibration, resolution, and cluster count.

6. Conclusion

We introduced DecoupAttn, a decoupled sparse attention that keeps compensation when the selection is fixed offline: semantic clusters compensate every key, exact attention corrects their residual on a few fixed blocks, Distortion Budget prices that residual, and Reassign Cache keeps the clusters current. On ten video generation, controlled generation, and neural rendering models, DecoupAttn is the most accurate method faster than dense attention, at up to 1.65× its speed. Our results argue for organizing exact computation and approximation separately: position blocks for the kernel, semantic clusters for the approximation error.

LLM Usage Statement

We used LLMs (Claude) to polish the grammar and word choice of the paper and to refine the plotting scripts of the figures. The authors designed the method and the experiments, ran all experiments, and checked every statement, number, and reference in the paper.

References

[1] Beidi Chen, Tri Dao, Eric Winsor, Zhao Song, Atri Rudra, and Christopher Ré. Scatterbrain: Unifying sparse and low-rank attention approximation. In *NeurIPS*, 2021. 2, 4, 5

[2] Jiarui Chen, Zeqiang Lai, Jiangshan Wang, Ziheng Ouyang, Ye Huang, Xiangyu Yue, Cewu Lu, and Chunchao Guo. MC-Sparse: Deconstructing and closing the dense-sparse attention gap in diffusion transformers. *arXiv preprint arXiv:2610.06801*, 2026. 2

[3] Zhen Dong, Zhewei Yao, Amir Gholami, Michael W. Mahoney, and Kurt Keutzer. HAWQ: Hessian AWare quantization of neural networks with mixed-precision. In *ICCV*, 2019. 2, 5

[4] Hugh Everett III. Generalized Lagrange multiplier method for solving problems of optimum allocation of resources. *Operations Research*, 11(3):399–417, 1963. 5, 7

[5] Kai He, Ruofan Liang, Jacob Munkberg, Jon Hasselgren, Nandita Vijaykumar, Alexander Keller, Sanja Fidler, Igor Gilitschenski, Zan Gojcic, and Zian Wang. UniRelight: Learning joint decomposition and synthesis for video relighting. In *NeurIPS*, 2025. 6

[6] Xun Huang, Zhengqi Li, Guande He, Mingyuan Zhou, and Eli Shechtman. Self forcing: Bridging the train-test gap in autoregressive video diffusion. In *NeurIPS*, pages 167283–167308, 2025. 1

[7] Ziqi Huang, Yinan He, Jiashuo Yu, Fan Zhang, Chenyang Si, Yuming Jiang, Yuanhan Zhang, Tianxing Wu, Qingyang Jin, Nattapol Chanpaisit, Yaohui Wang, Xinyuan Chen, Limin Wang, Dahua Lin, Yu Qiao, and Ziwei Liu. VBench: Comprehensive benchmark suite for video generative models. In *CVPR*, 2024. 7, 22

[8] Zheng-Hui Huang, Zhixiang Wang, Jiaming Tan, Ruihan Yu, Yidan Zhang, Bo Zheng, Yu-Lun Liu, Yung-Yu Chuang, and Kaipeng Zhang. Generative world renderer. *arXiv preprint arXiv:2604.02329*, 2026. 6

[9] Weijie Kong, Qi Tian, Zijian Zhang, Rox Min, Zuozhuo Dai, Jin Zhou, Jiangfeng Xiong, Xin Li, Bo Wu, Jianwei Zhang, et al. HunyuanVideo: A systematic framework for large video generative models. *arXiv preprint arXiv:2412.03603*, 2024. 5

[10] Dongyeun Lee, Amir Zandieh, Vahab Mirrokni, Junmo Kim, and Insu Han. HyperVAttention: Efficient sparse attention with spatio-temporal clustering for video diffusion. *arXiv preprint arXiv:2607.03012*, 2026. 2

[11] Haopeng Li, Yitong Li, Junsong Chen, Tian Ye, Haozhe Liu, Jincheng Yu, Duomin Wang, Ruihua Zhang, Zeke Xie, Enze Xie, and Song Han. Sol-Attn: Accelerating video generation inference via on-the-fly attention sparsification. *arXiv preprint arXiv:2607.24027*, 2026. 2, 6

[12] Haopeng Li, Shitong Shao, Wenliang Zhong, Zikai Zhou, Lichen Bai, Hui Xiong, and Zeke Xie. PISA: Piecewise sparse attention is wiser for efficient diffusion transformers. In *NeurIPS*, 2026. arXiv:2602.01077. 1, 2, 3, 6

[13] Xingyang Li, Muyang Li, Tianle Cai, Haocheng Xi, Shuo Yang, Yujun Lin, Lvmin Zhang, Songlin Yang, Jinbo Hu, Kelly Peng, Maneesh Agrawala, Ion Stoica, Kurt Keutzer, and Song Han. Radial attention: $O(n \log n)$ sparse attention with energy decay for long video generation. In *NeurIPS*, 2025. 1, 2

[14] Yu Li, Binxu Li, and Tian Lan. MomentKV: Closing the directional gap in KV cache eviction for long-context inference. *arXiv preprint arXiv:2606.01563*, 2026. 2

[15] Ruofan Liang, Zan Gojcic, Huan Ling, Jacob Munkberg, Jon Hasselgren, Chih-Hao Lin, Jun Gao, Alexander Keller, Nandita Vijaykumar, Sanja Fidler, and Zian Wang. Diffusion renderer: Neural inverse and forward rendering with video diffusion models. In *CVPR*, pages 26069–26080, 2025. 1, 6

[16] Enhuai Liu, Yunke Wang, Yutong Wang, Changming Sun, and Chang Xu. LoSA: Near-lossless sparse attention for training-free video diffusion acceleration. *arXiv preprint arXiv:2608.12032*, 2026. 2

[17] Xuewen Liu, Zhikai Li, Jing Zhang, Mengjuan Chen, and Qingyi Gu. Rectified SpaAttn: Revisiting attention sparsity for efficient video generation. *arXiv preprint arXiv:2511.19835*, 2025. 2

[18] Yuxi Liu, Yipeng Hu, Zekun Zhang, Kunze Jiang, and Kun Yuan. Mixture of distributions matters: Dynamic sparse attention for efficient video diffusion transformers. *arXiv preprint arXiv:2601.11641*, 2026. 2

[19] Stuart P. Lloyd. Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2):129–137, 1982. 5, 8

[20] Jiayi Luo, Jiayu Chen, Jiankun Wang, Cong Wang, Hanxin Zhu, Qingyun Sun, Chen Gao, Zhibo Chen, and Jianxin Li. Attention sparsity is input-stable: Training-free sparse attention for video generation via offline sparsity profiling and online QK co-clustering. In *ICML*, pages 82177–82195, 2026. 2, 6

[21] Tan M. Nguyen, Vai Suliafu, Stanley J. Osher, Long Chen, and Bao Wang. FMMformer: Efficient and flexible transformer via decomposed near-field and far-field attention. In *NeurIPS*, 2021. 2

[22] NVIDIA. Cosmos 3: Omnimodal world models for physical AI. *arXiv preprint arXiv:2606.02800*, 2026. 5

[23] NVIDIA, Niket Agarwal, Arslan Ali, Maciej Bala, Yogesh Balaji, Erik Barker, Tiffany Cai, Prithvijit Chattopadhyay, Yongxin Chen, Yin Cui, et al. Cosmos world foundation model platform for physical AI. *arXiv preprint arXiv:2501.03575*, 2025. 5

[24] NVIDIA, Arslan Ali, Junjie Bai, Maciej Bala, Yogesh Balaji, Aaron Blakeman, Tiffany Cai, Jiaxin Cao, Tianshi Cao, et al. World simulation with video foundation models for physical AI. *arXiv preprint arXiv:2511.00062*, 2025. 1, 5, 6

[25] Aurko Roy, Mohammad Saffar, Ashish Vaswani, and David Grangier. Efficient content-based sparse attention with routing transformers. *Transactions of the Association for Computational Linguistics*, 9:53–68, 2021. 2

[26] Carl-Erik Särndal, Bengt Swensson, and Jan Wretman. *Model Assisted Survey Sampling*. Springer-Verlag, New York, 1992. 3, 5

[27] Leonard Schuchman. Dither signals and their effect on quantization noise. *IEEE Transactions on Communication Technology*, 12(4):162–165, 1964. 2, 4, 5

[28] Jay Shah, Ganesh Bikshandi, Ying Zhang, Vijay Thakkar, Pradeep Ramani, and Tri Dao. FlashAttention-3: Fast and accurate attention with asynchrony and low-precision. In *NeurIPS*, 2024. 7

[29] Dor Shmilovich, Tony Wu, Aviad Dahan, and Yuval Domb. LiteAttention: A temporal sparse attention for diffusion transformers. *arXiv preprint arXiv:2511.11062*, 2025. 2

[30] Yair Shoham and Allen Gersho. Efficient bit allocation for an arbitrary set of quantizers. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 36(9):1445–1453, 1988. 2, 4

[31] Gary J. Sullivan and Thomas Wiegand. Rate-distortion optimization for video compression. *IEEE Signal Processing Magazine*, 15(6):74–90, 1998. 2, 4

[32] Pardis Taghavi, Reza Langari, and Gaurav Pandey. Partition the support, reconstruct the residual: Training-free sparse attention for video generation and world models. *arXiv preprint arXiv:2608.18484*, 2026. 1, 2, 6

[33] Haoyue Tan, Shengnan Wang, Yulin Qiao, Juncheng Zhang, Youhui Bai, Ping Gong, Zewen Jin, and Cheng Li. AdaCluster: Adaptive query-key clustering for sparse attention in video generation. In *CVPR*, 2026. 1, 2

[34] Team Wan, Ang Wang, Baole Ai, Bin Wen, Chaojie Mao, Chen-Wei Xie, Di Chen, Feiwu Yu, Haiming Zhao, Jianxiao Yang, et al. Wan: Open and advanced large-scale video generative models. *arXiv preprint arXiv:2503.20314*, 2025. 1, 5

[35] Apoorv Vyas, Angelos Katharopoulos, and François Fleuret. Fast transformers with clustered attention. In *NeurIPS*, 2020. 2

[36] Zhou Wang, Alan C. Bovik, Hamid R. Sheikh, and Eero P. Simoncelli. Image quality assessment: from error visibility to structural similarity. *IEEE Transactions on Image Processing*, 13(4):600–612, 2004. 7

[37] Haocheng Xi, Shuo Yang, Yilong Zhao, Chenfeng Xu, Muyang Li, Xiuyu Li, Yujun Lin, Han Cai, Jintao Zhang, Dacheng Li, Jianfei Chen, Ion Stoica, Kurt Keutzer, and Song Han. Sparse VideoGen: Accelerating video diffusion transformers with spatial-temporal sparsity. In *ICML*, 2025. 2

[38] Tianze Xia, Xiaoling Zhou, Ruya Lan, Yiqiao Liao, Chengru Song, and Shiming Xiang. Sparsity should be ALLOTted by output: Training-free allocation of block-sparse attention for video diffusion transformers. *arXiv preprint arXiv:2610.29441*, 2026. 2, 3

[39] Ruyi Xu, Guangxuan Xiao, Haofeng Huang, Junxian Guo, and Song Han. XAttention: Block sparse attention with antidiagonal scoring. In *ICML*, 2025. 1, 2

[40] Tiange Xu, Kai Liu, Linghe Kong, and Yulun Zhang. RECAP: Reciprocal sparse attention and hierarchical caching for efficient video diffusion. *arXiv preprint arXiv:2610.26244*, 2026. 2

[41] Shuo Yang, Haocheng Xi, Yilong Zhao, Muyang Li, Jintao Zhang, Han Cai, Yujun Lin, Xiuyu Li, Chenfeng Xu, Kelly Peng, Jianfei Chen, Song Han, Kurt Keutzer, and Ion Stoica. Sparse VideoGen2: Accelerate video generation with sparse attention via semantic-aware permutation. In *NeurIPS*, 2025. 1, 2, 6

[42] Zihao Ye, Lequn Chen, Ruihang Lai, Wuwei Lin, Yineng Zhang, Stephanie Wang, Tianqi Chen, Baris Kasikci, Vinod Grover, Arvind Krishnamurthy, and Luis Ceze. FlashInfer: Efficient and customizable attention engine for LLM inference serving. In *MLSys*, 2025. 3, 7

[43] Shai Yehezkel, Shahar Yadin, Noam Elata, Yaron Ostrovsky-Berman, and Bahjat Kawar. Accelerating text-to-video generation with calibrated sparse attention. *arXiv preprint arXiv:2603.05503*, 2026. 1, 2, 6, 13, 14

[44] Tianwei Yin, Qiang Zhang, Richard Zhang, William T. Freeman, Frédo Durand, Eli Shechtman, and Xun Huang. From slow bidirectional to fast autoregressive video diffusion models. In *CVPR*, pages 22963–22974, 2025. 1

[45] Jintao Zhang, Chendong Xiang, Haofeng Huang, Jia Wei, Haocheng Xi, Jun Zhu, and Jianfei Chen. SpargeAttention: Accurate and training-free sparse attention accelerating any model inference. In *ICML*, 2025. 1, 2

[46] Peiyuan Zhang, Yongqi Chen, Runlong Su, Hangliang Ding, Ion Stoica, Zhengzhong Liu, and Hao Zhang. Fast video generation with sliding tile attention. In *ICML*, 2025. 2

[47] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *CVPR*, 2018. 7

[48] Wentai Zhang, Ronghui Xi, Shiyao Peng, Jiayu Huang, Haoran Luo, Zichen Tang, and Haihong E. Ride the wave: Precision-allocated sparse attention for smooth video generation. *arXiv preprint arXiv:2604.12219*, 2026. 2, 14, 16

[49] Qi Zhao, Qirui Li, Hanlin Tang, Yiduo Li, Zhen Guo, Cuifeng Shen, Chao Xu, Zhaosheng Chi, Xiaojin Lu, Kan Liu, Tao Lan, Lin Qu, and Xi Li. SCOPE: Subspace clustering with online per-head top-k estimation for sparse video attention. *arXiv preprint arXiv:2608.12780*, 2026. 2

[50] Lin Zheng, Jianbo Yuan, Chong Wang, and Lingpeng Kong. Efficient attention via control variates. In *ICLR*, 2023. 2, 4, 5

[51] Xuzhe Zheng, Yuexiao Ma, Jing Xu, Xiawu Zheng, Rongrong Ji, and Fei Chao. HEART: Exploiting head heterogeneity in sparse attention for video diffusion. *arXiv preprint arXiv:2605.14513*, 2026. 2, 3

[52] Ruiliang Zhou, Xuecheng Wu, Kang He, Guangyun Han, Bin Liu, Qinqin Chen, Wende Xu, Qingjie Zhao, and Chengru Song. ScalingAttention: Discovering intrinsic sparse attention topology for video diffusion transformers. *arXiv preprint arXiv:2606.23019*, 2026. 2, 3

[53] Xuanyi Zhou, Qiuyang Mang, Shuo Yang, Haocheng Xi, Jintao Zhang, Huanzhi Mao, Joseph E. Gonzalez, Kurt Keutzer, Ion Stoica, and Alvin Cheung. SVG-EAR: Parameter-free linear compensation for sparse video generation via error-aware routing. In *ECCV*, pages 167–185, 2026. 1, 2, 3, 6

DecoupAttn: Decoupled Sparse Attention for Training-Free Acceleration of Video Diffusion

Supplementary Material

A. Theory of Decoupled Sparse Attention

This appendix develops the theory behind the three components of Section 4 and proves its three propositions. Section A.1 fixes the notation and the assumptions. Section A.2 shows that the decoupled estimator is an exact rearrangement of attention (Proposition 1) whose error is additive over blocks. Section A.3 expands the centroid compensation to second order and identifies what each correction removes. Section A.4 explains why fixed blocks suffice: compensation subtracts exactly the part of the block residuals that aligns, and for blocks fixed independently of the cluster membership the remaining deviation part is no more coherent than random errors. Section A.5 turns this into a rule for choosing the exact blocks, and Section A.6 relates the estimator to difference estimators. Section A.7 analyzes the plan objective, the sensitivity estimator, and the Lagrangian relaxation (Proposition 2). Section A.8 proves Proposition 3 and the monotone chain behind Reassign Cache, Section A.9 counts the cost of one call, and Section A.10 traces how errors travel through the sampler. Algebraic identities hold on every attention call, the expansions hold under the bounded-deviation assumption (A3), their orders in the deviations under its joint form (A3′), and the statements about expectations under the exchangeability model (A4).

A.1. Setting and Notation

Fix one self-attention call and one query i, and drop the indices of step, layer, and head. The logits are $s_{ij} = q_i^\top k_j/\sqrt{d}$, and the dense quantities N_{ib}, Z_{ib}, N_i, Z_i, and $y_i = N_i/Z_i$ are those of Eq. (1). The key blocks b and the clusters c are two partitions of the same L keys, and $n_{bc} = |b \cap c|$ counts the keys they share, so that

$$\textstyle\sum_b n_{bc} = |c|, \qquad \sum_c n_{bc} = |b|.$$

The query tile g of query i has exact blocks $\mathcal{S}_g$ and skipped blocks $\mathcal{K}_g$, the complement of $\mathcal{S}_g$. Every member j of cluster c is written as a deviation from the cluster means,

$$k_j = \bar{k}_c + \delta_j, \qquad v_j = \bar{v}_c + \nu_j,$$
$$\textstyle\sum_{j\in c} \delta_j = 0, \qquad \sum_{j\in c} \nu_j = 0,$$

and $x_{ij} = q_i^\top \delta_j/\sqrt{d}$ is the logit deviation, so that $e^{s_{ij}} = \phi_{ic}\, e^{x_{ij}}$ with $\phi_{ic} = e^{q_i^\top \bar{k}_c/\sqrt{d}}$. The key covariance and the value–key cross-covariance of a cluster are

$$C_c = \frac{1}{|c|}\sum_{j\in c} \delta_j\delta_j^\top, \qquad B_c = \frac{1}{|c|}\sum_{j\in c} \nu_j\delta_j^\top,$$

so the per-dimension variance σ_c^2 of Eq. (4) is the diagonal of C_c and the variance correction reads

$$\gamma_{ic} = \frac{1}{2d}\, q_i^\top \operatorname{diag}(C_c)\, q_i.$$

The per-member predictions of Eq. (4) are $\omega_{ic} = \phi_{ic}(1 + \gamma_{ic})$ and $a_{ic} = \omega_{ic}\bar{v}_c + \phi_{ic} H q_i/\sqrt{d}$. We use four assumptions, the third in a weak and a joint form. Table 3 collects the symbols used in the proofs.

Symbol	Meaning
s_{ij}, N_{ib}, Z_{ib}, y_i	logit, block numerator and mass, dense output (Eq. (1))
b, c, n_{bc}	key block, cluster, keys they share
$\mathcal{S}_g$, $\mathcal{K}_g$, n_{gc}	exact and skipped blocks of query tile g, members of c outside $\mathcal{S}_g$
$\bar{k}_c$, $\bar{v}_c$, δ_j, ν_j	cluster means and member deviations
x_{ij}, ϕ_{ic}	logit deviation $q_i^\top\delta_j/\sqrt{d}$, centroid weight
C_c, B_c, H	key covariance, value–key cross-covariance, pooled B_c
γ_{ic}, ω_{ic}, a_{ic}	variance correction, per-member mass and numerator prediction
u_{ij}, r_{ib}, U_{ic}	error of one compensated key, of a block, of a whole cluster
$\bar{z}_{ic}$, w_{ij}	cluster-mean contribution, member deviation
D_b, β_{ib}	deviation term and cluster bias of a block
η, V	deviation bound, largest value norm $\max_j \|v_j\|_2$
κ_i, R_i, ζ_i	coherence ratio, residual energy, mass ratio error
ε_u, J_u, w_u, X_u	distortion, sensitivity, step weight, density of a unit
μ, ρ, F, $\mathcal{L}_\mu$	multiplier, density target, priced distortion, Lagrangian

Table 3. Notation of Appendix A.

(A1) Count-based compensation. Every member of cluster c that lies outside the exact blocks contributes a_{ic} to the numerator and ω_{ic} to the denominator, and these terms depend only on the cluster and the query.
(A2) Current statistics. $\bar{k}_c$, $\bar{v}_c$, σ_c^2, and H are computed from the current members of each cluster, at every call.
(A3) Bounded deviations. $|x_{ij}| \le \eta$ for every query and every member of every cluster.
(A3′) Small joint deviations. $\|q_i\|_2\|\delta_j\|_2/\sqrt{d} \le \eta$ and $\|\nu_j\|_2 \le \eta V$, with $V = \max_j \|v_j\|_2$, for every query and every member of every cluster.
(A4) Exchangeable membership. Within each cluster, the members are distributed over the blocks by a uniformly random permutation that keeps the counts n_{bc}, independently across clusters and of the members' keys and values.

(A1) and (A2) describe the implementation and hold exactly. (A3) is the regime in which the expansions of Section A.3 are accurate, and clustering makes η small because

k-means groups nearby keys. (A3′) implies (A3) and adds that the values of a cluster vary on the same small scale, as when values move with their keys. Under it every correction term is quadratic and every remainder cubic in η. (A4) is a statistical model of the membership under which Section A.4 computes expectations: it states that which block a member falls into carries no information about the direction of its deviation once its cluster is known. Blocks fixed offline never read the membership of a new input, so the skipped set is independent of the membership, and Figure 9 tests the exchangeability itself on saved attention inputs.

A.2. The Decoupled Estimator (Proposition 1)

Proof of Proposition 1. The key blocks partition the keys, so $|c| = \sum_b n_{bc}$ and $n_{gc} = \sum_{b\notin\mathcal{S}_g} n_{bc}$. Substituting into Eq. (5) gives

$$\begin{aligned}\hat{N}_i &= \sum_{b\in\mathcal{S}_g} N_{ib} + \sum_{b\notin\mathcal{S}_g}\sum_c n_{bc}\, a_{ic} \\ &= \sum_c |c|\, a_{ic} + \sum_{b\in\mathcal{S}_g}\Big(N_{ib} - \sum_c n_{bc}\, a_{ic}\Big),\end{aligned}$$

which is the first identity of Eq. (6), and the denominator follows in the same way with Z_{ib} and ω_{ic}. Since $N_i = \sum_b N_{ib}$ and $Z_i = \sum_b Z_{ib}$ (Eq. (1)), the same rearrangement gives $N_i - \hat{N}_i = \sum_{b\notin\mathcal{S}_g}\Delta N_{ib}$ and $Z_i - \hat{Z}_i = \sum_{b\notin\mathcal{S}_g}\Delta Z_{ib}$. Using $N_i = y_i Z_i$,

$$\hat{y}_i - y_i = \frac{\hat{N}_i - y_i\hat{Z}_i}{\hat{Z}_i} = -\frac{1}{\hat{Z}_i}\sum_{b\notin\mathcal{S}_g}\big(\Delta N_{ib} - y_i\Delta Z_{ib}\big).$$

The identity holds for any fixed blocks and any clusters. □

The residual has a per-member form that the rest of this appendix uses. For a member j of cluster c, let

$$u_{ij} = e^{s_{ij}}\big(v_j - y_i\big) - \big(a_{ic} - y_i\,\omega_{ic}\big)$$

be the error that compensating this one key leaves. Expanding ΔN_{ib} and ΔZ_{ib} over the members of block b shows that the residual of a block is the sum of the errors of its keys:

$$r_{ib} = \sum_{j\in b} u_{ij}.$$

Lemma A.1 (Additive block removal). *Let $E_i(\mathcal{S}) = \hat{N}_i - y_i\hat{Z}_i$ be the unnormalized error of the estimator with exact blocks $\mathcal{S}$. Then $E_i(\mathcal{S}) = -\sum_{b\notin\mathcal{S}} r_{ib}$, and for every block $b\notin\mathcal{S}$,*

$$E_i(\mathcal{S}\cup\{b\}) = E_i(\mathcal{S}) + r_{ib}, \qquad \hat{Z}_i(\mathcal{S}\cup\{b\}) = \hat{Z}_i(\mathcal{S}) + \Delta Z_{ib}.$$

Proof. The first claim is the last display of the proof of Proposition 1 multiplied by $\hat{Z}_i$. Making b exact removes its term from the sum over $b\notin\mathcal{S}$, and replaces its compensated mass $\sum_c n_{bc}\omega_{ic}$ by Z_{ib} in the denominator. □

Lemma A.1 says that every block owns a fixed share r_{ib} of the unnormalized error, whatever the other blocks do. This is what makes a residual a property of a block that can be measured once and priced offline, while the coherence ratio and the mass $\hat{Z}_i$ decide how the shares combine into the output error (Corollary A.2).

Corollary A.2 (Three factors of the error). *With $R_i = \sum_{b\notin\mathcal{S}_g}\|r_{ib}\|_2^2$ and the coherence ratio κ_i of Eq. (7), set to one when $R_i = 0$ and between 0 and $|\mathcal{K}_g|$ by Cauchy–Schwarz,*

$$\|\hat{y}_i - y_i\|_2^2 = \frac{\kappa_i\, R_i}{\hat{Z}_i^2}, \qquad \frac{\|\hat{y}_i - y_i\|_2}{\|y_i\|_2} \le \frac{1}{1-\zeta_i}\,\frac{\sqrt{\kappa_i R_i}}{Z_i\|y_i\|_2}$$

whenever $\zeta_i = |\hat{Z}_i/Z_i - 1| < 1$ and $y_i \neq 0$. Proof. The equality combines Proposition 1 with the definition of κ_i. For the bound, $\hat{Z}_i \ge (1-\zeta_i)Z_i$. □

The three factors are the residual energy R_i, its coherence κ_i, and the mass ratio $\hat{Z}_i/Z_i$, which are the columns of Table 24. Only the variance correction acts on ζ_i, because the value correction enters $\hat{N}_i$ but not $\hat{Z}_i$: on DiffusionRenderer the median $|\hat{Z}_i/Z_i - 1|$ is 0.27 without compensation, 0.100 with plain centroids or the value correction alone, and 0.035 with the variance correction (Table 24).

A.3. Second-Order Analysis of the Compensation

Lemma A.3 (A centroid underestimates the mass). *For every cluster and every query, $\sum_{j\in c} e^{s_{ij}} \ge |c|\,\phi_{ic}$, with equality only if all logit deviations x_{ij} of the cluster are zero.*

Proof. The exponential is strictly convex and the members' keys average to $\bar{k}_c$, so Jensen's inequality gives

$$\frac{1}{|c|}\sum_{j\in c} e^{s_{ij}} \ge \exp\Big(\frac{1}{|c|}\sum_{j\in c} s_{ij}\Big) = \exp\big(q_i^\top\bar{k}_c/\sqrt{d}\big) = \phi_{ic},$$

with equality exactly when the s_{ij} are equal, that is, when the x_{ij} vanish. □

Lemma A.3 needs no expansion. It shows that a centroid underestimates the mass of a spread-out cluster and, under (A4), in expectation that of the members the exact blocks leave, since each member is skipped with equal probability.

Lemma A.4 (The variance correction is the diagonal second-order term). *Under (A3),*

$$\sum_{j\in c} e^{s_{ij}} = |c|\,\phi_{ic}\Big(1 + \frac{1}{2d}\,q_i^\top C_c\, q_i\Big) + \phi_{ic}\,\rho_{ic}, \qquad |\rho_{ic}| \le \frac{|c|\,\eta^3 e^\eta}{6},$$

and therefore the corrected prediction misses the mass of the cluster by

$$\sum_{j\in c} e^{s_{ij}} - |c|\,\omega_{ic} = \frac{|c|\,\phi_{ic}}{2d}\,q_i^\top\big(C_c - \mathrm{diag}(C_c)\big)q_i + \phi_{ic}\,\rho_{ic}.$$

Proof. Taylor's theorem gives $e^x = 1 + x + x^2/2 + \varrho(x)$ with $|\varrho(x)| \le |x|^3 e^{|x|}/6$. Summing $e^{s_{ij}} = \phi_{ic} e^{x_{ij}}$ over the cluster, the linear term vanishes because $\sum_{j\in c} x_{ij} = q_i^\top \sum_{j\in c} \delta_j/\sqrt{d} = 0$, and the quadratic term is

$$\frac{1}{2}\sum_{j\in c} x_{ij}^2 = \frac{1}{2d}\, q_i^\top \Big(\sum_{j\in c} \delta_j \delta_j^\top\Big) q_i = \frac{|c|}{2d}\, q_i^\top C_c\, q_i.$$

The remainders add up to $\rho_{ic} = \sum_{j\in c} \varrho(x_{ij})$, which (A3) bounds term by term. Subtracting $|c|\,\omega_{ic} = |c|\,\phi_{ic}(1+\gamma_{ic})$ with $\gamma_{ic} = q_i^\top \operatorname{diag}(C_c)\, q_i/(2d)$ gives the second display. □

The CENTROID estimator therefore leaves the whole second-order term $|c|\,\phi_{ic}\, q_i^\top C_c q_i/(2d) \ge 0$, in agreement with Lemma A.3. The variance correction removes its diagonal part and leaves only the covariance between different key dimensions within a cluster, of order η^2 under (A3′), plus a remainder of order η^3.

Lemma A.5 (The value correction is the pooled cross-covariance). *(i) The matrix of Eq. (4) is the size-weighted average of the clusters' cross-covariances,*

$$H = \sum_c \frac{|c|}{L}\, B_c.$$

(ii) Under (A3), with $V = \max_j \|v_j\|_2$ and $V_\nu = \max_j \|\nu_j\|_2$,

$$\sum_{j\in c} e^{s_{ij}}\, v_j = |c|\,\phi_{ic}\Big[\bar v_c\Big(1 + \frac{q_i^\top C_c q_i}{2d}\Big) + \frac{B_c\, q_i}{\sqrt{d}}\Big] + \phi_{ic}\,\rho'_{ic},$$

with $\|\rho'_{ic}\|_2 \le |c|\big(\eta^2 V_\nu/2 + \eta^3 e^\eta V/6\big)$.

Proof. (i) Within one cluster the cross terms vanish because the deviations have zero mean:

$$\sum_{j\in c} v_j k_j^\top = |c|\,\bar v_c \bar k_c^\top + \sum_{j\in c} \nu_j \delta_j^\top = |c|\,\bar v_c \bar k_c^\top + |c|\,B_c.$$

Summing over clusters and dividing by L gives $H = \frac{1}{L}\big(\sum_j v_j k_j^\top - \sum_c |c|\,\bar v_c \bar k_c^\top\big) = \sum_c \frac{|c|}{L} B_c$. (ii) Multiply the expansion of $e^{x_{ij}}$ by $v_j = \bar v_c + \nu_j$ and sum over the cluster. The terms $\sum_j \nu_j$ and $\bar v_c \sum_j x_{ij}$ vanish, $\sum_j x_{ij}\nu_j = \big(\sum_j \nu_j \delta_j^\top\big) q_i/\sqrt{d} = |c|\,B_c q_i/\sqrt{d}$, and $\frac{1}{2}\bar v_c \sum_j x_{ij}^2$ is the variance term of Lemma A.4. What remains is $\frac{1}{2}\sum_j x_{ij}^2 \nu_j + \sum_j \varrho(x_{ij})\, v_j$, which (A3) bounds by the stated amount. □

Proposition A.6 (The residual of a whole cluster is second order). *Let $U_{ic} = \sum_{j\in c} u_{ij}$ be the error that compensation leaves on all members of cluster c together. Under (A3),*

$$U_{ic} = |c|\,\phi_{ic}\Big[(\bar v_c - y_i)\,\frac{q_i^\top\big(C_c - \operatorname{diag}(C_c)\big) q_i}{2d} + \frac{(B_c - H)\, q_i}{\sqrt{d}}\Big] + \phi_{ic}\big(\rho'_{ic} - y_i\,\rho_{ic}\big).$$

Under (A3′), moreover, $\|U_{ic}\|_2 \le |c|\,\phi_{ic}\, V\big(3\eta^2 + \eta^3 e^\eta\big)$.

Proof. $U_{ic} = \big(\sum_j e^{s_{ij}} v_j - |c|\,a_{ic}\big) - y_i\big(\sum_j e^{s_{ij}} - |c|\,\omega_{ic}\big)$. Insert Lemma A.5(ii) and $|c|\,a_{ic} = |c|\,\phi_{ic}\big[\bar v_c(1+\gamma_{ic}) + H q_i/\sqrt{d}\big]$ in the first bracket, and Lemma A.4 in the second. For the bound, (A3′) gives $|x_{ij}| \le \eta$ and $\sum_p q_{ip}^2 \delta_{jp}^2 \le \|q_i\|_2^2 \|\delta_j\|_2^2 \le d\eta^2$, so $q_i^\top C_c q_i/(2d)$ and $q_i^\top \operatorname{diag}(C_c) q_i/(2d)$ both lie in $[0, \eta^2/2]$. It also gives $\|B_c q_i\|_2/\sqrt{d} \le \eta^2 V$ and $\|H q_i\|_2/\sqrt{d} \le \eta^2 V$, because $B_c q_i/\sqrt{d}$ averages $x_{ij}\nu_j$ and H averages the B_c. With $\|\bar v_c - y_i\|_2 \le 2V$ and $\|y_i\|_2 \le V$, the bracket is at most $3\eta^2 V$, and the remainders of Lemmas A.4 and A.5 add at most $|c|\,\eta^3 V(1/2 + e^\eta/3) \le |c|\,\eta^3 e^\eta V$. □

Proposition A.6 says what the two corrections achieve. On every cluster, compensation is exact up to three terms: the off-diagonal key covariance, the gap between the cluster's own cross-covariance B_c and the head average H, and a remainder, of orders η^2, η^2, and η^3 under (A3′). The two corrections of Eq. (4) are therefore the second-order terms that a diagonal, head-shared model can capture. A full covariance per cluster and a per-cluster H_c remove the two second-order terms, and Table 18 measures what these additions gain on saved attention inputs.

A.4. Why Fixed Blocks Suffice

Lemma A.7 (First-order cancellation). *Under (A3′), the per-member error splits as $u_{ij} = u_{ij}^{(1)} + u_{ij}^{(2)}$ with $\|u_{ij}^{(2)}\|_2 \le (3 + e^\eta)\,\phi_{ic} V \eta^2$ and the first-order part*

$$u_{ij}^{(1)} = \phi_{ic}\big(\nu_j + (\bar v_c - y_i)\, x_{ij}\big),$$

which sums to zero over every cluster. Consequently the first-order parts $r_{ib}^{(1)} = \sum_{j\in b} u_{ij}^{(1)}$ of the block residuals sum to zero over all blocks, and

$$\sum_{b\notin \mathcal{S}_g} r_{ib}^{(1)} = -\sum_{b\in \mathcal{S}_g} r_{ib}^{(1)}.$$

Proof. With $e^{s_{ij}} = \phi_{ic} e^{x_{ij}}$, $v_j = \bar v_c + \nu_j$, and $a_{ic} - y_i\omega_{ic} = \phi_{ic}(1+\gamma_{ic})(\bar v_c - y_i) + \phi_{ic} H q_i/\sqrt{d}$, the zeroth-order terms $\phi_{ic}(\bar v_c - y_i)$ cancel and

$$u_{ij}^{(2)} = \phi_{ic}\Big[x_{ij}\nu_j + \big(e^{x_{ij}} - 1 - x_{ij}\big)(v_j - y_i) - \gamma_{ic}(\bar v_c - y_i) - H q_i/\sqrt{d}\Big].$$

Under (A3′) the four terms are at most $\eta^2 V$, $\eta^2 e^\eta V$, $\eta^2 V$, and $\eta^2 V$, by $|e^x - 1 - x| \le x^2 e^{|x|}/2$, $\|v_j - y_i\|_2 \le 2V$, and the bounds in the proof of Proposition A.6. The first-order part sums to zero because $\sum_{j\in c} \nu_j = 0$ and $\sum_{j\in c} x_{ij} = 0$. Since $r_{ib} = \sum_{j\in b} u_{ij}$ and the blocks partition the keys, $\sum_b r_{ib}^{(1)} = \sum_c \sum_{j\in c} u_{ij}^{(1)} = 0$. □

Dropping a block has an analogous identity: the dropped residuals $r_{ib}^{\text{drop}} = \sum_{j\in b} e^{s_{ij}}(v_j - y_i)$ also sum to zero over

all blocks, because $\sum_j e^{s_{ij}}(v_j - y_i) = N_i - y_i Z_i = 0$. What separates the two is not the sum over all blocks but how the residuals of the skipped blocks align with each other. The next proposition makes this precise without any expansion.

Proposition A.8 (Compensation subtracts the aligned part). *Let $\bar{z}_{ic} = \frac{1}{|c|}\sum_{j\in c} e^{s_{ij}}(v_j - y_i)$ be the mean contribution of a member of cluster c, and $w_{ij} = e^{s_{ij}}(v_j - y_i) - \bar{z}_{ic}$ its deviation, which sums to zero over the cluster. Then for every block,*

$$r_{ib}^{\text{drop}} = \underbrace{\sum_c n_{bc}\,\bar{z}_{ic}}_{\text{cluster means}} + \underbrace{\sum_{j\in b} w_{ij}}_{\text{deviations}}, \qquad r_{ib} = \underbrace{\sum_{j\in b} w_{ij}}_{\text{deviations}} + \underbrace{\sum_c n_{bc}\,\frac{U_{ic}}{|c|}}_{\text{cluster bias } \beta_{ib}}.$$

Proof. The first identity splits each member's contribution into $\bar{z}_{ic}$ and w_{ij} and groups the means by cluster. For the second, $r_{ib} = r_{ib}^{\text{drop}} - \sum_c n_{bc}(a_{ic} - y_i\omega_{ic})$, and averaging u_{ij} over the cluster gives $a_{ic} - y_i\omega_{ic} = \bar{z}_{ic} - U_{ic}/|c|$. □

The cluster-mean term of a dropped block is a sum of the same vectors $\bar{z}_{ic}$ for every block that contains members of c. Blocks with similar cluster composition therefore push the output in the same direction. Compensation removes this term and replaces it by the cluster bias β_{ib}, built from the $U_{ic}/|c|$ that Proposition A.6 shows to be of second order under (A3′) and that the variance correction shrinks in practice (Appendix D.5). What remains is the deviation term, whose behavior the next lemma computes.

Lemma A.9 (Sampling identities). *Let a finite population $w_1, \dots, w_n$ of vectors, $n \geq 2$, have $\sum_j w_j = 0$ and $\tau^2 = \frac{1}{n}\sum_j \|w_j\|_2^2$, and let disjoint subsets A and A' of sizes a and a' be drawn by a uniformly random permutation. Then*

$$\mathbb{E}\sum_{j\in A} w_j = 0, \qquad \mathbb{E}\Big\|\sum_{j\in A} w_j\Big\|_2^2 = \frac{a(n-a)}{n-1}\,\tau^2,$$

$$\mathbb{E}\Big\langle \sum_{j\in A} w_j, \sum_{j\in A'} w_j \Big\rangle = -\frac{a\,a'}{n-1}\,\tau^2.$$

Proof. Each w_j belongs to A with probability a/n, so the mean is $\frac{a}{n}\sum_j w_j = 0$. For the second moment, $\mathbb{E}\|\sum_A w_j\|^2 = \sum_j \Pr[j \in A]\,\|w_j\|^2 + \sum_{j\neq j'} \Pr[j, j' \in A]\,\langle w_j, w_{j'}\rangle$ with $\Pr[j \in A] = a/n$ and $\Pr[j, j' \in A] = \frac{a(a-1)}{n(n-1)}$. Because $\sum_{j\neq j'}\langle w_j, w_{j'}\rangle = \|\sum_j w_j\|^2 - \sum_j \|w_j\|^2 = -n\tau^2$, this equals $a\tau^2 - \frac{a(a-1)}{n-1}\tau^2 = \frac{a(n-a)}{n-1}\tau^2$. Applying the second identity to A, A', and $A \cup A'$ and expanding $\|\sum_{A\cup A'} w_j\|^2$ gives the cross term:

$$\begin{aligned} 2\,\mathbb{E}\big\langle \textstyle\sum_A w_j, \sum_{A'} w_j\big\rangle &= \frac{\tau^2}{n-1}\big[(a+a')(n-a-a') \\ &\quad - a(n-a) - a'(n-a')\big] \\ &= -\frac{2aa'}{n-1}\,\tau^2. \qquad \square \end{aligned}$$

Theorem A.10 (Deviation terms add like random errors). *Under (A4), let $\mathcal{K}$ be a set of skipped blocks fixed independently of the permutation, let $\tau_{ic}^2 = \frac{1}{|c|}\sum_{j\in c}\|w_{ij}\|_2^2$ for clusters with $|c| \geq 2$ (a singleton has $w_{ij} = 0$ and drops out), let $m_c = \sum_{b\in\mathcal{K}} n_{bc}$ count the members of c in $\mathcal{K}$, and let $D_b = \sum_{j\in b} w_{ij}$ be the deviation term of Proposition A.8. Then $\mathbb{E}\,D_b = 0$ and*

$$\begin{aligned} \mathbb{E}\Big\|\sum_{b\in\mathcal{K}} D_b\Big\|_2^2 &= \sum_c \frac{m_c(|c| - m_c)}{|c|-1}\,\tau_{ic}^2 \\ &\leq \sum_c \sum_{b\in\mathcal{K}} \frac{n_{bc}(|c| - n_{bc})}{|c|-1}\,\tau_{ic}^2 \\ &= \sum_{b\in\mathcal{K}} \mathbb{E}\,\|D_b\|_2^2. \end{aligned}$$

The expected coherence ratio $\bar{\kappa}_D = \mathbb{E}\|\sum_{b\in\mathcal{K}} D_b\|_2^2 / \sum_{b\in\mathcal{K}} \mathbb{E}\|D_b\|_2^2$ is therefore at most one, and distinct blocks are negatively correlated: $\mathbb{E}\langle D_b, D_{b'}\rangle = -\sum_c n_{bc}n_{b'c}\tau_{ic}^2/(|c|-1) \leq 0$. With the cluster bias β_{ib} of Proposition A.8, the full residual satisfies

$$\mathbb{E}\Big\|\sum_{b\in\mathcal{K}} r_{ib}\Big\|_2^2 = \mathbb{E}\Big\|\sum_{b\in\mathcal{K}} D_b\Big\|_2^2 + \Big\|\sum_c \frac{m_c}{|c|}\,U_{ic}\Big\|_2^2.$$

Proof. By (A4), the members of each cluster that fall into each block form disjoint random subsets of sizes n_{bc}, and different clusters are independent. Writing $D_b = \sum_c D_{bc}$ with $D_{bc} = \sum_{j\in b\cap c} w_{ij}$, every D_{bc} has zero mean (Lemma A.9), so cross terms between different clusters vanish. Within a cluster, $\sum_{b\in\mathcal{K}} D_{bc}$ is the sum over a random subset of size m_c, and Lemma A.9 gives the left side. Each block alone is a subset of size n_{bc}, which gives the right side and the cross term. The inequality holds cluster by cluster because $m_c \geq n_{bc}$ for every $b \in \mathcal{K}$:

$$m_c(|c| - m_c) = \sum_{b\in\mathcal{K}} n_{bc}(|c| - m_c) \leq \sum_{b\in\mathcal{K}} n_{bc}(|c| - n_{bc}).$$

For the full residual, $r_{ib} = D_b + \beta_{ib}$, and $\sum_{b\in\mathcal{K}}\beta_{ib} = \sum_c m_c U_{ic}/|c|$ depends only on the counts because U_{ic} sums over the whole cluster, so the cross term vanishes with $\mathbb{E}\,D_b = 0$. □

Corollary A.11 (The selection enters only through counts). *Under (A4), for exact sets fixed independently of the permutation, the expected deviation error $\mathbb{E}\|\sum_{b\notin\mathcal{S}_g} D_b\|_2^2$ depends on the exact set $\mathcal{S}_g$ only through the numbers $m_c = \sum_{b\notin\mathcal{S}_g} n_{bc}$ of skipped members of each cluster. Two exact sets that skip the same number of members of every cluster have the same expected deviation error, whichever blocks they contain.*

Proof. The left side of Theorem A.10 is $\sum_c m_c(|c| - m_c)\tau_{ic}^2/(|c|-1)$, which contains $\mathcal{S}_g$ only through the m_c. □

Blocks fixed offline never read the keys of a new input, so they meet this independence by construction. Among such sets, a change of blocks under the same budget changes the deviation error only through the m_c, and with hundreds of clusters spread over many blocks each m_c moves by little. A selection made anew at every call reads the current keys and can correlate with the deviations, which moves the error at equal counts in either direction, and Table 16 measures the net effect: under one compensator, routing by attention mass at every call reaches 24.65 dB against 24.70 dB for fixed blocks on Cosmos-Predict2.5.

Corollary A.12 (Dropped residuals align). *Under (A4), the dropped residuals have mean $\mu_b = \sum_c n_{bc}\bar{z}_{ic}$, and for every set $\mathcal{K}$ fixed independently of the permutation*

$$\mathbb{E}\Big\|\sum_{b\in\mathcal{K}} r_{ib}^{\mathrm{drop}}\Big\|_2^2 = \Big\|\sum_{b\in\mathcal{K}}\mu_b\Big\|_2^2 + \mathbb{E}\Big\|\sum_{b\in\mathcal{K}} D_b\Big\|_2^2.$$

If the skipped blocks share one cluster composition, $n_{bc} = \alpha_b p_c$ with $\sum_c p_c\bar{z}_{ic} \neq 0$, the cluster-mean terms are parallel and their coherence ratio $\|\sum_b \mu_b\|^2 / \sum_b \|\mu_b\|^2 = (\sum_b \alpha_b)^2 / \sum_b \alpha_b^2$ reaches $|\mathcal{K}|$ for blocks of equal size.

Proof. By Proposition A.8, $r_{ib}^{\mathrm{drop}} = \mu_b + D_b$ with $\mathbb{E}\, D_b = 0$, so the cross term between the two parts vanishes in expectation. Under a shared composition, $\mu_b = \alpha_b \sum_c p_c\bar{z}_{ic}$ for every b. □

Theorem A.10 and Corollary A.12 give the two regimes of the coherence ratio. Dropped blocks carry a cluster-mean term that adds coherently over the blocks that share clusters, so its coherence ratio grows with the number of such blocks. Compensated blocks carry only the deviation term, whose expected coherence ratio $\bar{\kappa}_D$ is at most one, plus the cluster bias, of second order under (A3′) and the coherent part of the residual in practice (Appendix D.5). This is the mechanism of subtractive dither in quantization [27]: subtracting a known prediction leaves zero-mean errors that add like random ones. The measured medians follow these regimes, from 4.69, 5.03, and 10.37 without compensation to 1.61, 1.65, and 2.69 with both corrections on DiffusionRenderer, Wan2.1-14B, and Cosmos-Predict2.5 (Table 11). Theorem A.10 holds for every set of blocks fixed independently of the membership, so blocks fixed offline by position keep $\bar{\kappa}_D$ at most one, as observed in Section 4.1.

A.5. Choosing the Exact Blocks

Proposition A.13 (Residual ranking). *Fix a query tile and a budget of K exact blocks besides its diagonal block b_g, and take expectations over calibration inputs. If the residuals of distinct blocks are orthogonal in expectation, $\mathbb{E}\langle r_{ib}, r_{ib'}\rangle = 0$ for $b \neq b'$, the set of K blocks that minimizes the expected unnormalized error $\mathbb{E}\|\sum_{b\notin\mathcal{S}} r_{ib}\|_2^2$, summed over the queries of the tile, consists of the K blocks with the largest $\sum_{i\in g}\mathbb{E}\|r_{ib}\|_2^2$. If instead $\mathbb{E}\langle r_{ib}, r_{ib'}\rangle \leq 0$, as for the deviation terms of Theorem A.10, the same set minimizes an upper bound on this error.*

Proof. With orthogonal residuals, $\mathbb{E}\|\sum_{b\notin\mathcal{S}} r_{ib}\|^2 = \sum_b \mathbb{E}\|r_{ib}\|^2 - \sum_{b\in\mathcal{S}}\mathbb{E}\|r_{ib}\|^2$ for every query, and summing over the query tile leaves a total that the largest K terms minimize. With nonpositive cross moments, $\mathbb{E}\|\sum_{b\notin\mathcal{S}} r_{ib}\|^2 \leq \sum_{b\notin\mathcal{S}}\mathbb{E}\|r_{ib}\|^2$, and the same choice minimizes the right side. □

Equation (9) implements this rule with each query weighted by $1/(Z_i^2\|y_i\|_2^2)$, for queries with $y_i \neq 0$. The weight is the scale of the relative error of Corollary A.2 with the dense mass Z_i in place of $\hat{Z}_i$, which depends on the exact set, so the ranking minimizes a residual-energy proxy of the relative error. The rule ranks blocks by how badly compensation handles them, not by the attention they hold. A block with large attention mass that compensation already predicts well has a small residual and stays skipped.

A.6. Relation to Difference Estimators

The decoupled estimator has the form of a difference estimator from survey sampling [26]: a model predicts every term cheaply, and the exact values correct the prediction on a chosen subset. Scatterbrain [1] and control-variate attention [50] use the same structure with a support chosen for each input, whereas DecoupAttn keeps the clusters and the exact blocks as two groupings joined by counts and fixes the blocks offline. This view explains why compensation changes what the selection has to capture.

Proposition A.14 (Random exact blocks). *Let every block be made exact independently with probability $\pi_b \in (0, 1]$, let I_b indicate the event, and weight each correction of the numerator and the mass by $1/\pi_b$:*

$$(\hat{N}_i^{\mathrm{HT}}, \hat{Z}_i^{\mathrm{HT}}) = \sum_c |c|\,(a_{ic}, \omega_{ic}) + \sum_b \frac{I_b}{\pi_b}\,(\Delta N_{ib}, \Delta Z_{ib}).$$

Then $\mathbb{E}\,\hat{N}_i^{\mathrm{HT}} = N_i$, $\mathbb{E}\,\hat{Z}_i^{\mathrm{HT}} = Z_i$, the output $\hat{y}_i^{\mathrm{HT}} = \hat{N}_i^{\mathrm{HT}}/\hat{Z}_i^{\mathrm{HT}}$ has the error $\hat{y}_i^{\mathrm{HT}} - y_i = \sum_b (I_b/\pi_b - 1)\, r_{ib}/\hat{Z}_i^{\mathrm{HT}}$, and with the mass at its mean Z_i

$$\mathbb{E}\Big\|\frac{1}{Z_i}\sum_b \Big(\frac{I_b}{\pi_b} - 1\Big) r_{ib}\Big\|_2^2 = \frac{1}{Z_i^2}\sum_b \frac{1-\pi_b}{\pi_b}\,\|r_{ib}\|_2^2.$$

Proof. By Proposition 1 with every block exact, $N_i = \sum_c |c|\, a_{ic} + \sum_b \Delta N_{ib}$ and $Z_i = \sum_c |c|\,\omega_{ic} + \sum_b \Delta Z_{ib}$. Hence $\hat{N}_i^{\mathrm{HT}} - N_i = \sum_b (I_b/\pi_b - 1)\Delta N_{ib}$ and $\hat{Z}_i^{\mathrm{HT}} - Z_i = \sum_b (I_b/\pi_b - 1)\Delta Z_{ib}$, and with $N_i = y_i Z_i$ the error $(\hat{N}_i^{\mathrm{HT}} - y_i\hat{Z}_i^{\mathrm{HT}})/\hat{Z}_i^{\mathrm{HT}}$ collects $\Delta N_{ib} - y_i\Delta Z_{ib} = r_{ib}$. The weights $I_b/\pi_b - 1$ have zero mean, are independent, and have $\mathbb{E}(I_b/\pi_b - 1)^2 = (1-\pi_b)/\pi_b$. □

Without compensation the same estimator has $a_{ic} = \omega_{ic} = 0$, so r_{ib} becomes the dropped residual $N_{ib} - y_i Z_{ib}$ and the error is driven by the attention a block holds. With

compensation it is driven by the residual r_{ib} that the clusters leave, which is the quantity that Eq. (9) ranks and Distortion Budget prices.

Proposition A.15 (Optimal inclusion follows the residual). *Let every $\|r_{ib}\|_2$ be positive. Among inclusion probabilities with $0 < \pi_b \le 1$ and $\sum_b \pi_b = K$ exact blocks in expectation, the error of Proposition A.14 is minimized by $\pi_b = \min\big(1, \lambda\|r_{ib}\|_2\big)$, with λ chosen to meet the budget.*

Proof. For the blocks with $\pi_b < 1$, Cauchy–Schwarz gives $\big(\sum_b \|r_{ib}\|\big)^2 \le \sum_b \pi_b \cdot \sum_b \|r_{ib}\|^2/\pi_b$, with equality exactly when $\pi_b \propto \|r_{ib}\|$. The error is $Z_i^{-2}\big(\sum_b \|r_{ib}\|^2/\pi_b - \sum_b \|r_{ib}\|^2\big)$, so it is minimized by the proportional rule. When the rule asks for $\pi_b > 1$, the constraint is active, those blocks become certain, and the rule applies to the remaining budget, which is the stated form. □

A block with $r_{ib} = 0$ adds no error at any $\pi_b > 0$, so with such blocks the optimum is approached as their $\pi_b \downarrow 0$ and the rule applies to the others.

The blocks with the largest residuals become certain, which is the residual ranking of Proposition A.13. DecoupAttn makes this rule deterministic, $\pi_b \in \{0, 1\}$ without the weights $1/\pi_b$: the blocks with the largest residuals are exact, and the others are left to compensation. Its error is then the deterministic sum of the remaining residuals (Proposition 1), and Theorem A.10 shows that the deviation part of this sum adds like random errors.

A.7. The Plan Objective (Proposition 2)

Lemma A.16 (Projected distortion). *For a head output error $A = \hat{O}_u - O_u$ and the head's slice $W_{o,h}$ of the output projection,*

$$\big\|A\,W_{o,h}^\top\big\|_F^2 = \big\langle W_{o,h}^\top W_{o,h},\, A^\top A\big\rangle,$$

so ε_u of Eq. (10) is the expected squared error that the head adds to the residual stream, and it vanishes exactly when every row of the error lies in the null space of $W_{o,h}$ almost surely.

Proof. $\|AW^\top\|_F^2 = \mathrm{tr}(WA^\top AW^\top) = \mathrm{tr}(A^\top A\,W^\top W)$ by the cyclic property of the trace. The value is zero exactly when $AW^\top = 0$. □

The projected errors of one call's heads add in the residual stream, with cross terms as in Lemma A.17.

Lemma A.17 (First-order additivity). *Let the final latent be $\hat{x}_0 = \Phi(e)$, where $e = (e_u)_u$ collects the projected head errors injected at the units and $\Phi(0)$ is the dense output. If Φ is twice continuously differentiable with $\|\nabla^2\Phi\| \le \beta$ on the segment from 0 to e, then with G_u the Jacobian of Φ with respect to e_u at zero,*

$$\hat{x}_0 - x_0 = \sum_u G_u e_u + R, \qquad \|R\|_2 \le \frac{\beta}{2}\|e\|_2^2.$$

If the propagated errors of distinct units, and of the sampler steps that one unit stands for, are orthogonal in expectation, $\mathbb{E}\langle G_u e_u, G_{u'} e_{u'}\rangle = 0$ for $u \ne u'$, and the joint Jacobian $G = (G_u)_u$ has $\|G\| \le M$, then

$$\Big|\mathbb{E}\|\hat{x}_0 - x_0\|_2^2 - \sum_u \mathbb{E}\|G_u e_u\|_2^2\Big| \le M\beta\,\mathbb{E}\|e\|_2^3 + \tfrac{\beta^2}{4}\,\mathbb{E}\|e\|_2^4.$$

Proof. The first display is Taylor's theorem with the integral remainder $R = \int_0^1 (1-s)\,\nabla^2\Phi(se)[e,e]\,ds$. Squaring it gives $\|Ge\|^2 + 2\langle Ge, R\rangle + \|R\|^2$, the cross terms between units in $\|Ge\|^2$ vanish in expectation by assumption, and $2|\langle Ge, R\rangle| \le M\beta\|e\|^3$ and $\|R\|^2 \le \beta^2\|e\|^4/4$. □

Writing $\mathbb{E}\|G_u e_u\|^2 = J_u(o,\eta)\,w_u\,\varepsilon_u(o,\eta)$ defines the sensitivity of a unit under each option as its gain from projected error to final-latent error, with w_u counting the sampler steps the unit stands for. The gain depends on the option only through the direction of the error, and Distortion Budget prices all options of a unit with the one gain $J_u = J_u(\textsc{Static}_{\theta_0}, \textsc{Full})$ measured at the reference option of Eq. (11), so the gaps $J_u(o,\eta) - J_u$ enter the modeling error δ_{m} below. Lemma A.17 then turns the final-latent error of Eq. (3) into the first line of Eq. (12). Table 21 measures directly the cross terms that the lemma sets to zero.

Lemma A.18 (The sensitivity estimator). *Suppose the units of a group $\mathcal{G}$ share one gain $J_{\mathcal{G}}$, their propagated errors are orthogonal as in Lemma A.17, the run and the dense output differ, besides the injected errors, by a zero-mean numerical error ξ, the run-to-run difference of the same dense attention backend, independent of the injected errors and with $\mathbb{E}\|\xi\|_2^2 = \mathrm{MSE}_0$, and the group's summed distortion is positive. Then, to first order, Eq. (11) returns $J_{\mathcal{G}}$.*

Proof. The run of the group gives $\hat{x}_0 - x_0 = \sum_{u\in\mathcal{G}} G_u e_u + \xi$ up to second order. Independence and zero mean remove the cross terms with ξ, so

$$\mathrm{MSE}_{\mathcal{G}} = \sum_{u\in\mathcal{G}} \mathbb{E}\|G_u e_u\|_2^2 + \mathrm{MSE}_0 = J_{\mathcal{G}} \sum_{u\in\mathcal{G}} w_u\,\varepsilon_u + \mathrm{MSE}_0.$$

Subtracting MSE_0 and dividing by the summed distortion leaves $J_{\mathcal{G}}$. □

Without the subtraction, the estimate would carry the bias $\mathrm{MSE}_0/\sum_u w_u\varepsilon_u$, which is largest exactly where two dense attention backends already disagree, on the prompt-driven few-step models. In this model $\mathrm{MSE}_{\mathcal{G}} \ge \mathrm{MSE}_0$ and the estimate is nonnegative, so a negative cell of Table 22 marks a group whose injected error lies below the run-to-run variation of the two measured errors.

Lemma A.19 (Measuring one unit at a time). *Let the projected error of a unit be $e_u = E_u(z)$, a function of the unit's input z, and let the distortion be measured at the dense input z_{d}, as in Section 4.2. If E_u is ℓ_u-Lipschitz on the inputs that*

occur, and a sparse run perturbs the input to $z = z_{\mathrm{d}} + \Delta z$, then

$$\left| \left(\mathbb{E}\|E_u(z)\|_2^2\right)^{1/2} - \varepsilon_u^{1/2} \right| \le \ell_u \Delta, \qquad \Delta = \left(\mathbb{E}\|\Delta z\|_2^2\right)^{1/2},$$

and therefore $\left|\mathbb{E}\|E_u(z)\|_2^2 - \varepsilon_u\right| \le 2\,\varepsilon_u^{1/2}\ell_u\Delta + \ell_u^2\Delta^2$.

Proof. $\|E_u(z) - E_u(z_{\mathrm{d}})\| \le \ell_u\|\Delta z\|$ pointwise, and the triangle inequality in L^2 gives the first bound. The second follows from $a^2 - b^2 = (a-b)(a+b)$ with $a+b \le 2b + \ell_u\Delta$. □

The deviation of the measured distortion is thus led by the product of the unit's own error $\varepsilon_u^{1/2}$ and the upstream error $\ell_u\Delta$, where Δz is the accumulated error of the units upstream and ℓ_u bounds how fast the unit's error changes with its input. Measuring every unit at the dense input therefore captures its distortion up to a product of two errors whenever ℓ_u is bounded. It also keeps the measurement of a unit independent of the plan, so the distortion table is built once and every candidate plan is priced from it.

Cost of solving. For one multiplier, the per-call minimization visits every unit and option once under each of the two estimators, which costs $O(|\mathcal{U}|\,|\mathcal{O}|)$ operations on a table that is already measured. Bisection repeats this B_μ times and the greedy pass sorts the candidate changes once, for $O\big(|\mathcal{U}|\,|\mathcal{O}|\,B_\mu + |\mathcal{U}|\log|\mathcal{U}|\big)$ in total. This is why plan solving takes under a minute in Table 30, while measuring ε_u and J_u takes the GPU hours.

Proof of Proposition 2. Plans π (Definition 2) range over a finite set, $F(\pi) = \sum_u w_u J_u \varepsilon_u(o_u, \eta_{tl})$ is the priced distortion, and $X(\pi) = \sum_u w_u X_u(o_u)/w_\Sigma$ the density of Eq. (2), with $w_\Sigma = \sum_u w_u$. For a multiplier $\mu \ge 0$, let π_μ be a global minimizer of the Lagrangian of Eq. (12), $\mathcal{L}_\mu(\pi) = F(\pi) + \mu\, w_\Sigma X(\pi)$.

Exact minimization. $\mathcal{L}_\mu$ is a sum over calls (t,l) of $\sum_h w_u\big(J_u\varepsilon_u(o_u,\eta_{tl}) + \mu X_u(o_u)\big)$, and the estimator couples only the heads of one call. Minimizing each call over η_{tl} of $\sum_h \min_{o_u}(\cdot)$ therefore minimizes $\mathcal{L}_\mu$ exactly, which is the procedure of Section 4.2. Adding the two minimality inequalities shows that any minimizers at $\mu_1 < \mu_2$ have $X(\pi_{\mu_2}) \le X(\pi_{\mu_1})$, whichever ties are taken, so bisection on μ is well defined.

Optimality at its own density. Let π be any plan with $X(\pi) \le X(\pi_\mu)$. Because π_μ minimizes $\mathcal{L}_\mu$,

$$F(\pi_\mu) \le F(\pi) + \mu\, w_\Sigma\big(X(\pi) - X(\pi_\mu)\big) \le F(\pi),$$

so π_μ has the lowest priced distortion among all plans at most as dense as itself [4]. Bisection returns the feasible π_μ of largest density, and the greedy pass then lowers F within the remaining slack $\rho - X(\pi_\mu)$, which Lemma A.20(i) covers. □

Lemma A.20 (Gap of the relaxation). *Let $F^\star(\rho)$ be the lowest priced distortion over plans with $X(\pi) \le \rho$. (i) If bisection returns a feasible exact minimizer π_{b} of $\mathcal{L}_\mu$ for some $\mu \ge 0$, then*

$$F(\pi_{\mathrm{b}}) - F^\star(\rho) \le \mu\, w_\Sigma\big(\rho - X(\pi_{\mathrm{b}})\big),$$

and every plan that the greedy pass reaches from π_{b} by feasible changes that do not raise F keeps this bound. (ii) Let $W_{\max}$ be the largest total weight $\sum_h w_u$ of one call and $\mu^\star$ the multiplier at which the minimizers of $\mathcal{L}_\mu$ cross the budget. If, at $\mu^\star$, the minimizers are switched one call at a time, in any order, from an infeasible one toward a feasible one and the first feasible plan is kept, the bound of (i) is below $\mu^\star W_{\max}$.

Proof. (i) For every feasible π, minimality of π_{b} gives $F(\pi) \ge F(\pi_{\mathrm{b}}) + \mu\, w_\Sigma\big(X(\pi_{\mathrm{b}}) - X(\pi)\big) \ge F(\pi_{\mathrm{b}}) - \mu\, w_\Sigma\big(\rho - X(\pi_{\mathrm{b}})\big)$, since $X(\pi) \le \rho$. Taking the minimum over feasible π gives the bound, and a later plan with no larger F inherits it. (ii) If a minimizer of $\mathcal{L}_0 = F$ is feasible, it attains $F^\star(\rho)$ and the gap is zero. Otherwise, as a function of μ, the set of minimizers of $\mathcal{L}_\mu$ changes only at finitely many breakpoints, and at the breakpoint $\mu^\star$ it contains a plan π_- with $X(\pi_-) > \rho$ and a plan π_+ with $X(\pi_+) \le \rho$. Because $\mathcal{L}_{\mu^\star}$ separates over calls, every plan that takes a minimizing choice in each call is again a minimizer, so π_- turns into π_+ one call at a time through minimizers. The first feasible plan of this chain follows an infeasible one that differs from it in one call, so its slack $\rho - X$ is below $W_{\max}/w_\Sigma$, and (i) bounds its gap by less than $\mu^\star W_{\max}$. □

The bound of (i) is known once the plan is returned, and (ii) caps it at the price of a single call's density for a solver that switches tied minimizers one call at a time. Our solver keeps the whole minimizer at the feasible end of bisection, and on the deployed plans of the eight models it solves the bound of (i) is at most 0.54% of the priced distortion (Cosmos-Predict2.5) and at most 0.094% on the others, with at most 0.026% left after the greedy pass. With every step, layer, and head of a video in the plan, one call holds a small share of w_Σ, so the relaxation is nearly tight.

The gap concerns the priced distortion F, and it carries over to the final-latent error $D(\pi) = \mathbb{E}_x\|\hat{x}_0(\pi) - x_0\|_2^2$ of Eq. (3) through the modeling error. If $|D(\pi) - F(\pi)| \le \delta_{\mathrm{m}}$ on every feasible plan and $F(\hat{\pi}) \le F^\star(\rho) + \epsilon_{\mathrm{o}}$, then for every feasible π

$$D(\hat{\pi}) \le F(\hat{\pi}) + \delta_{\mathrm{m}} \le F(\pi) + \epsilon_{\mathrm{o}} + \delta_{\mathrm{m}} \le D(\pi) + \epsilon_{\mathrm{o}} + 2\delta_{\mathrm{m}},$$

so $\hat{\pi}$ is within $\epsilon_{\mathrm{o}} + 2\delta_{\mathrm{m}}$ of the best feasible plan in final-latent error. Lemmas A.17–A.19 trace where δ_{m} comes from: the cross terms between units, the shared gain of a group and of the options of a unit, and the measurement at the dense input.

A.8. Reassign Cache (Proposition 3)

Proof of Proposition 3. With the centroids m_c fixed, the total squared distance $\sum_j \|k_j^{(t)} - m_{c(j)}\|_2^2$ is a sum of one

term per key, and Eq. (13) minimizes each term separately. Hence, for any other assignment c', including the frozen assignment of the first sparse step,

$$\sum_j \left\|k_j^{(t)} - m_{c_t(j)}\right\|_2^2 \le \sum_j \left\|k_j^{(t)} - m_{c'(j)}\right\|_2^2,$$

which is the assignment step of Lloyd's algorithm [19]. The exact blocks $\mathcal{S}_g$ come from the offline ranking of Eq. (9), and the options and estimators from the offline solution of Eq. (12). Neither reads the assignments at inference, so reassignment leaves both unchanged. □

Lemma A.21 (Monotone chain with recomputed statistics). *Let c_t be the assignment after reassignment at step t, $\bar{k}_c^{(t)}$ the means of the reassigned clusters computed from their current keys, and c_f the frozen assignment. Then*

$$\sum_j \left\|k_j^{(t)} - \bar{k}_{c_t(j)}^{(t)}\right\|_2^2 \le \sum_j \left\|k_j^{(t)} - m_{c_t(j)}\right\|_2^2 \le \sum_j \left\|k_j^{(t)} - m_{c_f(j)}\right\|_2^2.$$

Proof. Within every cluster the mean minimizes the sum of squared distances, which gives the first inequality, and Proposition 3 gives the second. □

The left side is the within-cluster spread $\sum_c |c| \operatorname{tr}(C_c)$ that enters the compensation error. By Lemma A.4, the second-order mass term of a cluster is $|c|\,\phi_{ic}\,q_i^\top C_c q_i/(2d)$, and for Gaussian queries $q_i \sim \mathcal{N}(0, I)$ its mean is $|c|\,e^{\|a_c\|^2/2}\big(\operatorname{tr} C_c + a_c^\top C_c a_c\big)/(2d)$ with $a_c = \bar{k}_c/\sqrt{d}$, which grows with the spread $\operatorname{tr} C_c$ and with its part along the centroid. Reassignment thus lowers an upper bound on the spread, the distance to the cached centroids, at the cost of one product between the keys and the cached centroids, LCd multiply–adds, followed by an argmin. A cluster that loses all of its members has $|c| = n_{gc} = 0$ and drops out of the estimator and of H, and its cached centroid stays available to later assignments. Reassignment with recomputed statistics is one step of k-means whose centroids for the next assignment stay at the cached m_c, so the assignments cannot drift away from the clusters that the plan was priced on. Proposition 3 concerns the distance to the cached centroids, and the Frozen ablation, which keeps the assignments and recomputes the cluster statistics, measures the quality gain of reassignment directly (Tables 2 and 25).

A.9. Cost of One Call

Count multiply–adds per head for L tokens, head dimension d, C clusters, and density X (block padding included). Dense attention costs $2L^2d$, for $QK^\top$ and for the weighted sum of values. The exact part costs $2XL^2d$. The compensation computes the centroid logits $q_i^\top \bar{k}_c$ for every query and cluster (LCd), the variance corrections $\sum_p q_{ip}^2 \sigma_{cp}^2$ (LCd), the numerator $\sum_c n_{gc} a_{ic}$ (LCd), and the products Hq_i (Ld^2). The statistics $\bar{k}_c$, $\bar{v}_c$, σ_c^2, and H cost $O(Ld + Ld^2)$, reassignment costs LCd, and the counts n_{gc} are bookkeeping over the $N_B \times C$ table of n_{bc}. Relative to dense attention, one call therefore costs

$$r = X + \big(2 + \mathbf{1}_\mathrm{F} + \mathbf{1}_\mathrm{R}\big)\frac{C}{2L} + \mathbf{1}_\mathrm{F}\,\frac{d}{L},$$

where $\mathbf{1}_\mathrm{F} = 1$ for the FULL estimator, which adds the variance corrections and H, and $\mathbf{1}_\mathrm{R} = 1$ when Reassign Cache reassigns. On Cosmos-Predict2.5, the latent has $L = 24 \cdot 44 \cdot 80 = 84{,}480$ tokens, with $C = 880$ and $d = 128$, so for the full method the two overhead terms add about 0.021 and 0.0015 to X. In multiply–adds, compensation, counts, and reassignment thus cost about one tenth of the exact part at $X = 0.20$, and the k-means of the first sparse step adds $I \cdot LCd$ once per video for I rounds. If attention takes a share f of the latency and runs r_time times as long as dense attention, the speedup is

$$S = \frac{1}{(1-f) + f\,r_\mathrm{time}},$$

where the measured time ratio r_time adds to r the memory traffic and kernel launches that multiply–adds leave out. The share f differs across models, which is why one density gives different speedups in Table 1. The measured components of Table 27 map onto these terms: *Exact* is $2XL^2d$, *Statistics*, *Counts*, and *Merge* are the compensation, *Assign* is reassignment, and *k-means* is the one-time clustering.

A.10. Error Propagation Through the Sampler

The sensitivity J_u measures how much of a unit's error reaches the final latent. This section shows where its spread across steps and models comes from, and why the disagreement between two dense attention backends probes it.

Lemma A.22 (Amplification along the sampler). *Let the sampler update the latent by $x_{s-1} = \Psi_s(x_s)$ for $s = T, \dots, 1$, where Ψ_s contains the network evaluation of step s and is L_s-Lipschitz on the latents that occur. An error e added to the latent after step t changes the final latent by at most*

$$\left\|\hat{x}_0 - x_0\right\|_2 \le \Big(\prod_{s=1}^{t-1} L_s\Big)\|e\|_2.$$

If the Ψ_s are also differentiable along the dense trajectory, the change is, to first order, $\big(\prod_{s=1}^{t-1} \partial\Psi_s\big)e$, the product of the remaining step Jacobians applied to e.

Proof. Let $\hat{x}_{t-1} = x_{t-1} + e$. For every later step, $\|\hat{x}_{s-1} - x_{s-1}\| = \|\Psi_s(\hat{x}_s) - \Psi_s(x_s)\| \le L_s\|\hat{x}_s - x_s\|$, and induction down to $s = 1$ gives the bound. The first-order form follows from the chain rule. □

An attention error e_u at step t enters the latent through the network evaluation inside Ψ_t, to first order as $B_u e_u$ with B_u the Jacobian of Ψ_t with respect to the unit's projected output, and then follows the same products, so $G_u =$

$\left(\prod_{s=1}^{t-1} \partial\Psi_s\right) B_u$ and J_u is the squared gain of G_u in the direction of the unit's projected error, averaged over inputs and normalized as in Lemma A.17. Few-step samplers take large steps, and each Jacobian of a large step can stretch some directions strongly, so the products, and with them the sensitivities, differ by orders of magnitude across steps and models. A per-pixel condition anchors the output to the input, so the step maps contract the directions the condition determines and the products stay small.

Corollary A.23 (The backend floor probes the same products). *Let two dense attention backends differ at step s by a numerical error ξ_s, the difference of their updates of x_{s-1}, which given x_s has zero mean and covariance Σ_s and is uncorrelated across steps. Then, to first order and with $P_s = \prod_{s'=1}^{s-1} \partial\Psi_{s'}$ the Jacobian product from x_{s-1} to the output along the dense trajectory,*

$$\mathbb{E}\left\|x_0^{(1)} - x_0^{(2)}\right\|_2^2 = \sum_s \operatorname{tr}\left(P_s \Sigma_s P_s^\top\right) \leq \sum_s \Big(\prod_{s'=1}^{s-1} L_{s'}\Big)^2 \operatorname{tr}\Sigma_s.$$

Proof. By Lemma A.22, each ξ_s reaches the output as $P_s\xi_s$ to first order, and $\mathbb{E}\|P_s\xi_s\|^2 = \operatorname{tr}(P_s\Sigma_s P_s^\top)$. The terms between steps vanish, and $\|P_s\| \leq \prod_{s'<s} L_{s'}$ bounds each term. □

The floor weights the Jacobian products by the covariance of the numerical noise, and the sensitivities weight the same products by the directions of the sparse errors, so the floor probes the sensitivities. The Floor of Table 1 reads this disagreement on the decoded frames, as a PSNR averaged over clips. Figure 4d–f and Table 12 measure this: the cost of removing the sensitivity falls as the floor rises, with rank correlation -0.61 on ten models.

B. Constants and Implementation

Quantity	Value
Clusters per head, k-means rounds	880, 10 (first sparse step)
Density levels Θ	0.05, 0.10, 0.15, 0.20, 0.30, 0.45
Target density ρ	0.200 (reported band 0.195–0.205)
Kernel block size for density accounting	128×96, dense units count one
Always dense	first step, first layer
Reassignment threshold τ_s	1.0 dB
Speedup threshold τ_g	1.10
Reference level θ_0 for J	0.20
Block ranking $\mathbf{D}_u$	16 calibration clips
Distortion ε_u, sensitivity J_u	4 calibration clips
Block shapes tried (frames×rows×cols)	2×4×16, 1×8×16, 2×8×8, 4×4×8, 1×16×8
Seed	1000

Table 4. Constants of DecoupAttn, shared by all models.

Shared constants. All ten models use the constants of Table 4, and what differs between them is measured on their own clips: the block shape, the block ranking $\mathbf{D}_u$, the distortion table ε_u, the sensitivity J_u, and the assignment rule of Eq. (14). The per-model plans come from Distortion Budget's measurements, not tuning, so one procedure serves every family and sampler (Section 5.2).

Block shape. Five block shapes are scored by the attention a density-0.20 static mask keeps on eight calibration clips; the default and the best challenger are compared on eight validation clips, a gap of at least 0.1 dB decides, and otherwise the faster shape is kept.

Sensitivity groups. J is measured on groups of sampler steps (one step for four-step models) and one third of the layers, run at level θ_0 with full compensation while the rest stays dense. Every unit in a group shares its J_u, and J is normalized to mean one, so $J_u > 1$ marks a unit whose errors reach the final latent more than the average unit's. Cells that the floor subtraction leaves negative are set to zero before this normalization, and for the models measured by injection a group without a positive excess stops the plan build. No cell of the deployed plans is affected under the reference backend. DiffusionRenderer uses one weight per step from an earlier sensitivity map of the same quantity.

Speed check. The implementation also measures $S_{\max}$, the dense latency divided by the lowest latency any sparse plan reaches, and runs dense attention when $S_{\max} < \tau_g$. The experiments in this paper do not use this check.

Kernels. The exact blocks run on FlashInfer's variable-block FlashAttention-3 kernel. Each call computes the cluster means, the per-dimension key variances, and the matrix H (two GEMMs) from the current keys and values, and a fused Triton kernel adds the cluster terms of Eq. (5) and merges them with the exact part through the log-sum-exp. Centroid-only calls skip the variance pass and H.

C. Experimental Details

Model	Steps	Video	Test clips
Cosmos-Predict1	35 + CFG	1280×704, 121 fr.	39
Cosmos-Predict2.5	4	720p, 93 fr.	32
Wan2.1-14B	4	832×480, 81 fr.	32
Wan2.2-A14B	4	480p, 81 fr.	32
HunyuanVideo	50	1280×720, 129 fr.	8
Cosmos3-Nano (I2W)	35 + CFG	720p, 189 fr.	32
Cosmos-Transfer2.5	4	edge control, 93 fr.	39
DiffusionRenderer	4	1280×704, 57 fr.	39
AlayaRenderer	50 + CFG	latent 21×30×52	39
UniRelight	35 + CFG	latent 24×30×53	39

Table 5. Models, samplers, and test splits. Calibration, validation, and test clips are pairwise disjoint.

Models. Table 5 covers five four-step distilled samplers and five samplers with 35 or 50 steps, four of them with classifier-free guidance, with 8 to 39 test clips per model. Both sampler types run with the constants of Table 4, so what adapts to each model is its calibrated plan.

Density accounting and baseline settings. Density counts every computed block after padding to the kernel block size, with dense calls counted as one. Baselines run from their official code with their density knob set to 0.20; where the official configuration sets a higher per-head minimum, the baseline runs at that minimum (SVG2 and SVG-EAR at 0.225 and 0.224 on Wan2.2-A14B). On Cosmos-Predict1 and UniRelight, which have no official CalibAtt schedule, its energy schedule follows the official 50-step Wan schedule, scaled to density 0.20 on the calibration set.

Timing. Each timing process runs dense attention, the method, and dense attention again, with one warm-up and one timed generation each, measured with CUDA events; latency is the wall-clock time of the denoising loop, and the speedup is the mean dense latency divided by that of the method. Plan building, clustering, and reassignment are all counted inside the measured latency.

Ablation variants. Tied groups compensation by the selection blocks; Centroid drops both corrections; U sets every J_u to one; D replaces ε_u by the error of the same option without compensation; Frozen keeps the assignments of the first sparse step on the nine models where Eq. (14) reassigns them; Rebuild runs k-means at every sparse step (Table 9). Tied† applies the same grouping swap under our plan without re-solving, and every other variant is re-solved at density 0.20. DiffusionRenderer already keeps its assignments under the rule of Eq. (14), where reassigning them changes PSNR by $+0.05$ dB. In Table 6 the Full column gives the full method's PSNR with its standard deviation across clips.

Ablation results. Table 6 adds the spread and the clip count behind each cell of Table 2. The decoupling variants lose most consistently: Tied and Tied† lose on at least six of the eight clips of every model and Centroid on at least seven, since all three enlarge the residual that every query carries (Proposition 1). The cost of U depends more on the model, because it follows how unevenly the sampler amplifies the errors of different units (Table 12): U costs 4.10 dB on HunyuanVideo but 0.55 dB on Cosmos-Transfer2.5, whose edges damp the error growth (Figure 4d,e). Frozen costs at most 1.14 dB on the video models and most on AlayaRenderer and UniRelight (2.59 and 4.45 dB).

Speed of the variants. Table 8 shows what each component costs in time. U and D change only the prices, so their speedup stays within 0.03 of the full method's, and Distortion Budget gains its 1.18–1.68 dB at no cost in speed. Tied and Tied† group by the selection blocks and skip the clustering, which makes them slightly faster (0.04–0.15 in speedup) but costs 0.85–8.59 dB (Table 6), so the semantic clusters of Section 4.1 are worth their time. Frozen saves the per-step assignment, at most 0.08 in speedup, and Rebuild, which runs k-means at every sparse step, falls below dense speed on five of the ten models.

Re-clustering and reassignment. Rebuild changes PSNR by at most 0.95 dB (UniRelight) and by less than 0.4 dB on the other nine models, while adding 0.8–96.1 s per video (Table 9). Reassignment already moves every key to its nearest cached centroid (Eq. (13)), which never increases the spread (Proposition 3), so fresh centroids gain little, and the full method wins on five of the eight clips of HunyuanVideo and Cosmos3-Nano. In the timing session of Table 10, reassignment adds 108–227 ms (2–3%) over frozen assignments, whereas Rebuild adds 0.8–1.8 s (14–21%) and makes DiffusionRenderer and Wan2.1-14B slower than dense attention, which is why Reassign Cache keeps the cached centroids instead.

VBench. Table 7 scores each video on its own, without the dense output. DecoupAttn stays within 0.007 of dense attention on every model and is the best or tied-best sparse method on four of them, and on Cosmos3-Nano, where SVOO and SVG-EAR fall to 0.7988 and 0.8100, it keeps the dense score (0.8616 against 0.8612).

Coherence measurement. Table 11 takes the (clip, step, layer, head) units of the deployed plans on four calibration clips with seed 1000 (conditional branch, sparse steps, layers after the first): 10,040 units on DiffusionRenderer, 17,968 on Wan2.1-14B, and 4,948 on Cosmos-Predict2.5. Each unit pools the sums over its queries, $\kappa = \sum_i \| \sum_b r_{ib} \|_2^2 / \sum_i \sum_b \| r_{ib} \|_2^2$, in value space with dense probabilities ($Z_i = 1$), so high-energy queries weigh most. The normalized ratio $\zeta = (\kappa - 1)/(K_g - 1)$ divides out the number K_g of compensated key blocks $b \notin \mathcal{S}_g$ per query (medians 198, 319, and 594) and equals the mean pairwise cosine of the residuals when their norms are equal. Compensation lowers κ on all three models, and the corrections lower it further: the median falls 2.9–3.9-fold, from 4.69–10.37 without compensation to 2.67–5.12 with plain centroids and 1.61–2.69 with DecoupAttn, and the median ζ falls to 0.002–0.003, so the residuals of a few hundred compensated blocks are nearly orthogonal. This is the cancellation of Proposition 1: a compensated residual comes from member deviations that sum to zero over each cluster,

Model (steps)	Full	Decoupling (§4.1)			Distortion Budget (§4.2)		Reassign Cache (§4.3)
		Tied	Tied†	Centroid	U	D	Frozen
Video generation							
Cosmos-Predict1 (35)	23.42±3.91	−1.95±2.14 (6)	−2.45±1.60 (8)	−1.60±1.64 (8)	−1.93±2.97 (7)	−0.34±2.21 (5)	−0.95±0.90 (7)
Cosmos-Predict2.5 (4)	24.07±3.97	−1.63±0.71 (8)	−1.90±0.84 (8)	−1.03±0.67 (7)	−2.30±0.94 (8)	−0.96±0.39 (8)	−0.38±0.18 (8)
Wan2.1-14B, 480p (4)	25.31±2.80	−3.72±0.75 (8)	−4.19±1.04 (8)	−1.24±0.50 (8)	−2.14±0.79 (8)	−1.65±0.41 (8)	−1.14±0.31 (8)
Wan2.2-A14B (4)	17.21±2.87	−1.48±0.53 (8)	−1.66±0.36 (8)	−0.53±0.24 (8)	−1.06±0.56 (8)	−0.65±0.51 (8)	−0.36±0.18 (8)
HunyuanVideo (50)	25.84±4.57	−1.61±1.53 (6)	−0.85±1.05 (7)	−1.06±1.31 (7)	−4.10±2.43 (8)	−0.87±0.81 (8)	−0.38±0.68 (6)
Cosmos3-Nano I2W (35)	20.92±2.94	−1.71±2.28 (7)	−2.17±1.73 (8)	−0.60±1.93 (7)	−2.05±1.22 (8)	−1.06±2.25 (7)	−0.12±1.37 (5)
Controlled generation							
Cosmos-Transfer2.5 (4)	33.34±2.01	−2.21±0.78 (8)	−2.67±0.90 (8)	−1.20±0.62 (8)	−0.55±0.51 (6)	−1.31±0.68 (8)	−0.69±0.65 (8)
Neural rendering							
DiffusionRenderer (4)	50.20±1.44	−2.75±0.64 (8)	−5.29±2.58 (8)	−2.48±0.86 (8)	−0.77±0.47 (7)	−1.71±0.91 (8)	–
AlayaRenderer (50)	27.97±2.81	−4.07±1.51 (8)	−5.43±1.81 (8)	−1.43±0.90 (7)	−1.06±1.37 (7)	−0.86±0.95 (6)	−2.59±1.07 (8)
UniRelight (35)	41.54±4.39	−7.39±3.89 (8)	−8.59±4.42 (8)	−3.77±1.79 (8)	−0.85±3.89 (3)	−2.41±0.69 (8)	−4.45±2.09 (8)

Table 6. **Component ablations with standard deviations and clip counts** (full version of Table 2). Each cell is the PSNR of the variant minus that of the full method (dB) on eight validation clips, with the standard deviation of this difference (small ±) and the number of clips on which the full method is better. Variants are defined in Appendix C, their latencies are in Table 2, and their speedups are in Table 8.

Model (steps)	Floor	CalibAtt	SVOO	SVG2	SVG-EAR	Sol-Attn	PISA	SparsePR	**Ours**
Video generation									
Cosmos-Predict1 (35)	0.8375	0.8177	0.8273	0.8331	0.8360	0.8342	0.8339	**0.8364**	0.8313
Cosmos-Predict2.5 (4)	0.8623	0.8407	0.8633	0.8600	0.8621	**0.8646**	0.8597	0.8624	**0.8646**
Wan2.1-14B, 480p (4)	0.8668	0.8589	0.8664	0.8655	0.8675	**0.8689**	0.8677	0.8665	0.8673
Wan2.2-A14B (4)	0.8636	0.8638	0.8642	0.8622	0.8634	**0.8669**	0.8662	0.8628	0.8639
HunyuanVideo (50)	0.8500	0.8470	0.8469	0.8445	0.8499	**0.8548**	0.8544	0.8517	0.8468
Cosmos3-Nano I2W (35)	0.8612	0.8457	0.7988	0.8273	0.8100	0.8456	0.8585	0.8499	**0.8616**
Controlled generation									
Cosmos-Transfer2.5 (4)	0.8320	0.8107	0.8173	0.8208	0.8292	0.8264	0.8179	0.8310	**0.8335**
Neural rendering									
DiffusionRenderer (4)	0.8342	0.8315	0.8279	0.8306	0.8341	**0.8347**	0.8334	0.8344	0.8343
AlayaRenderer (50)	0.8269	0.8082	0.8136	0.8161	0.8244	0.8213	0.8232	0.8172	**0.8268**
UniRelight (35)	0.8330	0.8297	0.8284	0.8309	0.8331	**0.8336**	0.8330	0.8325	0.8333

Table 7. **VBench** for the models of Table 1: mean of subject consistency, background consistency, temporal flickering, motion smoothness, imaging quality, and aesthetic quality. SSIM and LPIPS with their standard deviations are in Table 1. Bold: best sparse method.

Model (steps)	Full	Tied	Tied†	Centroid	U	D	Frozen	Rebuild
Video generation								
Cosmos-Predict1 (35)	1.47	1.56	1.55	1.49	1.45	1.48	1.53	1.18
Cosmos-Predict2.5 (4)	1.48	1.63	1.62	1.50	1.47	1.48	1.51	1.33
Wan2.1-14B, 480p (4)	1.06	1.17	1.16	1.08	1.06	1.06	1.08	0.95
Wan2.2-A14B (4)	1.07	1.17	1.17	1.08	1.07	1.07	1.09	0.95
HunyuanVideo (50)	1.65	1.71	1.69	1.71	1.62	1.62	1.67	1.42
Cosmos3-Nano I2W (35)	1.21	1.28	1.28	1.25	1.20	1.21	1.26	0.97
Controlled generation								
Cosmos-Transfer2.5 (4)	1.47	1.60	1.60	1.48	1.46	1.47	1.49	1.32
Neural rendering								
DiffusionRenderer (4)	1.11	1.23	1.24	1.07	1.11	1.11	–	0.92
AlayaRenderer (50)	1.24	1.35	1.34	1.27	1.24	1.24	1.32	0.87
UniRelight (35)	1.30	1.39	1.38	1.33	1.29	1.30	1.36	1.01

Table 8. **Speedup over dense attention of every ablation variant** (validation clips, same process as its dense anchor). DiffusionRenderer already freezes its cluster assignments.

Model (steps)	ΔPSNR (dB)	Clips	Δms
Cosmos-Predict1 (35)	+0.35±0.30	2	+42,257
Cosmos-Predict2.5 (4)	+0.07±0.05	1	+944
Wan2.1-14B, 480p (4)	+0.20±0.10	0	+1,543
Wan2.2-A14B (4)	+0.07±0.09	1	+1,526
HunyuanVideo (50)	+0.31±0.86	5	+96,081
Cosmos3-Nano I2W (35)	−0.36±0.61	5	+47,008
Cosmos-Transfer2.5 (4)	0.00±0.24	2	+1,089
DiffusionRenderer (4)	+0.09±0.21	3	+825
AlayaRenderer (50)	+0.03±1.05	4	+19,443
UniRelight (35)	+0.95±0.89	2	+31,747

Table 9. **Re-clustering at every step** (Rebuild) on the validation clips of Table 2. Δ is Rebuild minus the full method with its standard deviation over the clips (small ±), and Clips counts the clips (of 8) on which the full method is better.

Model	Frozen	Reassign	Rebuild	Rule
DiffusionRenderer	3,904.0	4,012.2	4,729.7	frozen
Cosmos-Predict2.5	8,053.8	8,219.5	9,165.9	reassign
Wan2.1-14B	12,710.6	12,937.7	14,508.7	reassign

Table 10. Latency (ms) of the three assignment schemes in one timing session; dense is 4,329, 12,170, and 13,747 ms.

so the skipped blocks share almost no common direction, and blocks fixed offline suffice (Section 5.3).

Sensitivity spread. Table 12 relates the cost of U to how unevenly J spreads over the cells a plan uses. On the four-step video models the cost tracks the spread (coefficient of variation 1.11–2.23, cost 1.06–2.30 dB). A flat J misprices units when their true sensitivities differ. Cosmos-Transfer2.5 and UniRelight spread J widely yet lose less than 0.9 dB without it, and both have high dense-backend floors (41.93 and 55.49 dB): their conditions damp the error growth along the sampler, which leaves less for the sensitivity to steer (Figure 4d–f).

Model	Compensation	κ (25/50/75%)	median ζ
DiffusionRenderer	none	2.83 / 4.69 / 7.71	0.021
	centroid only	1.87 / 2.69 / 4.24	0.010
	ours	1.26 / 1.61 / 2.12	0.003
Wan2.1-14B	none	2.15 / 5.03 / 8.96	0.014
	centroid only	1.86 / 2.67 / 3.66	0.006
	ours	1.20 / 1.65 / 2.18	0.002
Cosmos-Predict2.5	none	6.29 / 10.37 / 16.96	0.017
	centroid only	3.22 / 5.12 / 7.36	0.008
	ours	2.02 / 2.69 / 3.42	0.003

Table 11. **Coherence ratio** κ of Eq. (7) as quartiles over the units of the deployed plans, with the median of the normalized ratio ζ.

Model	CV of J	Cost of U (dB)	Floor (dB)
Cosmos-Predict2.5	2.23	2.30	27.78
Wan2.1-14B	1.99	2.14	24.93
Wan2.2-A14B	1.11	1.06	21.10
Cosmos-Transfer2.5	1.80	0.55	41.93
DiffusionRenderer	0.62	0.77	54.70
AlayaRenderer	1.02	1.06	39.13
UniRelight	1.76	0.85	55.49

Table 12. Spread of the sensitivity (coefficient of variation over the cells the plan uses), the cost of removing it (Table 2), and the dense-backend Floor column (Table 1).

D. Supplementary Analyses

This appendix isolates each mechanism of Section 4 and tests whether its gain carries over to matched latency, unseen inputs, and task metrics. Every comparison shares prompts, seeds, conditioning inputs, and the dense reference, and runs on clips that took no part in calibration or validation, after every variant is fixed; Table 13 lists the models that each analysis uses. Cosmos-Predict2.5 is in every analysis, and the other models add another video family (Wan2.1-14B), many- and four-step rendering (UniRelight, DiffusionRenderer), and edge control (Cosmos-Transfer2.5), and the task metrics cover nine of the ten models, all but UniRelight.

Statistics. Mechanism diagnostics record the queries, keys, and values of selected (step, layer, head) units on eight held-out clips without storing attention maps, plan screening uses eight held-out clips, and the final comparisons use 8 to 39 prompts or conditions per model with one seed each. We report paired differences with 95% bootstrap confidence intervals, resampling prompts, together with the win rate and the worst-decile difference; frames, heads, and layers are never counted as independent samples.

Model	Analyses (Appendix)
All ten models	D.1, D.6 (rule)
Cosmos-Predict2.5	D.2–D.9
Wan2.1-14B	D.2, D.3, D.5, D.7–D.9
UniRelight	D.6, D.7
DiffusionRenderer	D.5, D.6, D.9
Cosmos-Transfer2.5	D.4, D.9
Other five models	D.9

Table 13. **Models of each supplementary analysis**, by the subsections of Appendix D that use them.

D.1. Density Accounting

Table 14 separates the always-dense first step and first layer from the units the plan sparsifies: f_{dense} is their share of the attention work, X_{sp} is the density of the remaining units, and the whole run computes

$$\rho_{\text{global}} = f_{\text{dense}} + (1 - f_{\text{dense}})\, X_{\text{sp}}.$$

Both guidance branches of the CFG models, the first step, the first layer, and every diagonal block count as computed, and every block is padded to the kernel block size. All ten plans spend density 0.200 on the units they sparsify, and the global density differs only through the dense share: the first step and the first layer take 27–28% of the attention work on the four-step models but 3.6–6.3% on the many-step models, giving global densities of 0.415–0.421 and 0.229–0.251. All methods keep these parts dense (Section 5.1), so this share bounds every method's speedup.

D.2. Decoupling

Grouping under a fixed mask. Table 15 holds the exact mask, the per-unit exact budget, the block size, the assignment scheme, and the precision fixed, and changes only how the skipped keys are summarized: by nothing, by position (the selection blocks), or by semantic clusters, each with plain centroids or with both corrections of Eq. (4). Figure 6 varies the number of groups of both schemes and plots the projected attention error against the compensation time. With the mask fixed and plain centroids, semantic clusters raise PSNR over position groups by 5.11 dB on Cosmos-Predict2.5 and 3.72 dB on Wan2.1-14B (Table 15), and with both corrections they lower the projected error at every group count in the diagnostic implementation of Figure 6, 3.6-fold and 7.1-fold at the deployed 880 clusters. The corrections help both groupings, yet position groups with both corrections (22.33 and 20.66 dB) remain below semantic clusters without them (23.47 and 23.41 dB): the corrections repair the mass and the first-order value of each group but cannot pull far-apart members together, and the spread of the members is what sets the residual (Proposition A.6). The semantic clusters cost 0.7–1.4 s more DiT time than position groups, which buys 2.37 and 4.04 dB with both corrections on, so decoupling the grouping from the selection pays on both models.

Selection under a fixed compensator. Table 16 keeps the semantic compensator, the cluster assignments, the per-unit exact budget, and the kernels fixed and changes only how the exact blocks are chosen. Online residual selection ranks the blocks of each call by their residual under that call's dense attention, information that no deployed method reads. The last row attaches the same compensator to the masks of CalibAtt [43]. Under one compensator, ranking blocks by their residual beats ranking them by attention mass by 0.32 and 0.33 dB at the same cost (Table 16), because the plan keeps exact the blocks that compensation handles worst, while a block with high attention mass may be one the centroids already predict well. Routing by mass at every call gains 0.27 dB over offline mass ranking on Cosmos-Predict2.5 for 650 ms of routing, loses 0.42 dB on Wan2.1-14B for 1,288 ms, and stays below our offline blocks on both, since with every key compensated the residuals cancel, and the remaining error depends on the blocks that compensation handles worst, which calibration measures once (Proposition 1). The blocks themselves matter: random blocks lose 7.44 and 8.34 dB, and the masks of CalibAtt reach 21.22 and 20.79 dB under our compensator, 3.48 and 3.91 dB below our blocks. The diagnostic online residual gains 2.33–3.27 dB at 86–143 times the latency, while our fixed blocks add no routing time.

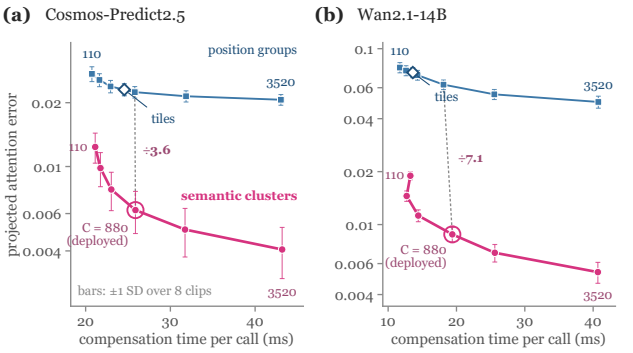

Figure 6. **Group count.** Projected attention error against compensation time as the number of position or semantic groups per head varies, with both corrections on (Cosmos-Predict2.5, Wan2.1-14B). Time: CUDA time per attention call of the diagnostic implementation, exact part and compensation, without the assignment refresh. Bars: ±1 standard deviation over eight clips.

D.3. Compensation Corrections

Factorial ablation. Table 17 switches the variance correction γ_{ic} and the first-order value correction H of Eq. (4) on and off. Under a fixed plan, the exact mask, the per-unit options, the assignments, and the set of FULL calls are held, so the solver cannot route around a removed term; under a re-solved plan, every variant is re-solved at the same density. The mass error $|\tilde{Z}_i/Z_i - 1|$ isolates what a centroid misses in the denominator, and the projected error of Eq. (10) measures the resulting output error. Under a fixed plan, each correction alone moves PSNR by -0.16 to $+0.28$ dB, but together they add 1.23 and 1.29 dB (Table 17). The variance correction restores the mass that a centroid misses and cuts the median mass error about threefold (0.0512 to 0.0162 and 0.0641 to 0.0228), and the value correction fixes the numerator; alone, each changes the latent error little (0.0541 to 0.0560 and 0.0546 on Cosmos-Predict2.5, 0.0364 to 0.0362 and 0.0325 on Wan2.1-14B), while together they lower it to 0.0412 and 0.0260. Under re-solved plans each correction alone already recovers 0.76–1.16 dB, because the solver moves exact blocks to the units that the missing term hurts most, so the fixed-plan rows measure the terms themselves, and both terms of Eq. (4) are needed.

Where each correction acts. Figure 7 bins clusters by key variance and by key–value covariance and plots the residual

Model	Steps	f_{dense}	X_{sp}	ρ_{global}	Computed blocks (M)
Cosmos-Predict1	35 + CFG	0.0633	0.200	0.251	4,057
Cosmos-Predict2.5	4	0.2768	0.200	0.421	439
Wan2.1-14B	4	0.2687	0.200	0.415	232
Wan2.2-A14B	4	0.2687	0.200	0.415	232
HunyuanVideo	50	0.0363	0.200	0.229	18,943
Cosmos3-Nano (I2W)	35 + CFG	0.0556	0.200	0.244	3,128
Cosmos-Transfer2.5	4	0.2734	0.200	0.419	498
DiffusionRenderer	4	0.2768	0.200	0.421	97
AlayaRenderer	50 + CFG	0.0527	0.200	0.242	761
UniRelight	35 + CFG	0.0633	0.200	0.251	1,863

Table 14. **Density accounting of the deployed plans.** Shares of attention work are weighted by sampler steps; computed blocks count 128×96 kernel blocks per video, in millions.

			Cosmos-Predict2.5					Wan2.1-14B				
Variant	Grouping	Corrections	Attn. MSE	Latent MSE	PSNR	LPIPS	DiT ms	Attn. MSE	Latent MSE	PSNR	LPIPS	DiT ms
A	–	none	0.415	0.214	17.57$_{\pm 2.33}$	0.385	7,176	0.577	0.186	17.08$_{\pm 1.10}$	0.300	11,291
B	Position	Centroid	0.0386	0.174	18.36$_{\pm 2.72}$	0.297	7,390	0.109	0.0886	19.69$_{\pm 1.46}$	0.155	11,444
C	Semantic	Centroid	0.0108	0.0541	23.47$_{\pm 3.68}$	0.124	8,127	0.0147	0.0364	23.41$_{\pm 2.05}$	0.078	12,756
D	Position	Both	0.0231	0.069	22.33$_{\pm 3.26}$	0.152	7,471	0.0732	0.0685	20.66$_{\pm 1.62}$	0.123	11,552
E (ours)	Semantic	Both	0.00623	0.0412	24.70$_{\pm 3.78}$	0.102	8,212	0.00880	0.0260	24.70$_{\pm 2.29}$	0.057	12,902

Table 15. **Compensation grouping under a fixed exact mask.** Only the summary of the skipped keys changes, A computes the exact blocks alone, and Both adds the variance and value corrections of Eq. (4). Small $\pm$: standard deviation of PSNR over the test prompts.

		Cosmos-Predict2.5				Wan2.1-14B			
Selection	Ranking	PSNR	LPIPS	Routing ms	DiT ms	PSNR	LPIPS	Routing ms	DiT ms
Offline	Attention mass	24.38$_{\pm 3.66}$	0.107	0	8,220	24.37$_{\pm 2.24}$	0.061	0	12,905
Offline	Residual (ours)	24.70$_{\pm 3.78}$	0.102	0	8,212	24.70$_{\pm 2.29}$	0.057	0	12,902
Online	Attention mass	24.65$_{\pm 3.46}$	0.098	650	8,862	23.95$_{\pm 2.08}$	0.066	1,288	14,192
Online	Residual*	27.97$_{\pm 3.96}$	0.058	1,162,863	1,171,074	27.03$_{\pm 2.48}$	0.036	1,092,072	1,104,968
Offline	Random	17.26$_{\pm 2.61}$	0.415	0	8,219	16.36$_{\pm 1.42}$	0.289	0	12,868
CalibAtt masks [43]		21.22$_{\pm 3.39}$	0.183	0	8,206	20.79$_{\pm 1.73}$	0.116	0	12,875

Table 16. **Block selection under a fixed semantic compensator.** The starred row reads the dense attention of the current call (diagnostic), and Random keeps every diagonal block. Small $\pm$: standard deviation of PSNR over the test prompts of each model.

energy each correction removes from the clusters of each bin. The variance correction removes 10–25% of the centroid error in every bin. The value correction applies the head average H to every cluster, so on the clusters whose own B_c is smallest the gap $B_c - H$ of Proposition A.6 exceeds B_c and the correction adds error, and elsewhere it changes the per-cluster error by less than 1%. Over the whole output it lowers the projected attention error from 0.0108 to 0.00916 under a fixed plan (Table 17), and a per-cluster H_c closes the gap (Table 18). Table 18 replaces the diagonal variance by the full covariance and the shared H by one H_c per cluster on saved attention inputs. Full covariance lowers the attention error by only 5% (0.0059 to 0.00559), so the diagonal variance already captures nearly all the mass a centroid misses, whereas one H_c per cluster lowers it by 27% (to 0.00428), in line with Figure 7b. The deployed compensator keeps the diagonal variance and the shared H of Eq. (4) at two GEMMs per call, and H_c extends the same correction.

D.4. Distortion Budget

Pricing and sensitivity. Table 19 crosses the error that prices a unit, measured before or after compensation, with the sensitivity, set to one or measured, at one compensator, block candidate set, assignment scheme, and density. Shuffling J across units while keeping its values tests whether the gain comes from locating sensitive units, and pricing the head output without the W_o projection of Eq. (10) tests the projection. The sensitivity also pays at the density 0.10 of Figure 5b–d, where setting it to one costs 1.54 dB on Cosmos-Predict2.5. Two time-only rows compare the curvature schedule of PASA [48], ported into DecoupAttn,

Plan	γ_{ic}	H	$\lvert\hat{Z}/Z-1\rvert$ (50%)	$\lvert\hat{Z}/Z-1\rvert$ (99%)	Attn. MSE	Latent MSE	PSNR	LPIPS	DiT ms
Cosmos-Predict2.5									
Fixed	off	off	0.0512	0.479	0.0108	0.0541	23.47$\pm$3.68	0.124	8,127
	on	off	0.0162	0.33	0.00832	0.056	23.44$\pm$3.71	0.129	8,206
	off	on	0.0512	0.479	0.00916	0.0546	23.31$\pm$3.52	0.128	8,171
	on	on	0.0162	0.33	0.00623	0.0412	24.70$\pm$3.78	0.102	8,212
Re-solved	off	off	–	–	–	0.0539	23.51$\pm$3.67	0.124	8,138
	on	off	–	–	–	0.0439	24.44$\pm$3.77	0.106	8,222
	off	on	–	–	–	0.0423	24.55$\pm$3.65	0.105	8,225
	on	on	0.0162	0.33	0.00623	0.0412	24.70$\pm$3.78	0.102	8,212
Wan2.1-14B									
Fixed	off	off	0.0641	0.433	0.0147	0.0364	23.41$\pm$2.05	0.078	12,756
	on	off	0.0228	0.311	0.0123	0.0362	23.47$\pm$2.09	0.079	12,836
	off	on	0.0641	0.433	0.0117	0.0325	23.69$\pm$2.13	0.069	12,808
	on	on	0.0228	0.311	0.00880	0.0260	24.70$\pm$2.29	0.057	12,902
Re-solved	off	off	–	–	–	0.0365	23.41$\pm$2.03	0.078	12,733
	on	off	–	–	–	0.0293	24.17$\pm$2.24	0.063	12,907
	off	on	–	–	–	0.0262	24.57$\pm$2.15	0.059	12,887
	on	on	0.0228	0.311	0.00880	0.0260	24.70$\pm$2.29	0.057	12,902

Table 17. **Variance and value corrections,** 2×2**.** Fixed holds the exact mask, options, assignments, and FULL calls, re-solved re-solves each variant at equal density, and mass errors are query medians and 99th percentiles. Small ±: standard deviation over test prompts.

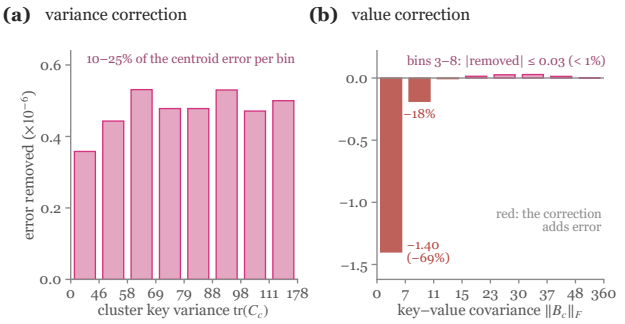

Figure 7. **Where the corrections act** (Cosmos-Predict2.5). Per-cluster residual energy removed by the variance correction, binned by cluster key variance (a), and by the value correction, binned by key–value covariance (b). Each bin holds one eighth of the clusters, and a negative bar means the correction adds error.

Variance	Value	Attn. MSE	Extra ms
Diagonal	Shared H (ours)	0.0059	0
Full covariance	Shared H	0.00559	–
Diagonal	Per-cluster H_c	0.00428	–

Table 18. **Richer corrections on saved attention inputs** (Cosmos-Predict2.5). Extra ms: added time per call (–: offline implementation).

with Distortion Budget restricted to one weight per step, and the density of every row is re-measured after padding. Both factors of the price $J_u\varepsilon_u$ matter, and they pay most together (Table 19): on Cosmos-Predict2.5, pricing the compensated error gains 0.41 dB over the uncompensated error with $J = 1$, the measured J gains 1.94 dB, and both together gain 2.99 dB (24.70 against 21.71 dB), with the same ordering on Cosmos-Transfer2.5. Shuffling J keeps its values but not their places and falls to 20.05 dB, below $J = 1$, so the gain comes from locating the units whose errors reach the final latent. Pricing the head output without W_o changes PSNR by at most 0.14 dB, and the two time-only schedules reach 21.79 and 22.01 dB, near $J = 1$, so most of the gain comes from pricing layers and heads, not steps alone.

Predicted versus actual distortion. For each model, 21 candidate plans at one density (the deployed plan, uniform allocation, compensation error alone, a time-only schedule, small budget swaps of the deployed plan, and the ablation plans) are priced on the calibration set by $\hat{D}(\pi) = \sum_u w_u J_u \varepsilon_u(o_u, \eta_{tl})$, and their actual distortion $D_{\text{actual}}(\pi) = \mathbb{E}_x\|\hat{x}_0(\pi) - x_0\|_2^2$ is measured on held-out clips (Figure 8). Table 20 reports the rank correlations within the budget, the actual rank of the plan the objective selects, and its regret against the best candidate. On both models the price ranks the 21 plans nearly as the held-out error does (Spearman 0.944 and 0.917), and the plan the objective selects is the second best, 0.6% and 1.2% above the best candidate, as Proposition 2 expects: the solver minimizes the price, and the price follows the actual distortion.

Additivity. Eq. (12) sums the priced distortion over units. For two groups of units A and B, sparse error is injected into A alone, B alone, and both, giving final-latent er-

Error priced	Sensitivity	Cosmos-Predict2.5				Cosmos-Transfer2.5			
		Latent MSE	PSNR	LPIPS	Density	Latent MSE	PSNR	LPIPS	Density
Before compensation	$J = 1$	0.0813	21.71	0.163	0.200	0.0180	31.62	0.074	0.200
After compensation	$J = 1$ (U)	0.0743	22.12	0.151	0.200	0.0146	32.70	0.060	0.200
Before compensation	Measured J (D)	0.052	23.65	0.122	0.200	0.0124	33.16	0.059	0.200
After compensation	Measured J (ours)	0.0412	24.70	0.102	0.200	0.0111	33.90	0.050	0.200
After compensation	Shuffled J	0.112	20.05	0.206	0.200	0.0180	31.68	0.072	0.200
Head output (no W_o)	Measured J	0.041	24.77	0.100	0.200	0.0114	33.76	0.051	0.200
Time only, PASA schedule [48]		0.0784	21.79	0.160	0.200	0.0197	30.95	0.082	0.200
Time only, Distortion Budget		0.0744	22.01	0.154	0.200	0.0179	31.35	0.079	0.200

Table 19. **Pricing error and sensitivity.** U and D are the ablations of Table 2 under this protocol, and density is measured after padding. Cosmos-Transfer2.5 comes from one rerun with saved latents, in which the full method reaches 33.90 dB (33.81 dB in Table 1).

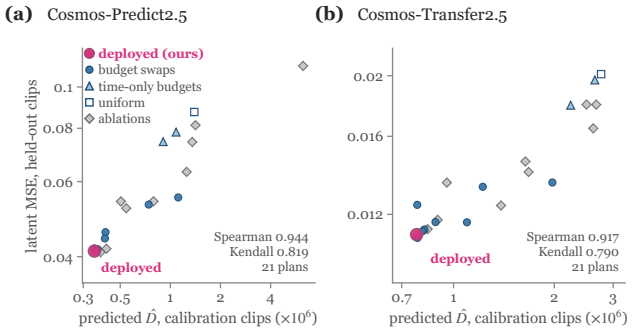

Figure 8. **Predicted versus actual distortion** of candidate plans at one density; predicted on calibration clips, measured on held-out clips (Cosmos-Predict2.5, Cosmos-Transfer2.5).

Model	Plans	Spearman	Kendall	Rank	Regret (%)
CP2.5	21	0.944	0.819	2	0.6
T2.5	21	0.917	0.790	2	1.2

Table 20. **How well the priced distortion ranks plans.** CP2.5: Cosmos-Predict2.5; T2.5: Cosmos-Transfer2.5. Rank: actual rank of the selected plan among the candidates; regret: its excess actual distortion over the best candidate.

rors e_A, e_B, and e_{AB}; Table 21 reports the nonlinearity $\|e_{AB} - e_A - e_B\|_2/(\|e_A\|_2 + \|e_B\|_2)$ and the cross term $2e_A^\top e_B/(\|e_A\|_2^2 + \|e_B\|_2^2)$ that summed squared errors omit. Errors of different unit groups interact moderately: the cross term is 0.12–0.45 of the summed energies and smallest between the first and the last sparse step (Table 21). The summed price ranks the candidate plans as their held-out error does (Table 20), all the plan choice of Eq. (12) needs.

Reference backend. The sensitivity of Eq. (11) subtracts the run-to-run floor of the same dense attention backend; Table 22 recomputes J and the method ranking against a second dense attention backend as reference and counts the cells where the subtraction turns negative. J keeps its order under the second reference (Spearman 0.884 and 0.967), the method ranking stays the same (Kendall 1.000), and at most

Model	Groups A, B	Nonlinearity	Cross term
CP2.5	Adjacent steps	$0.436_{\pm 0.050}$	$0.252_{\pm 0.055}$
	Early vs. late layers	$0.400_{\pm 0.057}$	$0.165_{\pm 0.062}$
	First vs. last sparse step	$0.326_{\pm 0.044}$	$0.122_{\pm 0.038}$
T2.5	Adjacent steps	$0.507_{\pm 0.087}$	$0.394_{\pm 0.124}$
	Early vs. late layers	$0.607_{\pm 0.104}$	$0.448_{\pm 0.215}$
	First vs. last sparse step	$0.496_{\pm 0.103}$	$0.294_{\pm 0.135}$

Table 21. **Interaction between unit groups** in the final latent (CP2.5: Cosmos-Predict2.5; T2.5: Cosmos-Transfer2.5). Mean over the test clips (32 and 39), small $\pm$: standard deviation.

6 of 81 cells turn negative (Table 22), so Distortion Budget measures the sensitivity of the model itself.

Model	Spearman of J	Kendall of ranking	Negative cells
CP2.5	0.884	1.000	6/81
T2.5	0.967	1.000	0/9

Table 22. **Stability under the dense reference** (CP2.5: Cosmos-Predict2.5; T2.5: Cosmos-Transfer2.5). Correlations between the two references; negative cells: J cells whose floor-subtracted error is below zero under the second reference.

D.5. Residual Structure

Exchangeability of the cluster members. Assumption (A4) of Appendix A.1 treats the block a cluster member falls into as independent of its deviation. Figure 9 tests it by the mean pairwise cosine of two skipped blocks' residuals, and of their deviation terms D_b of Proposition A.8, against their distance. Theorem A.10 bounds the coherence of the deviation terms by one under (A4), and the measured deviation terms add like random errors: their median coherence ratio is 0.97–1.72, against 4.86–8.25 for the residuals without compensation (Table 23), and their correlation decays with the distance between the blocks (Figure 9). The cluster bias $\beta_{ib} = r_{ib} - D_b$ is the coherent part of the residual, and the corrections shrink it from 0.73–1.34 to 0.54–0.83 of the residual norm, with the variance correction giving nearly all of this reduction (Table 23).

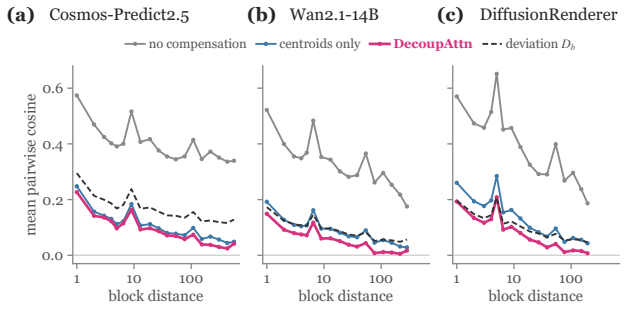

Figure 9. **Residual correlation against block distance.** Mean pairwise cosine of the residuals of two skipped blocks, binned by the distance between the blocks, without compensation, with plain centroids, and with full DecoupAttn, and of the deviation terms D_b of Proposition A.8 (dashed), which are the same under every compensator.

Quantity	Compensation	CP2.5	Wan2.1	DR
Coherence of D_b	–	1.72	0.97	1.04
Coherence of r_{ib}	None	8.25	4.86	5.67
	Centroid	4.48	2.76	3.36
	Variance only	2.99	2.12	2.96
	Value only	4.20	2.55	2.62
	Both (ours)	2.30	1.64	1.67
$\|\beta_{ib}\|_2/\|r_{ib}\|_2$	Centroid	1.34	0.73	0.97
	Variance only	0.82	0.53	0.71
	Value only	1.34	0.75	0.97
	Both (ours)	0.83	0.54	0.69

Table 23. **Coherence ratio and cluster bias** on saved attention inputs: medians over queries of the coherence ratio of D_b and of the residual r_{ib}, and over skipped blocks of $\|\beta_{ib}\|_2/\|r_{ib}\|_2$.

Energy, coherence, and mass. Eqs. (6) and (7) give $\|\hat{y}_i - y_i\|_2^2 = \kappa_i R_i/\hat{Z}_i^2$ with $R_i = \sum_{b\notin\mathcal{S}_g} \|r_{ib}\|_2^2$, so the output error splits into residual energy, coherence, and estimated mass. Table 24 reports the three factors and the error for five compensators on the same held-out inputs and masks, next to the gap between fixed blocks and online routing of Table 16. By Eq. (7), κ_i grows with the pairwise cosine of the skipped-block residuals, and Figure 10 shows that compensation moves the median cosine from 0.31–0.36 to 0.012–0.027. Unlike Table 11, κ_i is taken per query through the deployed bf16 kernel with the error normalized by the estimated mass $\hat{Z}_i$, the residual energy is normalized as $\tilde{R}_i = R_i/(Z_i^2\|y_i\|_2^2)$, and the error is $\|\hat{y}_i - y_i\|_2^2/\|y_i\|_2^2$. Compensation lowers all three factors, and the corrections act mainly on coherence and mass (Table 24): from plain centroids to both corrections, the residual energy falls by at most 10%, while κ_i falls from 2.73–4.43 to 1.56–2.29 and the mass error from 0.051–0.100 to 0.016–0.035, so the error falls 2.2–2.6-fold. The variance correction fixes the mass and also lowers κ_i, the value correction lowers κ_i without changing the mass, and both together reach the lowest error on all three models, as Eq. (7) predicts. Figure 10 shows these shifts across the whole distributions: with both corrections, κ_i and $\hat{Z}_i/Z_i$ concentrate closer to one, the error falls at every quantile from the 25th to the 99th, and the distribution of the residual energy stays in place.

Compensation	$\tilde{R}_i$	κ_i	$\lvert\hat{Z}_i/Z_i - 1\rvert$	Error	Gap (dB)
DiffusionRenderer					
None	0.00191	4.71	0.270	0.0222	–
Centroid	0.00113	3.00	0.100	0.00478	–
Variance only	0.00104	2.59	0.0351	0.0034	–
Value only	0.00112	2.32	0.100	0.00354	–
Both (ours)	0.00102	1.56	0.0351	0.00186	1.87
Wan2.1-14B					
None	0.00319	4.50	0.301	0.0570	–
Centroid	0.000973	2.73	0.0641	0.00295	–
Variance only	0.000924	2.03	0.0228	0.00166	–
Value only	0.000994	2.46	0.0641	0.00271	–
Both (ours)	0.000941	1.57	0.0228	0.00134	2.33
Cosmos-Predict2.5					
None	0.000378	8.1	0.107	0.00503	–
Centroid	0.00028	4.43	0.0512	0.00158	–
Variance only	0.00027	2.92	0.0162	0.000873	–
Value only	0.00028	4.27	0.0512	0.00154	–
Both (ours)	0.000267	2.29	0.0162	0.000674	3.27

Table 24. **Factors of the output error** (medians over queries on held-out clips, deployed bf16 kernel). Columns are defined in the text above, and the gap is the PSNR of online routing minus fixed blocks under the deployed compensator.

D.6. Reassign Cache

Assignment update schemes. Every scheme in Table 25 recomputes the cluster statistics from the current keys and values (Eq. (4)) and runs one exact plan; the schemes differ only in the cluster assignments. Frozen keeps the assignments of the first sparse step, Reassign assigns keys to the cached centroids (Eq. (13)), Rebuild runs k-means at every sparse step, and Periodic rebuild runs it every 5 steps on the many-step models. Figure 11 follows Frozen, Reassign, and Rebuild across the sparse steps: the distance to the cached centroids (the quantity of Proposition 3), the within-cluster error to the current means, the projected attention error, and the share of keys whose assignment changes. Reassignment keeps the distance to the cached centroids at or below its frozen value, as Proposition 3 states, and keeps each cluster tight around its current mean, so its attention error stays close to Rebuild's and far below Frozen's. Reassign recovers most of what Rebuild gains over Frozen at a fraction of its cost (Table 25): on UniRelight it lifts PSNR from 39.13 to 42.75 dB against 43.18 dB for Rebuild, changes 22.5% of the assignments per step against 37.0%, and adds 4.8 s of DiT time instead of 36.7 s; on Cosmos-Predict2.5 it recovers 0.47 of Rebuild's 0.52 dB for 0.15 s instead of 1.09 s. Periodic rebuild reaches Rebuild's accuracy (43.16 dB) for

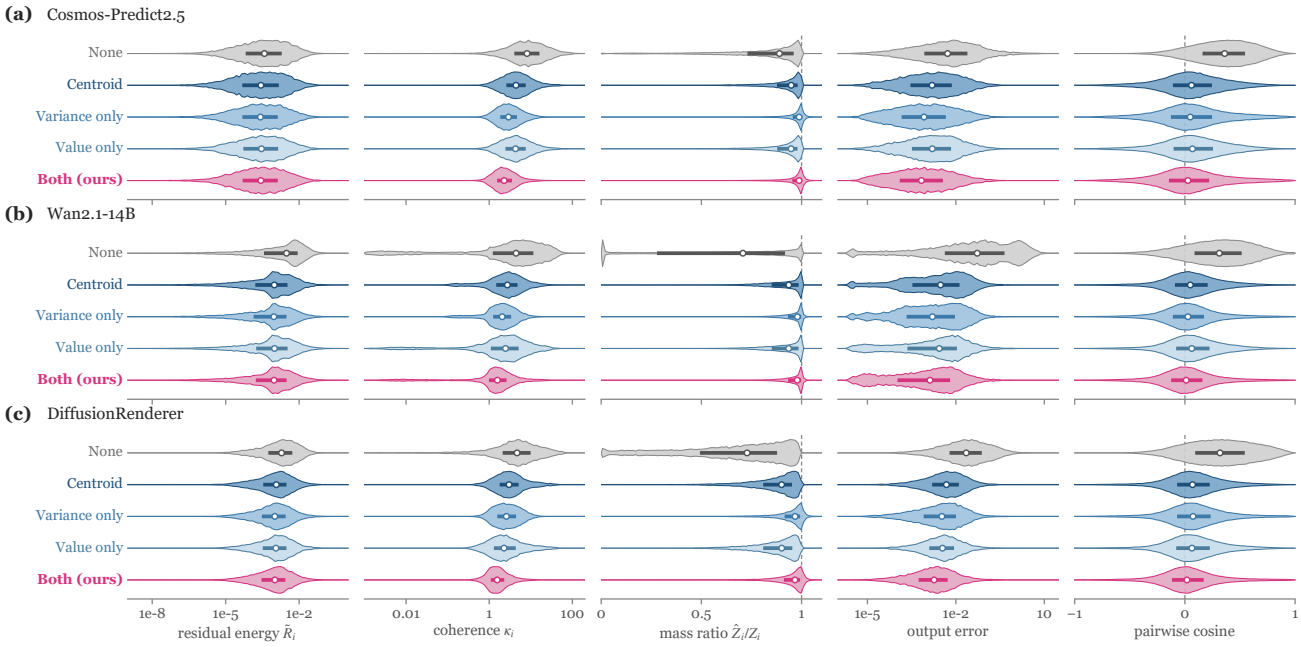

Figure 10. **Distributions of the error factors** of Table 24 for its five compensators: residual energy $\tilde{R}_i$, coherence κ_i, mass ratio $\hat{Z}_i/Z_i$, output error, and the cosine between the residuals of two skipped blocks of the same query tile (held-out clips, 10,000 queries and 100,000 random pairs per compensator). Bars: interquartile range, dots: medians.

10.0 s, still twice the cost of Reassign, and on DiffusionRenderer, where stale assignments cost only 0.46 dB on calibration clips (Table 26), the three schemes lie within 0.07 dB of each other.

Deployment rule. Table 26 checks the choices of Eq. (14) on test clips: the staleness measured on calibration clips, the scheme the rule selects, the test PSNR gain and latency cost of that scheme over the other, and whether the staleness measured on test clips falls on the same side of τ_s. On test clips the selected scheme gains 0.14–3.63 dB on all nine models that reassign, and DiffusionRenderer, which keeps its assignments, gives up 0.06 dB for 102 ms. The calibration staleness overstates the test gain, so the two fall on the same side of τ_s on five of the ten models.

D.7. Efficiency

Quality–latency curves. Each method runs at four or five operating points of its own density knob, with the density of every point measured. Figure 5b–d plots PSNR against DiT speedup, and Figure 12 plots LPIPS against end-to-end latency for the same points. In LPIPS the picture is the same: at every latency of DecoupAttn all baselines lie above its curve, and the lower LPIPS of SVG-EAR and SVG2 at top-p 0.99 costs 1.5–1.8 times the dense end-to-end time. Table 29 reads two numbers off them: the PSNR gain at the latency of the strongest baseline, and the speedup at a target PSNR fixed before testing. At the latency of the strongest baseline, DecoupAttn is 5.74–8.40 dB more accurate, and the strongest baseline does not reach our deployed PSNR at any of its operating points, so its speedup at equal PSNR is undefined (Table 29).

Time breakdown. Table 27 splits the time of one video into attention (exact blocks, first k-means, assignment, cluster statistics, block–cluster counts, and the fused merge), the DiT, and the full generation (encoders, DiT, VAE decoding, and transfers). Each entry is the median of five interleaved runs on one NVIDIA H800 GPU with batch size 1 and bf16 precision, without compilation or cold start; Table 30 accounts offline calibration, plan solving, and plan loading separately and gives the break-even count

$$N_{\text{break-even}} = \left\lceil \frac{T_{\text{calibration}}}{T_{\text{dense, E2E}} - T_{\text{ours, E2E}}} \right\rceil,$$

the smallest number of videos whose savings repay the calibration, for $T_{\text{ours, E2E}} < T_{\text{dense, E2E}}$.

Where the time goes. The gap to SVG-EAR comes mostly from the per-step work (Table 27): SVG-EAR re-clusters at every step and rebuilds its error map and top-p selection, which costs 4.7–18.4 s of k-means and 0.8–17.6 s of selection per video, while DecoupAttn clusters once (0.6–1.2 s) and later only reassigns, counts, and merges. Its attention therefore takes 1.1–1.7 times less time than dense attention and 1.9–2.8 times less than SVG-EAR's, and every video finishes faster end to end. Calibration takes 5.3–19.6 GPU-

Scheme	Cosmos-Predict2.5			DiffusionRenderer			UniRelight		
	PSNR	Changed (%)	DiT ms	PSNR	Changed (%)	DiT ms	PSNR	Changed (%)	DiT ms
Frozen	24.23	0.0	8,057	50.87	0.0	3,913	39.13	0.0	103,836
Reassign	24.70	36.5	8,208	50.91	22.1	4,020	42.75	22.5	108,671
Rebuild	24.75	52.9	9,148	50.94	45.1	4,736	43.18	37.0	140,497
Periodic rebuild	–	–	–	–	–	–	43.16	22.1	113,839

Table 25. **Assignment update schemes under one exact plan.** All schemes recompute cluster statistics from current keys and values, and Changed is the share of keys changing assignment per sparse step. DiT ms: median of interleaved runs (periodic rebuild: mean of two).

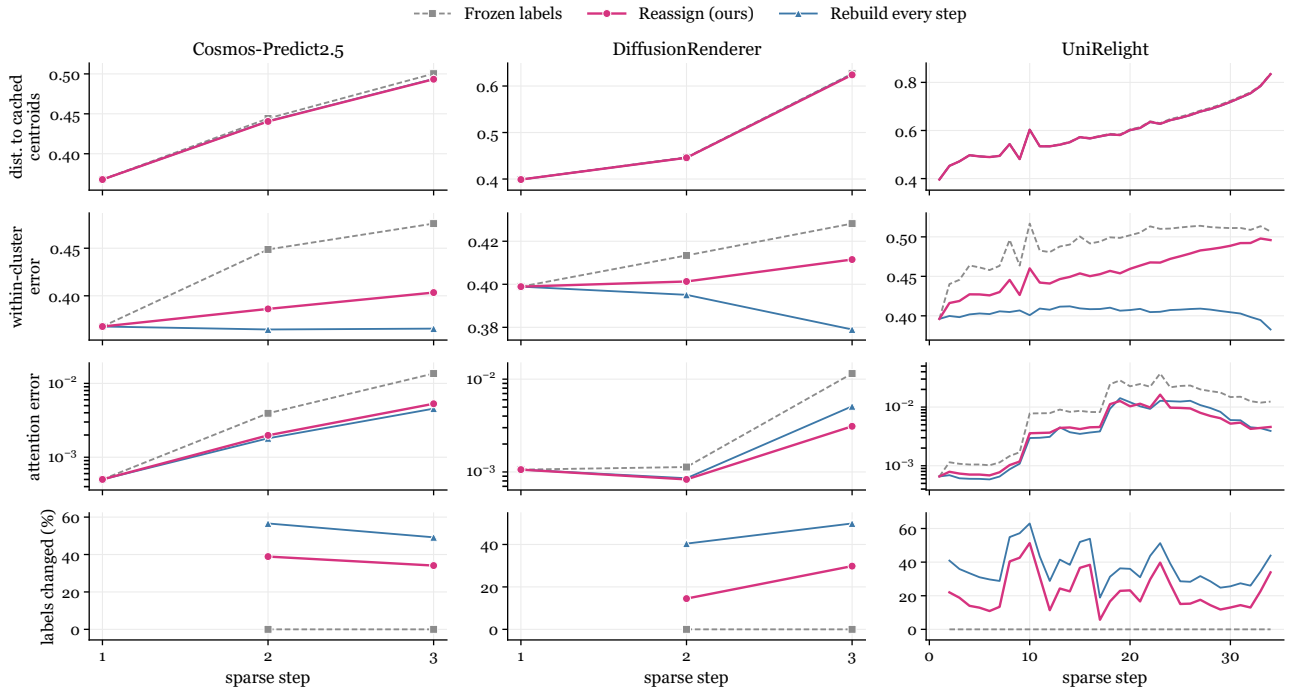

Figure 11. **Assignment schemes across sparse steps** (Cosmos-Predict2.5, DiffusionRenderer, UniRelight). Reassignment stays at or below frozen assignments in distance to the cached centroids and cuts the within-cluster and attention errors that frozen assignments let grow. Step 1 builds the clusters.

Model	δ_{stale} (dB)	Selected	Test gain (dB)	Time cost (ms)	Agrees
Cosmos-Predict1	2.35	reassign	1.28	7,484	yes
Cosmos-Predict2.5	3.00	reassign	0.47	150	no
Wan2.1-14B	2.73	reassign	1.09	218	yes
Wan2.2-A14B	4.35	reassign	0.41	255	no
HunyuanVideo	2.86	reassign	0.76	26,463	no
Cosmos3-Nano (I2W)	2.86	reassign	0.31	8,050	no
Cosmos-Transfer2.5	2.13	reassign	0.14	167	no
DiffusionRenderer	0.46	frozen	−0.06	−102	yes
AlayaRenderer	2.43	reassign	2.40	2,850	yes
UniRelight	1.87	reassign	3.63	4,780	yes

Table 26. **Assignment rule on test clips.** Gain and cost of the selected scheme over the other; agrees: whether the test-clip staleness falls on the same side of τ_s as the calibration staleness. For HunyuanVideo and Cosmos3-Nano this table re-runs the calibration from scratch, with a reassign PSNR of 27.25 and 21.95 dB.

Method	Attention							DiT	E2E	Peak GiB
	Exact	k-means	Assign	Statistics	Counts	Merge	Total			
Cosmos-Predict2.5										
Dense	9,352	–	–	–	–	–	9,352	12,146	19,682	25.4
SVG-EAR	2,294	4,730	133	23	762	–	10,687	13,347	20,964	27.3
Ours	1,291	620	171	96	46	260	5,484	8,201	15,706	38.2
Wan2.1-14B										
Dense	4,842	–	–	–	–	–	4,842	13,649	16,399	35.9
SVG-EAR	1,822	7,142	188	28	1,110	–	11,905	20,520	23,272	37.6
Ours	843	923	240	150	62	460	4,266	12,908	15,658	59.4
UniRelight										
Dense	68,424	–	–	–	–	–	68,424	141,099	144,398	34.1
SVG-EAR	29,514	18,435	2,845	444	17,592	–	77,650	144,762	147,921	36.0
Ours	14,157	1,218	5,403	2,534	864	6,466	39,694	108,432	111,718	41.1

Table 27. **Time breakdown of one video** (ms, median over interleaved runs). k-means is the first clustering, Counts the block–cluster counts, and E2E covers encoders, DiT, VAE decoding, and transfers. SVG-EAR has no merge, and its Counts column holds its error map and top-p block selection. Peak GiB is the maximum allocated GPU memory over the whole generation, measured with one method per process and the text embeddings precomputed.

Model	Method	Subject	Background	Flickering	Smoothness	Dynamic	Imaging	Aesthetic	Overall
Cosmos-Predict1	Dense	0.9451	0.9541	0.9872	0.9945	0.2821	0.6796	0.4644	0.2224
	SVG-EAR	0.9387	0.9514	0.9860	0.9944	0.2308	0.6835	0.4620	0.2269
	Ours	0.9325	0.9494	0.9870	0.9945	0.2051	0.6775	0.4469	0.2230
Cosmos-Predict2.5	Dense	0.9430	0.9574	0.9722	0.9878	0.5312	0.7087	0.6047	0.2864
	SVG-EAR	0.9445	0.9563	0.9715	0.9875	0.5312	0.7124	0.6005	0.2885
	Ours	0.9451	0.9588	0.9732	0.9884	0.5000	0.7168	0.6054	0.2864
Wan2.1-14B	Dense	0.9334	0.9537	0.9633	0.9853	0.5625	0.7052	0.6601	0.2586
	SVG-EAR	0.9335	0.9554	0.9636	0.9852	0.5312	0.7067	0.6604	0.2593
	Ours	0.9343	0.9536	0.9634	0.9854	0.5312	0.7082	0.6588	0.2583
Wan2.2-A14B	Dense	0.9162	0.9457	0.9552	0.9788	0.7500	0.7092	0.6787	0.2671
	SVG-EAR	0.9181	0.9482	0.9552	0.9787	0.7500	0.7044	0.6759	0.2671
	Ours	0.9172	0.9492	0.9563	0.9795	0.7500	0.7093	0.6718	0.2679
HunyuanVideo	Dense	0.9165	0.9636	0.9875	0.9946	0.6250	0.6753	0.5578	0.2569
	SVG-EAR	0.9241	0.9664	0.9881	0.9944	0.6250	0.6654	0.5610	0.2587
	Ours	0.9182	0.9583	0.9881	0.9946	0.6250	0.6647	0.5567	0.2585
Cosmos3-Nano (I2W)	Dense	0.9322	0.9466	0.9800	0.9936	0.5625	0.7182	0.5968	0.2406
	SVG-EAR	0.8440	0.8975	0.9663	0.9852	0.5312	0.6606	0.5064	0.2374
	Ours	0.9318	0.9466	0.9804	0.9938	0.5625	0.7178	0.5992	0.2428

Table 28. **VBench dimensions.** Subject and background consistency, temporal flickering, motion smoothness, dynamic degree, imaging and aesthetic quality, and overall consistency (text–video alignment).

Model	Baseline	ΔPSNR at equal latency	Speedup at equal PSNR	Target PSNR
CP2.5	Sol-Attn	+6.40	–	24.70
Wan2.1	PISA	+5.74	–	24.70
UniRelight	PISA	+8.40	–	42.75

Table 29. **Matched-latency and matched-quality comparison** against the strongest baseline of Figure 12 (CP2.5: Cosmos-Predict2.5; Wan2.1: Wan2.1-14B). –: the baseline does not reach the target PSNR at any operating point.

hours and the plan solve under a minute (Table 30), and the break-even count follows the end-to-end saving per video: UniRelight, which saves the most, breaks even after 2.1k videos, and Wan2.1-14B, which saves the least, after 23.7k.

D.8. Generalization of the Offline Plan

Calibration set size. Table 31 calibrates on 1, 2, 4, 8, and 16 training clips with three draws each and reports the test PSNR, the intersection over union of the exact blocks with those of the 16-clip plan, and the calibration time. PSNR rises from 24.59 to 24.87 dB as the set grows from 1 to 16 clips, and its spread over draws shrinks from 0.22 dB with one clip to 0.04 dB with eight (Table 31); with the deployed four clips, the exact blocks already overlap 94.8% with the 16-clip plan after 0.33 GPU-hours of error calibration.

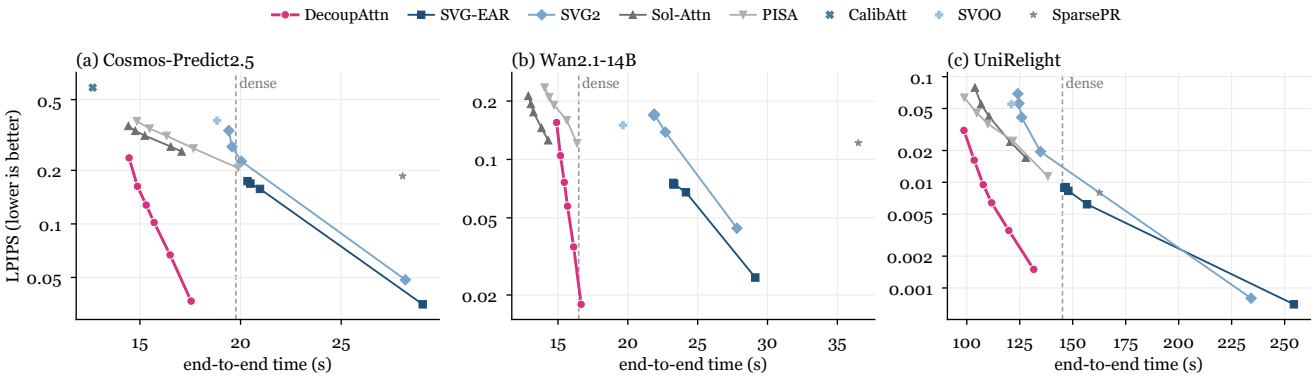

Figure 12. **LPIPS against end-to-end latency** for the operating points of Figure 5b–d; dashed: dense attention. End-to-end time covers text encoding, DiT, and VAE decoding (UniRelight: the whole generation).

Model	Calib. (GPU-h)	Peak mem. (GiB)	Plan solve (min)	Saving (s / video)	$N_{\text{break-even}}$
CP2.5	5.93	38.7	0.60	4.06	5,263
Wan2.1	5.28	39.3	0.47	0.80	23,664
UniRelight	19.62	54.4	0.80	33.66	2,099

Table 30. **Offline cost and break-even count** (CP2.5: Cosmos-Predict2.5; Wan2.1: Wan2.1-14B). Calib.: block ranking, distortion table, and error calibration of the deployed plan, without the sensitivity measurement. Peak memory: GPU memory during calibration. Saving: end-to-end time saved per video.

Clips	PSNR	Block overlap	Calibration (GPU-h)
1	24.59 ± 0.22	0.875	0.08
2	24.67 ± 0.18	0.922	0.17
4	24.80 ± 0.09	0.948	0.33
8	24.87 ± 0.04	0.968	0.67
16	24.87	1.000	1.34

Table 31. **Calibration set size** (Cosmos-Predict2.5; mean $\pm$ standard deviation over three draws). Only the error calibration uses the drawn clips, and the sensitivity and the distortion table are those of the deployed plan. The 4-clip draws are training clips 0–3, 4–7, and 8–11, and the first is the deployed plan. Block overlap: intersection over union of the exact blocks with the 16-clip plan. Calibration counts error calibration alone (301 s per clip), while Table 30 adds block ranking and the distortion table.

Unseen inputs and settings. With the plan frozen, Table 32 groups held-out prompts by content and reports the mean and the worst decile. The plan is tied to the token grid of one resolution, so a new resolution runs a new calibration, and Table 33 reports its cost for Wan2.1-14B at 720p. The frozen plan leads SVG-EAR on every content group by 1.22–3.79 dB, including categories absent from calibration (+3.39 dB) and complex lighting (+3.79 dB, Table 32). Its worst decile stays at 19.4–27.6 dB in every group (Table 32). The plan fixes only where exact attention runs, while each video clusters its own keys at its first sparse step, so new content changes the compensation, not the plan. At 720p the recalibrated plan reaches 23.89 dB with the same procedure and constants after 18.70 GPU-hours of calibration, 3.5 times that of the 480p plan (Table 33).

Content	Clips	Mean	Worst 10%	vs. SVG-EAR
Unseen categories	6	24.77	22.68	+3.39
Fast motion	7	23.35	19.39	+3.25
Several subjects	7	24.95	19.39	+3.26
Camera motion	7	22.22	19.39	+1.22
Fine texture	6	24.64	20.64	+1.52
Complex lighting	6	29.72	27.62	+3.79

Table 32. **Frozen plan on held-out content** (Cosmos-Predict2.5; PSNR, dB). vs. SVG-EAR: mean paired difference.

Resolution	PSNR (dB)	Calibration (GPU-h)
480p (deployed)	24.70	5.28
720p (recalibrated)	23.89	18.70

Table 33. **Recalibration cost for a new resolution** (Wan2.1-14B). 480p: the deployed plan of Table 1, with its cost from Table 30.

Cluster count and k-means rounds. Table 34 varies the clusters per head over 220, 440, 880, and 1,760 and the k-means rounds of the first sparse step. More clusters per head raise PSNR steadily, from 23.88 to 25.18 dB, because smaller clusters have members closer to their centroid and leave smaller residuals (Proposition A.6), while k-means time grows from 404 to 927 ms and DiT time from 7.78 to 8.78 s; 880 clusters keep all but 0.48 dB of the 1,760-cluster gain at 0.59 s less DiT time. The k-means rounds matter little, 3 rounds losing 0.06 dB and 20 rounds gaining 0.01 dB, since every later call recomputes the cluster statistics.

C	Rounds	PSNR	LPIPS	k-means ms	DiT ms
220	10	23.88	0.116	404	7,778
440	10	24.24	0.109	468	7,901
880 (ours)	10	24.70	0.102	625	8,187
1,760	10	25.18	0.095	927	8,778
880	3	24.64	0.103	278	7,879
880	20	24.71	0.101	1,090	8,633

Table 34. **Clusters per head and k-means rounds** (Cosmos-Predict2.5). k-means ms: the first clustering per video.

D.9. Task Quality

VBench dimensions. Table 28 reports each VBench dimension [7] together with dynamic degree, the amount of motion, and overall consistency, the text–video alignment score. DecoupAttn tracks dense attention in every dimension on the four-step models (Table 28), and on Cosmos3-Nano it keeps every dimension within 0.003 of dense attention, while SVG-EAR drops subject consistency from 0.932 to 0.844 and imaging quality from 0.718 to 0.661. Dynamic degree stays within one video of the dense value on every model but Cosmos-Predict1, where it drops by three of 39 videos (SVG-EAR: two), so the sparse videos keep nearly all the motion of the dense outputs.

Conditions and ground truth. Table 35 measures edge adherence on Cosmos-Transfer2.5, the F1 score between the input edges and the edges of the output, and the error against ground truth where the rendering test sets provide it. DecoupAttn keeps edge adherence at the dense level (0.423 against 0.421) and matches dense attention against ground truth on both renderers (27.48 against 27.49 dB and 12.08 against 12.17 dB), above SVG-EAR on all three (Table 35), so its fidelity to the dense output carries over to the input conditions and to the ground truth.

	Edge F1	PSNR to ground truth	
Method	T2.5	DR	Alaya
Dense	0.421	27.49	12.17
SVG-EAR	0.410	27.43	11.94
Ours	0.423	27.48	12.08

Table 35. **Condition adherence and ground-truth error.** T2.5: Cosmos-Transfer2.5; DR: DiffusionRenderer; Alaya: AlayaRenderer.

E. Case Studies

Four clips show where the method is stressed: two are preselected by a measurable property before generation (fast motion with fine texture, and occlusion), and two are the clips on which DecoupAttn scores lowest. Each figure shows the same frames from dense attention, SVG-EAR, and DecoupAttn, generated with the main-table configuration at density 0.20 and seed 1000.

Case 1: Fast motion (UniRelight, syn_test_0004). It has the largest mean absolute frame difference of the dense output among the 39 test clips, and the largest mean Laplacian variance, so it combines fast motion with fine texture. Figure 13 shows frames 0, 28, and 56 of this clip.

Key Observations: Both sparse outputs are indistinguishable from the dense frames: DecoupAttn scores 38.4–38.9 dB per frame and SVG-EAR 39.9–40.2 dB (38.76 and 39.94 dB over the clip). The remaining error of both methods lies on the edges of the cobblestones and the furniture, the fine texture that the motion carries across blocks. SSIM stays at 0.990–0.992 on every frame for DecoupAttn and 0.993–0.995 for SVG-EAR, so both keep the structure of the dense frames. The error of DecoupAttn has no block shape although its exact blocks are fixed: the compensated residuals of the skipped blocks cancel like dithered quantization errors (Proposition 1, Figure 4a–c), so what remains spreads as fine noise over the texture instead of gathering at block borders. UniRelight is also the model on which keeping the clusters current matters most (Table 25), and on its fastest clip DecoupAttn, which reassigns on this model, keeps every frame above 38 dB.

Case 2: Occlusion (UniRelight, syn_test_0006). It has the largest share of pixels whose G-buffer depth changes by more than 10% of its range between frames, so objects keep covering and uncovering each other. Figure 14 shows frames 0, 28, and 56 of this clip.

Key Observations: SVG-EAR stays closest to dense here (45.7–46.3 dB per frame against 37.7–38.4 dB for DecoupAttn, 46.02 and 38.18 dB over the clip), the largest gap of the four cases. The difference of DecoupAttn is a faint brightness offset spread evenly over the floor, together with thin lines along the outlines of the stool, the cans, and the box where the objects occlude one another. The objects and the floor texture match the dense output in every frame. The brightness offset is the visible form of an error in the estimated mass: the exact and the compensated parts share one softmax denominator (Eq. (5)), so a small error in $\hat{Z}_i$ scales a whole region evenly instead of leaving a local artifact, and SSIM stays at 0.993–0.994 on every frame. The variance correction of Eq. (4) targets exactly this mass error and cuts it about threefold on the three models of Table 24, so the offset stays faint.

Case 3: Worst clip of a few-step model (Cosmos-Predict2.5, test_19). It is the clip on which DecoupAttn has its lowest PSNR among the 32 test clips (16.39 dB). Figure 15 shows frames 0, 46, and 92 of this clip.

Key Observations: SVG-EAR scores lower still on this clip (15.85 dB over the clip, 15.6–15.9 dB per frame against 16.2–16.4 dB for DecoupAttn). Both sparse outputs keep the greenhouse, the hanging baskets, and the person bend-

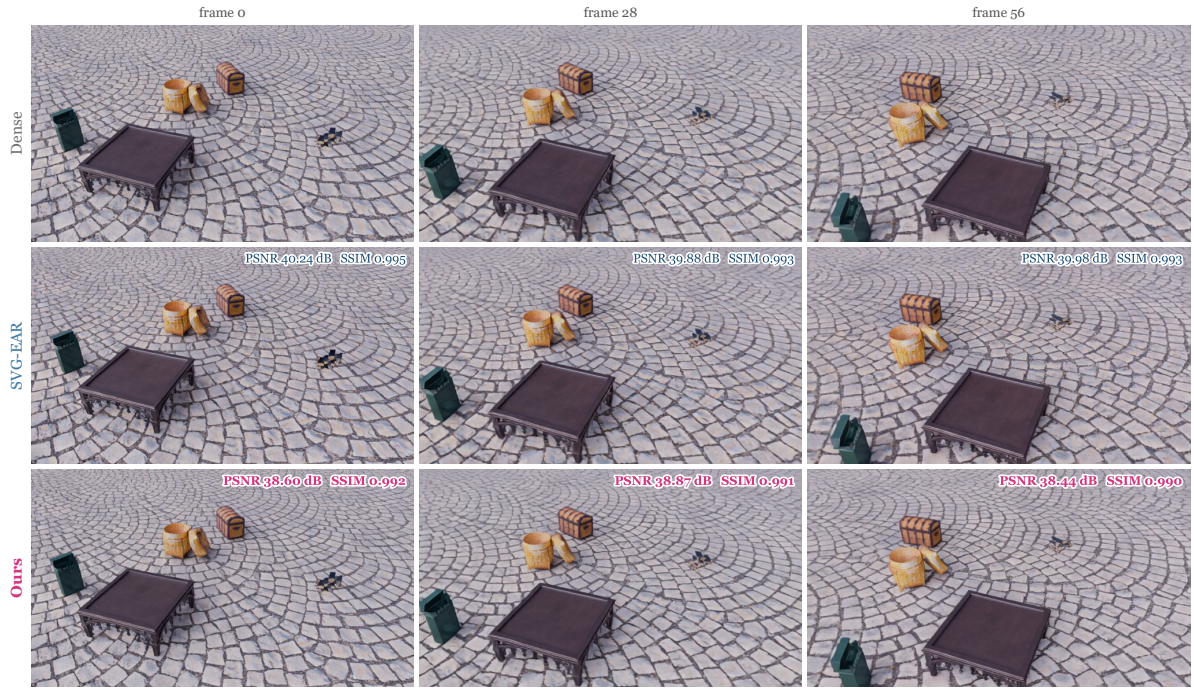

Figure 13. **Case 1: fast motion** (UniRelight, syn_test_0004, frames 0, 28, and 56). Rows: dense attention, SVG-EAR, and DecoupAttn, with per-frame PSNR and SSIM against dense attention.

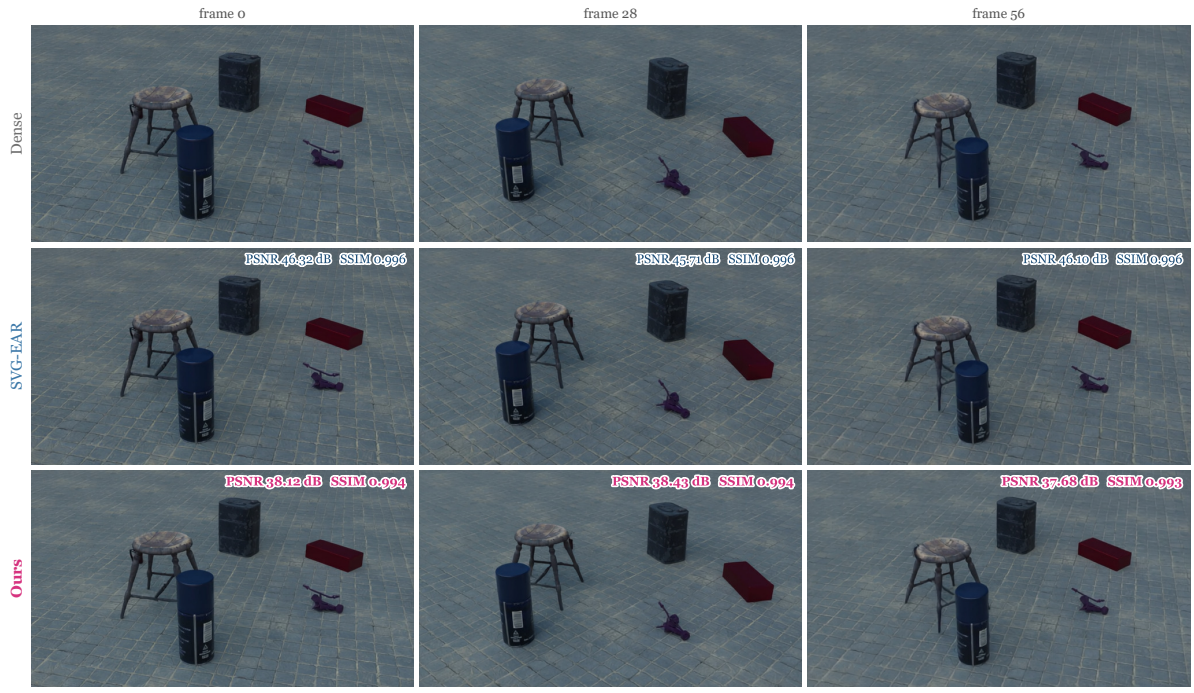

Figure 14. **Case 2: occlusion** (UniRelight, syn_test_0006, frames 0, 28, and 56). Rows: dense attention, SVG-EAR, and DecoupAttn, with per-frame PSNR and SSIM against dense attention.

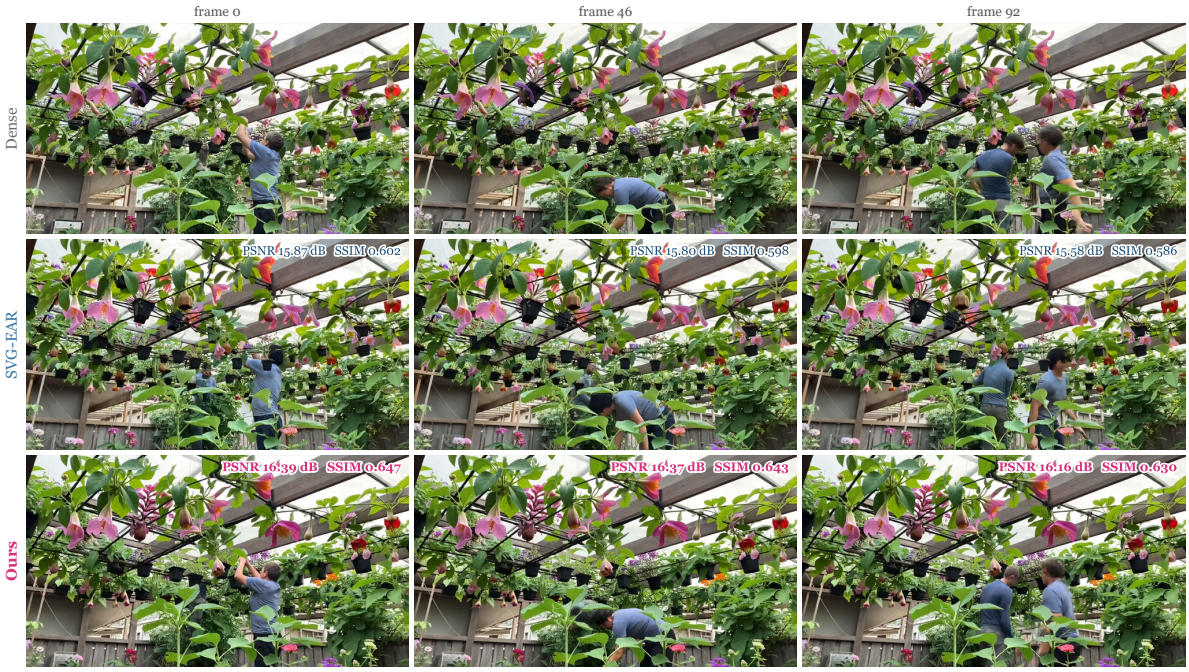

Figure 15. **Case 3: worst clip of a few-step model** (Cosmos-Predict2.5, test_19, frames 0, 46, and 92). Rows: dense attention, SVG-EAR, and DecoupAttn, with per-frame PSNR and SSIM against dense attention.

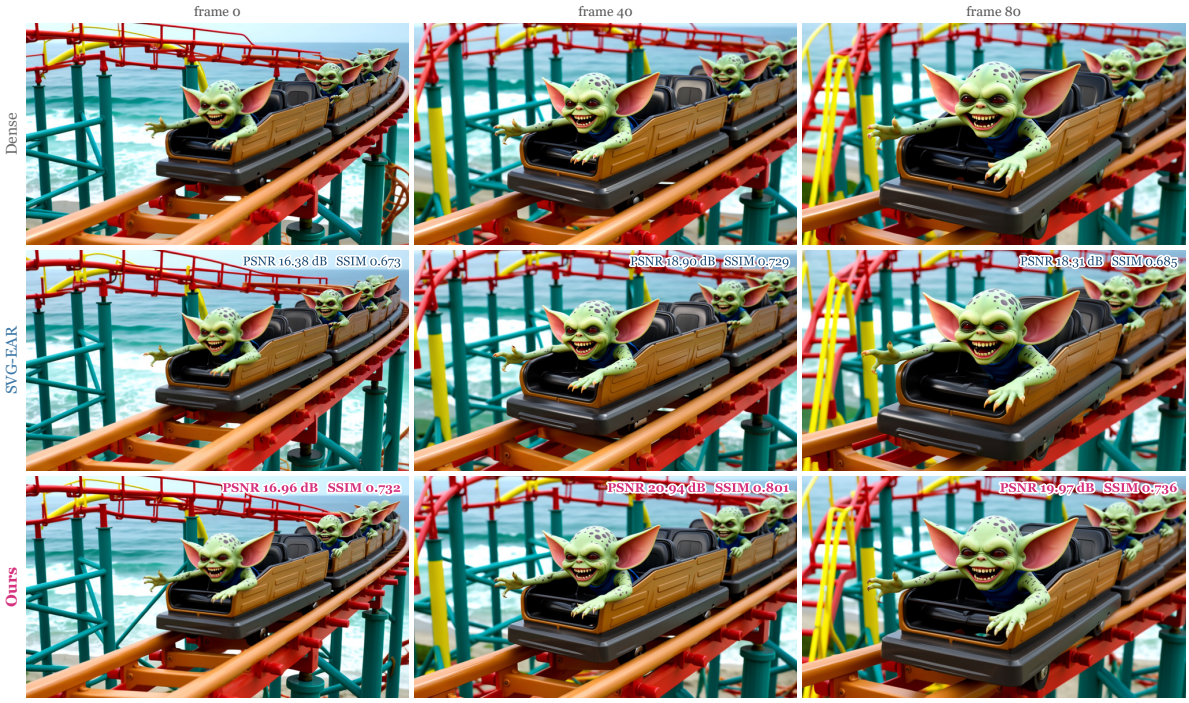

Figure 16. **Case 4: worst clip of a distilled 14B model** (Wan2.1-14B, test_26, frames 0, 40, and 80). Rows: dense attention, SVG-EAR, and DecoupAttn, with per-frame PSNR and SSIM against dense attention.

ing over the plants, but they redraw the dense canopy of small leaves and flowers at shifted positions, so the error covers the whole canopy evenly. The four-step sampler amplifies early differences the most (Figure 4d), and a prompt with this much fine texture turns them into a low PSNR for every sparse method. DecoupAttn keeps a higher SSIM on every frame (0.630–0.647 against 0.586–0.602), so even where both methods redraw the leaves it keeps more of the dense layout. The even spread of its error is what the cancellation of Proposition 1 predicts for compensated residuals, and it keeps the worst clip of Cosmos-Predict2.5 free of local failures.

Case 4: Worst clip of a distilled 14B model (Wan2.1-14B, test_26). It is the clip on which DecoupAttn has its lowest PSNR among the 32 test clips (19.95 dB). Figure 16 shows frames 0, 40, and 80 of this clip.

Key Observations: DecoupAttn again stays ahead of SVG-EAR (19.95 against 18.24 dB over the clip). The first frame is the hardest for both (16.96 and 16.38 dB), and later frames recover (20.94 and 18.90 dB at frame 40). Both outputs keep the goblin, the car, and the coaster and shift the thin rails and struts of the background, while SVG-EAR also departs more on the goblin's face and the front car, which DecoupAttn keeps close to the dense frames. Per frame, DecoupAttn leads by 0.58, 2.04, and 1.66 dB and keeps SSIM at 0.732–0.801 against 0.673–0.729, so its lead holds on every frame of the clip.

Summary. Across the four cases, from two video models and a renderer, the error of DecoupAttn is spread, not local: a fine texture shift under fast motion, an even brightness offset under occlusion, and redrawn leaves on the hardest prompt, never a broken object or a block-shaped artifact. DecoupAttn leads SVG-EAR on every frame of both four-step cases, and on both UniRelight cases its frames exceed 37.6 dB with SSIM $\geq$ 0.99 and match the dense frames. These are the behaviors that decoupling predicts: the residuals of the fixed blocks cancel and spread (Section 4.1), and the clusters of each video follow its own content as objects move and occlude each other (Appendix D.8).

Components at work. The cases show the three components of DecoupAttn in the frames. Decoupling lets fixed exact blocks coexist with semantic clusters that cover every key, so the textures of Cases 1 and 3 change as spread noise, not at block borders (Section 4.1). One offline plan per model sets the exact blocks of all four clips and each clip clusters its own keys, so a single frozen plan serves fast motion, occlusion, and dense texture, as it serves held-out content (Table 32). Distortion Budget matters on the four-step models of Cases 3 and 4, where removing its sensitivity costs 2.30 and 2.14 dB on average (Table 2). On UniRelight, the model of Cases 1 and 2, reassignment gains 3.63 dB over frozen assignments on test clips, the largest gain of the ten models (Table 26), so the clusters follow the moving and occluding objects.

Stability along the clip. The per-frame PSNR of DecoupAttn varies by at most 0.75 dB within Cases 1–3 (38.44–38.87, 37.68–38.43, and 16.16–16.39 dB), so its error does not build up from first to last frame, and in Case 4 it is highest after the first frame (20.94 dB, frame 40). SVG-EAR is as flat (39.88–40.24, 45.71–46.32, and 15.58–15.87 dB), so the per-frame scores follow the prompt, not the position of the frame in the clip.

Hard prompts. The prompts of Cases 3 and 4 are hard for every sparse method: SVG-EAR falls on them from its test means of 21.85 and 22.95 dB to 15.85 and 18.24 dB (Table 1), and DecoupAttn leads on both (16.39 and 19.95 dB). Their dense canopy of small leaves and their thin rails and struts are fine, repeated structures, and on them DecoupAttn keeps the layout, the people, and the vehicles of the dense video: the gardener and the goblin keep their poses and positions in every frame (Figures 15 and 16).

Structure and quality. DecoupAttn leads in SSIM on every frame of Cases 3 and 4, by 0.04–0.07, so it keeps more of the dense layout where both methods redraw fine detail, and its LPIPS is the lowest of all sparse methods on the three models of the cases (0.102, 0.057, and 0.006, Table 1). VBench stays within 0.007 of dense attention on every model (Table 7), and the error against ground truth matches that of dense attention on both renderers that provide it (Table 35), so the fidelity seen in the four cases holds across the test sets of all ten models.