arXiv:2611.08359v1 [math.OC] 5 Nov 2026

A Theoretical Analysis of a Bilevel Optimization Formulation for Combined Local-Global Optical Flow Models

El Mostafa Kalmoun[1*]

[1*]School of Science and Engineering, Al Akhawayn University in Ifrane, Ifrane, 53000, Morocco.

Corresponding author(s). E-mail(s): E.Kalmoun@aui.ma;

Abstract

This paper analyzes a bilevel formulation for learning the regularization parameter of combined local-global (CLG) optical flow models. The lower-level problem computes the flow from a Lucas–Kanade data term and a quadratic or total variation (TV) regularizer, and the upper-level problem selects the parameter that minimizes the squared distance to a ground-truth flow. The setting is finite-dimensional, and a full-rank assumption makes the lower-level solution unique. For the quadratic model we derive sufficient conditions for the existence of an optimal parameter, for every optimal parameter to be positive, and for strict local optimality. For TV regularization an optimal parameter always exists, because the lower-level solution vanishes for large parameters, and we give a sufficient condition for positivity that accounts for the coupling of the two flow components through the local structure tensors. Synthetic examples illustrate the conditions, including cases in which they fail.

Keywords: Bilevel Optimization, Optical Flow, Regularization Parameter Learning, Total Variation, Local-Global Models

1 Introduction

Variational methods estimate the optical flow between two frames by minimizing an energy that combines a data term, derived from brightness constancy, with a regularizer weighted by a parameter $\alpha > 0$. The quality of the estimate depends strongly on α. Too small a value leaves the flow noisy, and too large a value oversmooths it. This

paper studies the choice of α by bilevel learning from ground-truth flow and asks, in a fully discrete setting, when an optimal parameter exists and when it is positive.

The linearized brightness constancy constraint gives one equation per pixel for the two flow components, so the flow is not determined locally (the aperture problem). Lucas and Kanade [1] assume that the flow is constant in a window around each pixel and solve the resulting local least-squares problems. Horn and Schunck [2] instead add a global quadratic penalty on the flow gradient, which yields dense flow fields but smooths motion boundaries. Total variation (TV) regularization, introduced for image denoising [3], preserves such boundaries. Shulman and Hervé [4] proposed an early non-quadratic regularizer for discontinuous flow fields, and TV is widely used for optical flow [5–11]. We use its isotropic form, applied to each flow component. TV is convex but not differentiable, and it is usually minimized by primal–dual methods [7, 9, 12] or replaced by a smooth approximation such as the Huber function [13]. The combined local-global (CLG) model of Bruhn et al. [14] joins the Lucas–Kanade data term with a global regularizer. It inherits the robustness of the local method and the density of the global one, and it can be minimized efficiently, for instance by multigrid methods [15]. Spatially varying parameters $\alpha(p)$ adapted to image features have also been proposed [16–18].

Given two frames and a ground-truth flow $w_g \in \mathbb{R}^{2N}$ on a grid Ω of N pixels, we consider

$$\min_{\alpha \geq 0} \; \|w_\alpha - w_g\|_2^2 \quad \text{subject to} \quad w_\alpha = \operatorname*{arg\,min}_{w \in \mathbb{R}^{2N}} \; \frac{1}{2} \sum_{p \in \Omega} \|A_p w(p) - b_p\|_2^2 + \alpha R(w), \tag{1.1}$$

where $A_p w(p) = b_p$ is the Lucas–Kanade system of the window around the pixel p, and R is either a quadratic penalty on the discrete flow gradient or the TV seminorm. Section 2 defines all objects precisely.

Problem (1.1) rests on four modeling choices. First, the parameter is learned from ground truth. Classical parameter-choice rules, such as the discrepancy principle [19], the L-curve [20], generalized cross-validation [21] and Stein's unbiased risk estimate [22], need no ground truth. They rely instead on knowledge of the noise or on heuristic criteria, and they do not target the flow error itself. Bilevel learning uses training data with known ground truth and minimizes the reconstruction error directly. Second, the upper-level loss is the squared Euclidean distance, which equals the sum over all pixels of the squared endpoint errors, a standard accuracy measure for optical flow [23]. It is smooth in w, whereas the endpoint error itself and the angular error [24] are not differentiable where the error vanishes. Third, a single scalar parameter is the natural first case. It is the setting in which existence and positivity can be characterized by explicit conditions, and spatially varying parameters, as in the bilevel framework of [25], can build on it. Fourth, the CLG data term makes the lower-level problem well posed. When every window contains two pixels with linearly independent gradients, the data term is strongly convex, and the lower-level solution is unique for every $\alpha \geq 0$, including $\alpha = 0$ (Lemma 1). The Horn–Schunck data term alone is never strongly convex, because its matrix $J^\top J$ in (2.1) has a null space of dimension at least N.

Bilevel learning of regularization parameters in imaging was established by Kunisch and Pock [26] for variational denoising and by De los Reyes and Schönlieb [27] for learning noise models in TV denoising. Its theory includes existence results for a broad class of regularizers [28], refined stationarity systems for TV [29], stationarity and second-order conditions through a reformulation as a mathematical program with complementarity constraints [30], and spatially dependent parameters [25]. For finite-dimensional linear inverse problems with a continuously differentiable regularizer, Ehrhardt, Gazzola and Scott [31] give a sufficient condition for a positive optimal parameter. Our positivity result for the quadratic model (Theorem 3) is a special case of their condition. Applications include blind deconvolution [32] and the learning of fractional differential operators [33], and Crockett and Fessler [34] survey the field. Computational work addresses inexact lower-level solves through derivative-free methods [35], inexact hypergradients [36] and single-loop schemes [37]. Learning through truncated or approximate inference [38, 39] and learned variational inference layers [40] aim at empirical performance rather than at the analysis of a regularization weight. For optical flow, Jiang et al. [41] use a bilevel scheme that couples flow and camera motion in unsupervised deep learning. Closest to the present setting, El Hakoume et al. [42] learn the two Lamé coefficients of a fluid optical flow registration model by bilevel supervised learning, within a multi-frame super-resolution method. Positivity is imposed through a fixed box, uniqueness of the lower-level solution is assumed, and the bilevel problem "is not theoretically studied" because their model lacks the required regularity. Related work by Laghrib and coauthors learns spatially varying weights for nonlocal deblurring [43]. To our knowledge, no previous work establishes existence, positivity and second-order conditions for learning the regularization parameter of a variational optical flow model from ground-truth flow. Here all objects are finite-dimensional, the proofs are direct, and the conditions are explicit and computable from the data for both regularizers.

The main contributions are the following.

- A finite-dimensional bilevel formulation for CLG optical flow in which a scalar regularization parameter is learned from ground-truth flow.
- Existence of an optimal parameter. For any convex regularizer that vanishes only at zero, a sufficient condition is that the Lucas–Kanade estimate be closer to the ground truth than the zero flow (Proposition 2). For TV with zero boundary values this condition is not needed, because the lower-level solution vanishes beyond a finite value of α (Corollary 5).
- Sufficient conditions under which $\alpha = 0$ is not optimal, so that every optimal parameter is positive, for the quadratic model (Theorem 3, an instance of the condition of [31]) and for TV (Theorem 8), and second-order sufficient conditions for strict local optimality in the quadratic model (Theorem 4). The TV condition is not covered by [31], because TV is not differentiable, and it differs structurally from the analyses of scalar data terms [28, 29]. The Lucas–Kanade data term couples the two flow components through 2×2 blocks that vary from pixel to pixel. This coupling requires a four-way partition of the pixels, a term $\mathcal{R}_{\mathrm{mix}}$ for the pixels at which only one flow component has a nonzero gradient, and a term $\mathcal{R}_{\mathrm{var}}$ for the spatial variation of the blocks. These terms vanish for spatially uniform scalar data terms.

- Synthetic examples that illustrate the conditions, including cases in which the sufficient conditions fail.

Section 2 introduces the discrete setting and states the bilevel problem. Section 3 analyzes quadratic regularization and Section 4 total variation regularization. Section 5 presents the examples, and Section 6 discusses the limitations of the analysis and concludes.

2 Discrete Setting and Problem Formulation

All objects in this paper are finite-dimensional. The image domain Ω is a rectangular grid of $N = n_y n_x$ pixels with unit grid spacing, ordered row by row. Scalar fields on Ω, such as images, are vectors in $\mathbb{R}^N$. A flow field is $w = (u; v) \in \mathbb{R}^{2N}$ with $u, v \in \mathbb{R}^N$, and its value at a pixel p is $w(p) = (u(p), v(p))^\top \in \mathbb{R}^2$.

For two frames $I_1, I_2 \in \mathbb{R}^N$, let $I_x, I_y \in \mathbb{R}^N$ be the spatial derivatives and $I_t := I_2 - I_1$ the temporal derivative. The computation of I_x and I_y is described at the beginning of Section 5. With $\nabla I(p) := (I_x(p), I_y(p))^\top$ and $J := \begin{bmatrix}\operatorname{diag}(I_x) & \operatorname{diag}(I_y)\end{bmatrix} \in \mathbb{R}^{N \times 2N}$, the linearized brightness constancy constraint reads

$$Jw + I_t = 0 \in \mathbb{R}^N, \qquad \text{that is,} \qquad \nabla I(p)^\top w(p) + I_t(p) = 0 \quad \text{for every } p \in \Omega. \tag{2.1}$$

The Lucas–Kanade data term assigns to each pixel p a window $\Omega_p \subset \Omega$ that contains p, and $n_p := |\Omega_p|$. Imposing (2.1) at every $q \in \Omega_p$ with the common value $w(p)$ gives the local system $A_p w(p) = b_p$ with

$$A_p := \big[\nabla I(q)^\top\big]_{q \in \Omega_p} \in \mathbb{R}^{n_p \times 2}, \qquad b_p := -\big[I_t(q)\big]_{q \in \Omega_p} \in \mathbb{R}^{n_p}. \tag{2.2}$$

The data term is

$$D(w) := \frac{1}{2} \sum_{p \in \Omega} \|A_p w(p) - b_p\|_2^2 = \frac{1}{2}\, w^\top \mathcal{M}_{\mathrm{LK}} w - f_{\mathrm{LK}}^\top w + c, \tag{2.3}$$

where $f_{\mathrm{LK}} \in \mathbb{R}^{2N}$ collects the vectors $A_p^\top b_p$ and $c := \frac{1}{2} \sum_{p \in \Omega} \|b_p\|_2^2$. Let Π be the permutation matrix that orders the unknowns pixel by pixel:

$$\Pi w = (u(p_1), v(p_1), \ldots, u(p_N), v(p_N))^\top.$$

Stacking all local systems gives the global system

$$Aw = b, \tag{2.4}$$

whose matrix A has $n := \sum_p n_p$ rows and $2N$ columns, and $A\Pi^\top = \operatorname{diag}(A_{p_1}, \ldots, A_{p_N})$ is block diagonal with rectangular $n_p \times 2$ blocks. Hence

$$\mathcal{M}_{\mathrm{LK}} = A^\top A = \Pi^\top \operatorname{diag}\big(A_{p_1}^\top A_{p_1}, \ldots, A_{p_N}^\top A_{p_N}\big)\, \Pi.$$

In the ordering $w = (u; v)$, $\mathcal{M}_{\mathrm{LK}}$ consists of four diagonal $N \times N$ blocks formed from the entries of the matrices $A_p^\top A_p$.

Throughout the paper we assume

$$\textbf{(A1)} \quad \operatorname{rank} A_p = 2 \text{ for every } p \in \Omega,$$

that is, every window contains two pixels with linearly independent image gradients. Then $A_p^\top A_p \succ 0$ for every p, so $\mathcal{M}_{\mathrm{LK}} \succ 0$, and D has the unique minimizer $w_0 = \mathcal{M}_{\mathrm{LK}}^{-1} f_{\mathrm{LK}}$, the Lucas–Kanade solution

$$\begin{gathered} w_0(p) = (A_p^\top A_p)^{-1} A_p^\top b_p, \\ A_p^\top A_p = \sum_{q \in \Omega_p} \nabla I(q) \nabla I(q)^\top, \qquad A_p^\top b_p = -\sum_{q \in \Omega_p} I_t(q)\, \nabla I(q). \end{gathered} \tag{2.5}$$

The two regularizers use two difference operators, and their boundary conditions are part of the model. For the quadratic regularizer, K_D takes forward differences between horizontally and vertically adjacent pixels, with zero ghost values on all four sides of the grid, so $K_D \in \mathbb{R}^{m \times N}$ with $m = n_y(n_x + 1) + n_x(n_y + 1)$. The matrix $K_D^\top K_D$ is the five-point Laplacian with homogeneous Dirichlet conditions. On flows we set $\mathbf{K}_D := \operatorname{diag}(K_D, K_D)$ and

$$L := \mathbf{K}_D^\top \mathbf{K}_D = \operatorname{diag}\big(K_D^\top K_D,\, K_D^\top K_D\big) \in \mathbb{R}^{2N \times 2N}.$$

Since K_D is injective, $L \succ 0$. The matrix $\mathcal{L}_{\mathrm{glob}} := -L$ is the discrete vector Laplacian.

For the TV regularizer, the gradient of a scalar field s at a pixel p is

$$\nabla s(p) := (Ks)(p) := \big(\bar{s}(p + e_x) - s(p),\ \bar{s}(p + e_y) - s(p)\big)^\top \in \mathbb{R}^2,$$

where e_x and e_y are the shifts to the next column and the next row, and $\bar{s}$ equals s on Ω and 0 outside Ω. Thus the ghost values beyond the last column and the last row are zero. The map $K : \mathbb{R}^N \to \mathbb{R}^{2N}$ is injective. On flows, $\mathbf{K} := \operatorname{diag}(K, K) : \mathbb{R}^{2N} \to \mathbb{R}^{4N}$, and a dual field is $z = (z_u, z_v) \in \mathbb{R}^{4N}$ with $z_u(p), z_v(p) \in \mathbb{R}^2$. The discrete divergence is $\operatorname{div} := -K^\top$ for scalar fields and $\mathbf{div} := -\mathbf{K}^\top : \mathbb{R}^{4N} \to \mathbb{R}^{2N}$ for flows, so that

$$\sum_{p \in \Omega} \langle \nabla s(p), z(p) \rangle = -\sum_{p \in \Omega} s(p)\, (\operatorname{div} z)(p) \tag{2.6}$$

holds exactly, with no boundary terms. Forward differences of other fields on Ω, such as residuals, are also taken with K. The two operators differ at the border, since K has no differences across the left and the top sides, and $K^\top K \neq K_D^\top K_D$. The analysis of the quadratic model uses only $L \succ 0$. For TV, existence uses only the injectivity of K, and the positivity analysis also uses the forward-difference form of K (Lemma 6).

With these operators, the quadratic (CLG) energy and the TV energy are

$$E(w, \alpha) := D(w) + \frac{\alpha}{2} \|\mathbf{K}_D w\|_2^2 = D(w) + \frac{\alpha}{2}\, w^\top L w, \tag{2.7}$$

$$E(w,\alpha) := D(w) + \alpha\,\mathcal{R}_{TV}(w), \qquad \mathcal{R}_{TV}(w) := \sum_{p\in\Omega} \big(\|\nabla u(p)\|_2 + \|\nabla v(p)\|_2\big). \tag{2.8}$$

The Horn–Schunck energy [2]

$$\tfrac{1}{2}\|Jw + I_t\|_2^2 + \tfrac{\alpha}{2}\|\mathbf{K}_D w\|_2^2 \tag{2.9}$$

is the special case $\Omega_p = \{p\}$ of (2.7), for which (A1) fails.

Given a ground-truth flow $w_g \in \mathbb{R}^{2N}$, we consider

$$\min_{\alpha\geq 0}\; Q(\alpha) := \|w_\alpha - w_g\|_2^2 \quad \text{subject to} \quad w_\alpha = \operatorname*{arg\,min}_{w\in\mathbb{R}^{2N}} E(w,\alpha), \tag{2.10}$$

with E given by (2.7) in Section 3 and by (2.8) in Section 4.

Lemma 1 *Assume (A1), and let $R : \mathbb{R}^{2N} \to [0,\infty)$ be convex. For every $\alpha \geq 0$, the function $D+\alpha R$ is strongly convex with modulus $\mu := \lambda_{\min}(\mathcal{M}_{\mathrm{LK}}) > 0$ and has exactly one minimizer. In particular, w_α in (2.10) is uniquely defined for both energies and every $\alpha \geq 0$, including $\alpha = 0$, and Q is a single-valued function on $[0,\infty)$.*

Proof Since $\mathcal{M}_{\mathrm{LK}} - \mu I \succeq 0$, the function $D - \frac{\mu}{2}\|\cdot\|_2^2$ is convex, so $D+\alpha R$ is strongly convex. A finite strongly convex function on $\mathbb{R}^{2N}$ has exactly one minimizer. □

Without (A1), the lower-level problem may have many solutions, for instance at $\alpha = 0$ because of the aperture problem. Problem (2.10) would then require an optimistic or a pessimistic formulation, which we do not consider.

Remark 1 When the ground truth is only partially available or carries confidence weights, the loss can be replaced by $Q_W(\alpha) := \|W(w_\alpha - w_g)\|_2^2$ with $W := \operatorname{diag}(\omega,\omega) \in \mathbb{R}^{2N\times 2N}$ and weights $\omega \in \mathbb{R}^N$, $\omega \geq 0$. The lower-level solutions do not depend on W. The corresponding changes to the conditions of Sections 3 and 4 are stated in Remark 9.

3 Quadratic Regularization

In this section the lower-level energy is the quadratic CLG energy (2.7),

$$E(w,\alpha) = \frac{1}{2}\sum_{p\in\Omega} \|A_p w(p) - b_p\|_2^2 + \frac{\alpha}{2}\, w^\top L w. \tag{3.1}$$

At the minimizer the gradient of $E(\cdot,\alpha)$ vanishes. For each pixel this gives

$$A_p^\top \big(A_p w(p) - b_p\big) + \alpha\,(Lw)(p) = 0, \tag{3.2}$$

where $(Lw)(p) \in \mathbb{R}^2$ collects the two entries of Lw at p. Globally,

$$(\mathcal{M}_{\mathrm{LK}} + \alpha L)\, w = f_{\mathrm{LK}}. \tag{3.3}$$

Under (A1), $\mathcal{M}_{\mathrm{LK}} \succ 0$, and $L \succ 0$. The matrix $\mathcal{M}_{\mathrm{LK}} + \alpha L$ is therefore positive definite for every $\alpha \geq 0$. It remains positive definite for small negative α: since $w^\top(\mathcal{M}_{\mathrm{LK}} + \alpha L)w \geq \big(\lambda_{\min}(\mathcal{M}_{\mathrm{LK}}) - |\alpha|\,\lambda_{\max}(L)\big)\|w\|_2^2$, this holds for every $\alpha > -\alpha_0$ with $\alpha_0 := \lambda_{\min}(\mathcal{M}_{\mathrm{LK}})/\lambda_{\max}(L) > 0$. Hence

$$w_\alpha = X(\alpha) f_{\mathrm{LK}}, \qquad X(\alpha) := (\mathcal{M}_{\mathrm{LK}} + \alpha L)^{-1}, \tag{3.4}$$

and $\alpha \mapsto w_\alpha$ and Q are infinitely differentiable on the open interval $(-\alpha_0, \infty)$. In particular, $Q'(0)$ is an ordinary two-sided derivative, although only $\alpha \geq 0$ is admissible in the bilevel problem. From $X'(\alpha) = -X(\alpha)LX(\alpha)$ and $X(\alpha)Lw_\alpha = -w'_\alpha$,

$$w'_\alpha = -X(\alpha)Lw_\alpha, \qquad w''_\alpha = -2X(\alpha)Lw'_\alpha, \tag{3.5}$$

and therefore

$$Q'(\alpha) = 2\langle w'_\alpha,\, w_\alpha - w_g\rangle, \qquad Q''(\alpha) = 2\|w'_\alpha\|_2^2 - 4\langle X(\alpha)Lw'_\alpha,\, w_\alpha - w_g\rangle. \tag{3.6}$$

Substituting (3.4) reduces (2.10) to the single-level problem

$$\min_{\alpha \geq 0}\; Q(\alpha) = \|X(\alpha) f_{\mathrm{LK}} - w_g\|_2^2. \tag{3.7}$$

If every minimizer satisfies $\alpha^* > 0$, then some amount of regularization strictly improves on the unregularized Lucas–Kanade estimate.

Remark 2 In practice $X(\alpha)$ is never formed. The symmetric positive definite system (3.3) is solved by preconditioned conjugate gradients or multigrid, and w'_α requires one more solve with the same matrix. Since α is a scalar, Q can be minimized by a bounded one-dimensional search, which is what the examples of Section 5 use, or by a gradient-based method with the analytical gradient (3.6).

The proof of the next existence result uses the minimizing-sequence argument of [26, Proposition 3.1] in finite dimension.

Proposition 2 *Assume (A1). Let $R : \mathbb{R}^{2N} \to [0, \infty)$ be convex with $R(w) = 0$ if and only if $w = 0$, and let w_α be the minimizer of $D + \alpha R$. Then:*

(i) $\alpha \mapsto w_\alpha$ is continuous on $[0, \infty)$, $\alpha R(w_\alpha) \leq \frac{1}{2} w_0^\top \mathcal{M}_{\mathrm{LK}} w_0$ for all $\alpha \geq 0$, and $w_\alpha \to 0$ as $\alpha \to \infty$;

(ii) Q is continuous on $[0, \infty)$, and $Q(\alpha) \to \|w_g\|_2^2$ as $\alpha \to \infty$;

(iii) if

$$\|w_0 - w_g\|_2 < \|w_g\|_2, \tag{3.8}$$

then Q has a minimizer $\alpha^* \in [0, \infty)$, and $Q(\alpha^*) < \|w_g\|_2^2$.

The regularizers $\frac{1}{2}\|\mathbf{K}_D w\|_2^2$ and $\mathcal{R}_{TV}$ satisfy these hypotheses, because K_D and K are injective.

Proof Let $\|x\|_{\mathcal{M}}^2 := x^\top \mathcal{M}_{\mathrm{LK}} x$, so that $D(w) - D(w_0) = \frac{1}{2}\|w - w_0\|_{\mathcal{M}}^2$.

(i) Since $R(0) = 0$, we have $D(w_\alpha) + \alpha R(w_\alpha) \le D(0)$, and $D(0) - D(w_0) = \frac{1}{2}\|w_0\|_{\mathcal{M}}^2$. Hence $\frac{1}{2}\|w_\alpha - w_0\|_{\mathcal{M}}^2 + \alpha R(w_\alpha) \le \frac{1}{2}\|w_0\|_{\mathcal{M}}^2$. This gives the bound on $\alpha R(w_\alpha)$ and shows that $\{w_\alpha : \alpha \ge 0\}$ is bounded. Let $\alpha_k \to \infty$. Then $R(w_{\alpha_k}) \to 0$, and every accumulation point $\bar{w}$ of (w_{α_k}) satisfies $R(\bar{w}) = 0$ by continuity of R, so $\bar{w} = 0$. Hence $w_{\alpha_k} \to 0$. Now let $\alpha_k \to \hat{\alpha} \in [0, \infty)$, and let $\bar{w}$ be an accumulation point of (w_{α_k}). For every w, $D(w_{\alpha_k}) + \alpha_k R(w_{\alpha_k}) \le D(w) + \alpha_k R(w)$. Passing to the limit along the corresponding subsequence gives $D(\bar{w}) + \hat{\alpha} R(\bar{w}) \le D(w) + \hat{\alpha} R(w)$, so $\bar{w} = w_{\hat{\alpha}}$ by Lemma 1. Hence $w_{\alpha_k} \to w_{\hat{\alpha}}$.

(ii) This follows from (i).

(iii) Since $Q(\alpha) \to \|w_g\|_2^2 > Q(0)$, there is $A > 0$ with $Q(\alpha) > Q(0)$ for all $\alpha \ge A$. The continuous function Q attains its minimum on $[0, A]$ at some α^*. This point minimizes Q on $[0, \infty)$, and $Q(\alpha^*) \le Q(0) < \|w_g\|_2^2$. □

Condition (3.8) is sufficient but not necessary, since a finite minimizer can exist when it fails. It cannot, however, be dropped from Proposition 2. Consider an abstract instance with a single pixel, $A_p = I_2$ and the quadratic regularizer. Then $L = 4I_2$ and $w_\alpha = w_0/(1 + 4\alpha)$. With $w_0 = (1, 0)^\top$ and $w_g = (-1, 0)^\top$, $Q(\alpha) = \left(1 + (1 + 4\alpha)^{-1}\right)^2$ decreases strictly to $\|w_g\|_2^2 = 1$ and never attains its infimum. For the quadratic energy, w_α never vanishes at a finite α when $f_{\mathrm{LK}} \neq 0$, because $X(\alpha)$ is invertible. The TV case behaves differently (Section 4).

Theorem 3 *Assume (A1). Then $Q'(0) = -2\langle \mathcal{M}_{\mathrm{LK}}^{-1} L w_0,\, w_0 - w_g \rangle$. If*

$$\left\langle \mathcal{M}_{\mathrm{LK}}^{-1} L w_0,\; w_0 - w_g \right\rangle > 0, \tag{3.9}$$

then $Q(\alpha) < Q(0)$ for all sufficiently small $\alpha > 0$, and every minimizer of Q on $[0, \infty)$ satisfies $\alpha^* > 0$.

Proof At $\alpha = 0$, (3.5) gives $w_0' = -\mathcal{M}_{\mathrm{LK}}^{-1} L w_0$, and (3.6) then gives the formula for $Q'(0)$. Under (3.9), $Q'(0) < 0$, so Q decreases on a right neighborhood of 0 and $\alpha = 0$ is not a minimizer. □

Theorem 3 is the instance, for the CLG data term and the quadratic regularizer, of a general sufficient condition for positivity due to Ehrhardt, Gazzola and Scott [31, Corollary 1], which covers convex, continuously differentiable regularizers. For an identity forward operator and a quadratic regularizer, their condition is the one of Kunisch and Pock [26, Proposition 3.1].

Remark 3 Condition (3.9) has a first-order meaning. Since $w_\alpha = w_0 + \alpha w_0' + O(\alpha^2)$ with $w_0' = -\mathcal{M}_{\mathrm{LK}}^{-1} L w_0$,

$$Q(\alpha) = Q(0) + 2\alpha \,\langle w_0',\, w_0 - w_g \rangle + O(\alpha^2).$$

Condition (3.9) therefore states that the first-order smoothing correction w_0' has a negative component along the residual $w_0 - w_g$. At each pixel, $w_0'(p) = -(A_p^\top A_p)^{-1}(Lw_0)(p)$, so the correction can be large where $A_p^\top A_p$ is nearly singular. This condition concerns the sum over all pixels, not each pixel separately.

Remark 4 Suppose the regularizer uses Neumann-type differences (forward differences set to zero in the last column and the last row), in either the quadratic or the TV case. Its kernel is then the space $\mathcal{Z} := \operatorname{range} Z$ of constant flows, where $Z := \begin{bmatrix} \mathbf{1} & 0 \\ 0 & \mathbf{1} \end{bmatrix} \in \mathbb{R}^{2N\times 2}$ and $\mathbf{1} \in \mathbb{R}^N$ is the vector of ones. Lemma 1 still applies. As $\alpha \to \infty$, w_α converges to the minimizer of D on $\mathcal{Z}$,

$$w_\alpha \to P_{\mathcal{M}} w_0, \qquad P_{\mathcal{M}} := Z\big(Z^\top \mathcal{M}_{\mathrm{LK}} Z\big)^{-1} Z^\top \mathcal{M}_{\mathrm{LK}},$$

which is the constant flow $\big(\sum_p A_p^\top A_p\big)^{-1} \sum_p A_p^\top b_p$, the least-squares solution of all local systems together. The proof is the argument of Proposition 2, with 0 replaced by this minimizer. A sufficient condition for a finite minimizer is therefore

$$\|w_0 - w_g\|_2 < \|P_{\mathcal{M}} w_0 - w_g\|_2. \tag{3.10}$$

The 2×2 matrix $Z^\top \mathcal{M}_{\mathrm{LK}} Z = \sum_p A_p^\top A_p$ couples the two flow components. The operator $P_{\mathcal{M}}$ is the $\mathcal{M}_{\mathrm{LK}}$-orthogonal projector onto $\mathcal{Z}$, and it coincides with the Euclidean projector (componentwise averaging) if and only if all matrices $A_p^\top A_p$ are equal.

Theorem 4 *Assume (A1), and let $\alpha^* \geq 0$ satisfy the first-order necessary conditions for problem (3.7),*

$$\begin{cases} Q'(0) \geq 0 & \text{if } \alpha^* = 0, \\ Q'(\alpha^*) = 0 & \text{if } \alpha^* > 0. \end{cases} \tag{3.11}$$

Then α^* is a strict local minimizer of Q on $[0, \infty)$ if either:

(a) $\alpha^* = 0$ and $Q'(0) > 0$; or
(b) $\alpha^* > 0$, $Lw_{\alpha^*} \neq 0$, and

$$\|w_{\alpha^*} - w_g\|_2 < \frac{\|w'_{\alpha^*}\|_2^2}{2\,\|X(\alpha^*) L w'_{\alpha^*}\|_2}. \tag{3.12}$$

Proof (a) If $\alpha^* = 0$ and $Q'(0) > 0$, then Q is increasing on a right neighborhood of 0.
(b) By (3.6), $Q''(\alpha^*) = H_1 + H_2$ with

$$H_1 := 2\|w'_{\alpha^*}\|_2^2, \qquad H_2 := -4\langle X(\alpha^*) L w'_{\alpha^*},\, w_{\alpha^*} - w_g\rangle.$$

Since $X(\alpha^*)$ and L are invertible, $Lw_{\alpha^*} \neq 0$ gives $w'_{\alpha^*} \neq 0$ and $X(\alpha^*) L w'_{\alpha^*} \neq 0$. By the Cauchy–Schwarz inequality, $|H_2| \leq 4\|X(\alpha^*) L w'_{\alpha^*}\|_2\, \|w_{\alpha^*} - w_g\|_2$, which is smaller than H_1 under (3.12). Hence $Q''(\alpha^*) > 0$, and together with $Q'(\alpha^*) = 0$ this makes α^* a strict local minimizer. □

Remark 5 The bound (3.12) is sufficient but not necessary: if $H_2 \geq 0$, then $Q''(\alpha^*) \geq H_1 > 0$ without it. In the examples of Section 5, strict local optimality is checked by evaluating $Q''(\alpha^*)$ directly.

Proposition 2 and Theorems 3 and 4 do not determine whether Q has a unique global minimizer (see Section 6.1). Because w_α depends nonlinearly on α, Q need not be convex, and multiple local minima cannot be excluded a priori.

4 Total Variation Regularization

This section treats the TV energy (2.8) with the gradient $\nabla = K$ of Section 2. The TV term is not differentiable, so the lower-level optimality condition is an inclusion rather than a linear system, and no closed form such as (3.4) is available.

Write $\mathcal{R}_{TV} = h \circ \mathbf{K}$ with $h(y) := \sum_p \big(\|y_u(p)\|_2 + \|y_v(p)\|_2\big)$ for $y = (y_u, y_v) \in \mathbb{R}^{4N}$, where $y_u(p), y_v(p) \in \mathbb{R}^2$ are the entries at pixel p. Since h is finite and convex, $\partial \mathcal{R}_{TV}(w) = \mathbf{K}^\top \partial h(\mathbf{K}w)$. Hence w_α minimizes $E(\cdot, \alpha)$ in (2.8) if and only if there is a dual field $z_\alpha = (z_{\alpha,u}, z_{\alpha,v}) \in \mathbb{R}^{4N}$ such that, for every pixel p,

$$A_p^\top \big(A_p w_\alpha(p) - b_p\big) = \alpha\, (\mathbf{div}\, z_\alpha)(p), \qquad (\mathbf{div}\, z_\alpha)(p) = \big((\operatorname{div} z_{\alpha,u})(p),\, (\operatorname{div} z_{\alpha,v})(p)\big)^\top, \tag{4.1}$$

$$z_{\alpha,u}(p) \in \partial \|\cdot\|_2\big(\nabla u_\alpha(p)\big), \qquad z_{\alpha,v}(p) \in \partial \|\cdot\|_2\big(\nabla v_\alpha(p)\big). \tag{4.2}$$

Here $\partial \|\cdot\|_2(g) = \{g/\|g\|_2\}$ if $g \neq 0$, and $\partial \|\cdot\|_2(0)$ is the closed unit ball of $\mathbb{R}^2$. In global form,

$$\mathcal{M}_{\mathrm{LK}} w_\alpha - \alpha\, \mathbf{div}\, z_\alpha = f_{\mathrm{LK}}, \qquad z_\alpha \in \partial h(\mathbf{K} w_\alpha). \tag{4.3}$$

The upper-level problem is

$$\min_{\alpha \geq 0}\ Q(\alpha) = \|w_\alpha - w_g\|_2^2 \quad \text{subject to} \quad w_\alpha = \operatorname*{arg\,min}_{w \in \mathbb{R}^{2N}} E(w, \alpha), \tag{4.4}$$

and its constraint holds if and only if (4.3) holds for some z_α. The dual field z_α need not be unique, while w_α is.

Corollary 5 *Assume (A1), and let $\bar\alpha := \|f_{\mathrm{LK}}\|_2 / \sigma_{\min}(K)$, where $\sigma_{\min}(K) > 0$ is the smallest singular value of K. Then $w_\alpha = 0$ for every $\alpha \geq \bar\alpha$. Consequently $Q(\alpha) = \|w_g\|_2^2$ for $\alpha \geq \bar\alpha$, and problem (4.4) has a solution $\alpha^* \in [0, \bar\alpha]$. If in addition*

$$\|w_0 - w_g\|_2 < \|w_g\|_2, \tag{4.5}$$

then $Q(\alpha^*) < \|w_g\|_2^2$, $\alpha^* < \bar\alpha$ and $w_{\alpha^*} \neq 0$. The threshold $\bar\alpha$ is sufficient and need not be the smallest α with $w_\alpha = 0$.

Proof Let $B := \partial h(0) = \{y :\ \|y_u(p)\|_2 \leq 1 \text{ and } \|y_v(p)\|_2 \leq 1 \text{ for all } p\}$. Since D is differentiable with $\nabla D(0) = -f_{\mathrm{LK}}$, the point 0 minimizes $E(\cdot, \alpha)$ if and only if $f_{\mathrm{LK}} \in \alpha\, \mathbf{K}^\top B$. If $f_{\mathrm{LK}} = 0$, then $w_\alpha = 0$ for all $\alpha \geq 0$. Otherwise let $\alpha \geq \bar\alpha > 0$, and put $y := \alpha^{-1} \mathbf{K} (\mathbf{K}^\top \mathbf{K})^{-1} f_{\mathrm{LK}}$, which is well defined because $\mathbf{K}$ is injective. Then $\mathbf{K}^\top y = f_{\mathrm{LK}}/\alpha$. By a singular value decomposition, $\|\mathbf{K}(\mathbf{K}^\top \mathbf{K})^{-1}\| = 1/\sigma_{\min}(\mathbf{K}) = 1/\sigma_{\min}(K)$, so $\|y\|_2 \leq \|f_{\mathrm{LK}}\|_2 / (\alpha\, \sigma_{\min}(K)) \leq 1$. Hence $y \in B$, and $w_\alpha = 0$ by uniqueness.

By Proposition 2(ii), Q is continuous on the compact interval $[0, \bar\alpha]$, so it attains its minimum there at some α^*. Since $Q(\alpha) = \|w_g\|_2^2 = Q(\bar\alpha) \geq Q(\alpha^*)$ for $\alpha \geq \bar\alpha$, the point

α^* solves (4.4). Under (4.5), $Q(\alpha^*) \le Q(0) < \|w_g\|_2^2$, which excludes both $\alpha^* \ge \bar{\alpha}$ and $w_{\alpha^*} = 0$. □

In the TV case, condition (4.5) is therefore not needed for existence. It guarantees that the optimal parameter improves on the zero flow, which is the TV solution for all large α. Example 5 shows the difference: there (4.5) fails, yet an optimal parameter exists and improves on both the unregularized estimate and the zero flow.

Remark 6 With Neumann-type differences, $\mathcal{R}_{TV}$ vanishes on all constant flows. As $\alpha \to \infty$, w_α then converges to the constant flow $P_{\mathcal{M}} w_0$ of Remark 4, and condition (3.10) gives a finite optimal parameter.

We now give sufficient conditions under which $\alpha = 0$ is not optimal, that is, under which some $\alpha > 0$ gives $Q(\alpha) < Q(0)$, so that TV regularization strictly improves on the unregularized Lucas–Kanade estimate. The grid is assumed to have at least two rows and two columns.

For each pixel p, let $M_p := (A_p^\top A_p)^{-1}$, a symmetric positive definite 2×2 matrix with entries $m_{ij}(p)$, $i, j \in \{u, v\}$, where u indexes the first row and column. Let $r(p) = (r_u(p), r_v(p))^\top := w_0(p) - w_g(p)$ be the residual of the unregularized estimate. The spatial variation of M_p is measured by $\partial_d M_p := M_{p+e_d} - M_p$ for $d \in \{x, y\}$ when $p + e_d \in \Omega$, and by the backward difference $\partial_d M_p := M_p - M_{p-e_d}$ when p lies in the last column ($d = x$) or the last row ($d = y$). With the spectral norm $\|\cdot\|_{\mathrm{op}}$, let

$$\rho(p) := \big(\|\partial_x M_p\|_{\mathrm{op}}^2 + \|\partial_y M_p\|_{\mathrm{op}}^2\big)^{1/2}.$$

The pixels are partitioned according to which components of w_0 have a nonzero gradient:

$$\Omega_{11} = \{p : \nabla u_0(p) \neq 0,\ \nabla v_0(p) \neq 0\}, \qquad \Omega_{10} = \{p : \nabla u_0(p) \neq 0,\ \nabla v_0(p) = 0\},$$

$$\Omega_{01} = \{p : \nabla u_0(p) = 0,\ \nabla v_0(p) \neq 0\}, \qquad \Omega_{00} = \{p : \nabla u_0(p) = 0,\ \nabla v_0(p) = 0\}.$$

The data-alignment sums, the variation bound and the mixed-pixel bound are

$$S_u := \sum_{p:\, \nabla u_0(p) \neq 0} \Big[m_{11}(p) \Big\langle \nabla r_u(p), \tfrac{\nabla u_0(p)}{\|\nabla u_0(p)\|_2} \Big\rangle + m_{12}(p) \Big\langle \nabla r_v(p), \tfrac{\nabla u_0(p)}{\|\nabla u_0(p)\|_2} \Big\rangle \Big],$$

$$S_v := \sum_{p:\, \nabla v_0(p) \neq 0} \Big[m_{12}(p) \Big\langle \nabla r_u(p), \tfrac{\nabla v_0(p)}{\|\nabla v_0(p)\|_2} \Big\rangle + m_{22}(p) \Big\langle \nabla r_v(p), \tfrac{\nabla v_0(p)}{\|\nabla v_0(p)\|_2} \Big\rangle \Big],$$

$$\mathcal{R}_{\mathrm{var}} := \sum_{p \in \Omega} \rho(p) \big(|r_u(p)| + |r_v(p)| + \|\nabla r_u(p)\|_2 + \|\nabla r_v(p)\|_2 \big),$$

$$\mathcal{R}_{\mathrm{mix}} := \sum_{p \in \Omega_{01}} |m_{12}(p)|\, \|\nabla r_v(p)\|_2 + \sum_{p \in \Omega_{10}} |m_{12}(p)|\, \|\nabla r_u(p)\|_2.$$

Finally, let

$$D^+Q(0) := \limsup_{\alpha \downarrow 0} \frac{Q(\alpha) - Q(0)}{\alpha}$$

be the upper right Dini derivative of Q at 0.

For a vector field $\Phi = (\Phi_u, \Phi_v)$ on Ω we write $\partial_d \Phi(p) := \big((\partial_d \Phi_u)(p), (\partial_d \Phi_v)(p)\big)^\top$, where ∂_x and ∂_y are the two components of K, and $\bar{\Phi}$ denotes the extension of Φ by zero outside Ω.

Lemma 6 *Let $\Phi(p) := M_p r(p)$ for $p \in \Omega$. For every $p \in \Omega$ and $d \in \{x, y\}$,*

$$\partial_d \Phi(p) = M_p\, \partial_d r(p) + (\partial_d M_p)\, \bar{r}(p + e_d).$$

Proof If $p + e_d \in \Omega$, then

$$\partial_d \Phi(p) = M_{p+e_d} r(p + e_d) - M_p r(p) = M_p \big(r(p + e_d) - r(p)\big) + (M_{p+e_d} - M_p)\, r(p + e_d).$$

If $p + e_d \notin \Omega$, then $\partial_d \Phi(p) = 0 - M_p r(p) = M_p\, \partial_d r(p)$, and the second term vanishes because $\bar{r}(p + e_d) = 0$. □

Lemma 7 *For each $p \in \Omega$ let $\zeta_u(p), \zeta_v(p) \in \mathbb{R}^2$ satisfy $\|\zeta_u(p)\|_2 \le 1$ and $\|\zeta_v(p)\|_2 \le 1$, and set $\eta_d(p) := \big([\zeta_u(p)]_d, [\zeta_v(p)]_d\big)^\top \in \mathbb{R}^2$ for $d \in \{x, y\}$. Then*

$$\Big| \sum_{p \in \Omega} \sum_{d \in \{x,y\}} \big\langle (\partial_d M_p)\, \bar{r}(p + e_d),\, \eta_d(p) \big\rangle \Big| \le \sqrt{2}\, \mathcal{R}_{\mathrm{var}}.$$

Proof Fix p. Since $\|\eta_x(p)\|_2^2 + \|\eta_y(p)\|_2^2 = \|\zeta_u(p)\|_2^2 + \|\zeta_v(p)\|_2^2 \le 2$, the Cauchy–Schwarz inequality in $\mathbb{R}^2$ and then in the index d gives

$$\begin{aligned} \Big| \sum_d \big\langle (\partial_d M_p) \bar{r}(p + e_d), \eta_d(p) \big\rangle \Big| &\le \sum_d \|\partial_d M_p\|_{\mathrm{op}} \|\bar{r}(p + e_d)\|_2 \|\eta_d(p)\|_2 \\ &\le \sqrt{2}\, \rho(p) \max_d \|\bar{r}(p + e_d)\|_2. \end{aligned}$$

In both cases of the definition, $\bar{r}(p + e_d) = r(p) + \partial_d r(p)$, so

$$\begin{aligned} \|\bar{r}(p + e_d)\|_2 &\le |r_u(p)| + |r_v(p)| + |\partial_d r_u(p)| + |\partial_d r_v(p)| \\ &\le |r_u(p)| + |r_v(p)| + \|\nabla r_u(p)\|_2 + \|\nabla r_v(p)\|_2. \end{aligned}$$

Summing over p gives the claim. □

Theorem 8 *Assume (A1) and*

$$\nabla u_0(p) = 0 \implies \nabla u_g(p) = 0, \quad \nabla v_0(p) = 0 \implies \nabla v_g(p) = 0 \quad \forall p \in \Omega, \tag{4.6a}$$

$$S_u + S_v - \sqrt{2}\, \mathcal{R}_{\mathrm{var}} - \mathcal{R}_{\mathrm{mix}} > 0. \tag{4.6b}$$

Then $D^+Q(0) \le -2\big(S_u + S_v - \sqrt{2}\, \mathcal{R}_{\mathrm{var}} - \mathcal{R}_{\mathrm{mix}}\big) < 0$. Consequently $Q(\alpha) < Q(0)$ for all sufficiently small $\alpha > 0$, and $\alpha = 0$ is not a solution of (4.4). Every solution, which exists by Corollary 5, satisfies $\alpha^* > 0$.

Proof Subtracting (4.1) at $\alpha = 0$, where $A_p^\top (A_p w_0(p) - b_p) = 0$, gives

$$w_\alpha(p) - w_0(p) = \alpha M_p(\mathbf{div}\, z_\alpha)(p).$$

With $\Phi_\alpha(p) := M_p\big(w_\alpha(p) + w_0(p) - 2w_g(p)\big)$ and the symmetry of M_p,

$$\frac{Q(\alpha) - Q(0)}{\alpha} = \frac{1}{\alpha}\sum_p \big\langle w_\alpha(p) - w_0(p),\, w_\alpha(p) + w_0(p) - 2w_g(p)\big\rangle = \sum_p \big\langle (\mathbf{div}\, z_\alpha)(p), \Phi_\alpha(p)\big\rangle,$$

and the adjoint identity (2.6) for each component yields

$$\frac{Q(\alpha) - Q(0)}{\alpha} = -\sum_p \Big[\big\langle \nabla\Phi_{\alpha,u}(p), z_{\alpha,u}(p)\big\rangle + \big\langle \nabla\Phi_{\alpha,v}(p), z_{\alpha,v}(p)\big\rangle\Big]. \tag{4.7}$$

By Proposition 2(i), $w_\alpha \to w_0$ as $\alpha \downarrow 0$, and w_α is bounded for $\alpha \in (0, 1]$. Since the dual fields take values in the unit ball, the right-hand side of (4.7) is bounded, and $D^+Q(0)$ is finite. Choose $\alpha_n \downarrow 0$ such that the quotients converge to $D^+Q(0)$, and pass to a subsequence along which $z_{\alpha_n,u} \to z_{0,u}$ and $z_{\alpha_n,v} \to z_{0,v}$. Since $\nabla u_{\alpha_n}(p) \to \nabla u_0(p)$ and the graph of $\partial\|\cdot\|_2$ is closed, $z_{0,u}(p) \in \partial\|\cdot\|_2(\nabla u_0(p))$, and likewise for v. Moreover $\Phi_{\alpha_n} \to 2\Phi$ with $\Phi(p) := M_p r(p)$. With $\eta_d(p) := \big([z_{0,u}(p)]_d, [z_{0,v}(p)]_d\big)^\top$, the limit in (4.7) is

$$D^+Q(0) = -2\sum_p \Big[\langle \nabla\Phi_u(p), z_{0,u}(p)\rangle + \langle \nabla\Phi_v(p), z_{0,v}(p)\rangle\Big] = -2\sum_p\sum_d \big\langle \partial_d\Phi(p), \eta_d(p)\big\rangle.$$

By Lemma 6, the last double sum equals $T_1 + T_2$ with

$$T_1 := \sum_p\sum_d \big\langle M_p\, \partial_d r(p), \eta_d(p)\big\rangle, \qquad T_2 := \sum_p\sum_d \big\langle (\partial_d M_p)\, \bar r(p + e_d), \eta_d(p)\big\rangle,$$

and Lemma 7 with $\zeta_u = z_{0,u}$ and $\zeta_v = z_{0,v}$ gives $|T_2| \le \sqrt{2}\,\mathcal{R}_{\mathrm{var}}$. The expansion of the matrix product in T_1 gives

$$T_1 = \sum_p \sum_{i\in\{u,v\}} \Big\langle \sum_{j\in\{u,v\}} m_{ij}(p)\, \nabla r_j(p),\, z_{0,i}(p)\Big\rangle.$$

On Ω_{11} both subgradients are determined:

$$z_{0,u}(p) = \nabla u_0(p)/\|\nabla u_0(p)\|_2, \quad z_{0,v}(p) = \nabla v_0(p)/\|\nabla v_0(p)\|_2,$$

and the pixel contributes exactly its summands of S_u and S_v. On Ω_{10} the part with $i = u$ is the S_u summand. In the part with $i = v$, condition (4.6a) and the linearity of ∇ give $\nabla r_v(p) = \nabla v_0(p) - \nabla v_g(p) = 0$, so this part equals $m_{12}(p)\langle \nabla r_u(p), z_{0,v}(p)\rangle \ge -|m_{12}(p)|\,\|\nabla r_u(p)\|_2$. The set Ω_{01} is symmetric, with the roles of u and v exchanged. On Ω_{00}, condition (4.6a) gives $\nabla r_u(p) = \nabla r_v(p) = 0$, so the contribution is 0. Hence $T_1 \ge S_u + S_v - \mathcal{R}_{\mathrm{mix}}$.

Consequently $D^+Q(0) = -2(T_1 + T_2) \le -2\big(S_u + S_v - \mathcal{R}_{\mathrm{mix}} - \sqrt{2}\,\mathcal{R}_{\mathrm{var}}\big)$, which is negative by (4.6b). Since the lim sup is negative, there is $\delta > 0$ with $\big(Q(\alpha) - Q(0)\big)/\alpha < 0$ for every $\alpha \in (0, \delta)$. □

Remark 7 In the TV case the dual fields z_α need not be unique. The proof does not assume that $\alpha \mapsto w_\alpha$ is differentiable at $\alpha = 0$. The upper Dini derivative is always finite here (see the proof of Theorem 8), and its negativity is all the argument needs.

Remark 8 Condition (4.6a) is an idealized structural assumption. It asks that the ground truth be flat wherever a component of w_0 is flat. At such pixels the residual gradient of that component vanishes, so the undetermined subgradient does not enter the corresponding same-component term. Its remaining contributions, through the off-diagonal entry $m_{12}(p)$

and through the variation of M_p, are bounded by $\mathcal{R}_{\mathrm{mix}}$ and $\sqrt{2}\,\mathcal{R}_{\mathrm{var}}$. The condition is not expected to hold for real data in which w_0 has flat regions caused by the aperture problem. When w_0 is computed from noisy images, exact zeros of ∇u_0 or ∇v_0 are not expected, and (4.6a) then holds trivially. This is the case in Example 5, where all pixels lie in Ω_{11}.

Condition (4.6b) compares the aggregate alignment $S_u + S_v$ between the weighted error gradients and the flow gradients of w_0 with two worst-case bounds: $\sqrt{2}\,\mathcal{R}_{\mathrm{var}}$ for the spatial variation of M_p, and $\mathcal{R}_{\mathrm{mix}}$ for the pixels where exactly one component of ∇w_0 vanishes. Both bounds replace the unknown dual vectors by their largest possible values, so the condition is conservative. It can be evaluated from w_0, the ground truth w_g and the matrices M_p.

Corollary 9 *Under the assumptions of Theorem 8, suppose in addition that* $M_p = (A_p^\top A_p)^{-1} \equiv M$ *is spatially constant with* $m_{12} = 0$*, and that no mixed pixels occur, that is,* $\Omega_{10} = \Omega_{01} = \varnothing$*. Then* $\mathcal{R}_{\mathrm{var}} = \mathcal{R}_{\mathrm{mix}} = 0$*, and condition (4.6b) reduces to*

$$\sum_{p:\,\nabla u_0(p)\neq 0} m_{11}\Big\langle \nabla(u_0-u_g)(p),\, \tfrac{\nabla u_0(p)}{\|\nabla u_0(p)\|_2}\Big\rangle + \sum_{p:\,\nabla v_0(p)\neq 0} m_{22}\Big\langle \nabla(v_0-v_g)(p),\, \tfrac{\nabla v_0(p)}{\|\nabla v_0(p)\|_2}\Big\rangle > 0.$$

If this inequality holds, then $\alpha = 0$ *is not optimal.*

Proof If M_p is constant, both the forward and the backward differences $\partial_d M_p$ vanish, so $\rho = 0$ and $\mathcal{R}_{\mathrm{var}} = 0$. Without mixed pixels, $\mathcal{R}_{\mathrm{mix}} = 0$. With $m_{12} = 0$ the displayed inequality is condition (4.6b), and the conclusion follows from Theorem 8. □

Remark 9 For the weighted loss Q_W of Remark 1, the results of Sections 3 and 4 hold with the following changes, and the proofs are otherwise unchanged. In Proposition 2, the limit of Q_W is $\|Ww_g\|_2^2$ and condition (3.8) becomes $\|W(w_0 - w_g)\|_2 < \|Ww_g\|_2$. The additional condition of Corollary 5 changes in the same way. In Theorem 3, condition (3.9) becomes $\langle \mathcal{M}_{\mathrm{LK}}^{-1} L w_0,\, W^2(w_0 - w_g)\rangle > 0$. Theorem 4 applies with the first-order conditions (3.11) written for Q_W. In part (b), $Lw_{\alpha^*} \neq 0$ is replaced by $Ww'_{\alpha^*} \neq 0$, which follows from $Lw_{\alpha^*} \neq 0$ when $\omega > 0$ at every pixel but not for a mask, and (3.12) is replaced by $2\,\|WX(\alpha^*)Lw'_{\alpha^*}\|_2\,\|W(w_{\alpha^*} - w_g)\|_2 < \|Ww'_{\alpha^*}\|_2^2$. For TV, let $\widetilde{M}_p := \omega(p)^2 M_p$. Theorem 8 holds with $\widetilde{M}_p$ in place of M_p in S_u, S_v, ρ, $\mathcal{R}_{\mathrm{var}}$ and $\mathcal{R}_{\mathrm{mix}}$, and Corollary 9 then requires $\widetilde{M}_p$, not M_p, to be spatially constant, with m_{ij} the entries of $\widetilde{M}_p$.

5 Illustrative Examples

We present numerical examples that illustrate the theoretical conditions, including cases in which the sufficient conditions fail. These examples are not intended as a performance evaluation or comparison with state-of-the-art methods. Two scales are considered: a 3×3 grid (and a 4×4 grid in Example 6), which allows direct inspection of all conditions, and a 50×50 grid, which exercises the framework at a more realistic scale with noise and subpixel motion.

In the examples built from images, frame I_2 is produced by bilinear backward warping of I_1 with a constant global translation. Gaussian noise is then added to I_2, and the intensities are clipped to $[0, 1]$. The spatial derivatives I_x, I_y are central differences at interior pixels and one-sided differences at the border, computed on

each frame and averaged over the two frames. Each window Ω_p is the 5×5 block of pixels centered at p, truncated at the image border. All reported numerical values of these examples are illustrative and depend on the particular noise realization. In the examples on small grids, the image derivatives or the matrices $A_p^\top A_p$ are prescribed directly, and (A1) holds without modification.

The theoretical results (Proposition 2, Theorems 3, 4 and 8, and Corollary 5) assume $\mathcal{M}_{\mathrm{LK}} \succ 0$, which follows from full-rank LK patches. In the 50×50 experiments, the stabilization $A_p^\top A_p \mapsto A_p^\top A_p + 10^{-2} I$ ensures positive definiteness in low-texture regions where the rank condition may not hold strictly. The theorems apply to the stabilized $\mathcal{M}_{\mathrm{LK}}$.

Quadratic regularization uses the matrix L and TV regularization uses the gradient K of Section 2.

For quadratic regularization, the lower-level solution $w_\alpha = (\mathcal{M}_{\mathrm{LK}} + \alpha L)^{-1} f_{\mathrm{LK}}$ is computed by one linear solve. For TV, the lower-level problem is solved by the Chambolle–Pock primal–dual algorithm [9] with the exact proximal map of the data term, which is a 2×2 linear solve per pixel. The step sizes start from $\tau = \sigma = 0.99/\sqrt{8}$, so that $\tau\sigma\|K\|^2 < 1$, and follow the accelerated rule for a strongly convex data term. Each solve stops when the relative primal–dual gap falls below a prescribed tolerance or at an iteration limit, which is 10^5 for the sweep solves. The sweep tolerances are given in each example.

In Examples 1 and 2 the optimal parameter is located by bounded one-dimensional minimization (Brent's method), and sweeps of Q over a grid of α values produce the plotted curves. In Examples 4 and 5, Q is sampled on a grid of α values with warm starts and then on a finer grid around the coarse minimum, and a confirmation run at α^* uses the tolerance 10^{-12}. Example 3 prescribes the parameter and checks it by a grid sweep, and Example 6 evaluates the positivity condition without a parameter search.

5.1 Quadratic Regularization

Examples 1 and 2 satisfy the existence condition (3.8) and have finite minimizers. Remark 10 gives a case in which (3.8) fails and a finite minimizer still exists, and the one-pixel instance after Proposition 2 shows that without (3.8) the infimum need not be attained. A third, constructed example illustrates the second-order condition of Theorem 4(b).

Example 1 We consider a synthetic 3×3 image pair with per-pixel spatial gradients and temporal derivatives given by

$$\nabla I = \begin{bmatrix} (1.0, 0.0) & (0.0, 1.0) & (1.0, 1.0) \\ (0.5, 0.5) & (0.7, 0.3) & (0.0, 0.8) \\ (0.2, 0.9) & (1.0, -0.5) & (0.4, 0.6) \end{bmatrix}, \qquad I_t = \begin{bmatrix} -0.1 & -0.2 & -0.3 \\ -0.4 & -0.5 & -0.6 \\ -0.7 & -0.8 & -0.9 \end{bmatrix},$$

and a spatially varying ground-truth flow field

$$w_g = \begin{bmatrix} (0.1, 0.1) & (0.2, 0.1) & (0.3, 0.2) \\ (0.2, 0.3) & (0.4, 0.3) & (0.5, 0.4) \\ (0.3, 0.5) & (0.6, 0.5) & (0.7, 0.6) \end{bmatrix}.$$

For this toy example, the entire 3×3 grid is used as the local neighborhood for each pixel, so that all blocks $A_p^\top A_p$ and vectors $A_p^\top b_p$ are identical. We obtain

$$A_p^\top A_p = \begin{bmatrix} 3.940 & 1.380 \\ 1.380 & 4.400 \end{bmatrix}, \qquad A_p^\top b_p = \begin{bmatrix} 2.250 \\ 2.100 \end{bmatrix}.$$

The constant Lucas–Kanade solution is therefore $w_0(p) \approx (0.454,\, 0.335)$ at every pixel.

We first verify the conditions of Proposition 2 and Theorem 3. For this example,

$$Q(0) = \|w_0 - w_g\|_2^2 \approx 0.648 \quad \text{and} \quad \|w_g\|_2^2 \approx 2.79,$$

so $Q(0) < \|w_g\|_2^2$. Hence a finite optimal regularization parameter exists under Dirichlet boundary conditions.

The inner product

$$\left\langle \mathcal{M}_{\mathrm{LK}}^{-1} L w_0,\, w_0 - w_g \right\rangle \approx 0.110 > 0,$$

so condition (3.9) is satisfied and $\alpha^* > 0$. Although w_0 is spatially uniform, $Lw_0 \neq 0$. Since w_0 is constant and the grid spacing is 1, $(Lw_0)(p) = (4 - k_p)\, w_0(p)$, where k_p is the number of in-grid neighbors of p, that is, $k_p = 2$, 3 and 4 at the corners, the edge midpoints and the center. The inner product is therefore $\sum_p (4-k_p)\, \langle (A_p^\top A_p)^{-1} w_0(p),\, w_0(p) - w_g(p) \rangle \approx 0.1102$.

Numerical minimization of $Q(\alpha) = \|w_\alpha - w_g\|_2^2$ gives

$$\alpha^* \approx 0.585, \qquad Q(\alpha^*) \approx 0.597,$$

which is an error reduction of approximately 7.96%. Figure 1 shows $Q(\alpha)$.

At the computed minimizer, the first-order condition is confirmed: $|Q'(\alpha^*)| < 10^{-7}$. Moreover, $Lw_{\alpha^*} \neq 0$ and direct numerical evaluation gives $Q''(\alpha^*) \approx 0.204 > 0$. Hence α^* is a strict local minimizer, although the residual bound (3.12) of Theorem 4(b) is not satisfied here ($0.772 > 0.521$) (see Remark 11). Table 1 summarizes the key numerical values.

Table 1 Key numerical values for Example 1.

Quantity	Value
$Q(0) < \|w_g\|_2^2$ (Dirichlet existence, (3.8))	Yes
$\langle \mathcal{M}_{\mathrm{LK}}^{-1} L w_0,\, w_0 - w_g \rangle$	0.110
Condition (3.9) satisfied	Yes
Optimal α^*	0.585
$Q(0) = \|w_0 - w_g\|_2^2$	0.648
$\|w_g\|_2^2$	2.79
$Q(\alpha^*)$	0.597
Error reduction	7.96%
$\|Q'(\alpha^*)\|$	$< 10^{-7}$
$Q''(\alpha^*)$	0.204
$Lw_{\alpha^*} \neq 0$	Yes
Strict local minimizer ($Q''(\alpha^*) > 0$, direct evaluation)	Yes

Example 2 We consider a 50×50 scene where frame I_1 is a smooth random texture obtained by applying a Gaussian blur (standard deviation 3 pixels) to white noise and normalizing to

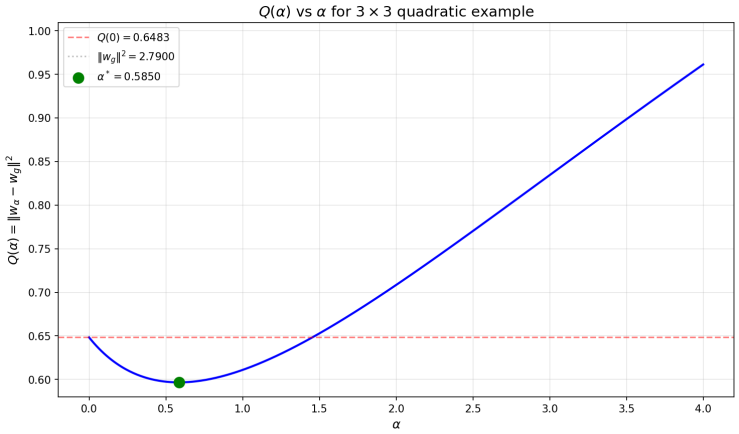

Fig. 1 Objective $Q(\alpha) = \|w_\alpha - w_g\|_2^2$ versus the regularization parameter α for the 3×3 quadratic example. The curve decreases for small $\alpha > 0$, which confirms the prediction $\alpha^* > 0$ from condition (3.9)

$[0, 1]$. Frame I_2 is the warp of I_1 by the constant global translation $w_g = (0.7, -0.3)$ at every pixel, with border values replicated, and the noise level is $\sigma = 0.10$. Both random fields come from NumPy's `default_rng` with seed 154.

Under Dirichlet boundary conditions,

$$\lim_{\alpha \to \infty} Q(\alpha) \;=\; \|w_g\|_2^2 \;=:\; Q_\infty \;=\; 1450.$$

We verify the conditions of Proposition 2 and Theorem 3. For the Dirichlet finite-existence condition (3.8),

$$Q(0) = \|w_0 - w_g\|_2^2 \approx 1294.8 \;<\; \|w_g\|_2^2 = 1450.0 = Q_\infty,$$

so condition (3.8) is *satisfied* and Proposition 2 guarantees a finite optimal regularization parameter.

For the positivity test (condition (3.9)),

$$\left\langle \mathcal{M}_{\mathrm{LK}}^{-1} L\, w_0,\; w_0 - w_g \right\rangle \approx 1.24 \times 10^4 \;>\; 0,$$

so condition (3.9) is satisfied and $Q'(0) < 0$, which guarantees $\alpha^* > 0$.

A log-scale sweep over $\alpha \in [10^{-2}, 10^2]$ brackets the minimizer, and the bounded search locates it at

$$\alpha^* \approx 0.532, \qquad Q(\alpha^*) \approx 896.4,$$

a 30.8% reduction versus $Q(0)$. Figure 2(d)–(f) shows the curve $Q(\alpha)$ and the pointwise error maps $|w_0 - w_g|$ and $|w(\alpha^*) - w_g|$.

At the computed minimizer: $|Q'(\alpha^*)| \approx 0$, $Lw_{\alpha^*} \neq 0$, and $Q''(\alpha^*) \approx 230 > 0$. The curvature decomposes as $Q''(\alpha^*) = H_1(\alpha^*) + H_2(\alpha^*)$ with $H_1(\alpha^*) \approx 104$ and $H_2(\alpha^*) \approx 126 \geq 0$, so $Q''(\alpha^*) \geq H_1(\alpha^*) > 0$ holds without invoking the residual bound (3.12). Six further realizations of the same generator (seeds 0–3, 42 and 99) give $Q(0)$ between 1092 and 1323, α^* between 0.52 and 0.68, and reductions between 30.0% and 38.2%. At every seed $H_2(\alpha^*) > 0$ and the residual bound (3.12) fails. This small sample describes the variability

Table 2 Key numerical values for Example 2.

Quantity	Value
$Q(0) = \|w_0 - w_g\|_2^2$	1294.8
$\|w_g\|_2^2 = Q_\infty$	1450.0
$Q(0) < \|w_g\|_2^2$ (Dirichlet existence, (3.8))	Yes
$\langle \mathcal{M}_{\mathrm{LK}}^{-1} L w_0,\, w_0 - w_g \rangle$ (positivity, (3.9))	1.24×10^4
α^* (interior minimum, guaranteed by Prop. 2 and Thm. 3)	0.532
$Q(\alpha^*)$	896.4
Reduction vs. $\alpha = 0$	30.8%
$\|Q'(\alpha^*)\|$	≈ 0
$Q''(\alpha^*) = H_1(\alpha^*) + H_2(\alpha^*)$	230
$H_2(\alpha^*) = -4\langle X(\alpha^*) L w'_{\alpha^*},\, w_{\alpha^*} - w_g \rangle$	126
$H_2 \geq 0 \Rightarrow Q'' \geq H_1 > 0$ (no bound needed)	Yes
Residual bound (3.12) satisfied	No
$L w_{\alpha^*} \neq 0$	Yes
Strict local minimizer ($Q''(\alpha^*) > 0$, $H_2 \geq 0$ route)	Yes

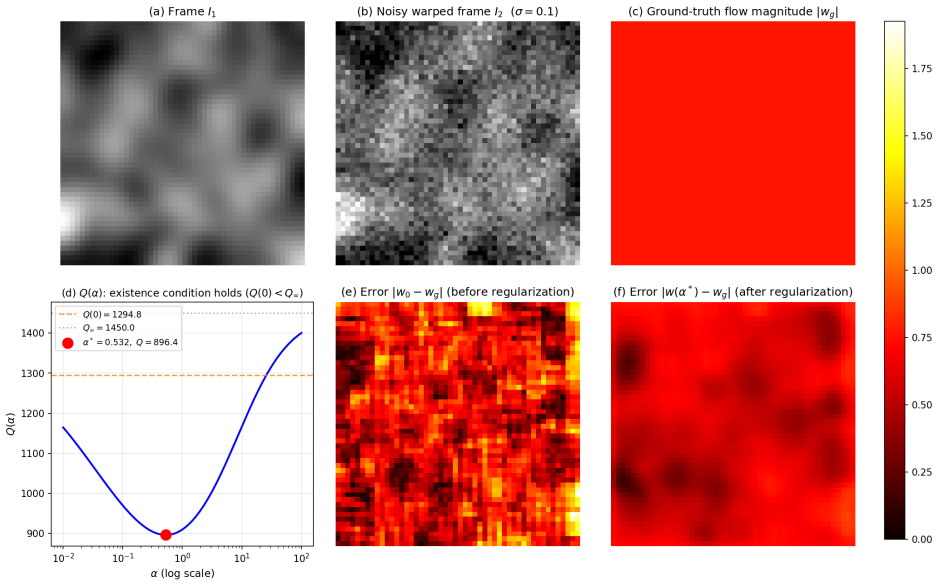

Fig. 2 Example 2: (a) Frame I_1 (smooth random texture). (b) Noisy warped frame I_2 ($\sigma = 0.10$). (c) Ground-truth flow magnitude $|w_g|$ (uniform: global translation $(0.7, -0.3)$, magnitude ≈ 0.76). (d) $Q(\alpha)$ on a log scale, with the minimizer $\alpha^* \approx 0.532$ (red marker). (e) Pointwise error $|w_0 - w_g|$ before regularization. (f) Pointwise error $|w(\alpha^*) - w_g|$ after regularization (30.8% reduction in Q)

of the realization and is not a robustness study. Hence $Q''(\alpha^*) > 0$ directly confirms that α^* is a strict local minimizer (see Remark 11). Table 2 summarizes the key values.

Figure 3 shows the noise sensitivity of the quadratic bilevel improvement. The improvement ratio $(Q(0) - Q(\alpha^*))/Q(0)$ rises from $\approx$18% at $\sigma = 0.02$ to $\approx$37% at $\sigma = 0.10$ and then stabilizes near 36%. The figure reports medians over five further seeds (42–46) of the same

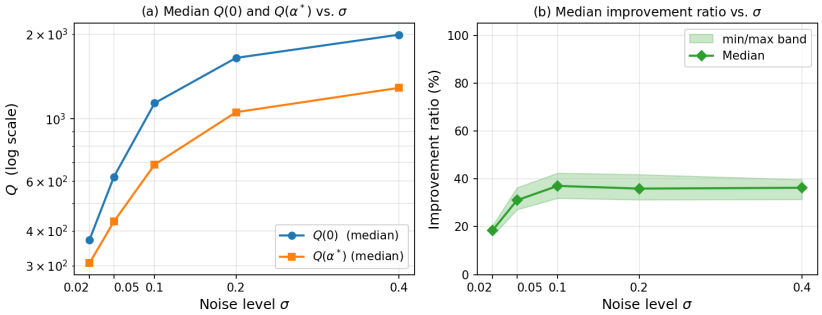

Fig. 3 Noise-sensitivity of the bilevel improvement for Example 2 (50×50 textured global-translation scene, quadratic regularization, five random seeds). (a) Median $Q(0)$ and $Q(\alpha^*)$ versus noise level σ (log y-scale). (b) Median improvement ratio $(Q(0) - Q(\alpha^*))/Q(0)$ with min/max band over seeds. The band width remains below 11 percentage points over the five seeds

generator, so its value at $\sigma = 0.10$ differs from the 30.8% of the single realization above. For this scene and these seeds, the benefit persists over the tested range of noise levels.

Remark 10 The existence condition (3.8) is sufficient but not necessary. We retain the textured scene of Example 2 (same generator and seed) and raise the noise to $\sigma = 0.50$, at which 880 of the 2500 pixels of I_2 are clipped. This gives $Q(0) \approx 2089.0 > \|w_g\|_2^2 = 1450.0$, so condition (3.8) *fails*. Nevertheless, numerical optimization finds a finite interior minimizer at $\alpha^* \approx 13.0$ with $Q(\alpha^*) \approx 1425.4 < Q_\infty = 1450.0$, a 31.8% reduction. In the realizations computed here, condition (3.8) holds at $\sigma = 0.10$ (seed 154 and the six further seeds of Example 2). It fails for all five seeds of the noise study (Figure 3) at $\sigma = 0.20$ and at $\sigma = 0.40$, where $Q(0)$ lies between 1609 and 2126, and at $\sigma = 0.50$. In all these cases a finite minimizer still reduces the error by 31% to 42%. Noise inflates $Q(0)$, and in these realizations the condition already fails at $\sigma = 0.20$, although regularization still reduces the error.

Example 3 We reuse the 3×3 quadratic setup of Example 1 (same $\mathcal{M}_{\mathrm{LK}}$, L, f_{LK}) and fix $\alpha^* = 1$. The flow w^* and its sensitivity $d = w'_{\alpha^*}$ are computed from the lower-level system, and w_g is then chosen so that α^* satisfies both the first-order condition and the residual bound of Theorem 4(b).

The key quantities are defined by

$$X^* := (\mathcal{M}_{\mathrm{LK}} + \alpha^* L)^{-1}, \qquad w^* := X^* f_{\mathrm{LK}},$$
$$d := w'_{\alpha^*} = -X^* L w^*, \qquad z := X^* L d,$$
$$v := \frac{P_{d^\perp} z}{\|P_{d^\perp} z\|_2}, \qquad \varepsilon := \frac{\|d\|_2^2}{4\|z\|_2}, \qquad w_g := w^* - \varepsilon v,$$

where $P_{d^\perp} z := z - \frac{\langle z, d\rangle}{\|d\|_2^2} d$ is the component of z orthogonal to d. Since $Lw^* \neq 0$ we have $d \neq 0$. The assumption $P_{d^\perp} z \neq 0$ ensures v is well-defined, with $\|v\|_2 = 1$, $v \perp d$, and $\langle z, v\rangle > 0$.

Because $w^* - w_g = \varepsilon v$ and $v \perp d$, the first-order condition gives $Q'(\alpha^*) = 2\langle d, \varepsilon v\rangle = 0$, so α^* is a critical point (cf. (3.11)). By construction $\varepsilon = \frac{\|d\|_2^2}{4\|z\|_2} < \frac{\|d\|_2^2}{2\|z\|_2}$, so the residual bound (3.12) holds. Moreover $H_2(\alpha^*) = -4\varepsilon\langle z, v\rangle < 0$. The example therefore lies in the regime where the curvature contribution H_2 is negative, and Theorem 4(b) guarantees that α^* is a strict local minimizer.

For the 3×3 setup with $\alpha^* = 1$, numerical verification gives $\|d\|_2 \approx 0.262$ and $\|z\|_2 \approx 0.0626$. A sweep of Q over $[0, 6]$ with step 10^{-3} attains its smallest value at $\alpha = 1$. Conditions (3.8) and (3.9) both hold for this constructed w_g. Table 3 collects the remaining quantities.

Table 3 Key numerical values for the auxiliary Example 3.

Quantity	Value
$Q(0) < \|w_g\|_2^2$ (Dirichlet existence, (3.8))	Yes
Condition (3.9) satisfied	Yes
Prescribed α^*	1.000
$Q(0) = \|w_0 - w_g\|_2^2$	0.166
$\|w_g\|_2^2$	1.764
$Q(\alpha^*)$	0.075
Error reduction	54.9%
$Q'(\alpha^*)$	≈ 0
$H_1(\alpha^*) = 2\|d\|_2^2$	0.137
$H_2(\alpha^*) = -4\langle z, w^* - w_g\rangle$	-0.026
$Q''(\alpha^*) = H_1 + H_2$	0.111
$\|w^* - w_g\|_2$ (residual ε)	0.274
$\|d\|_2^2/(2\|z\|_2)$ (SOSC bound)	0.548
Residual $<$ bound (3.12)	Yes
Strict local minimizer ($Q''(\alpha^*) > 0$, residual bound)	Yes

Remark 11 Examples 1–3 clarify the status of the second-order condition in Theorem 4. In Example 1 the residual bound (3.12) is not satisfied, yet direct evaluation confirms $Q''(\alpha^*) > 0$. The bound is conservative at this instance. In Example 2 the coupling term satisfies $H_2(\alpha^*) \geq 0$ at all seven computed realizations, so $Q''(\alpha^*) = H_1(\alpha^*) + H_2(\alpha^*) \geq H_1(\alpha^*) > 0$ follows automatically. The residual bound, which also fails here, is not needed to certify strict local optimality. The auxiliary Example 3 provides the complementary picture: the bound *is* satisfied and the coupling term is adverse ($H_2 < 0$), so the residual bound certifies $H_1 > |H_2|$ without direct evaluation of Q''.

5.2 Total Variation

Example 4 We consider a 3×3 grid with a ground truth flow that exhibits a step pattern, shared by both components:

$$w_g = \begin{bmatrix} (0,0) & (0,0) & (0,0) \\ (0,0) & (1,1) & (1,1) \\ (0,0) & (1,1) & (1,1) \end{bmatrix}.$$

To construct an unregularized solution w_0 that satisfies the structural assumptions of Theorem 8, we introduce oscillatory perturbations around the step:

$$w_0 = \begin{bmatrix} (-1,\ -1) & (0.5,\ 0.5) & (-0.5,\ -0.5) \\ (0.5,\ 0.5) & (0.5,\ 0.5) & (1.5,\ 1.5) \\ (-0.5,\ -0.5) & (1.5,\ 1.5) & (0.5,\ 0.5) \end{bmatrix}.$$

In this example $u_0 = v_0$ and $u_g = v_g$ at every pixel. With the gradient K, $\nabla u_0(p) \neq 0$ at all nine pixels, so every pixel lies in Ω_{11} and condition (4.6a) holds trivially.

This example uses the specialization $\mathcal{M}_{\mathrm{LK}} = I$ (identity data term), so the lower-level energy reduces to

$$E(w, \alpha) = \frac{1}{2}\|w - w_0\|_2^2 + \alpha\,\mathrm{TV}(w).$$

With $A_p^\top A_p = I$ at every pixel, Corollary 9 applies ($M_p = I$, $m_{12} = 0$, no mixed pixels), and condition (4.6b) reduces to $S_u + S_v > 0$.

We first verify the conditions of Corollary 5. Numerically, $Q(0) = \|w_0 - w_g\|_2^2 = 6.00$ and $\|w_g\|_2^2 = 8.00$. An optimal parameter exists by Corollary 5. Since $Q(0) < \|w_g\|_2^2$, condition (4.5) is satisfied, so the optimum also improves on the zero flow. The computed flow is zero for every sampled $\alpha \geq 0.72$, well below the sufficient threshold $\bar{\alpha} = \|w_0\|_2/\sigma_{\min}(K) \approx 5.94$ of the corollary.

The alignment sum $S_u + S_v \approx 17.91 > 0$ satisfies condition (4.6b), which by Theorem 8 implies $D^+Q(0) < 0$ and hence $\alpha^* > 0$. The forward-difference estimate $D^+Q(0) \approx -35.75 < 0$ (step 10^{-4}) confirms this result, consistent with $-2(S_u + S_v) \approx -35.81$ from the theoretical formula.

A sweep of 101 uniform points on $[0, 2]$, followed by 81 points around the coarse minimum, gives $\alpha^* \approx 0.211$ with $Q(\alpha^*) \approx 2.558$, an error reduction of approximately 57.4%. The sweep tolerance is 10^{-10}. The largest final relative gap was 1.1×10^{-9}. A confirmation run at α^* agrees to five digits. The behavior of $Q(\alpha)$ is shown in Figure 4. TV regularization reduces the oscillations of w_0. On this small grid it recovers the step of w_g only partly, and the zero ghost values pull the flow toward zero next to the last row and column. Table 4 summarizes the key numerical values.

Table 4 Key numerical values for Example 4.

Quantity	Value
$Q(0) = \|w_0 - w_g\|_2^2$	6.00
$\|w_g\|_2^2$	8.00
$Q(0) < \|w_g\|_2^2$ (improvement over zero flow, (4.5))	Yes
Condition (4.6a) satisfied	Yes
$\mathcal{R}_{\mathrm{var}} = 0$, $\mathcal{R}_{\mathrm{mix}} = 0$ ($M_p = I$)	—
Alignment $S_u + S_v$ in (4.6b)	17.91
$D^+Q(0)$ (forward difference; $-2(S_u + S_v) \approx -35.81$)	-35.75
α^*	0.211
$Q(\alpha^*)$	2.558
Error reduction	57.4%
Computed $w_\alpha = 0$ for all sampled $\alpha \geq$	0.72

Example 5 We consider a 50×50 scene containing a centered 10×10 square of intensity 1.0 on a zero background. The translation is $(0.7, -0.3)$, so that the square moves while the background stays static, and the noise has $\sigma = 0.10$ and seed 42. The warp uses zero values outside the image, and the noise is generated with NumPy's `RandomState`. The ground-truth flow w_g equals $(0.7, -0.3)$ on the pixels of the square in I_1 and zero elsewhere, so that $\|w_g\|_2^2 = 58.0$. The accelerated step sizes use the strong convexity modulus $\lambda_{\min}(\mathcal{M}_{\mathrm{LK}}) \approx 1.1 \times 10^{-2}$ of the data term. The sweep tolerance is 10^{-9}. The largest final relative gap over all sweep solves was 6×10^{-8}. The derivative estimates use the tolerance 10^{-13}. Both the derivative estimates and the confirmation run at α^* use at most 4×10^5 iterations. The parameter is

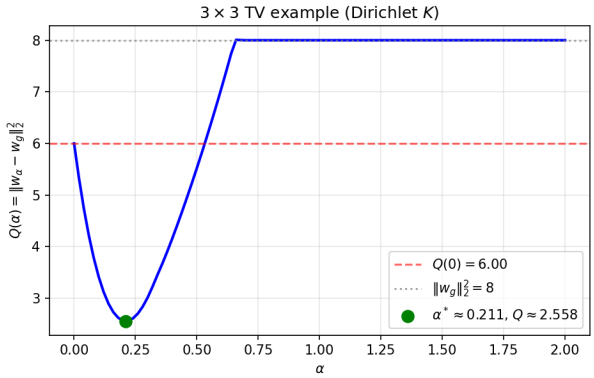

Fig. 4 Objective $Q(\alpha) = \|w_\alpha - w_g\|_2^2$ versus α for the 3×3 TV example. The minimum is at $\alpha^* \approx 0.211$. For $\alpha \geq 0.72$ the computed flow is zero and $Q(\alpha) = \|w_g\|_2^2 = 8$

sampled at 51 log-spaced points over $[10^{-4}, 10]$, followed by a 31-point linear refinement in the bracket surrounding the coarse minimum.

We first check the conditions of Corollary 5. Numerically, $Q(0) = \|w_0 - w_g\|_2^2 \approx 381.9$. Since $Q(0) > \|w_g\|_2^2$, the sufficient condition (4.5) *fails*: the unregularized Lucas–Kanade solution w_0 fits noise in background pixels and produces spurious flow that is farther from w_g than the zero-flow baseline. By Corollary 5, an optimal parameter exists whether or not condition (4.5) holds. The condition is sufficient, but not necessary, for the optimum to improve on the zero flow.

The computed w_0 satisfies $\|\nabla u_0(p)\|_2 > 0$ and $\|\nabla v_0(p)\|_2 > 0$ at all 2500 grid points. Hence $\Omega_{11} = \Omega$, and condition (4.6a) is satisfied.

Condition (4.6b) is verified directly with the stabilized per-pixel inverse $M_p = (A_p^\top A_p + 10^{-2} I)^{-1}$. Since $\Omega_{10} = \Omega_{01} = \varnothing$, the mixed-pixel coupling $\mathcal{R}_{\text{mix}} = 0$ structurally, independent of the off-diagonal entries $m_{12}(p)$. The remaining sums evaluate to $S_u \approx 24177$, $S_v \approx 24832$, $\mathcal{R}_{\text{var}} \approx 31638$, so that $S_u + S_v - \sqrt{2}\,\mathcal{R}_{\text{var}} - \mathcal{R}_{\text{mix}} \approx 4267 > 0$. Condition (4.6b) therefore holds, and Theorem 8 guarantees $\alpha^* > 0$. Forward-difference estimates give $D^+Q(0) \approx -9.6 \times 10^4$ with step 10^{-5} and $D^+Q(0) \approx -9.3 \times 10^4$ with step 10^{-4}. Both lie below the bound

$$-2\big[(S_u + S_v) - \sqrt{2}\,\mathcal{R}_{\text{var}} - \mathcal{R}_{\text{mix}}\big] \approx -8.5 \times 10^3$$

obtained in the proof of Theorem 8 (which is consistent with that theorem).

The sweep then locates the minimum at $\alpha^* \approx 1.68$ with $Q(\alpha^*) \approx 29.3$, an error reduction of approximately 92.3%. A confirmation solve at α^* reproduces $Q(\alpha^*)$ to four decimal places. For every sampled $\alpha \geq 3.2$ the computed flow is zero to solver accuracy ($\max_p |w_\alpha(p)| \leq 5.3 \times 10^{-12}$), and $Q(\alpha) = \|w_g\|_2^2$ to six digits. A dual certificate shows that the exact lower-level solution vanishes as well. By (4.3), $w_\alpha = 0$ if and only if $f_{\text{LK}} = -\alpha\,\mathbf{div}\,z$ for a dual field z with $\|z_u(p)\|_2 \leq 1$ and $\|z_v(p)\|_2 \leq 1$ at every pixel. A least-norm correction of the solver's dual field at the sample $\alpha \approx 3.16$ gives a field y with $-\mathbf{div}\,y = f_{\text{LK}}$ (relative residual 1.3×10^{-15}) and

$$\alpha_{\text{cert}} := \max_p \max\{\|y_u(p)\|_2, \|y_v(p)\|_2\} \approx 3.16054.$$

The field $z = y/\alpha$ is therefore admissible for every $\alpha \geq \alpha_{\text{cert}}$, and $w_\alpha = 0$ for all these α. No parameter beyond α_{cert} can therefore improve on the zero flow. The reported α^* gives the smallest loss found by the coarse and refined sweeps, and these samples do not certify global

optimality on $[0, \alpha_{\mathrm{cert}}]$. In particular $Q(\alpha^*) < \|w_g\|_2^2$, so the learned regularization improves on the zero flow as well as on the unregularized estimate.

TV produces a single block of nearly uniform flow around the moving square. Away from the square and a two-pixel band around it, the squared error falls from 259.9 to 1.7, and 27.6 of the remaining 29.3 lies on the square and that band. The block extends a few pixels beyond the square, and the motion inside it is underestimated: the horizontal component is about 0.40 instead of 0.7, and the vertical component is suppressed. This loss of contrast is characteristic of TV regularization [44]. On the same scene, quadratic regularization with its parameter selected in the same way attains only $Q(\alpha^*) \approx 43.7$, at $\alpha^* \approx 170$. Its mean flow magnitude on the square is about 0.19, and the flow decays only slowly into the background. Table 5 summarizes the key numerical values. The full curve $Q(\alpha)$ is shown in Figure 5(d); the pointwise error maps before and after TV regularization are displayed in Figure 5(e)–(f).

Table 5 Key numerical values for Example 5.

Quantity	Value
$Q(0) = \|w_0 - w_g\|_2^2$	381.9
$\|w_g\|_2^2$	58.0
$Q(0) < \|w_g\|_2^2$ (improvement over zero flow, (4.5))	No
Condition (4.6a) satisfied	Yes
$\|\Omega_{11}\|/\|\Omega_{10}\|/\|\Omega_{01}\|/\|\Omega_{00}\|$	2500/0/0/0
Condition (4.6b) verified	Yes
S_u	24177
S_v	24832
$\mathcal{R}_{\mathrm{var}}$	31638
$\mathcal{R}_{\mathrm{mix}}$ (structural zero, $\Omega_{10} = \Omega_{01} = \varnothing$)	0
$S_u + S_v - \sqrt{2}\,\mathcal{R}_{\mathrm{var}} - \mathcal{R}_{\mathrm{mix}}$	4267
$D^+Q(0)$ (numerical, step 10^{-5})	-9.6×10^4
α^*	1.68
$Q(\alpha^*)$	29.3
Reduction vs. $\alpha = 0$	92.3%
$Q(\alpha^*) < \|w_g\|_2^2$ (improves on zero flow)	Yes
$w_\alpha = 0$ (dual certificate) for all $\alpha \geq$	3.161

The next example is designed to exercise the full condition (4.6b), including the mixed-pixel coupling term $\mathcal{R}_{\mathrm{mix}}$ involving the off-diagonal entry m_{12}. Examples 4 and 5 both have $\Omega_{10} = \Omega_{01} = \varnothing$, so $\mathcal{R}_{\mathrm{mix}} = 0$ structurally regardless of m_{12}. Here we construct a 4×4 scene in which $|\Omega_{10}| = 9$ and $m_{12} \not\equiv 0$, so that $\mathcal{R}_{\mathrm{mix}} > 0$.

Example 6 We set the inverse structure tensor to be spatially uniform:

$$M_p = (A_p^\top A_p)^{-1} = \begin{pmatrix} 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 \end{pmatrix} \quad \text{for every pixel } p,$$

so $m_{12} = -\frac{1}{2} \neq 0$. Because M_p is spatially uniform, $\rho(p) = 0$ for every pixel and $\mathcal{R}_{\mathrm{var}} = 0$ exactly.

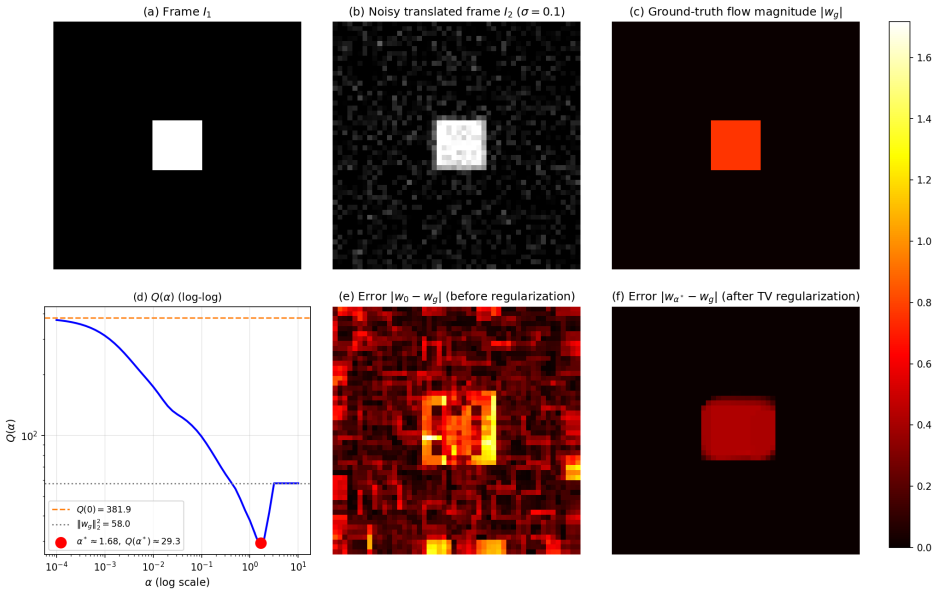

Fig. 5 Example 5: (a) Frame I_1: centered 10×10 square of intensity 1 (white) on a zero (black) background. (b) Noisy translated frame I_2 ($\sigma = 0.10$). (c) Ground-truth flow magnitude $|w_g|$: displacement $(0.7, -0.3)$ on the square, zero elsewhere (magnitude ≈ 0.76). (d) $Q(\alpha)$ on log–log axes: $Q(0) \approx 381.9 > \|w_g\|_2^2 = 58.0$ (condition (4.5) fails). The minimum $Q(\alpha^*) \approx 29.3$ at $\alpha^* \approx 1.68$ (red marker) lies below $\|w_g\|_2^2$. For $\alpha \geq 3.161$ the lower-level solution is zero, by a dual certificate, and $Q(\alpha) = \|w_g\|_2^2$. (e) Pointwise error $|w_0 - w_g|$ before regularization: large spurious errors in the static background. (f) Pointwise error $|w_{\alpha^*} - w_g|$ after TV regularization (92.3% reduction in Q): the background error is largely removed, and the remaining error is a nearly uniform block over the square and a narrow band around it, which reflects the loss of contrast of TV

Pixel coordinates are (i, j) with $i, j \in \{0, 1, 2, 3\}$. We prescribe piecewise-constant flows with a horizontal step at column $j = 2$:

$$u_0(i,j) = \begin{cases} 0 & j \leq 1, \\ \frac{3}{2} & j \geq 2, \end{cases} \quad v_0 \equiv 0; \qquad u_g(i,j) = \begin{cases} 0 & j \leq 1, \\ 1 & j \geq 2, \end{cases} \quad v_g \equiv 0.$$

The residual is $r_u(i,j) = \frac{1}{2}$ for $j \geq 2$ and zero otherwise; $r_v \equiv 0$. Hence $Q(0) = \|w_0 - w_g\|_2^2 = 2$ and $\|w_g\|_2^2 = 8$, so $Q(0) < \|w_g\|_2^2$ and condition (4.5) is satisfied (Corollary 5).

With the gradient K, $\nabla u_0(p) \neq 0$ at nine pixels. These are the four pixels $(i, 1)$, where $\nabla u_0 = (\frac{3}{2}, 0)^\top$; the three pixels $(i, 3)$ with $i \leq 2$, where $\nabla u_0 = (-\frac{3}{2}, 0)^\top$; the pixel $(3, 2)$, where $\nabla u_0 = (0, -\frac{3}{2})^\top$; and the pixel $(3, 3)$, where $\nabla u_0 = (-\frac{3}{2}, -\frac{3}{2})^\top$. Since $v_0 \equiv 0$, $\nabla v_0 \equiv 0$ throughout. The partition is $|\Omega_{11}| = 0$, $|\Omega_{10}| = 9$, $|\Omega_{01}| = 0$, $|\Omega_{00}| = 7$.

The step in u_g has the same support as u_0, so wherever $\nabla u_0(p) = 0$ one also has $\nabla u_g(p) = 0$. Since $v_g \equiv 0$, the v-implication holds trivially. Condition (4.6a) is satisfied.

Since $r_u = u_0 - u_g = \frac{1}{3} u_0$ and $r_v \equiv 0$, we have $\nabla r_u = \frac{1}{3} \nabla u_0$ and $\nabla r_v \equiv 0$. At every $p \in \Omega_{10}$ the gradient $\nabla r_u(p)$ is therefore aligned with $\hat{z}_u(p) = \nabla u_0(p) / \|\nabla u_0(p)\|_2$, and $\langle \nabla r_u(p), \hat{z}_u(p) \rangle = \|\nabla r_u(p)\|_2$. This norm equals $\frac{1}{2}$ at eight of the nine pixels and $\frac{\sqrt{2}}{2}$ at $(3, 3)$. Hence

$$S_u = \sum_{p \in \Omega_{10}} \left[m_{11} \langle \nabla r_u, \hat{z}_u \rangle + m_{12} \langle \nabla r_v, \hat{z}_u \rangle \right] = 8 \cdot \tfrac{1}{2} + \tfrac{\sqrt{2}}{2} \approx 4.707,$$

$$
\begin{aligned}
S_v &= 0 \qquad (\Omega_{11} = \Omega_{01} = \varnothing),\\
\mathcal{R}_{\mathrm{var}} &= 0 \qquad (\rho \equiv 0),\\
\mathcal{R}_{\mathrm{mix}} &= \sum_{p \in \Omega_{10}} |m_{12}(p)|\, \|\nabla r_u(p)\|_2 = \tfrac{1}{2}\, S_u \approx 2.354.
\end{aligned}
$$

Therefore $S_u + S_v - \sqrt{2}\,\mathcal{R}_{\mathrm{var}} - \mathcal{R}_{\mathrm{mix}} = \frac{1}{2}\, S_u \approx 2.354 > 0$, and Theorem 8 guarantees $\alpha^* > 0$. Table 6 summarizes the computed quantities.

Table 6 Key numerical values for Example 6.

Quantity	Value
M_p (spatially uniform)	$\left[\begin{smallmatrix} 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 \end{smallmatrix}\right]$
$Q(0) = \|w_0 - w_g\|_2^2$	2
$\|w_g\|_2^2$	8
$Q(0) < \|w_g\|_2^2$ (improvement over zero flow, (4.5))	Yes
Condition (4.6a) satisfied	Yes
$\|\Omega_{11}\|/\|\Omega_{10}\|/\|\Omega_{01}\|/\|\Omega_{00}\|$	0/9/0/7
S_u	4.707
S_v	0
$\mathcal{R}_{\mathrm{var}}$	0
$\mathcal{R}_{\mathrm{mix}}$ ($m_{12} \neq 0$, $\|\Omega_{10}\| = 9$)	2.354
$S_u + S_v - \sqrt{2}\,\mathcal{R}_{\mathrm{var}} - \mathcal{R}_{\mathrm{mix}}$	2.354
Condition (4.6b) verified	Yes
$\alpha^* > 0$ (Theorem 8)	Guaranteed

6 Limitations and Conclusion

6.1 Limitations

The results of this paper hold under assumptions that restrict their scope, and some of the conditions are conservative. We summarize these restrictions below.

The analysis requires assumption (A1), that is, full-rank local data matrices. This assumption fails in textureless regions. The 50×50 examples therefore use the stabilization $A_p^\top A_p + 10^{-2} I$, which makes $\mathcal{M}_{\mathrm{LK}}$ positive definite, and their results apply to the stabilized data term.

A single scalar parameter is learned from one pair of frames with an exact ground-truth flow. We do not analyze the effect of errors in the reference flow, nor whether a parameter learned on one pair remains suitable for other pairs. The examples are small synthetic scenes chosen to illustrate the conditions, and they do not evaluate performance on benchmark data.

The existence, positivity and second-order conditions are sufficient but not necessary, and they can be conservative. Condition (3.8) fails for the noisier realizations of Example 2, at the tested noise levels $\sigma \geq 0.20$ (Remark 10 and Figure 3). Condition (4.5), which guarantees an improvement over the zero flow, fails in Example 5. In all these cases an optimal parameter still exists and improves on w_0. Likewise, the

residual bound (3.12) of Theorem 4(b) fails in Examples 1 and 2, although $Q''(\alpha^*) > 0$ there, and only the constructed Example 3 satisfies it.

The results for TV are weaker than those for the quadratic model. We do not derive a second-order condition for the TV model, and positivity of the optimal parameter is obtained from a Dini-derivative test at $\alpha = 0$, on grids with at least two rows and two columns. Condition (4.6a) is an idealized assumption, and $\sqrt{2}\,\mathcal{R}_{\mathrm{var}}$ and $\mathcal{R}_{\mathrm{mix}}$ are worst-case bounds.

Finally, for neither model do we show that Q has a unique global minimizer. The examples locate the minimizers numerically, by Brent's method for the quadratic model and by grid refinement for TV, and these searches do not certify global optimality.

6.2 Conclusion

This paper analyzed a bilevel formulation for learning the scalar regularization parameter of CLG optical flow models from ground-truth flow, in a finite-dimensional setting in which the lower-level solution is unique.

For quadratic regularization, Proposition 2 gives existence under condition (3.8), which asks that the unregularized Lucas–Kanade estimate be closer to the ground truth than the zero flow. Theorem 3 shows that every minimizer, if one exists, satisfies $\alpha^* > 0$ when the first-order smoothing correction $-\mathcal{M}_{\mathrm{LK}}^{-1} L w_0$ has a negative component along the residual $w_0 - w_g$. Theorem 4 gives second-order sufficient conditions for strict local optimality through a residual bound derived from a Cauchy–Schwarz estimate. The examples show that this bound is conservative.

For TV regularization, Corollary 5 shows that an optimal parameter always exists, because the TV solution vanishes beyond a finite value of α. Condition (4.5) then guarantees that the optimum improves on the zero flow. Theorem 8 ensures $\alpha^* > 0$ under two conditions. The structural condition (4.6a) asks that the ground truth be flat wherever the unregularized estimate is flat. The quantitative condition (4.6b) asks that the weighted alignment $S_u + S_v$ exceed the worst-case bounds $\sqrt{2}\,\mathcal{R}_{\mathrm{var}}$, for the spatial variation of $(A_p^\top A_p)^{-1}$, and $\mathcal{R}_{\mathrm{mix}}$, for the pixels at which only one component of ∇w_0 vanishes. Corollary 9 simplifies the condition when the local structure tensors are spatially constant. In Example 5, TV regularization removes most of the background error, and the motion on the square is underestimated.

Several questions remain open. Spatially varying parameters $\alpha(p)$ can adapt the regularization to the image content. The scalar analysis suggests that such an extension will involve the same quantities: the residual of the unregularized estimate, its gradients, and the local structure tensors. Second-order conditions for TV would require a description of the nonsmooth dependence of w_α on α. Smoothing the TV term, for instance with the Huber function [13], is one possible route. The Huber function is only once continuously differentiable, however, so the second-order argument of Section 3 does not carry over directly.

Supplementary information. The Python code that reproduces every example, table and figure of Section 5 is provided as ancillary files of this arXiv submission

(directory `anc/`). Its README lists, for each example, the script, its inputs, its outputs and its location in the paper.

Declarations

Funding. This work was supported by the SSE Seed Funding program of Al Akhawayn University in Ifrane (AUI).

Competing interests. The author has no relevant financial or non-financial interests to disclose.

Data availability. All data are synthetic and are generated by the code in the ancillary files. No external data sets were used.

Code availability. The code is provided as ancillary files of this arXiv submission.

References

[1] Lucas, B.D., Kanade, T.: An iterative image registration technique with an application to stereo vision. In: Proceedings of the 7th International Joint Conference on Artificial Intelligence (IJCAI), vol. 2, pp. 674–679 (1981)

[2] Horn, B.K.P., Schunck, B.G.: Determining optical flow. Artif. Intell. **17**(1–3), 185–203 (1981) https://doi.org/10.1016/0004-3702(81)90024-2

[3] Rudin, L.I., Osher, S., Fatemi, E.: Nonlinear total variation based noise removal algorithms. Physica D **60**(1–4), 259–268 (1992) https://doi.org/10.1016/0167-2789(92)90242-F

[4] Shulman, D., Hervé, J.-Y.: Regularization of discontinuous flow fields. In: Proceedings of the IEEE Workshop on Visual Motion, Irvine, CA, USA, pp. 81–86 (1989). https://doi.org/10.1109/WVM.1989.47097

[5] Rouchouze, B., Mathieu, P., Gaidon, T., Barlaud, M.: Motion estimation based on Markov random fields. In: Proceedings of the 1st IEEE International Conference on Image Processing, Austin, TX, USA, vol. 3, pp. 270–274 (1994). https://doi.org/10.1109/ICIP.1994.413845

[6] Deriche, R., Kornprobst, P., Aubert, G.: Optical-flow estimation while preserving its discontinuities: A variational approach. In: Li, S.Z., Mital, D.P., Teoh, E.K., Wang, H. (eds.) Recent Developments in Computer Vision (ACCV'95). Lecture Notes in Computer Science, vol. 1035, pp. 69–80. Springer, Berlin, Heidelberg (1996). https://doi.org/10.1007/3-540-60793-5_63

[7] Zach, C., Pock, T., Bischof, H.: A duality based approach for realtime TV-L1 optical flow. In: Hamprecht, F.A., Schnörr, C., Jähne, B. (eds.) Pattern Recognition (DAGM 2007). Lecture Notes in Computer Science, vol. 4713, pp. 214–223.

Springer, Berlin, Heidelberg (2007). https://doi.org/10.1007/978-3-540-74936-3_22

[8] Wedel, A., Pock, T., Zach, C., Bischof, H., Cremers, D.: An improved algorithm for TV-L1 optical flow. In: Cremers, D., Rosenhahn, B., Yuille, A.L., Schmidt, F.R. (eds.) Statistical and Geometrical Approaches to Visual Motion Analysis. Lecture Notes in Computer Science, vol. 5604, pp. 23–45. Springer, Berlin, Heidelberg (2009). https://doi.org/10.1007/978-3-642-03061-1_2

[9] Chambolle, A., Pock, T.: A first-order primal-dual algorithm for convex problems with applications to imaging. J. Math. Imaging Vis. **40**(1), 120–145 (2011) https://doi.org/10.1007/s10851-010-0251-1

[10] Ranftl, R., Bredies, K., Pock, T.: Non-local total generalized variation for optical flow estimation. In: Fleet, D., Pajdla, T., Schiele, B., Tuytelaars, T. (eds.) Computer Vision – ECCV 2014. Lecture Notes in Computer Science, vol. 8689, pp. 439–454. Springer, Cham (2014). https://doi.org/10.1007/978-3-319-10590-1_29

[11] Ali, S., Daul, C., Galbrun, E., Blondel, W.: Illumination invariant optical flow using neighborhood descriptors. Comput. Vis. Image Underst. **145**, 95–110 (2016) https://doi.org/10.1016/j.cviu.2015.12.003

[12] Chambolle, A.: An algorithm for total variation minimization and applications. J. Math. Imaging Vis. **20**(1–2), 89–97 (2004) https://doi.org/10.1023/B:JMIV.0000011325.36760.1e

[13] Kalmoun, E.M.: An investigation of smooth TV-like regularization in the context of the optical flow problem. J. Imaging **4**(2), 31 (2018) https://doi.org/10.3390/jimaging4020031

[14] Bruhn, A., Weickert, J., Schnörr, C.: Lucas/Kanade meets Horn/Schunck: Combining local and global optic flow methods. Int. J. Comput. Vis. **61**(3), 211–231 (2005) https://doi.org/10.1023/B:VISI.0000045324.43199.43

[15] Kalmoun, E.M., Garrido, L.: Trust region versus line search for computing the optical flow. Multiscale Model. Simul. **11**(3), 890–906 (2013) https://doi.org/10.1137/120897213

[16] Wedel, A., Cremers, D., Pock, T., Bischof, H.: Structure- and motion-adaptive regularization for high accuracy optic flow. In: 2009 IEEE 12th International Conference on Computer Vision (ICCV), Kyoto, Japan, pp. 1663–1668 (2009). https://doi.org/10.1109/ICCV.2009.5459375

[17] Belhachmi, Z., Hecht, F.: Control of the effects of regularization on variational optic flow computations. J. Math. Imaging Vis. **40**(1), 1–19 (2011) https://doi.org/10.1007/s10851-010-0239-x

[18] Belhachmi, Z., Hecht, F.: An adaptive approach for the segmentation and the TV-filtering in the optic flow estimation. J. Math. Imaging Vis. **54**(3), 358–377 (2016) https://doi.org/10.1007/s10851-015-0608-6

[19] Engl, H.W., Hanke, M., Neubauer, A.: Regularization of Inverse Problems. Mathematics and Its Applications, vol. 375. Kluwer Academic Publishers, Dordrecht (1996). https://doi.org/10.1007/978-94-009-1740-8

[20] Hansen, P.C.: Analysis of discrete ill-posed problems by means of the L-curve. SIAM Rev. **34**(4), 561–580 (1992) https://doi.org/10.1137/1034115

[21] Golub, G.H., Heath, M., Wahba, G.: Generalized cross-validation as a method for choosing a good ridge parameter. Technometrics **21**(2), 215–223 (1979) https://doi.org/10.1080/00401706.1979.10489751

[22] Stein, C.M.: Estimation of the mean of a multivariate normal distribution. Ann. Stat. **9**(6), 1135–1151 (1981) https://doi.org/10.1214/aos/1176345632

[23] Baker, S., Scharstein, D., Lewis, J.P., Roth, S., Black, M.J., Szeliski, R.: A database and evaluation methodology for optical flow. Int. J. Comput. Vis. **92**(1), 1–31 (2011) https://doi.org/10.1007/s11263-010-0390-2

[24] Barron, J.L., Fleet, D.J., Beauchemin, S.S.: Performance of optical flow techniques. Int. J. Comput. Vis. **12**(1), 43–77 (1994) https://doi.org/10.1007/BF01420984

[25] Pagliari, V., Papafitsoros, K., Raiţă, B., Vikelis, A.: Bilevel training schemes in imaging for total variation–type functionals with convex integrands. SIAM J. Imaging Sci. **15**(4), 1690–1728 (2022) https://doi.org/10.1137/21M1467328

[26] Kunisch, K., Pock, T.: A bilevel optimization approach for parameter learning in variational models. SIAM J. Imaging Sci. **6**(2), 938–983 (2013) https://doi.org/10.1137/120882706

[27] De los Reyes, J.C., Schönlieb, C.-B.: Image denoising: learning the noise model via nonsmooth PDE-constrained optimization. Inverse Probl. Imaging **7**(4), 1183–1214 (2013) https://doi.org/10.3934/ipi.2013.7.1183

[28] De los Reyes, J.C., Schönlieb, C.-B., Valkonen, T.: The structure of optimal parameters for image restoration problems. J. Math. Anal. Appl. **434**(1), 464–500 (2016) https://doi.org/10.1016/j.jmaa.2015.09.023

[29] De los Reyes, J.C., Villacís, D.: Optimality conditions for bilevel imaging learning problems with total variation regularization. SIAM J. Imaging Sci. **15**(4), 1646–1689 (2022) https://doi.org/10.1137/21M143412X

[30] De los Reyes, J.C.: Bilevel imaging learning problems as mathematical programs

with complementarity constraints: Reformulation and theory. SIAM J. Imaging Sci. **16**(3), 1655–1686 (2023) https://doi.org/10.1137/21M1450744

[31] Ehrhardt, M.J., Gazzola, S., Scott, S.J.: On optimal regularization parameters via bilevel learning. In: Bubba, T.A. (ed.) Data-driven Models in Inverse Problems. Radon Series on Computational and Applied Mathematics, vol. 31, pp. 1–38. De Gruyter, Berlin, Boston (2025). https://doi.org/10.1515/9783111251233-001

[32] Hintermüller, M., Wu, T.: Bilevel optimization for calibrating point spread functions in blind deconvolution. Inverse Probl. Imaging **9**(4), 1139–1169 (2015) https://doi.org/10.3934/ipi.2015.9.1139

[33] Bartels, S., Weber, N.: Parameter learning and fractional differential operators: Applications in regularized image denoising and decomposition problems. Math. Control Relat. Fields **13**(1), 35–62 (2023) https://doi.org/10.3934/mcrf.2021048

[34] Crockett, C., Fessler, J.A.: Bilevel methods for image reconstruction. Found. Trends Signal Process. **15**(2–3), 121–289 (2022) https://doi.org/10.1561/2000000111

[35] Ehrhardt, M.J., Roberts, L.: Inexact derivative-free optimization for bilevel learning. J. Math. Imaging Vis. **63**(5), 580–600 (2021) https://doi.org/10.1007/s10851-021-01020-8

[36] Ehrhardt, M.J., Roberts, L.: Analyzing inexact hypergradients for bilevel learning. IMA J. Appl. Math. **89**(1), 254–278 (2024) https://doi.org/10.1093/imamat/hxad035

[37] Suonperä, E., Valkonen, T.: Single-loop methods for bilevel parameter learning in inverse imaging. J. Nonsmooth Anal. Optim. **5**, 15577 (2025) https://doi.org/10.46298/jnsao-2025-15577

[38] Domke, J.: Generic methods for optimization-based modeling. In: Lawrence, N.D., Girolami, M. (eds.) Proceedings of the Fifteenth International Conference on Artificial Intelligence and Statistics (AISTATS). Proceedings of Machine Learning Research, vol. 22, pp. 318–326 (2012). https://proceedings.mlr.press/v22/domke12.html

[39] Domke, J.: Learning graphical model parameters with approximate marginal inference. IEEE Trans. Pattern Anal. Mach. Intell. **35**(10), 2454–2467 (2013) https://doi.org/10.1109/TPAMI.2013.31

[40] Ranftl, R., Pock, T.: A deep variational model for image segmentation. In: Jiang, X., Hornegger, J., Koch, R. (eds.) Pattern Recognition (GCPR 2014). Lecture Notes in Computer Science, vol. 8753, pp. 107–118. Springer, Cham (2014). https://doi.org/10.1007/978-3-319-11752-2_9

[41] Jiang, S., Campbell, D., Liu, M., Gould, S., Hartley, R.: Joint unsupervised learning of optical flow and egomotion with bi-level optimization. In: 2020 International Conference on 3D Vision (3DV), pp. 682–691 (2020). https://doi.org/10.1109/3DV50981.2020.00078

[42] El Hakoume, A., Laghrib, A., Hadri, A., Afraites, L.: An optimal fluid optical flow registration for super-resolution with Lamé parameters learning. J. Optim. Theory Appl. **197**(2), 508–538 (2023) https://doi.org/10.1007/s10957-023-02186-4

[43] El Malki, I., Jauberteau, F., Laghrib, A., Nachaoui, M.: A bilevel learning approach for nonlocal image deblurring with variable weights parameter. Digit. Signal Process. **149**, 104505 (2024) https://doi.org/10.1016/j.dsp.2024.104505

[44] Strong, D., Chan, T.: Edge-preserving and scale-dependent properties of total variation regularization. Inverse Probl. **19**(6), S165–S187 (2003) https://doi.org/10.1088/0266-5611/19/6/059